

The Projection Spectral Layer Theory (PSLT): A Rank-2 Computable Closure for the Three-Generation Structure and Higgs Signal Strength

Bo-Yu Chen¹

¹*Independent Researcher*

(Dated: February 2, 2026)

We present a comprehensive, computable EFT-level “closed chain” for the Projection Spectral Layer Theory (PSLT) in which a self-consistent three-generation structure emerges on the baseline closure from competing spectral entropy, kinetics, and visibility. The framework is built on a minimal Einstein–Yang–Mills–Higgs (EYMH) effective action with non-minimal curvature coupling, where the background geometry induces a layer-indexed spectral scale ω_N . The theoretical closure is achieved by unifying three key modules: (i) a **Micro-degeneracy module** (g_N) implemented through a bounded two-dimensional phase-space profile (microcanonical low- N window + energy-gap tail prescription, phase-space normalized at $N = 3$); (ii) a **Rank-2 Computable Kinetic module** (Γ_N) defined as the largest positive eigenvalue of a 2×2 growth matrix with explicit WKB tunneling suppression; and (iii) a **Visibility module** (B_N) implemented as a fully normalized EFT-operator profile, where overlap-extracted flavor-layer couplings and mediator scales feed finite one-loop + LL-RG-normalized operator kernels, with neutral high- N saturation ($B_{N>3} = 1$) so that no extra visibility-side damping is added on top of the g_N regulator; a risk-weighted runtime-direct visibility parity branch is also available, but it is not interpreted as a strict all-direct closure. The resulting normalized layer probability $P_N(t_{\text{coh}})$ demonstrates, within this baseline closure, a stable “winner” phase diagram with clear dominance of layers $N = 1, 2, 3$ over a significant region of the geometric parameter space (D, η) . We verify the closure against two benchmarks: (1) internal stability of the three-generation hierarchy (Generation Ratio $> 90\%$ over 92.7% of the sampled (D, η) grid), and (2) a scan-level EFT/Wilson-matched compatibility test using the ATLAS Run-2+Run-3 combined $H \rightarrow \mu\mu$ signal-strength input ($\mu_{\mu\mu}^{\text{obs}} = 1.4 \pm 0.4$). In the current UV+LL-RG baseline, the illustrative acceptance band is about 15.0% of the scanned grid for $\chi^2 < 4$, providing a sharper falsifiable constraint on (D, η) .

CONTENTS

	2
II. Theoretical Framework	6
A. The Closed Chain Equation	6
B. Action-Derived Chain at a Glance	6
C. Geometric Foundation: Minimal EYMH + Curvature	6
D. Model-Chain Assumption Status	8
III. Geometry-to-Spectrum Worked Example	10
A. Two-Center Harmonic Conformal Factor	10
B. Action-Derived Effective Potential	11
C. Numerical Solution	11
D. Single-Track Results	12
E. Summary: Unified Geometry-to-Kinetics Pipeline	12
IV. Module 1: The Three Routes to Micro-Degeneracy g_N	14
A. Route A: Mock-Modular Counting	14
B. Route B: ER=EPR and Entanglement Entropy	14
C. Route C: Quantum Gravity Template	15
D. Unified Cardy-Controlled Envelope with High-N Regulator	15
E. Cardy Regime Validity and Finite-N Corrections	15
F. Assessment of a Phase-Space First-Principles Candidate	16
V. Module 2: Rank-2 Computable Kinetics Γ_N	16
A. Dual-Center Tunneling Suppression	16
B. Dimensionality and Units	17
C. Derivation of the Rank-2 Growth Matrix	17
D. Rank-2 Computable Closure	17
E. Physical Interpretation: Quasi-Bound Leakage and Width Proxy	19
VI. Module 3: EFT-Normalized Visibility B_N	20
A. High- N Neutral Saturation	21
B. Roadmap: Action-Derived Effective Yukawa Operator	22
VII. Verification Results	23
A. Configuration and Parameter Table	23
B. Three-Generation Phase Diagram	23
C. $H \rightarrow \mu\mu$ Signal Strength (EFT/Wilson-Matched Mapping, UV+LL-RG Baseline)	27

D. Projected Parent-Kernel Derivation	46
E. EFT/Wilson-Matched Extension to $H \rightarrow ee$ and $H \rightarrow \tau\tau$	58
VIII. Discussion and Conclusion	59
Acknowledgments	62
A. Action-Derived Chain Reproducibility SOP	62
1. Single Entry Point	62
2. Step Map and Deliverables	63
3. Acceptance Conditions	63
B. Action-Derived Effective Potential and WKB	63
1. Two-Center Harmonic Conformal Factor	63
2. Veff Derivation from KG Equation	64
3. Minimal Dirac Coupling in the Conformal Background	65
4. 1D On-Axis Reduction	66
5. 2D Axisymmetric Validation	66
6. WKB Action Convention	66
7. Splitting-Action Consistency Check	66
8. Action-Derived Numerical Results	67
9. Bound-State omega Convergence Benchmark	67
10. Fixed-dz Convergence	68
C. First-Principles Symmetry-Channel Test for χ	69
D. Localized-Channel First-Principles Extraction of χ	69
E. Preliminary First-Principles Replacement Candidate Checks	72
1. 1D Phase-Space Candidate for g_N	72
2. 2D Axisymmetric Benchmark for $g_N^{(\text{ps})}$	73
3. Baseline Cross-Check Under $g_{\text{mode}} = \text{fp-2d_full}$	75
4. Geometry-Informed and Microscopic Lindblad Scan-Ready Module for Open-System χ	76
5. Cross-Module Migration Snapshot	96
6. Reference-Normalization Sensitivity Audit	96
7. Dephasing-Based First-Principles Candidate for t_{coh}	97
8. Splitting-Prefactor First-Principles Candidate for η	98

9. Action-Derived First-Principles Candidate for Barrier-Leakage Channel Scaling	99
F. Ansatz Ablation Study	101
1. Inverse Yukawa Ansatz (Ablation A)	101
2. Comparison	101
References	102

I. INTRODUCTION

The existence of exactly three generations of fermions is one of the most persistent puzzles in the Standard Model (SM) of particle physics. While the SM accommodates three families through the CKM and PMNS mixing matrices, it offers no dynamical reason for why the number is three, nor why the mass hierarchy spans such a vast range. Traditional attempts to explain this structure often rely on complex discrete symmetries or string logic compactifications that, while elegant, can be difficult to falsify directly.

The Projection Spectral Layer Theory (PSLT) proposes a different paradigm: *Generation structure is a spectral consequence of the underlying projection geometry of spacetime.* In this framework, "generations" are not ad-hoc copies but rather distinct *spectral layers* ($N = 1, 2, 3, \dots$) of vacuum excitations, selected by a competition between microstate degeneracy (entropy) and geometric formation rates (kinetics). This structural viewpoint is complementary to recent Clifford-algebra work where three families emerge from an intrinsic S_3 organization rather than manual generation copying [1]; the emphasis here is different: a map-level computable closure with explicit scan-level falsifiability.

In this work, we consolidate the complete theoretical framework of PSLT into a single, falsifiable "closed chain". We upgrade earlier effective descriptions to a **computable** rank-2 dynamical system with an action-derived extraction subchain and an action-informed global scan chain, where $\chi_{LR}(D)$ and $\tilde{A}_\ell(D)$ are now injected on the full scan grid from direct action-derived profiles. The core deliverable is a map from geometric control parameters (D, η) —representing dual-center separation and overlap—to an experimentally observable layer probability P_N .

Our primary contributions are:

1. **Theoretical Motivation:** We present three convergent arguments for the micro-degeneracy g_N (Route A: Mock-modular, Route B: ER=EPR, Route C: QG Template), and implement a scan-ready baseline through the bounded two-dimensional phase-space profile with controlled high- N tail prescription.
2. **Rank-2 Computable Kinetics:** We replace heuristic growth rates with a rigorous eigenvalue problem for the formation matrix $\mathbf{M}_N$, incorporating numerical WKB tunneling suppression.
3. **Fully Normalized EFT Visibility:** We promote B_N to an EFT-operator-normalized baseline built from overlap-extracted flavor-layer couplings and mediator scales with finite one-loop + LL-RG normalization; overlap-only and Yukawa laws are retained only as

legacy comparators/fallbacks, while runtime-direct visibility extraction is release-qualified only as a risk-weighted parity check, not as a strict all-direct closure.

4. **Verification:** We demonstrate stable three-generation dominance at $\mathcal{R}_3 > 90\%$ and a constrained EFT/Wilson-matched-compatible band for LHC $H \rightarrow \mu\mu$.

II. THEORETICAL FRAMEWORK

A. The Closed Chain Equation

The central output of the PSLT framework is the normalized layer occupancy probability $P_N(t_{\text{coh}})$, defined by the competition between entropic weight (W_{entropy}) and kinetic formation (W_{kinetic}). The master equation is:

$$P_N(t_{\text{coh}}; D, \eta) = \frac{W_N(t_{\text{coh}})}{\sum_K W_K(t_{\text{coh}})}, \quad W_N = B_N \cdot g_N \cdot \left(1 - e^{-\Gamma_N(D, \eta)t_{\text{coh}}}\right). \quad (1)$$

This equation unifies the three critical modules:

- g_N : **Micro-degeneracy** (Entropy). The number of available microstates in layer N .
- Γ_N : **Dynamical Rate** (Kinetics). The rate at which these states can form/tunnel from the vacuum.
- B_N : **Visibility Factor** (Observation). The coupling strength of layer N to the observable sector.

Figure 1 illustrates the logical flow of the PSLT framework.

B. Action-Derived Chain at a Glance

For implementation-level clarity, Figure 2 summarizes the action-derived operator chain and its interfaces to scan-level diagnostics in one page. The executable, clone-level Standard Operating Procedure (SOP) is provided in Appendix A.

C. Geometric Foundation: Minimal EYMH + Curvature

We start from a minimal Einstein–Yang–Mills–Higgs (EYMH) effective action supplemented by a non-minimal curvature coupling $\xi R|\Phi|^2$:

$$S = \int d^4x \sqrt{-g} \left[\frac{M_{\text{Pl}}^2}{2} R - \frac{1}{4} F_{\mu\nu}^a F^{a\mu\nu} - (D_\mu \Phi)^\dagger (D^\mu \Phi) - \lambda(|\Phi|^2 - v^2)^2 - \xi R|\Phi|^2 \right]. \quad (2)$$




FIG. 1. Schematic overview of the PSLT framework. Geometric inputs (D, η) feed into three modules (micro-degeneracy g_N , kinetic rate Γ_N , visibility B_N), which combine in the master equation to yield the observable layer probabilities P_N and the three-generation ratio R_3 . In the baseline chain, high- N control is implemented by the bounded two-dimensional phase-space-tail prescription, while $B_{N>3}$ neutrally saturates to unity.

PSLT posits that the background geometry is a *projection* from a higher-dimensional manifold, inducing a locally conformally flat metric:

$$g_{\mu\nu}(x) = \Omega^2(x)\eta_{\mu\nu}. \quad (3)$$

Under conformal rescaling $\Phi \rightarrow \Omega^{-1}\tilde{\Phi}$, the scalar equation of motion reduces to a Schrödinger-like eigenproblem:

$$[-\nabla^2 + V_{\text{eff}}(x)]\psi_N = \omega_N^2\psi_N, \quad (4)$$

where the effective potential V_{eff} is determined by the conformal factor $\Omega(x)$ and the specific projection geometry (e.g., stereographic projection of dual centers).




FIG. 2. One-page action-derived chain overview used in this manuscript: geometry input $\rightarrow$ operator extraction $\rightarrow$ scan diagnostics $\rightarrow$ observable EFT/Wilson checks. This diagram is aligned with the repository reproducibility entrypoint and Appendix A.

The layer index N emerges naturally as the principal quantum number of this spectrum. In Section III, we derive the spectral scale $\mu(D)$ explicitly from a two-center geometry; the result supports a power-law scaling $\mu(D) \propto D^{-\gamma}$ with $\gamma \approx 0.1$ in the demonstrator regime. For analytic tractability, we adopt the **hydrogenic proxy**:

$$\omega_N(D) \simeq \mu(D) \left(1 - \frac{1}{2N^2} \right), \quad \mu(D) = \mu_0 D^{-\gamma}, \quad (5)$$

where μ_0 and γ are determined numerically from the geometry (Section III). This form captures the essential N -dependence from the hydrogenic spectrum while allowing the D -dependence to be set by explicit calculation. For transparency, we additionally provide a bound-state exact benchmark for (ω_1, ω_2) at $D = \{6, 12, 18\}$ with coarse/mid/fine convergence from the same action-derived operator chain (Appendix B9, Table X). These benchmark values use $E = \omega^2 - m_0^2 < 0$ and should not be mixed with the generalized localized-extraction eigenvalues λ in Appendix D.

D. Model-Chain Assumption Status

For transparency, Table I separates action-derived components from EFT-level surrogates currently used in the global scan.

TABLE I. Status map of the PSLT closed chain in the present manuscript.

Component	Current Status	Used in Global Scan	Next Upgrade Path
$\Omega(\rho, z; D)$ two-center geometry	Projected-Poisson two-source Green solution	Yes	Derive the same source from an explicit higher-dimensional parent
V_{eff} , low modes, WKB action	Action-derived and numerically converged	Yes	Extend the same solver chain to the full (D, η, N) map
g_N phase-space profile	Bounded two-dimensional phase-space baseline; direct spectral parity and risk-weighted visibility parity checks available	Yes	Treat the strict all-direct broader-grid rule as parked at the current case-aware endpoint; reopen only for a genuinely new observable-side mechanism class
$\chi_N^{(LR)}(D)$ localized channel	Full D-grid localized profile from 2D fine extraction; release-production parity path available	Yes	Keep the current localized-direct parity path as the production support line; the longer-horizon upgrade is a fully localized extraction over (D, η, N) rather than another profile-blend retune
$\Gamma_{N,\ell}$ channel normalization (A_ℓ)	Action-derived $\tilde{A}_\ell(D)$ baseline with stabilized localized-direct support; legacy constant- A_ℓ kept only as comparator	Yes	Move to fully localized channel-resolved direct extraction over (D, η, N)
Open-system χ_{eff} (Lindblad)	Locally closed projected effective-action diagnostic: scalar modulus κ_{env} , projected σ_z/σ_x jet, and single-pole gap-locked response family are now explicit on the audited knot set	Diagnostic only	Promote the current completion to a parent-side Schur statement by proving first-mode dominance of the projected resolvent on $\omega \in \{0, \Delta E_a\}$ with $\Omega_{1,a} = \omega_{1,a}$
B_N visibility profile from EFT normalization	EFT-operator-normalized baseline; runtime-direct operator extraction promoted only in the risk-weighted parity branch	Yes	Keep the release-qualified profile-blend branch separate from the parked strict all-direct endpoint; the real remaining upgrade is fully localized channel-resolved extraction, not another local blend retune
$H \rightarrow \mu\mu$ mapping [Eq. (56)]	EFT/Wilson-matched baseline with UV-tree, finite-one-loop, and LL-RG audits	Yes	Extend the current closure to fuller UV-to-EFT matching from the EYMH action

III. GEOMETRY-TO-SPECTRUM WORKED EXAMPLE

This section demonstrates the complete derivation chain $\Omega(x; D) \rightarrow V_{\text{eff}}(x) \rightarrow \omega_N(D)$, establishing the geometry-to-spectrum correspondence as a computational reality rather than an ansatz.

A. Two-Center Harmonic Conformal Factor

The geometric origin of the dual-center conformal factor is formulated as a projected Poisson constraint on the static spatial slice:

$$\nabla^2 \Omega(\mathbf{x}) = -4\pi \sigma(\mathbf{x}), \quad \sigma(\mathbf{x}) \equiv a[\rho_\varepsilon(\mathbf{x} - \mathbf{x}_+) + \rho_\varepsilon(\mathbf{x} - \mathbf{x}_-)]. \quad (6)$$

Here $\mathbf{x}_\pm = \pm(D/2)\hat{z}$ are the two projected centers and ρ_ε is a normalized smeared source. With Green's function $G(\mathbf{x}) = (4\pi|\mathbf{x}|)^{-1}$ for $-\nabla^2$, the unique asymptotically flat solution ($\Omega \rightarrow 1$ at $|\mathbf{x}| \rightarrow \infty$) is

$$\Omega(\mathbf{x}) = 1 + \int d^3x' \frac{\sigma(\mathbf{x}')}{|\mathbf{x} - \mathbf{x}'|}. \quad (7)$$

Choosing the Plummer-type kernel

$$\rho_\varepsilon(\mathbf{r}) = \frac{3\varepsilon^2}{4\pi(|\mathbf{r}|^2 + \varepsilon^2)^{5/2}}, \quad \int d^3r \rho_\varepsilon(\mathbf{r}) = 1, \quad (8)$$

one obtains the closed form

$$\boxed{\Omega(\rho, z; D) = 1 + a \left(\frac{1}{r_+} + \frac{1}{r_-} \right), \quad r_\pm = \sqrt{\rho^2 + (z \mp D/2)^2 + \varepsilon^2}} \quad (9)$$

in cylindrical coordinates (ρ, z) (axisymmetric, $m = 0$). The corresponding source identity follows from

$$\nabla^2 \left(\frac{1}{\sqrt{r^2 + \varepsilon^2}} \right) = -\frac{3\varepsilon^2}{(r^2 + \varepsilon^2)^{5/2}}, \quad (10)$$

hence

$$\nabla^2 \Omega = -4\pi a [\rho_\varepsilon(x - x_+) + \rho_\varepsilon(x - x_-)], \quad \rho_\varepsilon(x) = \frac{3\varepsilon^2}{4\pi(|x|^2 + \varepsilon^2)^{5/2}}. \quad (11)$$

Equation (9) therefore is not an arbitrary fitting form: once Eq. (6), two-center source support, and asymptotic flatness are fixed, it is the corresponding Green-function solution.

B. Action-Derived Effective Potential

We derive V_{eff} directly from the Klein-Gordon equation in curved spacetime. Consider a scalar field with non-minimal curvature coupling:

$$(\square_g - m_0^2 - \xi R)\Phi = 0. \quad (12)$$

For time-harmonic modes $\Phi = e^{-i\omega t}\phi(\mathbf{x})$ in the background $g_{\mu\nu} = \Omega^2\eta_{\mu\nu}$, we perform the conformal field rescaling $\phi = \Omega^{-1}\psi$. After explicit calculation of $\square_g$ and cancellation of first-derivative terms (see Appendix B 2), the equation reduces to:

$$\boxed{[-\nabla^2 + V_{\text{eff}}(\mathbf{x})]\psi = \omega^2\psi} \quad (13)$$

with the **action-derived** effective potential:

$$\boxed{V_{\text{eff}} = m_0^2\Omega^2 + (1 - 6\xi)\Omega^{-1}\nabla^2\Omega} \quad (14)$$

This is the **only form** consistent with the action—no engineered barrier or box terms appear.

a. Key properties.

1. At $\xi = \xi_c = 1/6$ (conformal coupling in 4D), the derivative term vanishes: $V_{\text{eff}} \rightarrow m_0^2\Omega^2$.
2. As $r \rightarrow \infty$: $\Omega \rightarrow 1$, $\nabla^2\Omega \rightarrow 0$, so $V_{\text{eff}} \rightarrow m_0^2$ (continuum threshold).
3. Near each center: $\nabla^2\Omega < 0$ (smeared source), so for $\xi < 1/6$ the potential develops attractive wells.

We define the shifted potential $U \equiv V_{\text{eff}} - m_0^2$ so that $U \rightarrow 0$ at infinity. Bound states satisfy $E = \omega^2 - m_0^2 < 0$. Figure 3 shows the decomposition.

C. Numerical Solution

We solve the 2D axisymmetric eigenproblem using finite differences on a (ρ, z) grid with Eq. (14):

$$\left[-\frac{\partial^2}{\partial \rho^2} - \frac{1}{\rho} \frac{\partial}{\partial \rho} - \frac{\partial^2}{\partial z^2} + V_{\text{eff}}(\rho, z; D) \right] \psi_N = \omega_N^2 \psi_N. \quad (15)$$

Boundary conditions: $\partial_\rho \psi|_{\rho=0} = 0$ (axis regularity), $\psi|_{\text{boundary}} = 0$ (Dirichlet). Parameters: $a = 1.0$, $\varepsilon = 0.2$, $m_0 = 1.0$, $\xi = 0.0$. Grid: $(n_\rho, n_z) = (50, 500)$ with fine resolution $\Delta z \ll \varepsilon$ near the cores.

Critical numerical requirement: The grid spacing must satisfy $\Delta z \ll \varepsilon$ to resolve the smeared source structure. This is essential for capturing the deep negative wells in U .




FIG. 3. Action-derived potential decomposition at $D = 12$. Top: mass term $m_0^2(\Omega^2 - 1)$ and curvature term $(1 - 6\xi)\Omega^{-1}\nabla^2\Omega$. Bottom: total $U = V_{\text{eff}} - m_0^2$ showing the double-well structure. The 4 turning points confirm the tunneling geometry.

D. Single-Track Results

Table II shows the unified results. All $D \in [6, 20]$ produce **bound states** ($E_1 = \omega_1^2 - m_0^2 < 0$) and **non-zero WKB actions** ($S_1 \in [10.6, 22.8]$). The 4 turning points confirm the double-well tunneling geometry.

E. Summary: Unified Geometry-to-Kinetics Pipeline

This worked example closes the logic gap between the conformal geometry and the spectral/kinetic structure. **The same V_{eff} determines both the layer frequencies and the tunneling rates:**




FIG. 4. Single-track results from action-derived V_{eff} . Top-left: Bound state energy $E_1 < 0$. Top-right: Layer frequency ω_1 . Bottom-left: WKB tunneling action S_1 . Bottom-right: Tunneling probability $r_1 = e^{-2S_1}$. All quantities are computed from the same $V_{\text{eff}}(\rho, z; D)$.

TABLE II. Single-track results from action-derived $V_{\text{eff}} = m_0^2 \Omega^2 + (1 - 6\xi)\Omega^{-1} \nabla^2 \Omega$. All quantities are computed from the same potential. $E_1 = \omega_1^2 - m_0^2 < 0$ indicates bound states. Parameters: $a = 1$, $\varepsilon = 0.2$, $m_0 = 1$, $\xi = 0$.

D	E_1	ω_1	S_1	$r_1 = e^{-2S_1}$	tp	n_{bound}
6	-0.24	0.87	10.65	4.9×10^{-10}	4	2
8	-0.46	0.73	13.15	3.2×10^{-12}	4	3
10	-0.56	0.66	15.45	2.5×10^{-14}	4	3
12	-0.92	0.28	18.85	1.3×10^{-17}	4	1
14	-0.41	0.77	18.52	3.4×10^{-17}	4	2
16	-0.21	0.89	18.81	1.6×10^{-17}	4	3
18	-0.07	0.96	18.97	6.8×10^{-18}	4	3
20	-0.34	0.81	22.79	1.5×10^{-20}	4	1

1. **Conformal factor:** $\Omega(\rho, z; D)$ is explicitly specified (Eq. (9)).
2. **Effective potential:** V_{eff} is derived from Ω (Eq. (14)), with explicit D -dependence.
3. **Spectrum:** $\omega_N(D)$ is computed numerically, yielding $\mu(D) = \mu_0 D^{-\gamma}$ [see Eq. (5)].
4. **Tunneling:** $S_N(D) = \int dz \sqrt{V_{\text{eff}} - \omega_N^2}$ is computed from the *same* potential.
5. **Kinetic rates:** $r_N = e^{-2S_N}$ directly follows from the geometry.

This is a **single-track derivation** for the extraction subchain: there is no separate "toy potential" inside this subsection, and the quantities $\Omega \rightarrow V_{\text{eff}} \rightarrow \omega_N \rightarrow S_N$ are computed from one specified geometry. In the present paper, this action-derived extraction is then propagated into a faster surrogate kinetic scan for the global (D, η) mapping (Section VII and item 3 of the limitations).

IV. MODULE 1: THE THREE ROUTES TO MICRO-DEGENERACY g_N

The factor g_N counts the effective degrees of freedom in layer N . We justify its form via three convergent theoretical routes.

A. Route A: Mock-Modular Counting

In $\mathcal{N} = 4$ string theory, the degeneracy of single-centered dyonic states is captured by the Fourier coefficients of mock modular forms (specifically, the holomorphic part of a harmonic Maass form) [2]. Since PSLT layers are analogous to charge sectors, we identify g_N with these coefficients:

$$g_N^{(\text{A})} = |c_N|, \quad \text{where} \quad \sum c_N q^N = \frac{1}{\eta(\tau)^{\chi}} \times (\text{Mock Piece}). \quad (16)$$

B. Route B: ER=EPR and Entanglement Entropy

Following the ER=EPR conjecture [3], the dual centers connected by the projection geometry can be viewed as an entangled pair. The layer index N corresponds to the excitation level of the wormhole connection. The entropy $S(N) = \ln g_N$ must scale extensively with the effective "area" of the excitation. Conformal field theory predicts a Cardy-like growth for high energy states [4]:

$$S(N) \sim 2\pi \sqrt{\frac{c_{\text{eff}} N}{6}}. \quad (17)$$

C. Route C: Quantum Gravity Template

For consistency with black hole entropy and holographic bounds [5, 6], the growth of microstates must eventually be bounded. This suggests that while the Cardy growth dominates asymptotically, it must be regulated.

D. Unified Cardy-Controlled Envelope with High- N Regulator

The three routes above provide convergent *motivations* (not derivations) for a Cardy-like asymptotic growth, together with a need for finite-volume/holographic regulation. We therefore keep an EFT-level reference envelope with the minimum number of free parameters:

$$g_N(c_{\text{eff}}, \nu, \kappa_g) = \frac{\exp\left(2\pi\sqrt{\frac{c_{\text{eff}}N}{6}}\right)}{N^\nu} \exp\left[-\kappa_g(N-1)^2\right]. \quad (18)$$

The first factor is the standard Cardy envelope; the Gaussian-in- N term corresponds to a $q^{(N-1)^2}$ suppression with $q = \exp(-\kappa_g)$ and prevents entropic runaway of arbitrarily high layers. In this manuscript, Eq. (18) is retained as a legacy comparator and uncertainty reference, while the reported scan baseline uses the bounded two-dimensional phase-space profile.

Summary of motivational roles. Route A (Mock-modular) and Route B (ER=EPR) support the Cardy envelope's plausibility; Route C (QG Template) supports the necessity of a high- N regulator. Equation (18) remains an *EFT-level comparator ansatz* in this version.

E. Cardy Regime Validity and Finite- N Corrections

The asymptotic Cardy formula is valid when the dimensionless entropy $S(N) = 2\pi\sqrt{c_{\text{eff}}N/6}$ satisfies $S(N) \gg 1$. For our demonstrator parameters:

$$S(N) = 2\pi\sqrt{\frac{0.5 \cdot N}{6}} \approx 1.81\sqrt{N}. \quad (19)$$

Thus $S(1) \approx 1.81$, $S(2) \approx 2.57$, $S(3) \approx 3.14$. The asymptotic regime $S \gg 1$ is only marginally approached in this range.

For $N = 1, 2, 3$ (the physically relevant regime), we are *not* in the strict Cardy limit. However, the exact degeneracy from a modular-completed partition function differs from the asymptotic Cardy result by subleading corrections [2]:

$$g_N^{\text{exact}} = g_N^{\text{Cardy}} \left[1 + \mathcal{O}\left(\frac{1}{\sqrt{N}}\right) \right]. \quad (20)$$

We estimate the finite- N error as:

$$\frac{\delta g_N}{g_N} \lesssim \frac{1}{S(N)} \approx \frac{0.55}{\sqrt{N}}. \quad (21)$$

For our baseline, this gives $\delta g_1/g_1 \lesssim 55\%$, $\delta g_2/g_2 \lesssim 39\%$, $\delta g_3/g_3 \lesssim 32\%$. These uncertainties are absorbed into the effective parameters (c_{eff}, ν) , which are calibrated against the phase diagram rather than derived from first principles. (The Gaussian regulator κ_g has negligible effect for $N \leq 3$.)

Key point: We do not claim that Eq. (18) is exact for $N = 1, 2, 3$. Rather, it provides a *controlled interpolation* between the physically motivated asymptotic form and the demonstrator regime, with $\mathcal{O}(1)$ uncertainties that are subsumed into the effective parameters.

F. Assessment of a Phase-Space First-Principles Candidate

We also evaluated a semiclassical phase-space candidate to replace Eq. (18),

$$\rho_{\text{WKB}}(E) \propto \int_{U(z) < E} \frac{dz}{\sqrt{E - U(z)}}, \quad g_N^{(\text{ps})} \sim 1 + \int \rho_{\text{WKB}}(E) dE, \quad (22)$$

with $U(z) = V_{\text{eff}}(z) - m_0^2$ from the same action-derived geometry. The naive 1D implementation is not a drop-in baseline (for bound layers $E_N < 0$, using $(0 \rightarrow E_N)$ is inconsistent with the bound-state threshold convention and can produce nonphysical normalization behavior, including $g_N < 1$ in low layers). We therefore use the converged two-dimensional microcanonical extraction as baseline, while Eq. (18) is kept as a legacy comparator.

V. MODULE 2: RANK-2 COMPUTABLE KINETICS Γ_N

The kinetic rate Γ_N determines how strictly the geometry selects specific layers. We upgrade previous heuristic models to a rigorous Rank-2 system.

A. Dual-Center Tunneling Suppression

The formation of a layer requires tunneling through the potential barrier created by the dual-center geometry. We model this via a WKB action integral:

$$S_N(D) = \int_{x_-}^{x_+} dx \sqrt{V_{\text{eff}}(x; D) - \omega_N^2}, \quad (23)$$

The tunneling probability is then:

$$r_N(D, \eta) = \eta e^{-2S_N(D)}. \quad (24)$$

Here, η is a dimensionless *overlap amplitude* (not a probability), representing the effective strength of the dual-center interaction; values $\eta > 1$ are permitted, corresponding to collective or multi-channel enhancement. An action-derived prefactor candidate $\eta_{\text{fp}}(D)$ extracted from splitting-action data is reported in Appendix E 8 as a conservative replacement direction; it is not used in the baseline scan.

B. Dimensionality and Units

We adopt natural units $\hbar = c = 1$. The theory is defined relative to a fundamental mass scale M_* (set to unity). The dual-center separation is parameterized by a dimensionless ratio D (representing physical separation in units of the characteristic length scale M_*^{-1}). The field frequencies scale as $\omega_N \sim M_*/D$.

C. Derivation of the Rank-2 Growth Matrix

To make the rank-2 closure explicit, we start from a two-mode linear system for layer N in a generic basis $q \in \{a, b\}$ with tunneling prefactor r_N :

$$\frac{d}{dt} \begin{pmatrix} n_{N,a} \\ n_{N,b} \end{pmatrix} = r_N \begin{pmatrix} \Gamma_{N,a} & \epsilon_{\text{mix}} \\ \epsilon_{\text{mix}} & \Gamma_{N,b} \end{pmatrix} \begin{pmatrix} n_{N,a} \\ n_{N,b} \end{pmatrix}, \quad (25)$$

where $n_{N,q}$ is the coarse-grained occupation in channel q . Equation (25) is the probability-level (coarse-grained) form of a coupled-mode system after absorbing convention-dependent amplitude factors into (A_ℓ, χ) . The matrix multiplying $(n_{N,a}, n_{N,b})^T$ is therefore the growth matrix $\mathbf{M}_N$ used below.

The off-diagonal term is modeled as

$$\epsilon_{\text{mix}} = \chi \sqrt{\Gamma_{N,a} \Gamma_{N,b}}, \quad (26)$$

which is consistent with a weak-coupling overlap estimate $\epsilon_{\text{mix}} \sim \int \psi_{N,1}^* \delta V \psi_{N,2} d^3x$ and guarantees the correct rate dimension.

D. Rank-2 Computable Closure

The formation rate Γ_N is derived from the positive eigenvalue of the rank-2 interaction matrix $\mathbf{M}_N$:

$$\Gamma_N = \max(0, \lambda_+) \quad (27)$$

where λ_+ is the largest real eigenvalue of:

$$\mathbf{M}_N = r_N \begin{pmatrix} \Gamma_{N,a} & \epsilon_{\text{mix}} \\ \epsilon_{\text{mix}} & \Gamma_{N,b} \end{pmatrix} \quad (28)$$

Here, $r_N = \eta e^{-2S_N}$ is the tunneling suppression factor. The mixing term is parameterized by a dimensionless coupling χ :

$$\epsilon_{\text{mix}} = \chi \sqrt{\Gamma_{N,a} \Gamma_{N,b}} \quad (29)$$

The channel rates $\Gamma_{N,\ell}$ are modeled by a Kerr-inspired barrier-leakage width proxy scaling [7]. In the present static two-center background this is used as an EFT-level channel-rate template (not a literal ergoregion superradiance derivation). To avoid dimensional ambiguity, we define

$$\hat{\omega}_N \equiv \frac{\omega_N}{M_*}, \quad \hat{\Gamma}_{N,\ell} \equiv \frac{\Gamma_{N,\ell}}{M_*}, \quad (30)$$

and write

$$\hat{\Gamma}_{N,\ell} = A_\ell(D) \hat{\omega}_N^{4\ell+5}, \quad (31)$$

which is equivalent to

$$\Gamma_{N,\ell} = A_\ell(D) \omega_N \left(\frac{\omega_N}{M_*} \right)^{4\ell+4}, \quad (32)$$

The legacy demonstrator corresponds to $A_\ell(D) = 1$, while the reported baseline uses the action-derived profile factors $A_\ell(D) = \tilde{A}_\ell(D)$ from Appendix E9. In the two-channel truncation, we map the two-mode basis to the two lowest leakage channels:

$$\Gamma_{N,a} \equiv \Gamma_{N,1}, \quad \Gamma_{N,b} \equiv \Gamma_{N,2}, \quad \bar{\Gamma}_N \equiv \sqrt{\Gamma_{N,a} \Gamma_{N,b}}. \quad (33)$$

Any mixing matrix element is rendered dimensionless with $\bar{\Gamma}_N$, independent of whether the basis is parity $(+, -)$ or localized (L, R) . Higher- ℓ modes are parametrically suppressed by the $\hat{\omega}_N^{4\ell+4}$ scaling and are deferred to future extensions. All baseline scans are reported in **dimensionless units** ($M_* = 1$), where Eq. (31) and Eq. (32) are numerically identical. Constant- A_ℓ results are retained only as comparators. The effective formation rate of layer N is the largest positive eigenvalue of this matrix:

$$\boxed{\Gamma_N(D, \eta) = \max(0, \lambda_+(\mathbf{M}_N(D, \eta)))}. \quad (34)$$

For completeness, the analytic expression for λ_+ is:

$$\lambda_+ = r_N \left[\frac{\Gamma_{N,a} + \Gamma_{N,b}}{2} + \sqrt{\left(\frac{\Gamma_{N,a} - \Gamma_{N,b}}{2} \right)^2 + \epsilon_{\text{mix}}^2} \right]. \quad (35)$$

This ‘‘Rank-2 Closure’’ ensures that Γ_N is not a fitted parameter but a derived quantity dependent explicitly on (D, η) .

E. Physical Interpretation: Quasi-Bound Leakage and Width Proxy

The channel scaling in Eq. (32) is interpreted as a quasi-bound leakage-width proxy in the two-center geometry. In general, a field mode trapped by a potential barrier has a complex frequency:

$$\omega_N = \omega_N^{(R)} - i\omega_N^{(I)}, \quad \Gamma_N^{(\text{width})} \equiv 2\omega_N^{(I)}, \quad (36)$$

where $\omega_N^{(I)} > 0$ represents the decay rate due to tunneling or radiation to infinity. This width proxy motivates the channel factors $\Gamma_{N,\ell}$; it is distinct from the final matrix-derived formation rate $\Gamma_N(D, \eta)$ in Eq. (34).

For Kerr superradiance [7], the scaling $\Gamma \propto (\omega M)^{4\ell+5}$ arises from a rotating-background setup. In the static dual-center geometry considered here (no ergoregion), we use the same power as a leakage-template ansatz: modes must tunnel through the geometric barrier (captured by WKB factor e^{-2S_N}), while $\Gamma_{N,\ell}$ encodes the residual channel-dependent escape width. A full non-Hermitian width derivation for this static geometry is left as a dedicated upgrade path.

Reviewer-facing status of the static-width upgrade: In the strict static formulation, the relevant object is the resonance problem for the same action-derived operator with outgoing boundary conditions, so that $\Gamma_{N,\ell}^{(\text{strict})} = -2\Im\omega_{N,\ell}$. We implemented a minimal complex-absorbing-potential pilot on the same 1D operator chain as Appendix E9 and scanned the benchmark set $D = \{6, 12, 18\}$ for the representative $(N, \ell) = (3, 1)$ mode. The numerical chain is stable as infrastructure, but the extracted widths remain absorber-dominated over the tested (η, z_{cap}) windows and do not yet exhibit a clean plateau. For that reason, the CAP outputs are not promoted into the reported baseline. At the current stage, the most conservative first-principles static surrogate is therefore the geometry-driven rate $\Gamma_{N,\ell}^{(\text{geo})} = \omega_N e^{-2S_{N,\ell}}$ from Appendix E9, while the full complex-pole extraction remains an explicit next-step upgrade rather than an over-claimed result.

Rank-2 truncation justification: The $\hat{\omega}_N^{4\ell+4}$ scaling ensures that higher- ℓ modes are exponentially suppressed for $\hat{\omega}_N < 1$ (our regime). Concretely, for $\hat{\omega}_N \sim 0.1$:

$$\frac{\Gamma_{N,3}}{\Gamma_{N,2}} \sim \hat{\omega}_N^4 \sim 10^{-4}. \quad (37)$$

Thus the $\ell = 1, 2$ truncation still captures the dominant contribution in the demonstrator regime; higher- ℓ channels are subleading and deferred to future work.

Mixing term interpretation: The off-diagonal term ϵ_{mix} can be interpreted as the mode-overlap integral between the $\ell = 1$ and $\ell = 2$ wavefunctions via the non-spherical part of the

potential:

$$\epsilon_{\text{mix}} \sim \int \psi_{N,1}^* \delta V \psi_{N,2} d^3x, \quad (38)$$

where δV is the deviation of V_{eff} from spherical symmetry. With the explicit spherical-average definition used in Eq. (73), this overlap is symmetry-suppressed for the parity-even/odd pair and numerically yields $\chi_N^{(\text{sym})} \approx 0$ (Appendix C). Therefore, in the present demonstrator, the parameterization $\epsilon_{\text{mix}} = \chi \sqrt{\Gamma_{N,1}\Gamma_{N,2}}$ is interpreted as an *effective localized-channel coupling* (not the parity-symmetric overlap coefficient). The basis transformation between parity modes and localized modes is a unitary change of representation and is mathematically legitimate; it does not change the underlying operator spectrum. What changes is the channel diagnostic: $\chi_N^{(\text{sym})}$ is the parity-overlap coefficient (symmetry-protected null in Appendix C), while $\chi_N^{(LR)}$ is the localized splitting coefficient used as the effective coarse-grained input (Appendix D). We report both quantities explicitly to avoid basis-mixing ambiguity. Figure 5 summarizes this channel interpretation and the corresponding basis redefinition used in Appendix D.




FIG. 5. Physical interpretation of the mixing channel. In the parity-symmetric basis, overlap cancellation yields $\chi_N^{(\text{sym})} \approx 0$. In the localized basis, the level splitting ΔE defines the effective coupling $\chi_N^{(LR)}(D)$ used in the global scan.

VI. MODULE 3: EFT-NORMALIZED VISIBILITY B_N

In early versions of the theory, B_N was an arbitrary matching factor. In the present EFT-operator-normalized baseline, $B_N(D)$ is defined by fully normalized EFT operator kernels built from overlap-extracted flavor-layer couplings and mediator scales from the same 2D localized operator chain as Appendix D; finite one-loop matching and LL-RG normalization are applied before layer-wise normalization. Mode tracking and microcanonical-window averaging suppress

branch-switch artifacts in the overlap extraction stage, while high- N regularization remains isolated in Eq. (18).

For comparison with the earlier visibility closure, we retain a legacy Yukawa-normalized fallback based on the **charged-lepton Yukawa** couplings:

$$Y_1 \equiv y_e, \quad Y_2 \equiv y_\mu, \quad Y_3 \equiv y_\tau, \quad y_\ell = \sqrt{2} m_\ell / v, \quad (39)$$

and define a cumulative effective strength:

$$\tilde{Y}_N \equiv \sum_{k=1}^{\min(N,3)} Y_k, \quad N \geq 1. \quad (40)$$

This cumulative definition restricts the legacy visibility normalization to the three observable generations. As a **legacy comparator/fallback**, we keep the Yukawa-derived visibility law with sublinear compression exponent $0 < p_B < 1$:

$$B_N = \left(\frac{\tilde{Y}_N}{\tilde{Y}_3} \right)^{p_B}, \quad N = 1, 2, 3, \quad B_3 \equiv 1. \quad (41)$$

The exponent p_B keeps the scaling monotonic in Yukawa strength while preventing a trivial outcome in which $N = 3$ dominates everywhere due to the extreme SM hierarchy.¹ Numerically, using PDG 2024 inputs [8] and $p_B = 0.30$, we obtain representative values $B_1 \simeq 0.085$, $B_2 \simeq 0.42$, $B_3 = 1$ (normalized).

A. High- N Neutral Saturation

For layers $N > 3$, we set $B_N = B_3 = 1$ (neutral saturation). In the baseline scan, high- N control is provided by the bounded two-dimensional phase-space tail prescription and finite-volume clipping in the g_N module; the κ_g factor in Eq. (18) is retained only for legacy Cardy-comparator runs. This keeps the visibility block from introducing an additional high- N damping factor on top of the phase-space regulator.² This saturation should not be read as an ad hoc removal of a fourth layer: it is deliberately neutral, in the sense that once $N > 3$ the visibility block inserts no additional suppression beyond $B_3 = 1$. In the current action-derived benchmark, no audited D point supports a fourth bound level on the single-track spectrum, and the companion one-dimensional kinetic proxy keeps the candidate $N = 4$ excitation above the continuum threshold, $E_4 = \omega_4^2 - m_0^2 > 0$, throughout the tested window. Moreover, forcing the

¹ An alternative using up-type quark Yukawas was explored; using lepton Yukawas removes sector ambiguity and enables a direct one-to-one mapping between the $N = 2$ layer and the muon-generation observable.

² An alternative approach with explicit B_N tail suppression ($B_N \propto e^{-\beta_B (N-3)^2}$) was explored; see Appendix F for comparison.

same quasi-channel continuation to $N = 4$ does not produce an additional small-width hierarchy relative to $N = 3$. The present evidence therefore favors a threshold/binding-loss interpretation of the high- N decoupling over an engineered visibility cutoff or an already-established $\Gamma_4 \rightarrow 0$ theorem; a sharper mode-resolved derivation remains a next-step first-principles target rather than a claimed result here.

B. Roadmap: Action-Derived Effective Yukawa Operator

With the EFT-operator-normalized baseline now in place, the release-production chain separates direct spectral-selection closure from visibility closure. The runtime direct-spectral branch recomputes $g_N(D)$, $\chi_{LR}(D)$, and $\tilde{A}_\ell(D)$ at scan runtime while keeping the baseline $B_N(D)$ profile closure; the runtime-direct visibility path is release-qualified only as a risk-weighted profile-blend parity branch, not as a strict all-direct Higgs-map closure. The strict no-cache variant is retained only as a stress diagnostic. The next upgrade remains an explicit lepton Yukawa sector in the same conformal background:

$$\mathcal{L}_Y = -y_0 \bar{L} H \ell_R + \text{h.c.}, \quad g_{\mu\nu} = \Omega^2 \eta_{\mu\nu}. \quad (42)$$

In the minimal conformal setup, we use the vierbein form

$$S_\Psi = \int d^4x \sqrt{-g} \, i \bar{\Psi} e^\mu_a \gamma^a \left(\partial_\mu + \frac{1}{4} \omega_{\mu bc} \gamma^{bc} \right) \Psi, \quad e_\mu^a = \Omega \delta_\mu^a, \quad (43)$$

with $\sqrt{-g} = \Omega^4$ and $e^\mu_a = \Omega^{-1} \delta^\mu_a$. The canonical Weyl rescaling

$$\Psi = \Omega^{-3/2} \psi, \quad H = \Omega^{-1} \varphi \quad (44)$$

removes the spin-connection derivative term and yields canonical flat-frame kinetic terms (Appendix B 3). The Yukawa block then scales as

$$\sqrt{-g} \bar{\Psi} H \Psi \propto \Omega^4 \cdot \Omega^{-3/2} \cdot \Omega^{-1} \cdot \Omega^{-3/2} = \Omega^0, \quad (45)$$

so the minimal baseline has no residual conformal prefactor in the overlap kernel. After fixing this normalization convention, we decompose the Higgs fluctuation on the same action-derived spatial modes used in Appendix D,

$$H(x, \rho, z) = \sum_N h_N(x) u_N(\rho, z; D), \quad (46)$$

and define a mode-resolved effective coupling by overlap:

$$y_N^{\text{eff}}(D) = y_0 \int 2\pi \rho \, d\rho \, dz \, u_N(\rho, z; D) f_L(\rho, z) f_R(\rho, z) \mathcal{W}_{\text{frame}}(\rho, z; D). \quad (47)$$

Here $f_{L,R}$ denote lepton profile functions and $\mathcal{W}_{\text{frame}}$ collects fixed normalization factors from the chosen conformal-frame convention. In the minimal Dirac-conformal baseline implied by Eqs. (44)–(45), $\mathcal{W}_{\text{frame}} = 1$. Nonzero conformal-frame powers are kept only as diagnostic sensitivity directions and are not used in baseline figures. In this roadmap, the same 2D localized solver and convergence criteria as Appendix D are reused; scan-level runs now separate release-production per-cell direct $g_N + \chi + A$ closure from the promoted risk-weighted profile-blend runtime-direct visibility parity branch (canonical D-only profile blend $\alpha(D)$) and the strict no-cache stress branch.

VII. VERIFICATION RESULTS

We implement the full closure numerically and scan the parameter space ($D \in [4, 20], \eta \in [0.2, 4.0]$).

A. Configuration and Parameter Table

Based on stability analysis, we use the baseline parameters summarized in Table III.

B. Three-Generation Phase Diagram

The winner map $N^*(D, \eta)$ (Fig. 6) reveals a robust structure:

- For $D \lesssim 8$, Layer 3 ($N = 3$) dominates.
- For $D \gtrsim 8$, Layer 2 ($N = 2$) dominates.
- Layer 1 ($N = 1$) is enhanced by B_1 but is kinetically suppressed at large D .

Crucially, we define the **Generation Ratio**:

$$\mathcal{R}_3 \equiv \sum_{N=1}^3 P_N = \frac{\sum_{N=1}^3 W_N}{\sum_{K=1}^{N_{\text{max}}} W_K}. \quad (48)$$

With full scan-grid action-derived profiles for $\chi_N^{(LR)}(D)$ and $\tilde{A}_\ell(D)$ from Appendix D, our numerical scan gives $\mathcal{R}_3 > 90\%$ over 92.7% of the sampled (D, η) grid, while $f(\mathcal{R}_3 > 0.95) = 0.2167$ in this setup. The profile support is anchored by the benchmark convergence set (Table XII) and the auxiliary full-scan support points (Table XIII) in Appendix D. This still supports a dominant three-layer sector; the current baseline combines action-derived g_N , $\chi_N^{(LR)}(D)$, and $\tilde{A}_\ell(D)$ profiles with EFT-operator-normalized $B_N(D)$. To quantify transfer bias from this surrogate

TABLE III. Baseline parameters for the PSLT demonstrator. All dimensionful quantities are in units of M_* .

Module	Parameter	Value	Physical Role
Micro-degeneracy g_N	Phase-space	Bounded 2D	Baseline mode: bounded low- N
	profile		window + energy-gap tail profile
	(α, β, s)	$(0.8, 1.1, 0.0)$	low- N window blend, tail damping, shell-power control
	$c_{\text{eff}}, \nu, \kappa_g$	$(0.5, 5.0, 0.03)$	Legacy Cardy-comparator parameters (not used for the bounded 2D baseline normalization)
Visibility B_N	Visibility normalization	EFT operator	Baseline fully normalized EFT visibility from overlap-extracted operator inputs (g_{iN}, λ_N) with finite one-loop + LL-RG normalization
	p_B	0.30	Legacy Yukawa comparator/fallback parameter (inactive in EFT-operator baseline)
Kinetics Γ_N	$\chi_N^{(LR)}(D)$	App. D	Localized-channel profile (action-derived extraction; baseline now uses full scan-grid profile injection, early demonstrator used constant $\chi = 0.2$)
	a_0	0.02	Geometric perturbation
	ϵ	0.2	Core regularization
	$\tilde{A}_1(D), \tilde{A}_2(D)$	App. E 9	Action-derived channel-normalization profile (constant A_ℓ kept as legacy comparator)
Dynamics	t_{coh}	1.0	Coherence time (M_*^{-1})

interpolation step, we performed a point-level audit with direct localized fine extraction injected at each queried D (recomputing χ_{LR} by the same 2D solver instead of using interpolated knots). The reproducible script is `scan_surrogate_vs_action_points.py`, with source table `paper/surrogate_vs_action_points.csv`. As summarized in Table IV, the direct-injection path changes χ by up to $\sim 31\%$ at representative points, while winner labels remain unchanged (3/3 matches) and $\mathcal{R}_3$ stays numerically stable ($\max |\Delta \mathcal{R}_3| = 5.27 \times 10^{-11}$); under EFT/Wilson-matched observables, $\mu_{\mu\mu}^{\text{pred}}$ can vary strongly at off-band points, so acceptance-level robustness is additionally tracked by the map summaries and Tables V, XXII.

To propagate finite- N surrogate uncertainty into map-level systematics, we also perform




FIG. 6. Left: Winner phase diagram showing regions dominated by Gen 2 and Gen 3. Right: Detailed probability distributions showing robust three-generation dominance ($\mathcal{R}_3 > 90\%$) in relevant regions.

TABLE IV. Point-level surrogate-vs-direct audit for the mixing-profile transfer. “Surrogate” uses $\chi_{LR}(D)$ from $D = \{6, 12, 18\}$ interpolation; “direct” recomputes fine-grid localized extraction at each listed D . Source CSV: `paper/surrogate_vs_action_points.csv`.

(D, η)	winner (sur/direct)	χ_{interp}	χ_{direct}	$ \Delta\mathcal{R}_3 $	$ \Delta\mu_{\mu\mu}^{\text{pred}} $	match
(4.0, 0.2)	2/2	4.02×10^{-4}	5.27×10^{-4}	5.27×10^{-11}	1.58×10^3	yes
(5.0, 2.0)	2/2	4.02×10^{-4}	3.31×10^{-4}	1.02×10^{-14}	9.91×10^2	yes
(11.0, 1.4)	2/2	2.51×10^{-4}	3.06×10^{-4}	0	2.85×10^{-2}	yes

a one-at-a-time local-parameter robustness scan around the baseline closure values: $c_{\text{eff}} \in [0.45, 0.55]$, $\nu \in [4.80, 5.20]$, and the legacy visibility-comparator exponent $p_B \in [0.28, 0.32]$ (with the same 60×60 map and full-grid action-derived $\chi_{LR}(D)/\tilde{A}_\ell(D)$ profiles). Results are summarized in Table V. Since p_B is inactive in the EFT-operator-normalized baseline, that row is retained as a legacy-comparator sensitivity check rather than an active closure knob. Across this local window, the key fractions remain unchanged: $f(\mathcal{R}_3 > 0.90) = 0.9269$, $f(\chi_{\mu\mu}^2 < 4) = 0.1500$, and $f(N_{\text{win}} > 3) \approx 2.78 \times 10^{-4}$.

To make the full direct chain auditable at release level, we compare the strict baseline map with two localized-direct validation surfaces and visibility-parity checks; Table VI summarizes the result. The direct spectral branch recomputes the two-dimensional phase-space spectra for $g_N(D)$ and evaluates $\chi_{LR}(D)$ and $\tilde{A}_\ell(D)$ on the active validation grids. On both audit grids it matches the full direct release map exactly: $\Delta f(\chi_{\mu\mu}^2 < 4) = 0$, mismatch fraction = 0, and $\max|\Delta\mu_{\mu\mu}| = 0$. The tuned visibility parity branch keeps a D-only risk-weighted profile blend $\alpha(D)$ with $(\alpha_{\min}, \alpha_{\max}, p) = (0.96, 0.99, 1.0)$, mean $\alpha = 0.9637$, and $p90(\alpha) = 0.9734$. It passes both validation surfaces (small complete grid: $\max|\Delta\mu_{\mu\mu}| = 0.0925$; large spot-check grid: $\max|\Delta\mu_{\mu\mu}| = 0.8373$), but remains explicitly blend-weighted in the observable sector. We therefore read it as a **risk-weighted profile-blend visibility parity branch**, not as a strict all-direct Higgs-map closure. A no-profile stress branch is retained only for diagnostics. The earlier de-profiled visibility stress branch remains useful as a historical neg-

TABLE V. Local one-at-a-time robustness scan around baseline closure parameters, with the p_B row retained only as a legacy visibility-comparator sensitivity because p_B is inactive in the EFT-operator-normalized baseline. Fractions are measured on the same 60×60 map as the baseline scan. Source CSV: `paper/core_param_robustness.csv`.

Parameter	Window	$f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	Max drift
c_{eff}	0.45/0.50/0.55	0.9269/0.9269/0.9269	0.1500/0.1500/0.1500	$\Delta f_{\mathcal{R}_3} = 0, \Delta f_{\chi^2} = 0$
ν	4.80/5.00/5.20	0.9269/0.9269/0.9269	0.1500/0.1500/0.1500	$\Delta f_{\mathcal{R}_3} = 0, \Delta f_{\chi^2} = 0$
p_B	0.28/0.30/0.32	0.9269/0.9269/0.9269	0.1500/0.1500/0.1500	$\Delta f_{\mathcal{R}_3} = 0, \Delta f_{\chi^2} = 0$

TABLE VI. Release-mode full direct map summary; exact source artifacts are listed in the reproducibility registry.

scenario	grid	n_{points}	$f(\chi_{\mu\mu}^2 < 4)$	best $\chi_{\mu\mu}^2$	frac winner mismatch	$\max \Delta\mu_{\mu\mu} $
full direct baseline	D60×E60	3600	0.1428	1.56×10^{-3}	—	—
small-surface direct bias	D21×E41	861	0.1905	1.45×10^{-5}	0	3.37×10^{-3}
large-surface direct bias	D60×E21	1260	0.1500	5.10×10^{-4}	0	0
runtime parity (small)	D21×E41	861	0.0952	2.11×10^{-1}	0	0
tuned parity (small)	D21×E41	861	0.0952	2.11×10^{-1}	0	9.25×10^{-2}
runtime parity (large)	D60×E21	1260	0.1429	1.56×10^{-3}	0	0
tuned parity (large)	D60×E21	1260	0.1429	1.56×10^{-3}	0	8.37×10^{-1}
extreme parity (large)	D60×E21	1260	0.1429	1.56×10^{-3}	0.0373	59.3

ative mechanism audit: it showed that removing the visibility profile too aggressively reopens large observable drift even when some local mismatch gates improve. That negative result has now been superseded operationally by two cleaner endpoints. First, the strict broader-grid all-direct validator is parked at its current case-aware endpoint, with zero acceptance mismatch and $p95|\Delta\mu_{\mu\mu}| = (0.0560, 0.1097)$ on the two validation surfaces. Second, the canonical reviewer observable branch is now the two-lobe full-width/reference-amplitude validator, which has moved from promotion work to routine surveillance. We therefore treat additional local $\alpha(D)$ retunes or de-profiled stress variants as diagnostics only unless a genuinely new observable-side mechanism class is introduced.

For the UV+LL-RG observable controls in the same baseline mode, we also scan one-at-a-time windows in $(\mu_{\text{low}}, \gamma_{\text{diag}}, \gamma_{\text{offdiag}}, \kappa_{\text{diag}}, \kappa_{\text{offdiag}})$. Table VII summarizes the result (source CSV: `paper/hll_rge_sensitivity.csv`); across all tested windows, $f(\chi_{\mu\mu}^2 < 4)$ stays fixed at 0.1500. To capture residual map-level drift in a nonzero finite-match setting, we additionally run a full-direct UV envelope scan with `code/scan_hll_uv_envelope.py`. On the D21×E41 map the acceptance span remains zero ($f_{\text{min}} = f_{\text{max}} = 0.1243$), while the pointwise half-width has mean 3.03×10^{-4} , $p95 = 1.50 \times 10^{-3}$, and max 3.89×10^{-3} (source files: `paper/hll_uv_e`

TABLE VII. One-at-a-time robustness scan for UV+LL-RG observable controls. Source CSV: `paper/hll_rge_sensitivity.csv`.

parameter	window (low/base/high)	$f(\chi_{\mu\mu}^2 < 4)$ low/base/high	max abs drift
μ_{low}	0.5/1.0/2.0	0.1500/0.1500/0.1500	0
γ_{diag}	1.0/2.0/3.0	0.1500/0.1500/0.1500	0
γ_{offdiag}	0.5/1.0/1.5	0.1500/0.1500/0.1500	0
κ_{diag}	-1.0/0.0/1.0	0.1500/0.1500/0.1500	0
κ_{offdiag}	-1.0/0.0/1.0	0.1500/0.1500/0.1500	0

`nvelope_summary.csv` and `paper/hll_uv_envelope_map.csv`).

C. $H \rightarrow \mu\mu$ Signal Strength (EFT/Wilson-Matched Mapping, UV+LL-RG Baseline)

We confront the theory with the ATLAS Run-2+Run-3 combined $H \rightarrow \mu\mu$ signal-strength input, $\mu_{\mu\mu}^{\text{obs}} = 1.4 \pm 0.4$ (combined uncertainty) [9]. The same ATLAS evidence result reports a Run-3-only signal strength near $\mu_{\mu\mu} = 1.6 \pm 0.6$; all scan maps in this manuscript use the combined input above.

Within the PSLT closure, we construct a scan-level EFT/Wilson-matched map (evaluated in baseline `eft_wilson_uv_rge` mode) by combining kinetic occupancy with overlap-extracted raw couplings:

$$\tilde{q}_N(D, \eta) = g_N(D) \left(1 - e^{-\Gamma_N(D, \eta) t_{\text{coh}}}\right), \quad P_N^{(\text{kin})}(D, \eta) = \frac{\tilde{q}_N(D, \eta)}{\sum_{K=1}^{N_{\text{max}}} \tilde{q}_K(D, \eta)}. \quad (49)$$

$$\epsilon_{\text{mix}}(D, \eta) = \text{clip} \left[\kappa_{\text{mix}} \chi_{\text{eff}}(D) \left(\frac{\eta}{\eta_{\text{ref}}} \right)^{p_\eta}, 0, \epsilon_{\text{max}} \right], \quad (50)$$

$$\Pi(\epsilon_{\text{mix}}) = \begin{pmatrix} 1 - \epsilon_{\text{mix}} & \epsilon_{\text{mix}} & 0 \\ \epsilon_{\text{mix}}/2 & 1 - \epsilon_{\text{mix}} & \epsilon_{\text{mix}}/2 \\ 0 & \epsilon_{\text{mix}} & 1 - \epsilon_{\text{mix}} \end{pmatrix}, \quad Y_{iN}(D, \eta) = \sqrt{y_N^{\text{eff, raw}}(D)} \Pi_{iN}(\epsilon_{\text{mix}}), \quad (51)$$

The bounded projector in Eqs. (50)–(51) is retained for matched-mode diagnostics; the baseline `eft_wilson_uv_rge` maps use the UV-tree closure below.

$$C_{eH}^{\text{tree, ij}}(D, \eta) = \sum_{N=1}^3 g_{iN}(D) \frac{P_N^{(\text{kin})}(D, \eta)}{M_N^2(D)} g_{jN}(D), \quad (52)$$

$$C_{eH}^{\text{match, ij}} = \begin{cases} \left(1 + \frac{\kappa_{\text{diag}}}{16\pi^2}\right) C_{eH}^{\text{tree, ij}}, & i = j, \\ \left(1 + \frac{\kappa_{\text{offdiag}}}{16\pi^2}\right) C_{eH}^{\text{tree, ij}}, & i \neq j, \end{cases} \quad (53)$$

$$C_{eH}^{\text{IR}} = C_{eH}^{\text{match}} + \frac{\ln(\mu_{\text{match}}/\mu_{\text{low}})}{16\pi^2} \gamma \odot C_{eH}^{\text{match}}, \quad (54)$$

where $\gamma \odot C$ denotes blockwise application of $(\gamma_{\text{diag}}, \gamma_{\text{offdiag}})$ to diagonal and off-diagonal matrix entries, respectively. In the baseline `eft.wilson-uv-rge` mode used for figures, we evaluate observables with $C_{eH} \equiv C_{eH}^{\text{IR}}$.

$$\frac{\Gamma_{\text{tot}}(D, \eta)}{\Gamma_{\text{tot}}(D_0, \eta_0)} = 1 + s_\Gamma \sum_{\ell=e, \mu, \tau} \text{BR}_\ell^{\text{SM}} \left[\left| \frac{C_{eH}^{\ell\ell}(D, \eta)}{C_{eH}^{\ell\ell}(D_0, \eta_0)} \right|^2 - 1 \right], \quad (55)$$

$$\mu_{\mu\mu}^{\text{pred}}(D, \eta) = \left| \frac{C_{eH}^{\mu\mu}(D, \eta)}{C_{eH}^{\mu\mu}(D_0, \eta_0)} \right|^2 \left[\frac{\Gamma_{\text{tot}}(D, \eta)}{\Gamma_{\text{tot}}(D_0, \eta_0)} \right]^{-1}. \quad (56)$$

In the baseline figures we keep a fixed reference point $(D_0, \eta_0) = (10, 1)$ for comparability across scan variants. This fixed point should be read as a *reference-normalization convention* for the relative scan-level observable in Eq. (56), not as a physical anchor of the Higgs sector itself. The stronger normalization-side information now comes from the absolute and loop-improved parent-action comparators below together with the reference-normalization-invariant Wilson-diagonal ratios reported later in the section. Equation (56) is implemented in `pslt-lib.py` as the UV+LL-RG Higgs observable; it remains a scan-level EFT module rather than a full loop-level EYMH→EFT derivation. Writing

$$x_\mu \equiv |C_{eH}^{\mu\mu}|^2, \quad R_{e/\mu} \equiv \frac{|C_{eH}^{ee}|^2}{|C_{eH}^{\mu\mu}|^2}, \quad R_{\tau/\mu} \equiv \frac{|C_{eH}^{\tau\tau}|^2}{|C_{eH}^{\mu\mu}|^2},$$

Eq. (56) may also be rewritten exactly as

$$\mu_{\mu\mu}^{\text{pred}} = \frac{A_\star x_\mu}{1 + s_\Gamma \left[A_\star x_\mu \left(\text{BR}_\mu^{\text{SM}} + \text{BR}_e^{\text{SM}} \frac{R_{e/\mu}}{R_{e/\mu}^\star} + \text{BR}_\tau^{\text{SM}} \frac{R_{\tau/\mu}}{R_{\tau/\mu}^\star} \right) - \sum_{\ell=e, \mu, \tau} \text{BR}_\ell^{\text{SM}} \right]}, \quad (57)$$

with

$$A_\star = x_{\mu, \star}^{-1}, \quad R_{e/\mu}^\star = \frac{|C_{eH, \star}^{ee}|^2}{|C_{eH, \star}^{\mu\mu}|^2}, \quad R_{\tau/\mu}^\star = \frac{|C_{eH, \star}^{\tau\tau}|^2}{|C_{eH, \star}^{\mu\mu}|^2}.$$

So the current reference-normalization arbitrariness is not an unconstrained map-level rescaling; it is exactly equivalent to three scalar bridge constants. The audit `code/audit_hll_absolute_normalization_bridge.py` on the canonical parented D21×E21 map (`output/hll_absolute_normalization/hll_absolute_normalization_bridge_summary.csv`) verifies Eq. (57) to machine precision ($\max |\Delta \mu_{\mu\mu}| = 5.56 \times 10^{-8}$). More importantly, the accepted-region invariant median already fixes the electron-to-muon ratio essentially uniquely: if $R_{e/\mu}^\star$ is frozen to the invariant audit median, a two-parameter fit in $(A_\star, R_{\tau/\mu}^\star)$ still closes the canonical map at the same 5.56×10^{-8} level with zero acceptance mismatch. At this bridge-reduction

stage the observable-side problem is therefore already narrower than a generic reference-anchor ambiguity: the live scalar constants are reduced to the absolute amplitude $A_\star$ and the tau-to-muon ratio $R_{\tau/\mu}^\star$, while $R_{e/\mu}^\star$ is effectively fixed by the invariant-ratio audit itself. The next source audit sharpens this split further. Because the present finite-match and LL-RG layers apply the same diagonal dressing to all three charged-lepton diagonals,

$$C_{eH}^{\ell\ell, \text{IR}} = Z_{\text{diag}} C_{eH}^{\ell\ell, \text{tree}}, \quad Z_{\text{diag}} = \left(1 + \frac{\kappa_{\text{diag}}^{\text{eff}}}{16\pi^2}\right) \left(1 + \frac{\gamma_{\text{diag}} \log(\mu_{\text{match}}/\mu_{\text{low}})}{16\pi^2}\right),$$

the flavor ratios $R_{e/\mu}$ and $R_{\tau/\mu}$ are exact UV-tree invariants of the current diagonal chain rather than additional loop-level normalization freedoms. The audits `code/audit_hll_absolute_amp_source.py` and `code/audit_hll_tau_ratio_source.py` (canonical outputs under `output/hll_absolute_normalization/`) verify this statement on the D21×E21 parented map: on active rows the residuals between UV and IR ratios remain at machine level ($\max |\Delta R_{e/\mu}| = 6.05 \times 10^{-13}$, $\max |\Delta R_{\tau/\mu}| = 3.23 \times 10^{-14}$), and at the live reference point the tau-to-muon constant is identical across tree/match/IR ($R_{\tau/\mu}^\star = 8.24039 \times 10^{-3}$). The same audit also shows that the absolute amplitude itself factorizes exactly as

$$A_\star = A_\star^{(\text{tree})} Z_{\text{diag}, \star}^{-2}, \quad A_\star^{(\text{tree})} \equiv |C_{eH, \star}^{\mu\mu, \text{tree}}|^{-2},$$

with map-level factorization residual $\max |\Delta C_{\mu\mu}^{\text{fact}}| = 9.32 \times 10^{-17}$. Numerically the remaining gap is now concentrated almost entirely in the reference tree amplitude: $A_\star^{(\text{tree})} = 5.26537 \times 10^8$, $A_\star = 5.20717 \times 10^8$, so the full diagonal dressing shifts the absolute normalization by only about 1.11%. Using the live ratio constants, a tree-only replacement for $A_\star$ already preserves the accepted region exactly, while restoring only the universal LL-RG factor reduces the map-level bridge error to $\max |\Delta \mu_{\mu\mu}| = 1.08 \times 10^{-3}$. Within the current diagonal UV+LL-RG chain, the outstanding parent-side normalization target is therefore primarily the absolute amplitude $A_\star$ itself rather than an unresolved flavor-ratio constant. The next source audit tightens this target once more. In the canonical parented baseline one has `uv_blend` = 0, so the UV-tree coupling matrix reduces to a diagonal overlap form,

$$g_{iN}^{\text{UV}}(D) = \delta_{iN} \sqrt{y_i^{\text{raw}}(D)},$$

and therefore the muon tree coefficient collapses exactly to a single layer:

$$C_{eH}^{\mu\mu, \text{tree}}(D, \eta) = y_2^{\text{raw}}(D) \frac{P_2^{\text{kin}}(D, \eta)}{M_2^2(D)}. \quad (58)$$

The audit `code/audit_hll_tree_mumu_parent_source.py` verifies Eq. (58) on the canonical D21×E21 map with $\max |\Delta C_{\mu\mu}^{\text{tree}}| = 9.76 \times 10^{-17}$ and exact layer support (μ -channel share

from $N = 2$ equals 1 on every active row, while the $N = 1, 3$ shares remain below 10^{-26} in magnitude). At the live reference point this gives

$$y_2^{\text{raw}} = 1.65981 \times 10^{-4}, \quad \frac{P_2^{\text{kin}}}{M_2^2} = 2.62559 \times 10^{-1}, \quad C_{eH,\star}^{\mu\mu,\text{tree}} = 4.35798 \times 10^{-5},$$

so the current absolute-normalization target may be written as

$$A_\star = \left(y_2^{\text{raw}}(D_\star) \frac{P_2^{\text{kin}}(D_\star, \eta_\star)}{M_2^2(D_\star)} \right)^{-2} Z_{\text{diag},\star}^{-2}.$$

Across active rows the log-variation of $C_{eH}^{\mu\mu,\text{tree}}$ is dominated by the overlap factor rather than by the kinetic-mass coefficient ($\sigma[\log y_2^{\text{raw}}] = 4.07$, $\sigma[\log(P_2^{\text{kin}}/M_2^2)] = 1.97 \times 10^{-1}$), with corresponding correlations 0.9998 and 0.9100 to $\log C_{eH}^{\mu\mu,\text{tree}}$. This suggests a sharper parent-side task ordering: first derive the reference overlap block $y_2^{\text{raw}}(D_\star)$ from the EYMH/projected-Yukawa side, then lift the smaller residual coefficient P_2^{kin}/M_2^2 and the already mild universal diagonal dressing. The next source audit identifies this overlap block more precisely. In the canonical overlap extractor the flavor kernels differ only through the chirality-width rescaling $\sigma_{\text{scale}}^{(\ell)}$, and for the muon one has $\sigma_{\text{scale}}^{(\mu)} = 1$ exactly. Therefore the raw overlap and the mu-flavor overlap coincide pointwise:

$$y_2^{\text{raw}}(D) = y_{\mu,2}^{\text{flavor}}(D) = (g_{\mu 2}^{\text{UV}}(D))^2.$$

More explicitly, if $u_k(D)$ are the tracked low-lying modes and $K_\mu(D)$ denotes the mu-flavor left-right overlap kernel, then

$$y_2^{\text{raw}}(D) = \sum_k w_{2k}(D) |\langle u_k(D), K_\mu(D) \rangle|,$$

with

$$K_\mu(D) = f_L^{(\mu)}(D) f_R^{(\mu)}(D) W_{\text{frame}}(D),$$

and where $w_{2k}(D)$ are the normalized Gaussian weights of the local microcanonical window around the tracked second branch. The audit `code/audit_h11_y2raw_parent_source.py` verifies both identities on the extracted overlap profile to machine precision ($\max |y_2^{\text{raw}} - y_{\mu,2}^{\text{flavor}}| = 0$, $\max |(g_{\mu 2}^{\text{UV}})^2 - y_2^{\text{raw}}| = 9.87 \times 10^{-17}$). The same audit also shows that collapsing this block to the center-mode overlap is not adequate: on the audited D4–20 detail grid, the center-only candidate $y_{\text{center},2}$ differs from y_2^{raw} by a mean relative error 8.61×10^{-1} and a 95% relative error 1.14. So the remaining normalization target is better understood as a *windowed projected-Yukawa block* rather than as a single-mode overlap. One more parent-side simplification is already available for the kernel itself. In the canonical baseline the overlap extractor keeps

$\sigma_L = \sigma_R \equiv \sigma_\mu = 2.5$ and $W_{\text{frame}} = 1$, so the muon kernel is not merely approximately but exactly a symmetric midplane bridge:

$$K_\mu(\rho, z; D) = A_\mu^{(\text{disc})}(D) \exp\left[-\frac{\rho^2 + z^2}{\sigma_\mu^2}\right].$$

The audit `code/audit_hll_kmu_parent_candidate.py` verifies this pointwise factorization on the audited D4–20 overlap grid with relative sup residual 3.05×10^{-15} . Replacing the discrete normalization by the corresponding finite-box continuum normalization gives a closed-form parent-side candidate

$$A_\mu^{(\text{box})}(D) = \frac{\exp[-D^2/(4\sigma_\mu^2)]}{I_\rho(\sigma_\mu, \rho_{\text{max}}) I_z(\sigma_\mu, D, z_{\text{margin}})},$$

where

$$I_\rho = \pi\sigma_\mu^2 \left(1 - e^{-\rho_{\text{max}}^2/\sigma_\mu^2}\right), \quad I_z = \frac{\sqrt{\pi}\sigma_\mu}{2} \left[\text{erf}\left(\frac{D + z_{\text{margin}}}{\sigma_\mu}\right) + \text{erf}\left(\frac{z_{\text{margin}}}{\sigma_\mu}\right) \right].$$

This finite-box bridge already matches the canonical discrete kernel at the 10^{-4} level (max relative sup error 9.45×10^{-5}). By contrast, the naive infinite-volume amplitude

$$A_\mu^{(\infty)}(D) = \pi^{-3/2} \sigma_\mu^{-3} \exp[-D^2/(4\sigma_\mu^2)]$$

still misses the canonical kernel by about 2.37×10^{-1} . At the live reference point $(D_\star, \eta_\star) = (9.6, 1.0)$ the exact discrete and finite-box amplitudes are

$$A_\mu^{(\text{disc})}(D_\star) = 3.77586 \times 10^{-4}, \quad A_\mu^{(\text{box})}(D_\star) = 3.77621 \times 10^{-4},$$

so the remaining parent-side task is no longer the entire kernel shape. It is better read as the origin of the effective width σ_μ (or its first curvature-controlled correction) inside this midplane bridge block. That effective width now admits a cleaner three-level split. First, because the canonical kernel already closes exactly as

$$\log K_\mu(\rho, z; D) = c_0(D) - \frac{\rho^2 + z^2}{\sigma_\mu^2},$$

a direct log-Hessian fit on the audited D4–20 overlap grid recovers $\sigma_\mu = 2.5$ to machine precision ($\max |\sigma_\mu^{(\text{Hess})}/\sigma_\mu - 1| = 2.49 \times 10^{-15}$, max fit residual 1.85×10^{-13}). Second, a box-normalized moment inversion gives an independent finite-volume check: solving

$$\langle \rho^2 + z^2 \rangle_{K_\mu} = M_2^{(\text{box})}(\sigma_\mu; \rho_{\text{max}}, z_{\text{max}})$$

returns $\sigma_\mu^{(\text{box})} = 2.4999$ with $\max |\sigma_\mu^{(\text{box})}/\sigma_\mu - 1| = 4.09 \times 10^{-5}$. Third, one step closer to the EYMH side, write the action-derived shifted potential near either core as

$$U(\rho, \delta z; D) = U_0(D) + \frac{1}{2} \kappa_{\rho, \mu}(D) \rho^2 + \frac{1}{2} \kappa_{z, \mu}(D) \delta z^2 + \cdots, \quad \delta z = z \mp D/2,$$

and define the isotropized core curvature

$$\bar{\kappa}_\mu(D) = \frac{2\kappa_{\rho,\mu}(D) + \kappa_{z,\mu}(D)}{3}.$$

In a harmonic-ground-state reading one expects a Gaussian width law $\sigma_\mu \propto \bar{\kappa}_\mu^{-1/4}$. Calibrating the proportionality constant at the live reference point therefore gives the current best parent-side candidate

$$\sigma_\mu^{(\text{curv})}(D) = c_\sigma \bar{\kappa}_\mu(D)^{-1/4}, \quad c_\sigma = 22.3027.$$

The audit `code/audit_hll_sigma_mu_parent_candidate.py` shows that the core Hessian is already almost isotropic on the tested D4–20 window ($\kappa_{\rho,\mu}/\kappa_{z,\mu} \in [0.999996, 0.99999997]$), the resulting curvature-controlled width stays within 1.01×10^{-3} of the canonical $\sigma_\mu = 2.5$, and the induced finite-box kernel deformation remains small (max relative sup error 1.11×10^{-2}). So the next parent-side target is now even narrower: not the entire kernel width profile, but the single calibration constant c_σ or its first curvature correction. That remaining gap also now has a clean first-correction audit. Define the bridge constant

$$c_\sigma(D) = \sigma_\mu \bar{\kappa}_\mu(D)^{1/4},$$

which on the audited D4–20 grid is already almost D -invariant: the audit `code/audit_hll_sigma_mu_curvature_correction.py` finds

$$c_\sigma^{\text{ref}} = 22.3027, \quad \frac{\max c_\sigma - \min c_\sigma}{\langle c_\sigma \rangle} = 1.38 \times 10^{-3}.$$

To probe the first non-harmonic correction, expand the same core potential one order further,

$$U(\rho, \delta z; D) = U_0(D) + \frac{1}{2}\kappa_{\rho,\mu}(D)\rho^2 + \frac{1}{2}\kappa_{z,\mu}(D)\delta z^2 + \lambda_{\rho,\mu}(D)\rho^4 + \lambda_{z,\mu}(D)\delta z^4 + \dots,$$

and introduce the isotropized quartic invariant

$$\lambda_\mu(D) = \frac{2\lambda_{\rho,\mu}(D) + \lambda_{z,\mu}(D)}{3}, \quad q_{4,\mu}(D) = \frac{\lambda_\mu(D)}{\bar{\kappa}_\mu(D)^{3/2}}.$$

The current best corrected EYMH-side candidate is then

$$\sigma_\mu^{(\text{curv},4)}(D) = c_{\sigma,*} \bar{\kappa}_\mu(D)^{-1/4} \left[1 + b_4 (q_{4,\mu}(D) - q_{4,\mu}(D_\star)) \right], \quad b_4 = 0.4244.$$

Numerically this single quartic correction collapses the remaining width drift from 1.01×10^{-3} to 8.47×10^{-7} , and pushes the finite-box kernel defect down from 1.11×10^{-2} to 9.47×10^{-5} , essentially saturating the residual box-normalization error. In that sense the current bridge problem is now very narrow: the surviving EYMH-side ambiguity is no longer the width profile,

but the source of the nearly constant bridge coefficient c_σ and, at the next order, the quartic-response coefficient b_4 .

Proposition U1 (UV-to-EFT witness). On the audited map, the scan-level chain

$$C_{eH}^{\text{tree}} \longrightarrow C_{eH}^{\text{match}} \longrightarrow C_{eH}^{\text{IR}}$$

admits an explicit operator-basis reconstruction with numerically controlled deformation relative to the UV-tree baseline. *Verified numerics.* The reproducible audit `scan_hll_uv_to_eft_matching.py` (canonical outputs `paper/hll_uv_to_eft_map.csv` and `paper/hll_uv_to_eft_summary.csv`) shows that in the baseline UV settings (`uv_blend` = 0, `uv_m2_power` = 1, $\kappa_{\text{diag}} = \kappa_{\text{offdiag}} = 0$) the finite-match step is numerically null at map level ($\max |\Delta C_{\mu\mu}^{\text{match}}| = 10^{-30}$), while LL-RG drift remains bounded: $\langle |\Delta \mu_{\mu\mu}| \rangle \simeq 6.19 \times 10^{-4}$ and $\max |\Delta \mu_{\mu\mu}| \simeq 8.93 \times 10^{-3}$; $f(\chi_{\mu\mu}^2 < 4)$ is unchanged between UV-tree and UV+LL-RG. To make that witness explicit at the operator level, we also export a layer-resolved basis audit with `scan_hll_uv_operator_basis_audit.py`; the corresponding map and summary are stored under `paper/hll_uv_operator_basis_*`. There the UV-tree closure is reconstructed as

$$C_{eH}^{\text{tree}} = \sum_{N=1}^3 \left(\frac{P_N^{\text{kin}}}{M_N^2} \right) \mathcal{B}_N^{(eH)}, \quad \mathcal{B}_N^{(eH)} \equiv \mathbf{g}_N^{(eH)} \mathbf{g}_N^{(eH)T}, \quad (59)$$

where $\mathbf{g}_N^{(eH)}$ is the flavor-layer coupling vector entering Eq. (52); this notation avoids collision with the micro-degeneracy scalar g_N and the visibility factor B_N used elsewhere. With explicit blockwise witnesses for ΔC^{match} and ΔC^{RGE} , the tree/match/IR reconstruction residuals are numerically zero at map level, while the reconstructed $\mu_{\mu\mu}$ agrees with the native UV+LL-RG observable to $\max |\Delta \mu_{\mu\mu}^{\text{recon}}| = 7.31 \times 10^{-8}$. This does not yet constitute a full EYM $\rightarrow$ SMEFT derivation, but it upgrades the current closure from a black-box Wilson-map audit to an explicit basis-level EFT witness. As a next structured step beyond fixed $(\kappa_{\text{diag}}, \kappa_{\text{offdiag}})$, we test a finite-match comparator whose effective shifts are built from local UV-basis invariants (shell spread, coefficient variation of P_N^{kin}/M_N^2 , and a normalized off-diagonal mixing norm). We treat this as a structured comparator, not as a replacement for the current constant finite-match baseline. On the canonical D21×E21 grid, the basis witness remains exact ($\max |\Delta \mu_{\mu\mu}^{\text{recon}}| = 5.55 \times 10^{-8}$, reconstruction residuals zero), and the refreshed comparator stays in the small-deformation regime relative to the constant baseline ($\langle |\Delta \mu_{\mu\mu}| \rangle = 5.06 \times 10^{-4}$, $p95 = 1.03 \times 10^{-3}$, $\max |\Delta \mu_{\mu\mu}| = 7.87 \times 10^{-3}$) with zero acceptance mismatch. The off-diagonal component is numerically inactive in the current UV basis, while the diagonal-threshold scale simply controls the strength of a small diagonal threshold deformation. To quantify this interpretation, we run a Phase-2 diagonal-threshold window audit. On the refreshed D21×E21 scan, both gates

require zero acceptance mismatch and select

$$\begin{aligned}\kappa_{\text{diag}}^{(\text{th})} &\in [0.25, 1.0] \quad (\text{conservative comparator window}), \\ \kappa_{\text{diag}}^{(\text{th})} &\in [0.25, 1.5] \quad (\text{extended stable window}).\end{aligned}$$

The canonical witness choice $\kappa_{\text{diag}}^{(\text{th})} = 1.0$ therefore sits at the upper edge of the conservative window: it keeps the diagonal threshold feed-through explicit while remaining in the small-deformation regime. This sharpens the current interpretation of $\kappa_{\text{diag}}^{(\text{th})}$ as a diagonal threshold susceptibility. The remaining missing ingredient is not off-diagonal closure, which is numerically absent in the present basis, but a parent-action normalization for this susceptibility. We now push this normalization one step closer to the parent action by introducing an action-normalized finite-match comparator. The basis witness remains exact ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.54 \times 10^{-8}$, zero reconstruction residuals), and the comparator keeps zero acceptance mismatch with only small deformation relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 1.00 \times 10^{-3}$, $p95 = 1.86 \times 10^{-3}$, $\max |\Delta\mu_{\mu\mu}| = 1.59 \times 10^{-2}$). We now probe the remaining gap directly with an absolute parent-action comparator. In this mode the external diagonal-scale choice is removed and replaced by an absolute witness built from parent-action-side invariants and coefficient alignment:

$$\kappa_{\text{abs}}^{(d)} \equiv S_{\text{kin}}^{(P)} A_{\text{coeff}}, \quad A_{\text{coeff}} \equiv \frac{\|c_N\|_2}{\|c_N\|_1},$$

so that the effective diagonal threshold is generated locally rather than chosen externally. On the canonical D21×E21 grid, this absolute parent-action comparator remains exact at the basis level ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.55 \times 10^{-8}$, zero reconstruction residuals), and the map again keeps zero acceptance mismatch relative to the refreshed constant baseline with only small deformation ($\langle |\Delta\mu_{\mu\mu}| \rangle = 7.17 \times 10^{-4}$, $p95 = 1.17 \times 10^{-3}$, $\max |\Delta\mu_{\mu\mu}| = 1.14 \times 10^{-2}$). To test whether this remaining prefactor can be pushed closer to a genuine curved-background loop origin, we next introduced a heat-kernel comparator family. The direct lesson from the raw and flat-space-subtracted local witnesses is negative: a purely local a_2 -type diagonal density remains dominated by the well region and does not align with the absolute parent-action normalization. The useful signal appears only after promoting the loop witness to a *contrast-based* non-local quantity that compares well and barrier curvature structure. On the canonical D21×E21 grid this contrast-based loop comparator remains exact at the basis level, keeps zero acceptance mismatch, and has only small deformation relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 4.88 \times 10^{-4}$, $p95 = 8.76 \times 10^{-4}$, $\max |\Delta\mu_{\mu\mu}| = 7.67 \times 10^{-3}$). We then add one more loop-improved comparator layer by modulating the parent-action absolute witness with a contrast-based loop prefactor,

$$\kappa_{\text{diag}}^{\text{eff}} = \kappa_{\text{abs}}^{(d)} \Pi_{\text{hk}}^{(d)} s_{\text{sh}} (1 + c_{\text{cv}}) a_{\text{norm}}^{(d)}, \quad (60)$$

$$\kappa_{\text{offdiag}}^{\text{eff}} = \kappa_{\text{abs}}^{(o)} \Pi_{\text{hk}}^{(o)} s_{\text{sh}} m_{\text{off}} a_{\text{norm}}^{(o)}. \quad (61)$$

Here $(\kappa_{\text{abs}}^{(d/o)}, \Pi_{\text{hk}}^{(d/o)}, s_{\text{sh}}, c_{\text{cv}}, m_{\text{off}}, a_{\text{norm}}^{(d/o)})$ denote the absolute parent-action witness, heat-kernel loop prefactor, shell-spread witness, coefficient variation, off-diagonal mixing norm, and normalized action weight. This loop-improved absolute comparator remains exact at the basis level ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.56 \times 10^{-8}$, zero reconstruction residuals) and keeps zero acceptance mismatch, with only small deformation relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 3.37 \times 10^{-4}$, $p95 = 6.92 \times 10^{-4}$, $\max |\Delta\mu_{\mu\mu}| = 5.33 \times 10^{-3}$). We now add one further EYMH-side screening layer, which keeps the same contrast-based loop family but modulates it by local mass-access and curvature-screen factors together with shell-dispersion screening,

$$A_X = \sqrt{\frac{X_{\text{well}}}{1 + X_{\text{well}}}}, \quad A_R = (1 + R_{\text{bar}})^{-1/2}, \quad (62)$$

$$\Pi_{\text{EYMH}} = \Pi_{\text{hk}} A_X A_R \sqrt{\frac{s_{\text{sh}}}{1 + s_{\text{sh}}}} \sqrt{c_{\text{align}}} (1 + g_{\text{cv}} + c_{\text{tree}})^{-1/2}. \quad (63)$$

Here $(X_{\text{well}}, R_{\text{bar}}, s_{\text{sh}}, c_{\text{align}}, g_{\text{cv}}, c_{\text{tree}})$ denote the well-mass, barrier-curvature, shell-spread, coefficient-alignment, gap-variation, and tree-diagonal variation witnesses. This EYMH-screened loop comparator remains exact at the basis level ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.56 \times 10^{-8}$, zero reconstruction residuals) and keeps zero acceptance mismatch, with only small deformation relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 4.17 \times 10^{-5}$, $p95 = 7.63 \times 10^{-5}$, $\max |\Delta\mu_{\mu\mu}| = 6.46 \times 10^{-4}$). We therefore read it as the current best *EYMH-side absolute loop-prefactor comparator*. We then promote the two dominant source factors already isolated by the EYMH-side audits into a refreshed source-informed EYMH comparator. This mode again remains exact at the basis level, keeps zero acceptance mismatch, and stays in the same small-deformation class relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 5.2 \times 10^{-5}$, $p95 = 1.09 \times 10^{-4}$, $\max |\Delta\mu_{\mu\mu}| = 7.81 \times 10^{-4}$). The promoted source-informed prefactor remains nonzero across the whole D21×E21 map, so this is the current best *source-informed EYMH comparator*.

We now rewrite that same source-informed block in explicit parent-action participation/compressibility language. This parented rewrite remains exact at the basis level ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.56 \times 10^{-8}$, zero reconstruction residuals), keeps zero acceptance mismatch, and on the canonical D21×E21 fix grid is map-identical to the source-informed comparator. Relative to the refreshed constant baseline it therefore inherits the same small-deformation scale ($\langle |\Delta\mu_{\mu\mu}| \rangle = 5.2 \times 10^{-5}$, $p95 = 1.09 \times 10^{-4}$, $\max |\Delta\mu_{\mu\mu}| = 7.81 \times 10^{-4}$). The numerical role of this mode is not to introduce a new baseline, but to show that the same comparator can be written directly in parent-action participation/coherence and shell-background compressibility language without changing the canonical map. Figure 7 summarizes this normalization-side endpoint visually: the parented rewrite is map-identical to the source-informed comparator




FIG. 7. EYMH-side parented finite-match comparator on the canonical D21×E21 fix grid. Top row: map-level deformation against the refreshed constant baseline and the source-informed comparator; the source-informed and parented maps coincide to numerical precision. Bottom row: the corresponding EYMH source prefactor, parented prefactor, and tree-diagonal compressibility witness.

while the remaining deformation relative to the constant finite-match baseline is small and structured.

Proposition U2 (Factorized EYMH prefactor). The canonical EYMH-side normalization block admits an exact factorization into local loop, shell-access, participation, and dispersion-screen sectors. The exported audit under `paper/hll_uv_eymh_prefactor_decomposition*_D21E21.*` verifies at numerical precision that

$$\Pi_{\text{EYMH}} = \Pi_{\text{loc}} A_{\text{sh}} A_{\text{part}} S_{\text{disp}}, \quad (64)$$

with $\max |\Delta_{\text{recon}}| = 9.71 \times 10^{-17}$ and $\max |\Delta_{\log_{10}}| = 1.48 \times 10^{-15}$. Here $(\Pi_{\text{loc}}, A_{\text{sh}}, A_{\text{part}}, S_{\text{disp}})$ denote the exported local loop, shell-access, participation/alignment, and dispersion-screen factors. The decomposition shows that the strongest positive correlation is carried by the participation block (0.9655), while the leading suppression comes from the dispersion screen (-0.8774); the local loop factor remains secondary (0.4675). The remaining UV-to-EFT gap is therefore no longer the existence or stability of an EYMH-side absolute prefactor, but the parent-action origin of the participation and dispersion-screen factors that dominate it.

To sharpen that statement further, we resolved the two dominant factors one level deeper in a source audit on the same canonical map. The alignment block is not an opaque extra screen:

it is exactly the coefficient-participation access

$$A_{\text{part}} = N_{\text{eff}}^{-1/4}, \quad N_{\text{eff}} = \left(\frac{c_{l1}}{c_{l2}} \right)^2,$$

with source-level reconstruction again closed to machine precision, $\max |\Delta_{\text{source}}| = 9.71 \times 10^{-17}$.

Likewise, the diagonal screening factor splits as

$$S_{\text{disp}} = S_{\text{gap}} S_{\text{tree}},$$

where the shell-gap contribution is numerically weak (corr = 0.0907 with the EYMH prefactor), while the tree-diagonal screen carries the dominant suppression (corr = -0.8226). Thus the remaining parent-action gap is now more specific: why coefficient participation controls the alignment block, and why tree-diagonal dispersion dominates the screening of the canonical prefactor. We can sharpen the second statement one step further by rewriting

$$S_{\text{tree}} = (1 + \chi_{\text{tree}})^{-1/2}, \quad \chi_{\text{tree}} \equiv \frac{c_{\text{tree}}^{\text{cv}}}{1 + g_{\text{cv}}},$$

so that the canonical screening factor is exactly the compressibility witness of a shell-background-normalized diagonal susceptibility. The corresponding audit closes identically ($\max |\Delta_{\text{tree}}| = 0$) and shows that the susceptibility itself is positively correlated with the EYMH prefactor (corr = 0.7409), while the compressibility carries the leading suppression (corr = -0.8226). A complementary tree-diagonal pressure fraction also correlates strongly with the prefactor (corr = 0.8449). Hence the unresolved parent-action question is now even narrower: one must explain not merely why tree-diagonal dispersion enters, but why a shell-normalized diagonal susceptibility/compressibility controls the dominant screening of the canonical EYMH loop prefactor. We can sharpen the source-side interpretation one step further by rewriting the participation block in the language of a two-mode loop trace,

$$N_{\text{trace}} = \frac{1}{p_1^2 + p_2^2}, \quad S_{\text{trace}}^{\text{norm}} = \frac{-p_1 \ln p_1 - p_2 \ln p_2}{\ln 2}, \quad p_{1,2} = \frac{c_{l1,l2}}{c_{l1} + c_{l2}}.$$

The dedicated parent-source audit (summary under `paper/h11_uv_*parent_source_model*`) closes both rewrites to machine precision, with $\max |\Delta_{\text{part}}| = 2.22 \times 10^{-16}$ and $\max |\Delta_{\text{tree}}| = 1.11 \times 10^{-16}$. Numerically, the participation witness is strongly controlled by the effective loop-trace participation number and its entropy, $\text{corr}(N_{\text{trace}}, A_{\text{part}}) = 0.9608$ and $\text{corr}(S_{\text{trace}}^{\text{norm}}, A_{\text{part}}) = -0.9560$, while the shell-background-normalized pressure/compressibility block remains the leading source-side screen, $\text{corr}(P_{\text{tree}}, \Pi_{\text{src}}) = -0.7213$ and $\text{corr}(S_{\text{tree}}, \Pi_{\text{src}}) = 0.7066$. This narrows the outstanding EYMH normalization problem again: the unresolved parent-action physics is no longer a generic alignment/dispersion question, but specifically why the loop trace

concentrates into the observed two-mode participation structure and why a shell-background-normalized tree-diagonal pressure response fixes the surviving source-informed prefactor. We can now close the participation side one step further with the exact audit stored under `paper/hll_uv_*participation_exact_audit*_D21E21_fix.*`. Writing the projected two-mode participation number as

$$N_{\text{eff}} = \frac{1}{p_1^2 + p_2^2}, \quad d = \sqrt{\frac{2}{N_{\text{eff}}} - 1}, \quad A_{\text{part}}^{\text{exact}} = \sqrt{\frac{1-d}{1+d}},$$

we reconstruct the parented participation factor to machine precision, with $\max |\Delta_{\text{part}}^{\text{exact}}| = 9.99 \times 10^{-16}$. The participation block is therefore no longer just strongly correlated with a projected response; it is exactly fixed by the two-mode loop-trace imbalance implied by the parented map. At this stage the remaining EYMH normalization gap has narrowed again: what remains to be explained from the parent action is why the fluctuation operator dynamically selects this exact two-mode participation closure together with the shell-background-normalized tree-diagonal pressure/compressibility response. We can sharpen this parent-action reading once more by eliminating the intermediate two-mode probabilities altogether and working directly with the projected coefficient vector $c_N = P_N^{\text{kin}}/M_N^2$. The coefficient-norm audit stored under `paper/hll_uv_*participation_norm_audit*_D21E21_fix.*` shows that

$$Q_2 = \frac{\|c\|_2^2}{\|c\|_1^2}, \quad N_{\text{eff}}^{\text{norm}} = \frac{\|c\|_1^2}{\|c\|_2^2}, \quad A_{\text{part}}^{\text{norm}} = Q_2^{1/4} = \sqrt{\frac{\|c\|_2}{\|c\|_1}}$$

reconstructs the parented participation factor to machine precision, with $\max |\Delta_{\text{part}}^{\text{norm}}| = 2.22 \times 10^{-16}$. The participation block can therefore be read directly as an exact norm-ratio coherence of the projected parent-action coefficient block. Numerically, this norm coherence remains correlated with the canonical parented prefactor, $\text{corr}(A_{\text{part}}^{\text{norm}}, \Pi_{\text{parent}}) = 0.7129$, while staying essentially orthogonal to the residual map deformation relative to the refreshed constant baseline. The remaining normalization gap is now narrower once again: explain why the parent EYMH fluctuation operator fixes this exact projected coefficient-norm participation closure together with the shell-background-normalized tree-diagonal pressure/compressibility response. The new coefficient-norm tilt audit stored under `paper/hll_uv_*participation_tilt_audit*_D21E21_fix.*` removes even the explicit norm-ratio intermediate and rewrites the same participation block as an exact projected free-energy tilt,

$$\Delta F_{\text{norm}} = \log \frac{\|c\|_1}{\|c\|_2}, \quad A_{\text{part}}^{\text{tilt}} = e^{-\Delta F_{\text{norm}}/2} = \sqrt{\frac{\|c\|_2}{\|c\|_1}}.$$

On the canonical D21×E21 fix grid this tilt coherence reconstructs the parented participation factor with $\max |\Delta_{\text{part}}^{\text{tilt}}| = 2.22 \times 10^{-16}$ and, together with the shell access and tree-diagonal

compressibility blocks, closes the full parented prefactor exactly as

$$\Pi_{\text{parent}} = \Pi_{\text{hk,local}} A_{\text{shell}} A_{\text{part}}^{\text{tilt}} S_{\text{tree}},$$

with $\max |\Delta \Pi_{\text{parent}}| = 9.02 \times 10^{-17}$. In this language the surviving EYMH normalization gap is narrower again: what remains is to explain why the fluctuation operator dynamically fixes this exact projected coefficient-norm tilt coherence together with the shell-background-normalized tree-diagonal pressure/compressibility response. **Proposition U3 (Projected response action)**. The participation tilt and the tree-diagonal compressibility combine additively into a single projected response action. The response-action audit stored under `paper/hll_uv_*response_action_audit*_D21E21_fix.*` defines

$$S_{\text{resp}} = \Delta F_{\text{norm}} + \log(1 + \chi_{\text{tree}}), \quad A_{\text{resp}} = e^{-S_{\text{resp}}/2},$$

so that the full canonical parented prefactor closes exactly as

$$\Pi_{\text{parent}} = \Pi_{\text{hk,local}} A_{\text{shell}} A_{\text{resp}},$$

with $\max |\Delta A_{\text{resp}}| = 2.22 \times 10^{-16}$ and $\max |\Delta \Pi_{\text{parent}}^{\text{resp}}| = 9.71 \times 10^{-17}$. In other words, the coefficient-norm tilt coherence and the shell-background-normalized tree-diagonal compressibility are not merely parallel witnesses; they enter additively in the same projected response action whose exponential fixes the canonical parented prefactor. The remaining EYMH normalization gap is therefore narrower again: explain why the fluctuation operator dynamically selects this exact projected response action rather than only its separate tilt and compressibility components. We can sharpen this once more by identifying the two terms in S_{resp} with the most direct projected kernel objects currently available in the audited projection. The log-det / Schur audit stored under `paper/hll_uv_*logdet_schur_audit*_D21E21_fix.*` defines

$$K_{\text{part}} = \frac{\|c\|_1}{\|c\|_2}, \quad G_{\text{Schur}} = \frac{1 + \text{gap}_{\text{cv}} + c_{\text{tree,diag,cv}}}{1 + \text{gap}_{\text{cv}}} = 1 + \chi_{\text{tree}},$$

so that

$$S_{\text{resp}} = \log \det K_{\text{part}} + \log G_{\text{Schur}}.$$

On the canonical D21×E21 fix grid this log-det / Schur rewriting again reconstructs the full response weight and parented prefactor to machine precision, with $\max |\Delta A_{\text{resp}}^{\text{logdet}}| = 2.22 \times 10^{-16}$ and $\max |\Delta \Pi_{\text{parent}}^{\text{logdet}}| = 9.71 \times 10^{-17}$. The surviving EYMH normalization block can therefore be read as an exact projected log-det participation kernel together with a shell-background-normalized Schur response. The remaining gap is now narrower still: explain why the fluctuation

operator dynamically selects precisely this projected log-det / Schur structure. **Proposition**

U4 (Kernel-selection principle). The remaining UV-side question can be phrased as a genuine kernel-selection statement rather than as bookkeeping. The audit stored under `paper/hll_uv_*kernel_selection*_D21E21_fix.*` probes the minimal deformed projected family

$$K_{\text{sel}} = \begin{pmatrix} e^{\alpha S_{\text{part}}} & \lambda \sqrt{(e^{\alpha S_{\text{part}}} - 1)(e^{\beta S_{\text{schur}}} - 1)} \\ \lambda \sqrt{(e^{\alpha S_{\text{part}}} - 1)(e^{\beta S_{\text{schur}}} - 1)} & e^{\beta S_{\text{schur}}} \end{pmatrix},$$

whose canonical parent point is $(\alpha, \beta, \lambda) = (1, 1, 0)$. On the D21×E21 fix grid the response-weight and prefactor residuals are uniquely minimized there: the first nontrivial runner-up occurs at $(1, 1, -0.1)$ but already opens an RMSE gap of 2.67×10^{-5} in the parent prefactor. A finite-difference stationarity test at the canonical point yields nearly vanishing gradients, $\partial_{\alpha} J = -3.00 \times 10^{-9}$, $\partial_{\beta} J = -6.34 \times 10^{-9}$, $\partial_{\lambda} J = 0$, together with non-negative Hessian eigenvalues (2.99×10^{-9} , 1.33×10^{-3} , 7.59×10^{-2}). Within this minimal projected kernel family, the fluctuation operator therefore selects unit log-det / Schur weights and suppresses explicit participation–tree cross-coupling. We can sharpen that statement once more by replacing the finite-difference scan with an exact local stationarity principle. The stationarity audit stored under `paper/hll_uv_*stationarity_audit*_D21E21_fix.*` defines the projected mismatch functional

$$J(\alpha, \beta, \lambda) = \left\langle (A(\alpha, \beta, \lambda) - A_{\text{ref}})^2 \right\rangle.$$

At the canonical point $(\alpha, \beta, \lambda) = (1, 1, 0)$ all first-variation moments vanish exactly: $\partial_{\alpha} J = \partial_{\beta} J = \partial_{\lambda} J = 0$. The (α, β) sector is then controlled by an exact quadratic stationarity matrix with positive eigenvalues (1.33×10^{-3} , 7.59×10^{-2}), while the cross-coupling direction is even under $\lambda \rightarrow -\lambda$ and therefore begins only at quartic order,

$$J(1, 1, \lambda) = C_4 \lambda^4 + \mathcal{O}(\lambda^6), \quad C_4 = 1.4974 \times 10^{-3}.$$

For the concrete probe $\lambda = 0.1$, this quartic law predicts an RMSE of 3.8696×10^{-4} , matching the directly evaluated 3.8734×10^{-4} . The projected EYMH selection principle can therefore be stated more sharply than before: the canonical log-det / Schur kernel is not only the best nearby fit, it is an exact stationary point with quadratic stability in the log-det / Schur weights and quartic suppression of explicit participation–tree mixing. We can package the same statement into a local projected effective-action gap. The variational-selection audit stored under `paper/hll_uv_*variational_selection*_D21E21_fix.*` defines

$$\Delta \Gamma_{\text{sel}}(\delta\alpha, \delta\beta, \lambda) = \frac{1}{2} \delta\theta^T H \delta\theta + C_4 \lambda^4, \quad \delta\theta = (\delta\alpha, \delta\beta),$$

using the exact quadratic stationarity matrix H from the (α, β) block and the exact quartic coefficient C_4 from the mixing direction. On a local grid around the canonical point, this variational gap and the exact mismatch functional have the same minimum at $(1, 1, 0)$, and the agreement is already very strong: $\text{corr}(J_{\text{exact}}, \Delta\Gamma_{\text{sel}}) = 0.9894$, the pure- λ slice is reproduced to within 1.88×10^{-8} in absolute gap, and even the $\lambda = 0$ (α, β) plane remains controlled at p95 absolute gap 2.97×10^{-5} . This is the cleanest projected effective-action statement so far: near the canonical point, the fluctuation operator selects the log-det / Schur kernel by minimizing a local effective action whose quadratic log-det/Schur sector and quartic mixing sector reproduce the observed selection landscape. We can push the same statement one step closer to a parent-action kernel form. The audit stored under `paper/hll_uv_*parent_kernel_statement*_D21E21_fix.*` rewrites the deformed projected kernel family as

$$K_{11} = e^{\alpha S_{\text{part}}}, \quad K_{22} = e^{\beta S_{\text{schur}}}, \quad K_{12} = \lambda \sqrt{(K_{11} - 1)(K_{22} - 1)},$$

and then expresses the deformed response weight exactly as

$$A(\alpha, \beta, \lambda) = A_{\text{ref}} \exp \left[-\frac{1}{2} \Delta S_{\text{kernel}} \right],$$

with

$$\Delta S_{\text{kernel}} = (\alpha - 1)S_{\text{part}} + (\beta - 1)S_{\text{schur}} + \log(1 - \lambda^2 \xi_{\text{cross}}).$$

The corresponding parent-kernel objective reproduces the direct mismatch functional to machine precision on the full local scan, with $\max |J_{\text{direct}} - J_{\text{parent}}| = 1.30 \times 10^{-18}$ and $\max |A_{\text{direct}} - A_{\text{parent}}| = 3.33 \times 10^{-16}$. The same exact functional is minimized again at the canonical point $(1, 1, 0)$. This is the strongest mother-action statement so far: the canonical log-det / Schur kernel is not only stationary and variationally preferred in a local surrogate sense, but is selected by an exact projected parent-kernel excess functional in which log-det and Schur deformations enter linearly while explicit participation-tree mixing appears only through an even determinant factor. We can now make the block structure behind that statement explicit. The audit stored under `paper/hll_uv_*block_split*_D21E21_fix.*` writes the canonical projected fluctuation kernel in terms of a participation block

$$K_{\text{part}} = \frac{\|c\|_1}{\|c\|_2},$$

a shell-background/tree block

$$G_{\text{Schur}} = \frac{1 + \text{gap}_{\text{cv}} + c_{\text{tree,diag,cv}}}{1 + \text{gap}_{\text{cv}}},$$

and a mixed scale

$$C_{\text{mix}} = \sqrt{(K_{\text{part}} - 1)(G_{\text{Schur}} - 1)}.$$

On the canonical map the response action closes exactly as $S_{\text{part}} + S_{\text{Schur}}$, with $\text{corr}(S_{\text{split}}, S_{\text{resp}}) = 1.0$ and $\max |S_{\text{split}} - S_{\text{resp}}| = 7.77 \times 10^{-16}$. For the deformed family

$$K_{\text{sel}} = \begin{pmatrix} K_{\text{part}} & \lambda C_{\text{mix}} \\ \lambda C_{\text{mix}} & G_{\text{Schur}} \end{pmatrix},$$

the determinant factorization

$$\det K_{\text{sel}} = K_{\text{part}} G_{\text{Schur}} (1 - \lambda^2 \xi_{\text{cross}})$$

also closes to machine precision on the full scan, with $\max |\Delta \det| = 4.44 \times 10^{-16}$ and $\max |\Delta S| = 3.33 \times 10^{-16}$. This is the clearest structural statement so far: the canonical projected fluctuation operator is block-diagonal in the participation and shell-background/tree sectors, and explicit participation–tree mixing survives only as an even determinant-level penalty. We can sharpen the same statement one step further by embedding it into a background-normalized parent block determinant. The audit stored under `paper/h11_uv_*parent_blockdet*_D21E21_fix.*` defines

$$K_{11} = e^{\alpha S_{\text{part}}}, \quad K_{\text{bg}} = 1 + \text{gap}_{\text{cv}}, \quad K_{22} = K_{\text{bg}} e^{\beta S_{\text{Schur}}},$$

and it is convenient to abbreviate the deformed normalized tree/background block as

$$G_{\beta} := \frac{K_{22}}{K_{\text{bg}}} = e^{\beta S_{\text{Schur}}}, \quad G_{\text{Schur}} = G_{\beta}|_{\beta=1}.$$

with parent mixing scale

$$C_{\text{parent}} = \sqrt{(K_{11} - 1)(K_{22} - K_{\text{bg}})}, \quad \mathcal{K}_{\text{parent}} = \begin{pmatrix} K_{11} & \lambda C_{\text{parent}} \\ \lambda C_{\text{parent}} & K_{22} \end{pmatrix}.$$

The normalized determinant and Schur factor then satisfy the exact identities

$$\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} = K_{11} G_{\beta} (1 - \lambda^2 \xi_{\text{cross}}), \quad \widehat{G}_{\text{Schur}} = \frac{K_{22} - K_{12}^2 / K_{11}}{K_{\text{bg}}} = G_{\beta} (1 - \lambda^2 \xi_{\text{cross}}),$$

so the response action is literally the block-determinant / Schur-complement factorization

$$S_{\text{parent}} = \log K_{11} + \log \widehat{G}_{\text{Schur}}.$$

On the D21×E21 fix grid this closes to machine precision, with $\max |S_{\text{parent}} - S_{\text{resp}}| = 5.55 \times 10^{-16}$, $\max |\Delta(\det / K_{\text{bg}})| = 1.33 \times 10^{-15}$, $\max |\Delta \widehat{G}_{\text{Schur}}| = 6.66 \times 10^{-16}$, and

$\text{corr}(J_{\text{direct}}, J_{\text{blockdet}}) = 1.0$. This is the strongest derivation statement so far: the canonical response weight is the inverse square root of a background-normalized projected parent block determinant, rather than merely a convenient response-action rewrite. The only remaining structural ambiguity in that parent block is the mixed entry. To test whether the geometric-mean choice is merely convenient or actually selected, the audit stored under `paper/hll_uv_*parent_mix_geomean*_D21E21_fix.*` scans the smallest symmetric-excess family

$$C_{\text{gen}} = \kappa (K_{11} - 1)^u (K_{22} - K_{\text{bg}})^v.$$

On the local D21×E21 fix grid the unique exact point is

$$(u, v, \kappa) = \left(\frac{1}{2}, \frac{1}{2}, 1\right), \quad C_{\text{parent}} = \sqrt{(K_{11} - 1)(K_{22} - K_{\text{bg}})},$$

for which the determinant, Schur, weight, and normalized cross-ratio residuals all collapse to machine precision:

$$\begin{aligned} \max |\Delta(\det / K_{\text{bg}})| &= 4.44 \times 10^{-16}, & \max |\Delta \hat{G}_{\text{Schur}}| &= 4.44 \times 10^{-16}, \\ \max |\Delta A| &= 1.11 \times 10^{-16}, & \max |\Delta \xi| &= 1.11 \times 10^{-16}. \end{aligned}$$

The first nontrivial runner-up, $(u, v, \kappa) = (0.625, 0.625, 1.1)$, already opens visible residuals, with $\max |\Delta(\det / K_{\text{bg}})| = 1.32 \times 10^{-3}$ and $\max |\Delta A| = 2.18 \times 10^{-4}$. This is the strongest naturality statement so far: once the projected parent kernel is required to couple the participation and tree/background sectors through a minimal symmetric excess family, the geometric-mean mixed block is uniquely selected. We can push that one step further by testing the nearest non-minimal extension: a ratio-dependent warp of the geometric mean. The audit stored under `paper/hll_uv_*parent_ratio_warp*_D21E21_fix.*` probes

$$C_{\text{warp}} = \kappa C_{\text{parent}} \exp[\delta L + \nu L^2], \quad L = \frac{1}{2} \log(E_{\text{part}}/E_{\text{tree}}),$$

with $E_{\text{part}} = K_{11} - 1$ and $E_{\text{tree}} = K_{22} - K_{\text{bg}}$. On the same D21×E21 fix grid the unique exact point is again the unwarped kernel, $(\kappa, \delta, \nu) = (1, 0, 0)$, for which the determinant, Schur, weight, and normalized cross-ratio residuals all vanish exactly. The first nontrivial runner-up, $(1, 0, -0.05)$, already opens visible residuals, with $\max |\Delta(\det / K_{\text{bg}})| = 7.18 \times 10^{-4}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 5.39 \times 10^{-4}$, $\max |\Delta A| = 7.14 \times 10^{-5}$, and $\max |\Delta \xi| = 6.08 \times 10^{-3}$. This is the strongest minimality statement so far: the low-mode parent block not only selects the geometric-mean mixed sector inside the minimal symmetric-excess family, it also rejects the first ratio-warped extension of that family. We can then reparameterize the same local family in the

coordinates that most directly match the remaining proof obligation, namely an overall normalization shift m , a symmetric homogeneity degree s , and an antisymmetric participation/tree tilt a :

$$C_{\text{gen}} = \exp(m) \exp\left[\frac{s}{2} \log(E_{\text{part}} E_{\text{tree}}) + \frac{a}{2} \log(E_{\text{part}}/E_{\text{tree}})\right].$$

The audit stored under `paper/h11\_uv\_parent\_symnorm\_D21E21\_fix.*` shows that on the same D21×E21 fix grid the unique exact point is

$$(m, s, a) = (0, 1, 0),$$

for which the determinant, Schur, weight, and normalized cross-ratio residuals again collapse to machine precision. The first nontrivial runner-up, $(0.05, 1.125, 0)$, already opens visible residuals, with $\max |\Delta(\det/K_{\text{bg}})| = 6.61 \times 10^{-4}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 4.72 \times 10^{-4}$, $\max |\Delta A| = 1.09 \times 10^{-4}$, and $\max |\Delta \xi| = 7.73 \times 10^{-3}$. This is the cleanest structural statement so far: once the parent block is written in symmetry/normalization coordinates, the low-mode projected parent action selects the canonical family by enforcing zero normalization shift, unit symmetric degree, and zero antisymmetric tilt. The last nearby structural ambiguity is whether the parent action truly selects a locally affine generator in these excess coordinates, or whether small log-curvature terms survive. The audit stored under `paper/h11\_uv\_parent\_generator\_affinity\_D21E21\_fix.*` probes the first non-affine local extension,

$$\log C_{\text{gen}} = \frac{1}{2} L_{\text{sum}} + q_{ss} L_{\text{sum}}^2 + q_{dd} L_{\text{diff}}^2 + q_{sd} L_{\text{sum}} L_{\text{diff}},$$

with $L_{\text{sum}} = \log(E_{\text{part}} E_{\text{tree}})$ and $L_{\text{diff}} = \log(E_{\text{part}}/E_{\text{tree}})$. On the same D21×E21 fix grid the unique exact point is

$$(q_{ss}, q_{dd}, q_{sd}) = (0, 0, 0),$$

for which the determinant, Schur, weight, and normalized cross-ratio residuals again collapse to machine precision. The first nontrivial runner-up, $(0, -0.0125, 0.0125)$, already opens visible residuals, with $\max |\Delta(\det/K_{\text{bg}})| = 2.76 \times 10^{-4}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 1.98 \times 10^{-4}$, $\max |\Delta A| = 2.62 \times 10^{-5}$, and $\max |\Delta \xi| = 2.23 \times 10^{-3}$. This is the narrowest structural statement so far: once symmetry and normalization have been fixed, the low-mode projected parent action also rejects the first local log-curvature corrections, so the surviving canonical parent block lies in the locally affine multiplicative excess class itself. The remaining coordinate-level ambiguity is whether the projected parent action genuinely uses the canonical excess variables themselves, or whether another local subtraction point could play the same role. The audit stored under

`paper/hll\_uv\_*excess\_coordinate*_D21E21\_fix.*` probes the minimal reference-offset family

$$E_{\text{part}}^{(r)} = K_{11} - r_{\text{part}}, \quad E_{\text{tree}}^{(r)} = K_{22} - r_{\text{tree}} K_{\text{bg}}.$$

On the same D21×E21 fix grid the unique exact point is

$$(r_{\text{part}}, r_{\text{tree}}) = (1, 1),$$

for which the determinant, Schur, weight, normalized cross-ratio, and strict anchor-leakage residuals all vanish identically. The first nontrivial runner-up, (1, 1.05), already opens visible errors, with $\max |\Delta(\det / K_{\text{bg}})| = 1.08 \times 10^{-3}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 7.02 \times 10^{-4}$, $\max |\Delta A| = 2.48 \times 10^{-4}$, and $\max |\Delta \xi| = 1.61 \times 10^{-2}$. More importantly, it distorts the first nonzero-response slices even though the strict anchor leakage still vanishes, with onset mismatches 4.53×10^{-2} in the participation direction and 9.78×10^{-2} in the tree/background direction. This is the sharpest fixed-point statement so far: the projected parent action naturally organizes the mixed block as a deviation from the identity participation block and the shell/background tree block themselves, rather than from an arbitrary nearby subtraction point. We then tightened the same statement by allowing the excess coordinates to vary inside the smallest smooth family that preserves both the fixed points and the tangent normalization there,

$$E_{\text{part}}^{(p)} = \text{BC}_p(K_{11}), \quad E_{\text{tree}}^{(q)} = K_{\text{bg}} \text{BC}_q(K_{22}/K_{\text{bg}}),$$

where $\text{BC}_p(x) = (x^p - 1)/p$ for $p \neq 0$ and $\text{BC}_0(x) = \log x$, so every member obeys $\text{BC}_p(1) = 0$ and $\text{BC}'_p(1) = 1$. On the same D21×E21 fix grid the unique exact point is again the linear additive excess choice

$$(p_{\text{part}}, p_{\text{tree}}) = (1, 1),$$

for which the determinant, Schur, weight, cross-ratio, anchor, and onset residuals all vanish to machine precision. The first nontrivial runner-up, (0.75, 1.0), already opens visible first-slice distortions, with onset mismatches 2.02×10^{-3} and 8.23×10^{-3} together with $\max |\Delta(\det / K_{\text{bg}})| = 8.25 \times 10^{-4}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 5.51 \times 10^{-4}$, $\max |\Delta A| = 9.80 \times 10^{-5}$, and $\max |\Delta \xi| = 7.85 \times 10^{-3}$. This is the narrowest coordinate statement so far: even after fixing the same low-mode fixed points and the same tangent normalization, the projected parent action still uniquely selects the linear excess variables themselves. We can sharpen that result into positive normal-coordinate language by probing the first nonlinear local jet family that preserves the same fixed points and the same unit tangent normalization,

$$E_{\text{part}}^{(\zeta_p)} = E_{\text{part}} + \zeta_p E_{\text{part}}^2, \quad E_{\text{tree}}^{(\zeta_t)} = E_{\text{tree}} + \zeta_t E_{\text{tree}}^2 / K_{\text{bg}}.$$

On the same D21×E21 fix grid the unique exact point is

$$(\zeta_p, \zeta_t) = (0, 0),$$

for which the determinant, Schur, weight, cross-ratio, anchor, and onset residuals again vanish identically. The first nontrivial runner-up, $(0.125, 0)$, already opens visible distortions, with onset mismatches 2.10×10^{-3} and 9.63×10^{-3} together with $\max |\Delta(\det / K_{\text{bg}})| = 9.90 \times 10^{-4}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 6.59 \times 10^{-4}$, $\max |\Delta A| = 1.18 \times 10^{-4}$, and $\max |\Delta \xi| = 9.47 \times 10^{-3}$. This is the strongest positive coordinate statement so far: the projected parent action naturally uses the zero-second-jet normal coordinates around the identity/background fixed points, and those normal coordinates are exactly the linear excess variables.

D. Projected Parent-Kernel Derivation

The preceding exact audits leave only one conceptual task: explain, directly from the projected EYMH parent action, why the localized low-mode response closes into the particular background-normalized 2×2 kernel isolated numerically above. In other words, the remaining gap is no longer a family-search problem but a derivation problem. The three lemmas below separate that derivation into its natural pieces: reduction of the parent Hessian to a two-mode block, uniqueness of the mixed entry inside the anchored positive family, and uniqueness of the linear excess variables as normal coordinates near the identity/background fixed point.

To keep the notation tight, we reserve K_{part} and G_{Schur} for the canonical undeformed split blocks already identified in the earlier exact audits, G_β for the one-parameter deformed normalized tree/background block of the parent-block determinant family, and K_{11} , K_{22} , and $\hat{G}_{\text{Schur}}$ for the entries and Schur complement of the generic projected parent kernel. At the canonical point these objects coincide through

$$K_{11} = K_{\text{part}} = e^{S_{\text{part}}}, \quad K_{22} = K_{\text{bg}} G_{\text{Schur}} = K_{\text{bg}} e^{S_{\text{Schur}}}, \quad \hat{G}_{\text{Schur}} = G_{\text{Schur}}.$$

a. Assumptions and scope. The theorem chain below is formulated on a localized stationary background Ψ_0 of the projected EYMH parent action, under four explicit assumptions: first, the low-mode response admits a dominant two-dimensional projected truncation; second, the selected projected sector is the one whose diagonal responses reproduce the canonical participation block and the shell/tree block already isolated by the earlier block-split closure; third, that projected sector is locally stable so that the induced quadratic form is positive there; and fourth, local coordinate statements are restricted to the audited coordinate families in a neighborhood of the identity/background anchor. The exact audits are therefore used as supporting evidence

for the canonical identification of the projected sectors and for the closure of the resulting formulas, but not as hidden additional premises inside the theorem statements themselves.

In this subsection, “dominant two-dimensional projected truncation” is meant in that same localized and audited sense: on the parented map, the determinant / Schur / weight closures already isolate two projected sectors whose diagonal responses reproduce the participation block and the shell/tree block to numerical exactness, and the theorem chain is only claiming that the parent action closes on that same low-mode sector. We are therefore not asserting a global completeness theorem for the full fluctuation spectrum; the statement is local to the same low-mode neighborhood on which the exact bridge audits close.

Lemma 1 (Projected 2×2 Hessian reduction). Let Ψ_0 be a localized stationary background of the projected EYMH parent action, and let $\mathcal{H}_{\text{EYMH}}$ denote its quadratic fluctuation Hessian. Assume that the localized low-mode response admits a dominant two-mode truncation by Gram-normalizable directions u_{part} and u_{tree} , selected so that their diagonal projected responses reconstruct the participation block and the shell/tree block isolated earlier, and assume that the quadratic form induced by $\mathcal{H}_{\text{EYMH}}$ is locally positive on that projected sector. After restriction to that sector, and after Gram normalization of the projected subspace, the quadratic fluctuation form can be written as

$$\delta^2 S_{\text{EYMH}}|_{\text{proj}} = \frac{1}{2} \begin{pmatrix} a_{\text{part}} & a_{\text{tree}} \end{pmatrix} \mathcal{K}_{\text{parent}} \begin{pmatrix} a_{\text{part}} \\ a_{\text{tree}} \end{pmatrix},$$

with positive background-normalized kernel

$$\mathcal{K}_{\text{parent}} = \begin{pmatrix} K_{11} & K_{12} \\ K_{12} & K_{22} \end{pmatrix}, \quad K_{\text{bg}} = 1 + \text{gap}_{\text{cv}},$$

and with response weight

$$A_{\text{resp}} \propto \left(\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} \right)^{-1/2} = \left(K_{11} \hat{G}_{\text{Schur}} \right)^{-1/2}, \quad \hat{G}_{\text{Schur}} = \frac{K_{22} - K_{12}^2/K_{11}}{K_{\text{bg}}}.$$

Proof sketch. Writing the collective fields as $\Psi = \Psi_0 + \delta\Psi$, the parent action expands as

$$S[\Psi_0 + \delta\Psi] = S[\Psi_0] + \delta S[\Psi_0; \delta\Psi] + \frac{1}{2} \langle \delta\Psi, \mathcal{H}_{\text{EYMH}} \delta\Psi \rangle + \mathcal{O}(\delta\Psi^3).$$

Stationarity of the localized background removes the linear term. Restricting to

$$\delta\Psi = a_{\text{part}} u_{\text{part}} + a_{\text{tree}} u_{\text{tree}}$$

and introducing the column matrix $U = (u_{\text{part}}, u_{\text{tree}})$ together with the amplitude vector $a = (a_{\text{part}}, a_{\text{tree}})^T$, define the orthogonal projector onto the localized two-mode sector by

$$P := UU^\dagger, \quad V_2 := \text{im } P,$$

so that after Gram normalization of the projected basis one has

$$U^\dagger U = \mathbf{1}_2, \quad P^\dagger = P, \quad P^2 = P.$$

On that projected sector, $\delta\Psi = P\delta\Psi$, and the projected Hessian operator is

$$H_P := P\mathcal{H}_{\text{EYMH}}P|_{V_2}.$$

Its matrix representation in the basis $(u_{\text{part}}, u_{\text{tree}})$ is

$$\mathcal{K}_{\text{parent}} := U^\dagger \mathcal{H}_{\text{EYMH}} U.$$

Thus $\mathcal{K}_{\text{parent}}$ is the projected Hessian block itself, not an inverse propagator inserted by hand. After Gram normalization this reduces to the matrix of projected Hessian entries

$$H_{ij} = \langle u_i, \mathcal{H}_{\text{EYMH}} u_j \rangle, \quad i, j \in \{\text{part}, \text{tree}\},$$

and therefore to a positive 2×2 projected kernel on the localized low-mode sector. Here positivity comes from the projected local-stability assumption, not from stationarity alone. At the canonical point the diagonal entries reduce to the canonical split blocks,

$$K_{11} = K_{\text{part}} = e^{S_{\text{part}}}, \quad \frac{K_{22}}{K_{\text{bg}}} = G_{\text{Schur}} = e^{S_{\text{schur}}}.$$

The content of this identification is that the two-mode truncation is not an arbitrary basis choice but the proposition-level sector-selection hypothesis of the EYMH bridge: the earlier block-split and parent-kernel closures already separate one projected direction whose diagonal response reconstructs the participation block and a second projected direction whose background-normalized diagonal response reconstructs the shell/tree block. In that sense u_{part} and u_{tree} are selected by the previously closed projected sectors rather than inserted ad hoc, and Lemma 1 promotes that exact numerical split into a projected-Hessian statement about the parent action itself. The projected Gaussian fluctuation integral is therefore

$$Z_{\text{proj}} \propto \int_{\mathbb{R}^2} d^2 a \exp \left[-\frac{1}{2} a^T \mathcal{K}_{\text{parent}} a \right] \propto (\det \mathcal{K}_{\text{parent}})^{-1/2}.$$

Dividing out the shell-background reference determinant $Z_{\text{bg}} \propto K_{\text{bg}}^{-1/2}$ produces the normalized response factor

$$A_{\text{resp}} \propto \frac{Z_{\text{proj}}}{Z_{\text{bg}}} \propto \left(\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} \right)^{-1/2},$$

and the Schur identity

$$\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} = K_{11} \frac{K_{22} - K_{12}^2/K_{11}}{K_{\text{bg}}} = K_{11} \hat{G}_{\text{Schur}}$$

immediately yields

$$S_{\text{parent}} = \log K_{11} + \log \widehat{G}_{\text{Schur}}.$$

Audit closure. The exact audits verify each algebraic step appearing in the theorem statement and support the canonical sector assignment used above: the earlier block-split closure identifies the participation and shell/tree diagonals themselves as the canonical projected sectors; the ‘logdet + Schur’ witness reconstructs the response weight and parented prefactor to machine precision, with $\max |\Delta A_{\text{resp}}^{\log \det}| = 2.22 \times 10^{-16}$ and $\max |\Delta \Pi_{\text{parent}}^{\log \det}| = 9.71 \times 10^{-17}$; the background-normalized parent block-determinant / Schur audit gives $\max |S_{\text{parent}} - S_{\text{resp}}| = 5.55 \times 10^{-16}$, $\max |\Delta(\det / K_{\text{bg}})| = 1.33 \times 10^{-15}$, and $\max |\Delta \widehat{G}_{\text{Schur}}| = 6.66 \times 10^{-16}$; and the exact parent-kernel excess functional reproduces the direct mismatch functional with $\max |J_{\text{direct}} - J_{\text{parent}}| = 1.30 \times 10^{-18}$. Lemma 1 therefore reduces the remaining work to making the projected-Hessian construction from the parent action explicit enough that this already-exact numerical structure can be read as a genuine Gaussian fluctuation theorem on the localized two-mode sector.

Local truncation remark. What is still *not* claimed here is a global spectral-completeness theorem for the full EYMH fluctuation operator. The current statement is local to the audited low-mode sector and assumes that the complement does not reorganize the projected closure at leading order. A stronger theorem would require an explicit control of the complement remainder, for example through a Feshbach/Schur estimate on the off-sector couplings H_{PQ}, H_{QP} relative to the complement block H_{QQ} , or an equivalent local spectral-separation bound. At present we use the exact bridge residuals as the empirical evidence that this remainder is negligible on the localized parented map, while leaving the fully quantified truncation bound as part of the remaining proof-level tightening.

Corollary 1 (Local complement remainder form). Let $Q := I - P$, and write the quadratic fluctuation operator in P/Q block form

$$\mathcal{H}_{\text{EYMH}} = \begin{pmatrix} H_{PP} & H_{PQ} \\ H_{QP} & H_{QQ} \end{pmatrix}, \quad H_{PP} = P \mathcal{H}_{\text{EYMH}} P|_{V_2}.$$

Assume in the same localized regime that H_{QQ} is invertible on the complement and that the off-sector coupling is small enough that

$$A := H_{PP}^{-1/2} H_{PQ} H_{QQ}^{-1} H_{QP} H_{PP}^{-1/2}$$

satisfies $\|A\| < 1$. Then integrating out the complement produces the effective low-mode operator

$$H_{\text{eff}} = H_{PP} - H_{PQ} H_{QQ}^{-1} H_{QP},$$

and its determinant differs from the projected determinant by

$$\det H_{\text{eff}} = \det H_{PP} \det(I - A), \quad |\log \det H_{\text{eff}} - \log \det H_{PP}| \leq \frac{2\|A\|}{1 - \|A\|}.$$

Proof sketch. The Schur/Feshbach reduction on the complement gives the exact identity

$$H_{\text{eff}} = H_{PP}^{1/2}(I - A)H_{PP}^{1/2},$$

so

$$\det H_{\text{eff}} = \det H_{PP} \det(I - A).$$

Because H_{PP} is 2×2 , the eigenvalues α_1, α_2 of A obey $|\alpha_i| \leq \|A\| < 1$, hence

$$\log \det(I - A) = \log(1 - \alpha_1) + \log(1 - \alpha_2),$$

and the elementary bound $|\log(1 - x)| \leq |x|/(1 - |x|)$ for $|x| < 1$ yields

$$|\log \det(I - A)| \leq \frac{|\alpha_1|}{1 - |\alpha_1|} + \frac{|\alpha_2|}{1 - |\alpha_2|} \leq \frac{2\|A\|}{1 - \|A\|}.$$

If, in addition, H_{QQ} is positive on the complement, then A is positive semidefinite and $\|A\| < 1$ implies $I - A > 0$; hence H_{eff} is positive on the projected low-mode sector as well. Thus the remaining truncation error is controlled by the product of the off-sector couplings with the complement resolvent, which is exactly the quantity that the current theorem chain leaves as an audited local remainder rather than a global spectral statement.

Once that projected block structure is fixed, all diagonal ambiguity is gone and only one matrix-level choice survives: the off-diagonal coupling between the participation sector and the shell/background sector. Lemma 2 addresses exactly that remaining freedom. Its role is narrower than Lemma 1, but also sharper: after the block decomposition has been fixed at the parent-action level, the mixed entry must still be chosen, and the claim is that only one positive homogeneous anchored coupling remains compatible with the same local symmetry and normalization structure.

Lemma 2 (Geometric-mean mixed entry). With the diagonal sectors fixed as in Lemma 1, define the anchored excess variables

$$E_{\text{part}} = K_{11} - 1, \quad E_{\text{tree}} = K_{22} - K_{\text{bg}}.$$

On the local positive cone $E_{\text{part}} \geq 0$, $E_{\text{tree}} \geq 0$, with the positive square-root branch fixed by continuity from the canonical point, consider positive mixed couplings that vanish at either fixed point, remain symmetric after shell-background normalization, and are locally affine in

the logarithms of the excess variables. Within that class, the projected parent block uniquely selects

$$C_{\text{parent}} = \sqrt{E_{\text{part}} E_{\text{tree}}} = \sqrt{(K_{11} - 1)(K_{22} - K_{\text{bg}})}.$$

Proof sketch. Once the diagonal blocks are fixed, the mixed entry is the only remaining structural freedom. The generator-class hypothesis is not inserted ad hoc at this point: on the interior of the local positive cone, any positive mixed entry may be written locally as $C = \exp g(\log E_{\text{part}}, \log E_{\text{tree}})$, and the resulting expression is extended to the anchored boundary by continuity from the canonical point. Local multiplicative reweightings of the two anchored excess amplitudes,

$$(E_{\text{part}}, E_{\text{tree}}) \mapsto (\rho E_{\text{part}}, \sigma E_{\text{tree}}), \quad (\rho, \sigma) > 0,$$

act by translations in logarithmic variables,

$$(\log E_{\text{part}}, \log E_{\text{tree}}) \mapsto (\log E_{\text{part}} + \log \rho, \log E_{\text{tree}} + \log \sigma),$$

or equivalently

$$(\Sigma, \Delta) \mapsto \left(\Sigma + \frac{1}{2}(\log \rho + \log \sigma), \Delta + \frac{1}{2}(\log \rho - \log \sigma) \right).$$

Writing $z = (\Sigma, \Delta)$ and $\tau = (\tau_\Sigma, \tau_\Delta)$ for a small local translation in these logarithmic coordinates, Taylor's theorem gives

$$g(z + \tau) = g(z) + Dg(z)\tau + \frac{1}{2}\tau^T [D^2g(z + \theta\tau)]\tau, \quad 0 < \theta < 1.$$

Thus the first obstruction to pure 1-jet translation stability is exactly the quadratic jet D^2g . Within the audited local family, those quadratic corrections are parametrized by the generator-curvature coefficients (q_{ss}, q_{dd}, q_{sd}) , so the curvature audit is testing the first non-affine term in the local logarithmic normal form rather than an unrelated ad hoc deformation. Consequently one may write

$$g(\Sigma, \Delta) = m + s\Sigma + a\Delta + R_2(\Sigma, \Delta),$$

with the local remainder bounded by

$$|R_2(\Sigma, \Delta)| \leq \frac{1}{2} \sup_{\zeta \in \mathcal{N}} \|D^2g(\zeta)\| \|\Sigma, \Delta\|^2$$

on a sufficiently small neighborhood $\mathcal{N}$ of the canonical point. The natural residual class is therefore the one whose generator is stable under these local translations to first nontrivial order, namely the locally affine class already isolated by the earlier generator-affinity closure.

Within that class, the smallest positive multiplicative excess family compatible with the anchor constraints is

$$C_{\text{gen}} = \kappa E_{\text{part}}^u E_{\text{tree}}^v.$$

Equivalently, in logarithmic excess variables

$$\Sigma = \frac{1}{2} \log(E_{\text{part}} E_{\text{tree}}), \quad \Delta = \frac{1}{2} \log(E_{\text{part}}/E_{\text{tree}}),$$

one may write the most general nearby positive homogeneous coupling as

$$\log C_{\text{gen}} = m + s \Sigma + a \Delta + \mathcal{O}(\Sigma^2, \Sigma \Delta, \Delta^2),$$

where m is an overall normalization shift, s is the symmetric homogeneity degree, and a is the antisymmetric tilt. Homogeneity of degree one in the excess amplitudes forces $u + v = 1$, equivalently $s = 1$ in the logarithmic generator, while symmetry under interchange of the two anchored excess sectors forces $u = v$, equivalently $a = 0$. The remaining freedom is the overall normalization $\kappa = e^m$, which is fixed by the canonical anchor normalization. The exact audit under `paper/hll_uv_*parent_mix_geomean*_D21E21_fix.*` closes uniquely at

$$(u, v, \kappa) = \left(\frac{1}{2}, \frac{1}{2}, 1\right),$$

which is precisely the geometric mean. In particular, the closure forces the symmetry condition $u = v$ rather than assuming it a priori. The nearby deformations then show that this is structural, not accidental: the first ratio-warped extension is rejected at $(\kappa, \delta, \nu) = (1, 0, 0)$, the symmetry/normalization audit closes only at $(m, s, a) = (0, 1, 0)$, and the first local generator-curvature corrections are rejected at $(q_{ss}, q_{dd}, q_{sd}) = (0, 0, 0)$. The combined effect is to set the logarithmic generator exactly to

$$\log C_{\text{parent}} = \Sigma = \frac{1}{2} \log(E_{\text{part}} E_{\text{tree}}),$$

and hence

$$C_{\text{parent}} = \sqrt{E_{\text{part}} E_{\text{tree}}}.$$

Thus, once the diagonal sectors, anchors, and local multiplicative structure are fixed, only one positive homogeneous mixed coupling survives. At this stage we deliberately do *not* use the explicit determinant law of the corollary as an external input, because in the current audit chain that determinant identity is only established after the geometric-mean mixed entry has already been inserted into the parent block. The theorem therefore remains non-circular by taking the generator-class closure as the primary uniqueness statement and postponing the explicit determinant formula to the corollary. *Audit closure.* The first nontrivial runner-ups already open visible residuals at every stage: the geomean runner-up gives

$\max |\Delta(\det / K_{\text{bg}})| = 1.32 \times 10^{-3}$ and $\max |\Delta A| = 2.18 \times 10^{-4}$; the first ratio warp gives $\max |\Delta(\det / K_{\text{bg}})| = 7.18 \times 10^{-4}$ and $\max |\Delta \xi| = 6.08 \times 10^{-3}$; the first symmetry/normalization deformation gives $\max |\Delta(\det / K_{\text{bg}})| = 6.61 \times 10^{-4}$ and $\max |\Delta A| = 1.09 \times 10^{-4}$; and the first local generator-curvature deformation gives $\max |\Delta(\det / K_{\text{bg}})| = 2.76 \times 10^{-4}$ and $\max |\Delta \xi| = 2.23 \times 10^{-3}$. In particular, the generator-affinity audit shows that once the diagonal sectors and anchors are fixed, the projected parent action does not preserve any first local log-curvature correction to the mixed entry. After Lemma 1 has fixed the diagonal sectors, Lemma 2 therefore leaves a single symmetry-respecting, anchor-vanishing mixed coupling: the geometric mean.

Consistency remark. The link back to Lemma 1 is local but direct. In logarithmic coordinates the affine generator is exactly the part of g that survives local translations without producing new curvature data, whereas the quadratic jet is the first non-affine obstruction. After reweighting, that obstruction generates not only pure $\tau_\Sigma^2, \tau_\Sigma \tau_\Delta, \tau_\Delta^2$ terms but also mixed coordinate-dependent corrections once the translated jet is re-expanded about the original canonical point. On the projected parent block these corrections first enter through distortions of the mixed entry K_{12} , and therefore through the Schur response

$$\hat{G}_{\text{Schur}} = \frac{K_{22} - K_{12}^2 / K_{11}}{K_{\text{bg}}}.$$

This is why the generator-curvature audit is the numerically natural check of the affine-generator claim: it is probing the first local deformation that can alter the Schur term while leaving the diagonal sectors of Lemma 1 fixed.

Non-circular bridge. The generic off-diagonal follow-up audits support this conclusion without replacing the theorem proof. A direct determinant-inversion argument is currently circular because the exact parent-block determinant audit already hard-codes the geometric mean, so the meaningful independent question is whether generic off-diagonal witnesses can reproduce the even λ -curvature implied by the canonical mixed response. On the canonical $D21 \times E21$ fix map, the best global quartic proxy is a small blend of curvature-screened and action-absolute off-diagonal witnesses, but blocked holdout shows that those blend coefficients drift substantially across D -blocks. By contrast, the low- D regime is stably captured by the single normalized witness `hk_abs_offdiag/diag`, while the dense $D41 \times E21$ replay shows that the remaining complexity lives primarily in the tree-sector regime structure rather than in the participation envelope. We therefore use the generic off-diagonal program as supporting evidence that the geometric mean is the least arbitrary surviving global coupling, while keeping the theorem itself anchored to the non-circular generator-family closure stated above.

At that point the matrix entries are fixed, but the local variables in which those entries are expressed are still logically separate data. Lemma 3 closes that last freedom by showing that

the linear excess variables isolated numerically are exactly the local normal-form coordinates of the projected parent block.

Lemma 3 (Linear excess variables as normal coordinates). Near the fixed point

$$(K_{11}, K_{22}) = (1, K_{\text{bg}}),$$

the natural local coordinates of the projected parent block are exactly

$$E_{\text{part}} = K_{11} - 1, \quad E_{\text{tree}} = K_{22} - K_{\text{bg}},$$

within the audited local coordinate families, in the sense that they are simultaneously the unique anchor-preserving coordinates, the unique unit-tangent-normalized coordinates, and the unique zero-second-jet coordinates compatible with exact determinant / Schur / weight closure.

Proof sketch. Let $E = (E_{\text{part}}, E_{\text{tree}})$ and let $\tilde{E} = \Phi(E)$ be any local reparameterization of the projected parent block near the fixed point $E = 0$. In this language a normal-form coordinate system is characterized by

$$\Phi(0) = 0, \quad D\Phi(0) = I, \quad D^2\Phi(0) = 0,$$

namely exact anchoring at the fixed point, unit tangent normalization there, and vanishing quadratic jet. Within the audited axis-preserving local coordinate families, these are also the first three independent departures from the identity chart, because any such local chart admits an expansion of the form

$$\tilde{E} = b + LE + Q(E, E) + \mathcal{O}(\|E\|^3),$$

with b testing anchor shifts, L testing tangent normalization, and Q testing the first nonlinear jet. The three exact coordinate audits therefore probe the normal-form conditions in the natural order in which local coordinate freedom appears. The anchor family

$$E_{\text{part}}^{(r)} = K_{11} - r_{\text{part}}, \quad E_{\text{tree}}^{(r)} = K_{22} - r_{\text{tree}}K_{\text{bg}}$$

already closes uniquely at $(r_{\text{part}}, r_{\text{tree}}) = (1, 1)$, so $\Phi(0) = 0$ is forced at the identity participation block and the shell/background tree block. Among coordinate families that preserve those fixed points, the fixed-point Box-Cox audit then closes uniquely at

$$(p_{\text{part}}, p_{\text{tree}}) = (1, 1),$$

which is precisely the linear additive excess choice and therefore fixes $D\Phi(0) = I$. Finally, the nonlinear-jet family

$$E_{\text{part}}^{(\zeta_p)} = E_{\text{part}} + \zeta_p E_{\text{part}}^2, \quad E_{\text{tree}}^{(\zeta_t)} = E_{\text{tree}} + \zeta_t E_{\text{tree}}^2 / K_{\text{bg}}$$

closes uniquely at

$$(\zeta_p, \zeta_t) = (0, 0),$$

so the same coordinates are also zero-second-jet normal coordinates, i.e. $D^2\Phi(0) = 0$. The three exact closures therefore read together as a local normal-form statement: moving the anchor, changing the unit tangent normalization, or turning on a nonzero second jet each destroys exact determinant / Schur / weight closure. Equivalently, within the audited coordinate families, the identity map in the excess variables is the unique local chart that agrees with the projected parent block to quadratic order. Lemma 3 is thus the coordinate counterpart of Lemmas 1 and 2: once the block structure and mixed entry are fixed, the local variables are fixed as well. *Audit closure.* The first nontrivial deviation from each of $\Phi(0) = 0$, $D\Phi(0) = I$, and $D^2\Phi(0) = 0$ already opens visible residuals. In the anchor audit the first runner-up produces onset residuals 4.53×10^{-2} and 9.78×10^{-2} together with $\max|\Delta(\det/K_{\text{bg}})| = 1.08 \times 10^{-3}$ and $\max|\Delta\xi| = 1.61 \times 10^{-2}$. The first tangent-normalized Box-Cox runner-up produces onset residuals 2.02×10^{-3} and 8.23×10^{-3} with $\max|\Delta(\det/K_{\text{bg}})| = 8.25 \times 10^{-4}$ and $\max|\Delta\xi| = 7.85 \times 10^{-3}$. The first nonlinear-jet runner-up produces onset residuals 2.10×10^{-3} and 9.63×10^{-3} with $\max|\Delta(\det/K_{\text{bg}})| = 9.90 \times 10^{-4}$ and $\max|\Delta\xi| = 9.47 \times 10^{-3}$. Hence the natural local variables are exactly the zero-second-jet linear excess coordinates.

Together, Lemmas 1–3 successively eliminate the three remaining ambiguities of the parent-kernel bridge: first the projected block structure, then the mixed coupling, and finally the local coordinate chart. At that point there is no remaining freedom in the local 2×2 parent block, because the diagonal sectors are fixed to $(1 + E_{\text{part}}, K_{\text{bg}} + E_{\text{tree}})$, the off-diagonal sector is fixed to $\sqrt{E_{\text{part}}E_{\text{tree}}}$, and the variables themselves are fixed to the linear zero-second-jet excess chart. The corollary is therefore not an informal restatement of the exact audits, but the compact parent-action statement obtained by substituting those three outputs back into the projected kernel of Lemma 1.

Corollary (Projected EYMH parent kernel). Combining Lemmas 1–3, and still working on the same local positive cone and within the same audited local coordinate families, the localized low-mode EYMH fluctuation operator projects to a background-normalized 2×2 parent kernel whose entries are now fixed term by term:

$$K_{11} = 1 + E_{\text{part}}, \quad K_{22} = K_{\text{bg}} + E_{\text{tree}}, \quad K_{12} = \sqrt{E_{\text{part}}E_{\text{tree}}}.$$

Substituting these three identities into the generic projected block of Lemma 1 gives the explicit

canonical parent kernel

$$\mathcal{K}_{\text{parent}} = \begin{pmatrix} 1 + E_{\text{part}} & \sqrt{E_{\text{part}} E_{\text{tree}}} \\ \sqrt{E_{\text{part}} E_{\text{tree}}} & K_{\text{bg}} + E_{\text{tree}} \end{pmatrix}.$$

Its normalized determinant and response weight are therefore

$$\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} = \frac{(1 + E_{\text{part}})(K_{\text{bg}} + E_{\text{tree}}) - E_{\text{part}} E_{\text{tree}}}{K_{\text{bg}}} = 1 + E_{\text{part}} + \frac{E_{\text{tree}}}{K_{\text{bg}}},$$

$$A_{\text{resp}} \propto \left(\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} \right)^{-1/2} = \left(1 + E_{\text{part}} + \frac{E_{\text{tree}}}{K_{\text{bg}}} \right)^{-1/2}.$$

Equivalently, the same determinant may be read in Schur form as

$$\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} = K_{11} \hat{G}_{\text{Schur}}, \quad \hat{G}_{\text{Schur}} = \frac{K_{22} - K_{12}^2 / K_{11}}{K_{\text{bg}}} = 1 + \frac{E_{\text{tree}}}{K_{\text{bg}}(1 + E_{\text{part}})}.$$

This corollary is the compact parent-action statement that remains after the exact audit closures: the parent block is fixed by its two canonical diagonal sectors, the unique geometric-mean mixed coupling, and the unique linear zero-second-jet excess coordinates.

b. Reader-facing interpretation. A denser diagnostic replay of the canonical parented target on a $D41 \times E21$ grid sharpens the qualitative reading of the same kernel. On that dense grid the target

$$\xi_{\text{target}} = \frac{E_{\text{part}} E_{\text{tree}}}{K_{11} K_{22}}$$

is nearly η -independent at fixed D , so the dominant structure is genuinely D -driven. The two excess sectors then play visibly different roles. The participation excess E_{part} is a slow envelope: it stays broad, varies monotonically over long D -ranges, and reaches its largest values only in the late- D tail. By contrast, the tree/Schur excess E_{tree} carries the multi-peak structure: its strongest crest sits near $D \approx 6$, it exhibits the alternating low- and mid- D shoulders that generate the non-monotone target profile, and it is also the first sector to decline before the final high- D collapse. The target maximum at $D \approx 9.2$ is therefore not the peak of either sector separately, but the point where the smooth participation envelope and the oscillatory tree sector overlap most efficiently under the kernel normalization.

This decomposition is useful for interpretation because it explains why a single global phenomenological ansatz for $\xi_{\text{target}}(D)$ is only partially successful. The dense-grid diagnostics support a more structured reading:

$$\begin{aligned} E_{\text{part}}(D) &\approx \text{slow envelope,} \\ E_{\text{tree}}(D) &\approx \text{low/mid-}D \text{ multi-frequency oscillation} \\ &\quad \times \text{high-}D \text{ cutoff,} \end{aligned}$$

with ξ_{target} reconstructed from their normalized product rather than from a single primitive oscillatory law. In particular, a smooth-envelope model already describes E_{part} well, while E_{tree} prefers either a two-frequency global fit or, more cleanly, a piecewise regime fit with a low/mid- D oscillatory branch and a high- D cutoff branch. We therefore regard the phenomenological regime decomposition as a reader-facing diagnostic of the corollary, not as a replacement for the parent-kernel statement itself.

$$(D_0, \eta_0)_{\chi^2\text{-best}} \equiv \arg \min_{(D, \eta) \in \mathcal{G}} \chi^2(D, \eta), \quad (65)$$

$$(D_0, \eta_0)_{\text{robust}} \equiv \arg \max_{(D, \eta) \in \mathcal{A}} d_{\partial\mathcal{A}}(D, \eta), \quad \mathcal{A} \equiv \{(D, \eta) \in \mathcal{G} \mid \chi^2 \leq 4\}, \quad (66)$$

where $\mathcal{G}$ is the scan grid and $d_{\partial\mathcal{A}}$ is the Euclidean distance to the acceptance-boundary complement on that grid. We use $\mathcal{R}_3$ (higher first) and then χ^2 (lower first) as tie-breakers for Eq. (66). A full reference-normalization sensitivity audit is reported in Appendix E6: over the tested normalization grid, $f(\chi_{\mu\mu}^2 < 4)$ spans $[0.0333, 0.2411]$, while the baseline fixed reference gives 0.1500, the dynamic χ^2 -best reference gives 0.1333, and the robust-center reference gives 0.1381. To reduce reference-normalization arbitrariness in reporting, we additionally track Wilson-diagonal ratios that are invariant under the same re-normalization:

$$\mathcal{R}_{e/\mu}^{(\text{inv})}(D, \eta) \equiv \left| \frac{C_{eH}^{ee, \text{IR}}(D, \eta)}{C_{eH}^{\mu\mu, \text{IR}}(D, \eta)} \right|^2, \quad \mathcal{R}_{\tau/\mu}^{(\text{inv})}(D, \eta) \equiv \left| \frac{C_{eH}^{\tau\tau, \text{IR}}(D, \eta)}{C_{eH}^{\mu\mu, \text{IR}}(D, \eta)} \right|^2. \quad (67)$$

Using `code/scan_anchor_invariant_ratios.py` on `paper/hll_uv_to_eft_map.csv`, we obtain (all-grid) $p10/p50/p90$ bands $\mathcal{R}_{e/\mu}^{(\text{inv})} = (1.60 \times 10^{-2}, 1.26 \times 10^{-1}, 8.10 \times 10^{-1})$ and $\mathcal{R}_{\tau/\mu}^{(\text{inv})} = (2.59 \times 10^{-3}, 1.01 \times 10^{-2}, 3.84 \times 10^{-2})$; restricted to the illustrative acceptance band $\chi_{\mu\mu}^2 < 4$, they become $\mathcal{R}_{e/\mu}^{(\text{inv})} = (1.67 \times 10^{-2}, 6.39 \times 10^{-2}, 3.54 \times 10^{-1})$ and $\mathcal{R}_{\tau/\mu}^{(\text{inv})} = (5.23 \times 10^{-3}, 9.72 \times 10^{-3}, 5.17 \times 10^{-2})$. These diagnostics are anchor-independent and are reported as complementary constraints to Eq. (56).

$$\chi^2(D, \eta) = \frac{(\mu_{\mu\mu}^{\text{pred}} - \mu_{\mu\mu}^{\text{obs}})^2}{\sigma_{\mu\mu}^2}, \quad (68)$$

where $\sigma_{\mu\mu} = 0.4$. For one degree of freedom, we also report the local p -value

$$p_{\text{local}}(D, \eta) = 1 - F_{\chi_1^2}(\chi^2(D, \eta)) = \text{erfc}\left(\sqrt{\chi^2(D, \eta)/2}\right). \quad (69)$$

Figure 8 shows the **EFT/Wilson-matched acceptance region** in two equivalent languages,

$$\chi^2 < 4 \quad \Longleftrightarrow \quad p_{\text{local}} > 4.55 \times 10^{-2} \quad (\text{dof} = 1). \quad (70)$$

We use this threshold only as an *illustrative acceptance band* for map-level comparison and do not claim a formal 95% CL exclusion. Because the (D, η) scan is a multiple-comparison procedure

(look-elsewhere) and no global trials correction is applied, this band is interpreted only as a falsifiable target region. In the present full-grid action-derived profile scan, the accepted fraction is

$$f(\chi^2 < 4) = f(p_{\text{local}} > 4.55 \times 10^{-2}) = 0.1500, \quad (71)$$

with best grid point $\chi^2 \simeq 2.21 \times 10^{-5}$ at $(D, \eta) \approx (4.00, 0.264)$. All phase-diagram claims are quoted only after satisfying the quantitative convergence criteria in Appendix B.




FIG. 8. Compatibility with $H \rightarrow \mu\mu$ under the EFT/Wilson-matched UV+LL-RG map. Left: χ^2 map showing a restricted illustrative accepted band ($\chi^2 < 4$, equivalently $p_{\text{local}} > 4.55 \times 10^{-2}$ for dof=1). Right: predicted $\mu_{\mu\mu}^{\text{pred}}$.

E. EFT/Wilson-Matched Extension to $H \rightarrow ee$ and $H \rightarrow \tau\tau$

Using the same map-level EFT/Wilson-matched definition, we additionally evaluate lepton-channel signal strengths by layer assignment $N = 1 \rightarrow ee$, $N = 2 \rightarrow \mu\mu$, and $N = 3 \rightarrow \tau\tau$:

$$\mu_{\ell\ell}^{\text{pred}}(D, \eta) = \left| \frac{C_{eH}^{\ell\ell}(D, \eta)}{C_{eH}^{\ell\ell}(D_0, \eta_0)} \right|^2 \left[\frac{\Gamma_{\text{tot}}(D, \eta)}{\Gamma_{\text{tot}}(D_0, \eta_0)} \right]^{-1}. \quad (72)$$

Figure 9 summarizes the corresponding maps on the same (D, η) grid; tabulated statistics are stored in `paper/hll_signal_strength_summary.csv`. This extension is still a scan-level EFT map (not full UV matching), while the visibility sector in the global scan uses baseline fully normalized EFT operator visibility built from overlap-extracted inputs with mode tracking and microcanonical windowing. The repository keeps the three-channel observation fields in `data/pdg_leptons.json` as a reproducible interface. The corresponding runtime-direct visibility branch is release-qualified only in the risk-weighted profile-blend form; the reported baseline maps remain full-direct and EFT-operator normalized unless stated otherwise.




FIG. 9. PSLT EFT/Wilson-matched extension from $H \rightarrow \mu\mu$ to three lepton channels on the same scan grid: $H \rightarrow ee$ (left), $H \rightarrow \mu\mu$ (middle), and $H \rightarrow \tau\tau$ (right). Each panel shows $\mu_{\ell\ell}^{\text{pred}}$ from Eq. (56) with the same reference-normalization point $(D_0, \eta_0) = (10, 1)$.

VIII. DISCUSSION AND CONCLUSION

We have presented a computable EFT-level demonstrator of the Projection Spectral Layer Theory. By closing the loop between geometric inputs and observable layer probabilities, we have shown that:

1. **Generations are spectral within the baseline closure:** The “three generation” structure is not imposed as an input; it emerges on the current release baseline as a dynamical output of competing entropy (g_N), kinetics (Γ_N), and visibility (B_N).
2. **Stability without ad-hoc cutoffs:** The infinite tower of layers becomes numerically and physically stable once micro-degeneracy is regulated by the `fp_2d_full` microcanonical-window + tail prescription and the observable sector uses EFT-operator-normalized visibility profiles $B_N(D)$ with neutral high- N saturation rather than an extra visibility-side cutoff.
3. **Falsifiability:** The closure yields concrete predictions for the winner phase diagram. With full-grid action-derived $\chi_N^{(LR)}(D)$ and $\tilde{A}_\ell(D)$ profiles, the $H \rightarrow \mu\mu$ illustrative EFT/Wilson-matched acceptance band ($\chi^2 < 4 \Leftrightarrow p_{\text{local}} > 4.55 \times 10^{-2}$, dof= 1) covers about 15.0% of the scan and defines a falsifiable target region in (D, η) .

At the same time, the present implementation keeps two EFT-level approximations explicit: map-level profile closure remains active in the baseline $B_N(D)$ visibility sector (while release-production and risk-weighted runtime-direct parity branches are both available) and a bounded nearest-neighbor flavor projector in the matched $H \rightarrow \mu\mu$ map of Eq. (56). Therefore, the present manuscript supports a robust statement about spectral generation selection on the release baseline, and also about a release-qualified runtime-direct visibility parity path, but

not yet a strict claim that the observable sector is closed by per-cell all-direct extraction at production level.

Limitations and outlook. The present closure is internally consistent but remains an EFT-level demonstrator in several places: The use of mixed-resolution ingredients is a deliberate design choice: action-derived extraction fixes low-mode profiles (`fp_2d_full` for g_N , EFT-operator-normalized $B_N(D)$ from overlap-extracted inputs, $\chi_N^{(LR)}(D)$, and $\tilde{A}_\ell(D)$), while global-map evaluation is still an EFT-level closure rather than a full direct localized recomputation at every (D, η, N) point.

1. **Observable mapping:** the $H \rightarrow \mu\mu$ comparison now uses a scan-level EFT/Wilson-matched map with bounded flavor mixing and leptonic-width closure [Eqs. (50)–(57)]. The fixed point in Eq. (56) should be read only as a reference-normalization convention for the relative observable, while the stronger normalization-side evidence now comes from the absolute and loop-improved EYMH parent-action comparators and the Wilson-diagonal re-normalization invariants. The new absolute-bridge audit shows that this arbitrariness reduces exactly to three scalar bridge constants $(A_\star, R_{e/\mu}^\star, R_{\tau/\mu}^\star)$ and that the accepted-region invariant median already fixes $R_{e/\mu}^\star$ to practical numerical precision on the canonical map. A full parent-side derivation of the remaining absolute amplitude $A_\star$ and tau-to-muon ratio $R_{\tau/\mu}^\star$ is still pending.
2. **Mixing channel:** under the parity-symmetric first-principles definition [Eq. (73)], we obtain $\chi_N^{(\text{sym})} \sim 10^{-19}$ for $D = \{6, 12, 18\}$ (Appendix C), i.e., symmetry-protected cancellation. We therefore redefine the channel in localized basis, extract $\chi_N^{(LR)}$ (Appendix D), and inject its full scan-grid action-derived profile in the global map. The remaining gap is a full (D, η, N) localized projection.
3. **Open-system extension:** a Lindblad-type χ_{eff} mechanism is now implemented as a scan-ready diagnostic mode with profiled $\gamma_\phi(D), \gamma_{\text{mix}}(D)$ (Appendix E4). The original multi-anchor + holdout calibration remains in place, but the line is now much narrower than a generic reservoir ansatz: on the audited knot set, κ_{env} is isolated as a scalar amplitude modulus, the localized two-level projection leaves only the leading σ_z/σ_x directions, and the effective kernels are on-shell samples of a single gap-locked single-pole response family with $\tau_{\text{env},a} = \omega_{1,a}^{-1}$ to machine precision. The sharp remaining theorem target is parent-side: show that the projected EYMH Schur resolvent admits a first-mode-dominant truncation on $\omega \in \{0, \Delta E_a\}$ with leading pole scale $\Omega_{1,a} = \omega_{1,a}$ and uniformly small higher-mode remainder. Until that step is taken, this block remains outside baseline

figures.

4. **Model-chain unification:** $\chi_N^{(LR)}(D)$ and $\tilde{A}_\ell(D)$ are injected from action-derived full-grid profiles in the global (D, η) map, while the release baseline keeps strict D-grid lookup and the automatic chain remains a comparator. The release-production path now pairs active-grid direct two-dimensional spectra for $g_N(D)$ with per-cell direct evaluation of $\chi_{LR}(D)$ and $\tilde{A}_\ell(D)$ while keeping baseline B_N profile closure; on both release audits it matches the release baseline exactly ($\Delta f(\chi_{\mu\mu}^2 < 4) = 0$, mismatch = 0, $\max |\Delta\mu_{\mu\mu}| = 0$). The tuned visibility path should therefore be read as a release-qualified risk-weighted visibility parity check, not as a strict all-direct closure of B_N . Separately, the broader-grid strict all-direct validator is parked at its current case-aware endpoint on the two validation surfaces, so this line should be reopened only for a genuinely new observable-side mechanism class rather than another local blend retune. For the UV-to-EFT map, the refreshed diagonal-threshold comparator and the absolute/loop-improved EYMH comparators show the same pattern: zero acceptance mismatch, exact basis reconstruction, and only small deformation. The remaining UV-to-EFT gap is therefore no longer the existence of a stable local absolute witness, but the derivation of its prefactor from a bona fide parent EYMH loop calculation.
5. **Spectral tower:** the baseline validation is strongest for the low-lying bound sector ($N = 1, 2$ in the physical-gap scan), while the high- N contribution is treated at EFT level through (g_N, B_N) regularization. The audited single-track benchmark does not exhibit a fourth bound layer on the tested D window, and the current one-dimensional kinetic proxy favors threshold loss of binding over a proven exponentially small- Γ_4 mechanism. A mode-resolved no-fourth-layer theorem is therefore not claimed here.
6. **Finite-time dynamics:** t_{coh} is treated as a control parameter in the baseline scan; a geometry-closed dephasing candidate $t_{\text{coh}}^{(\text{deph})}(D) = \pi/\Delta\omega_{12}(D)$ is benchmarked in Appendix E 7 but is not yet propagated into the reported baseline figures.
7. **Overlap amplitude:** η is scanned as an external control in the baseline maps. A first-principles prefactor candidate $\eta_{\text{fp}}(D)$ from splitting-action data is benchmarked in Appendix E 8; map-level impact is mild in profile-scaled mode but this closure is not yet propagated into baseline figures.
8. **Barrier-leakage channel normalization:** the baseline maps now use action-derived profile factors $(A_1, A_2) = (\tilde{A}_1(D), \tilde{A}_2(D))$ from Appendix E 9 (historical “superrad” arti-

fact filenames retained for reproducibility). The remaining gap is to replace profile-level insertion with a fully localized, channel-resolved extraction across the full scan grid.

Status and redefinition of the mixing channel. Using Eq. (33), we fix a basis-independent normalization scale $\bar{\Gamma}_N$ and define the parity-symmetric overlap channel by

$$\chi_N^{(\text{sym})}(D) = \frac{|M_{12}^{(\text{sym})}(D, N)|}{\bar{\Gamma}_N(D)}, \quad (73)$$

$$M_{12}^{(\text{sym})}(D, N) = 2\pi \int_0^{\rho_{\text{max}}} \rho d\rho \int_{-z_{\text{max}}}^{z_{\text{max}}} dz \psi_{N,1}^*(\rho, z) \delta V(\rho, z; D) \psi_{N,2}(\rho, z),$$

and Eq. (38) with the explicit spherical-average δV gives $\chi_N^{(\text{sym})} \approx 0$ (Appendix C). To avoid mixing two inequivalent notions, we redefine the phenomenological mixing channel in a localized-well basis:

$$\psi_{N,L} = \frac{\psi_{N,+} + \psi_{N,-}}{\sqrt{2}}, \quad \psi_{N,R} = \frac{\psi_{N,+} - \psi_{N,-}}{\sqrt{2}}, \quad (74)$$

and define the channel mixing from the corresponding two-state Hamiltonian element:

$$M_{LR}^{(H)}(D, N) \equiv \frac{\lambda_{N,2}(D) - \lambda_{N,1}(D)}{2}, \quad \chi_N^{(LR)}(D) \equiv \frac{|M_{LR}^{(H)}(D, N)|}{\bar{\Gamma}_N(D)}. \quad (75)$$

The baseline mixing entry in Table III is interpreted as the localized-channel coefficient. First-principles localized extraction results are given in Appendix D, and its full scan-grid action-derived profile is used in the global (D, η) scan of Section VII. In terminology, $M_{LR}^{(H)}$ is a *closed-system localized-basis splitting* quantity; it is not itself a decoherence rate. Open-system decoherence is introduced separately through Lindblad rates $(\gamma_\phi, \gamma_{\text{mix}})$ in Appendix E 4.

ACKNOWLEDGMENTS

The author acknowledges the use of PDG 2024 data and standard Python scientific stacks for verification.

Appendix A: Action-Derived Chain Reproducibility SOP

This appendix provides a clone-level SOP for reproducing the action-derived extraction chain and the paper-level Fig/Table artifacts.

1. Single Entry Point

From repository root:

TABLE VIII. Compact SOP stage map for action-derived reproducibility. The one-click wrapper executes these stages in a fixed order, logs each stage, and records exact script/output names in the artifact registry.

Stage group	Deliverables and checks
Production maps and observables	Regenerates phase diagrams, Higgs signal-strength maps, and UV-to-EFT matching summaries.
Release/direct validation	Runs the full-direct release audit, localized-direct validation surfaces, chain-mode parity checks, and artifact-status registry.
First-principles profile diagnostics	Recomputes localized χ , bounded phase-space g_N , open-system diagnostics, and the t_{coh} , η , and channel-normalization candidate profiles.
Packaging and provenance	Collects figures, tables, logs, checksums, manifest metadata, and the per-artifact map under the reproducibility run directory.

```
bash scripts/repro/reproduce_paper.sh
```

Optional modes:

```
bash scripts/repro/reproduce_paper.sh --with-paper
```

```
bash scripts/repro/reproduce_paper.sh --package-only
```

2. Step Map and Deliverables

3. Acceptance Conditions

A run is considered reproducible for this manuscript if:

1. the one-click command exits with return code 0;
2. `repro/latest/manifest.json` reports `missing_required = 0`;
3. `repro/latest/checksums.sha256` and `repro/latest/artifact_index.csv` are generated and non-empty.

Appendix B: Action-Derived Effective Potential and WKB

1. Two-Center Harmonic Conformal Factor

We define the conformal factor by a projected Poisson problem on the static slice,

$$\nabla^2 \Omega(\mathbf{x}) = -4\pi a [\rho_\varepsilon(\mathbf{x} - \mathbf{x}_+) + \rho_\varepsilon(\mathbf{x} - \mathbf{x}_-)], \quad \Omega \rightarrow 1 \quad (|\mathbf{x}| \rightarrow \infty), \quad (\text{B1})$$

with $\mathbf{x}_\pm = \pm(D/2)\hat{z}$. Using the Green representation

$$\Omega(\mathbf{x}) = 1 + a \sum_{s=\pm} \int d^3x' \frac{\rho_\varepsilon(\mathbf{x}' - \mathbf{x}_s)}{|\mathbf{x} - \mathbf{x}'|}, \quad (\text{B2})$$

and the normalized Plummer kernel

$$\rho_\varepsilon(\mathbf{r}) = \frac{3\varepsilon^2}{4\pi(|\mathbf{r}|^2 + \varepsilon^2)^{5/2}}, \quad \int d^3r \rho_\varepsilon(\mathbf{r}) = 1, \quad (\text{B3})$$

the integral is analytic and gives

$$\Omega(\rho, z; D) = 1 + a \left(\frac{1}{\sqrt{\rho^2 + (z - D/2)^2 + \varepsilon^2}} + \frac{1}{\sqrt{\rho^2 + (z + D/2)^2 + \varepsilon^2}} \right) \quad (\text{B4})$$

where a is the geometric strength, ε is the core regularization scale, and D is the dual-center separation. Away from the cores ($|\mathbf{x} - \mathbf{x}_\pm| \gg \varepsilon$), this satisfies $\nabla^2 \Omega \approx 0$. The radial identity

$$\nabla^2 \left(\frac{1}{\sqrt{r^2 + \varepsilon^2}} \right) = -\frac{3\varepsilon^2}{(r^2 + \varepsilon^2)^{5/2}} \quad (\text{B5})$$

gives the smeared Laplacian:

$$\begin{aligned} \nabla^2 \Omega &= -3a\varepsilon^2 \left[(r_+^2 + \varepsilon^2)^{-5/2} + (r_-^2 + \varepsilon^2)^{-5/2} \right] \\ &< 0. \end{aligned} \quad (\text{B6})$$

which is equivalent to the source form in Eq. (6). This negative contribution is essential for forming the attractive wells in V_{eff} . As a reproducibility check, `verify_omega_geometric_origin.py` confirms $\int d^3r \rho_\varepsilon = 1$ with absolute error 3.75×10^{-7} and validates the radial Laplacian identity with max relative error 1.39×10^{-4} in the non-asymptotic band; run outputs are listed in the artifact registry.

2. Veff Derivation from KG Equation

Starting from $(\square_g - m_0^2 - \xi R)\Phi = 0$ with $g_{\mu\nu} = \Omega^2 \eta_{\mu\nu}$:

a. Step 1: $\square_g$ in conformal coordinates.

$$\square_g \Phi = \Omega^{-2}(-\partial_t^2 + \nabla^2)\Phi + 2\Omega^{-3}\nabla\Omega \cdot \nabla\Phi \quad (\text{B7})$$

b. Step 2: Time separation. For $\Phi = e^{-i\omega t}\phi(\mathbf{x})$:

$$\nabla^2 \phi + 2\Omega^{-1}\nabla\Omega \cdot \nabla\phi + [\omega^2 - \Omega^2(m_0^2 + \xi R)]\phi = 0 \quad (\text{B8})$$

c. Step 3: Conformal field rescaling. Setting $\phi = \Omega^{-1}\psi$ eliminates the first-derivative term. Using $R = -6\Omega^{-3}\nabla^2\Omega$ for 4D conformally flat space:

$$\boxed{[-\nabla^2 + V_{\text{eff}}]\psi = \omega^2\psi, \quad V_{\text{eff}} = m_0^2\Omega^2 + (1 - 6\xi)\Omega^{-1}\nabla^2\Omega} \quad (\text{B9})$$

d. Double-well formation condition. For $\xi < 1/6$ and $\nabla^2\Omega < 0$ near cores, the curvature term contributes **negatively** to V_{eff} , creating attractive wells. The mass term $m_0^2\Omega^2$ provides the continuum threshold at infinity.

e. Self-consistency check. At $\xi = \xi_c = 1/6$ (conformal coupling): $V_{\text{eff}} \rightarrow m_0^2\Omega^2$, and the curvature contribution vanishes identically. $\checkmark$

3. Minimal Dirac Coupling in the Conformal Background

To make the visibility block explicitly compatible with the fermionic Yukawa origin, we summarize the minimal conformal Dirac reduction used by the baseline overlap kernel.

For $g_{\mu\nu} = \Omega^2\eta_{\mu\nu}$, choose

$$e_\mu^a = \Omega \delta_\mu^a, \quad e_a^\mu = \Omega^{-1} \delta_a^\mu, \quad \sqrt{-g} = \Omega^4. \quad (\text{B10})$$

The curved-space Dirac action is

$$S_\Psi = \int d^4x \sqrt{-g} \, i \bar{\Psi} e_a^\mu \gamma^a \left(\partial_\mu + \frac{1}{4} \omega_{\mu bc} \gamma^{bc} \right) \Psi, \quad (\text{B11})$$

and the canonical Weyl rescaling

$$\Psi = \Omega^{-3/2} \psi \quad (\text{B12})$$

removes the first-derivative conformal piece from the spin connection, yielding the flat-frame kinetic normalization. For the Higgs mode we use

$$H = \Omega^{-1} \varphi. \quad (\text{B13})$$

The Yukawa density therefore carries net conformal power

$$\Omega^4 \Omega^{-3/2} \Omega^{-1} \Omega^{-3/2} = \Omega^0, \quad (\text{B14})$$

so the minimal overlap kernel has

$$\mathcal{W}_{\text{frame}} = 1. \quad (\text{B15})$$

This is the baseline convention in the three-channel effective-Yukawa extraction; nonzero frame power is treated only as a diagnostic deformation.

As a reproducibility check, the minimal Dirac-conformal audit exports a machine-readable CSV and paper copy verifying the zero net conformal power in both kinetic and Yukawa factors.

4. 1D On-Axis Reduction

For efficient scanning, we use the on-axis ($\rho = 0$) reduction:

$$\left[-\frac{d^2}{dz^2} + U(z)\right]\psi_n(z) = E_n\psi_n(z), \quad U(z) = V_{\text{eff}}(\rho = 0, z) - m_0^2 \quad (\text{B16})$$

where $E_n = \omega_n^2 - m_0^2$. Bound states satisfy $E_n < 0$; stable modes additionally require $\omega_n^2 = m_0^2 + E_n > 0$.

5. 2D Axisymmetric Validation

We validated the 1D reduction by comparing with the full 2D axisymmetric Laplacian:

$$\nabla^2\Omega = \frac{\partial^2\Omega}{\partial\rho^2} + \frac{1}{\rho}\frac{\partial\Omega}{\partial\rho} + \frac{\partial^2\Omega}{\partial z^2} \quad (\text{B17})$$

For the physical-gap parameters ($a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$), the relative error at $\rho = 0$ is:

D	$\text{Max } U_{2\text{D}} - U_{1\text{D}} / \text{max } U $
6	4.0×10^{-4}
12	4.0×10^{-4}
18	4.0×10^{-4}

The 1D 3D-identity approximation is accurate to $< 0.1\%$, validating its use for phase scans.

6. WKB Action Convention

The tunneling action through the central barrier is:

$$S_N = \int_{z_1}^{z_2} \sqrt{U(z) - E_N} dz \quad (\text{B18})$$

where $z_{1,2}$ are the turning points satisfying $U(z_{1,2}) = E_N$. The tunneling suppression factor is:

$$r_N = \eta \cdot e^{-2S_N} \quad (\text{B19})$$

Convention note: For rates/probabilities we use e^{-2S} , while the level-splitting amplitude scales as e^{-S} . The empirical splitting relation below therefore tests the amplitude-level law.

7. Splitting–Action Consistency Check

The most stringent internal consistency check is the relation between level splitting and tunneling action. For symmetric double wells, WKB theory predicts $\Delta E \propto e^{-S}$. From our

action-derived scan:

$$\boxed{\ln \Delta E \approx -1.01 S_1 + 0.69, \quad R^2 = 0.9999} \quad (\text{B20})$$

This near-perfect linear relation (Fig. 10) confirms that the **same** V_{eff} controls both the spectrum and the tunneling kinetics—not an engineered coincidence.




FIG. 10. Splitting-action consistency check. The level splitting $\Delta E = E_2 - E_1$ follows $\exp(-S_1)$ to $R^2 = 0.9999$, confirming the action-derived V_{eff} governs both spectrum and tunneling.

8. Action-Derived Numerical Results

Table IX shows the key spectral and tunneling quantities from the action-derived calculation. Grid convergence: $|E_1|$ variation $< 0.2\%$ for $dz \in [0.01, 0.02]$ at fixed $z_{\text{max}} = 80$.

9. Bound-State omega Convergence Benchmark

To make the surrogate-vs-exact distinction explicit, we report a direct bound-state benchmark for ω_N from the same action-derived 1D operator chain used in Table IX. The settings are

TABLE IX. Action-derived spectral and tunneling data. Parameters: $a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$.

D	E_1	E_2	ω_1	S_1	ΔE
4	-0.505	-0.390	0.704	2.92	1.15×10^{-1}
8	-0.476	-0.469	0.724	5.68	6.32×10^{-3}
12	-0.478	-0.478	0.722	8.51	3.69×10^{-4}
16	-0.481	-0.481	0.721	11.33	2.19×10^{-5}
20	-0.482	-0.482	0.720	14.16	1.30×10^{-6}

fixed to the physical-gap parameter set ($a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$), $z_{\max} = 80$, and three resolutions: coarse ($N_z = 4001$), mid ($N_z = 6001$), fine ($N_z = 8001$). The reproducible script is `code/extract_omega_exact_convergence.py`.

TABLE X. Bound-state ω_N convergence benchmark at $D = \{6, 12, 18\}$ from the action-derived operator chain. Relative errors are quoted versus fine grid.

D	level	N_z	ω_1	ω_2	$\Delta\omega_{12}$	max rel. ω vs fine
6	coarse	4001	0.720883	0.738989	1.8107×10^{-2}	8.65×10^{-4}
6	mid	6001	0.721345	0.739468	1.8123×10^{-2}	2.18×10^{-4}
6	fine	8001	0.721500	0.739629	1.8129×10^{-2}	0
12	coarse	4001	0.721781	0.722035	2.5438×10^{-4}	8.84×10^{-4}
12	mid	6001	0.722258	0.722513	2.5533×10^{-4}	2.23×10^{-4}
12	fine	8001	0.722418	0.722674	2.5565×10^{-4}	0
18	coarse	4001	0.719404	0.719408	3.6700×10^{-6}	8.92×10^{-4}
18	mid	6001	0.719885	0.719888	3.6943×10^{-6}	2.25×10^{-4}
18	fine	8001	0.720046	0.720050	3.7025×10^{-6}	0

This benchmark is the operator-level reference for ω_N in the bound-state convention. In contrast, Appendix D uses generalized localized-extraction eigenvalues λ for mixing-channel extraction; the two quantities serve different roles and are not identified.

10. Fixed-dz Convergence

We verified numerical stability using fixed grid spacing dz (rather than fixed N_z):

- For $dz \in \{0.04, 0.02, 0.01\}$ and $z_{\max} \in \{60, 80\}$: E_1 variation $< 0.3\%$, S_1 variation $< 0.2\%$.
- Results are independent of $z_{\max}$ for $z_{\max} > 60$ (domain truncation error negligible).

Appendix C: First-Principles Symmetry-Channel Test for χ

Using the explicit definition in Eq. (73), we computed M_{12} and $\chi_N^{(\text{sym})}$ on 2D axisymmetric grids for $D = \{6, 12, 18\}$ with coarse/mid/fine resolutions. The modes are normalized with the cylindrical measure $2\pi\rho d\rho dz$, and orthogonality satisfies $|\langle\psi_1, \psi_2\rangle| < 10^{-15}$ in all runs.

TABLE XI. Fine-grid symmetry-channel extraction using Eq. (73). Parameters: $a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$.

D	$ M_{12}^{(\text{sym})} $	$\chi_N^{(\text{sym})}$	$ \langle\psi_1, \psi_2\rangle $
6	2.09×10^{-16}	1.66×10^{-19}	3.12×10^{-17}
12	6.03×10^{-16}	5.69×10^{-19}	3.47×10^{-18}
18	6.24×10^{-16}	6.20×10^{-19}	2.11×10^{-16}

These values are numerically consistent with symmetry-protected cancellation of the parity-symmetric overlap channel. Across coarse/mid/fine grids, $M_{12}^{(\text{sym})}$ changes sign but remains at $\mathcal{O}(10^{-16})$, indicating no resolved nonzero mixing in this channel.

For this reason, the phenomenological Rank-2 coefficient in the main text is interpreted as a localized-channel effective coupling (Section VIII), not as $\chi_N^{(\text{sym})}$ from Eq. (73).

Appendix D: Localized-Channel First-Principles Extraction of χ

We compute the nonzero mixing channel directly in localized basis using

$$M_{LR}^{(H)} = \frac{\lambda_2 - \lambda_1}{2}, \quad \chi_N^{(LR)} = \frac{|M_{LR}^{(H)}|}{\bar{\Gamma}_N}, \quad (\text{D1})$$

with $\bar{\Gamma}_N$ from Eq. (33). Here $\lambda_{1,2}$ are generalized operator eigenvalues from $K\psi = \lambda M\psi$ in the localized extraction solver; they are not the bound-state energy convention $E = \omega^2 - m_0^2 < 0$ used in Appendix B. The extraction uses `code/extract_chi_localized_2d.py` with fixed settings:

- Domain: $\rho_{\text{max}} = 3.0$, $z_{\text{max}} = D/2 + 6.0$.
- Grids: coarse $(d\rho, dz) = (0.12, 0.06)$, mid $(0.08, 0.04)$, fine $(0.06, 0.03)$.
- Solver: generalized eigensystem (shift-invert low-mode targeting, $\sigma = 2.5$), tolerance 10^{-8} , maxiter 3×10^4 .
- Acceptance criteria: $\delta_{\Delta E} \equiv |\Delta E_{\text{level}} - \Delta E_{\text{fine}}|/|\Delta E_{\text{fine}}| < 5\%$, $\delta_\chi \equiv |\chi_{\text{level}} - \chi_{\text{fine}}|/|\chi_{\text{fine}}| < 5\%$, and null-channel absolute bound $|M_{12}^{(\text{sym})}| < 10^{-12}$.

TABLE XII. Localized-channel extraction for $D = \{6, 12, 18\}$ (coarse/mid/fine). Parameters: $a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$.

D	level	$(d\rho, dz)$	(N_ρ, N_z)	λ_1	λ_2	$M_{LR}^{(H)}$	$\chi_N^{(LR)}$	$ M_{12}^{(\text{sym})} $
6	coarse	(0.12, 0.06)	(25, 300)	2.14765	2.60073	2.26541×10^{-1}	4.13261×10^{-4}	5.29×10^{-17}
6	mid	(0.08, 0.04)	(38, 450)	2.14362	2.59798	2.27178×10^{-1}	4.17350×10^{-4}	2.44×10^{-19}
6	fine	(0.06, 0.03)	(50, 600)	2.16611	2.62106	2.27474×10^{-1}	4.01827×10^{-4}	1.56×10^{-18}
12	coarse	(0.12, 0.06)	(25, 400)	2.36032	2.71730	1.78492×10^{-1}	2.27261×10^{-4}	2.89×10^{-17}
12	mid	(0.08, 0.04)	(38, 600)	2.35598	2.71374	1.78883×10^{-1}	2.29384×10^{-4}	6.96×10^{-17}
12	fine	(0.06, 0.03)	(50, 800)	2.37821	2.73632	1.79054×10^{-1}	2.21414×10^{-4}	1.03×10^{-17}
18	coarse	(0.12, 0.06)	(25, 500)	2.22204	2.49466	1.36311×10^{-1}	2.18683×10^{-4}	1.94×10^{-17}
18	mid	(0.08, 0.04)	(38, 750)	2.21686	2.49001	1.36575×10^{-1}	2.21053×10^{-4}	1.54×10^{-15}
18	fine	(0.06, 0.03)	(50, 1000)	2.23864	2.51202	1.36694×10^{-1}	2.13187×10^{-4}	2.19×10^{-16}

The complete pipeline, scripts, and raw tables used here are available in the project repository: <https://github.com/boypatrick/PSLT>.

Across all non-fine runs in Table XII, we obtain

$$\max \delta_{\Delta E} = 4.10 \times 10^{-3}, \quad \max \delta_\chi = 3.86 \times 10^{-2}, \quad \max |M_{12}^{(\text{sym})}| = 1.54 \times 10^{-15}, \quad (\text{D2})$$

which satisfies the stated convergence and absolute-null criteria.

For interpolation support and reproducibility of scan-level diagnostics, we additionally ran a fine-grid-only auxiliary sweep on integer separations $D = 4, \dots, 20$ using the same solver settings (`--full-scan` in `extract_chi_localized_2d.py`). These points are used as profile-support data in the repository, while the formal acceptance criteria above remain defined by the benchmark convergence set $D = \{6, 12, 18\}$.

Figure 11 shows the fine-grid 2D parity eigenmodes used in the localized extraction at $D = \{6, 12, 18\}$. The expected even/odd nodal structure is manifest and remains stable across the three benchmark separations.

To test profile uncertainty in the global map, we benchmarked six $\chi_N^{(LR)}(D)$ choices built from the same fine-grid knots $(D, \chi_N^{(LR)}) = \{(6, 4.02 \times 10^{-4}), (12, 2.21 \times 10^{-4}), (18, 2.13 \times 10^{-4})\}$: linear interpolation, log-linear exponential fit $\ln \chi = -7.5956 - 0.05282 D$, and $\pm 20\%$ amplitude rescalings of each profile. Re-running the full 60×60 scan keeps the generation fractions fixed and yields a narrow $H \rightarrow \mu\mu$ acceptance spread:

$$f(\mathcal{R}_3 > 0.90) = 0.927, \quad f(\mathcal{R}_3 > 0.95) = 0.217, \quad f(\chi_{\mu\mu}^2 < 4) \in [0.1625, 0.1681]. \quad (\text{D3})$$

The best-fit point moves within $(D, \eta) \approx (8.07\text{--}9.42, 1.29\text{--}3.10)$ across these six profiles. Thus, at current scan resolution, profile-interpolation uncertainty is subdominant for $\mathcal{R}_3$ and controlled

TABLE XIII. Auxiliary fine-grid localized profile points from the full-scan run ($D = 4\text{--}20$). Source CSV: `output/chi_fp_2d/localized_chi_D4-5-...-20.csv`.

D	λ_1	λ_2	$\chi_N^{(LR)}$
4	2.06671	2.56696	5.266×10^{-4}
5	2.34635	2.85471	3.311×10^{-4}
6	2.16611	2.62106	4.018×10^{-4}
7	2.42273	2.88524	2.661×10^{-4}
8	2.24778	2.66479	3.202×10^{-4}
9	2.10111	2.48071	3.758×10^{-4}
10	2.32055	2.70644	2.623×10^{-4}
11	2.17620	2.52926	3.064×10^{-4}
12	2.37821	2.73632	2.214×10^{-4}
13	2.24440	2.57510	2.554×10^{-4}
14	2.42772	2.76173	1.906×10^{-4}
15	2.30010	2.61030	2.181×10^{-4}
16	2.47457	2.78817	1.661×10^{-4}
17	2.34884	2.64089	1.895×10^{-4}
18	2.23864	2.51202	2.132×10^{-4}
19	2.39519	2.67161	1.663×10^{-4}
20	2.28581	2.54509	1.867×10^{-4}

TABLE XIV. Extended profile-sensitivity test for $\chi_N^{(LR)}(D)$ in the 60×60 global scan. Source data file: `paper/chi_profile_robustness.csv`.

Profile	$f(\mathcal{R}_3 > 0.90)$	$f(\mathcal{R}_3 > 0.95)$	$f(\chi_{\mu\mu}^2 < 4)$	Best (D, η)
Linear interpolation	0.927	0.217	0.168	(9.42, 1.29)
Exponential fit	0.927	0.217	0.163	(9.15, 1.94)
Linear $\times 0.8$	0.927	0.217	0.167	(8.61, 3.03)
Linear $\times 1.2$	0.927	0.217	0.168	(8.07, 2.26)
Exp-fit $\times 0.8$	0.927	0.217	0.163	(8.61, 3.10)
Exp-fit $\times 1.2$	0.927	0.217	0.163	(8.88, 2.13)

for the $H \rightarrow \mu\mu$ acceptance metric.

As a stronger stress test, we further rescaled the localized profile by large factors and tracked boundary movement directly on the (D, η) grid. Figure 12 shows that, at the present 60×60 scan resolution, the $\mathcal{R}_3 > 0.90$ mask is stable up to $\times 10^4$ and first shifts at $\times 10^5$, while the $H \rightarrow \mu\mu$ acceptance mask already shifts for the user-requested $\times 0.5$ and $\times 2$ cases (changed-cell fractions 0.0331 and 0.0258, respectively).

Appendix C: 2D Axisymmetric Eigenmodes for Localized-Channel Extraction
 Fine grids with $(d\rho, dz) = (0.06, 0.03)$ for $D = \{6, 12, 18\}$




FIG. 11. Fine-grid 2D axisymmetric eigenmodes used in Appendix D. Each row corresponds to one benchmark separation ($D = 6, 12, 18$); columns show the lowest even/odd parity modes (ψ_1, ψ_2) entering the localized-channel extraction pipeline.

Appendix E: Preliminary First-Principles Replacement Candidate Checks

1. 1D Phase-Space Candidate for g_N

To test a first-principles replacement direction for Eq. (18), we evaluated the semiclassical 1D candidate on the same action-derived axis potential:

$$\rho_{\text{WKB}}(E; D) = \frac{1}{2\pi} \int_{U(z; D) < E} \frac{dz}{\sqrt{E - U(z; D)}}, \quad (\text{E1})$$




FIG. 12. Strong χ -amplitude stress test on the global map. Left: overlay of $\chi^2_{\mu\mu} \leq 4$ boundary for selected scales, including the requested $\times 0.5$ and $\times 2$. Right: changed-cell fractions versus scale (relative to baseline $\times 1$), showing early movement for the $H \rightarrow \mu\mu$ mask and delayed movement for the $\mathcal{R}_3$ mask. Source data file: `output/chi_fp_2d/chi_scale_stress_test.csv`.

TABLE XV. Preliminary 1D phase-space candidate $g_N^{(\text{ps})}$ at $D = 12$.

level	dz	E_1	E_2	E_3	$g_1^{(\text{ps})}$	$g_2^{(\text{ps})}$	max rel. g vs fine
coarse	0.04	-0.636696	-0.608074	-0.575010	1.056725	1.054421	1.31×10^{-2}
mid	0.02	-0.637061	-0.608419	-0.575335	1.066813	1.074099	5.30×10^{-3}
fine	0.01	-0.637233	-0.608582	-0.575488	1.063919	1.068433	0

$$g_N^{(\text{ps})}(D) = 1 + \int_{E_{\min}(D)}^{E_N(D)} \rho_{\text{WKB}}(E'; D) dE', \quad E_{\min}(D) \equiv \min_z U(z; D). \quad (\text{E2})$$

Using the reproducible script `extract_gn_phase_space_candidate.py` (in the `code/` directory), we ran a coarse/mid/fine test at $D = 12$ (Table XV). The candidate is numerically stable at the few- 10^{-3} level versus fine grid.


2. 2D Axisymmetric Benchmark for $g_N^{(\text{ps})}$

To reduce projection ambiguity, we also ran a 2D axisymmetric benchmark on $D = \{6, 12, 18\}$ using the same localized extraction chain as Appendix D: identical generalized operator build, identical box/range convention, and the same shift-invert low-mode settings ($\sigma = 2.5$, $\text{tol} = 10^{-8}$, $\text{maxiter} = 3 \times 10^4$). We define

$$N_{\text{ps}}(E) = \frac{1}{4\pi} \int d^2x d^2p \Theta(E - U - p^2) = \frac{1}{2} \int \rho d\rho dz [E - U(\rho, z)]_+, \quad (\text{E3})$$

TABLE XVI. Fine-grid 2D phase-space profile candidate at $D = \{6, 12, 18\}$.

D	λ_1	λ_2	λ_3	$\hat{g}_1$	$\hat{g}_2$
6	0.719476	0.804126	0.945018	9.867×10^{-2}	4.370×10^{-1}
12	0.686186	0.747510	0.822830	1.193×10^{-1}	5.146×10^{-1}
18	0.667848	0.713807	0.767832	1.291×10^{-1}	5.294×10^{-1}

Low- N comparison: legacy Cardy reference vs 2D phase-space profileFIG. 13. Low- N comparison between legacy Cardy reference and 2D phase-space profile on $D = \{6, 12, 18\}$, shown as normalized ratios g_N/g_3 (log scale). The highlighted band marks $N = 1, 2, 3$.Source files: `paper/gn_cardy_vs_phase_space.png`, `paper/gn_cardy_vs_phase_space.csv`.

$$\mathcal{W}_N \equiv [\lambda_1, \lambda_N], \quad g_{N,\text{raw}}^{(\text{ps})} = 1 + \int_{\mathcal{W}_N} \rho_{\text{ps}}(E') dE' = 1 + N_{\text{ps}}(\lambda_N) - N_{\text{ps}}(\lambda_1), \quad \hat{g}_N \equiv \frac{g_{N,\text{raw}}^{(\text{ps})}}{g_{3,\text{raw}}^{(\text{ps})}}. \quad (\text{E4})$$

The two-dimensional phase-space extraction exports both the low- N profile table and the low-mode spectral table with explicit window bounds. For non-fine levels, the maximum profile deviation is

$$\max_{\text{non-fine}} \left(\max_N \frac{|\hat{g}_N - \hat{g}_N^{(\text{fine})}|}{|\hat{g}_N^{(\text{fine})}|} \right) = 2.35 \times 10^{-2}. \quad (\text{E5})$$

As in the 1D candidate check, this benchmark is reported as a replacement candidate. To test map-level impact, we connected the baseline, legacy Cardy comparator, one-dimensional phase-space candidate, and two-dimensional phase-space candidate to the main scan through the reproducible script `scan_gn_profile_impact.py`. Results in Table XVII show that the main conclusion is retained: relative to the bounded two-dimensional baseline, the largest absolute shift is 0.0669 in $f(\mathcal{R}_3 > 0.90)$ (for the raw two-dimensional phase-space candidate), while high- N runaway remains controlled with $f(N_{\text{win}} > 3) \simeq 2.78 \times 10^{-4}$ in all tested modes.

Across $N_{\text{max}} = 20, 30, 40$, both baseline and fp_2d scenarios are numerically stable at map level: the key fractions $f(\mathcal{R}_3 > 0.90)$, $f(\chi_{\mu\mu}^2 < 4)$, and $f(N_{\text{win}} > 3)$ remain unchanged in this

TABLE XVII. Map-level impact of first-principles g_N profiles on the 60×60 scan (baseline is `fp_2d_full`; legacy Cardy retained as comparator). Source CSV: `paper/gn_profile_impact.csv`.

case	$f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	$f(N_{\text{win}} > 3)$	mean tail prob.
baseline <code>fp_2d_full</code>	0.927	0.1500	0.0003	0.1267
legacy Cardy reference	0.877	0.1486	0.0003	0.1125
g_N fp_1d profile	0.860	0.1486	0.0003	0.1137
g_N fp_2d profile	0.860	0.1492	0.0003	0.1144

TABLE XVIII. N_{max} convergence check for baseline and g_N fp_2d map-level summaries. Source CSV: `paper/gn_nmax_convergence.csv`.

case	N_{max}	$f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	$f(N_{\text{win}} > 3)$	mean tail prob.
baseline <code>fp_2d_full</code>	20	0.927	0.1500	0.0003	0.126709
baseline <code>fp_2d_full</code>	30	0.927	0.1500	0.0003	0.126709
baseline <code>fp_2d_full</code>	40	0.927	0.1500	0.0003	0.126709
g_N fp_2d profile	20	0.860	0.1492	0.0003	0.114352
g_N fp_2d profile	30	0.860	0.1492	0.0003	0.114352
g_N fp_2d profile	40	0.860	0.1492	0.0003	0.114352

scan.

3. Baseline Cross-Check Under $g_{\text{mode}} = \text{fp_2d_full}$

To implement a strict same-framework upgrade, we set the bounded two-dimensional phase-space profile as the baseline. Instead of direct low- N substitution plus geometric continuation, this mode uses an explicit bounded microcanonical window for $N \leq 3$ and an energy-gap tail prescription for $N > 3$:

$$g_N^{(\text{full})}(D) = g_3^{(\text{ps})}(D) \hat{g}_N(D), \quad \hat{g}_n = \left(\hat{g}_n^{(\text{dir})} \right)^{1-\alpha} \left(\hat{g}_n^{(\text{win})} \right)^\alpha \quad (n = 1, 2, 3), \quad (\text{E6})$$

with

$$\hat{g}_1^{(\text{win})} = 1 + N_{\text{ps}}(\lambda_3) - N_{\text{ps}}(\lambda_1), \quad \hat{g}_2^{(\text{win})} = 1 + N_{\text{ps}}(\lambda_3) - N_{\text{ps}}(\lambda_2), \quad \hat{g}_3^{(\text{win})} = 1, \quad (\text{E7})$$

with $g_3^{(\text{ps})}(D) \equiv 1 + N_{\text{ps}}(\lambda_3) - N_{\text{ps}}(\lambda_1)$ from the same 2D spectrum table. For $N > 3$,

$$\hat{g}_N^{(\text{target})} = \left(\frac{\Delta N_{\text{ps}}^{(N)}}{\Delta N_{\text{ps}}^{(3)}} \right)^s \exp \left[-\beta \frac{\lambda_N - \lambda_3}{\lambda_3 - \lambda_2} \right], \quad (\text{E8})$$

followed by the same per-step finite-volume clipping used in `pslt_lib.py`. In the present baseline run, we use $(\alpha, \beta, s) = (0.8, 1.1, 0.0)$. Results are summarized in Table XIX (source CSV:

TABLE XIX. Map-level baseline cross-check for g_N : `fp_2d_full` baseline versus legacy Cardy reference on the same 60×60 grid. Source CSV: `paper/gn_baseline_replacement.csv`.

case	$f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	$f(N_{\text{win}} > 3)$	$\Delta f(\mathcal{R}_3 > 0.90)$ vs baseline
baseline <code>fp_2d_full</code>	0.9269	0.1500	0.0003	0.0000
legacy Cardy reference	0.8769	0.1486	0.0003	-0.0500

`paper/gn_baseline_replacement.csv`). Relative to baseline `fp_2d_full`, the legacy Cardy reference lowers $f(\mathcal{R}_3 > 0.90)$ by 0.0500 (from 0.9269 to 0.8769), lowers $f(\chi_{\mu\mu}^2 < 4)$ by 0.0014 (from 0.1500 to 0.1486), and keeps $f(N_{\text{win}} > 3) \simeq 2.78 \times 10^{-4}$. Therefore, the map-level main conclusion is retained under this baseline switch.

Across $N_{\text{max}} = 20, 30, 40$, both rows in Table XIX are numerically stable in this setup (source CSV: `paper/gn_baseline_replacement_nmax.csv`).

4. Geometry-Informed and Microscopic Lindblad Scan-Ready Module for Open-System

χ

As a methodological bridge (not a replacement of Appendix D), we consider a two-level open-system model in the standard GKSL/Lindblad form [10, 11]:

$$\dot{\rho} = -i[H, \rho] + \gamma_{\phi} \mathcal{D}[\sigma_z] \rho + \gamma_{\text{mix}} \mathcal{D}[\sigma_x] \rho, \quad H = \begin{pmatrix} 0 & \Delta/2 \\ \Delta/2 & 0 \end{pmatrix}. \quad (\text{E9})$$

As external motivation for this diagnostic direction, recent entanglement-based flavor analyses report that minimizing scattering-generated entanglement can qualitatively reproduce observed CKM/PMNS patterns [12]. In this manuscript we use that connection only as qualitative support for coherence-sensitive diagnostics, not as a baseline closure assumption. For this module we set $\gamma_{\text{mix}}(D) \equiv M_{LR}^{(H)}(D)$ from Appendix D, and define a geometry-informed dephasing proxy

$$\gamma_{\phi}(D) \equiv \sqrt{\left\langle (\delta V)^2 \right\rangle_{\rho}}, \quad \delta V = V_{\text{eff}} - \bar{V}_{\text{eff}}(r), \quad (\text{E10})$$

with cylindrical weighting $\langle \cdots \rangle_{\rho}$. The minimal Lindblad evolution and geometry-profile extraction output reproducible diagnostics such as

$$C_{\text{max}} \equiv \max_t |\rho_{LR}(t)|, \quad P_{\text{max}} \equiv \max_t \rho_{RR}(t), \quad (\text{E11})$$

and a normalized proxy $\chi_{\text{eff}}^{(\text{proxy})} = 2\gamma_{\text{mix}} C_{\text{max}} / \bar{\Gamma}_N$. Results for $D = \{6, 12, 18\}$ are summarized in Table XX. Extending the same profile to $D = 4, \dots, 20$ gives a stable geometry-informed ratio band $\chi_{\text{open}} / \chi_{LR} \in [0.0852, 0.2127]$ with mean 0.1269. This value belongs to the geometry-only

proxy closure in this subsection; the anchor-calibrated microscopic band is reported separately in Eq. (E13) and Table XXI.

To move this module toward microscopic closure, we also extract localized-basis couplings (g_z, g_x) directly from the same 2D eigenmodes and set

$$\gamma_\phi = \kappa_{\text{env}} g_z^2 S(0), \quad \gamma_{\text{mix}} = \kappa_{\text{env}} g_x^2 S(\Delta E), \quad S(\omega) = \frac{2\tau_{\text{env}}}{1 + (\omega\tau_{\text{env}})^2}, \quad (\text{E12})$$

from the localized microscopic extraction. To avoid map-level fitting, we calibrate κ_{env} with a multi-anchor rule: least-squares match to geometry-profile ratios on anchors $D = \{6, 9, 12, 15, 18\}$, then validate on holdout $D = \{4, 5, 7, 8, 10, 11, 13, 14, 16, 17, 19, 20\}$. This gives $\kappa_{\text{env}} = 4.31 \times 10^5$ with anchor RMSE 8.61×10^{-2} and holdout RMSE 6.95×10^{-2} (holdout max-abs error 1.25×10^{-1}). The corresponding scan-level band is

$$\chi_{\text{open}}^{(\text{micro})}/\chi_{LR} \in [0.0173, 0.3271], \quad \langle \chi_{\text{open}}^{(\text{micro})}/\chi_{LR} \rangle = 0.1380, \quad (\text{E13})$$

from the microscopic ratio-band audit. To pin down the microscopic bookkeeping, we additionally export a bridge audit with map, summary, and figure artifacts. On the calibrated $D = 4, \dots, 20$ knot set, this audit reconstructs

$$\gamma_\phi = \kappa_{\text{env}} g_z^2 S(0), \quad \gamma_{\text{mix}} = \kappa_{\text{env}} g_x^2 S(\Delta E), \quad \chi_{\text{eff}}^{(\text{micro})} = 2\gamma_{\text{mix}} C_{\text{max}}/\Gamma_{\text{ref}}$$

with numerically negligible residuals:

$$\max |\Delta\gamma_\phi| = 2.03 \times 10^{-20}, \quad \max |\Delta\gamma_{\text{mix}}| = 4.40 \times 10^{-14}, \quad \max |\Delta\chi_{\text{eff}}^{(\text{loader})}| = 5.58 \times 10^{-17}.$$

The corresponding bridge summary identifies the effective system as the localized two-level (L/R) subspace, the current bath witness as a Gaussian Markov bath with Lorentzian PSD, and the coupling operators as the diagonal/off-diagonal pair (σ_z, σ_x) . This means the open-system gap is no longer in the system-to-Lindblad bookkeeping, but in the microscopic EYMH origin and normalization of the bath itself.

We can therefore push the bath normalization one step further without over-claiming a full microscopic derivation. In the κ_{env} window audit, we treat

$$\kappa_{\text{env}} \rightarrow \kappa_{\text{scale}} \kappa_{\text{env}}^{(\text{calib})}$$

as a uniform bath-normalization susceptibility that rescales both γ_ϕ and γ_{mix} while leaving the localized two-level Hamiltonian $(\Delta E, \sigma_x)$ fixed. On the current $D = 4, \dots, 20$ knot set, the calibration-consistent candidate window is

$$\kappa_{\text{scale}} \in [0.5, 1.0],$$

while a broader scan-stable window

$$\kappa_{\text{scale}} \in [0.25, 1.5]$$

keeps the map-level fractions $f(\mathcal{R}_3 > 0.90)$ and $f(\chi_{\mu\mu}^2 \leq 4)$ unchanged and only relaxes the holdout tolerance slightly. In this sense, κ_{env} is no longer just a single-point fit parameter: within the present witness it acts as a bounded bath-normalization susceptibility. What still remains open is the EYMH-side derivation of its absolute normalization and of the microscopic bath degrees of freedom themselves.

The newer parent-bath audits sharpen this statement. First, they show that the witness already factorizes into a system block (g_z^2, g_x^2) , a bath-shape block $(S_{zz}(0), S_{xx}(\Delta E))$, and an amplitude block κ_{env} , with the canonical constant normalization remaining the unique exact amplitude choice. Second, they rewrite the rates as the projected bath block

$$K_{\text{bath}} = \kappa_{\text{env}} \sqrt{K_{\text{sys}}} K_{\text{spec}} \sqrt{K_{\text{sys}}},$$

and the subsequent family, log-coordinate, normal-coordinate, and generator-affinity audits show that the canonical projected bath class is uniquely selected on the current knot set: $(m, u, v) = (0, 0, 0)$, $(p_{\text{sys}}, p_{\text{spec}}) = (0, 0)$, $(\zeta_{\text{sys}}, \zeta_{\text{spec}}) = (0, 0)$, and $(q_{ss}, q_{bb}, q_{sb}) = (0, 0, 0)$. It is therefore natural to define the normalized projected bath block

$$B_a(z) := \kappa_{\text{env}}^{-1} K_{\text{sys},a}^{-1/2}(z) K_{\text{bath},a}(z) K_{\text{sys},a}^{-1/2}(z),$$

which the audited factorization already identifies as the bath-shape witness itself. Thus the remaining open-system gap is no longer bookkeeping or nearby-family ambiguity; it is the final parent-action statement for why the projected bath generator itself naturally lives in this affine log class. Finally, `scan_chi_open_system_parent_bath_cocycle_audit.py` turns this into a positive integrability statement: after dividing out κ_{env} , the normalized parent bath block defines an exact additive cocycle in the canonical log variables, with pairwise cocycle residuals at 8.95×10^{-16} on the identifiable ϕ subset and 8.64×10^{-13} on the mixing branch, while triangle flatness defects stay at 1.11×10^{-15} in both channels. So the affine log-generator is now supported not only by nearby-family exclusion, but also by exact cocycle/flatness closure. Pushing one step further, `scan_chi_open_system_parent_bath_potential_audit.py` shows that this cocycle is globally integrable on the current knot set: the normalized bath block admits an exact anchored scalar potential $\Phi_a = \log(B_a/B_a^{\text{ref}})$ whose direct decomposition, anchor-average recovery, and nearest-neighbor chain recovery agree to machine precision (max residuals $\sim 10^{-14}$ on the identifiable ϕ subset and $\sim 5 \times 10^{-13}$ or better on the mixing branch). Thus the projected bath witness is no longer characterized only by an affine local generator; it also carries a single-valued




FIG. 14. Projected parent-bath potential closure for the open-system module on the calibrated knot set. The normalized bath block admits a single-valued scalar primitive Φ_a whose direct reconstruction, anchor recovery, and chain recovery agree to machine precision: the maximum residuals are 1.42×10^{-14} on the identifiable ϕ subset and 5.50×10^{-13} on the mixing branch. Source files: `paper/chi_open_system_parent_bath_potential.*`.

anchored potential, so the remaining gap is specifically the parent-action origin of this potential itself.

Figure 14 gives the representative late-stage open-system diagnostic: the normalized parent-bath block has already been reduced to an integrable projected potential rather than an unconstrained reservoir fit.

At this point the cocycle/potential relation can be written explicitly on the audited knot set. Defining the discrete increment

$$G_a(z_i, z_j) := \log B_a(z_j) - \log B_a(z_i),$$

the cocycle audit shows that G_a is additively consistent:

$$G_a(z_i, z_k) = G_a(z_i, z_j) + G_a(z_j, z_k),$$

up to the machine-level residuals quoted above. Fixing an anchor z_* , the potential is then not

an extra ansatz but the integrated cocycle itself,

$$\Phi_a(z) := G_a(z_\star, z), \quad G_a(z_i, z_j) = \Phi_a(z_j) - \Phi_a(z_i).$$

So the anchored potential is simply the scalar primitive of the already-audited normalized bath block in canonical log variables.

This also gives the natural local parent-bath effective-action normal form on the audited knot set. Writing $z = (z_{\text{sys}}, z_{\text{spec}})$ for the canonical log variables and $B_a(z) = B_a^{\text{ref}} e^{\Phi_a(z)}$, the cocycle closure identifies the bath witness as an anchored exponential potential on the projected bath block,

$$K_{\text{bath},a}(z) = \kappa_{\text{env}} K_{\text{sys},a}^{1/2}(z) e^{\Phi_a(z)} K_{\text{sys},a}^{1/2}(z),$$

so, channelwise on the present positive knot set,

$$B_a(z) = e^{\Phi_a(z)}, \quad \log K_{\text{bath},a}(z) = \log \kappa_{\text{env}} + \log K_{\text{sys},a}(z) + \Phi_a(z).$$

while the family, log-coordinate, normal-coordinate, and generator-affinity audits isolate the canonical class as the one with no surviving first local log-curvature jet. Equivalently, on a sufficiently small neighborhood $\mathcal{N}$ of the canonical anchor,

$$\Phi_a(z + \tau) = \Phi_a(z) + \nabla \Phi_a(z) \cdot \tau + R_{a,2}(z, \tau), \quad |R_{a,2}(z, \tau)| \leq \frac{1}{2} \sup_{\zeta \in \mathcal{N}} \|D^2 \Phi_a(\zeta)\| \|\tau\|^2,$$

so the affine log-generator is just the first Taylor jet of Φ_a on the projected bath block, while the first non-affine correction is controlled by $R_{a,2}$. The audited canonical point is precisely the one where this first non-affine correction is rejected on the current knot set. In this sense, the open-system witness is already promoted from a free nearby-family fit to a locally anchored projected effective action; what remains open is not the existence or local normal form of Φ_a , but the microscopic EYMH derivation of its coefficients, of κ_{env} , and of the bath degrees of freedom themselves.

The implication chain here is non-circular. The primitive audited object is the normalized projected bath block itself, namely $K_{\text{bath},a}/\kappa_{\text{env}}$ on the canonical knot set. The nearby-family, log-coordinate, normal-coordinate, and generator-affinity audits act directly on that block and select the unique local affine generator class without assuming any potential. The cocycle audit is then applied to the same normalized block, upgrading this local affine closure to an exact additive consistency statement in the canonical log variables. The potential audit comes only after that step: it integrates the already-audited cocycle and checks path-independence / anchor recovery, rather than feeding any potential ansatz back into the family selection. Equation (E 4) together with the anchored form above therefore repackages the same projected bath witness

into a local effective-action normal form; it does not create new evidence by definition. The remaining open-system gap is still the microscopic parent-action origin of the coefficients and bath degrees of freedom, not the internal consistency of the projected bath block on the present audited knot set.

This also sharpens the remaining microscopic origin question into three narrower obligations. First, since

$$B_a(z) = \kappa_{\text{env}}^{-1} K_{\text{sys},a}^{-1/2}(z) K_{\text{bath},a}(z) K_{\text{sys},a}^{-1/2}(z),$$

one has

$$\partial_{\log \kappa_{\text{env}}} \log B_a(z) = 0, \quad \partial_{\log \kappa_{\text{env}}} \log K_{\text{bath},a}(z) = 1,$$

so κ_{env} already behaves as a pure scalar bath-susceptibility modulus rather than as an independent bath-shape degree of freedom. Second, the coefficients of Φ_a are no longer free local fit parameters on the audited knot set: the cocycle primitive $G_a(z_i, z_j) = \Phi_a(z_j) - \Phi_a(z_i)$ already fixes their discrete local jet, and the remaining task is only to identify the parent-action source of those coefficients. Third, the unresolved bath degrees of freedom have already been reduced to the two effective response channels appearing in the audited micro bridge,

$$\gamma_{\phi,a} = \kappa_{\text{env}} g_{z,a}^2 S_{zz,a}(0), \quad \gamma_{\text{mix},a} = \kappa_{\text{env}} g_{x,a}^2 S_{xx,a}(\Delta E_a),$$

rather than to an arbitrary reservoir family. Thus the next microscopic step, if taken at all, is not to re-justify the local projected effective-action form, but to explain why the parent EYMH dynamics projects specifically onto this scalar modulus, this anchored potential jet, and this minimal two-channel bath kernel.

One can also sharpen the second obligation without yet claiming a full parent-action derivation of the coefficients. Restrict to the identifiable audited subset

$$\mathcal{K}_\phi := \{D : \phi_a(D) \text{ is identifiable}\},$$

and let

$$G_a(D_i, D_j) := \Phi_a(D_j) - \Phi_a(D_i)$$

denote the already-audited cocycle primitive on that subset. For any local anchor $D_\star \in \mathcal{K}_\phi$ and any finite neighborhood $\mathcal{N}(D_\star) \subset \mathcal{K}_\phi$, the discrete first jet coefficient is fixed by the primitive itself, not by a new fit,

$$c_a(D_\star) := \arg \min_c \sum_{D_j \in \mathcal{N}(D_\star)} w_j \left[\Phi_a(D_j) - \Phi_a(D_\star) - c(D_j - D_\star) \right]^2 = \frac{\sum_j w_j (D_j - D_\star) G_a(D_\star, D_j)}{\sum_j w_j (D_j - D_\star)^2},$$

with local remainder

$$r_{a,\star}(D_j) := \Phi_a(D_j) - \Phi_a(D_\star) - c_a(D_\star)(D_j - D_\star), \quad \sum_j w_j(D_j - D_\star) r_{a,\star}(D_j) = 0.$$

So the local coefficient of Φ_a on the audited knot set is already an anchor-locked discrete jet of the cocycle primitive, rather than an unconstrained nearby-family parameter. The current audits then push one step further still: on the entire exported map, the observed parent-bath action coefficient is exactly the signed carrier of the same primitive,

$$A_a(D) := (\text{parent-bath action}_\phi)_a(D) = -\Phi_a(D), \quad c_a^{(A)}(D_\star) = -c_a^{(\Phi)}(D_\star).$$

Numerically, the potential reconstruction residual stays below 1.5×10^{-14} on the identifiable subset, the anchor and chain recoveries stay below 7.2×10^{-15} and 2.2×10^{-14} respectively, the audited cocycle / local-generator defect stays below 7.2×10^{-15} , and the generator/potential representatives agree pointwise at the same 2.2×10^{-14} level. Moreover, the canonical affine class is not merely acceptable but strongly selected: the generator-affinity and normal-coordinate selection gaps exceed the canonical objective by factors of 2.2×10^8 and 4.4×10^8 respectively. Thus the unresolved microscopic issue is no longer whether Φ_a admits a stable local first jet on the audited knot set; it is why the EYMH parent action generates precisely this signed discrete jet and not some nearby nonlinear deformation.

Within that first obligation, the present audits already push one step further than a generic “fitted constant” interpretation. If one allows a nearby absolute-normalization family

$$\kappa_{\text{eff}}(z) = \kappa_{\text{env}} s F(z),$$

with a single anchor-fitted amplitude s but a non-uniform profile $F(z)$, then the normalized projected bath block becomes

$$B_a^{(F)}(z) = \kappa_{\text{eff}}^{-1}(z) K_{\text{sys},a}^{-1/2}(z) K_{\text{bath},a}(z) K_{\text{sys},a}^{-1/2}(z) = (sF(z))^{-1} B_a(z).$$

So an exact amplitude-only renormalization preserves the normalized block if and only if $F(z)$ is constant on the audited knot set. This is exactly what the absolute-normalization audit finds: the canonical locked uniform family keeps the bridge-normalization residuals at machine level ($\sim 8 \times 10^{-13}$ on the observable/mixing channels), whereas the best non-uniform admissible family improves the holdout RMSE only marginally (6×10^{-5}) while inflating the same normalized residuals to $\sim 10^{-2}$. In that precise local sense, κ_{env} is already singled out as the unique exact amplitude modulus, and the remaining microscopic question is no longer whether it mixes with bath shape on the present knot set, but why the parent EYMH dynamics produces this modulus in the first place.

Proposition O1 (Projected two-channel interaction). Without invoking a full bath derivation, the parent perturbation projected to the audited parity doublet already fixes the leading localized interaction to a two-channel Pauli block. Let ΔV denote the parent perturbation projected to the parity doublet, so that on the audited knot set

$$V_{\text{par},a} = \begin{pmatrix} v_{11,a} & v_{12,a} \\ v_{12,a} & v_{22,a} \end{pmatrix}$$

is real and symmetric. Rotating to the localized basis with

$$U_{LR} = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$$

gives

$$V_{\text{loc},a} = U_{LR}^T V_{\text{par},a} U_{LR} = \frac{v_{11,a} + v_{22,a}}{2} I + v_{12,a} \sigma_z + \frac{v_{11,a} - v_{22,a}}{2} \sigma_x.$$

Hence every audited localized two-level block already lies in the span of I, σ_z, σ_x , with no surviving σ_y component: after removing the trace part, the leading projected interaction is at most two-channel. The extracted couplings are therefore not independent fit coordinates but exactly the Pauli coefficients of the localized block,

$$g_{z,a} = |v_{12,a}|, \quad g_{x,a} = \left| \frac{v_{11,a} - v_{22,a}}{2} \right|.$$

This remains an effective two-level statement rather than a completed EYMH bath theorem, but it shows that the present σ_z/σ_x bridge is already the most general leading projection compatible with the audited real localized doublet. Numerically, the full 17-point micro export verifies this decomposition at machine precision: the reconstructed localized entries (V_{LL}, V_{RR}, V_{LR}) agree with the exported map at $\leq 7 \times 10^{-18}$, while $g_z - |v_{12}|$ and $g_x - |(v_{11} - v_{22})/2|$ stay below 2.2×10^{-18} and 1.3×10^{-18} respectively. The remaining microscopic bath question is therefore narrower than a generic reservoir identification: one still has to explain why the parent EYMH environment induces the two on-shell kernels multiplying these projected σ_z and σ_x directions, not why those directions themselves appear in the audited two-level reduction.

Proposition O2 (Single gap-locked response family). Within the present minimal bridge, the dephasing and mixing kernels are not independent shape functions but on-shell evaluations of a single response family. The audited micro extraction uses

$$S_a(\omega) = \frac{2\tau_{\text{env},a}}{1 + (\omega\tau_{\text{env},a})^2},$$

so that

$$S_{zz,a}(0) = S_a(0) = 2\tau_{\text{env},a}, \quad S_{xx,a}(\Delta E_a) = S_a(\Delta E_a) = \frac{2\tau_{\text{env},a}}{1 + (\Delta E_a \tau_{\text{env},a})^2}.$$

Equivalently,

$$\frac{S_{xx,a}(\Delta E_a)}{S_{zz,a}(0)} = \frac{1}{1 + (\Delta E_a \tau_{\text{env},a})^2},$$

so the mixing channel is just the same minimal Lorentzian response sampled on shell at the transition energy, while the dephasing channel is the zero-frequency limit of that same family. On the current 17-point micro export this identification is exact at the CSV level: the reconstructed residuals for $S_{zz,a}(0) = 2\tau_{\text{env},a}$ and $S_{xx,a}(\Delta E_a) = 2\tau_{\text{env},a}/(1 + (\Delta E_a \tau_{\text{env},a})^2)$ vanish to machine precision, while the ratio relation above holds with maximal defect 1.1×10^{-16} . Thus the unresolved EYMH bath-origin problem is narrower still: beyond the already-audited two-channel projection, one only has to explain why the parent environment induces this single minimal response family and why its on-shell samples are taken at $\omega = 0$ and $\omega = \Delta E_a$.

One can push this “single response family” statement one notch further without yet claiming a full bath derivation. The same audited family is exactly the Fourier image of the minimal even single-pole correlation kernel

$$C_a(t) = e^{-|t|/\tau_{\text{env},a}},$$

for which

$$S_a(\omega) = \int_{-\infty}^{\infty} dt C_a(t) e^{i\omega t} = 2 \int_0^{\infty} e^{-t/\tau_{\text{env},a}} \cos(\omega t) dt = \frac{2\tau_{\text{env},a}}{1 + (\omega\tau_{\text{env},a})^2}.$$

Moreover, on the present audited knot set the correlation time is not a second free shape coordinate: the exported micro bridge uses

$$\tau_{\text{env},a} = \max\left(\frac{\tau_{\text{scale}}}{\omega_{1,a}}, \tau_{\text{floor}}\right),$$

and the floor never activates on the current 17-point map, so one has numerically

$$\tau_{\text{env},a} = \frac{1}{\omega_{1,a}}$$

to machine precision throughout the audited window. Hence the minimal response family is already gap-locked,

$$S_a(\omega) = \frac{2/\omega_{1,a}}{1 + (\omega/\omega_{1,a})^2}, \quad \frac{S_{xx,a}(\Delta E_a)}{S_{zz,a}(0)} = \frac{1}{1 + (\Delta E_a/\omega_{1,a})^2}.$$

On the current export the maximal defect in $\tau_{\text{env},a} - \omega_{1,a}^{-1}$ is 1.1×10^{-16} , the on-shell suppression law above again holds to 3.3×10^{-16} , and the resulting band $S_{xx,a}(\Delta E_a)/S_{zz,a}(0)$ stays in the narrow interval $[0.9246, 0.9800]$. In this precise effective sense, the remaining EYMH bath-origin question is no longer why two unrelated kernels appear, but why the parent environment

linearizes to a single dominant pole whose correlation time is locked to the same local gap scale $\omega_{1,a}^{-1}$ that enters the audited two-level reduction.

Proposition O3 (Schur-resolvent reformulation). The dominant-pole question can be stated directly in the same parent-kernel language used in the EYMH subsection. Let

$$R_a(\omega) := u_a^\top (M_a^2 + \omega^2 I)^{-1} u_a$$

denote the projected Schur-resolvent response of the parent bath/complement block in the same localized channel. Its spectral decomposition reads

$$R_a(\omega) = \sum_{n \geq 1} \frac{Z_{n,a}}{\Omega_{n,a}^2 + \omega^2}, \quad Z_{n,a} := |\langle u_a, n, a \rangle|^2 \geq 0.$$

If the first mode dominates on the audited on-shell window,

$$R_a(\omega) \approx \frac{Z_{1,a}}{\Omega_{1,a}^2 + \omega^2},$$

then, after absorbing the residue $Z_{1,a}$ into the overall amplitude block κ_{env} and the projected couplings $g_{z,a}, g_{x,a}$, one recovers exactly the same single-pole family used above. Identifying moreover $\Omega_{1,a} = \omega_{1,a}$ yields the observed gap-locked form

$$S_a(\omega) = \frac{2/\omega_{1,a}}{1 + (\omega/\omega_{1,a})^2}.$$

The present audits do not yet prove this first-mode dominance theorem, but they isolate it as the next genuinely microscopic target rather than leaving an undifferentiated reservoir ansatz. On the EYMH parent side, the log-det / Schur witness and the parent block-determinant / Schur audit already close to machine precision, with $\max |J_{\text{direct}} - J_{\text{parent}}| = 1.30 \times 10^{-18}$, $\max |S_{\text{parent}} - S_{\text{resp}}| = 5.55 \times 10^{-16}$, and $\max |\Delta \hat{G}_{\text{Schur}}| = 6.66 \times 10^{-16}$. On the open-system side, the same audited knot set already satisfies $\tau_{\text{env},a} = \omega_{1,a}^{-1}$ and the on-shell rate identities to machine precision. In this narrowed sense, the unresolved bath-origin problem is no longer why some effective Lorentzian fit works, but why the projected Schur resolvent is first-mode dominant and why its leading pole is locked to the same local gap scale $\omega_{1,a}$.

Theorem Target O1 (First-mode remainder bound). The remaining microscopic step is an explicit remainder estimate rather than a qualitative bath ansatz. Decompose

$$R_a(\omega) = R_a^{(1)}(\omega) + \mathcal{R}_a(\omega), \quad R_a^{(1)}(\omega) := \frac{Z_{1,a}}{\Omega_{1,a}^2 + \omega^2},$$

so that

$$\mathcal{R}_a(\omega) = \sum_{n \geq 2} \frac{Z_{n,a}}{\Omega_{n,a}^2 + \omega^2}.$$

Then the exact relative remainder satisfies

$$\frac{|\mathcal{R}_a(\omega)|}{R_a^{(1)}(\omega)} \leq \sum_{n \geq 2} \frac{Z_{n,a}}{Z_{1,a}} \frac{\Omega_{1,a}^2 + \omega^2}{\Omega_{n,a}^2 + \omega^2}.$$

Hence, if one can show on the audited on-shell window $\omega \in \{0, \Delta E_a\}$ that

$$\varepsilon_a := \sup_{\omega \in \{0, \Delta E_a\}} \sum_{n \geq 2} \frac{Z_{n,a}}{Z_{1,a}} \frac{\Omega_{1,a}^2 + \omega^2}{\Omega_{n,a}^2 + \omega^2} \ll 1,$$

then

$$R_a(\omega) = R_a^{(1)}(\omega) [1 + O(\varepsilon_a)]$$

throughout the same on-shell window. This is the cleanest proof-level formulation of the dominant-pole claim: it is enough to control the complement remainder in the projected Schur resolvent, rather than to derive a full microscopic reservoir from scratch. The current export already shows that this on-shell window is numerically narrow, with $\Delta E_a / \omega_{1,a} \in [0.1430, 0.2857]$ and therefore $S_{xx,a}(\Delta E_a) / S_{zz,a}(0) \in [0.9246, 0.9800]$, so the audited dephasing and mixing channels sample only the low-frequency shoulder of the same pole. What is still missing is exactly the parent-side bound that identifies $\Omega_{1,a} = \omega_{1,a}$ and proves that the higher-mode remainder is uniformly negligible on this window.

A useful quantitative sharpening is immediate once one separates the gap question from the tail-weight question. Define

$$q_a := \frac{R_a(\Delta E_a)}{R_a(0)}, \quad r_a := \frac{\Delta E_a}{\Omega_{1,a}}, \quad \nu_{n,a} := \frac{Z_{n,a} / \Omega_{n,a}^2}{R_a(0)},$$

so that $\sum_{n \geq 1} \nu_{n,a} = 1$ and

$$q_a = \sum_{n \geq 1} \frac{\nu_{n,a}}{1 + (\Delta E_a / \Omega_{n,a})^2}.$$

If the higher modes satisfy the conditional gap bound $\Omega_{n,a} \geq g_a \Omega_{1,a}$ for all $n \geq 2$ with some $g_a > 1$, then with $f(x) := (1 + x^2)^{-1}$ one has

$$q_a \geq \nu_{1,a} f(r_a) + (1 - \nu_{1,a}) f\left(\frac{r_a}{g_a}\right),$$

hence

$$1 - \nu_{1,a} \leq \frac{q_a - f(r_a)}{f(r_a/g_a) - f(r_a)}.$$

Since $\varepsilon_a(0) = \frac{1 - \nu_{1,a}}{\nu_{1,a}}$ and $\varepsilon_a(\Delta E_a) \leq \varepsilon_a(0) \frac{1 + r_a^2}{g_a^2 + r_a^2}$, the theorem target can be restated as a two-knob quantitative task: prove a parent-side lower bound on g_a and a small on-shell defect

$\delta_{q,a} := q_a - f(r_a)$. On the audited knot set, the exported gap-conditioned audit shows that for the representative benchmark $g_a = 2$ one has

$$f(r_a/2) - f(r_a) \in [1.50 \times 10^{-2}, 5.55 \times 10^{-2}], \quad \frac{1 + r_a^2}{4 + r_a^2} \in [2.54 \times 10^{-1}, 2.65 \times 10^{-1}],$$

so even a future parent-side match at the modest level $\delta_{q,a} \leq 10^{-3}$ would already force

$$1 - \nu_{1,a} \leq 6.68 \times 10^{-2}, \quad \varepsilon_a(\Delta E_a) \leq 1.82 \times 10^{-2}$$

uniformly across the current audited window. The present effective export, by contrast, has $\max |\delta_{q,a}| = 3.33 \times 10^{-16}$ because the fitted response family is exactly single-pole on shell. This does not yet prove the parent-side gap itself, but it converts the remaining Schur-resolvent problem into a sharply quantified gap-plus-defect statement rather than a diffuse reservoir ansatz.

One can now push this one notch further on the parent side without yet claiming a full spectral theorem. Write the background-normalized canonical complement block as

$$\tilde{K}_a := \begin{pmatrix} K_{11,a} & C_a/\sqrt{K_{\text{bg},a}} \\ C_a/\sqrt{K_{\text{bg},a}} & G_a \end{pmatrix}, \quad G_a := K_{22,a}/K_{\text{bg},a},$$

with the audited geometric-mean identity

$$\frac{C_a^2}{K_{\text{bg},a}} = (K_{11,a} - 1)(G_a - 1).$$

Hence

$$\tilde{K}_a = I + v_a v_a^\top, \quad v_a = \begin{pmatrix} \sqrt{K_{11,a} - 1} \\ \sqrt{G_a - 1} \end{pmatrix},$$

so the two proxy eigenvalues are fixed exactly by

$$\lambda_{1,a}^{(\text{proxy})} = 1, \quad \lambda_{2,a}^{(\text{proxy})} = 1 + \|v_a\|^2 = K_{11,a} + G_a - 1.$$

Calibrating the identity mode to the already-audited gap-locked pole then gives the calculable parent-side proxies

$$\begin{aligned} \Omega_{1,a}^{(\text{proxy})} &:= \omega_{1,a}, \\ \Omega_{2,a}^{(\text{proxy})} &:= \omega_{1,a} \sqrt{K_{11,a} + G_a - 1}, \\ g_a^{(\text{proxy})} &:= \frac{\Omega_{2,a}^{(\text{proxy})}}{\Omega_{1,a}^{(\text{proxy})}} = \sqrt{K_{11,a} + G_a - 1}. \end{aligned}$$

The corresponding audit closes to machine precision:

$$\max |C_a^2/K_{\text{bg},a} - (K_{11,a} - 1)(G_a - 1)| = 2.78 \times 10^{-16},$$

$$\max \left| \lambda_{1,a}^{(\text{proxy})} - 1 \right| = \max \left| \lambda_{2,a}^{(\text{proxy})} - (K_{11,a} + G_a - 1) \right| = 4.44 \times 10^{-16}.$$

On the nontrivial active rows of the D21xE21 fix grid this already yields

$$g_a^{(\text{proxy})} \in [1.3537, 1.5609], \quad \text{median } g_a^{(\text{proxy})} = 1.4692.$$

On the overlap with the currently audited open-system knot set, the induced separator band becomes

$$f(r_a/g_a^{(\text{proxy})}) - f(r_a) \in [1.43 \times 10^{-2}, 4.03 \times 10^{-2}],$$

so a future parent-side defect estimate as mild as $\delta_{q,a} \leq 10^{-3}$ would already imply

$$1 - \nu_{1,a} \leq 6.98 \times 10^{-2}, \quad \varepsilon_a(\Delta E_a) \leq 3.60 \times 10^{-2}.$$

This is still only a proxy, not a proof that the true Schur-resolvent spectrum has the same gap. But it does show that the canonical projected EYMH complement block already carries a concrete, calculable gap floor rather than leaving g_a completely unconstrained.

One can also turn this proxy into a local operator-family lower bound by combining it with the exact parent mismatch functional from Proposition U4. Let

$$\delta\theta := \begin{pmatrix} \alpha - 1 \\ \beta - 1 \end{pmatrix}, \quad Q_{\text{loc}}(\delta\theta, \lambda) := \frac{1}{2} \delta\theta^\top H \delta\theta + C_4 \lambda^4,$$

where the stationarity audit supplies the exact quadratic matrix H in the (α, β) sector and the exact quartic coefficient C_4 in the mixing direction. On the audited local box

$$|\alpha - 1| \leq 0.1, \quad |\beta - 1| \leq 0.1, \quad |\lambda| \leq 0.2,$$

the parent-kernel scan gives the empirical coercivity bound

$$J_{\text{parent}}(\alpha, \beta, \lambda) \geq c_{\text{loc}} Q_{\text{loc}}(\delta\theta, \lambda), \quad c_{\text{loc}} = 3.98 \times 10^{-1},$$

with $\lambda_{\min}(H) = 1.33 \times 10^{-3}$ and $C_4 = 1.497 \times 10^{-3}$. Hence any local operator-side candidate obeying $J_{\text{parent}} \leq J_*$ inside that box must satisfy

$$\|\delta\theta\| \leq \rho_* := \sqrt{\frac{2J_*}{c_{\text{loc}} \lambda_{\min}(H)}}.$$

For the same rowwise gap proxy

$$g_a(\alpha, \beta) = \sqrt{e^{\alpha S_{\text{part},a}} + e^{\beta S_{\text{schur},a}} - 1},$$

the corresponding audit gives a local Lipschitz bound $\|\nabla g_a\| \leq L_a^{(\text{box})}$, so

$$g_a(\alpha, \beta) \geq g_{0,a} - L_a^{(\text{box})} \rho_*, \quad g_{0,a} := \sqrt{K_{11,a} + G_a - 1}.$$

Taking $J_* = 10^{-6}$, the audited box has $L_{\max}^{(\text{box})} = 5.11 \times 10^{-1}$ and therefore yields the uniform local operator-family floor

$$g_a \geq 1.3375.$$

Direct enumeration of the admissible scan points with $J_{\text{parent}} \leq 10^{-6}$ is sharper still and gives

$$g_a \geq 1.3537$$

throughout the same local family. On the overlap with the current open-system knot set, the conservative analytic floor $g_a \geq 1.3375$ already implies

$$f(r_a/g_a) - f(r_a) \geq 1.19 \times 10^{-2},$$

so the same reference defect $\delta_{q,a} \leq 10^{-3}$ would force

$$1 - \nu_{1,a} \leq 8.37 \times 10^{-2}, \quad \varepsilon_a(\Delta E_a) \leq 5.28 \times 10^{-2}.$$

This is still not a proof of the exact complement spectrum, because it remains confined to the audited projected family. But it upgrades the previous proxy statement into a genuine local operator-side lower bound: within the canonical projected EYMH basin, a small exact parent mismatch already forbids the second-mode scale from collapsing back to the leading pole.

Corollary 1 also gives a minimal lifting lemma from this local family floor to the exact Schur-reduced operator. Let

$$\sigma_a := -\log \det(I - A_a),$$

with the same positive semidefinite complement contraction

$$A_a = H_{PP,a}^{-1/2} H_{PQ,a} H_{QQ,a}^{-1} H_{QP,a} H_{PP,a}^{-1/2}.$$

If the eigenvalues of A_a lie in $[0, 1)$, then the eigenvalues of $I - A_a$ also lie in $(0, 1]$, and therefore

$$I - A_a \succeq \det(I - A_a) I = e^{-\sigma_a} I.$$

Hence the exact Schur-reduced operator obeys the Loewner bound

$$H_{\text{eff},a} = H_{PP,a}^{1/2} (I - A_a) H_{PP,a}^{1/2} \succeq e^{-\sigma_a} H_{PP,a},$$

so if the projected family supplies any lower gap floor $g_a^{(\text{fam})}$, then the exact effective gap satisfies

$$g_a^{(\text{exact})} \geq e^{-\sigma_a/2} g_a^{(\text{fam})}.$$

Using the conservative audited family floor $g_a^{(\text{fam})} \geq 1.3375$ from the local operator-family bound above, the current overlap with the open-system knot set yields the conditional exact floors

$$g_a^{(\text{exact})} \geq 1.3368 \quad (\sigma_a \leq 10^{-3}), \quad g_a^{(\text{exact})} \geq 1.3308 \quad (\sigma_a \leq 10^{-2}).$$

The corresponding separator band remains

$$f(r_a/g_a^{(\text{exact})}) - f(r_a) \geq 1.19 \times 10^{-2} \quad (\sigma_a \leq 10^{-3}),$$

up to the displayed precision, and only weakens to

$$f(r_a/g_a^{(\text{exact})}) - f(r_a) \geq 1.18 \times 10^{-2} \quad (\sigma_a \leq 10^{-2}).$$

With the same reference defect $\delta_{q,a} \leq 10^{-3}$ this still implies

$$\varepsilon_a(\Delta E_a) \leq 5.29 \times 10^{-2} \quad (\sigma_a \leq 10^{-3}), \quad \varepsilon_a(\Delta E_a) \leq 5.41 \times 10^{-2} \quad (\sigma_a \leq 10^{-2}).$$

In this form the remaining theorem target becomes sharper still: it is enough to prove a small complement log-det loss (or, equivalently, a small $\|A_a\|$) on the audited localized window, because the already-established family gap floor then lifts automatically to the exact Schur-reduced operator.

There is also an auxiliary parent-side route to that loss parameter if one assumes that the complement contraction is itself first-channel dominant on the same local window. In the rank-one model

$$A_a \approx a_a u_a u_a^\top, \quad 0 \leq a_a < 1,$$

one has

$$\|A_a\| = a_a, \quad \sigma_a = -\log(1 - a_a).$$

The exact projected kernel family already carries the corresponding scalar contraction candidate, because

$$\det K_a(\alpha, \beta, \lambda) = K_{11,a} K_{22,a} (1 - \lambda^2 \xi_a(\alpha, \beta)), \quad \xi_a(\alpha, \beta) := \frac{(K_{11,a} - 1)(K_{22,a} - 1)}{K_{11,a} K_{22,a}},$$

so the minimal dominant-channel identification is

$$a_a^{(1)} := \lambda^2 \xi_a.$$

The local coercivity bound above already gives

$$|\lambda| \leq \left(\frac{J_*}{c_{\text{loc}} C_4} \right)^{1/4},$$

which yields $|\lambda| \leq 2.02 \times 10^{-1}$ at $J_* = 10^{-6}$. On the overlap with the current open-system knot set, the exact canonical factor satisfies

$$\max \xi_a = 1.3324 \times 10^{-1},$$

while the best raw off-diagonal envelope built directly from the parented EYMH witnesses,

$$\xi_a^{(\text{raw})} := 0.127472 \frac{hk_{\text{abs,offdiag}}}{\text{diag}} + 0.435064 \frac{\text{action}_{\text{abs,offdiag}}}{\text{diag}},$$

has overlap p95 residual 4.37×10^{-2} and therefore gives the conservative envelope

$$\xi_a \leq 1.797 \times 10^{-1}.$$

Hence the same local objective tolerance already implies the rank-one dominant-channel candidates

$$a_a^{(1)} \leq 5.46 \times 10^{-3}, \quad \sigma_a^{(1)} \leq 5.47 \times 10^{-3}$$

from the exact canonical ξ_a , and

$$a_a^{(1)} \leq 7.36 \times 10^{-3}, \quad \sigma_a^{(1)} \leq 7.39 \times 10^{-3}$$

from the raw-witness envelope. Feeding the latter into the conditional lift still leaves

$$g_a^{(\text{exact})} \geq 1.3326, \quad \varepsilon_a(\Delta E_a) \leq 5.37 \times 10^{-2}$$

at the same reference defect $\delta_{q,a} \leq 10^{-3}$. This is not yet a full operator theorem, because the rank-one dominant-channel hypothesis itself remains to be justified, but it shows that once the complement is reduced to one effective contraction channel, the audited parent-side off-diagonal witnesses already drive the remaining log-det loss into the small- 10^{-3} regime rather than leaving it order one. We therefore keep this line as an auxiliary source model rather than as the preferred theorem route.

What the same audit now adds is that, once any such bound on σ_a has been obtained, its consequences for the exact complement contraction are genuine operator inequalities rather than rank-one bookkeeping identities. If the eigenvalues of A_a are $\alpha_{1,a}, \alpha_{2,a} \in [0, 1)$, then

$$\sigma_a = \sum_{i=1}^2 -\log(1 - \alpha_{i,a}) \geq \sum_{i=1}^2 \alpha_{i,a} = \text{tr } A_a,$$

because $-\log(1 - x) \geq x$ on $[0, 1)$. Likewise,

$$e^{-\sigma_a} = \prod_{i=1}^2 (1 - \alpha_{i,a}) \leq 1 - \max_i \alpha_{i,a} = 1 - \|A_a\|,$$

so

$$\|A_a\| \leq 1 - e^{-\sigma_a}.$$

Thus the same conditional σ_a budget immediately implies the true operator bounds

$$\text{tr } A_a \leq \sigma_a, \quad \|A_a\| \leq 1 - e^{-\sigma_a}, \quad H_{\text{eff},a} \succeq e^{-\sigma_a} H_{PP,a},$$

and Corollary 1 then gives

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq \frac{2(1 - e^{-\sigma_a})}{e^{-\sigma_a}}.$$

At the same local objective tolerance $J_* = 10^{-6}$, the exact canonical route $\sigma_a \leq 5.47 \times 10^{-3}$ therefore yields

$$\text{tr } A_a \leq 5.47 \times 10^{-3}, \quad \|A_a\| \leq 5.46 \times 10^{-3}, \quad H_{\text{eff},a} \succeq 0.9945 H_{PP,a},$$

together with

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.10 \times 10^{-2}.$$

The conservative raw off-diagonal envelope $\sigma_a \leq 7.39 \times 10^{-3}$ still gives

$$\text{tr } A_a \leq 7.39 \times 10^{-3}, \quad \|A_a\| \leq 7.36 \times 10^{-3}, \quad H_{\text{eff},a} \succeq 0.9926 H_{PP,a},$$

with

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.48 \times 10^{-2}.$$

So the remaining gap is now narrower than before: the open theorem target is no longer to guess whether the complement is small, but to prove a parent-side hypothesis that supplies a comparably small loss budget. The determinant-deficit reformulation below is the preferred theorem route; the rank-one dominant-channel line above should be read only as an auxiliary source model.

Proposition O4 (Determinant-deficit control). One can remove the explicit rank-one step from the operator consequences once a determinant-deficit budget has been identified. For any 2×2 positive contraction A_a with eigenvalues $\alpha_{1,a}, \alpha_{2,a} \in [0, 1)$, define

$$u_a := 1 - \det(I - A_a) = \alpha_{1,a} + \alpha_{2,a} - \alpha_{1,a}\alpha_{2,a}.$$

Then

$$u_a - \alpha_{1,a} = \alpha_{2,a}(1 - \alpha_{1,a}) \geq 0, \quad u_a - \alpha_{2,a} = \alpha_{1,a}(1 - \alpha_{2,a}) \geq 0,$$

so

$$\|A_a\| = \max(\alpha_{1,a}, \alpha_{2,a}) \leq u_a.$$

Writing $s_a := \alpha_{1,a} + \alpha_{2,a} = \text{tr } A_a$ and using $\alpha_{1,a}\alpha_{2,a} \leq s_a^2/4$ gives

$$u_a = s_a - \alpha_{1,a}\alpha_{2,a} \geq s_a - \frac{s_a^2}{4},$$

hence

$$\text{tr } A_a \leq 2(1 - \sqrt{1 - u_a}).$$

Therefore any parent-side bound $u_a \leq u_* < 1$ already implies

$$\|A_a\| \leq u_*, \quad \text{tr } A_a \leq 2(1 - \sqrt{1 - u_*}), \quad H_{\text{eff},a} \succeq (1 - u_*)H_{PP,a},$$

and Corollary 1 becomes

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq \frac{2u_*}{1 - u_*}.$$

So the operator theorem target can be stated without any single-channel hypothesis: it is enough to control the determinant deficit u_a on the audited window.

Proposition O5 (Continuum neighborhood determinant-deficit bound). The exact projected kernel family already carries a canonical determinant-deficit candidate through its local factor

$$1 - \lambda^2 \xi_a, \quad u_a(\alpha, \beta, \lambda) := \lambda^2 \xi_a(\alpha, \beta).$$

Within the canonical local box

$$|\alpha - 1| \leq 0.1, \quad |\beta - 1| \leq 0.1, \quad |\lambda| \leq 0.2,$$

the same coercive audit used above gives

$$\frac{1}{2} \delta \theta^\top H \delta \theta + C_4 \lambda^4 \leq J_*/c_{\text{loc}} \quad \text{whenever } J_{\text{parent}} \leq J_*.$$

Hence the admissible family is reduced to a compact coercive neighborhood, and the determinant-deficit source is obtained by maximizing the smooth scalar $u_a(\alpha, \beta, \lambda)$ on that set. The audited optimum at $J_* = 10^{-6}$ occurs near $(\alpha, \beta, \lambda) \approx (1.01, 0.992, 0.2)$ and yields

$$u_a \leq 5.34 \times 10^{-3}$$

so Proposition O4 immediately gives

$$\text{tr } A_a \leq 5.35 \times 10^{-3}, \quad \|A_a\| \leq 5.34 \times 10^{-3}, \quad H_{\text{eff},a} \succeq 0.9947 H_{PP,a},$$

and

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.08 \times 10^{-2}.$$

Proof sketch. The coercive inequality controls both the quadratic (α, β) displacement and the quartic mixing direction λ , so the sublevel set $J_{\text{parent}} \leq J_*$ stays inside a small neighborhood where the exact projected factor $1 - \lambda^2 \xi_a$ remains valid and u_a is uniformly bounded. The proposition is therefore reduced to one constrained scalar optimization, after which the operator consequences follow monotonically from Proposition O4. Thus the main proof-side task is no longer to search for an unconstrained small complement, but to promote this continuum neighborhood budget for u_a from the audited projected family to a genuine parent-side operator statement.

Corollary O1 (Neighborhood budget on the on-shell window). If the continuum neighborhood budget of Proposition O5 is inherited by the exact Schur-reduced operator on the audited on-shell window, then at $J_* = 10^{-6}$ one has

$$g_a^{(\text{exact})} \geq 1.3339, \quad 1 - \nu_{1,a} \leq 8.43 \times 10^{-2}, \quad \varepsilon_a(\Delta E_a) \leq 5.35 \times 10^{-2}$$

under the same reference defect target $\delta_{q,a} \leq 10^{-3}$ used in Theorem Target O1. In other words, Proposition O5 already compresses the remaining first-mode statement to a single parent-side task: justify that the audited continuum neighborhood budget survives the lift from the projected family to the exact Schur complement.

Benchmark O1 (Dense-family sharpness check). The dense exact-family scan should now be read as a sharpness benchmark rather than as the theorem input. Resolving the same family on

$$\alpha, \beta \in [0.95, 1.05], \quad \lambda \in [-0.2, 0.2],$$

with step ‘0.005’ and imposing exact admissibility $J_{\text{parent}} \leq 10^{-6}$ shows that the extremal ridge sits near

$$(\alpha, \beta, \lambda) \approx (1.03, 0.99, \pm 0.2),$$

giving

$$u_a \leq 5.42 \times 10^{-3}, \quad \text{tr } A_a \leq 5.42 \times 10^{-3}, \quad \|A_a\| \leq 5.42 \times 10^{-3},$$

with

$$H_{\text{eff},a} \succeq 0.9946 H_{PP,a}, \quad |\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.09 \times 10^{-2}.$$

TABLE XX. Geometry-informed open-system demonstrator at $D = \{6, 12, 18\}$. Source CSV: `output/chi_open_system/chi_open_system_geometry_D6-12-18.csv`.

D	γ_ϕ	γ_{mix}	C_{max}	$\chi_{\text{eff}}^{(\text{proxy})}$	$\chi_N^{(LR)}$	ratio
6	0.656901	0.227474	0.077658	6.24×10^{-5}	4.02×10^{-4}	0.155
12	0.744389	0.179054	0.063765	2.82×10^{-5}	2.21×10^{-4}	0.128
18	0.805854	0.136694	0.051381	2.19×10^{-5}	2.13×10^{-4}	0.103

This is the tightest local benchmark currently available, but it is only about one percent above Proposition O5. Even the conservative raw off-diagonal envelope remains moderate on the same dense admissible window, giving only

$$u_a \leq 7.19 \times 10^{-3}, \quad |\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.45 \times 10^{-2}.$$

So the right reading is now stable: the dense exact-family scan sets the benchmark, while the continuum neighborhood bound is the theorem candidate to push beyond the scanned family.

For the reviewer, the microscopic remainder can therefore be summarized as a three-step narrowing chain rather than as a generic open-system ansatz:

$$\kappa_{\text{env}} \implies P\Delta VP \in \text{span}\{I, \sigma_z, \sigma_x\} \implies S_{zz,a}(0), S_{xx,a}(\Delta E_a) = S_a(0), S_a(\Delta E_a).$$

The absolute-normalization audit isolates κ_{env} as the unique exact amplitude modulus on the audited knot set; the localized two-level projection reduces the leading interaction block to the span of I, σ_z, σ_x , so that the nontrivial projected interaction is at most two-channel; and the micro bridge shows that the dephasing and mixing kernels are on-shell evaluations of the same minimal response family rather than unrelated bath shapes. Accordingly, the remaining EYMH task is no longer to justify the local projected open-system form itself, but to identify the parent origin of the scalar modulus, the projected two-level jet, and the single response family behind the two on-shell channels.

Across all sensitivity cases in Table XXI, the generation metrics remain unchanged on this grid ($f(\mathcal{R}_3 > 0.90) = 0.9269$, $f(\mathcal{R}_3 > 0.95) = 0.2167$, $f(N_{\text{win}} > 3) \simeq 2.78 \times 10^{-4}$), while the $H \rightarrow \mu\mu$ acceptance fraction remains fixed at $f(\chi_{\mu\mu}^2 < 4) = 0.1500$ in this setup. These values are model-dependent because the Lindblad structure is still an effective environment ansatz and not yet a microscopic EYMH bath derivation; therefore, the open-system microscopic branch is treated as scan-ready diagnostic only and is not propagated into baseline reported figures. The explicit gate file `output/chi_open_system/open_system_micro_baseline_candidate.csv` records that this diagnostic passes the multi-anchor + holdout candidate gates (including holdout RMSE/max-abs thresholds), but we keep it outside baseline until the bath-side microscopic derivation is completed.

TABLE XXI. Microscopic open-system sensitivity scan with profiled $(\gamma_\phi, \gamma_{\text{mix}})$ on the 60×60 map. Source CSV: `paper/chi_open_system_micro_sensitivity.csv`.

case	ratio min	ratio max	ratio mean
baseline localized	1.000	1.000	1.000
open-micro base (1, 1)	0.017	0.327	0.138
$\phi \times 0.5$, mix $\times 1$	0.017	0.327	0.138
$\phi \times 2$, mix $\times 1$	0.017	0.327	0.138
$\phi \times 1$, mix $\times 0.5$	0.009	0.220	0.082
$\phi \times 1$, mix $\times 2$	0.034	0.412	0.208

Map-Level Migration Snapshot: baseline vs first-principles switches
R3 drift (absolute): legacy_cardy=-0.0500, open_system_micro=+0.0000




FIG. 15. Map-level migration snapshot across three scenarios. Left to right: $f(\mathcal{R}_3 > 0.90)$, $f(\chi^2_{\mu\mu} < 4)$, and $f(N_{\text{win}} > 3)$.

5. Cross-Module Migration Snapshot

To compare the two main migration directions on the same map-level metrics under the updated baseline, we aggregate three reviewer-facing scenarios: the release baseline, a legacy Cardy comparator, and the open-system microscopic diagnostic. Exact mode flags and file-level provenance are retained in the artifact registry; Fig. 15 shows the resulting summary.

6. Reference-Normalization Sensitivity Audit

Because Eq. (56) is currently reported as a relative scan-level observable normalized to a reference point (D_0, η_0) , we explicitly audit the sensitivity of this reporting convention with a registered selector-and-rescan workflow:

1. build a fixed-reference baseline map at $(D_0, \eta_0) = (10, 1)$ on the same 60×60 grid;
2. extract three candidate reference points: fixed-grid point, χ^2 -best [Eq. (65)], and robust-center [Eq. (66)];

TABLE XXII. Cross-module migration comparison. Scenario names are reviewer-facing summaries; the source registry retains exact run-mode flags.

scenario	$f(\mathcal{R}_3 > 0.90)$	$\Delta f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	$f(N_{\text{win}} > 3)$
release baseline	0.927	0.0000	0.1500	0.0003
legacy Cardy comparator	0.877	-0.0500	0.1486	0.0003
open-system diagnostic	0.927	0.0000	0.1500	0.0003

TABLE XXIII. Reference-normalization candidates and re-normalized map metrics. Source CSV: `paper/reference_anchor_candidates.csv`; selector JSON: `paper/reference_anchor_choice.json` (historical artifact filenames retained for reproducibility).

mode	(D_0, η_0)	χ^2 at selector point	$f(\chi_{\mu\mu}^2 < 4)$ (re-normalized)	best χ^2 (re-normalized)
fixed-grid	(9.97, 0.97)	8.83×10^{-1}	0.1500	1.18×10^{-3}
χ^2 -best	(4.00, 0.264)	2.21×10^{-5}	0.1333	1.90×10^{-1}
robust-center	(9.42, 3.87)	1.39×10^{-3}	0.1381	4.74×10^{-3}

3. re-normalize Eq. (56) with each candidate and recompute map summaries.

Table XXIII summarizes the selector outputs, and Fig. 16 shows the full grid-level sensitivity maps. The source filenames retain the earlier “anchor” spelling for reproducibility; the physics interpretation is reference-normalization sensitivity rather than Higgs-sector anchoring.

Across the tested reference-normalization grid $D_0 \in \{7, \dots, 13\}$ and $\eta_0 \in \{0.6, 0.9, \dots, 2.4\}$, we find

$$f(\chi_{\mu\mu}^2 < 4) \in [0.0333, 0.2411], \quad \chi_{\min}^2 \in [1.52 \times 10^{-5}, 7.47 \times 10^{-1}]. \quad (\text{E14})$$

The baseline fixed-reference value used in main figures is $f(\chi_{\mu\mu}^2 < 4) = 0.1500$. For reporting discipline, we keep this fixed normalization point in headline plots and treat dynamic re-normalizations as sensitivity diagnostics rather than post-hoc optimization. The physical normalization story should instead be read from the absolute and loop-improved parent-action comparators in Sec. VII, together with the Wilson-diagonal ratios in Eq. (67).

7. Dephasing-Based First-Principles Candidate for t_{coh}

To close the finite-time module with the same action-derived spectral chain, we test the dephasing candidate

$$\Delta\omega_{12}(D) \equiv \omega_2(D) - \omega_1(D), \quad t_{\text{coh}}^{(\text{deph})}(D) \equiv \frac{\pi}{\Delta\omega_{12}(D)}, \quad (\text{E15})$$




FIG. 16. Systematic sensitivity over reference-normalization points (D_0, η_0) under the EFT/Wilson-matched UV+LL-RG map. Left: $f(\chi_{\mu\mu}^2 < 4)$; right: best χ^2 . Overlaid markers indicate baseline fixed, χ^2 -best, and robust-center selectors. Source files: `paper/reference_anchor_sensitivity.csv`, `paper/reference_anchor_sensitivity.png`.

where (ω_1, ω_2) are extracted from the bound-state convention $E_n = \omega_n^2 - m_0^2 < 0$ in the same 1D on-axis operator chain as Appendix B. Detailed coarse/mid/fine values and implementation provenance are retained in the artifact registry. On the benchmark set $D = \{6, 12, 18\}$, the maximum non-fine relative deviation is

$$\max \left(\frac{|t_{\text{coh}} - t_{\text{coh}}^{(\text{fine})}|}{|t_{\text{coh}}^{(\text{fine})}|} \right) = 8.86 \times 10^{-3}, \quad (\text{E16})$$

which is numerically stable at the sub- 10^{-2} level for this benchmark set. The profile-impact scan compares constant $t_{\text{coh}} = 1$, raw Eq. (E15), and the capped variant $\min(t_{\text{coh}}^{(\text{deph})}, 10^4)$. All three leave the headline map summaries unchanged in this diagnostic chain: $f(\mathcal{R}_3 > 0.90) = 0.983$, $f(\mathcal{R}_3 > 0.95) = 0.217$, $f(\chi_{\mu\mu}^2 < 4) = 0.0833$, and best $\chi^2 = 9.62 \times 10^{-2}$. Therefore, in this manuscript, Eq. (E15) is reported as a first-principles candidate benchmark only and is not propagated into the baseline reported figures.

8. Splitting-Prefactor First-Principles Candidate for η

To test a geometry-closed replacement direction for the overlap amplitude in Eq. (24), we use the same action-derived quantities $(\Delta E_{12}, S_1)$ and define the splitting prefactor

$$\Delta E_{12}(D) \approx A_{12}(D) e^{-S_1(D)}, \quad A_{12}(D) \equiv \Delta E_{12}(D) e^{S_1(D)}. \quad (\text{E17})$$

With a reference point $D_{\text{ref}} = 12$, we define two dimensionless candidates:

$$\eta_{\text{amp}}(D) \equiv \frac{A_{12}(D)}{A_{12}(D_{\text{ref}})}, \quad \eta_{\text{prob}}(D) \equiv \eta_{\text{amp}}(D)^2. \quad (\text{E18})$$

Detailed coarse/mid/fine values and implementation provenance are retained in the artifact registry. On the same $D = \{6, 12, 18\}$ benchmark set, we obtain

$$\max\left(\frac{|\eta_{\text{amp}} - \eta_{\text{amp}}^{(\text{fine})}|}{|\eta_{\text{amp}}^{(\text{fine})}|}\right) = 1.08 \times 10^{-2}, \quad \max\left(\frac{|\eta_{\text{prob}} - \eta_{\text{prob}}^{(\text{fine})}|}{|\eta_{\text{prob}}^{(\text{fine})}|}\right) = 2.18 \times 10^{-2}, \quad (\text{E19})$$

and the fine full-scan profile over $D = 4, \dots, 20$ gives $\eta_{\text{amp}} \in [0.9996, 1.1695]$ and $\eta_{\text{prob}} \in [0.9992, 1.3677]$. Profile-scaled and fully closed implementations leave the headline map-level summaries unchanged in this diagnostic chain; the residual movement is only in sub-table-precision best- χ^2 shifts. Therefore, Eq. (E18) is reported as a first-principles candidate benchmark and is not propagated into baseline reported figures.

9. Action-Derived First-Principles Candidate for Barrier-Leakage Channel Scaling

To test a geometry-closed replacement direction for the channel normalization in Eq. (32), we keep the same 1D action-derived chain and define channel-resolved barrier actions

$$U_\ell(z; D) \equiv U(z; D) + \frac{\ell(\ell+1)}{z^2 + \varepsilon^2}, \quad S_{N,\ell}(D) \equiv \int_{z_-}^{z_+} dz \sqrt{U_\ell(z; D) - E_N(D)}, \quad (\text{E20})$$

with $E_N = \omega_N^2 - m_0^2$ in bound-state convention. We then define geometry-driven channel rates

$$\Gamma_{N,\ell}^{(\text{geo})}(D) \equiv \omega_N(D) e^{-2S_{N,\ell}(D)}, \quad (\text{E21})$$

and infer effective channel normalizations by matching to Eq. (32):

$$A_\ell^{(\text{fp})}(D; N) \equiv \frac{\Gamma_{N,\ell}^{(\text{geo})}(D)}{\omega_N(D) (\omega_N(D) M_*)^{4\ell+4}}. \quad (\text{E22})$$

Using $(D_{\text{ref}}, N_{\text{ref}}) = (12, 2)$, we define profile factors

$$\tilde{A}_\ell(D) \equiv \frac{A_\ell^{(\text{fp})}(D; N_{\text{ref}})}{A_\ell^{(\text{fp})}(D_{\text{ref}}; N_{\text{ref}})}. \quad (\text{E23})$$

Detailed coarse/mid/fine values and implementation provenance are retained in the artifact registry. On the $D = \{6, 12, 18\}$ benchmark set, the maximum non-fine relative deviations are

$$\max\left(\frac{|\tilde{A}_1 - \tilde{A}_1^{(\text{fine})}|}{|\tilde{A}_1^{(\text{fine})}|}\right) = 3.22 \times 10^{-2}, \quad \max\left(\frac{|\tilde{A}_2 - \tilde{A}_2^{(\text{fine})}|}{|\tilde{A}_2^{(\text{fine})}|}\right) = 3.78 \times 10^{-2}. \quad (\text{E24})$$

The fine full-scan profile over $D = 4, \dots, 20$ yields $\tilde{A}_1 \in [8.54 \times 10^{-6}, 4.10 \times 10^5]$ and $\tilde{A}_2 \in [3.95 \times 10^{-6}, 4.92 \times 10^6]$, indicating a very stiff D -dependence in this preliminary closure.

As a strict-upgrade check, we also connected a minimal complex-absorbing-potential solver to the same 1D action-derived operator chain and scanned benchmark windows in the absorber

strength η and absorber onset z_{cap} . This pilot confirms that the non-Hermitian extraction pipeline is operational, but the benchmark mode remains in an absorber-linear or mixed regime rather than a plateau regime. Accordingly, the CAP outputs are retained as infrastructure diagnostics only; the reported baseline continues to use $\Gamma_{N,\ell}^{(\text{geo})}$ as the leading static surrogate while reserving complex-pole width extraction for a later dedicated release.

At current scan resolution, Table XXIV shows that the t_{coh} and η candidates are stable diagnostics whose map-level effects are below headline precision, while replacing the legacy constant- A_ℓ reference by the action-derived profile induces non-negligible shifts in $\mathcal{R}_3$ coverage. In the updated global scans of this manuscript, Eq. (E23) is used in baseline mode, and $(A_1, A_2) = (1, 1)$ is retained only as a comparator.

TABLE XXIV. Compressed late-stage diagnostic summary for the three first-principles profile candidates. Detailed coarse/mid/fine tables and exact run artifacts are retained in the artifact registry.

Candidate	Local benchmark	Full/profile range	Map-level impact
$t_{\text{coh}}^{(\text{deph})}$	non-fine drift 8.86×10^{-3} vs fine	raw and capped profiles tested	$f(\mathcal{R}_3 > 0.90) = 0.983$, $f(\chi_{\mu\mu}^2 < 4) = 0.0833$, best $\chi^2 = 9.62 \times 10^{-2}$ unchanged
$\eta_{\text{amp}}, \eta_{\text{prob}}$	max drifts 1.08×10^{-2} and 2.18×10^{-2}	$\eta_{\text{amp}} \in [0.9996, 1.1695]$, $\eta_{\text{prob}} \in [0.9992, 1.3677]$	profile-scaled and fully closed variants leave headline map summaries unchanged
$\tilde{A}_\ell(D)$	max drifts 3.22×10^{-2} and 3.78×10^{-2}	$\tilde{A}_1 \in$ $[8.54 \times 10^{-6}, 4.10 \times 10^5]$, $\tilde{A}_2 \in$ $[3.95 \times 10^{-6}, 4.92 \times 10^6]$	action profile shifts $f(\mathcal{R}_3 > 0.90)$ from 0.9833 to 0.9269; $H \rightarrow \mu\mu$ acceptance stays 0.0833

Appendix F: Ansatz Ablation Study

To demonstrate that the three-generation structure is robust and not an artifact of a particular ansatz choice, we compare the baseline EFT-operator-normalized $B_N(D)$ profile (this work) against the alternative “inverse Yukawa” ansatz explored in earlier versions.

1. Inverse Yukawa Ansatz (Ablation A)

The inverse Yukawa ansatz sets $B_N^{(\text{inv})} \propto 1/\sum_{f \in \text{Gen } N} y_f$, yielding:

$$B_1^{(\text{inv})} \approx 10^5, \quad B_2^{(\text{inv})} \approx 125, \quad B_3^{(\text{inv})} = 1. \quad (\text{F1})$$

This massive hierarchy ($B_1 \gg B_2 \gg B_3$) directly encodes the answer by over-enhancing layer 1. While it can produce $\mathcal{R}_3 > 95\%$, the mechanism is less transparent: the visibility factor, rather than the entropic/kinetic competition, drives the result.

2. Comparison

- **Baseline (EFT-operator-normalized profile):** $B_N(D)$ remains smooth and bounded across the scanned range, with shape inherited from tracked microcanonical overlap extraction and operator normalization. The three-generation structure emerges from the g_N – Γ_N competition modulated by these profile differences.
- **Ablation A (Inverse Yukawa):** $B_1 \gg B_2 \gg B_3$. Layer 1 is artificially boosted; the selection mechanism is less falsifiable.

We recommend the baseline as the primary closure due to its greater transparency and falsifiability.

-
- [1] N. Gresnigt, Electroweak structure and three fermion generations in clifford algebra with s_3 family symmetry (2026), arXiv:2601.07857 [hep-ph].
 - [2] A. Dabholkar, S. Murthy, and D. Zagier, Quantum black holes, wall crossing, and mock modular forms (2012), arXiv:1208.4074 [hep-th].
 - [3] J. Maldacena and L. Susskind, Fortschritte der Physik **61**, 781 (2013), arXiv:1306.0533 [hep-th].
 - [4] J. L. Cardy, Nucl. Phys. B **270**, 186 (1986).
 - [5] J. D. Bekenstein, Phys. Rev. D **7**, 2333 (1973).
 - [6] G. 't Hooft, Dimensional reduction in quantum gravity (1993), arXiv:gr-qc/9310026 [gr-qc].
 - [7] S. Detweiler, Phys. Rev. D **22**, 2323 (1980).
 - [8] S. Navas *et al.*, Phys. Rev. D **110**, 030001 (2024), Particle Data Group.
 - [9] ATLAS Collaboration, Phys. Rev. Lett. **135**, 231802 (2025), arXiv:2507.03595 [hep-ex].
 - [10] V. Gorini, A. Kossakowski, and E. C. G. Sudarshan, J. Math. Phys. **17**, 821 (1976).
 - [11] G. Lindblad, Commun. Math. Phys. **48**, 119 (1976).
 - [12] J. C. F. Thaler and S. Trifinopoulos, Phys. Rev. D **111**, 056021 (2025), arXiv:2410.23343 [hep-ph].