

The Projection Spectral Layer Theory (PSLT): A Rank-2 Computable Closure for the Three-Generation Structure and Higgs Signal Strength

Bo-Yu Chen¹

¹*Independent Researcher*

(Dated: February 2, 2026)

We present a comprehensive, computable EFT-level “closed chain” for the Projection Spectral Layer Theory (PSLT) in which a self-consistent three-generation structure emerges on the baseline closure from competing spectral entropy, kinetics, and visibility. The framework is built on a minimal Einstein–Yang–Mills–Higgs (EYMH) effective action with non-minimal curvature coupling, where the background geometry induces a layer-indexed spectral scale ω_N . The theoretical closure is achieved by unifying three key modules: (i) a **Micro-degeneracy module** (g_N) implemented through a bounded two-dimensional phase-space profile (microcanonical low- N window + energy-gap tail prescription, phase-space normalized at $N = 3$); (ii) a **Rank-2 Computable Kinetic module** (Γ_N) defined as the largest positive eigenvalue of a 2×2 growth matrix with explicit WKB tunneling suppression; and (iii) a **Visibility module** (B_N) implemented as a fully normalized EFT-operator profile, where overlap-extracted flavor-layer couplings and mediator scales feed finite one-loop + LL-RG-normalized operator kernels, with neutral high- N saturation ($B_{N>3} = 1$) so that no extra visibility-side damping is added on top of the g_N regulator; a risk-weighted runtime-direct visibility parity branch is also available, but it is not interpreted as a strict all-direct closure. The resulting normalized layer probability $P_N(t_{\text{coh}})$ demonstrates, within this baseline closure, a stable “winner” phase diagram with clear dominance of layers $N = 1, 2, 3$ over a significant region of the geometric parameter space (D, η) . We verify the closure against two benchmarks: (1) internal stability of the three-generation hierarchy (Generation Ratio $> 90\%$ over 92.7% of the sampled (D, η) grid), and (2) a scan-level EFT/Wilson-matched compatibility test using the ATLAS Run-2+Run-3 combined $H \rightarrow \mu\mu$ signal-strength input ($\mu_{\mu\mu}^{\text{obs}} = 1.4 \pm 0.4$). In the current UV+LL-RG baseline, the illustrative acceptance band is about 15.0% of the scanned grid for $\chi^2 < 4$, providing a sharper falsifiable constraint on (D, η) .

CONTENTS

	2
II. Theoretical Framework	6
A. The Closed Chain Equation	6
B. Action-Derived Chain at a Glance	6
C. Geometric Foundation: Minimal EYMH + Curvature	6
D. Model-Chain Assumption Status	9
III. Geometry-to-Spectrum Worked Example	10
A. Two-Center Harmonic Conformal Factor	10
B. Action-Derived Effective Potential	11
C. Numerical Solution	11
D. Single-Track Results	12
E. Summary: Unified Geometry-to-Kinetics Pipeline	12
IV. Module 1: The Three Routes to Micro-Degeneracy g_N	14
A. Route A: Mock-Modular Counting	14
B. Route B: ER=EPR and Entanglement Entropy	14
C. Route C: Quantum Gravity Template	15
D. Unified Cardy-Controlled Envelope with High-N Regulator	15
E. Cardy Regime Validity and Finite-N Corrections	15
F. Assessment of a Phase-Space First-Principles Candidate	16
V. Module 2: Rank-2 Computable Kinetics Γ_N	16
A. Dual-Center Tunneling Suppression	16
B. Dimensionality and Units	17
C. Derivation of the Rank-2 Growth Matrix	17
D. Rank-2 Computable Closure	17
E. Physical Interpretation: Quasi-Bound Leakage and Width Proxy	19
VI. Module 3: EFT-Normalized Visibility B_N	38
A. High- N Neutral Saturation	39
B. Roadmap: Action-Derived Effective Yukawa Operator	41
VII. Verification Results	43
A. Configuration and Parameter Table	43
B. Three-Generation Phase Diagram	44
C. $H \rightarrow \mu\mu$ Signal Strength (EFT/Wilson-Matched Mapping, UV+LL-RG Baseline)	47

D. Projected Parent-Kernel Derivation	75
E. EFT/Wilson-Matched Extension to $H \rightarrow ee$ and $H \rightarrow \tau\tau$	87
VIII. Discussion and Conclusion	88
Acknowledgments	92
A. Action-Derived Chain Reproducibility SOP	92
1. Single Entry Point	92
2. Step Map and Deliverables	92
3. Acceptance Conditions	92
B. Action-Derived Effective Potential and WKB	93
1. Two-Center Harmonic Conformal Factor	93
2. Veff Derivation from KG Equation	94
3. Minimal Dirac Coupling in the Conformal Background	94
4. 1D On-Axis Reduction	95
5. 2D Axisymmetric Validation	95
6. WKB Action Convention	96
7. Splitting-Action Consistency Check	96
8. Action-Derived Numerical Results	96
9. Bound-State omega Convergence Benchmark	96
10. Fixed-dz Convergence	98
C. First-Principles Symmetry-Channel Test for χ	98
D. Localized-Channel First-Principles Extraction of χ	99
E. Preliminary First-Principles Replacement Candidate Checks	102
1. 1D Phase-Space Candidate for g_N	102
2. 2D Axisymmetric Benchmark for $g_N^{(\text{ps})}$	103
3. Baseline Cross-Check Under $g_{\text{mode}} = \text{fp-2d_full}$	104
4. Geometry-Informed and Microscopic Lindblad Scan-Ready Module for Open-System χ	105
5. Cross-Module Migration Snapshot	239
6. Reference-Normalization Sensitivity Audit	240
7. Dephasing-Based First-Principles Candidate for t_{coh}	241
8. Splitting-Prefactor First-Principles Candidate for η	242

9. Action-Derived First-Principles Candidate for Barrier-Leakage Channel Scaling	242
F. Ansatz Ablation Study	245
1. Inverse Yukawa Ansatz (Ablation A)	245
2. Comparison	245
References	246

I. INTRODUCTION

The existence of exactly three generations of fermions is one of the most persistent puzzles in the Standard Model (SM) of particle physics. While the SM accommodates three families through the CKM and PMNS mixing matrices, it offers no dynamical reason for why the number is three, nor why the mass hierarchy spans such a vast range. Traditional attempts to explain this structure often rely on complex discrete symmetries or string logic compactifications that, while elegant, can be difficult to falsify directly.

The Projection Spectral Layer Theory (PSLT) proposes a different paradigm: *Generation structure is a spectral consequence of the underlying projection geometry of spacetime.* In this framework, "generations" are not ad-hoc copies but rather distinct *spectral layers* ($N = 1, 2, 3, \dots$) of vacuum excitations, selected by a competition between microstate degeneracy (entropy) and geometric formation rates (kinetics). This structural viewpoint is complementary to recent Clifford-algebra work where three families emerge from an intrinsic S_3 organization rather than manual generation copying [1]; the emphasis here is different: a map-level computable closure with explicit scan-level falsifiability.

In this work, we consolidate the complete theoretical framework of PSLT into a single, falsifiable "closed chain". We upgrade earlier effective descriptions to a **computable** rank-2 dynamical system with an action-derived extraction subchain and an action-informed global scan chain, where $\chi_{LR}(D)$ and $\tilde{A}_\ell(D)$ are now injected on the full scan grid from direct action-derived profiles. The core deliverable is a map from geometric control parameters (D, η) —representing dual-center separation and overlap—to an experimentally observable layer probability P_N .

Our primary contributions are:

1. **Theoretical Motivation:** We present three convergent arguments for the micro-degeneracy g_N (Route A: Mock-modular, Route B: ER=EPR, Route C: QG Template), and implement a scan-ready baseline through the bounded two-dimensional phase-space profile with controlled high- N tail prescription.
2. **Rank-2 Computable Kinetics:** We replace heuristic growth rates with a rigorous eigenvalue problem for the formation matrix $\mathbf{M}_N$, incorporating numerical WKB tunneling suppression.
3. **Fully Normalized EFT Visibility:** We promote B_N to an EFT-operator-normalized baseline built from overlap-extracted flavor-layer couplings and mediator scales with finite one-loop + LL-RG normalization; overlap-only and Yukawa laws are retained only as

legacy comparators/fallbacks, while runtime-direct visibility extraction is release-qualified only as a risk-weighted parity check, not as a strict all-direct closure.

4. **Verification:** We demonstrate stable three-generation dominance at $\mathcal{R}_3 > 90\%$ and a constrained EFT/Wilson-matched-compatible band for LHC $H \rightarrow \mu\mu$.

II. THEORETICAL FRAMEWORK

A. The Closed Chain Equation

The central output of the PSLT framework is the normalized layer occupancy probability $P_N(t_{\text{coh}})$, defined by the competition between entropic weight (W_{entropy}) and kinetic formation (W_{kinetic}). The master equation is:

$$P_N(t_{\text{coh}}; D, \eta) = \frac{W_N(t_{\text{coh}})}{\sum_K W_K(t_{\text{coh}})}, \quad W_N = B_N \cdot g_N \cdot \left(1 - e^{-\Gamma_N(D, \eta)t_{\text{coh}}}\right). \quad (1)$$

This equation unifies the three critical modules:

- g_N : **Micro-degeneracy** (Entropy). The number of available microstates in layer N .
- Γ_N : **Dynamical Rate** (Kinetics). The rate at which these states can form/tunnel from the vacuum.
- B_N : **Visibility Factor** (Observation). The coupling strength of layer N to the observable sector.

Figure 1 illustrates the logical flow of the PSLT framework.

B. Action-Derived Chain at a Glance

For implementation-level clarity, Figure 2 summarizes the action-derived operator chain and its interfaces to scan-level diagnostics in one page. The executable, clone-level Standard Operating Procedure (SOP) is provided in Appendix A.

C. Geometric Foundation: Minimal EYMH + Curvature

We start from a minimal Einstein–Yang–Mills–Higgs (EYMH) effective action supplemented by a non-minimal curvature coupling $\xi R|\Phi|^2$:

$$S = \int d^4x \sqrt{-g} \left[\frac{M_{\text{Pl}}^2}{2} R - \frac{1}{4} F_{\mu\nu}^a F^{a\mu\nu} - (D_\mu \Phi)^\dagger (D^\mu \Phi) - \lambda(|\Phi|^2 - v^2)^2 - \xi R|\Phi|^2 \right]. \quad (2)$$




FIG. 1. Schematic overview of the PSLT framework. Geometric inputs (D, η) feed into three modules (micro-degeneracy g_N , kinetic rate Γ_N , visibility B_N), which combine in the master equation to yield the observable layer probabilities P_N and the three-generation ratio $\mathcal{R}_3$. In the baseline chain, high- N control is implemented by the bounded two-dimensional phase-space-tail prescription, while $B_{N>3}$ neutrally saturates to unity.

PSLT posits that the background geometry is a *projection* from a higher-dimensional manifold, inducing a locally conformally flat metric:

$$g_{\mu\nu}(x) = \Omega^2(x)\eta_{\mu\nu}. \quad (3)$$

Under conformal rescaling $\Phi \rightarrow \Omega^{-1}\tilde{\Phi}$, the scalar equation of motion reduces to a Schrödinger-like eigenproblem:

$$[-\nabla^2 + V_{\text{eff}}(x)]\psi_N = \omega_N^2\psi_N, \quad (4)$$

where the effective potential V_{eff} is determined by the conformal factor $\Omega(x)$ and the specific projection geometry (e.g., stereographic projection of dual centers).




FIG. 2. One-page action-derived chain overview used in this manuscript: geometry input $\rightarrow$ operator extraction $\rightarrow$ scan diagnostics $\rightarrow$ observable EFT/Wilson checks. This diagram is aligned with the repository reproducibility entrypoint and Appendix A.

The layer index N emerges naturally as the discrete principal-like label of this spectrum. In Section III, we derive the spectral scale $\mu(D)$ explicitly from a two-center geometry; the result supports a power-law scaling $\mu(D) \propto D^{-\gamma}$ with $\gamma \approx 0.1$ in the demonstrator regime. In the reported baseline, $\omega_N(D)$ is taken from the action-derived spectrum (or its audited interpolation), while the older hydrogenic $1 - 1/(2N^2)$ form is retained only as a **legacy comparator** for early demonstrator bookkeeping. We therefore write

$$\omega_N(D) \equiv \omega_N^{(\text{exact})}(D), \quad \mu(D) = \mu_0 D^{-\gamma}, \quad (5)$$

where μ_0 and γ are determined numerically from the geometry (Section III). The envelope $\mu(D)$ remains useful for compact analytic summaries, but the mainline spectral input is the numerically extracted action-derived $\omega_N^{(\text{exact})}(D)$ rather than a hydrogenic $1/N^2$ law. For transparency, we additionally provide a bound-state exact benchmark for (ω_1, ω_2) at $D = \{6, 12, 18\}$ with coarse/mid/fine convergence from the same action-derived operator chain (Appendix B 9, Table X). These benchmark values use $E = \omega^2 - m_0^2 < 0$ and should not be mixed with the generalized localized-extraction eigenvalues λ in Appendix D.

D. Model-Chain Assumption Status

For transparency, Table I separates action-derived components from EFT-level surrogates currently used in the global scan.

TABLE I. Status map of the PSLT closed chain in the present manuscript.

Component	Current Status	Used in Global Scan	Next Upgrade Path
$\Omega(\rho, z; D)$ two-center geometry	Projected-Poisson two-source Green solution	Yes	Derive the same source from an explicit higher-dimensional parent
V_{eff} , low modes, WKB action	Action-derived and numerically converged	Yes	Extend the same solver chain to the full (D, η, N) map
g_N phase-space profile	Bounded two-dimensional phase-space baseline; direct spectral parity and risk-weighted visibility parity checks available	Yes	Treat the strict all-direct broader-grid rule as parked at the current case-aware endpoint; reopen only for a genuinely new observable-side mechanism class
$\chi_N^{(LR)}(D)$ localized channel	Full D-grid localized profile from 2D fine extraction; release-production parity path available	Yes	Keep the current localized-direct parity path as the production support line; the longer-horizon upgrade is a fully localized extraction over (D, η, N) rather than another profile-blend retune
$\Gamma_{N,\ell}$ channel normalization (A_ℓ)	Action-derived $\tilde{A}_\ell(D)$ baseline with stabilized localized-direct support; legacy constant- A_ℓ kept only as comparator	Yes	Move to fully localized channel-resolved direct extraction over (D, η, N)
Open-system χ_{eff} (Lindblad)	Locally closed projected effective-action diagnostic: scalar modulus κ_{env} , projected σ_z/σ_x jet, and single-pole gap-locked response family are now explicit on the audited knot set	Diagnostic only	Promote the current completion to a parent-side Schur statement by proving first-mode dominance of the projected resolvent on $\omega \in \{0, \Delta E_a\}$ with $\Omega_{1,a} = \omega_{1,a}$
B_N visibility profile from EFT normalization	EFT-operator-normalized baseline; runtime-direct operator extraction promoted only in the risk-weighted parity branch	Yes	Keep the release-qualified profile-blend branch separate from the parked strict all-direct endpoint; the real remaining upgrade is fully localized channel-resolved extraction, not another local blend retune
$H \rightarrow \mu\mu$ mapping [Eq. (56)]	EFT/Wilson-matched baseline with UV-tree, finite-one-loop, and LL-RG audits	Yes	Extend the current closure to fuller UV-to-EFT matching from the EYMH action

III. GEOMETRY-TO-SPECTRUM WORKED EXAMPLE

This section demonstrates the complete derivation chain $\Omega(x; D) \rightarrow V_{\text{eff}}(x) \rightarrow \omega_N(D)$, establishing the geometry-to-spectrum correspondence as a computational reality rather than an ansatz.

A. Two-Center Harmonic Conformal Factor

The geometric origin of the dual-center conformal factor is formulated as a projected Poisson constraint on the static spatial slice:

$$\nabla^2 \Omega(\mathbf{x}) = -4\pi \sigma(\mathbf{x}), \quad \sigma(\mathbf{x}) \equiv a[\rho_\varepsilon(\mathbf{x} - \mathbf{x}_+) + \rho_\varepsilon(\mathbf{x} - \mathbf{x}_-)]. \quad (6)$$

Here $\mathbf{x}_\pm = \pm(D/2)\hat{z}$ are the two projected centers and ρ_ε is a normalized smeared source. With Green's function $G(\mathbf{x}) = (4\pi|\mathbf{x}|)^{-1}$ for $-\nabla^2$, the unique asymptotically flat solution ($\Omega \rightarrow 1$ at $|\mathbf{x}| \rightarrow \infty$) is

$$\Omega(\mathbf{x}) = 1 + \int d^3x' \frac{\sigma(\mathbf{x}')}{|\mathbf{x} - \mathbf{x}'|}. \quad (7)$$

Choosing the Plummer-type kernel

$$\rho_\varepsilon(\mathbf{r}) = \frac{3\varepsilon^2}{4\pi(|\mathbf{r}|^2 + \varepsilon^2)^{5/2}}, \quad \int d^3r \rho_\varepsilon(\mathbf{r}) = 1, \quad (8)$$

one obtains the closed form

$$\boxed{\Omega(\rho, z; D) = 1 + a \left(\frac{1}{r_+} + \frac{1}{r_-} \right), \quad r_\pm = \sqrt{\rho^2 + (z \mp D/2)^2 + \varepsilon^2}} \quad (9)$$

in cylindrical coordinates (ρ, z) (axisymmetric, $m = 0$). The corresponding source identity follows from

$$\nabla^2 \left(\frac{1}{\sqrt{r^2 + \varepsilon^2}} \right) = -\frac{3\varepsilon^2}{(r^2 + \varepsilon^2)^{5/2}}, \quad (10)$$

hence

$$\nabla^2 \Omega = -4\pi a [\rho_\varepsilon(x - x_+) + \rho_\varepsilon(x - x_-)], \quad \rho_\varepsilon(x) = \frac{3\varepsilon^2}{4\pi(|x|^2 + \varepsilon^2)^{5/2}}. \quad (11)$$

Equation (9) therefore is not an arbitrary fitting form: once Eq. (6), two-center source support, and asymptotic flatness are fixed, it is the corresponding Green-function solution.

B. Action-Derived Effective Potential

We derive V_{eff} directly from the Klein-Gordon equation in curved spacetime. Consider a scalar field with non-minimal curvature coupling:

$$(\square_g - m_0^2 - \xi R)\Phi = 0. \quad (12)$$

For time-harmonic modes $\Phi = e^{-i\omega t}\phi(\mathbf{x})$ in the background $g_{\mu\nu} = \Omega^2\eta_{\mu\nu}$, we perform the conformal field rescaling $\phi = \Omega^{-1}\psi$. After explicit calculation of $\square_g$ and cancellation of first-derivative terms (see Appendix B 2), the equation reduces to:

$$\boxed{[-\nabla^2 + V_{\text{eff}}(\mathbf{x})]\psi = \omega^2\psi} \quad (13)$$

with the **action-derived** effective potential:

$$\boxed{V_{\text{eff}} = m_0^2\Omega^2 + (1 - 6\xi)\Omega^{-1}\nabla^2\Omega} \quad (14)$$

This is the **only form** consistent with the action—no engineered barrier or box terms appear.

a. Key properties.

1. At $\xi = \xi_c = 1/6$ (conformal coupling in 4D), the derivative term vanishes: $V_{\text{eff}} \rightarrow m_0^2\Omega^2$.
2. As $r \rightarrow \infty$: $\Omega \rightarrow 1$, $\nabla^2\Omega \rightarrow 0$, so $V_{\text{eff}} \rightarrow m_0^2$ (continuum threshold).
3. Near each center: $\nabla^2\Omega < 0$ (smeared source), so for $\xi < 1/6$ the potential develops attractive wells.

We define the shifted potential $U \equiv V_{\text{eff}} - m_0^2$ so that $U \rightarrow 0$ at infinity. Bound states satisfy $E = \omega^2 - m_0^2 < 0$. Figure 3 shows the decomposition.

C. Numerical Solution

We solve the 2D axisymmetric eigenproblem using finite differences on a (ρ, z) grid with Eq. (14):

$$\left[-\frac{\partial^2}{\partial \rho^2} - \frac{1}{\rho} \frac{\partial}{\partial \rho} - \frac{\partial^2}{\partial z^2} + V_{\text{eff}}(\rho, z; D) \right] \psi_N = \omega_N^2 \psi_N. \quad (15)$$

Boundary conditions: $\partial_\rho \psi|_{\rho=0} = 0$ (axis regularity), $\psi|_{\text{boundary}} = 0$ (Dirichlet). Parameters: $a = 1.0$, $\varepsilon = 0.2$, $m_0 = 1.0$, $\xi = 0.0$. Grid: $(n_\rho, n_z) = (50, 500)$ with fine resolution $\Delta z \ll \varepsilon$ near the cores.

Critical numerical requirement: The grid spacing must satisfy $\Delta z \ll \varepsilon$ to resolve the smeared source structure. This is essential for capturing the deep negative wells in U .




FIG. 3. Action-derived potential decomposition at $D = 12$. Top: mass term $m_0^2(\Omega^2 - 1)$ and curvature term $(1 - 6\xi)\Omega^{-1}\nabla^2\Omega$. Bottom: total $U = V_{\text{eff}} - m_0^2$ showing the double-well structure. The 4 turning points confirm the tunneling geometry.

D. Single-Track Results

Table II shows the unified results. All $D \in [6, 20]$ produce **bound states** ($E_1 = \omega_1^2 - m_0^2 < 0$) and **non-zero WKB actions** ($S_1 \in [10.6, 22.8]$). The 4 turning points confirm the double-well tunneling geometry.

E. Summary: Unified Geometry-to-Kinetics Pipeline

This worked example closes the logic gap between the conformal geometry and the spectral/kinetic structure. **The same V_{eff} determines both the layer frequencies and the tunneling rates:**




FIG. 4. Single-track results from action-derived V_{eff} . Top-left: Bound state energy $E_1 < 0$. Top-right: Layer frequency ω_1 . Bottom-left: WKB tunneling action S_1 . Bottom-right: Tunneling probability $r_1 = e^{-2S_1}$. All quantities are computed from the same $V_{\text{eff}}(\rho, z; D)$.

TABLE II. Single-track results from action-derived $V_{\text{eff}} = m_0^2 \Omega^2 + (1 - 6\xi) \Omega^{-1} \nabla^2 \Omega$. All quantities are computed from the same potential. $E_1 = \omega_1^2 - m_0^2 < 0$ indicates bound states. Parameters: $a = 1$, $\varepsilon = 0.2$, $m_0 = 1$, $\xi = 0$.

D	E_1	ω_1	S_1	$r_1 = e^{-2S_1}$	tp	n_{bound}
6	-0.24	0.87	10.65	4.9×10^{-10}	4	2
8	-0.46	0.73	13.15	3.2×10^{-12}	4	3
10	-0.56	0.66	15.45	2.5×10^{-14}	4	3
12	-0.92	0.28	18.85	1.3×10^{-17}	4	1
14	-0.41	0.77	18.52	3.4×10^{-17}	4	2
16	-0.21	0.89	18.81	1.6×10^{-17}	4	3
18	-0.07	0.96	18.97	6.8×10^{-18}	4	3
20	-0.34	0.81	22.79	1.5×10^{-20}	4	1

1. **Conformal factor:** $\Omega(\rho, z; D)$ is explicitly specified (Eq. (9)).
2. **Effective potential:** V_{eff} is derived from Ω (Eq. (14)), with explicit D -dependence.
3. **Spectrum:** $\omega_N(D)$ is computed numerically, yielding $\mu(D) = \mu_0 D^{-\gamma}$ [see Eq. (5)].
4. **Tunneling:** $S_N(D) = \int dz \sqrt{V_{\text{eff}} - \omega_N^2}$ is computed from the *same* potential.
5. **Kinetic rates:** $r_N = e^{-2S_N}$ directly follows from the geometry.

This is a **single-track derivation** for the extraction subchain: there is no separate "toy potential" inside this subsection, and the quantities $\Omega \rightarrow V_{\text{eff}} \rightarrow \omega_N \rightarrow S_N$ are computed from one specified geometry. In the present paper, this action-derived extraction is then propagated into a faster surrogate kinetic scan for the global (D, η) mapping (Section VII and item 3 of the limitations).

IV. MODULE 1: THE THREE ROUTES TO MICRO-DEGENERACY g_N

The factor g_N counts the effective degrees of freedom in layer N . We justify its form via three convergent theoretical routes.

A. Route A: Mock-Modular Counting

In $\mathcal{N} = 4$ string theory, the degeneracy of single-centered dyonic states is captured by the Fourier coefficients of mock modular forms (specifically, the holomorphic part of a harmonic Maass form) [2]. Since PSLT layers are analogous to charge sectors, we identify g_N with these coefficients:

$$g_N^{(\text{A})} = |c_N|, \quad \text{where} \quad \sum c_N q^N = \frac{1}{\eta(\tau)^{\chi}} \times (\text{Mock Piece}). \quad (16)$$

B. Route B: ER=EPR and Entanglement Entropy

Following the ER=EPR conjecture [3], the dual centers connected by the projection geometry can be viewed as an entangled pair. The layer index N corresponds to the excitation level of the wormhole connection. The entropy $S(N) = \ln g_N$ must scale extensively with the effective "area" of the excitation. Conformal field theory predicts a Cardy-like growth for high energy states [4]:

$$S(N) \sim 2\pi \sqrt{\frac{c_{\text{eff}} N}{6}}. \quad (17)$$

C. Route C: Quantum Gravity Template

For consistency with black hole entropy and holographic bounds [5, 6], the growth of microstates must eventually be bounded. This suggests that while the Cardy growth dominates asymptotically, it must be regulated.

D. Unified Cardy-Controlled Envelope with High- N Regulator

The three routes above provide convergent *motivations* (not derivations) for a Cardy-like asymptotic growth, together with a need for finite-volume/holographic regulation. We therefore keep an EFT-level reference envelope with the minimum number of free parameters:

$$g_N(c_{\text{eff}}, \nu, \kappa_g) = \frac{\exp\left(2\pi\sqrt{\frac{c_{\text{eff}}N}{6}}\right)}{N^\nu} \exp\left[-\kappa_g(N-1)^2\right]. \quad (18)$$

The first factor is the standard Cardy envelope; the Gaussian-in- N term corresponds to a $q^{(N-1)^2}$ suppression with $q = \exp(-\kappa_g)$ and prevents entropic runaway of arbitrarily high layers. In this manuscript, Eq. (18) is retained as a legacy comparator and uncertainty reference, while the reported scan baseline uses the bounded two-dimensional phase-space profile.

Summary of motivational roles. Route A (Mock-modular) and Route B (ER=EPR) support the Cardy envelope's plausibility; Route C (QG Template) supports the necessity of a high- N regulator. Equation (18) remains an *EFT-level comparator ansatz* in this version.

E. Cardy Regime Validity and Finite- N Corrections

The asymptotic Cardy formula is valid when the dimensionless entropy $S(N) = 2\pi\sqrt{c_{\text{eff}}N/6}$ satisfies $S(N) \gg 1$. For our demonstrator parameters:

$$S(N) = 2\pi\sqrt{\frac{0.5 \cdot N}{6}} \approx 1.81\sqrt{N}. \quad (19)$$

Thus $S(1) \approx 1.81$, $S(2) \approx 2.57$, $S(3) \approx 3.14$. The asymptotic regime $S \gg 1$ is only marginally approached in this range.

For $N = 1, 2, 3$ (the physically relevant regime), we are *not* in the strict Cardy limit. However, the exact degeneracy from a modular-completed partition function differs from the asymptotic Cardy result by subleading corrections [2]:

$$g_N^{\text{exact}} = g_N^{\text{Cardy}} \left[1 + \mathcal{O}\left(\frac{1}{\sqrt{N}}\right) \right]. \quad (20)$$

We estimate the finite- N error as:

$$\frac{\delta g_N}{g_N} \lesssim \frac{1}{S(N)} \approx \frac{0.55}{\sqrt{N}}. \quad (21)$$

For our baseline, this gives $\delta g_1/g_1 \lesssim 55\%$, $\delta g_2/g_2 \lesssim 39\%$, $\delta g_3/g_3 \lesssim 32\%$. These uncertainties are absorbed into the effective parameters (c_{eff}, ν) , which are calibrated against the phase diagram rather than derived from first principles. (The Gaussian regulator κ_g has negligible effect for $N \leq 3$.)

Key point: We do not claim that Eq. (18) is exact for $N = 1, 2, 3$. Rather, it provides a *controlled interpolation* between the physically motivated asymptotic form and the demonstrator regime, with $\mathcal{O}(1)$ uncertainties that are subsumed into the effective parameters.

F. Assessment of a Phase-Space First-Principles Candidate

We also evaluated a semiclassical phase-space candidate to replace Eq. (18),

$$\rho_{\text{WKB}}(E) \propto \int_{U(z) < E} \frac{dz}{\sqrt{E - U(z)}}, \quad g_N^{(\text{ps})} \sim 1 + \int \rho_{\text{WKB}}(E) dE, \quad (22)$$

with $U(z) = V_{\text{eff}}(z) - m_0^2$ from the same action-derived geometry. The naive 1D implementation is not a drop-in baseline (for bound layers $E_N < 0$, using $(0 \rightarrow E_N)$ is inconsistent with the bound-state threshold convention and can produce nonphysical normalization behavior, including $g_N < 1$ in low layers). We therefore use the converged two-dimensional microcanonical extraction as baseline, while Eq. (18) is kept as a legacy comparator.

V. MODULE 2: RANK-2 COMPUTABLE KINETICS Γ_N

The kinetic rate Γ_N determines how strictly the geometry selects specific layers. We upgrade previous heuristic models to a rigorous Rank-2 system.

A. Dual-Center Tunneling Suppression

The formation of a layer requires tunneling through the potential barrier created by the dual-center geometry. We model this via a WKB action integral:

$$S_N(D) = \int_{x_-}^{x_+} dx \sqrt{V_{\text{eff}}(x; D) - \omega_N^2}, \quad (23)$$

The effective rank-2 tunneling suppression factor is then:

$$r_N(D, \eta) = \eta e^{-2S_N(D)}. \quad (24)$$

Here, η is a dimensionless *overlap amplitude* (not a probability), representing the effective strength of the dual-center interaction; values $\eta > 1$ are permitted, corresponding to collective or multi-channel enhancement. Because r_N enters only as the coarse-grained prefactor multiplying the growth matrix in Eq. (25), it should not be interpreted as a literal single-particle transmission probability. An action-derived prefactor candidate $\eta_{\text{fp}}(D)$ extracted from splitting-action data is reported in Appendix E 8 as a conservative replacement direction; it is not used in the baseline scan.

B. Dimensionality and Units

We adopt natural units $\hbar = c = 1$. The theory is defined relative to a fundamental mass scale M_* (set to unity). The dual-center separation is parameterized by a dimensionless ratio D (representing physical separation in units of the characteristic length scale M_*^{-1}). The field frequencies scale as $\omega_N \sim M_*/D$.

C. Derivation of the Rank-2 Growth Matrix

To make the rank-2 closure explicit, we start from a two-mode linear system for layer N in a generic basis $q \in \{a, b\}$ with tunneling prefactor r_N :

$$\frac{d}{dt} \begin{pmatrix} n_{N,a} \\ n_{N,b} \end{pmatrix} = r_N \begin{pmatrix} \Gamma_{N,a} & \epsilon_{\text{mix}} \\ \epsilon_{\text{mix}} & \Gamma_{N,b} \end{pmatrix} \begin{pmatrix} n_{N,a} \\ n_{N,b} \end{pmatrix}, \quad (25)$$

where $n_{N,q}$ is the coarse-grained occupation in channel q . Equation (25) is the probability-level (coarse-grained) form of a coupled-mode system after absorbing convention-dependent amplitude factors into (A_ℓ, χ) . The matrix multiplying $(n_{N,a}, n_{N,b})^T$ is therefore the growth matrix $\mathbf{M}_N$ used below.

The off-diagonal term is modeled as

$$\epsilon_{\text{mix}} = \chi \sqrt{\Gamma_{N,a} \Gamma_{N,b}}, \quad (26)$$

which is consistent with a weak-coupling overlap estimate $\epsilon_{\text{mix}} \sim \int \psi_{N,1}^* \delta V \psi_{N,2} d^3x$ and guarantees the correct rate dimension.

D. Rank-2 Computable Closure

The formation rate Γ_N is derived from the positive eigenvalue of the rank-2 interaction matrix $\mathbf{M}_N$:

$$\Gamma_N = \max(0, \lambda_+) \quad (27)$$

where λ_+ is the largest real eigenvalue of:

$$\mathbf{M}_N = r_N \begin{pmatrix} \Gamma_{N,a} & \epsilon_{\text{mix}} \\ \epsilon_{\text{mix}} & \Gamma_{N,b} \end{pmatrix} \quad (28)$$

Here, $r_N = \eta e^{-2S_N}$ is the tunneling suppression factor. The mixing term is parameterized by a dimensionless coupling χ :

$$\epsilon_{\text{mix}} = \chi \sqrt{\Gamma_{N,a} \Gamma_{N,b}} \quad (29)$$

The channel rates $\Gamma_{N,\ell}$ are organized in the current global scan by a residual leakage-hierarchy template, historically written in the Kerr-inspired power-law form [7]. In the present static two-center background this should be read only as an EFT-level bookkeeping convention for the channel hierarchy, not as a literal ergoregion superradiance derivation. To avoid dimensional ambiguity, we define

$$\hat{\omega}_N \equiv \frac{\omega_N}{M_*}, \quad \hat{\Gamma}_{N,\ell} \equiv \frac{\Gamma_{N,\ell}}{M_*}, \quad (30)$$

and write

$$\hat{\Gamma}_{N,\ell} = A_\ell(D) \hat{\omega}_N^{4\ell+5}, \quad (31)$$

which is equivalent to

$$\Gamma_{N,\ell} = A_\ell(D) \omega_N \left(\frac{\omega_N}{M_*} \right)^{4\ell+4}, \quad (32)$$

The legacy demonstrator corresponds to $A_\ell(D) = 1$, while the reported baseline uses the action-derived profile factors $A_\ell(D) = \tilde{A}_\ell(D)$ from Appendix E9. In this sense the explicit D -dependence is carried by the action-derived profile factors, while the $\hat{\omega}_N^{4\ell+5}$ factor is retained only as a compact residual hierarchy convention. In the two-channel truncation, we map the two-mode basis to the two lowest leakage channels:

$$\Gamma_{N,a} \equiv \Gamma_{N,1}, \quad \Gamma_{N,b} \equiv \Gamma_{N,2}, \quad \bar{\Gamma}_N \equiv \sqrt{\Gamma_{N,a} \Gamma_{N,b}}. \quad (33)$$

Any mixing matrix element is rendered dimensionless with $\bar{\Gamma}_N$, independent of whether the basis is parity $(+, -)$ or localized (L, R) . Higher- ℓ modes are parametrically suppressed by the $\hat{\omega}_N^{4\ell+4}$ scaling and are deferred to future extensions. All baseline scans are reported in **dimensionless units** ($M_* = 1$), where Eq. (31) and Eq. (32) are numerically identical. Constant- A_ℓ results are retained only as comparators. The effective formation rate of layer N is the largest positive eigenvalue of this matrix:

$$\boxed{\Gamma_N(D, \eta) = \max(0, \lambda_+(\mathbf{M}_N(D, \eta)))}. \quad (34)$$

For completeness, the analytic expression for λ_+ is:

$$\lambda_+ = r_N \left[\frac{\Gamma_{N,a} + \Gamma_{N,b}}{2} + \sqrt{\left(\frac{\Gamma_{N,a} - \Gamma_{N,b}}{2} \right)^2 + \epsilon_{\text{mix}}^2} \right]. \quad (35)$$

This “Rank-2 Closure” ensures that Γ_N is not a fitted parameter but a derived quantity dependent explicitly on (D, η) .

E. Physical Interpretation: Quasi-Bound Leakage and Width Proxy

The channel scaling in Eq. (32) is interpreted as a quasi-bound leakage-width proxy in the two-center geometry. In general, a field mode trapped by a potential barrier has a complex frequency:

$$\omega_N = \omega_N^{(R)} - i \omega_N^{(I)}, \quad \Gamma_N^{(\text{width})} \equiv 2\omega_N^{(I)}, \quad (36)$$

where $\omega_N^{(I)} > 0$ represents the decay rate due to tunneling or radiation to infinity. This width proxy motivates the channel factors $\Gamma_{N,\ell}$; it is distinct from the final matrix-derived formation rate $\Gamma_N(D, \eta)$ in Eq. (34).

For Kerr superradiance [7], the scaling $\Gamma \propto (\omega M)^{4\ell+5}$ arises from a rotating-background setup. In the static dual-center geometry considered here (no ergoregion), we no longer treat that same power as the physical origin of the escape width. Instead, the current baseline should be read as follows: the geometry-driven static surrogate is $\Gamma_{N,\ell}^{(\text{geo})} = \omega_N e^{-2S_{N,\ell}}$, while Eq. (32) retains a Kerr-shaped power only as a legacy hierarchy convention absorbed into the action-derived profile factors $\tilde{A}_\ell(D)$. A full non-Hermitian width derivation for this static geometry is left as a dedicated upgrade path.

Reviewer-facing status of the static-width upgrade: In the strict static formulation, the relevant object is the resonance problem for the same action-derived operator with outgoing boundary conditions, so that $\Gamma_{N,\ell}^{(\text{strict})} = -2 \Im \omega_{N,\ell}$. We implemented a minimal complex-absorbing-potential pilot on the same 1D operator chain as Appendix E9 and scanned the benchmark set $D = \{6, 12, 18\}$ for the representative $(N, \ell) = (3, 1)$ mode. The numerical chain is stable as infrastructure, but the extracted widths remain absorber-dominated over the tested (η, z_{cap}) windows and do not yet exhibit a clean plateau. For that reason, the CAP outputs are not promoted into the reported baseline. At the current stage, the most conservative first-principles static surrogate is therefore the geometry-driven rate $\Gamma_{N,\ell}^{(\text{geo})} = \omega_N e^{-2S_{N,\ell}}$ from Appendix E9, while the full complex-pole extraction remains an explicit next-step upgrade rather than an over-claimed result.

Proposition S1 (Siegert/DtN gate for promoting the static width). For fixed D, ℓ , set

$$H_{D,\ell} = -\frac{d^2}{dz^2} + U_\ell(z; D), \quad U_\ell(z; D) = U(z; D) + \frac{\ell(\ell+1)}{z^2 + \varepsilon^2},$$

and $E = \omega^2 - m_0^2$. A strict static resonance is a non- L^2 solution of

$$(H_{D,\ell} - E)\psi = 0$$

with outgoing behavior at both ends. Let $f_+(z; E)$ be the outgoing exterior solution on $[R, \infty)$, and define the right Dirichlet-to-Neumann map

$$m_+(E; R) \equiv \frac{f'_+(R; E)}{f_+(R; E)}.$$

The left map is defined analogously. Eliminating the exterior gives the finite-interval Siegert problem

$$\psi'(R) = m_+(E; R)\psi(R), \quad \psi'(-R) = -m_-(E; R)\psi(-R).$$

Because the audited one-dimensional operator is even in z , the same condition can be written in parity form. Let ψ_e solve

$$\psi_e(0) = 1, \quad \psi'_e(0) = 0,$$

and let ψ_o solve

$$\psi_o(0) = 0, \quad \psi'_o(0) = 1.$$

Then the even/odd pole conditions are

$$F_{e/o}(E; R) \equiv \psi'_{e/o}(R; E) - m_+(E; R)\psi_{e/o}(R; E) = 0.$$

For any root $E_\star$, the strict width is

$$\Gamma_{\text{strict}} = -2\Im\sqrt{m_0^2 + E_\star}.$$

In the narrow-width limit,

$$E_\star = E_R - i\gamma_E/2, \quad \Gamma_{\text{strict}} = \frac{\gamma_E}{2\omega_R} + O(\gamma_E^2), \quad \omega_R = \Re\sqrt{m_0^2 + E_\star}.$$

Therefore a numerical static-width extraction can be promoted only if the complex pole is stable under outward motion of R and under replacement of the numerical exterior representation (CAP, exact DtN, ECS, or Siegert/Robin discretization). Plateau failure is not a small numerical nuisance: it means the computed imaginary part is still controlled by the chosen exterior closure rather than by an intrinsic pole.

Proof. The exterior problems on $(-\infty, -R]$ and $[R, \infty)$ are linear second-order ODEs. For non-threshold E , specifying the outgoing branch fixes each exterior solution up to a scalar multiple; hence each exterior half-line contributes only its logarithmic derivative at the interface. Substituting these logarithmic derivatives gives the finite-interval DtN boundary conditions above. Conversely, any nonzero interior solution satisfying those DtN conditions extends uniquely to outgoing exterior solutions, hence gives a Siegert state. The parity reduction follows from evenness of U_ℓ . Finally, the width formula is just the frequency-plane parametrization $\omega = \omega_R - i\Gamma/2$, while the narrow-width relation follows by expanding $E = \omega^2 - m_0^2$.

Verified numerics. The audit code `audit_static_width_siegert_extraction.py` exports `output/cap_resonance_1d/static_width_siegert_audit_detail.csv` and `output/cap_resonance_1d/static_width_siegert_audit_summary.csv`. It uses a finite-difference half-line parity recurrence and the local outgoing Robin approximation $m_R^{\text{loc}}(E) = i\sqrt{E - U_\ell(R; D)}$ as a first Siegert gate, not as the final exact Coulomb-DtN exterior. On $D = \{6, 12, 18\}$, $\ell = 1$, $N = 3$, and $R = \{60, 80, 100, 120\}$, the nonlinear boundary residuals are solved to

$$\max_D |F_{e/o}(E; R)| = 3.48 \times 10^{-13}.$$

The even/odd roots are essentially degenerate, with maximum relative width splitting 5.06×10^{-8} and maximum complex-energy splitting 5.96×10^{-10} . However, the roots fail the required R -stability test:

$$\max_D \text{span}_R(\Gamma_{\text{Siegert}}) = 1.598, \quad \min_D \frac{\Gamma_{\text{Siegert}}(R = 120)}{\Gamma_{\text{Siegert}}(R = 80)} = 0.463.$$

The real parts drift similarly, with maximum relative R -span 1.085. The same audit compresses the existing CAP negative control: the median η -scan width span is 0.4085, matching the imposed absorber-strength span, while Γ_{CAP}/η is constant to 3.48×10^{-7} ; meanwhile the smallest z_{cap} -scan width span is still 1.9608. Thus the CAP chain is absorber-linear and the finite-radius Siegert/Robin roots are boundary-sensitive. Finally, the local Siegert widths exceed the current geometry surrogate by at least 2.95×10^7 on this benchmark, so they cannot be read as a stable extraction of $\Gamma_{N,\ell}^{(\text{geo})}$.

The combined extractor `code/extract_static_width_physical.py` implements this promotion rule explicitly and exports `output/cap_resonance_1d/static_width_physical_extraction_detail.csv` and `output/cap_resonance_1d/static_width_physical_extraction_summary.csv`. It first asks for CAP stationarity in both η and z_{cap} , together with nonlinearity of Γ_{CAP}/η , then asks for local Robin/Siegert residual, parity, and R -stability gates. A physical static width is populated only if the two independent candidates also agree. On the

same benchmark, the CAP candidate widths are only

$$7.29 \times 10^{-7} \leq \Gamma_{\text{CAP}}^{\text{cand}} \leq 7.67 \times 10^{-7},$$

but fail the η - and z_{cap} -stationarity gates. The local Robin/Siegert candidate widths are

$$3.77 \times 10^{-3} \leq \Gamma_{\text{Siegert}}^{\text{cand}} \leq 3.90 \times 10^{-3},$$

but fail the R -stability gate. The maximum CAP-vs-Siegert candidate relative difference is 1.99923, so every D -row and the global row are marked `not_promoted`. The extraction is therefore implemented, but it correctly returns no promoted physical static width on the current finite-radius/local-exterior data.

Corollary S1 (Static-width roadmap status). The strict static-width line now has an executable physical-extraction gate, but it is not promoted. The baseline should continue to report $\Gamma_{N,\ell}^{(\text{geo})} = \omega_N e^{-2S_{N,\ell}}$ as a conservative static surrogate. The next valuable task is not another unconstrained CAP scan and not a new damping object; it is one of the following controlled exterior upgrades: exact Coulomb-tail DtN, exterior complex scaling, or a Siegert method whose pole passes an explicit R -stability/plateau test.

Proposition S2 (Riccati-DtN replacement audit). The first controlled replacement of the local Robin closure is obtained by propagating the exterior logarithmic derivative. If $m(z; E) = f'(z; E)/f(z; E)$ for the outgoing exterior solution, then

$$m'(z; E) + m(z; E)^2 = U_\ell(z; D) - E.$$

At a large exterior start R_{tail} , the outgoing WKB branch gives

$$m(R_{\text{tail}}; E) = iq(R_{\text{tail}}; E) + \frac{U'_\ell(R_{\text{tail}}; D)}{4q(R_{\text{tail}}; E)^2}, \quad q(z; E) = \sqrt{E - U_\ell(z; D)},$$

with the outgoing square-root branch. Integrating this Riccati equation inward to R gives the numerical DtN value $m_+^{\text{Ric}}(E; R, R_{\text{tail}})$, which replaces m_R^{loc} in the same parity residual $F_{\text{e/o}}(E; R) = 0$.

The audit code `extract_static_width_exact_dtn.py` implements this replacement and exports `output/cap_resonance_1d/static_width_exact_dtn_detail.csv` and `output/cap_resonance_1d/static_width_exact_dtn_summary.csv`. On $D = \{6, 12, 18\}$, $\ell = 1$, $N = 3$, $R = \{60, 80, 100, 120\}$, and $R_{\text{tail}} = \{400, 800\}$, the nonlinear residuals are solved with

$$\max |F_{\text{e/o}}| = 4.78 \times 10^{-8}.$$

The resulting Riccati-DtN widths lie in the range

$$5.35 \times 10^{-4} \leq \Gamma_{\text{DtN}}^{\text{cand}} \leq 2.11 \times 10^{-3},$$

which is below the local-Robin widths but still fails promotion:

$$\max_D \text{span}_R(\Gamma_{\text{DtN}}) = 1.083, \quad \max_D \text{span}_{R_{\text{tail}}}(\Gamma_{\text{DtN}}) = 0.979, \quad \min_D \frac{\Gamma_{\text{DtN}}(R=120)}{\Gamma_{\text{DtN}}(R=80)} = 0.716.$$

One branch-tracking stress point in the $R_{\text{tail}} = 400$ validation also gives a parity-width split 1.45×10^{-1} , confirming that the WKB-start Riccati closure is not yet an exact exterior pole extractor. Thus the exact-DtN infrastructure is in place, but the static width is still not promoted; the remaining upgrade is an analytic Coulomb-tail DtN normalization or an ECS cross-check that removes the R_{tail} sensitivity.

Proposition S3 (Analytic Coulomb-tail normalization is not yet a physical static-width extraction). The large- z tail of the audited one-dimensional operator has the universal Coulomb form

$$U_\ell(z; D) = \frac{A}{z} + \frac{B_\ell}{z^2} + O(z^{-3}), \quad A = 4am_0^2, \quad B_\ell = \ell(\ell+1) + 4a^2m_0^2.$$

For the analytic tail equation $m' + m^2 = A/z + B_\ell/z^2 - E$, write

$$m(z) = ik + \sum_{n=1}^M c_n z^{-n}, \quad k = \sqrt{E},$$

with the outgoing branch. Equating powers gives the triangular recurrence

$$2ik c_n + \sum_{j=1}^{n-1} c_j c_{n-j} - (n-1)c_{n-1} = r_n, \quad r_1 = A, \quad r_2 = B_\ell, \quad r_n = 0 \ (n \geq 3).$$

The audit uses this analytic Coulomb value at a matching radius R_C , then propagates the full Riccati equation with the actual $U_\ell(z; D)$ inward to the interface R . This removes the crude WKB start as an independent approximation while retaining the D -dependent $O(z^{-3})$ tail between R_C and R .

The implementation is the Coulomb-series mode of `code/extract_static_width_exact_dtn.py`,

```
--tail-mode coulomb_series --coulomb-order 6 --tail-validation --fallback-guesses,
```

and exports `output/cap_resonance_1d/static_width_coulomb_tail_dtn_o6_detail.csv` and `output/cap_resonance_1d/static_width_coulomb_tail_dtn_o6_summary.csv`. On the same $D = \{6, 12, 18\}$, $R = \{60, 80, 100, 120\}$, and $R_C = \{400, 800\}$ benchmark, fallback branch tracking gives

$$\max |F_{e/o}| = 5.41 \times 10^{-7}, \quad 1.67 \times 10^{-3} \leq \Gamma_{\text{Coul}}^{\text{cand}} \leq 1.87 \times 10^{-3}.$$

However the promotion gates still fail:

$$\max_D \text{span}_R(\Gamma_{\text{Coul}}) = 0.763, \quad \max_D \text{span}_{R_C}(\Gamma_{\text{Coul}}) = 0.617, \quad \min_D \frac{\Gamma_{\text{Coul}}(R=120)}{\Gamma_{\text{Coul}}(R=80)} = 0.673.$$

The largest parity split in this stress run is 9.19×10^{-2} . Thus the analytic Coulomb-tail normalization is useful as a negative control—it removes the purely ad hoc WKB start—but it does not yet produce an exterior-stable pole. The remaining static-width path should therefore be an ECS cross-check or a fully analytic Coulomb/Whittaker DtN implementation, not a further CAP scan.

Proposition S4 (Exterior complex scaling also rejects the current static-width promotion). The next independent exterior check is exterior complex scaling. For $x \in [-L, L]$, define the piecewise analytic contour

$$Z_{\theta, R_0}(x) = \begin{cases} x, & |x| \leq R_0, \\ \text{sgn}(x)\{R_0 + e^{i\theta}(|x| - R_0)\}, & |x| > R_0, \end{cases} \quad J_{\theta, R_0}(x) = \frac{dZ_{\theta, R_0}}{dx}.$$

If $\phi(x) = \psi(Z_{\theta, R_0}(x))$, then the exterior-scaled one-dimensional operator is discretized in flux form as

$$H_{\text{ECS}}(\theta, R_0, L)\phi = -J_{\theta, R_0}^{-1} \frac{d}{dx} \left(J_{\theta, R_0}^{-1} \frac{d\phi}{dx} \right) + U_\ell(Z_{\theta, R_0}(x); D)\phi.$$

This is a direct complex-contour replacement of the finite-box exterior. A genuine resonance pole should be insensitive to the artificial parameters θ , R_0 , and L , while the rotated continuum and box branches move with them.

The implementation `code/audit_static_width_ecs_crosscheck.py` evaluates the complex-scaled potential directly on Z_{θ, R_0} and therefore avoids the real-casting shortcut used by the finite-radius utilities. It exports `output/cap_resonance_1d/static_width_ecs_crosscheck_detail.csv` and `output/cap_resonance_1d/static_width_ecs_crosscheck_summary.csv`. On $D = \{6, 12, 18\}$, $\ell = 1$, $N = 3$, $\theta = \{0.08, 0.12, 0.16\}$, $R_0 = \{50, 60\}$, $L = \{100, 120\}$, and $h = 0.15$, the sparse eigen-residual is small:

$$\max \|(H_{\text{ECS}} - E_{\text{ECS}})\phi\|/\|\phi\| = 4.98 \times 10^{-14}.$$

The tracked seed branch has candidate widths

$$2.09 \times 10^{-4} \leq \Gamma_{\text{ECS}}^{\text{cand}} \leq 5.78 \times 10^{-4}, \quad \Gamma_{\text{ECS}}^{\text{med}} = 3.61 \times 10^{-4} \quad (\text{global}).$$

However the stability gates fail:

$$\max_D \text{span}_\theta(\Gamma_{\text{ECS}}) = 0.669, \quad \max_D \text{span}_{R_0}(\Gamma_{\text{ECS}}) = 0.223, \quad \max_D \text{span}_L(\Gamma_{\text{ECS}}) = 0.144.$$

The complex energy itself moves by

$$\max_D \text{span}_\theta(E_{\text{ECS}}) = 2.90 \times 10^{-4}, \quad \max_D \text{span}_L(E_{\text{ECS}}) = 1.15 \times 10^{-3}.$$

The overlap tracker remains high ($\min |\langle \psi_{\text{seed}}, \phi_{\text{ECS}} \rangle| = 0.862$), and the residuals are at roundoff scale, so this is not a solver failure. It is a clean diagnostic that the current tracked branch is still an exterior-scaling-sensitive box/continuum branch rather than an isolated resonance pole.

Corollary S4 (Static-width promotion remains blocked after ECS). The CAP, local Robin/Siegert, Riccati-DtN, analytic Coulomb-tail DtN, and ECS cross-checks now all agree on the same conservative conclusion: the present finite-radius/open-boundary extractions are useful infrastructure, but none supplies a parameter-stable physical pole. The paper should therefore continue to report $\Gamma_{N,\ell}^{(\text{geo})} = \omega_N e^{-2S_{N,\ell}}$ as the static-width surrogate. Further work, if pursued, should be a full Coulomb/Whittaker DtN or a more refined ECS branch-tracking theorem, not a new CAP scan and not a new static-width family.

Proposition S5 (Refined ECS branch tracking finds no resolved stable pole). The single-branch ECS audit can be sharpened without changing the operator family. For each ECS parameter

$$p = (\theta, R_0, L)$$

let

$$\Sigma_p(D) = \{E_j(p; D)\}_{j=1}^K$$

be the sparse spectrum returned near the Hermitian $N = 3, \ell = 1$ seed. For every anchor $a \in \Sigma_{p_a}(D)$, define the optimistic nearest-neighbor continuation

$$B_a(p) = \arg \min_{E \in \Sigma_p(D)} |E - a|.$$

This construction is deliberately favorable to promotion: it does not force a fixed overlap label and allows the nearest branch in the computed ECS spectrum to represent the same pole at each parameter point. Consequently, any resolved isolated pole in the sampled spectral window, separated from neighboring branches by more than the numerical residual and moving within the stability tolerances, would be detected by at least one such anchor branch.

The audit `code/audit_static_width_ecs_branch_tracking.py` implements this branch envelope and exports `output/cap_resonance_1d/static_width_ecs_branch_spectrum_detail.csv`, `output/cap_resonance_1d/static_width_ecs_branch_candidates.csv`, and `output/cap_resonance_1d/static_width_ecs_branch_summary.csv`. With $K = 24$, the

same $D = \{6, 12, 18\}$, $\theta = \{0.08, 0.12, 0.16\}$, $R_0 = \{50, 60\}$, $L = \{100, 120\}$, and $h = 0.15$ grid gives 864 resolved eigenvalues. The anchor filter produces 84 non-growing branch candidates, of which 4 are seed-qualified by the overlap gate. None passes all ECS stability gates:

$$N_{\text{pass}} = 0.$$

For the seed-qualified branches, the best candidates reproduce the same scaling drift as Proposition S4:

$$\max_D \text{span}_\theta(\Gamma_{\text{branch}}) = 0.669, \quad \max_D \text{span}_{R_0}(\Gamma_{\text{branch}}) = 0.223, \quad \max_D \text{span}_L(\Gamma_{\text{branch}}) = 0.144,$$

with

$$\max_D \text{span}_L(E_{\text{branch}}) = 1.15 \times 10^{-3}, \quad \min_D \min_p |\langle \psi_{\text{seed}}, B_a(p) \rangle| = 0.864.$$

The median seed-branch width is again on the 3.5×10^{-4} scale, but it is not stable:

$$3.50 \times 10^{-4} \leq \Gamma_{\text{branch}}^{\text{seed}} \leq 3.61 \times 10^{-4} \quad \text{across } D = \{6, 12, 18\}.$$

Thus the negative conclusion is not an artifact of choosing the overlap-maximal branch in Proposition S4. Even an optimistic nearest-neighbor continuation over the resolved ECS spectrum finds no branch that is both seed-connected and stable in (θ, R_0, L) .

Corollary S5 (Static-width mainline is now narrowed to full analytic DtN). The refined ECS theorem closes the current ECS route as a promotion mechanism on the audited grid. The static-width language should therefore remain conservative: $\Gamma_{N,\ell}^{(\text{geo})}$ is still the reported surrogate, and the non-Hermitian outputs remain diagnostics. The remaining mathematically distinct route is the full Coulomb/Whittaker outgoing DtN map audited next in Proposition S6; further CAP scans, truncated tail starts, and unconstrained ECS branch picking should not be treated as new support.

Proposition S6 (Full Coulomb/Whittaker DtN is audited but still not a physical static-width extraction). The final non-CAP exterior replacement in this line is obtained by solving the analytic Coulomb tail exactly. For the same one-dimensional operator, the large- z tail has

$$f''(z) + \left(E - \frac{A}{z} - \frac{B_\ell}{z^2} \right) f(z) = 0, \quad A = 4am_0^2, \quad B_\ell = \ell(\ell + 1) + 4a^2m_0^2.$$

Let $k = \sqrt{E}$ on the outgoing branch and set

$$x = -2ikz, \quad \kappa = -\frac{iA}{2k}, \quad \mu = \sqrt{B_\ell + \frac{1}{4}}.$$

Since $d/dz = (-2ik)d/dx$, the equation becomes the Whittaker equation

$$\frac{d^2 W}{dx^2} + \left(-\frac{1}{4} + \frac{\kappa}{x} + \frac{1/4 - \mu^2}{x^2} \right) W = 0.$$

The outgoing solution is $W_{\kappa,\mu}(x)$, because $W_{\kappa,\mu}(-2ikz) \sim e^{ikz}(-2ikz)^\kappa$ on the outgoing sheet.

Using the hypergeometric representation

$$W_{\kappa,\mu}(x) = e^{-x/2} x^{\mu+1/2} U(a, b, x), \quad a = \mu - \kappa + \frac{1}{2}, \quad b = 2\mu + 1,$$

and $\partial_x U(a, b, x) = -a U(a+1, b+1, x)$, the exact outgoing tail logarithmic derivative at the matching radius R_C is

$$m_W(R_C; E) = (-2ik) \left[-\frac{1}{2} + \frac{\mu + 1/2}{x} - a \frac{U(a+1, b+1, x)}{U(a, b, x)} \right]_{x=-2ikR_C}.$$

This is the requested full Coulomb/Whittaker DtN datum. The implementation is the `whittaker` mode of `code/extract_static_width_exact_dtn.py`; it evaluates the complex U-ratio with `mpmath` at 50 digits and then propagates the full Riccati equation for the actual $U_\ell(z; D)$ inward from R_C to R .

The audit exports `output/cap_resonance_1d/static_width_whittaker_dtn_dps50_detail.csv` and `output/cap_resonance_1d/static_width_whittaker_dtn_dps50_summary.csv`. On $D = \{6, 12, 18\}$, $\ell = 1$, $N = 3$, $R = \{60, 80, 100, 120\}$, and $R_C = \{400, 800\}$, the nonlinear residual gate is still met:

$$\max |F_{e/o}| = 9.82 \times 10^{-7}.$$

The candidate medians are on the same static-width scale as the Coulomb-series audit,

$$1.66 \times 10^{-3} \leq \Gamma_W^{\text{cand}} \leq 1.91 \times 10^{-3},$$

but the exterior-stability gates still fail:

$$\max_D \text{span}_R(\Gamma_W) = 0.763, \quad \max_D \text{span}_{R_C}(\Gamma_W) = 0.536,$$

$$0.673 \leq \frac{\Gamma_W(R=120)}{\Gamma_W(R=80)} \leq 0.681.$$

The largest parity-width split is 8.94×10^{-2} , concentrated in the same stress sector that already failed the earlier exact-DtN gates. A higher-precision spot check at 80 digits on the $D = 6$, $R = 60$, $R_C = 800$ row gives $\Gamma_{\text{even}} = 1.99471337 \times 10^{-3}$, $\Gamma_{\text{odd}} = 1.99471342 \times 10^{-3}$, and parity split 2.61×10^{-8} , identical at displayed precision to the 50-digit backend. Thus the failure is not a complex-special-function precision failure; it is the same absence of an exterior-stable pole on the current finite-radius branch.

Corollary S6 (Full analytic DtN closes the current static-width promotion route).

The full Coulomb/Whittaker tail removes the last ad hoc exterior datum from the finite-radius DtN chain. Since even this audited backend remains strongly R - and R_C -dependent, the strict static-width upgrade is still not promoted. The manuscript should continue to report $\Gamma_{N,\ell}^{(\text{geo})} = \omega_N e^{-2S_{N,\ell}}$ as the conservative static surrogate. Further work in this direction is no longer another CAP scan, another truncated tail normalization, or another unconstrained ECS branch search; it would have to be a new exterior-stable complex-pole theorem or branch-isolation argument, with a continued-fraction H'/H implementation serving only as an independent backend cross-check.

Proposition S7 (Nearest-neighbor Whittaker branch isolation finds no resolved stable branch). The Whittaker-DtN audit can be sharpened without changing the exterior family. For a fixed D , let

$$\mathcal{P} = \{(R_C, R, \pi) : R_C \in \{400, 800\}, R \in \{60, 80, 100, 120\}, \pi \in \{\text{e}, \text{o}\}\}$$

be the audited parameter set, and let $\mathcal{Z}_p(D)$ be the resolved Whittaker-DtN roots recorded at parameter point p . For every anchor $a \in \cup_{p \in \mathcal{P}} \mathcal{Z}_p(D)$, define the optimistic nearest-neighbor continuation

$$B_a(p) = \arg \min_{z \in \mathcal{Z}_p(D)} |z - a|.$$

This is deliberately favorable to promotion: it does not force the originally selected seed label to persist. Therefore, if the audited Whittaker candidate set already contained an exterior-stable pole whose displacement stayed inside the R -, R_C -, and parity-stability gates, at least one anchor continuation B_a would pass those gates.

The implementation `code/audit_static_width_whittaker_branch_isolation.py` applies this construction to the full S6 detail table and exports `output/cap_resonance_1d/static_width_whittaker_branch_isolation_candidates.csv`, `output/cap_resonance_1d/static_width_whittaker_branch_isolation_selected.csv`, `output/cap_resonance_1d/static_width_whittaker_branch_isolation_summary.csv`, and `output/cap_resonance_1d/static_width_whittaker_branch_isolation_roots.png`. Across $D = \{6, 12, 18\}$ it tests 48 anchor continuations. None passes:

$$N_{\text{pass}} = 0.$$

Per D , this is 0/16 branches. The best available branch at $D = 6$ still has

$$\text{span}_R(\Gamma) = 0.674, \quad \text{span}_{R_C}(\Gamma) = 0.536, \quad 0.680 \leq \Gamma(R = 120)/\Gamma(R = 80) \leq 0.681,$$

and it also retains the large stress-sector parity split 8.94×10^{-2} . At $D = 12$ and $D = 18$, the parity split is already small, but the exterior-stability obstruction persists:

D	$\text{span}_R(\Gamma)$	$\text{span}_{R_C}(\Gamma)$	$\Gamma(R = 120)/\Gamma(R = 80)$
12	0.751	0.526	0.6765–0.6766
18	0.763	0.522	0.6729–0.6741

The best normalized branch-stability scores are 8.94×10^3 , 10.53, and 10.43 for $D = 6, 12, 18$, respectively, where a passing branch would require score ≤ 1 . Thus the failure is not merely that the seed-selected row was an unlucky label: even the most permissive nearest-neighbor continuation over the resolved Whittaker-DtN candidate set finds no stable branch.

Corollary S7 (The remaining static-width task is a zero-count/conditioning theorem, not more branch picking). The current static-width evidence now closes the resolved-candidate loophole as well as the backend loophole. Continuing this line is still possible, but the task has changed: one must either give an argument-principle zero count for the finite-interval Whittaker-DtN residual on a specified complex contour, or prove a tail-remainder/condition-number theorem showing that the exact exterior pole survives the $R_C \rightarrow \infty$ and R -movement limits with a quantified displacement bound. Until such a theorem is supplied, the stable reviewer-facing statement remains conservative: $\Gamma_{N,\ell}^{(\text{geo})}$ is the reported static surrogate, while all finite-radius complex roots remain diagnostics.

Proposition S8 (Local argument-principle contours certify isolated Whittaker-DtN zeros). The next check separates two issues that were coupled in S6–S7: whether the computed roots are genuine local zeros of the finite-radius residual, and whether those local zeros define a single exterior-stable pole as R and R_C vary. For fixed $p = (D, R_C, R, \pi)$, write the Whittaker-DtN finite-interval residual as $F_p(E)$. Around each numerical root $E_0(p)$, define the rectangle

$$Q_p = \{E : |\Re(E - E_0)| \leq 10^{-4}, |\Im(E - E_0)| \leq 10^{-4}\}.$$

If F_p is pole-free in Q_p , the argument principle gives

$$N_{Q_p} = \frac{1}{2\pi i} \oint_{\partial Q_p} \frac{F'_p(E)}{F_p(E)} dE = \frac{1}{2\pi} \Delta_{\partial Q_p} \arg F_p(E).$$

The implemented contour is traversed clockwise, so a single simple zero gives argument index -1 . The audit therefore records $|\text{index}|$ as the local zero count, together with the boundary gap

$$g_p = \min_{E \in \partial Q_p} |F_p(E)|$$

and a Newton displacement proxy

$$\delta_p^{\text{Newt}} = \frac{|F_p(E_0)|}{|F'_p(E_0)|}.$$

The implementation `code/audit_static_width_whittaker_argument_principle.py` evaluates the same Whittaker-DtN residual used in S6 and exports `output/cap_resonance_1d/static_width_whittaker_argument_principle_detail.csv`, `output/cap_resonance_1d/static_width_whittaker_argument_principle_summary.csv`, and `output/cap_resonance_1d/static_width_whittaker_argument_principle_index.png`. On all 48 audited roots,

$$\text{index} = -1, \quad |\text{index}| = 1, \quad N_{\text{local pass}} = 48/48.$$

The numerical contour is well separated from zero:

$$\min_p g_p = 2.40 \times 10^{-2}, \quad \min_p \frac{g_p}{|F_p(E_0)|} = 8.35 \times 10^5,$$

and the largest Newton displacement proxy is tiny relative to the contour size:

$$\max_p \frac{\delta_p^{\text{Newt}}}{10^{-4}} = 7.66 \times 10^{-7}.$$

The Cauchy–Riemann directional derivative mismatch of the numerical residual stays below 9.53×10^{-3} , so the local index computation is not being driven by gross non-holomorphic numerical noise. Thus the S6 roots are locally isolated finite-radius Whittaker-DtN zeros. Their failure is not that they are solver ghosts; their failure is that the isolated zeros move too much with the exterior parameters.

Corollary S8 (The remaining theorem is a displacement bound, not local root existence). The local zero-count gate is now positive, while the resolved-branch stability gate remains negative. Consequently, a static-width promotion cannot be obtained by further local root polishing: the roots already exist locally. The missing theorem is a displacement estimate comparing F_{R,R_C} to a limiting exterior residual F_∞ , for example

$$|\Delta E| \leq \frac{\|F_{R,R_C} - F_\infty\|_{\partial Q}}{\inf_{\partial Q} |F'_\infty|} \quad \text{or an equivalent Rouché/condition-number bound,}$$

strong enough to beat the observed R - and R_C -drift. Until such a bound is proved, the static-width result remains diagnostic rather than promoted.

Proposition S9 (Parent-lift criterion for finite Whittaker-DtN zeros). The finite-radius zeros can be pushed to a parent static-width pole only through a Rouché/conditioning estimate, not through another local root solve. Let $m_\infty(R, E; D)$ denote the exact outgoing Weyl logarithmic derivative for the action-derived half-line tail, and let $m_{R_C}^W(R, E; D)$ denote the Whittaker-initialized Riccati DtN value used in Proposition S6. On the same finite-difference interior grid, define

$$F_{\infty,p}(E) := F_p(E) \text{ with } m_{R_C}^W(R, E; D) \text{ replaced by } m_\infty(R, E; D), \quad p = (D, R_C, R, \pi).$$

Since the boundary residual is affine in the DtN value,

$$F_p(E) - F_{\infty,p}(E) = -\frac{2}{h_p} (m_{R_C}^W(R, E; D) - m_{\infty}(R, E; D)).$$

Thus the whole parent-lift problem reduces to an exterior Weyl-function remainder. The clean way to read that remainder is two-step. First compare the exact action-derived Weyl value m_{∞} on $[R_C, \infty)$ with the pure Coulomb/Whittaker Weyl value m_C . For $\Delta_C = m_C - m_{\infty}$, the Riccati equations give the forced equation

$$\Delta'_C(z, E) + (m_C(z, E) + m_{\infty}(z, E))\Delta_C(z, E) = U_{\text{tail}}^C(z) - U_{\ell}(z; D).$$

Second, propagate the mismatch from R_C back to R through the same action-derived Riccati flow. For $\delta m(z, E) = m_{R_C}^W(z, E) - m_{\infty}(z, E)$ on $R \leq z \leq R_C$, both sides solve the same Riccati equation, hence

$$\delta m'(z, E) + (m_{R_C}^W(z, E) + m_{\infty}(z, E))\delta m(z, E) = 0, \quad \delta m(R_C, E) = \Delta_C(R_C, E).$$

Therefore any parent-side theorem must supply a tail-forcing estimate and a transfer-stability multiplier $\mathcal{S}_{R,R_C}(E)$, schematically

$$|\delta m(R, E)| \leq \mathcal{S}_{R,R_C}(E) (|\Delta_C(R_C, E)| + \mathcal{R}_{>R_C}(E)).$$

Combining this with the affine residual identity gives the auditable Rouché condition

$$\epsilon_p := \sup_{E \in \partial Q_p} |F_p(E) - F_{\infty,p}(E)| < g_p := \min_{E \in \partial Q_p} |F_p(E)|.$$

If this strict inequality holds and both residuals are holomorphic in Q_p , then $F_{\infty,p}$ has the same number of zeros in Q_p as F_p . Since Proposition S8 gives one finite zero in each Q_p , this would place one exact-parent grid pole inside that same local box.

There is also a geometry condition before any tail estimate can succeed branchwise. If a single parent pole $E_*(D)$ explains every finite zero in a candidate branch B , then $E_*(D)$ must lie in the intersection of the local boxes. Therefore the audited branch diameter must satisfy at least the square-contour necessary lower-bound test

$$\frac{\text{diam}_E(B)}{2\sqrt{2}\rho} \leq 1, \quad \rho = 10^{-4},$$

up to the harmless distinction between square and disk contours. The implementation `code/audit_static_width_whittaker_parent_lift.py` combines the S7 nearest-neighbor branches with the S8 contour data and exports `output/cap_resonance_1d/static_width_whittaker_parent_lift_detail.csv`, `output/cap_resonance_1d/static_width_whittaker_parent_`

lift_summary.csv, and output/cap_resonance_1d/static_width_whittaker_parent_lift_budget.png. It finds

$$N_{\text{local argument}} = 48/48, \quad N_{\text{local box cover}} = 0/48, \quad N_{\text{parent lift budget}} = 0/48.$$

The best branches still require a cover radius much larger than the local contour:

$$\min_B \frac{\text{diam}_E(B)}{2\rho} = 38.72, \quad \min_B \frac{\text{diam}_E(B)}{2\sqrt{2}\rho} = 27.38.$$

In residual units, the most favorable half-span displacement is still too large for the local Rouché budget:

$$\min_B \frac{|F'_p(E_0)| \text{diam}_E(B)/2}{\min(g_p, |F'_p(E_0)|\rho)} = 43.48.$$

Thus the current finite Whittaker-DtN zeros do not yet lift to a parent static-width object. The obstruction is now sharply localized: one would need an exterior Weyl-function remainder/conditioning theorem at least this strong, or a larger pole-free parent contour with a new zero-count proof.

Corollary S9 (Static-width promotion remains blocked, but the missing theorem is explicit). The parent-side target is no longer a special-function backend, a branch-labeling choice, or a local existence argument. Those gates are closed by S6–S8. What remains is the single inequality

$$\frac{2}{h_p} \sup_{\partial Q_p} |m_{R_C}^W(R, E; D) - m_\infty(R, E; D)| < g_p$$

on a contour large enough to contain a common parent pole, together with the corresponding zero-count/holomorphy statement. The present audited constants miss that lift by a factor of at least 27.38 geometrically and 43.48 in the local residual budget. Therefore the manuscript should continue to report $\Gamma_{N,\ell}^{(\text{geo})}$ as the strict static surrogate; future progress must be a parent Weyl transfer-stability/conditioning proof rather than another finite-radius root family.

Proposition S10 (The Coulomb-tail forcing is small; the Riccati transfer conditioning is the obstruction). The S9 inequality can be stress-tested without solving for a new root. At each already-audited S6 zero $E_0(p)$, evaluate the same Whittaker-DtN map twice, once with $R_C = 400$ and once with $R_C = 800$, but keep $E = E_0(p)$ fixed. This gives the finite transfer diagnostic

$$\Delta m_{400,800}(R, E_0; D) := m_{400}^W(R, E_0; D) - m_{800}^W(R, E_0; D).$$

The local S8 Rouché budget in DtN units is

$$B_p^{(m)} = \frac{h_p}{2} g_p, \quad g_p = \min_{\partial Q_p} |F_p(E)|,$$

because the boundary residual changes by $(2/h_p)\Delta m$. Thus a necessary local tail-conditioning pass is

$$|\Delta m_{400,800}(R, E_0; D)| \leq B_p^{(m)}.$$

The raw Coulomb-tail forcing is measured by

$$T_{R_C}(D) = \int_{R_C}^{\infty} \left| U_{\ell}(z; D) - \left(\frac{4am_0^2}{z} + \frac{\ell(\ell+1) + 4a^2m_0^2}{z^2} \right) \right| dz.$$

For the action-derived two-center potential, the first omitted term is $O(z^{-3})$; numerically the terminal coefficient satisfies the expected D -dependent z^{-3} scaling, so the audit completes the finite quadrature by a $C/(2Z_{\max}^2)$ tail.

The implementation `code/audit_static_width_whittaker_tail_conditioning.py` exports `output/cap_resonance_1d/static_width_whittaker_tail_conditioning_detail.csv`, `output/cap_resonance_1d/static_width_whittaker_tail_conditioning_summary.csv`, and `output/cap_resonance_1d/static_width_whittaker_tail_conditioning_budget.png`. On the 48 audited S6/S8 rows, the raw forcing is uniformly small relative to the local budget:

$$\max_p \frac{T_{R_C(p)}(D)}{B_p^{(m)}} = 1.64 \times 10^{-2}.$$

The returned DtN map is not uniformly small:

$$\min_p \frac{|\Delta m_{400,800}|}{B_p^{(m)}} = 0.9896, \quad \max_p \frac{|\Delta m_{400,800}|}{B_p^{(m)}} = 88.60.$$

Only 4/48 rows pass this tail-conditioning budget, namely the farthest-interface $R = 120$, $R_C = 800$ even/odd rows for $D = 12$ and $D = 18$. The worst failures occur at $R_C = 400$, with the smallest local budgets and a returned DtN displacement near 10^{-1} . Equivalently, the empirical transfer multiplier lower bound ranges from

$$3.61 \times 10^2 \leq \frac{|\Delta m_{400,800}|}{\max(T_{400}, T_{800})} \leq 4.53 \times 10^4,$$

while the local budget would require the effective multiplier to remain below the row-dependent allowed range, whose global minimum is 60.97. Hence the raw Coulomb-tail approximation is not the bottleneck; the missing theorem is a nontrivial Riccati/Weyl transfer-stability theorem.

Corollary S10 (The next theorem must control transfer amplification, not tail coefficients). The parent-lift problem has narrowed again. Improving the Whittaker special-function backend or adding more terms to the Coulomb tail cannot by itself close the static-width gate, because the direct tail remainder is already below 2% of the local DtN budget on every audited row. The obstruction is the inward transfer from R_C to R : it can amplify the tiny tail

mismatch by factors 10^2 – 10^4 . A future positive theorem must therefore prove a sharp outgoing Weyl transfer bound, or else use a larger pole-free contour whose Rouché budget survives this amplification. Until then, $\Gamma_{N,\ell}^{(\text{geo})}$ remains the reported strict static surrogate.

Proposition S11 (Outgoing Weyl transfer-stability attempt and condition-number audit). Let $T_{R,R_0}(E; D)$ be the Riccati transfer map that sends an outgoing logarithmic derivative $a = m(R_0)$ at R_0 to the value $m(R)$ obtained from

$$m'(z) = U_\ell(z; D) - E - m(z)^2, \quad m(R_0) = a.$$

The infinitesimal condition number of this map is obtained by differentiating with respect to a . If

$$\kappa(z) = \frac{\partial m(z)}{\partial a},$$

then

$$\kappa'(z) = -2m(z)\kappa(z), \quad \kappa(R_0) = 1,$$

and the local linear condition number is

$$K_{R,R_0}(E; D) = |\kappa(R)|.$$

A direct sufficient theorem for the S9 lift would be

$$K_{R,R_0}(E; D) |\Delta a(R_0, E; D)| < B_p^{(m)},$$

where $B_p^{(m)} = (h_p/2)g_p$. This is the most transparent parent-side condition-number route. The audit below tests it together with a finite-amplitude dyadic Cauchy version,

$$\frac{|T_{R,R_0}(m_{R_0}^C) - T_{R,2R_0}(m_{2R_0}^C)|}{B_p^{(m)}} \leq 1,$$

for the two dyadic pairs $R_0 \rightarrow 2R_0 = 400 \rightarrow 800$ and $800 \rightarrow 1600$.

The implementation `code/audit_static_width_whittaker_transfer_condition.py` exports `output/cap_resonance_1d/static_width_whittaker_transfer_condition_detail.csv`, `output/cap_resonance_1d/static_width_whittaker_transfer_condition_summary.csv`, and `output/cap_resonance_1d/static_width_whittaker_transfer_condition_budget.png`. The $400 \rightarrow 800$ finite transfer is not uniform:

$$N_{\text{emp}}^{400 \rightarrow 800} = 4/48, \quad \frac{\Delta m_{400,800}}{B^{(m)}} \in [0.9896, 88.60],$$

with median ratio 8.36. Moving the dyadic pair outward helps but still does not close the theorem:

$$N_{\text{emp}}^{800 \rightarrow 1600} = 24/48, \quad \frac{\Delta m_{800,1600}}{B^{(m)}} \in [1.39 \times 10^{-3}, 10.66],$$

with median ratio 2.16. Thus the parent tail is Cauchy-improving as R_C moves outward, but not uniformly enough on the current local contours.

The infinitesimal condition-number route is even less favorable. On $400 \rightarrow 800$,

$$K_{R,400} \in [1.02 \times 10^2, 2.31 \times 10^5],$$

and on $800 \rightarrow 1600$,

$$K_{R,800} \in [4.80 \times 10^2, 6.45 \times 10^6].$$

Consequently the linearized condition $K|\Delta a| < B^{(m)}$ is far more pessimistic than the finite dyadic displacement: for $800 \rightarrow 1600$ it passes 0/48 rows even though the finite dyadic test passes 24/48. The reason is that the Riccati map is a fractional/second-order transfer map; finite perturbations can saturate or rotate in phase, while the crude infinitesimal modulus bound discards that structure.

Corollary S11 (The viable theorem is finite-amplitude Weyl contraction, not a crude Gronwall bound). The attempted condition-number theorem identifies the right mathematical object but rejects the simplest proof route. A proof based only on $\exp \int 2|m| dz$ or on the raw infinitesimal modulus K_{R,R_0} is too pessimistic to promote the static width. The only promising parent-side route left in this family is a finite-amplitude outgoing Weyl contraction or a projective/Jost transfer estimate that exploits the fractional-linear structure of the Riccati flow. Numerically, moving from $400 \rightarrow 800$ to $800 \rightarrow 1600$ improves the dyadic pass count from 4/48 to 24/48, so the limit is plausibly Cauchy at sufficiently large R_C , but the current local boxes do not give a reviewer-facing uniform theorem. The strict static-width result therefore remains unpromoted.

Proposition S12 (Finite-amplitude projective/Jost transfer contraction audit). The Riccati flow in Proposition S11 is fractional-linear because it is induced by the second-order transfer matrix. Let u_a, u_b solve

$$u''(z) = (U_\ell(z; D) - E)u(z), \quad u_a(R_0) = u_b(R_0) = 1, \quad u'_a(R_0) = a, \quad u'_b(R_0) = b.$$

Then the transferred logarithmic derivatives are

$$T_{R,R_0}(a) = \frac{u'_a(R)}{u_a(R)}, \quad T_{R,R_0}(b) = \frac{u'_b(R)}{u_b(R)}.$$

The Wronskian is constant:

$$u_a u'_b - u'_a u_b = b - a.$$

Therefore the finite-amplitude transfer identity is exact:

$$T_{R,R_0}(a) - T_{R,R_0}(b) = \frac{a - b}{u_a(R)u_b(R)}.$$

Equivalently, if $M = \begin{pmatrix} A & B \\ C & D \end{pmatrix}$ is the transfer matrix, then

$$T(a) = \frac{C + Da}{A + Ba}, \quad T(a) - T(b) = \frac{\det(M)(a - b)}{(A + Ba)(A + Bb)}.$$

Thus the right parent-side theorem is not a crude Gronwall estimate for K , but a Jost-product lower bound:

$$\frac{|a - b|}{|u_a(R)u_b(R)|} < B_p^{(m)}.$$

The executable audit `code/audit_static_width_whittaker_projective_transfer.py` implements this identity with periodically rescaled Jost solutions and exports `output/cap_resonance_1d/static_width_whittaker_projective_transfer_detail.csv`, `output/cap_resonance_1d/static_width_whittaker_projective_transfer_summary.csv`, and `output/cap_resonance_1d/static_width_whittaker_projective_transfer_budget.png`. It keeps $\Gamma_{N,\ell}^{(\text{geo})}$ as the strict static surrogate and tests only the dyadic transfer contraction of the existing Whittaker roots. The exact projective endpoint budget reproduces S11 on the first two dyadic pairs:

$$N_{\text{proj}}^{400 \rightarrow 800} = 4/48, \quad \frac{\Delta m_{400,800}}{B^{(m)}} \in [0.9896, 88.60],$$

with median 8.36, and

$$N_{\text{proj}}^{800 \rightarrow 1600} = 24/48, \quad \frac{\Delta m_{800,1600}}{B^{(m)}} \in [1.39 \times 10^{-3}, 10.66],$$

with median 2.16. On the next dyadic pair, however, the finite-amplitude endpoint contraction closes on the audited grid:

$$N_{\text{proj}}^{1600 \rightarrow 3200} = 48/48, \quad \frac{\Delta m_{1600,3200}}{B^{(m)}} \leq 1.435 \times 10^{-3},$$

with median ratio 6×10^{-6} . The Wronskian/Jost identity is numerically stable at the budget level: the maximum absolute formula error divided by $B^{(m)}$ is 1.78×10^{-8} . An independent Riccati rerun at $1600 \rightarrow 3200$, exported as `output/cap_resonance_1d/static_width_whittaker_transfer_condition_r1600_r3200_summary.csv`, also gives 48/48 empirical passes with the same maximum ratio 1.435×10^{-3} .

The conservative convex-segment denominator bound is still too pessimistic. Replacing the endpoint product $|u_a(R)u_b(R)|$ by a uniform lower bound over the straight segment between a and b passes only 24/48 rows at $1600 \rightarrow 3200$, because the segment can approach the projective pole even when both endpoint Jost products are large. Hence the numerically correct finite-amplitude theorem should be an endpoint Jost-product or pole-exclusion theorem for the actual outgoing dyadic data, not an all-perturbations Lipschitz ball around the segment.

Corollary S12 (The remaining static-width gap is analytic Jost-product control, not backend design). The static-width obstruction is now sharply localized. Full Whittaker DtN roots are locally isolated, the finite-radius parent lift fails at small R_C , crude infinitesimal condition numbers are too pessimistic, but the finite-amplitude projective endpoint transfer is already contractive on the audited grid at $R_C = 1600 \rightarrow 3200$. This is strong numerical evidence for a parent outgoing-Weyl Cauchy limit in the same family. It is not yet a promoted static width, because the paper still lacks an analytic lower bound for $|u_a(R)u_b(R)|$, or equivalently a parent-side exclusion of the relevant Jost projective pole, uniform over the local argument contours. Until that analytic lift is supplied, the reported strict static surrogate remains $\Gamma_{N,\ell}^{(\text{geo})}$.

Rank-2 truncation justification: Within the present hierarchy convention, the $\hat{\omega}_N^{4\ell+4}$ factor keeps higher- ℓ channels parametrically small for $\hat{\omega}_N < 1$ (our regime). Concretely, for $\hat{\omega}_N \sim 0.1$:

$$\frac{\Gamma_{N,3}}{\Gamma_{N,2}} \sim \hat{\omega}_N^4 \sim 10^{-4}. \quad (37)$$

Thus the $\ell = 1, 2$ truncation still captures the dominant contribution in the demonstrator regime; higher- ℓ channels are subleading and deferred to future work.

Mixing term interpretation: The off-diagonal term ϵ_{mix} can be interpreted as the mode-overlap integral between the $\ell = 1$ and $\ell = 2$ wavefunctions via the non-spherical part of the potential:

$$\epsilon_{\text{mix}} \sim \int \psi_{N,1}^* \delta V \psi_{N,2} d^3x, \quad (38)$$

where δV is the deviation of V_{eff} from spherical symmetry. With the explicit spherical-average definition used in Eq. (73), this overlap is symmetry-suppressed for the parity-even/odd pair and numerically yields $\chi_N^{(\text{sym})} \approx 0$ (Appendix C). Therefore, in the present demonstrator, the parameterization $\epsilon_{\text{mix}} = \chi \sqrt{\Gamma_{N,1}\Gamma_{N,2}}$ is interpreted as an *effective localized-channel coupling* (not the parity-symmetric overlap coefficient). The basis transformation between parity modes and localized modes is a unitary change of representation and is mathematically legitimate; it does not change the underlying operator spectrum. What changes is the channel diagnostic: $\chi_N^{(\text{sym})}$ is the parity-overlap coefficient (symmetry-protected null in Appendix C), while $\chi_N^{(LR)}$

is the localized splitting coefficient used as the effective coarse-grained input (Appendix D). We report both quantities explicitly to avoid basis-mixing ambiguity. Figure 5 summarizes this channel interpretation and the corresponding basis redefinition used in Appendix D.




FIG. 5. Physical interpretation of the mixing channel. In the parity-symmetric basis, overlap cancellation yields $\chi_N^{(\text{sym})} \approx 0$. In the localized basis, the level splitting ΔE defines the effective coupling $\chi_N^{(LR)}(D)$ used in the global scan.

VI. MODULE 3: EFT-NORMALIZED VISIBILITY B_N

In early versions of the theory, B_N was an arbitrary matching factor. In the present EFT-operator-normalized baseline, $B_N(D)$ is defined by fully normalized EFT operator kernels built from overlap-extracted flavor-layer couplings and mediator scales from the same 2D localized operator chain as Appendix D; finite one-loop matching and LL-RG normalization are applied before layer-wise normalization. Mode tracking and microcanonical-window averaging suppress branch-switch artifacts in the overlap extraction stage, while high- N regularization remains isolated in Eq. (18).

For comparison with the earlier visibility closure, we retain a legacy Yukawa-normalized fallback based on the **charged-lepton Yukawa** couplings:

$$Y_1 \equiv y_e, \quad Y_2 \equiv y_\mu, \quad Y_3 \equiv y_\tau, \quad y_\ell = \sqrt{2} m_\ell / v, \quad (39)$$

and define a cumulative effective strength:

$$\tilde{Y}_N \equiv \sum_{k=1}^{\min(N,3)} Y_k, \quad N \geq 1. \quad (40)$$

This cumulative definition restricts the legacy visibility normalization to the three observable generations. As a **legacy comparator/fallback**, we keep the Yukawa-derived visibility law

with sublinear compression exponent $0 < p_B < 1$:

$$B_N = \left(\frac{\tilde{Y}_N}{\tilde{Y}_3} \right)^{p_B}, \quad N = 1, 2, 3, \quad B_3 \equiv 1. \quad (41)$$

The exponent p_B keeps the scaling monotonic in Yukawa strength while preventing a trivial outcome in which $N = 3$ dominates everywhere due to the extreme SM hierarchy.¹ Numerically, using PDG 2024 inputs [8] and $p_B = 0.30$, we obtain representative values $B_1 \simeq 0.085$, $B_2 \simeq 0.42$, $B_3 = 1$ (normalized).

A. High- N Neutral Saturation

For layers $N > 3$, we set $B_N = B_3 = 1$ (neutral saturation). In the baseline scan, high- N control is provided by the bounded two-dimensional phase-space tail prescription and finite-volume clipping in the g_N module; the κ_g factor in Eq. (18) is retained only for legacy Cardy-comparator runs. This keeps the visibility block from introducing an additional high- N damping factor on top of the phase-space regulator.² This saturation should not be read as an ad hoc removal of a fourth layer: it is deliberately neutral, in the sense that once $N > 3$ the visibility block inserts no additional suppression beyond $B_3 = 1$. In the current action-derived benchmark, no audited D point supports a fourth bound level on the single-track spectrum, and the companion one-dimensional kinetic proxy keeps the candidate $N = 4$ excitation above the continuum threshold, $E_4 = \omega_4^2 - m_0^2 > 0$, throughout the tested window. Moreover, forcing the same quasi-channel continuation to $N = 4$ does not produce an additional small-width hierarchy relative to $N = 3$. The theorem below is the current mode-resolved reading: high- N decoupling is a threshold/binding-loss statement on the audited operator, not an engineered visibility cutoff or a hidden $\Gamma_4 \rightarrow 0$ assumption.

Proposition N1 (Mode-resolved high- N threshold theorem on the audited kinetic chain). For fixed D , let

$$H_D = -\frac{d^2}{dz^2} + U(z; D)$$

denote the action-derived one-dimensional kinetic operator in bound-state convention, so that a mode is bound exactly when

$$E_j(D) = \omega_j(D)^2 - m_0^2 < 0.$$

¹ An alternative using up-type quark Yukawas was explored; using lepton Yukawas removes sector ambiguity and enables a direct one-to-one mapping between the $N = 2$ layer and the muon-generation observable.

² An alternative approach with explicit B_N tail suppression ($B_N \propto e^{-\beta_B(N-3)^2}$) was explored; see Appendix F for comparison.

Let $H_D^{(L,h)}$ be the Dirichlet finite-volume tridiagonal discretization on $[-L, L]$, with diagonal entries $a_i = 2h^{-2} + U(z_i; D)$ and off-diagonal entries $b_i = -h^{-2}$. Define the zero-threshold Sturm pivots

$$p_1 = a_1, \quad p_i = a_i - \frac{b_{i-1}^2}{p_{i-1}}, \quad i \geq 2,$$

with the usual limiting interpretation if a pivot vanishes. The number of negative pivots is exactly the number $n_-^{(L,h)}(D)$ of discrete eigenvalues below the continuum threshold $E = 0$. Therefore, if

$$n_-^{(L,h)}(D) \leq 3$$

on the audited D -window, the discretized kinetic chain supports no fourth bound level on that window. If, more sharply,

$$E_3^{(L,h)}(D) > 0,$$

then even the third and fourth entries of the one-dimensional proxy are continuum-threshold entries rather than additional bound states. In either case the neutral saturation $B_{N>3} = 1$ does not hide a fourth-generation visibility cutoff: a candidate $N = 4$ mode is absent already at the mode-resolved threshold level. *Proof.* The finite-volume operator $H_D^{(L,h)}$ is a real symmetric Jacobi matrix. By the min-max principle, its ordered eigenvalues satisfy

$$E_j^{(L,h)}(D) = \min_{\dim V=j} \max_{0 \neq \psi \in V} \frac{\langle \psi, H_D^{(L,h)} \psi \rangle}{\|\psi\|^2}.$$

Hence the number of bound modes below the continuum threshold is the inertia count of $H_D^{(L,h)}$ below zero. The tridiagonal LDL^T factorization gives

$$H_D^{(L,h)} = L \operatorname{diag}(p_1, \dots, p_M) L^T,$$

where the pivots obey the Sturm recursion above. Sylvester's law of inertia then implies that the number of negative pivots equals the number of negative eigenvalues. Equivalently, this is the finite-difference form of the Sturm oscillation theorem: the zero-energy Sturm index counts the modes below threshold. Thus $n_-^{(L,h)} \leq 3$ excludes a fourth bound mode, and $E_3^{(L,h)} > 0$ gives the stronger statement that the one-dimensional proxy has at most two bound modes.

Verified numerics. The certificate audit `code/audit_highN_threshold_theorem.py` exports `output/highN_decoupling/highN_threshold_theorem_detail.csv` and `output/highN_decoupling/highN_threshold_theorem_summary.csv`. On the primary finite-volume certificate

$(L, N_z) = (80, 8001)$, the audited integer window $D = 4, \dots, 20$ has

$$\begin{aligned}\max_D n_-^{(80,h)}(D) &= 2, \\ \min_D E_3^{(80,h)}(D) &= 5.6366674 \times 10^{-3}, \\ \min_D E_4^{(80,h)}(D) &= 5.6379261 \times 10^{-3},\end{aligned}$$

with the worst threshold margin at $D = 4$. The same Sturm count remains $\max n_- = 2$ on all tested box/grid variants $(L, N_z) = (60, 6001), (80, 4001), (80, 12001), (100, 10001)$, and the smallest audited E_4 over those variants is still positive, 4.1701167×10^{-3} . At fixed $L = 80$, refining the grid changes E_3 and E_4 by at most 9.44×10^{-9} and 1.01×10^{-8} , respectively. The independent single-track benchmark is weaker but still sufficient for the actual fourth-layer question: it has $\max n_{\text{bound}} = 3$ and $\max \mathbf{1}_{\{N=4 \text{ bound}\}} = 0$. Finally, the forced $N = 4$ quasi-channel confirms that the missing fourth layer is not a hidden tiny-width hierarchy: on the same one-dimensional chain,

$$\Gamma_{4,\ell=1}^{\text{geo}} / \Gamma_{3,\ell=1}^{\text{geo}} \in [1.0000036, 1.0000125],$$

so the forced $N = 4$ continuation is essentially a continuum continuation of the $N = 3$ rate rather than a separately suppressed bound channel.

Corollary N1 (Status of the high- N line). On the present audited domain, the fourth-layer question is closed as a mode-resolved threshold statement. The closed ingredients are:

1. neutral saturation $B_{N>3} = 1$, which contributes no extra high- N damping;
2. the single-track bound-count audit, which never supports $N = 4$;
3. the one-dimensional Sturm certificate, which gives at most two negative kinetic-proxy modes and positive E_3, E_4 margins on the audited window;
4. the forced quasi-channel check, which shows no additional Γ_4 suppression hierarchy.

What remains is not a new visibility object and not a new high- N damping ansatz. The only possible strengthening is a continuum-tail version of the same Sturm/min-max certificate, replacing the present finite-volume audit by a fully analytic bound on the half-line tails and the D -continuous interpolation between audited knots.

B. Roadmap: Action-Derived Effective Yukawa Operator

With the EFT-operator-normalized baseline now in place, the release-production chain separates direct spectral-selection closure from visibility closure. The runtime direct-spectral branch

recomputes $g_N(D)$, $\chi_{LR}(D)$, and $\tilde{A}_\ell(D)$ at scan runtime while keeping the baseline $B_N(D)$ profile closure; the runtime-direct visibility path is release-qualified only as a risk-weighted profile-blend parity branch, not as a strict all-direct Higgs-map closure. The strict no-cache variant is retained only as a stress diagnostic. The next upgrade remains an explicit lepton Yukawa sector in the same conformal background:

$$\mathcal{L}_Y = -y_0 \bar{L} H \ell_R + \text{h.c.}, \quad g_{\mu\nu} = \Omega^2 \eta_{\mu\nu}. \quad (42)$$

In the minimal conformal setup, we use the vierbein form

$$S_\Psi = \int d^4x \sqrt{-g} \, i \bar{\Psi} e^\mu_a \gamma^a \left(\partial_\mu + \frac{1}{4} \omega_{\mu bc} \gamma^{bc} \right) \Psi, \quad e_\mu^a = \Omega \delta_\mu^a, \quad (43)$$

with $\sqrt{-g} = \Omega^4$ and $e^\mu_a = \Omega^{-1} \delta^\mu_a$. The canonical Weyl rescaling

$$\Psi = \Omega^{-3/2} \psi, \quad H = \Omega^{-1} \varphi \quad (44)$$

removes the spin-connection derivative term and yields canonical flat-frame kinetic terms (Appendix B 3). The Yukawa block then scales as

$$\sqrt{-g} \bar{\Psi} H \Psi \propto \Omega^4 \cdot \Omega^{-3/2} \cdot \Omega^{-1} \cdot \Omega^{-3/2} = \Omega^0, \quad (45)$$

so the minimal baseline has no residual conformal prefactor in the overlap kernel. After fixing this normalization convention, we decompose the Higgs fluctuation on the same action-derived spatial modes used in Appendix D,

$$H(x, \rho, z) = \sum_N h_N(x) u_N(\rho, z; D), \quad (46)$$

and define a mode-resolved effective coupling by overlap:

$$y_N^{\text{eff}}(D) = y_0 \int 2\pi \rho \, d\rho \, dz \, u_N(\rho, z; D) f_L(\rho, z) f_R(\rho, z) \mathcal{W}_{\text{frame}}(\rho, z; D). \quad (47)$$

Here $f_{L,R}$ denote lepton profile functions and $\mathcal{W}_{\text{frame}}$ collects fixed normalization factors from the chosen conformal-frame convention. In the minimal Dirac-conformal baseline implied by Eqs. (44)–(45), $\mathcal{W}_{\text{frame}} = 1$. Nonzero conformal-frame powers are kept only as diagnostic sensitivity directions and are not used in baseline figures. In this roadmap, the same 2D localized solver and convergence criteria as Appendix D are reused; scan-level runs now separate release-production per-cell direct $g_N + \chi + A$ closure from the promoted risk-weighted profile-blend runtime-direct visibility parity branch (canonical D-only profile blend $\alpha(D)$) and the strict no-cache stress branch.

VII. VERIFICATION RESULTS

We implement the full closure numerically and scan the parameter space ($D \in [4, 20], \eta \in [0.2, 4.0]$).

A. Configuration and Parameter Table

Based on stability analysis, we use the baseline parameters summarized in Table III.

TABLE III. Baseline parameters for the PSLT demonstrator. All dimensionful quantities are in units of M_* .

Module	Parameter	Value	Physical Role
Micro-degeneracy g_N	Phase-space	Bounded 2D	Baseline mode: bounded low- N
	profile		window + energy-gap tail profile
	(α, β, s)	$(0.8, 1.1, 0.0)$	low- N window blend, tail damping, shell-power control
	$c_{\text{eff}}, \nu, \kappa_g$	$(0.5, 5.0, 0.03)$	Legacy Cardy-comparator parameters (not used for the bounded 2D baseline normalization)
Visibility B_N	Visibility normalization	EFT operator	Baseline fully normalized EFT visibility from overlap-extracted operator inputs (g_{iN}, λ_N) with finite one-loop + LL-RG normalization
	p_B	0.30	Legacy Yukawa comparator/fallback parameter (inactive in EFT-operator baseline)
Kinetics Γ_N	$\chi_N^{(LR)}(D)$	App. D	Localized-channel profile (action-derived extraction; baseline now uses full scan-grid profile injection, early demonstrator used constant $\chi = 0.2$)
	a_0	0.02	Geometric perturbation
	ϵ	0.2	Core regularization
	$\tilde{A}_1(D), \tilde{A}_2(D)$	App. E 9	Action-derived channel-normalization profile (constant A_ℓ kept as legacy comparator)
Dynamics	t_{coh}	1.0	Coherence time (M_*^{-1})

B. Three-Generation Phase Diagram

The winner map $N^*(D, \eta)$ (Fig. 6) reveals a robust structure:

- For $D \lesssim 8$, Layer 3 ($N = 3$) dominates.
- For $D \gtrsim 8$, Layer 2 ($N = 2$) dominates.
- Layer 1 ($N = 1$) is enhanced by B_1 but is kinetically suppressed at large D .

Crucially, we define the **Generation Ratio**:

$$\mathcal{R}_3 \equiv \sum_{N=1}^3 P_N = \frac{\sum_{N=1}^3 W_N}{\sum_{K=1}^{N_{\max}} W_K}. \quad (48)$$

With full scan-grid action-derived profiles for $\chi_N^{(LR)}(D)$ and $\tilde{A}_\ell(D)$ from Appendix D, our numerical scan gives $\mathcal{R}_3 > 90\%$ over 92.7% of the sampled (D, η) grid, while $f(\mathcal{R}_3 > 0.95) = 0.2167$ in this setup. The profile support is anchored by the benchmark convergence set (Table XII) and the auxiliary full-scan support points (Table XIII) in Appendix D. This still supports a dominant three-layer sector; the current baseline combines action-derived g_N , $\chi_N^{(LR)}(D)$, and $\tilde{A}_\ell(D)$ profiles with EFT-operator-normalized $B_N(D)$. To quantify transfer bias from this surrogate interpolation step, we performed a point-level audit with direct localized fine extraction injected at each queried D (recomputing χ_{LR} by the same 2D solver instead of using interpolated knots). The reproducible script is `scan_surrogate_vs_action_points.py`, with source table `paper/surrogate_vs_action_points.csv`. As summarized in Table IV, the direct-injection path changes χ by up to $\sim 31\%$ at representative points, while winner labels remain unchanged (3/3 matches) and $\mathcal{R}_3$ stays numerically stable ($\max |\Delta \mathcal{R}_3| = 5.27 \times 10^{-11}$); under EFT/Wilson-matched observables, $\mu_{\mu\mu}^{\text{pred}}$ can vary strongly at off-band points, so acceptance-level robustness is additionally tracked by the map summaries and Tables V, XXII.




FIG. 6. Left: Winner phase diagram showing regions dominated by Gen 2 and Gen 3. Right: Detailed probability distributions showing robust three-generation dominance ($\mathcal{R}_3 > 90\%$) in relevant regions.

To propagate finite- N surrogate uncertainty into map-level systematics, we also perform a one-at-a-time local-parameter robustness scan around the baseline closure values: $c_{\text{eff}} \in$

TABLE IV. Point-level surrogate-vs-direct audit for the mixing-profile transfer. “Surrogate” uses $\chi_{LR}(D)$ from $D = \{6, 12, 18\}$ interpolation; “direct” recomputes fine-grid localized extraction at each listed D . Source CSV: `paper/surrogate_vs_action_points.csv`.

(D, η)	winner (sur/direct)	χ_{interp}	χ_{direct}	$ \Delta\mathcal{R}_3 $	$ \Delta\mu_{\mu\mu}^{\text{pred}} $	match
(4.0, 0.2)	2/2	4.02×10^{-4}	5.27×10^{-4}	5.27×10^{-11}	1.58×10^3	yes
(5.0, 2.0)	2/2	4.02×10^{-4}	3.31×10^{-4}	1.02×10^{-14}	9.91×10^2	yes
(11.0, 1.4)	2/2	2.51×10^{-4}	3.06×10^{-4}	0	2.85×10^{-2}	yes

TABLE V. Local one-at-a-time robustness scan around baseline closure parameters, with the p_B row retained only as a legacy visibility-comparator sensitivity because p_B is inactive in the EFT-operator-normalized baseline. Fractions are measured on the same 60×60 map as the baseline scan. Source CSV: `paper/core_param_robustness.csv`.

Parameter	Window	$f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	Max drift
c_{eff}	0.45/0.50/0.55	0.9269/0.9269/0.9269	0.1500/0.1500/0.1500	$\Delta f_{\mathcal{R}_3} = 0, \Delta f_{\chi^2} = 0$
ν	4.80/5.00/5.20	0.9269/0.9269/0.9269	0.1500/0.1500/0.1500	$\Delta f_{\mathcal{R}_3} = 0, \Delta f_{\chi^2} = 0$
p_B	0.28/0.30/0.32	0.9269/0.9269/0.9269	0.1500/0.1500/0.1500	$\Delta f_{\mathcal{R}_3} = 0, \Delta f_{\chi^2} = 0$

$[0.45, 0.55]$, $\nu \in [4.80, 5.20]$, and the legacy visibility-comparator exponent $p_B \in [0.28, 0.32]$ (with the same 60×60 map and full-grid action-derived $\chi_{LR}(D)/\tilde{A}_\ell(D)$ profiles). Results are summarized in Table V. Since p_B is inactive in the EFT-operator-normalized baseline, that row is retained as a legacy-comparator sensitivity check rather than an active closure knob. Across this local window, the key fractions remain unchanged: $f(\mathcal{R}_3 > 0.90) = 0.9269$, $f(\chi_{\mu\mu}^2 < 4) = 0.1500$, and $f(N_{\text{win}} > 3) \approx 2.78 \times 10^{-4}$.

To make the full direct chain auditable at release level, we compare the strict baseline map with two localized-direct validation surfaces and visibility-parity checks; Table VI summarizes the result. The direct spectral branch recomputes the two-dimensional phase-space spectra for $g_N(D)$ and evaluates $\chi_{LR}(D)$ and $\tilde{A}_\ell(D)$ on the active validation grids. On both audit grids it matches the full direct release map exactly: $\Delta f(\chi_{\mu\mu}^2 < 4) = 0$, mismatch fraction = 0, and $\max|\Delta\mu_{\mu\mu}| = 0$. The tuned visibility parity branch keeps a D-only risk-weighted profile blend $\alpha(D)$ with $(\alpha_{\text{min}}, \alpha_{\text{max}}, p) = (0.96, 0.99, 1.0)$, mean $\alpha = 0.9637$, and $p90(\alpha) = 0.9734$. It passes both validation surfaces (small complete grid: $\max|\Delta\mu_{\mu\mu}| = 0.0925$; large spot-check grid: $\max|\Delta\mu_{\mu\mu}| = 0.8373$), but remains explicitly blend-weighted in the observable sector. We therefore read it as a **risk-weighted profile-blend visibility parity branch**, not as a strict all-direct Higgs-map closure. A no-profile stress branch is retained only for diagnostics. The earlier de-profiled visibility stress branch remains useful as a historical negative mechanism audit: it showed that removing the visibility profile too aggressively reopens

TABLE VI. Release-mode full direct map summary; exact source artifacts are listed in the reproducibility registry.

scenario	grid	n_{points}	$f(\chi_{\mu\mu}^2 < 4)$	best $\chi_{\mu\mu}^2$	frac winner mismatch	$\max \Delta\mu_{\mu\mu} $
full direct baseline	D60×E60	3600	0.1428	1.56×10^{-3}	—	—
small-surface direct bias	D21×E41	861	0.1905	1.45×10^{-5}	0	3.37×10^{-3}
large-surface direct bias	D60×E21	1260	0.1500	5.10×10^{-4}	0	0
runtime parity (small)	D21×E41	861	0.0952	2.11×10^{-1}	0	0
tuned parity (small)	D21×E41	861	0.0952	2.11×10^{-1}	0	9.25×10^{-2}
runtime parity (large)	D60×E21	1260	0.1429	1.56×10^{-3}	0	0
tuned parity (large)	D60×E21	1260	0.1429	1.56×10^{-3}	0	8.37×10^{-1}
extreme parity (large)	D60×E21	1260	0.1429	1.56×10^{-3}	0.0373	59.3

TABLE VII. One-at-a-time robustness scan for UV+LL-RG observable controls. Source CSV: `paper/hll_rge_sensitivity.csv`.

parameter	window (low/base/high)	$f(\chi_{\mu\mu}^2 < 4)$	low/base/high	max abs drift
μ_{low}	0.5/1.0/2.0	0.1500/0.1500/0.1500		0
γ_{diag}	1.0/2.0/3.0	0.1500/0.1500/0.1500		0
γ_{offdiag}	0.5/1.0/1.5	0.1500/0.1500/0.1500		0
κ_{diag}	−1.0/0.0/1.0	0.1500/0.1500/0.1500		0
κ_{offdiag}	−1.0/0.0/1.0	0.1500/0.1500/0.1500		0

large observable drift even when some local mismatch gates improve. That negative result has now been superseded operationally by two cleaner endpoints. First, the strict broader-grid all-direct validator is parked at its current case-aware endpoint, with zero acceptance mismatch and $p95|\Delta\mu_{\mu\mu}| = (0.0560, 0.1097)$ on the two validation surfaces. Second, the canonical reviewer observable branch is now the two-lobe full-width/reference-amplitude validator, which has moved from promotion work to routine surveillance. We therefore treat additional local $\alpha(D)$ retunes or de-profiled stress variants as diagnostics only unless a genuinely new observable-side mechanism class is introduced.

For the UV+LL-RG observable controls in the same baseline mode, we also scan one-at-a-time windows in $(\mu_{\text{low}}, \gamma_{\text{diag}}, \gamma_{\text{offdiag}}, \kappa_{\text{diag}}, \kappa_{\text{offdiag}})$. Table VII summarizes the result (source CSV: `paper/hll_rge_sensitivity.csv`); across all tested windows, $f(\chi_{\mu\mu}^2 < 4)$ stays fixed at 0.1500. To capture residual map-level drift in a nonzero finite-match setting, we additionally run a full-direct UV envelope scan with `code/scan_hll_uv_envelope.py`. On the D21×E41 map the acceptance span remains zero ($f_{\text{min}} = f_{\text{max}} = 0.1243$), while the pointwise half-width has mean 3.03×10^{-4} , $p95 = 1.50 \times 10^{-3}$, and max 3.89×10^{-3} (source files: `paper/hll_uv_envelope_summary.csv` and `paper/hll_uv_envelope_map.csv`).

C. $H \rightarrow \mu\mu$ Signal Strength (EFT/Wilson-Matched Mapping, UV+LL-RG Baseline)

We confront the theory with the ATLAS Run-2+Run-3 combined $H \rightarrow \mu\mu$ signal-strength input, $\mu_{\mu\mu}^{\text{obs}} = 1.4 \pm 0.4$ (combined uncertainty) [9]. The same ATLAS evidence result reports a Run-3-only signal strength near $\mu_{\mu\mu} = 1.6 \pm 0.6$; all scan maps in this manuscript use the combined input above.

Within the PSLT closure, we construct a scan-level EFT/Wilson-matched map (evaluated in baseline `eft_wilson_uv_rge` mode) by combining kinetic occupancy with overlap-extracted raw couplings:

$$\tilde{q}_N(D, \eta) = g_N(D) \left(1 - e^{-\Gamma_N(D, \eta) t_{\text{coh}}}\right), \quad P_N^{(\text{kin})}(D, \eta) = \frac{\tilde{q}_N(D, \eta)}{\sum_{K=1}^{N_{\text{max}}} \tilde{q}_K(D, \eta)}. \quad (49)$$

$$\epsilon_{\text{mix}}(D, \eta) = \text{clip} \left[\kappa_{\text{mix}} \chi_{\text{eff}}(D) \left(\frac{\eta}{\eta_{\text{ref}}} \right)^{p_\eta}, 0, \epsilon_{\text{max}} \right], \quad (50)$$

$$\Pi(\epsilon_{\text{mix}}) = \begin{pmatrix} 1 - \epsilon_{\text{mix}} & \epsilon_{\text{mix}} & 0 \\ \epsilon_{\text{mix}}/2 & 1 - \epsilon_{\text{mix}} & \epsilon_{\text{mix}}/2 \\ 0 & \epsilon_{\text{mix}} & 1 - \epsilon_{\text{mix}} \end{pmatrix}, \quad Y_{iN}(D, \eta) = \sqrt{y_N^{\text{eff, raw}}(D)} \Pi_{iN}(\epsilon_{\text{mix}}), \quad (51)$$

The bounded projector in Eqs. (50)–(51) is retained for matched-mode diagnostics; the baseline `eft_wilson_uv_rge` maps use the UV-tree closure below.

$$C_{eH}^{\text{tree, ij}}(D, \eta) = \sum_{N=1}^3 g_{iN}(D) \frac{P_N^{(\text{kin})}(D, \eta)}{M_N^2(D)} g_{jN}(D), \quad (52)$$

$$C_{eH}^{\text{match, ij}} = \begin{cases} \left(1 + \frac{\kappa_{\text{diag}}}{16\pi^2}\right) C_{eH}^{\text{tree, ij}}, & i = j, \\ \left(1 + \frac{\kappa_{\text{offdiag}}}{16\pi^2}\right) C_{eH}^{\text{tree, ij}}, & i \neq j, \end{cases} \quad (53)$$

$$C_{eH}^{\text{IR}} = C_{eH}^{\text{match}} + \frac{\ln(\mu_{\text{match}}/\mu_{\text{low}})}{16\pi^2} \gamma \odot C_{eH}^{\text{match}}, \quad (54)$$

where $\gamma \odot C$ denotes blockwise application of $(\gamma_{\text{diag}}, \gamma_{\text{offdiag}})$ to diagonal and off-diagonal matrix entries, respectively. In the baseline `eft_wilson_uv_rge` mode used for figures, we evaluate observables with $C_{eH} \equiv C_{eH}^{\text{IR}}$.

$$\frac{\Gamma_{\text{tot}}(D, \eta)}{\Gamma_{\text{tot}}(D_0, \eta_0)} = 1 + s_\Gamma \sum_{\ell=e, \mu, \tau} \text{BR}_\ell^{\text{SM}} \left[\left| \frac{C_{eH}^{\ell\ell}(D, \eta)}{C_{eH}^{\ell\ell}(D_0, \eta_0)} \right|^2 - 1 \right], \quad (55)$$

$$\mu_{\mu\mu}^{\text{pred}}(D, \eta) = \left| \frac{C_{eH}^{\mu\mu}(D, \eta)}{C_{eH}^{\mu\mu}(D_0, \eta_0)} \right|^2 \left[\frac{\Gamma_{\text{tot}}(D, \eta)}{\Gamma_{\text{tot}}(D_0, \eta_0)} \right]^{-1}. \quad (56)$$

In the baseline figures we keep a fixed reference point $(D_0, \eta_0) = (10, 1)$ for comparability across scan variants. This fixed point should be read as a *reference-normalization convention* for the relative scan-level observable in Eq. (56), not as a physical anchor of the Higgs sector itself. The stronger normalization-side information now comes from the absolute and loop-improved parent-action comparators below together with the reference-normalization-invariant Wilson-diagonal ratios reported later in the section. Equation (56) is implemented in `pslt_lib.py` as the UV+LL-RG Higgs observable; it remains a scan-level EFT module rather than a full loop-level EYMH→EFT derivation. Writing

$$x_\mu \equiv |C_{eH}^{\mu\mu}|^2, \quad R_{e/\mu} \equiv \frac{|C_{eH}^{ee}|^2}{|C_{eH}^{\mu\mu}|^2}, \quad R_{\tau/\mu} \equiv \frac{|C_{eH}^{\tau\tau}|^2}{|C_{eH}^{\mu\mu}|^2},$$

Eq. (56) may also be rewritten exactly as

$$\mu_{\mu\mu}^{\text{pred}} = \frac{A_\star x_\mu}{1 + s_\Gamma \left[A_\star x_\mu \left(\text{BR}_\mu^{\text{SM}} + \text{BR}_e^{\text{SM}} \frac{R_{e/\mu}}{R_{e/\mu}^\star} + \text{BR}_\tau^{\text{SM}} \frac{R_{\tau/\mu}}{R_{\tau/\mu}^\star} \right) - \sum_{\ell=e,\mu,\tau} \text{BR}_\ell^{\text{SM}} \right]}, \quad (57)$$

with

$$A_\star = x_{\mu,\star}^{-1}, \quad R_{e/\mu}^\star = \frac{|C_{eH,\star}^{ee}|^2}{|C_{eH,\star}^{\mu\mu}|^2}, \quad R_{\tau/\mu}^\star = \frac{|C_{eH,\star}^{\tau\tau}|^2}{|C_{eH,\star}^{\mu\mu}|^2}.$$

So the current reference-normalization arbitrariness is not an unconstrained map-level rescaling; it is exactly equivalent to three scalar bridge constants. The audit `code/audit_hll_absolute_normalization_bridge.py` on the canonical parented D21×E21 map (`output/hll_absolute_normalization/hll_absolute_normalization_bridge_summary.csv`) verifies Eq. (57) to machine precision ($\max|\Delta\mu_{\mu\mu}| = 5.56 \times 10^{-8}$). More importantly, the accepted-region invariant median already fixes the electron-to-muon ratio essentially uniquely: if $R_{e/\mu}^\star$ is frozen to the invariant audit median, a two-parameter fit in $(A_\star, R_{\tau/\mu}^\star)$ still closes the canonical map at the same 5.56×10^{-8} level with zero acceptance mismatch. At this bridge-reduction stage the observable-side problem is therefore already narrower than a generic reference-anchor ambiguity: the live scalar constants are reduced to the absolute amplitude $A_\star$ and the tau-to-muon ratio $R_{\tau/\mu}^\star$, while $R_{e/\mu}^\star$ is effectively fixed by the invariant-ratio audit itself. The next source audit sharpens this split further. Because the present finite-match and LL-RG layers apply the same diagonal dressing to all three charged-lepton diagonals,

$$C_{eH}^{\ell\ell, \text{IR}} = Z_{\text{diag}} C_{eH}^{\ell\ell, \text{tree}}, \quad Z_{\text{diag}} = \left(1 + \frac{\kappa_{\text{diag}}^{\text{eff}}}{16\pi^2} \right) \left(1 + \frac{\gamma_{\text{diag}} \log(\mu_{\text{match}}/\mu_{\text{low}})}{16\pi^2} \right),$$

the flavor ratios $R_{e/\mu}$ and $R_{\tau/\mu}$ are exact UV-tree invariants of the current diagonal chain rather than additional loop-level normalization freedoms. The audits `code/audit_hll_absolute_amp_source.py` and `code/audit_hll_tau_ratio_source.py` (canonical outputs under

`output/hll_absolute_normalization/`) verify this statement on the D21×E21 parented map: on active rows the residuals between UV and IR ratios remain at machine level ($\max |\Delta R_{e/\mu}| = 6.05 \times 10^{-13}$, $\max |\Delta R_{\tau/\mu}| = 3.23 \times 10^{-14}$), and at the live reference point the tau-to-muon constant is identical across tree/match/IR ($R_{\tau/\mu}^* = 8.24039 \times 10^{-3}$). The same audit also shows that the absolute amplitude itself factorizes exactly as

$$A_\star = A_\star^{(\text{tree})} Z_{\text{diag},\star}^{-2}, \quad A_\star^{(\text{tree})} \equiv |C_{eH,\star}^{\mu\mu,\text{tree}}|^{-2},$$

with map-level factorization residual $\max |\Delta C_{\mu\mu}^{\text{fact}}| = 9.32 \times 10^{-17}$. Numerically the remaining gap is now concentrated almost entirely in the reference tree amplitude: $A_\star^{(\text{tree})} = 5.26537 \times 10^8$, $A_\star = 5.20717 \times 10^8$, so the full diagonal dressing shifts the absolute normalization by only about 1.11%. Using the live ratio constants, a tree-only replacement for $A_\star$ already preserves the accepted region exactly, while restoring only the universal LL-RG factor reduces the map-level bridge error to $\max |\Delta \mu_{\mu\mu}| = 1.08 \times 10^{-3}$. Within the current diagonal UV+LL-RG chain, the outstanding parent-side normalization target is therefore primarily the absolute amplitude $A_\star$ itself rather than an unresolved flavor-ratio constant. The next source audit tightens this target once more. In the canonical parented baseline one has `uv_blend` = 0, so the UV-tree coupling matrix reduces to a diagonal overlap form,

$$g_{iN}^{\text{UV}}(D) = \delta_{iN} \sqrt{y_i^{\text{raw}}(D)},$$

and therefore the muon tree coefficient collapses exactly to a single layer:

$$C_{eH}^{\mu\mu,\text{tree}}(D, \eta) = y_2^{\text{raw}}(D) \frac{P_2^{\text{kin}}(D, \eta)}{M_2^2(D)}. \quad (58)$$

The audit `code/audit_hll_tree_mumu_parent_source.py` verifies Eq. (58) on the canonical D21×E21 map with $\max |\Delta C_{\mu\mu}^{\text{tree}}| = 9.76 \times 10^{-17}$ and exact layer support (μ -channel share from $N = 2$ equals 1 on every active row, while the $N = 1, 3$ shares remain below 10^{-26} in magnitude). At the live reference point this gives

$$y_2^{\text{raw}} = 1.65981 \times 10^{-4}, \quad \frac{P_2^{\text{kin}}}{M_2^2} = 2.62559 \times 10^{-1}, \quad C_{eH,\star}^{\mu\mu,\text{tree}} = 4.35798 \times 10^{-5},$$

so the current absolute-normalization target may be written as

$$A_\star = \left(y_2^{\text{raw}}(D_\star) \frac{P_2^{\text{kin}}(D_\star, \eta_\star)}{M_2^2(D_\star)} \right)^{-2} Z_{\text{diag},\star}^{-2}.$$

Across active rows the log-variation of $C_{eH}^{\mu\mu,\text{tree}}$ is dominated by the overlap factor rather than by the kinetic-mass coefficient ($\sigma[\log y_2^{\text{raw}}] = 4.07$, $\sigma[\log(P_2^{\text{kin}}/M_2^2)] = 1.97 \times 10^{-1}$), with corresponding correlations 0.9998 and 0.9100 to $\log C_{eH}^{\mu\mu,\text{tree}}$. This suggests a sharper parent-side task ordering: first derive the reference overlap block $y_2^{\text{raw}}(D_\star)$ from the EYMH/projected-Yukawa

side, then lift the smaller residual coefficient P_2^{kin}/M_2^2 and the already mild universal diagonal dressing. The next source audit identifies this overlap block more precisely. In the canonical overlap extractor the flavor kernels differ only through the chirality-width rescaling $\sigma_{\text{scale}}^{(\ell)}$, and for the muon one has $\sigma_{\text{scale}}^{(\mu)} = 1$ exactly. Therefore the raw overlap and the mu-flavor overlap coincide pointwise:

$$y_2^{\text{raw}}(D) = y_{\mu,2}^{\text{flavor}}(D) = (g_{\mu,2}^{\text{UV}}(D))^2.$$

More explicitly, if $u_k(D)$ are the tracked low-lying modes and $K_\mu(D)$ denotes the mu-flavor left-right overlap kernel, then

$$y_2^{\text{raw}}(D) = \sum_k w_{2k}(D) |\langle u_k(D), K_\mu(D) \rangle|,$$

with

$$K_\mu(D) = f_L^{(\mu)}(D) f_R^{(\mu)}(D) W_{\text{frame}}(D),$$

and where $w_{2k}(D)$ are the normalized Gaussian weights of the local microcanonical window around the tracked second branch. The audit `code/audit_hll_y2raw_parent_source.py` verifies both identities on the extracted overlap profile to machine precision ($\max |y_2^{\text{raw}} - y_{\mu,2}^{\text{flavor}}| = 0$, $\max |(g_{\mu,2}^{\text{UV}})^2 - y_2^{\text{raw}}| = 9.87 \times 10^{-17}$). The same audit also shows that collapsing this block to the center-mode overlap is not adequate: on the audited D4–20 detail grid, the center-only candidate $y_{\text{center},2}$ differs from y_2^{raw} by a mean relative error 8.61×10^{-1} and a 95% relative error 1.14. So the remaining normalization target is better understood as a *windowed projected-Yukawa block* rather than as a single-mode overlap. One more parent-side simplification is already available for the kernel itself. In the canonical baseline the overlap extractor keeps $\sigma_L = \sigma_R \equiv \sigma_\mu = 2.5$ and $W_{\text{frame}} = 1$, so the muon kernel is not merely approximately but exactly a symmetric midplane bridge:

$$K_\mu(\rho, z; D) = A_\mu^{(\text{disc})}(D) \exp\left[-\frac{\rho^2 + z^2}{\sigma_\mu^2}\right].$$

The audit `code/audit_hll_kmu_parent_candidate.py` verifies this pointwise factorization on the audited D4–20 overlap grid with relative sup residual 3.05×10^{-15} . Replacing the discrete normalization by the corresponding finite-box continuum normalization gives a closed-form parent-side candidate

$$A_\mu^{(\text{box})}(D) = \frac{\exp[-D^2/(4\sigma_\mu^2)]}{I_\rho(\sigma_\mu, \rho_{\text{max}}) I_z(\sigma_\mu, D, z_{\text{margin}})},$$

where

$$I_\rho = \pi \sigma_\mu^2 \left(1 - e^{-\rho_{\text{max}}^2/\sigma_\mu^2}\right), \quad I_z = \frac{\sqrt{\pi} \sigma_\mu}{2} \left[\text{erf}\left(\frac{D + z_{\text{margin}}}{\sigma_\mu}\right) + \text{erf}\left(\frac{z_{\text{margin}}}{\sigma_\mu}\right) \right].$$

This finite-box bridge already matches the canonical discrete kernel at the 10^{-4} level (max relative sup error 9.45×10^{-5}). By contrast, the naive infinite-volume amplitude

$$A_\mu^{(\infty)}(D) = \pi^{-3/2} \sigma_\mu^{-3} \exp[-D^2/(4\sigma_\mu^2)]$$

still misses the canonical kernel by about 2.37×10^{-1} . At the live reference point $(D_\star, \eta_\star) = (9.6, 1.0)$ the exact discrete and finite-box amplitudes are

$$A_\mu^{(\text{disc})}(D_\star) = 3.77586 \times 10^{-4}, \quad A_\mu^{(\text{box})}(D_\star) = 3.77621 \times 10^{-4},$$

so the remaining parent-side task is no longer the entire kernel shape. It is better read as the origin of the effective width σ_μ (or its first curvature-controlled correction) inside this midplane bridge block. That effective width now admits a cleaner three-level split. First, because the canonical kernel already closes exactly as

$$\log K_\mu(\rho, z; D) = c_0(D) - \frac{\rho^2 + z^2}{\sigma_\mu^2},$$

a direct log-Hessian fit on the audited D4–20 overlap grid recovers $\sigma_\mu = 2.5$ to machine precision ($\max |\sigma_\mu^{(\text{Hess})}/\sigma_\mu - 1| = 2.49 \times 10^{-15}$, max fit residual 1.85×10^{-13}). Second, a box-normalized moment inversion gives an independent finite-volume check: solving

$$\langle \rho^2 + z^2 \rangle_{K_\mu} = M_2^{(\text{box})}(\sigma_\mu; \rho_{\text{max}}, z_{\text{max}})$$

returns $\sigma_\mu^{(\text{box})} = 2.4999$ with $\max |\sigma_\mu^{(\text{box})}/\sigma_\mu - 1| = 4.09 \times 10^{-5}$. Third, one step closer to the EYMH side, write the action-derived shifted potential near either core as

$$U(\rho, \delta z; D) = U_0(D) + \frac{1}{2} \kappa_{\rho, \mu}(D) \rho^2 + \frac{1}{2} \kappa_{z, \mu}(D) \delta z^2 + \dots, \quad \delta z = z \mp D/2,$$

and define the isotropized core curvature

$$\bar{\kappa}_\mu(D) = \frac{2\kappa_{\rho, \mu}(D) + \kappa_{z, \mu}(D)}{3}.$$

In a harmonic-ground-state reading one expects a Gaussian width law $\sigma_\mu \propto \bar{\kappa}_\mu^{-1/4}$. Calibrating the proportionality constant at the live reference point therefore gives the current best parent-side candidate

$$\sigma_\mu^{(\text{curv})}(D) = c_\sigma \bar{\kappa}_\mu(D)^{-1/4}, \quad c_\sigma = 22.3027.$$

The audit code `audit_hll_sigma_mu_parent_candidate.py` shows that the core Hessian is already almost isotropic on the tested D4–20 window ($\kappa_{\rho, \mu}/\kappa_{z, \mu} \in [0.999996, 0.99999997]$), the resulting curvature-controlled width stays within 1.01×10^{-3} of the canonical $\sigma_\mu = 2.5$, and the induced finite-box kernel deformation remains small (max relative sup error 1.11×10^{-2}).

So the next parent-side target is now even narrower: not the entire kernel width profile, but the single calibration constant c_σ or its first curvature correction. One more narrowing step is already available for that constant itself. The audit `code/audit_hll_csigma_source.py` compares three global selectors for c_σ : direct width matching,

$$\arg \min_c \sum_D [\log \sigma_\mu^{\text{exact}}(D) - \log(c \bar{\kappa}_\mu(D)^{-1/4})]^2,$$

finite-box midplane amplitude matching,

$$\arg \min_c \sum_D [\log A_\mu^{\text{disc}}(D) - \log A_\mu^{\text{box}}(D; c)]^2,$$

and full-kernel matching,

$$\arg \min_c \max_D \text{reldsup}(K_\mu^{\text{exact}}(D), K_\mu^{\text{box}}(D; c)).$$

All three selectors collapse onto the same narrow interval:

$$c_\sigma^{\text{ref}} = 22.3033, \quad c_\sigma^{\text{width}} = 22.3025, \quad c_\sigma^{\text{amp}} = c_\sigma^{\text{kernel}} = 22.3100,$$

with relative drifts only 3.8×10^{-5} (width) and 3.0×10^{-4} (amplitude/kernel) from the reference calibration. The finite-box core-to-box choice $c_\sigma = 22.3100$ is the most useful parent-side reading at present: it improves the maximum amplitude mismatch from 1.02×10^{-2} to 1.88×10^{-3} and simultaneously lowers the maximum finite-box kernel defect from 1.02×10^{-2} to 1.88×10^{-3} . In that sense the surviving ambiguity is no longer whether a stable c_σ exists, but why this nearly unique finite-box bridge constant emerges from the EYMH / projected-Yukawa side. That bridge constant is now supported by two independent overlap-side checks rather than by kernel matching alone. First, the pointwise core-to-box inversion in `code/audit_hll_csigma_core_box_local.py` solves

$$c_\sigma^{(\text{box}, \text{local})}(D) = \arg \min_c \left| \log A_\mu^{\text{box}}(D; c) - \log A_\mu^{\text{disc}}(D) \right|$$

on the audited D4-20 overlap profile. Apart from the known near-merger tracking outlier at $D = 5$, the recovered local constants stay on the same narrow band, with mean 22.3033, minimum 22.2825, maximum 22.3100, and relative span 1.23×10^{-3} . Second, the projected-Yukawa benchmark `code/audit_hll_csigma_projected_overlap.py` re-solves the fine overlap block at the representative points $D = \{6, 12, 18\}$ using the same shift-invert target as the original overlap extractor and verifies that the canonical rebuild is exact to machine precision (max relative rebuild error 1.27×10^{-13}). On the same benchmark, the projected-overlap selector

$$\arg \min_c \sum_D [\log y_2^{\text{pred}}(D; c) - \log y_2^{\text{raw}}(D)]^2$$

lands at

$$c_{\sigma}^{(\text{projected})} = 22.3095,$$

with only 1.60×10^{-4} relative drift from c_{σ}^{ref} and 2.24×10^{-5} from c_{σ}^{amp} ; the corresponding pointwise selector values 22.2665, 22.3055, and 22.3100 reproduce y_2^{raw} with maximum relative error 1.27×10^{-4} . So the current evidence is stronger than a single-kernel fit: both the finite-box core-to-box inversion and the projected-Yukawa overlap benchmark now select the same $c_{\sigma} \approx 22.31$ bridge band. One more step now identifies the body of that bridge constant itself. The audit `code/audit_hll_csigma_eymh_core_source.py` expands the EYMH-side shifted potential

$$U(\rho, \delta z; D) = m_0^2(\Omega^2 - 1) + (1 - 6\xi) \frac{\nabla^2 \Omega}{\Omega}$$

directly around one core, first in the isolated one-center Plummer limit and then in the full two-center local jet. In the one-center limit,

$$\Omega_{\text{self}}(r) = 1 + \frac{a}{\sqrt{r^2 + \varepsilon^2}} = \Omega_{0,\text{self}} + \Omega_{2,\text{self}} r^2 + \dots, \quad \Omega_{0,\text{self}} = 1 + \frac{a}{\varepsilon}, \quad \Omega_{2,\text{self}} = -\frac{a}{2\varepsilon^3},$$

while

$$\nabla^2 \Omega_{\text{self}}(r) = L_{0,\text{self}} + L_{2,\text{self}} r^2 + \dots, \quad L_{0,\text{self}} = -\frac{3a}{\varepsilon^3}, \quad L_{2,\text{self}} = \frac{15a}{2\varepsilon^5}.$$

So the isotropic one-center curvature source is already explicit:

$$\kappa_{\text{self}} = 2 \left[m_0^2 (2\Omega_{0,\text{self}} \Omega_{2,\text{self}}) + (1 - 6\xi) \left(\frac{L_{2,\text{self}}}{\Omega_{0,\text{self}}} - \frac{L_{0,\text{self}} \Omega_{2,\text{self}}}{\Omega_{0,\text{self}}^2} \right) \right].$$

With the canonical EYMH parameters $(a, \varepsilon, m_0, \xi) = (0.04, 0.1, 1, 0.14)$ this gives

$$\kappa_{\text{self}} = 6353.3061, \quad c_{\sigma}^{(\text{self})} = \sigma_{\mu} \kappa_{\text{self}}^{1/4} = 22.3198,$$

already within 1.78×10^{-3} of the full extracted bridge band. The same audit then includes the second center analytically in the local jet by expanding the distant term around $\delta z = z - D/2 = 0$ and reconstructing the full quadratic coefficients of both Ω and $\nabla^2 \Omega / \Omega$. This produces an analytic two-center curvature $\bar{\kappa}_{\mu}^{(\text{analytic})}(D)$ that tracks the finite-difference EYMH Hessian to 1.68×10^{-4} relative error and therefore gives

$$c_{\sigma}^{(\text{analytic})}(D) = \sigma_{\mu} \bar{\kappa}_{\mu}^{(\text{analytic})}(D)^{1/4}$$

with maximum relative error only 4.21×10^{-5} on the audited D4–20 window. Equally importantly, the mirror correction remains parametrically small:

$$\frac{\delta \kappa_{\text{mirror}}(D)}{\kappa_{\text{self}}} = \frac{\bar{\kappa}_{\mu}^{(\text{analytic})}(D) - \kappa_{\text{self}}}{\kappa_{\text{self}}} \in [-6.91 \times 10^{-3}, -1.39 \times 10^{-3}],$$

so the surviving D-dependence is genuinely a subleading two-center correction to an essentially one-center EYMH core source. In that sense the mainline bridge problem is now narrower again: ‘why does c_σ exist?’ is no longer the right question; the more precise question is why the one-center Plummer/EYMH core sets $c_\sigma \approx 22.32$ and how the small mirror correction refines it to the observed 22.30 band.

Proposition U5 (Core-to-projected normalization bridge). The remaining c_σ ambiguity admits a direct normalization-side closure on the projected-Yukawa benchmark. Define the box-normalized projected overlap functional

$$\mathcal{N}_{\mu,2}^{(\text{proj-box})}(D; c) := \sum_k w_{2k}(D) \left| \left\langle u_k(D), K_\mu^{\text{box}}(D; c) \right\rangle \right|,$$

where the representative projected selector already gave

$$c_\sigma^{(\text{projected})} = 22.3095$$

on $D = \{6, 12, 18\}$. Then the one-center EYMH / Plummer source

$$c_\sigma^{(\text{self})} = 22.3198,$$

already fixes the bulk of the projected-Yukawa box normalization, while the pointwise analytic mirror refinement

$$c_\sigma^{(\text{analytic})}(D) = \sigma_\mu \bar{\kappa}_\mu^{(\text{analytic})}(D)^{1/4}$$

saturates the same benchmark. *Verified numerics.* The audit `code/audit_hll_csigma_projected_box_source.py` inserts these EYMH candidates directly into $\mathcal{N}_{\mu,2}^{(\text{proj-box})}$. The one-center source lies only 4.61×10^{-4} above the projected selector and already reproduces the overlap block with y_2 log-RMSE 6.15×10^{-3} , maximum relative y_2 error 9.26×10^{-3} , and mean relative y_2 error 4.93×10^{-3} . Once the analytic mirror correction is included pointwise through $c_\sigma^{(\text{analytic})}(D)$, the same functional improves to log-RMSE 6.29×10^{-4} with maximum relative y_2 error 9.89×10^{-4} , slightly outperforming the fixed projected selector itself (log-RMSE 9.08×10^{-4} , maximum relative error 1.45×10^{-3}). So the normalization statement is now sharper than “the one-center core sets 22.3198”: the one-center Plummer/EYMH core already fixes the bulk of the projected-Yukawa box normalization, and the small two-center mirror correction is enough to saturate the projected benchmark. That remaining gap also now has a clean first-correction audit. Define the bridge constant

$$c_\sigma(D) = \sigma_\mu \bar{\kappa}_\mu(D)^{1/4},$$

which on the audited D4–20 grid is already almost D -invariant: the audit `code/audit_hll_sigma_mu_curvature_correction.py` finds

$$c_\sigma^{\text{ref}} = 22.3027, \quad \frac{\max c_\sigma - \min c_\sigma}{\langle c_\sigma \rangle} = 1.38 \times 10^{-3}.$$

To probe the first non-harmonic correction, expand the same core potential one order further,

$$U(\rho, \delta z; D) = U_0(D) + \frac{1}{2}\kappa_{\rho,\mu}(D)\rho^2 + \frac{1}{2}\kappa_{z,\mu}(D)\delta z^2 + \lambda_{\rho,\mu}(D)\rho^4 + \lambda_{z,\mu}(D)\delta z^4 + \dots,$$

and introduce the isotropized quartic invariant

$$\lambda_\mu(D) = \frac{2\lambda_{\rho,\mu}(D) + \lambda_{z,\mu}(D)}{3}, \quad q_{4,\mu}(D) = \frac{\lambda_\mu(D)}{\bar{\kappa}_\mu(D)^{3/2}}.$$

The current best corrected EYMH-side candidate is then

$$\sigma_\mu^{(\text{curv},4)}(D) = c_{\sigma,*} \bar{\kappa}_\mu(D)^{-1/4} \left[1 + b_4 (q_{4,\mu}(D) - q_{4,\mu}(D_\star)) \right], \quad b_4 = 0.4244.$$

Numerically this single quartic correction collapses the remaining width drift from 1.01×10^{-3} to 8.47×10^{-7} , and pushes the finite-box kernel defect down from 1.11×10^{-2} to 9.47×10^{-5} , essentially saturating the residual box-normalization error. In that sense the current bridge problem is now very narrow: the surviving EYMH-side ambiguity is no longer the width profile, but the source of the nearly constant bridge coefficient c_σ and, at the next order, the quartic-response coefficient b_4 .

Proposition U6 (Absolute-normalization synthesis chain). On the current diagonal UV+LL-RG chain, the observable-side absolute normalization reduces exactly to the reference overlap block and a structurally fixed finite-box bridge:

$$A_\star = Z_{\text{diag},\star}^{-2} \left[y_2^{\text{raw}}(D_\star) \frac{P_2^{\text{kin}}(D_\star, \eta_\star)}{M_2^2(D_\star)} \right]^{-2},$$

with $R_{e/\mu}^\star$ already fixed by the accepted-region invariant audit for practical purposes and $R_{\tau/\mu}^\star$ an exact UV-tree invariant of the current diagonal chain. Moreover,

$$y_2^{\text{raw}}(D) = y_{\mu,2}^{\text{flavor}}(D) = (g_{\mu 2}^{\text{UV}}(D))^2,$$

and the corresponding overlap kernel closes exactly as

$$K_\mu(\rho, z; D) = A_\mu^{(\text{disc})}(D) \exp \left[-\frac{\rho^2 + z^2}{\sigma_\mu^2} \right].$$

Thus the only remaining parent-side source ambiguity is the EYMH origin of the nearly constant bridge coefficient c_σ in the structurally fixed width law

$$\sigma_\mu(D) = c_\sigma(D) \bar{\kappa}_\mu(D)^{-1/4},$$

or, after the first quartic tightening,

$$\sigma_\mu^{(\text{curv},4)}(D) = c_{\sigma,\star} \bar{\kappa}_\mu(D)^{-1/4} \left[1 + b_4(q_{4,\mu}(D) - q_{4,\mu}(D_\star)) \right].$$

On the representative projected-Yukawa benchmark $D = \{6, 12, 18\}$, this residual bridge constant is already fixed in bulk by the one-center source $c_\sigma^{(\text{self})}$ and saturated by the analytic two-center refinement $c_\sigma^{(\text{analytic})}(D)$. *Audit synthesis.* The new consolidation audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_hll_absolute_normalization_synthesis_source.py` reads the existing normalization-side summaries directly and confirms that the chain now separates into exact-closure stages and secondary-tightening stages. On the exact side, the diagonal flavor ratios remain UV-tree invariants to machine precision ($\max |\Delta R_{e/\mu}^{\text{UV} \rightarrow \text{IR}}| = 6.05 \times 10^{-13}$, $\max |\Delta R_{\tau/\mu}^{\text{UV} \rightarrow \text{IR}}| = 3.23 \times 10^{-14}$); the amplitude factorization $A_\star = A_\star^{(\text{tree})} Z_{\text{diag},\star}^{-2}$ closes with residual 9.32×10^{-17} ; the single-layer reduction $C_{\mu\mu}^{\text{tree}} = y_2^{\text{raw}}(P_2^{\text{kin}}/M_2^2)$ closes with residual 9.76×10^{-17} and exact $N = 2$ support; the overlap identity $y_2^{\text{raw}} = (g_{\mu 2}^{\text{UV}})^2$ closes with residual 9.87×10^{-17} ; and the Gaussian midplane kernel factorization closes with relative sup residual 3.05×10^{-15} . On the structural-but-not-final side, the finite-box Gaussian bridge already stays at 9.45×10^{-5} relative sup error, the curvature width law is within 1.01×10^{-3} of the canonical width, and the quartic correction collapses that drift to 8.47×10^{-7} while lowering the finite-box kernel defect to 9.47×10^{-5} . Excluding the known near-merger outlier at $D = 5$, the pointwise core-to-box inversion keeps c_σ on the narrow band $[22.2825, 22.3100]$ with relative span 1.23×10^{-3} . The one-center EYMH source itself gives $c_\sigma^{(\text{self})} = 22.3198$, only 7.37×10^{-4} above the reference selector and 4.61×10^{-4} above the projected selector, while the analytic two-center jet tracks the extracted $c_\sigma(D)$ with maximum relative error 4.21×10^{-5} . Finally, on the projected-box benchmark, the one-center source alone reproduces the overlap block with log-RMSE 6.15×10^{-3} and maximum relative y_2 error 9.26×10^{-3} , whereas the analytic mirror refinement improves this to log-RMSE 6.29×10^{-4} and maximum relative error 9.89×10^{-4} , slightly better than the fixed projected selector itself.

Corollary U6 (What is closed, and what remains only as secondary tightening, for $A_\star$). On the current audited domain, the following normalization-side blocks are already structurally closed:

1. the flavor-ratio side $(R_{e/\mu}^\star, R_{\tau/\mu}^\star)$, with $R_{e/\mu}^\star$ practically fixed by the invariant audit and $R_{\tau/\mu}^\star$ exact along the diagonal UV chain;
2. the exact amplitude factorization $A_\star = A_\star^{(\text{tree})} Z_{\text{diag},\star}^{-2}$;
3. the single-layer tree reduction $C_{\mu\mu}^{\text{tree}} = y_2^{\text{raw}}(P_2^{\text{kin}}/M_2^2)$;

4. the overlap identity $y_2^{\text{raw}} = y_{\mu,2}^{\text{flavor}} = (g_{\mu 2}^{\text{UV}})^2$;
5. the Gaussian midplane kernel shape $K_\mu = A_\mu^{(\text{disc})} \exp[-(\rho^2 + z^2)/\sigma_\mu^2]$;
6. the structural width profile $\sigma_\mu \propto \bar{\kappa}_\mu^{-1/4}$ together with its first quartic tightening.

Only secondary tightening remains in:

$$c_\sigma, \quad \delta c_\sigma^{(\text{mirror})}, \quad b_4,$$

namely in the one-center EYMH normalization that fixes the bridge coefficient, the already small analytic two-center mirror correction, and the first quartic-response coefficient. Therefore the mainline task is no longer to discover a new overlap object, a new kernel shape, or a new support reduction for $A_\star$. The chain is already structurally closed; what remains, if any further improvement is desired, is only to formalize the one-center EYMH source of c_σ and, secondarily, to tighten its mirror/quartic corrections.

Proposition U7 (c_σ closes as a one-center EYMH constant plus a tiny analytic mirror multiplier). Let

$$c_\sigma^{(\text{exact})}(D) = \sigma_\mu \bar{\kappa}_\mu^{(\text{FD})}(D)^{1/4}, \quad c_\sigma^{(\text{self})} = \sigma_\mu \kappa_{\text{self}}^{1/4},$$

and define the analytic two-center curvature refinement by

$$\delta \kappa_{\text{mirror}}(D) = \bar{\kappa}_\mu^{(\text{analytic})}(D) - \kappa_{\text{self}}.$$

Then the analytic EYMH bridge coefficient satisfies the exact factorization

$$c_\sigma^{(\text{analytic})}(D) = \sigma_\mu \bar{\kappa}_\mu^{(\text{analytic})}(D)^{1/4} = c_\sigma^{(\text{self})} \left(1 + \frac{\delta \kappa_{\text{mirror}}(D)}{\kappa_{\text{self}}} \right)^{1/4}.$$

Equivalently,

$$c_\sigma^{(\text{analytic})}(D) = c_\sigma^{(\text{self})} \left[1 + \frac{1}{4} \frac{\delta \kappa_{\text{mirror}}(D)}{\kappa_{\text{self}}} + O\left(\left(\frac{\delta \kappa_{\text{mirror}}(D)}{\kappa_{\text{self}}} \right)^2 \right) \right],$$

so the entire D -dependence of the bridge coefficient is carried by a parametrically small mirror multiplier on top of the one-center EYMH core constant. *Verified numerics.* The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_hll_csigma_tightening_synthesis.py` consolidates the current c_σ -side summaries and confirms that this is no longer merely a heuristic reading. The extracted bridge coefficient itself already stays on the narrow D4–20 band

$$22.2802 \leq c_\sigma^{(\text{exact})}(D) \leq 22.3111, \quad \frac{\max c_\sigma^{(\text{exact})} - \min c_\sigma^{(\text{exact})}}{\langle c_\sigma^{(\text{exact})} \rangle} = 1.38 \times 10^{-3}.$$

The one-center source is constant,

$$c_{\sigma}^{(\text{self})} = 22.3198,$$

and the analytic mirror factorization closes to machine precision,

$$\max_D \left| c_{\sigma}^{(\text{analytic})}(D) - c_{\sigma}^{(\text{self})} \left(1 + \frac{\delta\kappa_{\text{mirror}}(D)}{\kappa_{\text{self}}} \right)^{1/4} \right| = 3.55 \times 10^{-15}.$$

Moreover,

$$\frac{\delta\kappa_{\text{mirror}}(D)}{\kappa_{\text{self}}} \in [-6.91 \times 10^{-3}, -1.39 \times 10^{-3}],$$

so the mirror multiplier itself lies in the very narrow interval

$$0.998269 \leq \left(1 + \frac{\delta\kappa_{\text{mirror}}(D)}{\kappa_{\text{self}}} \right)^{1/4} \leq 0.999653,$$

with relative span 1.38×10^{-3} . Even the first-order linearization already captures that multiplier to

$$\max_D \text{relerr} \left(1 + \frac{1}{4} \frac{\delta\kappa_{\text{mirror}}}{\kappa_{\text{self}}}, \left(1 + \frac{\delta\kappa_{\text{mirror}}}{\kappa_{\text{self}}} \right)^{1/4} \right) = 4.50 \times 10^{-6}.$$

Against the extracted finite-difference bridge coefficient, the one-center constant alone stays within

$$\max_D \text{relerr} \left(c_{\sigma}^{(\text{self})}, c_{\sigma}^{(\text{exact})} \right) = 1.78 \times 10^{-3},$$

while the analytic mirror-refined coefficient improves this immediately to

$$\max_D \text{relerr} \left(c_{\sigma}^{(\text{analytic})}, c_{\sigma}^{(\text{exact})} \right) = 4.21 \times 10^{-5}.$$

The same tightening is visible on the projected-Yukawa box benchmark: the one-center source gives maximum relative y_2 error 9.26×10^{-3} and log-RMSE 6.15×10^{-3} , whereas the analytic mirror refinement lowers these to 9.89×10^{-4} and 6.29×10^{-4} , i.e. improvement factors 9.37 and 9.79, respectively.

Corollary U7 (What remains secondary after the c_{σ} tightening). On the audited normalization chain, the mainline c_{σ} problem is no longer an open support-selection or kernel-shape problem. It is already closed, to current numerical precision, by:

1. the one-center EYMH core constant $c_{\sigma}^{(\text{self})}$, which sets the bulk normalization scale;
2. the exact analytic mirror multiplier $(1 + \delta\kappa_{\text{mirror}}/\kappa_{\text{self}})^{1/4}$, which carries the residual D -dependence;

3. the projected-box benchmark, which is already saturated by the analytic refinement at sub- 10^{-3} maximum relative y_2 error.

Accordingly, the remaining c_σ -side work is now only secondary tightening:

$$b_4, \quad \text{and at most sub-}10^{-3} \text{ projected-box polish.}$$

No new overlap object, no new support reduction, and no new kernel ansatz is needed for the normalization mainline.

Proposition U8 (The quartic response is already a single stable scalar tightening).

Once the bridge coefficient has been reduced to the one-center EYMH core plus the tiny analytic mirror multiplier, the remaining width-side correction is still controlled by the original isotropized quartic invariant

$$q_{4,\mu}(D) = \frac{\lambda_\mu(D)}{\bar{\kappa}_\mu(D)^{3/2}}, \quad \lambda_\mu(D) = \frac{2\lambda_{\rho,\mu}(D) + \lambda_{z,\mu}(D)}{3},$$

with no need for any new overlap block or support refinement. Writing

$$\sigma_\mu^{(\text{curv},0)}(D) = c_{\sigma,\star} \bar{\kappa}_\mu(D)^{-1/4},$$

the first non-harmonic correction is the one-parameter family

$$\sigma_\mu^{(\text{curv},4)}(D) = \sigma_\mu^{(\text{curv},0)}(D) \left[1 + b_4 (q_{4,\mu}(D) - q_{4,\mu}(D_\star)) \right].$$

The optimal quartic-response coefficient is therefore the single scalar

$$b_4 = \arg \min_b \sum_D \left[\sigma_\mu^{\text{exact}}(D) - \sigma_\mu^{(\text{curv},0)}(D) \left(1 + b (q_{4,\mu}(D) - q_{4,\mu}(D_\star)) \right) \right]^2.$$

Verified numerics. The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_hll_b4_tightening_synthesis.py`, together with the underlying curvature-correction detail table, shows that this last correction is already exhausted by one global scalar. The fitted value is

$$b_4 = 0.4244007048874373,$$

with the quartic coordinate itself varying only over the narrow range

$$-1.0410744 \leq q_{4,\mu}(D) \leq -1.0378108, \quad \frac{\max q_{4,\mu} - \min q_{4,\mu}}{\langle q_{4,\mu} \rangle} = 3.14 \times 10^{-3}.$$

Using no quartic correction, the width candidate already has

$$\max_D \text{relerr}(\sigma_\mu^{(\text{curv},0)}, \sigma_\mu^{\text{exact}}) = 1.01 \times 10^{-3},$$

while the single fitted quartic correction lowers this immediately to

$$\max_D \text{relerr}(\sigma_\mu^{(\text{curv},4)}, \sigma_\mu^{\text{exact}}) = 8.47 \times 10^{-7},$$

an improvement factor 1.19×10^3 . The same scalar also lowers the finite-box kernel defect from

$$1.11 \times 10^{-2} \quad \text{to} \quad 9.47 \times 10^{-5},$$

an improvement factor 1.17×10^2 , essentially saturating the residual box floor.

This fitted scalar is also stable under coarse re-fits of the audited D-window. Re-fitting on the outer window $D \leq 8$ or $D \geq 12$ changes the coefficient by only

$$3.98 \times 10^{-6}$$

relative to the full fit, and even excluding the known near-merger point $D = 5$ changes it by only

$$6.27 \times 10^{-5}.$$

More aggressively splitting the window still leaves only permille-level drift:

$$\frac{|b_4^{\text{left}} - b_4|}{b_4} = 5.35 \times 10^{-4}, \quad \frac{|b_4^{\text{right}} - b_4|}{b_4} = 1.25 \times 10^{-3}.$$

So on the current audited domain, b_4 is already not a new functional degree of freedom, but a single stable scalar response coefficient on a fixed quartic coordinate.

Corollary U8 (After U7, only scalar-level quartic polish remains). On the normalization mainline, the b_4 problem is no longer a search for a new quartic observable, a new support choice, or a new kernel class. It has already reduced to one global scalar on the fixed coordinate $q_{4,\mu}(D) - q_{4,\mu}(D_\star)$, and that scalar already saturates the audited width-side residuals:

1. the uncorrected curvature width law is at the 10^{-3} level;
2. the single fitted quartic coefficient lowers that to 8.47×10^{-7} ;
3. the same coefficient simultaneously saturates the finite-box kernel defect at 9.47×10^{-5} ;
4. subset re-fits change b_4 only at the 10^{-6} to 10^{-3} relative level.

Therefore the remaining normalization-side work after U7 is truly only polish:

either keep $b_4 = 0.4244007$ as final on the audited domain, or pursue only sub- 10^{-3} benchmark cleanup.

No further structural expansion of the $A_\star$ chain is needed.

Proposition U9 (Reviewer-facing closure of the $A_\star$ normalization line). Combining Propositions U6–U8, the current diagonal UV+LL-RG normalization line has a closed parent-side reading on the audited domain:

$$\mu_{\mu\mu}^{\text{pred}} = \frac{A_\star x_\mu}{1 + s_\Gamma \left[A_\star x_\mu \left(\text{BR}_\mu^{\text{SM}} + \text{BR}_e^{\text{SM}} \frac{R_{e/\mu}}{R_{e/\mu}^\star} + \text{BR}_\tau^{\text{SM}} \frac{R_{\tau/\mu}}{R_{\tau/\mu}^\star} \right) - \sum_{\ell=e,\mu,\tau} \text{BR}_\ell^{\text{SM}} \right]},$$

where the two ratio constants are not independent normalization freedoms in the present diagonal chain: $R_{e/\mu}^\star$ is fixed by the invariant-ratio audit at the accepted-region level, and $R_{\tau/\mu}^\star$ is an exact UV-tree invariant. The remaining scalar amplitude is

$$A_\star = Z_{\text{diag},\star}^{-2} \left[y_2^{\text{raw}}(D_\star) \frac{P_2^{\text{kin}}(D_\star, \eta_\star)}{M_2^2(D_\star)} \right]^{-2},$$

with

$$y_2^{\text{raw}}(D) = y_{\mu,2}^{\text{flavor}}(D) = (g_{\mu 2}^{\text{UV}}(D))^2 = \sum_k w_{2k}(D) |\langle u_k(D), K_\mu(D) \rangle|.$$

The parent-side kernel entering this last overlap is already fixed to the finite-box Gaussian midplane bridge

$$K_\mu(\rho, z; D) = A_\mu^{(\text{box})}(D) \exp \left[-\frac{\rho^2 + z^2}{\sigma_\mu(D)^2} \right]$$

up to the audited box-normalization floor. Its width is fixed, to the current parent-side accuracy, by the EYMH curvature jet

$$\sigma_\mu(D) = c_\sigma^{(\text{analytic})}(D) \bar{\kappa}_\mu(D)^{-1/4} \left[1 + b_4(q_{4,\mu}(D) - q_{4,\mu}(D_\star)) \right],$$

where

$$c_\sigma^{(\text{analytic})}(D) = c_\sigma^{(\text{self})} \left(1 + \frac{\delta \kappa_{\text{mirror}}(D)}{\kappa_{\text{self}}} \right)^{1/4}, \quad c_\sigma^{(\text{self})} = \sigma_\mu \kappa_{\text{self}}^{1/4}.$$

Thus the $A_\star$ line no longer contains an open support choice, a new kernel family, or an independent observable-side reference anchor. It has reduced to the already-audited windowed mu-flavor overlap block, with the one-center EYMH core setting the bulk normalization and the two-center mirror and quartic terms acting only as secondary corrections. *Verified numerics.* The exact pieces close at machine or structural precision: $\max |\Delta R_{e/\mu}^{\text{UV} \rightarrow \text{IR}}| = 6.05 \times 10^{-13}$, $\max |\Delta R_{\tau/\mu}^{\text{UV} \rightarrow \text{IR}}| = 3.23 \times 10^{-14}$, the amplitude-factorization residual is 9.32×10^{-17} , the single-layer tree residual is 9.76×10^{-17} , the identity $y_2^{\text{raw}} = (g_{\mu 2}^{\text{UV}})^2$ closes with residual 9.87×10^{-17} , and the exact Gaussian kernel factorization has relative sup residual 3.05×10^{-15} . The finite-box bridge floor is 9.45×10^{-5} . The uncorrected curvature width is already within 1.01×10^{-3} , the quartic correction lowers this to 8.47×10^{-7} , and the corrected finite-box kernel defect is

9.47×10^{-5} . For the bridge coefficient, the one-center core gives $c_\sigma^{(\text{self})} = 22.3198$, the extracted local band excluding the known $D = 5$ outlier is $[22.2825, 22.3100]$, the analytic two-center jet tracks $c_\sigma(D)$ with maximum relative error 4.21×10^{-5} , and the projected-box benchmark improves from one-center maximum relative y_2 error 9.26×10^{-3} to 9.89×10^{-4} after the analytic mirror correction. The remaining quartic scalar is stable at $b_4 = 0.4244007048874373$, with subset refits changing it only from 10^{-6} to 10^{-3} in relative terms.

Corollary U9 (Roadmap status after the $A_\star$ synthesis). On the present audited domain, the $H \rightarrow \mu\mu$ absolute-normalization line should be treated as reviewer-facing closed at the synthesis level. The closed blocks are:

1. the reference-normalization rewrite into $(A_\star, R_{e/\mu}^\star, R_{\tau/\mu}^\star)$;
2. the diagonal-chain invariance of $R_{e/\mu}^\star$ and $R_{\tau/\mu}^\star$;
3. the exact factorization $A_\star = A_\star^{(\text{tree})} Z_{\text{diag}, \star}^{-2}$;
4. the exact single-layer reduction through $y_2^{\text{raw}}(P_2^{\text{kin}}/M_2^2)$;
5. the exact identity of y_2^{raw} with the mu-flavor UV overlap block;
6. the finite-box Gaussian midplane kernel shape and curvature/quartic width profile;
7. the one-center EYMH source plus analytic mirror refinement for the bridge coefficient.

Only secondary tightening remains: a cleaner prose derivation of the continuum one-center constant, a sharper analytic bound on the mirror remainder, and optional sub- 10^{-3} projected-box polish. None of these should be treated as a new support/object search. In roadmap terms, the normalization mainline is closed on the audited domain; the separate high- N threshold line is now addressed by Proposition N1 / Corollary N1, leaving only optional continuum-tail strengthening there.

Proposition U1 (UV-to-EFT witness). On the audited map, the scan-level chain

$$C_{eH}^{\text{tree}} \longrightarrow C_{eH}^{\text{match}} \longrightarrow C_{eH}^{\text{IR}}$$

admits an explicit operator-basis reconstruction with numerically controlled deformation relative to the UV-tree baseline. *Verified numerics.* The reproducible audit `scan_hll_uv_to_eft_matching.py` (canonical outputs `paper/hll_uv_to_eft_map.csv` and `paper/hll_uv_to_eft_summary.csv`) shows that in the baseline UV settings (`uv_blend = 0`, `uv_m2_power = 1`, $\kappa_{\text{diag}} = \kappa_{\text{offdiag}} = 0$) the finite-match step is numerically null at map level

($\max |\Delta C_{\mu\mu}^{\text{match}}| = 10^{-30}$), while LL-RG drift remains bounded: $\langle |\Delta\mu_{\mu\mu}| \rangle \simeq 6.19 \times 10^{-4}$ and $\max |\Delta\mu_{\mu\mu}| \simeq 8.93 \times 10^{-3}$; $f(\chi_{\mu\mu}^2 < 4)$ is unchanged between UV-tree and UV+LL-RG. To make that witness explicit at the operator level, we also export a layer-resolved basis audit with `scan_hll_uv_operator_basis_audit.py`; the corresponding map and summary are stored under `paper/hll_uv_operator_basis_*`. There the UV-tree closure is reconstructed as

$$C_{eH}^{\text{tree}} = \sum_{N=1}^3 \left(\frac{P_N^{\text{kin}}}{M_N^2} \right) \mathcal{B}_N^{(eH)}, \quad \mathcal{B}_N^{(eH)} \equiv \mathbf{g}_N^{(eH)} \mathbf{g}_N^{(eH)T}, \quad (59)$$

where $\mathbf{g}_N^{(eH)}$ is the flavor-layer coupling vector entering Eq. (52); this notation avoids collision with the micro-degeneracy scalar g_N and the visibility factor B_N used elsewhere. With explicit blockwise witnesses for ΔC^{match} and ΔC^{RGE} , the tree/match/IR reconstruction residuals are numerically zero at map level, while the reconstructed $\mu_{\mu\mu}$ agrees with the native UV+LL-RG observable to $\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 7.31 \times 10^{-8}$. This does not yet constitute a full EYMH→SMEFT derivation, but it upgrades the current closure from a black-box Wilson-map audit to an explicit basis-level EFT witness. As a next structured step beyond fixed $(\kappa_{\text{diag}}, \kappa_{\text{offdiag}})$, we test a finite-match comparator whose effective shifts are built from local UV-basis invariants (shell spread, coefficient variation of P_N^{kin}/M_N^2 , and a normalized off-diagonal mixing norm). We treat this as a structured comparator, not as a replacement for the current constant finite-match baseline. On the canonical D21×E21 grid, the basis witness remains exact ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.55 \times 10^{-8}$, reconstruction residuals zero), and the refreshed comparator stays in the small-deformation regime relative to the constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 5.06 \times 10^{-4}$, $p95 = 1.03 \times 10^{-3}$, $\max |\Delta\mu_{\mu\mu}| = 7.87 \times 10^{-3}$) with zero acceptance mismatch. The off-diagonal component is numerically inactive in the current UV basis, while the diagonal-threshold scale simply controls the strength of a small diagonal threshold deformation. To quantify this interpretation, we run a Phase-2 diagonal-threshold window audit. On the refreshed D21×E21 scan, both gates require zero acceptance mismatch and select

$$\begin{aligned} \kappa_{\text{diag}}^{(\text{th})} &\in [0.25, 1.0] \quad (\text{conservative comparator window}), \\ \kappa_{\text{diag}}^{(\text{th})} &\in [0.25, 1.5] \quad (\text{extended stable window}). \end{aligned}$$

The canonical witness choice $\kappa_{\text{diag}}^{(\text{th})} = 1.0$ therefore sits at the upper edge of the conservative window: it keeps the diagonal threshold feed-through explicit while remaining in the small-deformation regime. This sharpens the current interpretation of $\kappa_{\text{diag}}^{(\text{th})}$ as a diagonal threshold susceptibility. The remaining missing ingredient is not off-diagonal closure, which is numerically absent in the present basis, but a parent-action normalization for this susceptibility. We now push this normalization one step closer to the parent action by introducing an action-normalized

finite-match comparator. The basis witness remains exact ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.54 \times 10^{-8}$, zero reconstruction residuals), and the comparator keeps zero acceptance mismatch with only small deformation relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 1.00 \times 10^{-3}$, $p95 = 1.86 \times 10^{-3}$, $\max |\Delta\mu_{\mu\mu}| = 1.59 \times 10^{-2}$). We now probe the remaining gap directly with an absolute parent-action comparator. In this mode the external diagonal-scale choice is removed and replaced by an absolute witness built from parent-action-side invariants and coefficient alignment:

$$\kappa_{\text{abs}}^{(d)} \equiv S_{\text{kin}}^{(P)} A_{\text{coeff}}, \quad A_{\text{coeff}} \equiv \frac{\|c_N\|_2}{\|c_N\|_1},$$

so that the effective diagonal threshold is generated locally rather than chosen externally. On the canonical D21×E21 grid, this absolute parent-action comparator remains exact at the basis level ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.55 \times 10^{-8}$, zero reconstruction residuals), and the map again keeps zero acceptance mismatch relative to the refreshed constant baseline with only small deformation ($\langle |\Delta\mu_{\mu\mu}| \rangle = 7.17 \times 10^{-4}$, $p95 = 1.17 \times 10^{-3}$, $\max |\Delta\mu_{\mu\mu}| = 1.14 \times 10^{-2}$). To test whether this remaining prefactor can be pushed closer to a genuine curved-background loop origin, we next introduced a heat-kernel comparator family. The direct lesson from the raw and flat-space-subtracted local witnesses is negative: a purely local a_2 -type diagonal density remains dominated by the well region and does not align with the absolute parent-action normalization. The useful signal appears only after promoting the loop witness to a *contrast-based* non-local quantity that compares well and barrier curvature structure. On the canonical D21×E21 grid this contrast-based loop comparator remains exact at the basis level, keeps zero acceptance mismatch, and has only small deformation relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 4.88 \times 10^{-4}$, $p95 = 8.76 \times 10^{-4}$, $\max |\Delta\mu_{\mu\mu}| = 7.67 \times 10^{-3}$). We then add one more loop-improved comparator layer by modulating the parent-action absolute witness with a contrast-based loop prefactor,

$$\kappa_{\text{diag}}^{\text{eff}} = \kappa_{\text{abs}}^{(d)} \Pi_{\text{hk}}^{(d)} s_{\text{sh}} (1 + c_{\text{cv}}) a_{\text{norm}}^{(d)}, \quad (60)$$

$$\kappa_{\text{offdiag}}^{\text{eff}} = \kappa_{\text{abs}}^{(o)} \Pi_{\text{hk}}^{(o)} s_{\text{sh}} m_{\text{off}} a_{\text{norm}}^{(o)}. \quad (61)$$

Here $(\kappa_{\text{abs}}^{(d/o)}, \Pi_{\text{hk}}^{(d/o)}, s_{\text{sh}}, c_{\text{cv}}, m_{\text{off}}, a_{\text{norm}}^{(d/o)})$ denote the absolute parent-action witness, heat-kernel loop prefactor, shell-spread witness, coefficient variation, off-diagonal mixing norm, and normalized action weight. This loop-improved absolute comparator remains exact at the basis level ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.56 \times 10^{-8}$, zero reconstruction residuals) and keeps zero acceptance mismatch, with only small deformation relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 3.37 \times 10^{-4}$, $p95 = 6.92 \times 10^{-4}$, $\max |\Delta\mu_{\mu\mu}| = 5.33 \times 10^{-3}$). We now add one further EYMH-side screening layer, which keeps the same contrast-based loop family but modulates it by local

mass-access and curvature-screen factors together with shell-dispersion screening,

$$A_X = \sqrt{\frac{X_{\text{well}}}{1 + X_{\text{well}}}}, \quad A_R = (1 + R_{\text{bar}})^{-1/2}, \quad (62)$$

$$\Pi_{\text{EYMH}} = \Pi_{\text{hk}} A_X A_R \sqrt{\frac{s_{\text{sh}}}{1 + s_{\text{sh}}}} \sqrt{c_{\text{align}}} (1 + g_{\text{cv}} + c_{\text{tree}})^{-1/2}. \quad (63)$$

Here $(X_{\text{well}}, R_{\text{bar}}, s_{\text{sh}}, c_{\text{align}}, g_{\text{cv}}, c_{\text{tree}})$ denote the well-mass, barrier-curvature, shell-spread, coefficient-alignment, gap-variation, and tree-diagonal variation witnesses. This EYMH-screened loop comparator remains exact at the basis level ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.56 \times 10^{-8}$, zero reconstruction residuals) and keeps zero acceptance mismatch, with only small deformation relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 4.17 \times 10^{-5}$, $p95 = 7.63 \times 10^{-5}$, $\max |\Delta\mu_{\mu\mu}| = 6.46 \times 10^{-4}$). We therefore read it as the current best *EYMH-side absolute loop-prefactor comparator*. We then promote the two dominant source factors already isolated by the EYMH-side audits into a refreshed source-informed EYMH comparator. This mode again remains exact at the basis level, keeps zero acceptance mismatch, and stays in the same small-deformation class relative to the refreshed constant baseline ($\langle |\Delta\mu_{\mu\mu}| \rangle = 5.2 \times 10^{-5}$, $p95 = 1.09 \times 10^{-4}$, $\max |\Delta\mu_{\mu\mu}| = 7.81 \times 10^{-4}$). The promoted source-informed prefactor remains nonzero across the whole D21×E21 map, so this is the current best *source-informed EYMH comparator*.

We now rewrite that same source-informed block in explicit parent-action participation/compressibility language. This parented rewrite remains exact at the basis level ($\max |\Delta\mu_{\mu\mu}^{\text{recon}}| = 5.56 \times 10^{-8}$, zero reconstruction residuals), keeps zero acceptance mismatch, and on the canonical D21×E21 fix grid is map-identical to the source-informed comparator. Relative to the refreshed constant baseline it therefore inherits the same small-deformation scale ($\langle |\Delta\mu_{\mu\mu}| \rangle = 5.2 \times 10^{-5}$, $p95 = 1.09 \times 10^{-4}$, $\max |\Delta\mu_{\mu\mu}| = 7.81 \times 10^{-4}$). The numerical role of this mode is not to introduce a new baseline, but to show that the same comparator can be written directly in parent-action participation/coherence and shell-background compressibility language without changing the canonical map. Figure 7 summarizes this normalization-side endpoint visually: the parented rewrite is map-identical to the source-informed comparator while the remaining deformation relative to the constant finite-match baseline is small and structured.

Proposition U2 (Factorized EYMH prefactor). The canonical EYMH-side normalization block admits an exact factorization into local loop, shell-access, participation, and dispersion-screen sectors. The exported audit under `paper/h11_uv_eymh_prefactor_decomposition*_D21E21.*` verifies at numerical precision that

$$\Pi_{\text{EYMH}} = \Pi_{\text{loc}} A_{\text{sh}} A_{\text{part}} S_{\text{disp}}, \quad (64)$$




FIG. 7. EYMH-side parented finite-match comparator on the canonical D21x21 fix grid. Top row: map-level deformation against the refreshed constant baseline and the source-informed comparator; the source-informed and parented maps coincide to numerical precision. Bottom row: the corresponding EYMH source prefactor, parented prefactor, and tree-diagonal compressibility witness.

with $\max |\Delta_{\text{recon}}| = 9.71 \times 10^{-17}$ and $\max |\Delta_{\log_{10}}| = 1.48 \times 10^{-15}$. Here $(\Pi_{\text{loc}}, A_{\text{sh}}, A_{\text{part}}, S_{\text{disp}})$ denote the exported local loop, shell-access, participation/alignment, and dispersion-screen factors. The decomposition shows that the strongest positive correlation is carried by the participation block (0.9655), while the leading suppression comes from the dispersion screen (-0.8774); the local loop factor remains secondary (0.4675). The remaining UV-to-EFT gap is therefore no longer the existence or stability of an EYMH-side absolute prefactor, but the parent-action origin of the participation and dispersion-screen factors that dominate it.

To sharpen that statement further, we resolved the two dominant factors one level deeper in a source audit on the same canonical map. The alignment block is not an opaque extra screen: it is exactly the coefficient-participation access

$$A_{\text{part}} = N_{\text{eff}}^{-1/4}, \quad N_{\text{eff}} = \left(\frac{c_{l1}}{c_{l2}} \right)^2,$$

with source-level reconstruction again closed to machine precision, $\max |\Delta_{\text{source}}| = 9.71 \times 10^{-17}$. Likewise, the diagonal screening factor splits as

$$S_{\text{disp}} = S_{\text{gap}} S_{\text{tree}},$$

where the shell-gap contribution is numerically weak (corr = 0.0907 with the EYMH prefactor), while the tree-diagonal screen carries the dominant suppression (corr = -0.8226). Thus the

remaining parent-action gap is now more specific: why coefficient participation controls the alignment block, and why tree-diagonal dispersion dominates the screening of the canonical prefactor. We can sharpen the second statement one step further by rewriting

$$S_{\text{tree}} = (1 + \chi_{\text{tree}})^{-1/2}, \quad \chi_{\text{tree}} \equiv \frac{c_{\text{tree}}^{\text{cv}}}{1 + g_{\text{cv}}},$$

so that the canonical screening factor is exactly the compressibility witness of a shell-background-normalized diagonal susceptibility. The corresponding audit closes identically ($\max |\Delta_{\text{tree}}| = 0$) and shows that the susceptibility itself is positively correlated with the EYMH prefactor ($\text{corr} = 0.7409$), while the compressibility carries the leading suppression ($\text{corr} = -0.8226$). A complementary tree-diagonal pressure fraction also correlates strongly with the prefactor ($\text{corr} = 0.8449$). Hence the unresolved parent-action question is now even narrower: one must explain not merely why tree-diagonal dispersion enters, but why a shell-normalized diagonal susceptibility/compressibility controls the dominant screening of the canonical EYMH loop prefactor. We can sharpen the source-side interpretation one step further by rewriting the participation block in the language of a two-mode loop trace,

$$N_{\text{trace}} = \frac{1}{p_1^2 + p_2^2}, \quad S_{\text{trace}}^{\text{norm}} = \frac{-p_1 \ln p_1 - p_2 \ln p_2}{\ln 2}, \quad p_{1,2} = \frac{c_{l1,l2}}{c_{l1} + c_{l2}}.$$

The dedicated parent-source audit (summary under `paper/hll_uv_*parent_source_model*`) closes both rewrites to machine precision, with $\max |\Delta_{\text{part}}| = 2.22 \times 10^{-16}$ and $\max |\Delta_{\text{tree}}| = 1.11 \times 10^{-16}$. Numerically, the participation witness is strongly controlled by the effective loop-trace participation number and its entropy, $\text{corr}(N_{\text{trace}}, A_{\text{part}}) = 0.9608$ and $\text{corr}(S_{\text{trace}}^{\text{norm}}, A_{\text{part}}) = -0.9560$, while the shell-background-normalized pressure/compressibility block remains the leading source-side screen, $\text{corr}(P_{\text{tree}}, \Pi_{\text{src}}) = -0.7213$ and $\text{corr}(S_{\text{tree}}, \Pi_{\text{src}}) = 0.7066$. This narrows the outstanding EYMH normalization problem again: the unresolved parent-action physics is no longer a generic alignment/dispersion question, but specifically why the loop trace concentrates into the observed two-mode participation structure and why a shell-background-normalized tree-diagonal pressure response fixes the surviving source-informed prefactor. We can now close the participation side one step further with the exact audit stored under `paper/hll_uv_*participation_exact_audit*_D21E21_fix.*`. Writing the projected two-mode participation number as

$$N_{\text{eff}} = \frac{1}{p_1^2 + p_2^2}, \quad d = \sqrt{\frac{2}{N_{\text{eff}}}} - 1, \quad A_{\text{part}}^{\text{exact}} = \sqrt{\frac{1-d}{1+d}},$$

we reconstruct the parent participation factor to machine precision, with $\max |\Delta_{\text{part}}^{\text{exact}}| = 9.99 \times 10^{-16}$. The participation block is therefore no longer just strongly correlated with a projected response; it is exactly fixed by the two-mode loop-trace imbalance implied by the

parented map. At this stage the remaining EYMH normalization gap has narrowed again: what remains to be explained from the parent action is why the fluctuation operator dynamically selects this exact two-mode participation closure together with the shell-background-normalized tree-diagonal pressure/compressibility response. We can sharpen this parent-action reading once more by eliminating the intermediate two-mode probabilities altogether and working directly with the projected coefficient vector $c_N = P_N^{\text{kin}}/M_N^2$. The coefficient-norm audit stored under `paper/hll_uv_*participation_norm_audit*_D21E21_fix.*` shows that

$$Q_2 = \frac{\|c\|_2^2}{\|c\|_1^2}, \quad N_{\text{eff}}^{\text{norm}} = \frac{\|c\|_1^2}{\|c\|_2^2}, \quad A_{\text{part}}^{\text{norm}} = Q_2^{1/4} = \sqrt{\frac{\|c\|_2}{\|c\|_1}}$$

reconstructs the parent participation factor to machine precision, with $\max |\Delta_{\text{part}}^{\text{norm}}| = 2.22 \times 10^{-16}$. The participation block can therefore be read directly as an exact norm-ratio coherence of the projected parent-action coefficient block. Numerically, this norm coherence remains correlated with the canonical parented prefactor, $\text{corr}(A_{\text{part}}^{\text{norm}}, \Pi_{\text{parent}}) = 0.7129$, while staying essentially orthogonal to the residual map deformation relative to the refreshed constant baseline. The remaining normalization gap is now narrower once again: explain why the parent EYMH fluctuation operator fixes this exact projected coefficient-norm participation closure together with the shell-background-normalized tree-diagonal pressure/compressibility response. The new coefficient-norm tilt audit stored under `paper/hll_uv_*participation_tilt_audit*_D21E21_fix.*` removes even the explicit norm-ratio intermediate and rewrites the same participation block as an exact projected free-energy tilt,

$$\Delta F_{\text{norm}} = \log \frac{\|c\|_1}{\|c\|_2}, \quad A_{\text{part}}^{\text{tilt}} = e^{-\Delta F_{\text{norm}}/2} = \sqrt{\frac{\|c\|_2}{\|c\|_1}}.$$

On the canonical D21×E21 fix grid this tilt coherence reconstructs the parent participation factor with $\max |\Delta_{\text{part}}^{\text{tilt}}| = 2.22 \times 10^{-16}$ and, together with the shell access and tree-diagonal compressibility blocks, closes the full parented prefactor exactly as

$$\Pi_{\text{parent}} = \Pi_{\text{hk,local}} A_{\text{shell}} A_{\text{part}}^{\text{tilt}} S_{\text{tree}},$$

with $\max |\Delta \Pi_{\text{parent}}| = 9.02 \times 10^{-17}$. In this language the surviving EYMH normalization gap is narrower again: what remains is to explain why the fluctuation operator dynamically fixes this exact projected coefficient-norm tilt coherence together with the shell-background-normalized tree-diagonal pressure/compressibility response. **Proposition U3 (Projected response action).** The participation tilt and the tree-diagonal compressibility combine additively into a single projected response action. The response-action audit stored under `paper/hll_uv_*response_action_audit*_D21E21_fix.*` defines

$$S_{\text{resp}} = \Delta F_{\text{norm}} + \log(1 + \chi_{\text{tree}}), \quad A_{\text{resp}} = e^{-S_{\text{resp}}/2},$$

so that the full canonical parented prefactor closes exactly as

$$\Pi_{\text{parent}} = \Pi_{\text{hk,local}} A_{\text{shell}} A_{\text{resp}},$$

with $\max |\Delta A_{\text{resp}}| = 2.22 \times 10^{-16}$ and $\max |\Delta \Pi_{\text{parent}}^{\text{resp}}| = 9.71 \times 10^{-17}$. In other words, the coefficient-norm tilt coherence and the shell-background-normalized tree-diagonal compressibility are not merely parallel witnesses; they enter additively in the same projected response action whose exponential fixes the canonical parented prefactor. The remaining EYMH normalization gap is therefore narrower again: explain why the fluctuation operator dynamically selects this exact projected response action rather than only its separate tilt and compressibility components. We can sharpen this once more by identifying the two terms in S_{resp} with the most direct projected kernel objects currently available in the audited projection. The log-det / Schur audit stored under `paper/hll_uv_*logdet_schur_audit*_D21E21_fix.*` defines

$$K_{\text{part}} = \frac{\|c\|_1}{\|c\|_2}, \quad G_{\text{Schur}} = \frac{1 + \text{gap}_{\text{cv}} + c_{\text{tree,diag,cv}}}{1 + \text{gap}_{\text{cv}}} = 1 + \chi_{\text{tree}},$$

so that

$$S_{\text{resp}} = \log \det K_{\text{part}} + \log G_{\text{Schur}}.$$

On the canonical D21×E21 fix grid this log-det / Schur rewriting again reconstructs the full response weight and parented prefactor to machine precision, with $\max |\Delta A_{\text{resp}}^{\log \det}| = 2.22 \times 10^{-16}$ and $\max |\Delta \Pi_{\text{parent}}^{\log \det}| = 9.71 \times 10^{-17}$. The surviving EYMH normalization block can therefore be read as an exact projected log-det participation kernel together with a shell-background-normalized Schur response. The remaining gap is now narrower still: explain why the fluctuation operator dynamically selects precisely this projected log-det / Schur structure. **Proposition**

U4 (Kernel-selection principle). The remaining UV-side question can be phrased as a genuine kernel-selection statement rather than as bookkeeping. The audit stored under `paper/hll_uv_*kernel_selection*_D21E21_fix.*` probes the minimal deformed projected family

$$K_{\text{sel}} = \begin{pmatrix} e^{\alpha S_{\text{part}}} & \lambda \sqrt{(e^{\alpha S_{\text{part}}} - 1)(e^{\beta S_{\text{schur}}} - 1)} \\ \lambda \sqrt{(e^{\alpha S_{\text{part}}} - 1)(e^{\beta S_{\text{schur}}} - 1)} & e^{\beta S_{\text{schur}}} \end{pmatrix},$$

whose canonical parented point is $(\alpha, \beta, \lambda) = (1, 1, 0)$. On the D21×E21 fix grid the response-weight and prefactor residuals are uniquely minimized there: the first nontrivial runner-up occurs at $(1, 1, -0.1)$ but already opens an RMSE gap of 2.67×10^{-5} in the parented prefactor. A finite-difference stationarity test at the canonical point yields nearly vanishing gradients, $\partial_{\alpha} J = -3.00 \times 10^{-9}$, $\partial_{\beta} J = -6.34 \times 10^{-9}$, $\partial_{\lambda} J = 0$, together with non-negative Hessian eigenvalues $(2.99 \times 10^{-9}, 1.33 \times 10^{-3}, 7.59 \times 10^{-2})$. Within this minimal projected kernel family,

the fluctuation operator therefore selects unit log-det / Schur weights and suppresses explicit participation–tree cross-coupling. We can sharpen that statement once more by replacing the finite-difference scan with an exact local stationarity principle. The stationarity audit stored under `paper/hll_uv_*stationarity_audit*_D21E21_fix.*` defines the projected mismatch functional

$$J(\alpha, \beta, \lambda) = \left\langle (A(\alpha, \beta, \lambda) - A_{\text{ref}})^2 \right\rangle.$$

At the canonical point $(\alpha, \beta, \lambda) = (1, 1, 0)$ all first-variation moments vanish exactly: $\partial_\alpha J = \partial_\beta J = \partial_\lambda J = 0$. The (α, β) sector is then controlled by an exact quadratic stationarity matrix with positive eigenvalues $(1.33 \times 10^{-3}, 7.59 \times 10^{-2})$, while the cross-coupling direction is even under $\lambda \rightarrow -\lambda$ and therefore begins only at quartic order,

$$J(1, 1, \lambda) = C_4 \lambda^4 + \mathcal{O}(\lambda^6), \quad C_4 = 1.4974 \times 10^{-3}.$$

For the concrete probe $\lambda = 0.1$, this quartic law predicts an RMSE of 3.8696×10^{-4} , matching the directly evaluated 3.8734×10^{-4} . The projected EYMH selection principle can therefore be stated more sharply than before: the canonical log-det / Schur kernel is not only the best nearby fit, it is an exact stationary point with quadratic stability in the log-det / Schur weights and quartic suppression of explicit participation–tree mixing. We can package the same statement into a local projected effective-action gap. The variational-selection audit stored under `paper/hll_uv_*variational_selection*_D21E21_fix.*` defines

$$\Delta\Gamma_{\text{sel}}(\delta\alpha, \delta\beta, \lambda) = \frac{1}{2} \delta\theta^T H \delta\theta + C_4 \lambda^4, \quad \delta\theta = (\delta\alpha, \delta\beta),$$

using the exact quadratic stationarity matrix H from the (α, β) block and the exact quartic coefficient C_4 from the mixing direction. On a local grid around the canonical point, this variational gap and the exact mismatch functional have the same minimum at $(1, 1, 0)$, and the agreement is already very strong: $\text{corr}(J_{\text{exact}}, \Delta\Gamma_{\text{sel}}) = 0.9894$, the pure- λ slice is reproduced to within 1.88×10^{-8} in absolute gap, and even the $\lambda = 0$ (α, β) plane remains controlled at p95 absolute gap 2.97×10^{-5} . This is the cleanest projected effective-action statement so far: near the canonical point, the fluctuation operator selects the log-det / Schur kernel by minimizing a local effective action whose quadratic log-det/Schur sector and quartic mixing sector reproduce the observed selection landscape. We can push the same statement one step closer to a parent-action kernel form. The audit stored under `paper/hll_uv_*parent_kernel_statement*_D21E21_fix.*` rewrites the deformed projected kernel family as

$$K_{11} = e^{\alpha S_{\text{part}}}, \quad K_{22} = e^{\beta S_{\text{schur}}}, \quad K_{12} = \lambda \sqrt{(K_{11} - 1)(K_{22} - 1)},$$

and then expresses the deformed response weight exactly as

$$A(\alpha, \beta, \lambda) = A_{\text{ref}} \exp \left[-\frac{1}{2} \Delta S_{\text{kernel}} \right],$$

with

$$\Delta S_{\text{kernel}} = (\alpha - 1)S_{\text{part}} + (\beta - 1)S_{\text{schur}} + \log(1 - \lambda^2 \xi_{\text{cross}}).$$

The corresponding parent-kernel objective reproduces the direct mismatch functional to machine precision on the full local scan, with $\max |J_{\text{direct}} - J_{\text{parent}}| = 1.30 \times 10^{-18}$ and $\max |A_{\text{direct}} - A_{\text{parent}}| = 3.33 \times 10^{-16}$. The same exact functional is minimized again at the canonical point $(1, 1, 0)$. This is the strongest mother-action statement so far: the canonical log-det / Schur kernel is not only stationary and variationally preferred in a local surrogate sense, but is selected by an exact projected parent-kernel excess functional in which log-det and Schur deformations enter linearly while explicit participation-tree mixing appears only through an even determinant factor. We can now make the block structure behind that statement explicit. The audit stored under `paper/hll_uv_*block_split*_D21E21_fix.*` writes the canonical projected fluctuation kernel in terms of a participation block

$$K_{\text{part}} = \frac{\|c\|_1}{\|c\|_2},$$

a shell-background/tree block

$$G_{\text{Schur}} = \frac{1 + \text{gap}_{\text{cv}} + c_{\text{tree,diag,cv}}}{1 + \text{gap}_{\text{cv}}},$$

and a mixed scale

$$C_{\text{mix}} = \sqrt{(K_{\text{part}} - 1)(G_{\text{Schur}} - 1)}.$$

On the canonical map the response action closes exactly as $S_{\text{part}} + S_{\text{schur}}$, with $\text{corr}(S_{\text{split}}, S_{\text{resp}}) = 1.0$ and $\max |S_{\text{split}} - S_{\text{resp}}| = 7.77 \times 10^{-16}$. For the deformed family

$$K_{\text{sel}} = \begin{pmatrix} K_{\text{part}} & \lambda C_{\text{mix}} \\ \lambda C_{\text{mix}} & G_{\text{Schur}} \end{pmatrix},$$

the determinant factorization

$$\det K_{\text{sel}} = K_{\text{part}} G_{\text{Schur}} (1 - \lambda^2 \xi_{\text{cross}})$$

also closes to machine precision on the full scan, with $\max |\Delta \det| = 4.44 \times 10^{-16}$ and $\max |\Delta S| = 3.33 \times 10^{-16}$. This is the clearest structural statement so far: the canonical projected fluctuation operator is block-diagonal in the participation and shell-background/tree sectors, and explicit

participation–tree mixing survives only as an even determinant-level penalty. We can sharpen the same statement one step further by embedding it into a background-normalized parent block determinant. The audit stored under `paper/hll_uv_*parent_blockdet*_D21E21_fix.*` defines

$$K_{11} = e^{\alpha S_{\text{part}}}, \quad K_{\text{bg}} = 1 + \text{gap}_{\text{cv}}, \quad K_{22} = K_{\text{bg}} e^{\beta S_{\text{Schur}}},$$

and it is convenient to abbreviate the deformed normalized tree/background block as

$$G_\beta := \frac{K_{22}}{K_{\text{bg}}} = e^{\beta S_{\text{Schur}}}, \quad G_{\text{Schur}} = G_\beta|_{\beta=1}.$$

with parent mixing scale

$$C_{\text{parent}} = \sqrt{(K_{11} - 1)(K_{22} - K_{\text{bg}})}, \quad \mathcal{K}_{\text{parent}} = \begin{pmatrix} K_{11} & \lambda C_{\text{parent}} \\ \lambda C_{\text{parent}} & K_{22} \end{pmatrix}.$$

The normalized determinant and Schur factor then satisfy the exact identities

$$\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} = K_{11} G_\beta (1 - \lambda^2 \xi_{\text{cross}}), \quad \widehat{G}_{\text{Schur}} = \frac{K_{22} - K_{12}^2 / K_{11}}{K_{\text{bg}}} = G_\beta (1 - \lambda^2 \xi_{\text{cross}}),$$

so the response action is literally the block-determinant / Schur-complement factorization

$$S_{\text{parent}} = \log K_{11} + \log \widehat{G}_{\text{Schur}}.$$

On the D21×E21 fix grid this closes to machine precision, with $\max |S_{\text{parent}} - S_{\text{resp}}| = 5.55 \times 10^{-16}$, $\max |\Delta(\det / K_{\text{bg}})| = 1.33 \times 10^{-15}$, $\max |\Delta \widehat{G}_{\text{Schur}}| = 6.66 \times 10^{-16}$, and $\text{corr}(J_{\text{direct}}, J_{\text{blockdet}}) = 1.0$. This is the strongest derivation statement so far: the canonical response weight is the inverse square root of a background-normalized projected parent block determinant, rather than merely a convenient response-action rewrite. The only remaining structural ambiguity in that parent block is the mixed entry. To test whether the geometric-mean choice is merely convenient or actually selected, the audit stored under `paper/hll_uv_*parent_mix_geomean*_D21E21_fix.*` scans the smallest symmetric-excess family

$$C_{\text{gen}} = \kappa (K_{11} - 1)^u (K_{22} - K_{\text{bg}})^v.$$

On the local D21×E21 fix grid the unique exact point is

$$(u, v, \kappa) = \left(\frac{1}{2}, \frac{1}{2}, 1\right), \quad C_{\text{parent}} = \sqrt{(K_{11} - 1)(K_{22} - K_{\text{bg}})},$$

for which the determinant, Schur, weight, and normalized cross-ratio residuals all collapse to machine precision:

$$\begin{aligned} \max |\Delta(\det / K_{\text{bg}})| &= 4.44 \times 10^{-16}, & \max |\Delta \widehat{G}_{\text{Schur}}| &= 4.44 \times 10^{-16}, \\ \max |\Delta A| &= 1.11 \times 10^{-16}, & \max |\Delta \xi| &= 1.11 \times 10^{-16}. \end{aligned}$$

The first nontrivial runner-up, $(u, v, \kappa) = (0.625, 0.625, 1.1)$, already opens visible residuals, with $\max |\Delta(\det / K_{\text{bg}})| = 1.32 \times 10^{-3}$ and $\max |\Delta A| = 2.18 \times 10^{-4}$. This is the strongest naturality statement so far: once the projected parent kernel is required to couple the participation and tree/background sectors through a minimal symmetric excess family, the geometric-mean mixed block is uniquely selected. We can push that one step further by testing the nearest non-minimal extension: a ratio-dependent warp of the geometric mean. The audit stored under `paper/hll_uv_*parent_ratio_warp*_D21E21_fix.*` probes

$$C_{\text{warp}} = \kappa C_{\text{parent}} \exp[\delta L + \nu L^2], \quad L = \frac{1}{2} \log(E_{\text{part}}/E_{\text{tree}}),$$

with $E_{\text{part}} = K_{11} - 1$ and $E_{\text{tree}} = K_{22} - K_{\text{bg}}$. On the same D21×E21 fix grid the unique exact point is again the unwarped kernel, $(\kappa, \delta, \nu) = (1, 0, 0)$, for which the determinant, Schur, weight, and normalized cross-ratio residuals all vanish exactly. The first nontrivial runner-up, $(1, 0, -0.05)$, already opens visible residuals, with $\max |\Delta(\det / K_{\text{bg}})| = 7.18 \times 10^{-4}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 5.39 \times 10^{-4}$, $\max |\Delta A| = 7.14 \times 10^{-5}$, and $\max |\Delta \xi| = 6.08 \times 10^{-3}$. This is the strongest minimality statement so far: the low-mode parent block not only selects the geometric-mean mixed sector inside the minimal symmetric-excess family, it also rejects the first ratio-warped extension of that family. We can then reparameterize the same local family in the coordinates that most directly match the remaining proof obligation, namely an overall normalization shift m , a symmetric homogeneity degree s , and an antisymmetric participation/tree tilt a :

$$C_{\text{gen}} = \exp(m) \exp\left[\frac{s}{2} \log(E_{\text{part}} E_{\text{tree}}) + \frac{a}{2} \log(E_{\text{part}}/E_{\text{tree}})\right].$$

The audit stored under `paper/hll_uv\_parent\_symnorm*_D21E21\_fix.*` shows that on the same D21×E21 fix grid the unique exact point is

$$(m, s, a) = (0, 1, 0),$$

for which the determinant, Schur, weight, and normalized cross-ratio residuals again collapse to machine precision. The first nontrivial runner-up, $(0.05, 1.125, 0)$, already opens visible residuals, with $\max |\Delta(\det / K_{\text{bg}})| = 6.61 \times 10^{-4}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 4.72 \times 10^{-4}$, $\max |\Delta A| = 1.09 \times 10^{-4}$, and $\max |\Delta \xi| = 7.73 \times 10^{-3}$. This is the cleanest structural statement so far: once the parent block is written in symmetry/normalization coordinates, the low-mode projected parent action selects the canonical family by enforcing zero normalization shift, unit symmetric degree, and zero antisymmetric tilt. The last nearby structural ambiguity is whether the parent action truly selects a locally affine generator in these excess coordinates, or whether small log-curvature

terms survive. The audit stored under `paper/hll\_uv\_parent\_generator\_affinity\_D21E21\_fix.*` probes the first non-affine local extension,

$$\log C_{\text{gen}} = \frac{1}{2}L_{\text{sum}} + q_{ss}L_{\text{sum}}^2 + q_{dd}L_{\text{diff}}^2 + q_{sd}L_{\text{sum}}L_{\text{diff}},$$

with $L_{\text{sum}} = \log(E_{\text{part}}E_{\text{tree}})$ and $L_{\text{diff}} = \log(E_{\text{part}}/E_{\text{tree}})$. On the same D21×E21 fix grid the unique exact point is

$$(q_{ss}, q_{dd}, q_{sd}) = (0, 0, 0),$$

for which the determinant, Schur, weight, and normalized cross-ratio residuals again collapse to machine precision. The first nontrivial runner-up, $(0, -0.0125, 0.0125)$, already opens visible residuals, with $\max|\Delta(\det/K_{\text{bg}})| = 2.76 \times 10^{-4}$, $\max|\Delta\hat{G}_{\text{Schur}}| = 1.98 \times 10^{-4}$, $\max|\Delta A| = 2.62 \times 10^{-5}$, and $\max|\Delta\xi| = 2.23 \times 10^{-3}$. This is the narrowest structural statement so far: once symmetry and normalization have been fixed, the low-mode projected parent action also rejects the first local log-curvature corrections, so the surviving canonical parent block lies in the locally affine multiplicative excess class itself. The remaining coordinate-level ambiguity is whether the projected parent action genuinely uses the canonical excess variables themselves, or whether another local subtraction point could play the same role. The audit stored under `paper/hll\_uv\_excess\_coordinate\_D21E21\_fix.*` probes the minimal reference-offset family

$$E_{\text{part}}^{(r)} = K_{11} - r_{\text{part}}, \quad E_{\text{tree}}^{(r)} = K_{22} - r_{\text{tree}}K_{\text{bg}}.$$

On the same D21×E21 fix grid the unique exact point is

$$(r_{\text{part}}, r_{\text{tree}}) = (1, 1),$$

for which the determinant, Schur, weight, normalized cross-ratio, and strict anchor-leakage residuals all vanish identically. The first nontrivial runner-up, $(1, 1.05)$, already opens visible errors, with $\max|\Delta(\det/K_{\text{bg}})| = 1.08 \times 10^{-3}$, $\max|\Delta\hat{G}_{\text{Schur}}| = 7.02 \times 10^{-4}$, $\max|\Delta A| = 2.48 \times 10^{-4}$, and $\max|\Delta\xi| = 1.61 \times 10^{-2}$. More importantly, it distorts the first nonzero-response slices even though the strict anchor leakage still vanishes, with onset mismatches 4.53×10^{-2} in the participation direction and 9.78×10^{-2} in the tree/background direction. This is the sharpest fixed-point statement so far: the projected parent action naturally organizes the mixed block as a deviation from the identity participation block and the shell/background tree block themselves, rather than from an arbitrary nearby subtraction point. We then tightened the same statement by allowing the excess coordinates to vary inside the smallest smooth family that preserves both the fixed points and the tangent normalization there,

$$E_{\text{part}}^{(p)} = \text{BC}_p(K_{11}), \quad E_{\text{tree}}^{(q)} = K_{\text{bg}} \text{BC}_q(K_{22}/K_{\text{bg}}),$$

where $\text{BC}_p(x) = (x^p - 1)/p$ for $p \neq 0$ and $\text{BC}_0(x) = \log x$, so every member obeys $\text{BC}_p(1) = 0$ and $\text{BC}'_p(1) = 1$. On the same D21×E21 fix grid the unique exact point is again the linear additive excess choice

$$(p_{\text{part}}, p_{\text{tree}}) = (1, 1),$$

for which the determinant, Schur, weight, cross-ratio, anchor, and onset residuals all vanish to machine precision. The first nontrivial runner-up, $(0.75, 1.0)$, already opens visible first-slice distortions, with onset mismatches 2.02×10^{-3} and 8.23×10^{-3} together with $\max |\Delta(\det / K_{\text{bg}})| = 8.25 \times 10^{-4}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 5.51 \times 10^{-4}$, $\max |\Delta A| = 9.80 \times 10^{-5}$, and $\max |\Delta \xi| = 7.85 \times 10^{-3}$. This is the narrowest coordinate statement so far: even after fixing the same low-mode fixed points and the same tangent normalization, the projected parent action still uniquely selects the linear excess variables themselves. We can sharpen that result into positive normal-coordinate language by probing the first nonlinear local jet family that preserves the same fixed points and the same unit tangent normalization,

$$E_{\text{part}}^{(\zeta_p)} = E_{\text{part}} + \zeta_p E_{\text{part}}^2, \quad E_{\text{tree}}^{(\zeta_t)} = E_{\text{tree}} + \zeta_t E_{\text{tree}}^2 / K_{\text{bg}}.$$

On the same D21×E21 fix grid the unique exact point is

$$(\zeta_p, \zeta_t) = (0, 0),$$

for which the determinant, Schur, weight, cross-ratio, anchor, and onset residuals again vanish identically. The first nontrivial runner-up, $(0.125, 0)$, already opens visible distortions, with onset mismatches 2.10×10^{-3} and 9.63×10^{-3} together with $\max |\Delta(\det / K_{\text{bg}})| = 9.90 \times 10^{-4}$, $\max |\Delta \hat{G}_{\text{Schur}}| = 6.59 \times 10^{-4}$, $\max |\Delta A| = 1.18 \times 10^{-4}$, and $\max |\Delta \xi| = 9.47 \times 10^{-3}$. This is the strongest positive coordinate statement so far: the projected parent action naturally uses the zero-second-jet normal coordinates around the identity/background fixed points, and those normal coordinates are exactly the linear excess variables.

D. Projected Parent-Kernel Derivation

The preceding exact audits leave only one conceptual task: explain, directly from the projected EYMH parent action, why the localized low-mode response closes into the particular background-normalized 2×2 kernel isolated numerically above. In other words, the remaining gap is no longer a family-search problem but a derivation problem. The three lemmas below separate that derivation into its natural pieces: reduction of the parent Hessian to a two-mode block, uniqueness of the mixed entry inside the anchored positive family, and uniqueness of the linear excess variables as normal coordinates near the identity/background fixed point.

To keep the notation tight, we reserve K_{part} and G_{Schur} for the canonical undeformed split blocks already identified in the earlier exact audits, G_β for the one-parameter deformed normalized tree/background block of the parent-block determinant family, and K_{11} , K_{22} , and $\hat{G}_{\text{Schur}}$ for the entries and Schur complement of the generic projected parent kernel. At the canonical point these objects coincide through

$$K_{11} = K_{\text{part}} = e^{S_{\text{part}}}, \quad K_{22} = K_{\text{bg}} G_{\text{Schur}} = K_{\text{bg}} e^{S_{\text{Schur}}}, \quad \hat{G}_{\text{Schur}} = G_{\text{Schur}}.$$

a. Assumptions and scope. The theorem chain below is formulated on a localized stationary background Ψ_0 of the projected EYMH parent action, under four explicit assumptions: first, the low-mode response admits a dominant two-dimensional projected truncation; second, the selected projected sector is the one whose diagonal responses reproduce the canonical participation block and the shell/tree block already isolated by the earlier block-split closure; third, that projected sector is locally stable so that the induced quadratic form is positive there; and fourth, local coordinate statements are restricted to the audited coordinate families in a neighborhood of the identity/background anchor. The exact audits are therefore used as supporting evidence for the canonical identification of the projected sectors and for the closure of the resulting formulas, but not as hidden additional premises inside the theorem statements themselves.

In this subsection, “dominant two-dimensional projected truncation” is meant in that same localized and audited sense: on the parent map, the determinant / Schur / weight closures already isolate two projected sectors whose diagonal responses reproduce the participation block and the shell/tree block to numerical exactness, and the theorem chain is only claiming that the parent action closes on that same low-mode sector. We are therefore not asserting a global completeness theorem for the full fluctuation spectrum; the statement is local to the same low-mode neighborhood on which the exact bridge audits close.

Lemma 1 (Projected 2×2 Hessian reduction). Let Ψ_0 be a localized stationary background of the projected EYMH parent action, and let $\mathcal{H}_{\text{EYMH}}$ denote its quadratic fluctuation Hessian. Assume that the localized low-mode response admits a dominant two-mode truncation by Gram-normalizable directions u_{part} and u_{tree} , selected so that their diagonal projected responses reconstruct the participation block and the shell/tree block isolated earlier, and assume that the quadratic form induced by $\mathcal{H}_{\text{EYMH}}$ is locally positive on that projected sector. After restriction to that sector, and after Gram normalization of the projected subspace, the quadratic fluctuation form can be written as

$$\delta^2 S_{\text{EYMH}}|_{\text{proj}} = \frac{1}{2} \begin{pmatrix} a_{\text{part}} & a_{\text{tree}} \end{pmatrix} \mathcal{K}_{\text{parent}} \begin{pmatrix} a_{\text{part}} \\ a_{\text{tree}} \end{pmatrix},$$

with positive background-normalized kernel

$$\mathcal{K}_{\text{parent}} = \begin{pmatrix} K_{11} & K_{12} \\ K_{12} & K_{22} \end{pmatrix}, \quad K_{\text{bg}} = 1 + \text{gap}_{\text{cv}},$$

and with response weight

$$A_{\text{resp}} \propto \left(\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} \right)^{-1/2} = \left(K_{11} \hat{G}_{\text{Schur}} \right)^{-1/2}, \quad \hat{G}_{\text{Schur}} = \frac{K_{22} - K_{12}^2/K_{11}}{K_{\text{bg}}}.$$

Proof sketch. Writing the collective fields as $\Psi = \Psi_0 + \delta\Psi$, the parent action expands as

$$S[\Psi_0 + \delta\Psi] = S[\Psi_0] + \delta S[\Psi_0; \delta\Psi] + \frac{1}{2} \langle \delta\Psi, \mathcal{H}_{\text{EYMH}} \delta\Psi \rangle + \mathcal{O}(\delta\Psi^3).$$

Stationarity of the localized background removes the linear term. Restricting to

$$\delta\Psi = a_{\text{part}} u_{\text{part}} + a_{\text{tree}} u_{\text{tree}}$$

and introducing the column matrix $U = (u_{\text{part}}, u_{\text{tree}})$ together with the amplitude vector $a = (a_{\text{part}}, a_{\text{tree}})^T$, define the orthogonal projector onto the localized two-mode sector by

$$P := UU^\dagger, \quad V_2 := \text{im } P,$$

so that after Gram normalization of the projected basis one has

$$U^\dagger U = \mathbf{1}_2, \quad P^\dagger = P, \quad P^2 = P.$$

On that projected sector, $\delta\Psi = P\delta\Psi$, and the projected Hessian operator is

$$H_P := P\mathcal{H}_{\text{EYMH}}P|_{V_2}.$$

Its matrix representation in the basis $(u_{\text{part}}, u_{\text{tree}})$ is

$$\mathcal{K}_{\text{parent}} := U^\dagger \mathcal{H}_{\text{EYMH}} U.$$

Thus $\mathcal{K}_{\text{parent}}$ is the projected Hessian block itself, not an inverse propagator inserted by hand. After Gram normalization this reduces to the matrix of projected Hessian entries

$$H_{ij} = \langle u_i, \mathcal{H}_{\text{EYMH}} u_j \rangle, \quad i, j \in \{\text{part}, \text{tree}\},$$

and therefore to a positive 2×2 projected kernel on the localized low-mode sector. Here positivity comes from the projected local-stability assumption, not from stationarity alone. At the canonical point the diagonal entries reduce to the canonical split blocks,

$$K_{11} = K_{\text{part}} = e^{S_{\text{part}}}, \quad \frac{K_{22}}{K_{\text{bg}}} = G_{\text{Schur}} = e^{S_{\text{schur}}}.$$

The content of this identification is that the two-mode truncation is not an arbitrary basis choice but the proposition-level sector-selection hypothesis of the EYMH bridge: the earlier block-split and parent-kernel closures already separate one projected direction whose diagonal response reconstructs the participation block and a second projected direction whose background-normalized diagonal response reconstructs the shell/tree block. In that sense u_{part} and u_{tree} are selected by the previously closed projected sectors rather than inserted ad hoc, and Lemma 1 promotes that exact numerical split into a projected-Hessian statement about the parent action itself. The projected Gaussian fluctuation integral is therefore

$$Z_{\text{proj}} \propto \int_{\mathbb{R}^2} d^2 a \exp \left[-\frac{1}{2} a^T \mathcal{K}_{\text{parent}} a \right] \propto (\det \mathcal{K}_{\text{parent}})^{-1/2}.$$

Dividing out the shell-background reference determinant $Z_{\text{bg}} \propto K_{\text{bg}}^{-1/2}$ produces the normalized response factor

$$A_{\text{resp}} \propto \frac{Z_{\text{proj}}}{Z_{\text{bg}}} \propto \left(\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} \right)^{-1/2},$$

and the Schur identity

$$\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} = K_{11} \frac{K_{22} - K_{12}^2 / K_{11}}{K_{\text{bg}}} = K_{11} \hat{G}_{\text{Schur}}$$

immediately yields

$$S_{\text{parent}} = \log K_{11} + \log \hat{G}_{\text{Schur}}.$$

Audit closure. The exact audits verify each algebraic step appearing in the theorem statement and support the canonical sector assignment used above: the earlier block-split closure identifies the participation and shell/tree diagonals themselves as the canonical projected sectors; the ‘logdet + Schur’ witness reconstructs the response weight and parented prefactor to machine precision, with $\max |\Delta A_{\text{resp}}^{\text{logdet}}| = 2.22 \times 10^{-16}$ and $\max |\Delta \Pi_{\text{parent}}^{\text{logdet}}| = 9.71 \times 10^{-17}$; the background-normalized parent block-determinant / Schur audit gives $\max |S_{\text{parent}} - S_{\text{resp}}| = 5.55 \times 10^{-16}$, $\max |\Delta(\det / K_{\text{bg}})| = 1.33 \times 10^{-15}$, and $\max |\Delta \hat{G}_{\text{Schur}}| = 6.66 \times 10^{-16}$; and the exact parent-kernel excess functional reproduces the direct mismatch functional with $\max |J_{\text{direct}} - J_{\text{parent}}| = 1.30 \times 10^{-18}$. Lemma 1 therefore reduces the remaining work to making the projected-Hessian construction from the parent action explicit enough that this already-exact numerical structure can be read as a genuine Gaussian fluctuation theorem on the localized two-mode sector.

Local truncation remark. What is still *not* claimed here is a global spectral-completeness theorem for the full EYMH fluctuation operator. The current statement is local to the audited low-mode sector and assumes that the complement does not reorganize the projected closure

at leading order. A stronger theorem would require an explicit control of the complement remainder, for example through a Feshbach/Schur estimate on the off-sector couplings H_{PQ}, H_{QP} relative to the complement block H_{QQ} , or an equivalent local spectral-separation bound. At present we use the exact bridge residuals as the empirical evidence that this remainder is negligible on the localized parent map, while leaving the fully quantified truncation bound as part of the remaining proof-level tightening.

Corollary 1 (Local complement remainder form). Let $Q := I - P$, and write the quadratic fluctuation operator in P/Q block form

$$\mathcal{H}_{\text{EYMH}} = \begin{pmatrix} H_{PP} & H_{PQ} \\ H_{QP} & H_{QQ} \end{pmatrix}, \quad H_{PP} = P\mathcal{H}_{\text{EYMH}}P|_{V_2}.$$

Assume in the same localized regime that H_{QQ} is invertible on the complement and that the off-sector coupling is small enough that

$$A := H_{PP}^{-1/2} H_{PQ} H_{QQ}^{-1} H_{QP} H_{PP}^{-1/2}$$

satisfies $\|A\| < 1$. Then integrating out the complement produces the effective low-mode operator

$$H_{\text{eff}} = H_{PP} - H_{PQ} H_{QQ}^{-1} H_{QP},$$

and its determinant differs from the projected determinant by

$$\det H_{\text{eff}} = \det H_{PP} \det(I - A), \quad |\log \det H_{\text{eff}} - \log \det H_{PP}| \leq \frac{2\|A\|}{1 - \|A\|}.$$

Proof sketch. The Schur/Feshbach reduction on the complement gives the exact identity

$$H_{\text{eff}} = H_{PP}^{1/2} (I - A) H_{PP}^{1/2},$$

so

$$\det H_{\text{eff}} = \det H_{PP} \det(I - A).$$

Because H_{PP} is 2×2 , the eigenvalues α_1, α_2 of A obey $|\alpha_i| \leq \|A\| < 1$, hence

$$\log \det(I - A) = \log(1 - \alpha_1) + \log(1 - \alpha_2),$$

and the elementary bound $|\log(1 - x)| \leq |x|/(1 - |x|)$ for $|x| < 1$ yields

$$|\log \det(I - A)| \leq \frac{|\alpha_1|}{1 - |\alpha_1|} + \frac{|\alpha_2|}{1 - |\alpha_2|} \leq \frac{2\|A\|}{1 - \|A\|}.$$

If, in addition, H_{QQ} is positive on the complement, then A is positive semidefinite and $\|A\| < 1$ implies $I - A > 0$; hence H_{eff} is positive on the projected low-mode sector as well. Thus the

remaining truncation error is controlled by the product of the off-sector couplings with the complement resolvent, which is exactly the quantity that the current theorem chain leaves as an audited local remainder rather than a global spectral statement.

Once that projected block structure is fixed, all diagonal ambiguity is gone and only one matrix-level choice survives: the off-diagonal coupling between the participation sector and the shell/background sector. Lemma 2 addresses exactly that remaining freedom. Its role is narrower than Lemma 1, but also sharper: after the block decomposition has been fixed at the parent-action level, the mixed entry must still be chosen, and the claim is that only one positive homogeneous anchored coupling remains compatible with the same local symmetry and normalization structure.

Lemma 2 (Geometric-mean mixed entry). With the diagonal sectors fixed as in Lemma 1, define the anchored excess variables

$$E_{\text{part}} = K_{11} - 1, \quad E_{\text{tree}} = K_{22} - K_{\text{bg}}.$$

On the local positive cone $E_{\text{part}} \geq 0$, $E_{\text{tree}} \geq 0$, with the positive square-root branch fixed by continuity from the canonical point, consider positive mixed couplings that vanish at either fixed point, remain symmetric after shell-background normalization, and are locally affine in the logarithms of the excess variables. Within that class, the projected parent block uniquely selects

$$C_{\text{parent}} = \sqrt{E_{\text{part}} E_{\text{tree}}} = \sqrt{(K_{11} - 1)(K_{22} - K_{\text{bg}})}.$$

Proof sketch. Once the diagonal blocks are fixed, the mixed entry is the only remaining structural freedom. The generator-class hypothesis is not inserted ad hoc at this point: on the interior of the local positive cone, any positive mixed entry may be written locally as $C = \exp g(\log E_{\text{part}}, \log E_{\text{tree}})$, and the resulting expression is extended to the anchored boundary by continuity from the canonical point. Local multiplicative reweightings of the two anchored excess amplitudes,

$$(E_{\text{part}}, E_{\text{tree}}) \mapsto (\rho E_{\text{part}}, \sigma E_{\text{tree}}), \quad (\rho, \sigma) > 0,$$

act by translations in logarithmic variables,

$$(\log E_{\text{part}}, \log E_{\text{tree}}) \mapsto (\log E_{\text{part}} + \log \rho, \log E_{\text{tree}} + \log \sigma),$$

or equivalently

$$(\Sigma, \Delta) \mapsto \left(\Sigma + \frac{1}{2}(\log \rho + \log \sigma), \Delta + \frac{1}{2}(\log \rho - \log \sigma) \right).$$

Writing $z = (\Sigma, \Delta)$ and $\tau = (\tau_\Sigma, \tau_\Delta)$ for a small local translation in these logarithmic coordinates, Taylor's theorem gives

$$g(z + \tau) = g(z) + Dg(z)\tau + \frac{1}{2}\tau^T [D^2g(z + \theta\tau)]\tau, \quad 0 < \theta < 1.$$

Thus the first obstruction to pure 1-jet translation stability is exactly the quadratic jet D^2g . Within the audited local family, those quadratic corrections are parametrized by the generator-curvature coefficients (q_{ss}, q_{dd}, q_{sd}) , so the curvature audit is testing the first non-affine term in the local logarithmic normal form rather than an unrelated ad hoc deformation. Consequently one may write

$$g(\Sigma, \Delta) = m + s\Sigma + a\Delta + R_2(\Sigma, \Delta),$$

with the local remainder bounded by

$$|R_2(\Sigma, \Delta)| \leq \frac{1}{2} \sup_{\zeta \in \mathcal{N}} \|D^2g(\zeta)\| \|\Sigma, \Delta\|^2$$

on a sufficiently small neighborhood $\mathcal{N}$ of the canonical point. The natural residual class is therefore the one whose generator is stable under these local translations to first nontrivial order, namely the locally affine class already isolated by the earlier generator-affinity closure. Within that class, the smallest positive multiplicative excess family compatible with the anchor constraints is

$$C_{\text{gen}} = \kappa E_{\text{part}}^u E_{\text{tree}}^v.$$

Equivalently, in logarithmic excess variables

$$\Sigma = \frac{1}{2} \log(E_{\text{part}} E_{\text{tree}}), \quad \Delta = \frac{1}{2} \log(E_{\text{part}}/E_{\text{tree}}),$$

one may write the most general nearby positive homogeneous coupling as

$$\log C_{\text{gen}} = m + s\Sigma + a\Delta + \mathcal{O}(\Sigma^2, \Sigma\Delta, \Delta^2),$$

where m is an overall normalization shift, s is the symmetric homogeneity degree, and a is the antisymmetric tilt. Homogeneity of degree one in the excess amplitudes forces $u + v = 1$, equivalently $s = 1$ in the logarithmic generator, while symmetry under interchange of the two anchored excess sectors forces $u = v$, equivalently $a = 0$. The remaining freedom is the overall normalization $\kappa = e^m$, which is fixed by the canonical anchor normalization. The exact audit under `paper/hll_uv_*parent_mix_geomean*_D21E21_fix.*` closes uniquely at

$$(u, v, \kappa) = \left(\frac{1}{2}, \frac{1}{2}, 1\right),$$

which is precisely the geometric mean. In particular, the closure forces the symmetry condition $u = v$ rather than assuming it a priori. The nearby deformations then show that this is structural, not accidental: the first ratio-warped extension is rejected at $(\kappa, \delta, \nu) = (1, 0, 0)$, the symmetry/normalization audit closes only at $(m, s, a) = (0, 1, 0)$, and the first local generator-curvature corrections are rejected at $(q_{ss}, q_{dd}, q_{sd}) = (0, 0, 0)$. The combined effect is to set the logarithmic generator exactly to

$$\log C_{\text{parent}} = \Sigma = \frac{1}{2} \log(E_{\text{part}} E_{\text{tree}}),$$

and hence

$$C_{\text{parent}} = \sqrt{E_{\text{part}} E_{\text{tree}}}.$$

Thus, once the diagonal sectors, anchors, and local multiplicative structure are fixed, only one positive homogeneous mixed coupling survives. At this stage we deliberately do *not* use the explicit determinant law of the corollary as an external input, because in the current audit chain that determinant identity is only established after the geometric-mean mixed entry has already been inserted into the parent block. The theorem therefore remains non-circular by taking the generator-class closure as the primary uniqueness statement and postponing the explicit determinant formula to the corollary. *Audit closure.* The first nontrivial runner-ups already open visible residuals at every stage: the geomean runner-up gives $\max |\Delta(\det / K_{\text{bg}})| = 1.32 \times 10^{-3}$ and $\max |\Delta A| = 2.18 \times 10^{-4}$; the first ratio warp gives $\max |\Delta(\det / K_{\text{bg}})| = 7.18 \times 10^{-4}$ and $\max |\Delta \xi| = 6.08 \times 10^{-3}$; the first symmetry/normalization deformation gives $\max |\Delta(\det / K_{\text{bg}})| = 6.61 \times 10^{-4}$ and $\max |\Delta A| = 1.09 \times 10^{-4}$; and the first local generator-curvature deformation gives $\max |\Delta(\det / K_{\text{bg}})| = 2.76 \times 10^{-4}$ and $\max |\Delta \xi| = 2.23 \times 10^{-3}$. In particular, the generator-affinity audit shows that once the diagonal sectors and anchors are fixed, the projected parent action does not preserve any first local log-curvature correction to the mixed entry. After Lemma 1 has fixed the diagonal sectors, Lemma 2 therefore leaves a single symmetry-respecting, anchor-vanishing mixed coupling: the geometric mean.

Consistency remark. The link back to Lemma 1 is local but direct. In logarithmic coordinates the affine generator is exactly the part of g that survives local translations without producing new curvature data, whereas the quadratic jet is the first non-affine obstruction. After reweighting, that obstruction generates not only pure $\tau_{\Sigma}^2, \tau_{\Sigma} \tau_{\Delta}, \tau_{\Delta}^2$ terms but also mixed coordinate-dependent corrections once the translated jet is re-expanded about the original canonical point. On the projected parent block these corrections first enter through distortions of the mixed entry K_{12} , and therefore through the Schur response

$$\hat{G}_{\text{Schur}} = \frac{K_{22} - K_{12}^2 / K_{11}}{K_{\text{bg}}}.$$

This is why the generator-curvature audit is the numerically natural check of the affine-generator claim: it is probing the first local deformation that can alter the Schur term while leaving the diagonal sectors of Lemma 1 fixed.

Non-circular bridge. The generic off-diagonal follow-up audits support this conclusion without replacing the theorem proof. A direct determinant-inversion argument is currently circular because the exact parent-block determinant audit already hard-codes the geometric mean, so the meaningful independent question is whether generic off-diagonal witnesses can reproduce the even λ -curvature implied by the canonical mixed response. On the canonical $D21 \times E21$ fix map, the best global quartic proxy is a small blend of curvature-screened and action-absolute off-diagonal witnesses, but blocked holdout shows that those blend coefficients drift substantially across D -blocks. By contrast, the low- D regime is stably captured by the single normalized witness $\text{hk_abs_offdiag}/\text{diag}$, while the dense $D41 \times E21$ replay shows that the remaining complexity lives primarily in the tree-sector regime structure rather than in the participation envelope. We therefore use the generic off-diagonal program as supporting evidence that the geometric mean is the least arbitrary surviving global coupling, while keeping the theorem itself anchored to the non-circular generator-family closure stated above.

At that point the matrix entries are fixed, but the local variables in which those entries are expressed are still logically separate data. Lemma 3 closes that last freedom by showing that the linear excess variables isolated numerically are exactly the local normal-form coordinates of the projected parent block.

Lemma 3 (Linear excess variables as normal coordinates). Near the fixed point

$$(K_{11}, K_{22}) = (1, K_{\text{bg}}),$$

the natural local coordinates of the projected parent block are exactly

$$E_{\text{part}} = K_{11} - 1, \quad E_{\text{tree}} = K_{22} - K_{\text{bg}},$$

within the audited local coordinate families, in the sense that they are simultaneously the unique anchor-preserving coordinates, the unique unit-tangent-normalized coordinates, and the unique zero-second-jet coordinates compatible with exact determinant / Schur / weight closure.

Proof sketch. Let $E = (E_{\text{part}}, E_{\text{tree}})$ and let $\tilde{E} = \Phi(E)$ be any local reparameterization of the projected parent block near the fixed point $E = 0$. In this language a normal-form coordinate system is characterized by

$$\Phi(0) = 0, \quad D\Phi(0) = I, \quad D^2\Phi(0) = 0,$$

namely exact anchoring at the fixed point, unit tangent normalization there, and vanishing quadratic jet. Within the audited axis-preserving local coordinate families, these are also the

first three independent departures from the identity chart, because any such local chart admits an expansion of the form

$$\tilde{E} = b + LE + Q(E, E) + \mathcal{O}(\|E\|^3),$$

with b testing anchor shifts, L testing tangent normalization, and Q testing the first nonlinear jet. The three exact coordinate audits therefore probe the normal-form conditions in the natural order in which local coordinate freedom appears. The anchor family

$$E_{\text{part}}^{(r)} = K_{11} - r_{\text{part}}, \quad E_{\text{tree}}^{(r)} = K_{22} - r_{\text{tree}}K_{\text{bg}}$$

already closes uniquely at $(r_{\text{part}}, r_{\text{tree}}) = (1, 1)$, so $\Phi(0) = 0$ is forced at the identity participation block and the shell/background tree block. Among coordinate families that preserve those fixed points, the fixed-point Box-Cox audit then closes uniquely at

$$(p_{\text{part}}, p_{\text{tree}}) = (1, 1),$$

which is precisely the linear additive excess choice and therefore fixes $D\Phi(0) = I$. Finally, the nonlinear-jet family

$$E_{\text{part}}^{(\zeta_p)} = E_{\text{part}} + \zeta_p E_{\text{part}}^2, \quad E_{\text{tree}}^{(\zeta_t)} = E_{\text{tree}} + \zeta_t E_{\text{tree}}^2 / K_{\text{bg}}$$

closes uniquely at

$$(\zeta_p, \zeta_t) = (0, 0),$$

so the same coordinates are also zero-second-jet normal coordinates, i.e. $D^2\Phi(0) = 0$. The three exact closures therefore read together as a local normal-form statement: moving the anchor, changing the unit tangent normalization, or turning on a nonzero second jet each destroys exact determinant / Schur / weight closure. Equivalently, within the audited coordinate families, the identity map in the excess variables is the unique local chart that agrees with the projected parent block to quadratic order. Lemma 3 is thus the coordinate counterpart of Lemmas 1 and 2: once the block structure and mixed entry are fixed, the local variables are fixed as well. *Audit closure.* The first nontrivial deviation from each of $\Phi(0) = 0$, $D\Phi(0) = I$, and $D^2\Phi(0) = 0$ already opens visible residuals. In the anchor audit the first runner-up produces onset residuals 4.53×10^{-2} and 9.78×10^{-2} together with $\max|\Delta(\det/K_{\text{bg}})| = 1.08 \times 10^{-3}$ and $\max|\Delta\xi| = 1.61 \times 10^{-2}$. The first tangent-normalized Box-Cox runner-up produces onset residuals 2.02×10^{-3} and 8.23×10^{-3} with $\max|\Delta(\det/K_{\text{bg}})| = 8.25 \times 10^{-4}$ and $\max|\Delta\xi| = 7.85 \times 10^{-3}$. The first nonlinear-jet runner-up produces onset residuals 2.10×10^{-3} and 9.63×10^{-3}

with $\max |\Delta(\det/K_{\text{bg}})| = 9.90 \times 10^{-4}$ and $\max |\Delta\xi| = 9.47 \times 10^{-3}$. Hence the natural local variables are exactly the zero-second-jet linear excess coordinates.

Together, Lemmas 1–3 successively eliminate the three remaining ambiguities of the parent-kernel bridge: first the projected block structure, then the mixed coupling, and finally the local coordinate chart. At that point there is no remaining freedom in the local 2×2 parent block, because the diagonal sectors are fixed to $(1 + E_{\text{part}}, K_{\text{bg}} + E_{\text{tree}})$, the off-diagonal sector is fixed to $\sqrt{E_{\text{part}}E_{\text{tree}}}$, and the variables themselves are fixed to the linear zero-second-jet excess chart. The corollary is therefore not an informal restatement of the exact audits, but the compact parent-action statement obtained by substituting those three outputs back into the projected kernel of Lemma 1.

Corollary (Projected EYMH parent kernel). Combining Lemmas 1–3, and still working on the same local positive cone and within the same audited local coordinate families, the localized low-mode EYMH fluctuation operator projects to a background-normalized 2×2 parent kernel whose entries are now fixed term by term:

$$K_{11} = 1 + E_{\text{part}}, \quad K_{22} = K_{\text{bg}} + E_{\text{tree}}, \quad K_{12} = \sqrt{E_{\text{part}}E_{\text{tree}}}.$$

Substituting these three identities into the generic projected block of Lemma 1 gives the explicit canonical parent kernel

$$\mathcal{K}_{\text{parent}} = \begin{pmatrix} 1 + E_{\text{part}} & \sqrt{E_{\text{part}}E_{\text{tree}}} \\ \sqrt{E_{\text{part}}E_{\text{tree}}} & K_{\text{bg}} + E_{\text{tree}} \end{pmatrix}.$$

Its normalized determinant and response weight are therefore

$$\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} = \frac{(1 + E_{\text{part}})(K_{\text{bg}} + E_{\text{tree}}) - E_{\text{part}}E_{\text{tree}}}{K_{\text{bg}}} = 1 + E_{\text{part}} + \frac{E_{\text{tree}}}{K_{\text{bg}}},$$

$$A_{\text{resp}} \propto \left(\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} \right)^{-1/2} = \left(1 + E_{\text{part}} + \frac{E_{\text{tree}}}{K_{\text{bg}}} \right)^{-1/2}.$$

Equivalently, the same determinant may be read in Schur form as

$$\frac{\det \mathcal{K}_{\text{parent}}}{K_{\text{bg}}} = K_{11} \hat{G}_{\text{Schur}}, \quad \hat{G}_{\text{Schur}} = \frac{K_{22} - K_{12}^2/K_{11}}{K_{\text{bg}}} = 1 + \frac{E_{\text{tree}}}{K_{\text{bg}}(1 + E_{\text{part}})}.$$

This corollary is the compact parent-action statement that remains after the exact audit closures: the parent block is fixed by its two canonical diagonal sectors, the unique geometric-mean mixed coupling, and the unique linear zero-second-jet excess coordinates.

b. Reader-facing interpretation. A denser diagnostic replay of the canonical parented target on a $D41 \times E21$ grid sharpens the qualitative reading of the same kernel. On that dense grid the target

$$\xi_{\text{target}} = \frac{E_{\text{part}}E_{\text{tree}}}{K_{11}K_{22}}$$

is nearly η -independent at fixed D , so the dominant structure is genuinely D -driven. The two excess sectors then play visibly different roles. The participation excess E_{part} is a slow envelope: it stays broad, varies monotonically over long D -ranges, and reaches its largest values only in the late- D tail. By contrast, the tree/Schur excess E_{tree} carries the multi-peak structure: its strongest crest sits near $D \approx 6$, it exhibits the alternating low- and mid- D shoulders that generate the non-monotone target profile, and it is also the first sector to decline before the final high- D collapse. The target maximum at $D \approx 9.2$ is therefore not the peak of either sector separately, but the point where the smooth participation envelope and the oscillatory tree sector overlap most efficiently under the kernel normalization.

This decomposition is useful for interpretation because it explains why a single global phenomenological ansatz for $\xi_{\text{target}}(D)$ is only partially successful. The dense-grid diagnostics support a more structured reading:

$$\begin{aligned} E_{\text{part}}(D) &\approx \text{slow envelope,} \\ E_{\text{tree}}(D) &\approx \text{low/mid-}D \text{ multi-frequency oscillation} \\ &\quad \times \text{high-}D \text{ cutoff,} \end{aligned}$$

with ξ_{target} reconstructed from their normalized product rather than from a single primitive oscillatory law. In particular, a smooth-envelope model already describes E_{part} well, while E_{tree} prefers either a two-frequency global fit or, more cleanly, a piecewise regime fit with a low/mid- D oscillatory branch and a high- D cutoff branch. We therefore regard the phenomenological regime decomposition as a reader-facing diagnostic of the corollary, not as a replacement for the parent-kernel statement itself.

$$(D_0, \eta_0)_{\chi^2\text{-best}} \equiv \arg \min_{(D, \eta) \in \mathcal{G}} \chi^2(D, \eta), \quad (65)$$

$$(D_0, \eta_0)_{\text{robust}} \equiv \arg \max_{(D, \eta) \in \mathcal{A}} d_{\partial\mathcal{A}}(D, \eta), \quad \mathcal{A} \equiv \{(D, \eta) \in \mathcal{G} \mid \chi^2 \leq 4\}, \quad (66)$$

where $\mathcal{G}$ is the scan grid and $d_{\partial\mathcal{A}}$ is the Euclidean distance to the acceptance-boundary complement on that grid. We use $\mathcal{R}_3$ (higher first) and then χ^2 (lower first) as tie-breakers for Eq. (66). A full reference-normalization sensitivity audit is reported in Appendix E6: over the tested normalization grid, $f(\chi_{\mu\mu}^2 < 4)$ spans $[0.0333, 0.2411]$, while the baseline fixed reference gives 0.1500, the dynamic χ^2 -best reference gives 0.1333, and the robust-center reference gives 0.1381. To reduce reference-normalization arbitrariness in reporting, we additionally track Wilson-diagonal ratios that are invariant under the same re-normalization:

$$\mathcal{R}_{e/\mu}^{(\text{inv})}(D, \eta) \equiv \left| \frac{C_{eH}^{ee, \text{IR}}(D, \eta)}{C_{eH}^{\mu\mu, \text{IR}}(D, \eta)} \right|^2, \quad \mathcal{R}_{\tau/\mu}^{(\text{inv})}(D, \eta) \equiv \left| \frac{C_{eH}^{\tau\tau, \text{IR}}(D, \eta)}{C_{eH}^{\mu\mu, \text{IR}}(D, \eta)} \right|^2. \quad (67)$$

Using `code/scan_anchor_invariant_ratios.py` on `paper/hll_uv_to_eft_map.csv`, we obtain (all-grid) $p10/p50/p90$ bands $\mathcal{R}_{e/\mu}^{(\text{inv})} = (1.60 \times 10^{-2}, 1.26 \times 10^{-1}, 8.10 \times 10^{-1})$ and $\mathcal{R}_{\tau/\mu}^{(\text{inv})} = (2.59 \times 10^{-3}, 1.01 \times 10^{-2}, 3.84 \times 10^{-2})$; restricted to the illustrative acceptance band $\chi_{\mu\mu}^2 < 4$, they become $\mathcal{R}_{e/\mu}^{(\text{inv})} = (1.67 \times 10^{-2}, 6.39 \times 10^{-2}, 3.54 \times 10^{-1})$ and $\mathcal{R}_{\tau/\mu}^{(\text{inv})} = (5.23 \times 10^{-3}, 9.72 \times 10^{-3}, 5.17 \times 10^{-2})$. These diagnostics are anchor-independent and are reported as complementary constraints to Eq. (56).

$$\chi^2(D, \eta) = \frac{(\mu_{\mu\mu}^{\text{pred}} - \mu_{\mu\mu}^{\text{obs}})^2}{\sigma_{\mu\mu}^2}, \quad (68)$$

where $\sigma_{\mu\mu} = 0.4$. For one degree of freedom, we also report the local p -value

$$p_{\text{local}}(D, \eta) = 1 - F_{\chi_1^2}(\chi^2(D, \eta)) = \text{erfc}\left(\sqrt{\chi^2(D, \eta)/2}\right). \quad (69)$$

Figure 8 shows the **EFT/Wilson-matched acceptance region** in two equivalent languages,

$$\chi^2 < 4 \quad \Longleftrightarrow \quad p_{\text{local}} > 4.55 \times 10^{-2} \quad (\text{dof} = 1). \quad (70)$$

We use this threshold only as an *illustrative acceptance band* for map-level comparison and do not claim a formal 95% CL exclusion. Because the (D, η) scan is a multiple-comparison procedure (look-elsewhere) and no global trials correction is applied, this band is interpreted only as a falsifiable target region. In the present full-grid action-derived profile scan, the accepted fraction is

$$f(\chi^2 < 4) = f(p_{\text{local}} > 4.55 \times 10^{-2}) = 0.1500, \quad (71)$$

with best grid point $\chi^2 \simeq 2.21 \times 10^{-5}$ at $(D, \eta) \approx (4.00, 0.264)$. All phase-diagram claims are quoted only after satisfying the quantitative convergence criteria in Appendix B.

E. EFT/Wilson-Matched Extension to $H \rightarrow ee$ and $H \rightarrow \tau\tau$

Using the same map-level EFT/Wilson-matched definition, we additionally evaluate lepton-channel signal strengths by layer assignment $N = 1 \rightarrow ee$, $N = 2 \rightarrow \mu\mu$, and $N = 3 \rightarrow \tau\tau$:

$$\mu_{\ell\ell}^{\text{pred}}(D, \eta) = \left| \frac{C_{eH}^{\ell\ell}(D, \eta)}{C_{eH}^{\ell\ell}(D_0, \eta_0)} \right|^2 \left[\frac{\Gamma_{\text{tot}}(D, \eta)}{\Gamma_{\text{tot}}(D_0, \eta_0)} \right]^{-1}. \quad (72)$$

Figure 9 summarizes the corresponding maps on the same (D, η) grid; tabulated statistics are stored in `paper/hll_signal_strength_summary.csv`. This extension is still a scan-level EFT map (not full UV matching), while the visibility sector in the global scan uses baseline fully normalized EFT operator visibility built from overlap-extracted inputs with mode




FIG. 8. Compatibility with $H \rightarrow \mu\mu$ under the EFT/Wilson-matched UV+LL-RG map. Left: χ^2 map showing a restricted illustrative accepted band ($\chi^2 < 4$, equivalently $p_{\text{local}} > 4.55 \times 10^{-2}$ for dof= 1). Right: predicted $\mu_{\mu\mu}^{\text{pred}}$.

tracking and microcanonical windowing. The repository keeps the three-channel observation fields in `data/pdg_leptons.json` as a reproducible interface. The corresponding runtime-direct visibility branch is release-qualified only in the risk-weighted profile-blend form; the reported baseline maps remain full-direct and EFT-operator normalized unless stated otherwise.




FIG. 9. PSLT EFT/Wilson-matched extension from $H \rightarrow \mu\mu$ to three lepton channels on the same scan grid: $H \rightarrow ee$ (left), $H \rightarrow \mu\mu$ (middle), and $H \rightarrow \tau\tau$ (right). Each panel shows $\mu_{\ell\ell}^{\text{pred}}$ from Eq. (56) with the same reference-normalization point $(D_0, \eta_0) = (10, 1)$.

VIII. DISCUSSION AND CONCLUSION

We have presented a computable EFT-level demonstrator of the Projection Spectral Layer Theory. By closing the loop between geometric inputs and observable layer probabilities, we have shown that:

1. **Generations are spectral within the baseline closure:** The “three generation” structure is not imposed as an input; it emerges on the current release baseline as a

dynamical output of competing entropy (g_N), kinetics (Γ_N), and visibility (B_N).

2. **Stability without ad-hoc cutoffs:** The infinite tower of layers becomes numerically and physically stable once micro-degeneracy is regulated by the `fp_2d_full` microcanonical-window + tail prescription and the observable sector uses EFT-operator-normalized visibility profiles $B_N(D)$ with neutral high- N saturation rather than an extra visibility-side cutoff.
3. **Falsifiability:** The closure yields concrete predictions for the winner phase diagram. With full-grid action-derived $\chi_N^{(LR)}(D)$ and $\tilde{A}_\ell(D)$ profiles, the $H \rightarrow \mu\mu$ illustrative EFT/Wilson-matched acceptance band ($\chi^2 < 4 \Leftrightarrow p_{\text{local}} > 4.55 \times 10^{-2}$, dof= 1) covers about 15.0% of the scan and defines a falsifiable target region in (D, η) .

At the same time, the present implementation keeps two EFT-level approximations explicit: map-level profile closure remains active in the baseline $B_N(D)$ visibility sector (while release-production and risk-weighted runtime-direct parity branches are both available) and a bounded nearest-neighbor flavor projector in the matched $H \rightarrow \mu\mu$ map of Eq. (56). Therefore, the present manuscript supports a robust statement about spectral generation selection on the release baseline, and also about a release-qualified runtime-direct visibility parity path, but not yet a strict claim that the observable sector is closed by per-cell all-direct extraction at production level.

Limitations and outlook. The present closure is internally consistent but remains an EFT-level demonstrator in several places: The use of mixed-resolution ingredients is a deliberate design choice: action-derived extraction fixes low-mode profiles (`fp_2d_full` for g_N , EFT-operator-normalized $B_N(D)$ from overlap-extracted inputs, $\chi_N^{(LR)}(D)$, and $\tilde{A}_\ell(D)$), while global-map evaluation is still an EFT-level closure rather than a full direct localized recomputation at every (D, η, N) point.

1. **Observable mapping:** the $H \rightarrow \mu\mu$ comparison now uses a scan-level EFT/Wilson-matched map with bounded flavor mixing and leptonic-width closure [Eqs. (50)–(57)]. The fixed point in Eq. (56) should be read only as a reference-normalization convention for the relative observable, while the stronger normalization-side evidence now comes from the absolute and loop-improved EYMH parent-action comparators and the Wilson-diagonal re-normalization invariants. The new absolute-bridge audit shows that this arbitrariness reduces exactly to three scalar bridge constants ($A_\star, R_{e/\mu}^\star, R_{\tau/\mu}^\star$) and that the accepted-region invariant median already fixes $R_{e/\mu}^\star$ to practical numerical precision

on the canonical map. A full parent-side derivation of the remaining absolute amplitude $A_\star$ and tau-to-muon ratio $R_{\tau/\mu}^\star$ is still pending.

2. **Mixing channel:** under the parity-symmetric first-principles definition [Eq. (73)], we obtain $\chi_N^{(\text{sym})} \sim 10^{-19}$ for $D = \{6, 12, 18\}$ (Appendix C), i.e., symmetry-protected cancellation. We therefore redefine the channel in localized basis, extract $\chi_N^{(LR)}$ (Appendix D), and inject its full scan-grid action-derived profile in the global map. The remaining gap is a full (D, η, N) localized projection.
3. **Open-system extension:** a Lindblad-type χ_{eff} mechanism is now implemented as a scan-ready diagnostic mode with profiled $\gamma_\phi(D), \gamma_{\text{mix}}(D)$ (Appendix E 4). The original multi-anchor + holdout calibration remains in place, but the line is now much narrower than a generic reservoir ansatz: on the audited knot set, κ_{env} is isolated as a scalar amplitude modulus, the localized two-level projection leaves only the leading σ_z/σ_x directions, and the effective kernels are on-shell samples of a single gap-locked single-pole response family with $\tau_{\text{env},a} = \omega_{1,a}^{-1}$ to machine precision. The sharp remaining theorem target is parent-side: show that the projected EYMH Schur resolvent admits a first-mode-dominant truncation on $\omega \in \{0, \Delta E_a\}$ with leading pole scale $\Omega_{1,a} = \omega_{1,a}$ and uniformly small higher-mode remainder. Until that step is taken, this block remains outside baseline figures.
4. **Model-chain unification:** $\chi_N^{(LR)}(D)$ and $\tilde{A}_\ell(D)$ are injected from action-derived full-grid profiles in the global (D, η) map, while the release baseline keeps strict D-grid lookup and the automatic chain remains a comparator. The release-production path now pairs active-grid direct two-dimensional spectra for $g_N(D)$ with per-cell direct evaluation of $\chi_{LR}(D)$ and $\tilde{A}_\ell(D)$ while keeping baseline B_N profile closure; on both release audits it matches the release baseline exactly ($\Delta f(\chi_{\mu\mu}^2 < 4) = 0$, mismatch = 0, $\max |\Delta \mu_{\mu\mu}| = 0$). The tuned visibility path should therefore be read as a release-qualified risk-weighted visibility parity check, not as a strict all-direct closure of B_N . Separately, the broader-grid strict all-direct validator is parked at its current case-aware endpoint on the two validation surfaces, so this line should be reopened only for a genuinely new observable-side mechanism class rather than another local blend retune. For the UV-to-EFT map, the refreshed diagonal-threshold comparator and the absolute/loop-improved EYMH comparators show the same pattern: zero acceptance mismatch, exact basis reconstruction, and only small deformation. The remaining UV-to-EFT gap is therefore no longer the existence of a stable local absolute witness, but the derivation of its prefactor from a bona

fide parent EYMH loop calculation.

5. **Spectral tower:** the baseline validation is strongest for the low-lying bound sector ($N = 1, 2$ in the physical-gap scan), while the high- N contribution is treated at EFT level through (g_N, B_N) regularization. The audited single-track benchmark does not exhibit a fourth bound layer on the tested D window, and Proposition N1 / Corollary N1 promote the one-dimensional kinetic proxy to a mode-resolved finite-volume Sturm threshold theorem: on the audited domain the fourth-layer absence is a threshold/binding-loss statement, not an engineered visibility cutoff or a hidden exponentially small- Γ_4 mechanism. The only remaining strengthening is a continuum-tail / D -continuous interpolation version of the same certificate.
6. **Finite-time dynamics:** t_{coh} is treated as a control parameter in the baseline scan; a geometry-closed dephasing candidate $t_{\text{coh}}^{(\text{deph})}(D) = \pi/\Delta\omega_{12}(D)$ is benchmarked in Appendix E 7 but is not yet propagated into the reported baseline figures.
7. **Overlap amplitude:** η is scanned as an external control in the baseline maps. A first-principles prefactor candidate $\eta_{\text{fp}}(D)$ from splitting-action data is benchmarked in Appendix E 8; map-level impact is mild in profile-scaled mode but this closure is not yet propagated into baseline figures.
8. **Barrier-leakage channel normalization:** the baseline maps now use action-derived profile factors $(A_1, A_2) = (\tilde{A}_1(D), \tilde{A}_2(D))$ from Appendix E 9 (historical “superrad” artifact filenames retained for reproducibility). The remaining gap is to replace profile-level insertion with a fully localized, channel-resolved extraction across the full scan grid.

Status and redefinition of the mixing channel. Using Eq. (33), we fix a basis-independent normalization scale $\bar{\Gamma}_N$ and define the parity-symmetric overlap channel by

$$\chi_N^{(\text{sym})}(D) = \frac{|M_{12}^{(\text{sym})}(D, N)|}{\bar{\Gamma}_N(D)}, \quad (73)$$

$$M_{12}^{(\text{sym})}(D, N) = 2\pi \int_0^{\rho_{\text{max}}} \rho d\rho \int_{-z_{\text{max}}}^{z_{\text{max}}} dz \psi_{N,1}^*(\rho, z) \delta V(\rho, z; D) \psi_{N,2}(\rho, z),$$

and Eq. (38) with the explicit spherical-average δV gives $\chi_N^{(\text{sym})} \approx 0$ (Appendix C). To avoid mixing two inequivalent notions, we redefine the phenomenological mixing channel in a localized-well basis:

$$\psi_{N,L} = \frac{\psi_{N,+} + \psi_{N,-}}{\sqrt{2}}, \quad \psi_{N,R} = \frac{\psi_{N,+} - \psi_{N,-}}{\sqrt{2}}, \quad (74)$$

and define the channel mixing from the corresponding two-state Hamiltonian element:

$$M_{LR}^{(H)}(D, N) \equiv \frac{\lambda_{N,2}(D) - \lambda_{N,1}(D)}{2}, \quad \chi_N^{(LR)}(D) \equiv \frac{|M_{LR}^{(H)}(D, N)|}{\bar{\Gamma}_N(D)}. \quad (75)$$

The baseline mixing entry in Table III is interpreted as the localized-channel coefficient. First-principles localized extraction results are given in Appendix D, and its full scan-grid action-derived profile is used in the global (D, η) scan of Section VII. In terminology, $M_{LR}^{(H)}$ is a *closed-system localized-basis splitting* quantity; it is not itself a decoherence rate. Open-system decoherence is introduced separately through Lindblad rates $(\gamma_\phi, \gamma_{\text{mix}})$ in Appendix E 4.

ACKNOWLEDGMENTS

The author acknowledges the use of PDG 2024 data and standard Python scientific stacks for verification.

Appendix A: Action-Derived Chain Reproducibility SOP

This appendix provides a clone-level SOP for reproducing the action-derived extraction chain and the paper-level Fig/Table artifacts.

1. Single Entry Point

From repository root:

```
bash scripts/repro/reproduce_paper.sh
```

Optional modes:

```
bash scripts/repro/reproduce_paper.sh --with-paper
```

```
bash scripts/repro/reproduce_paper.sh --package-only
```

2. Step Map and Deliverables

3. Acceptance Conditions

A run is considered reproducible for this manuscript if:

1. the one-click command exits with return code 0;
2. `repro/latest/manifest.json` reports `missing_required = 0`;

TABLE VIII. Compact SOP stage map for action-derived reproducibility. The one-click wrapper executes these stages in a fixed order, logs each stage, and records exact script/output names in the artifact registry.

Stage group	Deliverables and checks
Production maps and observables	Regenerates phase diagrams, Higgs signal-strength maps, and UV-to-EFT matching summaries.
Release/direct validation	Runs the full-direct release audit, localized-direct validation surfaces, chain-mode parity checks, and artifact-status registry.
First-principles profile diagnostics	Recomputes localized χ , bounded phase-space g_N , open-system diagnostics, and the t_{coh} , η , and channel-normalization candidate profiles.
Packaging and provenance	Collects figures, tables, logs, checksums, manifest metadata, and the per-artifact map under the reproducibility run directory.

3. `repro/latest/checksums.sha256` and `repro/latest/artifact.index.csv` are generated and non-empty.

Appendix B: Action-Derived Effective Potential and WKB

1. Two-Center Harmonic Conformal Factor

We define the conformal factor by a projected Poisson problem on the static slice,

$$\nabla^2 \Omega(\mathbf{x}) = -4\pi a [\rho_\varepsilon(\mathbf{x} - \mathbf{x}_+) + \rho_\varepsilon(\mathbf{x} - \mathbf{x}_-)], \quad \Omega \rightarrow 1 \quad (|\mathbf{x}| \rightarrow \infty), \quad (\text{B1})$$

with $\mathbf{x}_\pm = \pm(D/2)\hat{z}$. Using the Green representation

$$\Omega(\mathbf{x}) = 1 + a \sum_{s=\pm} \int d^3 x' \frac{\rho_\varepsilon(\mathbf{x}' - \mathbf{x}_s)}{|\mathbf{x} - \mathbf{x}'|}, \quad (\text{B2})$$

and the normalized Plummer kernel

$$\rho_\varepsilon(\mathbf{r}) = \frac{3\varepsilon^2}{4\pi (|\mathbf{r}|^2 + \varepsilon^2)^{5/2}}, \quad \int d^3 r \rho_\varepsilon(\mathbf{r}) = 1, \quad (\text{B3})$$

the integral is analytic and gives

$$\Omega(\rho, z; D) = 1 + a \left(\frac{1}{\sqrt{\rho^2 + (z - D/2)^2 + \varepsilon^2}} + \frac{1}{\sqrt{\rho^2 + (z + D/2)^2 + \varepsilon^2}} \right) \quad (\text{B4})$$

where a is the geometric strength, ε is the core regularization scale, and D is the dual-center separation. Away from the cores ($|\mathbf{x} - \mathbf{x}_\pm| \gg \varepsilon$), this satisfies $\nabla^2 \Omega \approx 0$. The radial identity

$$\nabla^2 \left(\frac{1}{\sqrt{r^2 + \varepsilon^2}} \right) = -\frac{3\varepsilon^2}{(r^2 + \varepsilon^2)^{5/2}} \quad (\text{B5})$$

gives the smeared Laplacian:

$$\begin{aligned}\nabla^2\Omega &= -3a\varepsilon^2 \left[(r_+^2 + \varepsilon^2)^{-5/2} + (r_-^2 + \varepsilon^2)^{-5/2} \right] \\ &< 0.\end{aligned}\tag{B6}$$

which is equivalent to the source form in Eq. (6). This negative contribution is essential for forming the attractive wells in V_{eff} . As a reproducibility check, `verify_omega_geometric_origin.py` confirms $\int d^3r \rho_\varepsilon = 1$ with absolute error 3.75×10^{-7} and validates the radial Laplacian identity with max relative error 1.39×10^{-4} in the non-asymptotic band; run outputs are listed in the artifact registry.

2. Veff Derivation from KG Equation

Starting from $(\square_g - m_0^2 - \xi R)\Phi = 0$ with $g_{\mu\nu} = \Omega^2\eta_{\mu\nu}$:

a. Step 1: $\square_g$ in conformal coordinates.

$$\square_g\Phi = \Omega^{-2}(-\partial_t^2 + \nabla^2)\Phi + 2\Omega^{-3}\nabla\Omega \cdot \nabla\Phi\tag{B7}$$

b. Step 2: Time separation. For $\Phi = e^{-i\omega t}\phi(\mathbf{x})$:

$$\nabla^2\phi + 2\Omega^{-1}\nabla\Omega \cdot \nabla\phi + [\omega^2 - \Omega^2(m_0^2 + \xi R)]\phi = 0\tag{B8}$$

c. Step 3: Conformal field rescaling. Setting $\phi = \Omega^{-1}\psi$ eliminates the first-derivative term. Using $R = -6\Omega^{-3}\nabla^2\Omega$ for 4D conformally flat space:

$$\boxed{[-\nabla^2 + V_{\text{eff}}]\psi = \omega^2\psi, \quad V_{\text{eff}} = m_0^2\Omega^2 + (1 - 6\xi)\Omega^{-1}\nabla^2\Omega}\tag{B9}$$

d. Double-well formation condition. For $\xi < 1/6$ and $\nabla^2\Omega < 0$ near cores, the curvature term contributes **negatively** to V_{eff} , creating attractive wells. The mass term $m_0^2\Omega^2$ provides the continuum threshold at infinity.

e. Self-consistency check. At $\xi = \xi_c = 1/6$ (conformal coupling): $V_{\text{eff}} \rightarrow m_0^2\Omega^2$, and the curvature contribution vanishes identically. $\checkmark$

3. Minimal Dirac Coupling in the Conformal Background

To make the visibility block explicitly compatible with the fermionic Yukawa origin, we summarize the minimal conformal Dirac reduction used by the baseline overlap kernel.

For $g_{\mu\nu} = \Omega^2\eta_{\mu\nu}$, choose

$$e_\mu^a = \Omega \delta_\mu^a, \quad e_a^\mu = \Omega^{-1}\delta_a^\mu, \quad \sqrt{-g} = \Omega^4.\tag{B10}$$

The curved-space Dirac action is

$$S_\Psi = \int d^4x \sqrt{-g} \, i \bar{\Psi} e^\mu_a \gamma^a \left(\partial_\mu + \frac{1}{4} \omega_{\mu bc} \gamma^{bc} \right) \Psi, \quad (\text{B11})$$

and the canonical Weyl rescaling

$$\Psi = \Omega^{-3/2} \psi \quad (\text{B12})$$

removes the first-derivative conformal piece from the spin connection, yielding the flat-frame kinetic normalization. For the Higgs mode we use

$$H = \Omega^{-1} \varphi. \quad (\text{B13})$$

The Yukawa density therefore carries net conformal power

$$\Omega^4 \Omega^{-3/2} \Omega^{-1} \Omega^{-3/2} = \Omega^0, \quad (\text{B14})$$

so the minimal overlap kernel has

$$\mathcal{W}_{\text{frame}} = 1. \quad (\text{B15})$$

This is the baseline convention in the three-channel effective-Yukawa extraction; nonzero frame power is treated only as a diagnostic deformation.

As a reproducibility check, the minimal Dirac-conformal audit exports a machine-readable CSV and paper copy verifying the zero net conformal power in both kinetic and Yukawa factors.

4. 1D On-Axis Reduction

For efficient scanning, we use the on-axis ($\rho = 0$) reduction:

$$\left[-\frac{d^2}{dz^2} + U(z) \right] \psi_n(z) = E_n \psi_n(z), \quad U(z) = V_{\text{eff}}(\rho = 0, z) - m_0^2 \quad (\text{B16})$$

where $E_n = \omega_n^2 - m_0^2$. Bound states satisfy $E_n < 0$; stable modes additionally require $\omega_n^2 = m_0^2 + E_n > 0$.

5. 2D Axisymmetric Validation

We validated the 1D reduction by comparing with the full 2D axisymmetric Laplacian:

$$\nabla^2 \Omega = \frac{\partial^2 \Omega}{\partial \rho^2} + \frac{1}{\rho} \frac{\partial \Omega}{\partial \rho} + \frac{\partial^2 \Omega}{\partial z^2} \quad (\text{B17})$$

For the physical-gap parameters ($a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$), the relative error at $\rho = 0$ is:

D	$\text{Max } U_{2D} - U_{1D} / \text{max } U $
6	4.0×10^{-4}
12	4.0×10^{-4}
18	4.0×10^{-4}

The 1D 3D-identity approximation is accurate to $< 0.1\%$, validating its use for phase scans.

6. WKB Action Convention

The tunneling action through the central barrier is:

$$S_N = \int_{z_1}^{z_2} \sqrt{U(z) - E_N} dz \quad (\text{B18})$$

where $z_{1,2}$ are the turning points satisfying $U(z_{1,2}) = E_N$. The tunneling suppression factor is:

$$r_N = \eta \cdot e^{-2S_N} \quad (\text{B19})$$

Convention note: For rates/probabilities we use e^{-2S} , while the level-splitting amplitude scales as e^{-S} . The empirical splitting relation below therefore tests the amplitude-level law.

7. Splitting–Action Consistency Check

The most stringent internal consistency check is the relation between level splitting and tunneling action. For symmetric double wells, WKB theory predicts $\Delta E \propto e^{-S}$. From our action-derived scan:

$$\boxed{\ln \Delta E \approx -1.01 S_1 + 0.69, \quad R^2 = 0.9999} \quad (\text{B20})$$

This near-perfect linear relation (Fig. 10) confirms that the **same** V_{eff} controls both the spectrum and the tunneling kinetics—not an engineered coincidence.

8. Action-Derived Numerical Results

Table IX shows the key spectral and tunneling quantities from the action-derived calculation. Grid convergence: $|E_1|$ variation $< 0.2\%$ for $dz \in [0.01, 0.02]$ at fixed $z_{\text{max}} = 80$.

9. Bound-State omega Convergence Benchmark

To make the surrogate-vs-exact distinction explicit, we report a direct bound-state benchmark for ω_N from the same action-derived 1D operator chain used in Table IX. The settings are




FIG. 10. Splitting-action consistency check. The level splitting $\Delta E = E_2 - E_1$ follows $\exp(-S_1)$ to $R^2 = 0.9999$, confirming the action-derived V_{eff} governs both spectrum and tunneling.

TABLE IX. Action-derived spectral and tunneling data. Parameters: $a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$.

D	E_1	E_2	ω_1	S_1	ΔE
4	-0.505	-0.390	0.704	2.92	1.15×10^{-1}
8	-0.476	-0.469	0.724	5.68	6.32×10^{-3}
12	-0.478	-0.478	0.722	8.51	3.69×10^{-4}
16	-0.481	-0.481	0.721	11.33	2.19×10^{-5}
20	-0.482	-0.482	0.720	14.16	1.30×10^{-6}

fixed to the physical-gap parameter set ($a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$), $z_{\text{max}} = 80$, and three resolutions: coarse ($N_z = 4001$), mid ($N_z = 6001$), fine ($N_z = 8001$). The reproducible script is `code/extract_omega_exact_convergence.py`.

This benchmark is the operator-level reference for ω_N in the bound-state convention. In contrast, Appendix D uses generalized localized-extraction eigenvalues λ for mixing-channel extraction; the two quantities serve different roles and are not identified.

TABLE X. Bound-state ω_N convergence benchmark at $D = \{6, 12, 18\}$ from the action-derived operator chain. Relative errors are quoted versus fine grid.

D	level	N_z	ω_1	ω_2	$\Delta\omega_{12}$	max rel. ω vs fine
6	coarse	4001	0.720883	0.738989	1.8107×10^{-2}	8.65×10^{-4}
6	mid	6001	0.721345	0.739468	1.8123×10^{-2}	2.18×10^{-4}
6	fine	8001	0.721500	0.739629	1.8129×10^{-2}	0
12	coarse	4001	0.721781	0.722035	2.5438×10^{-4}	8.84×10^{-4}
12	mid	6001	0.722258	0.722513	2.5533×10^{-4}	2.23×10^{-4}
12	fine	8001	0.722418	0.722674	2.5565×10^{-4}	0
18	coarse	4001	0.719404	0.719408	3.6700×10^{-6}	8.92×10^{-4}
18	mid	6001	0.719885	0.719888	3.6943×10^{-6}	2.25×10^{-4}
18	fine	8001	0.720046	0.720050	3.7025×10^{-6}	0

10. Fixed-dz Convergence

We verified numerical stability using fixed grid spacing dz (rather than fixed N_z):

- For $dz \in \{0.04, 0.02, 0.01\}$ and $z_{\max} \in \{60, 80\}$: E_1 variation $< 0.3\%$, S_1 variation $< 0.2\%$.
- Results are independent of $z_{\max}$ for $z_{\max} > 60$ (domain truncation error negligible).

Appendix C: First-Principles Symmetry-Channel Test for χ

Using the explicit definition in Eq. (73), we computed M_{12} and $\chi_N^{(\text{sym})}$ on 2D axisymmetric grids for $D = \{6, 12, 18\}$ with coarse/mid/fine resolutions. The modes are normalized with the cylindrical measure $2\pi\rho d\rho dz$, and orthogonality satisfies $|\langle\psi_1, \psi_2\rangle| < 10^{-15}$ in all runs.

TABLE XI. Fine-grid symmetry-channel extraction using Eq. (73). Parameters: $a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$.

D	$ M_{12}^{(\text{sym})} $	$\chi_N^{(\text{sym})}$	$ \langle\psi_1, \psi_2\rangle $
6	2.09×10^{-16}	1.66×10^{-19}	3.12×10^{-17}
12	6.03×10^{-16}	5.69×10^{-19}	3.47×10^{-18}
18	6.24×10^{-16}	6.20×10^{-19}	2.11×10^{-16}

These values are numerically consistent with symmetry-protected cancellation of the parity-symmetric overlap channel. Across coarse/mid/fine grids, $M_{12}^{(\text{sym})}$ changes sign but remains at $\mathcal{O}(10^{-16})$, indicating no resolved nonzero mixing in this channel.

For this reason, the phenomenological Rank-2 coefficient in the main text is interpreted as a localized-channel effective coupling (Section VIII), not as $\chi_N^{(\text{sym})}$ from Eq. (73).

TABLE XII. Localized-channel extraction for $D = \{6, 12, 18\}$ (coarse/mid/fine). Parameters: $a = 0.04$, $\varepsilon = 0.1$, $m_0 = 1$, $\xi = 0.14$.

D	level	$(d\rho, dz)$	(N_ρ, N_z)	λ_1	λ_2	$M_{LR}^{(H)}$	$\chi_N^{(LR)}$	$ M_{12}^{(\text{sym})} $
6	coarse	(0.12, 0.06)	(25, 300)	2.14765	2.60073	2.26541×10^{-1}	4.13261×10^{-4}	5.29×10^{-17}
6	mid	(0.08, 0.04)	(38, 450)	2.14362	2.59798	2.27178×10^{-1}	4.17350×10^{-4}	2.44×10^{-19}
6	fine	(0.06, 0.03)	(50, 600)	2.16611	2.62106	2.27474×10^{-1}	4.01827×10^{-4}	1.56×10^{-18}
12	coarse	(0.12, 0.06)	(25, 400)	2.36032	2.71730	1.78492×10^{-1}	2.27261×10^{-4}	2.89×10^{-17}
12	mid	(0.08, 0.04)	(38, 600)	2.35598	2.71374	1.78883×10^{-1}	2.29384×10^{-4}	6.96×10^{-17}
12	fine	(0.06, 0.03)	(50, 800)	2.37821	2.73632	1.79054×10^{-1}	2.21414×10^{-4}	1.03×10^{-17}
18	coarse	(0.12, 0.06)	(25, 500)	2.22204	2.49466	1.36311×10^{-1}	2.18683×10^{-4}	1.94×10^{-17}
18	mid	(0.08, 0.04)	(38, 750)	2.21686	2.49001	1.36575×10^{-1}	2.21053×10^{-4}	1.54×10^{-15}
18	fine	(0.06, 0.03)	(50, 1000)	2.23864	2.51202	1.36694×10^{-1}	2.13187×10^{-4}	2.19×10^{-16}

Appendix D: Localized-Channel First-Principles Extraction of χ

We compute the nonzero mixing channel directly in localized basis using

$$M_{LR}^{(H)} = \frac{\lambda_2 - \lambda_1}{2}, \quad \chi_N^{(LR)} = \frac{|M_{LR}^{(H)}|}{\bar{\Gamma}_N}, \quad (\text{D1})$$

with $\bar{\Gamma}_N$ from Eq. (33). Here $\lambda_{1,2}$ are generalized operator eigenvalues from $K\psi = \lambda M\psi$ in the localized extraction solver; they are not the bound-state energy convention $E = \omega^2 - m_0^2 < 0$ used in Appendix B. The extraction uses `code/extract_chi_localized_2d.py` with fixed settings:

- Domain: $\rho_{\text{max}} = 3.0$, $z_{\text{max}} = D/2 + 6.0$.
- Grids: coarse $(d\rho, dz) = (0.12, 0.06)$, mid $(0.08, 0.04)$, fine $(0.06, 0.03)$.
- Solver: generalized eigensystem (shift-invert low-mode targeting, $\sigma = 2.5$), tolerance 10^{-8} , maxiter 3×10^4 .
- Acceptance criteria: $\delta_{\Delta E} \equiv |\Delta E_{\text{level}} - \Delta E_{\text{fine}}|/|\Delta E_{\text{fine}}| < 5\%$, $\delta_\chi \equiv |\chi_{\text{level}} - \chi_{\text{fine}}|/|\chi_{\text{fine}}| < 5\%$, and null-channel absolute bound $|M_{12}^{(\text{sym})}| < 10^{-12}$.

The complete pipeline, scripts, and raw tables used here are available in the project repository: <https://github.com/boypatrick/PSLT>.

Across all non-fine runs in Table XII, we obtain

$$\max \delta_{\Delta E} = 4.10 \times 10^{-3}, \quad \max \delta_\chi = 3.86 \times 10^{-2}, \quad \max |M_{12}^{(\text{sym})}| = 1.54 \times 10^{-15}, \quad (\text{D2})$$

which satisfies the stated convergence and absolute-null criteria.

TABLE XIII. Auxiliary fine-grid localized profile points from the full-scan run ($D = 4\text{--}20$). Source CSV: `output/chi_fp_2d/localized_chi_D4-5-...-20.csv`.

D	λ_1	λ_2	$\chi_N^{(LR)}$
4	2.06671	2.56696	5.266×10^{-4}
5	2.34635	2.85471	3.311×10^{-4}
6	2.16611	2.62106	4.018×10^{-4}
7	2.42273	2.88524	2.661×10^{-4}
8	2.24778	2.66479	3.202×10^{-4}
9	2.10111	2.48071	3.758×10^{-4}
10	2.32055	2.70644	2.623×10^{-4}
11	2.17620	2.52926	3.064×10^{-4}
12	2.37821	2.73632	2.214×10^{-4}
13	2.24440	2.57510	2.554×10^{-4}
14	2.42772	2.76173	1.906×10^{-4}
15	2.30010	2.61030	2.181×10^{-4}
16	2.47457	2.78817	1.661×10^{-4}
17	2.34884	2.64089	1.895×10^{-4}
18	2.23864	2.51202	2.132×10^{-4}
19	2.39519	2.67161	1.663×10^{-4}
20	2.28581	2.54509	1.867×10^{-4}

For interpolation support and reproducibility of scan-level diagnostics, we additionally ran a fine-grid-only auxiliary sweep on integer separations $D = 4, \dots, 20$ using the same solver settings (`--full-scan` in `extract_chi_localized_2d.py`). These points are used as profile-support data in the repository, while the formal acceptance criteria above remain defined by the benchmark convergence set $D = \{6, 12, 18\}$.

Figure 11 shows the fine-grid 2D parity eigenmodes used in the localized extraction at $D = \{6, 12, 18\}$. The expected even/odd nodal structure is manifest and remains stable across the three benchmark separations.

To test profile uncertainty in the global map, we benchmarked six $\chi_N^{(LR)}(D)$ choices built from the same fine-grid knots $(D, \chi_N^{(LR)}) = \{(6, 4.02 \times 10^{-4}), (12, 2.21 \times 10^{-4}), (18, 2.13 \times 10^{-4})\}$: linear interpolation, log-linear exponential fit $\ln \chi = -7.5956 - 0.05282 D$, and $\pm 20\%$ amplitude rescalings of each profile. Re-running the full 60×60 scan keeps the generation fractions fixed and yields a narrow $H \rightarrow \mu\mu$ acceptance spread:

$$f(\mathcal{R}_3 > 0.90) = 0.927, \quad f(\mathcal{R}_3 > 0.95) = 0.217, \quad f(\chi_{\mu\mu}^2 < 4) \in [0.1625, 0.1681]. \quad (\text{D3})$$

The best-fit point moves within $(D, \eta) \approx (8.07\text{--}9.42, 1.29\text{--}3.10)$ across these six profiles. Thus,




FIG. 11. Fine-grid 2D axisymmetric eigenmodes used in Appendix D. Each row corresponds to one benchmark separation ($D = 6, 12, 18$); columns show the lowest even/odd parity modes (ψ_1, ψ_2) entering the localized-channel extraction pipeline.

at current scan resolution, profile-interpolation uncertainty is subdominant for $\mathcal{R}_3$ and controlled for the $H \rightarrow \mu\mu$ acceptance metric.

As a stronger stress test, we further rescaled the localized profile by large factors and tracked boundary movement directly on the (D, η) grid. Figure 12 shows that, at the present 60×60 scan resolution, the $\mathcal{R}_3 > 0.90$ mask is stable up to $\times 10^4$ and first shifts at $\times 10^5$, while the $H \rightarrow \mu\mu$ acceptance mask already shifts for the user-requested $\times 0.5$ and $\times 2$ cases (changed-cell fractions 0.0331 and 0.0258, respectively).

TABLE XIV. Extended profile-sensitivity test for $\chi_N^{(LR)}(D)$ in the 60×60 global scan. Source data file: `paper/chi_profile_robustness.csv`.

Profile	$f(\mathcal{R}_3 > 0.90)$	$f(\mathcal{R}_3 > 0.95)$	$f(\chi_{\mu\mu}^2 < 4)$	Best (D, η)
Linear interpolation	0.927	0.217	0.168	(9.42, 1.29)
Exponential fit	0.927	0.217	0.163	(9.15, 1.94)
Linear $\times 0.8$	0.927	0.217	0.167	(8.61, 3.03)
Linear $\times 1.2$	0.927	0.217	0.168	(8.07, 2.26)
Exp-fit $\times 0.8$	0.927	0.217	0.163	(8.61, 3.10)
Exp-fit $\times 1.2$	0.927	0.217	0.163	(8.88, 2.13)




FIG. 12. Strong χ -amplitude stress test on the global map. Left: overlay of $\chi_{\mu\mu}^2 \leq 4$ boundary for selected scales, including the requested $\times 0.5$ and $\times 2$. Right: changed-cell fractions versus scale (relative to baseline $\times 1$), showing early movement for the $H \rightarrow \mu\mu$ mask and delayed movement for the $\mathcal{R}_3$ mask. Source data file: `output/chi_fp_2d/chi_scale_stress_test.csv`.

Appendix E: Preliminary First-Principles Replacement Candidate Checks

1. 1D Phase-Space Candidate for g_N

To test a first-principles replacement direction for Eq. (18), we evaluated the semiclassical 1D candidate on the same action-derived axis potential:

$$\rho_{\text{WKB}}(E; D) = \frac{1}{2\pi} \int_{U(z; D) < E} \frac{dz}{\sqrt{E - U(z; D)}}, \quad (\text{E1})$$

$$g_N^{(\text{ps})}(D) = 1 + \int_{E_{\min}(D)}^{E_N(D)} \rho_{\text{WKB}}(E'; D) dE', \quad E_{\min}(D) \equiv \min_z U(z; D). \quad (\text{E2})$$

Using the reproducible script `extract_gn_phase_space_candidate.py` (in the `code/` directory), we ran a coarse/mid/fine test at $D = 12$ (Table XV). The candidate is numerically stable at the few- 10^{-3} level versus fine grid.

TABLE XV. Preliminary 1D phase-space candidate $g_N^{(\text{ps})}$ at $D = 12$.

level	dz	E_1	E_2	E_3	$g_1^{(\text{ps})}$	$g_2^{(\text{ps})}$	max rel. g vs fine
coarse	0.04	-0.636696	-0.608074	-0.575010	1.056725	1.054421	1.31×10^{-2}
mid	0.02	-0.637061	-0.608419	-0.575335	1.066813	1.074099	5.30×10^{-3}
fine	0.01	-0.637233	-0.608582	-0.575488	1.063919	1.068433	0

2. 2D Axisymmetric Benchmark for $g_N^{(\text{ps})}$

To reduce projection ambiguity, we also ran a 2D axisymmetric benchmark on $D = \{6, 12, 18\}$ using the same localized extraction chain as Appendix D: identical generalized operator build, identical box/range convention, and the same shift-invert low-mode settings ($\sigma = 2.5$, $\text{tol} = 10^{-8}$, $\text{maxiter} = 3 \times 10^4$). We define

$$N_{\text{ps}}(E) = \frac{1}{4\pi} \int d^2x d^2p \Theta(E - U - p^2) = \frac{1}{2} \int \rho d\rho dz [E - U(\rho, z)]_+, \quad (\text{E3})$$

$$\mathcal{W}_N \equiv [\lambda_1, \lambda_N], \quad g_{N,\text{raw}}^{(\text{ps})} = 1 + \int_{\mathcal{W}_N} \rho_{\text{ps}}(E') dE' = 1 + N_{\text{ps}}(\lambda_N) - N_{\text{ps}}(\lambda_1), \quad \hat{g}_N \equiv \frac{g_{N,\text{raw}}^{(\text{ps})}}{g_{3,\text{raw}}^{(\text{ps})}}. \quad (\text{E4})$$

The two-dimensional phase-space extraction exports both the low- N profile table and the low-mode spectral table with explicit window bounds. For non-fine levels, the maximum profile deviation is


$$\max_{\text{non-fine}} \left(\max_N \frac{|\hat{g}_N - \hat{g}_N^{(\text{fine})}|}{|\hat{g}_N^{(\text{fine})}|} \right) = 2.35 \times 10^{-2}. \quad (\text{E5})$$

As in the 1D candidate check, this benchmark is reported as a replacement candidate. To test map-level impact, we connected the baseline, legacy Cardy comparator, one-dimensional phase-space candidate, and two-dimensional phase-space candidate to the main scan through the reproducible script `scan_gn_profile_impact.py`. Results in Table XVII show that the main conclusion is retained: relative to the bounded two-dimensional baseline, the largest absolute shift is 0.0669 in $f(\mathcal{R}_3 > 0.90)$ (for the raw two-dimensional phase-space candidate), while high- N runaway remains controlled with $f(N_{\text{win}} > 3) \simeq 2.78 \times 10^{-4}$ in all tested modes.

Across $N_{\text{max}} = 20, 30, 40$, both baseline and fp_2d scenarios are numerically stable at map level: the key fractions $f(\mathcal{R}_3 > 0.90)$, $f(\chi_{\mu\mu}^2 < 4)$, and $f(N_{\text{win}} > 3)$ remain unchanged in this scan.

TABLE XVI. Fine-grid 2D phase-space profile candidate at $D = \{6, 12, 18\}$.

D	λ_1	λ_2	λ_3	$\hat{g}_1$	$\hat{g}_2$
6	0.719476	0.804126	0.945018	9.867×10^{-2}	4.370×10^{-1}
12	0.686186	0.747510	0.822830	1.193×10^{-1}	5.146×10^{-1}
18	0.667848	0.713807	0.767832	1.291×10^{-1}	5.294×10^{-1}

Low- N comparison: legacy Cardy reference vs 2D phase-space profileFIG. 13. Low- N comparison between legacy Cardy reference and 2D phase-space profile on $D = \{6, 12, 18\}$, shown as normalized ratios g_N/g_3 (log scale). The highlighted band marks $N = 1, 2, 3$.Source files: `paper/gn_cardy_vs_phase_space.png`, `paper/gn_cardy_vs_phase_space.csv`.

3. Baseline Cross-Check Under $g_{\text{mode}} = \text{fp_2d_full}$

To implement a strict same-framework upgrade, we set the bounded two-dimensional phase-space profile as the baseline. Instead of direct low- N substitution plus geometric continuation, this mode uses an explicit bounded microcanonical window for $N \leq 3$ and an energy-gap tail prescription for $N > 3$:

$$g_N^{(\text{full})}(D) = g_3^{(\text{ps})}(D) \hat{g}_N(D), \quad \hat{g}_n = \left(\hat{g}_n^{(\text{dir})} \right)^{1-\alpha} \left(\hat{g}_n^{(\text{win})} \right)^\alpha \quad (n = 1, 2, 3), \quad (\text{E6})$$

with

$$\hat{g}_1^{(\text{win})} = 1 + N_{\text{ps}}(\lambda_3) - N_{\text{ps}}(\lambda_1), \quad \hat{g}_2^{(\text{win})} = 1 + N_{\text{ps}}(\lambda_3) - N_{\text{ps}}(\lambda_2), \quad \hat{g}_3^{(\text{win})} = 1, \quad (\text{E7})$$

with $g_3^{(\text{ps})}(D) \equiv 1 + N_{\text{ps}}(\lambda_3) - N_{\text{ps}}(\lambda_1)$ from the same 2D spectrum table. For $N > 3$,

$$\hat{g}_N^{(\text{target})} = \left(\frac{\Delta N_{\text{ps}}^{(N)}}{\Delta N_{\text{ps}}^{(3)}} \right)^s \exp \left[-\beta \frac{\lambda_N - \lambda_3}{\lambda_3 - \lambda_2} \right], \quad (\text{E8})$$

followed by the same per-step finite-volume clipping used in `pslt_lib.py`. In the present baseline run, we use $(\alpha, \beta, s) = (0.8, 1.1, 0.0)$. Results are summarized in Table XIX (source CSV:

TABLE XVII. Map-level impact of first-principles g_N profiles on the 60×60 scan (baseline is `fp_2d_full`; legacy Cardy retained as comparator). Source CSV: `paper/gn_profile_impact.csv`.

case	$f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	$f(N_{\text{win}} > 3)$	mean tail prob.
baseline <code>fp_2d_full</code>	0.927	0.1500	0.0003	0.1267
legacy Cardy reference	0.877	0.1486	0.0003	0.1125
g_N fp-1d profile	0.860	0.1486	0.0003	0.1137
g_N fp-2d profile	0.860	0.1492	0.0003	0.1144

TABLE XVIII. N_{max} convergence check for baseline and g_N fp-2d map-level summaries. Source CSV: `paper/gn_nmax_convergence.csv`.

case	N_{max}	$f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	$f(N_{\text{win}} > 3)$	mean tail prob.
baseline <code>fp_2d_full</code>	20	0.927	0.1500	0.0003	0.126709
baseline <code>fp_2d_full</code>	30	0.927	0.1500	0.0003	0.126709
baseline <code>fp_2d_full</code>	40	0.927	0.1500	0.0003	0.126709
g_N fp-2d profile	20	0.860	0.1492	0.0003	0.114352
g_N fp-2d profile	30	0.860	0.1492	0.0003	0.114352
g_N fp-2d profile	40	0.860	0.1492	0.0003	0.114352

`paper/gn_baseline_replacement.csv`). Relative to baseline `fp_2d_full`, the legacy Cardy reference lowers $f(\mathcal{R}_3 > 0.90)$ by 0.0500 (from 0.9269 to 0.8769), lowers $f(\chi_{\mu\mu}^2 < 4)$ by 0.0014 (from 0.1500 to 0.1486), and keeps $f(N_{\text{win}} > 3) \simeq 2.78 \times 10^{-4}$. Therefore, the map-level main conclusion is retained under this baseline switch.

Across $N_{\text{max}} = 20, 30, 40$, both rows in Table XIX are numerically stable in this setup (source CSV: `paper/gn_baseline_replacement_nmax.csv`).

4. Geometry-Informed and Microscopic Lindblad Scan-Ready Module for Open-System

χ

As a methodological bridge (not a replacement of Appendix D), we consider a two-level open-system model in the standard GKSL/Lindblad form [10, 11]:

$$\dot{\rho} = -i[H, \rho] + \gamma_{\phi} \mathcal{D}[\sigma_z] \rho + \gamma_{\text{mix}} \mathcal{D}[\sigma_x] \rho, \quad H = \begin{pmatrix} 0 & \Delta/2 \\ \Delta/2 & 0 \end{pmatrix}. \quad (\text{E9})$$

As external motivation for this diagnostic direction, recent entanglement-based flavor analyses report that minimizing scattering-generated entanglement can qualitatively reproduce observed CKM/PMNS patterns [12]. In this manuscript we use that connection only as qualitative support for coherence-sensitive diagnostics, not as a baseline closure assumption. For this module

TABLE XIX. Map-level baseline cross-check for g_N : `fp_2d_full` baseline versus legacy Cardy reference on the same 60×60 grid. Source CSV: `paper/gn_baseline_replacement.csv`.

case	$f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	$f(N_{\text{win}} > 3)$	$\Delta f(\mathcal{R}_3 > 0.90)$ vs baseline
baseline <code>fp_2d_full</code>	0.9269	0.1500	0.0003	0.0000
legacy Cardy reference	0.8769	0.1486	0.0003	-0.0500

we set $\gamma_{\text{mix}}(D) \equiv M_{LR}^{(H)}(D)$ from Appendix D, and define a geometry-informed dephasing proxy

$$\gamma_\phi(D) \equiv \sqrt{\langle (\delta V)^2 \rangle_\rho}, \quad \delta V = V_{\text{eff}} - \bar{V}_{\text{eff}}(r), \quad (\text{E10})$$

with cylindrical weighting $\langle \cdots \rangle_\rho$. The minimal Lindblad evolution and geometry-profile extraction output reproducible diagnostics such as

$$C_{\text{max}} \equiv \max_t |\rho_{LR}(t)|, \quad P_{\text{max}} \equiv \max_t \rho_{RR}(t), \quad (\text{E11})$$

and a normalized proxy $\chi_{\text{eff}}^{(\text{proxy})} = 2\gamma_{\text{mix}} C_{\text{max}} / \bar{\Gamma}_N$. Results for $D = \{6, 12, 18\}$ are summarized in Table XX. Extending the same profile to $D = 4, \dots, 20$ gives a stable geometry-informed ratio band $\chi_{\text{open}}/\chi_{LR} \in [0.0852, 0.2127]$ with mean 0.1269. This value belongs to the geometry-only proxy closure in this subsection; the anchor-calibrated microscopic band is reported separately in Eq. (E13) and Table XXI.

To move this module toward microscopic closure, we also extract localized-basis couplings (g_z, g_x) directly from the same 2D eigenmodes and set

$$\gamma_\phi = \kappa_{\text{env}} g_z^2 S(0), \quad \gamma_{\text{mix}} = \kappa_{\text{env}} g_x^2 S(\Delta E), \quad S(\omega) = \frac{2\tau_{\text{env}}}{1 + (\omega\tau_{\text{env}})^2}, \quad (\text{E12})$$

from the localized microscopic extraction. To avoid map-level fitting, we calibrate κ_{env} with a multi-anchor rule: least-squares match to geometry-profile ratios on anchors $D = \{6, 9, 12, 15, 18\}$, then validate on holdout $D = \{4, 5, 7, 8, 10, 11, 13, 14, 16, 17, 19, 20\}$. This gives $\kappa_{\text{env}} = 4.31 \times 10^5$ with anchor RMSE 8.61×10^{-2} and holdout RMSE 6.95×10^{-2} (holdout max-abs error 1.25×10^{-1}). The corresponding scan-level band is

$$\chi_{\text{open}}^{(\text{micro})}/\chi_{LR} \in [0.0173, 0.3271], \quad \langle \chi_{\text{open}}^{(\text{micro})}/\chi_{LR} \rangle = 0.1380, \quad (\text{E13})$$

from the microscopic ratio-band audit. To pin down the microscopic bookkeeping, we additionally export a bridge audit with map, summary, and figure artifacts. On the calibrated $D = 4, \dots, 20$ knot set, this audit reconstructs

$$\gamma_\phi = \kappa_{\text{env}} g_z^2 S(0), \quad \gamma_{\text{mix}} = \kappa_{\text{env}} g_x^2 S(\Delta E), \quad \chi_{\text{eff}}^{(\text{micro})} = 2\gamma_{\text{mix}} C_{\text{max}} / \Gamma_{\text{ref}}$$

with numerically negligible residuals:

$$\max |\Delta \gamma_\phi| = 2.03 \times 10^{-20}, \quad \max |\Delta \gamma_{\text{mix}}| = 4.40 \times 10^{-14}, \quad \max |\Delta \chi_{\text{eff}}^{(\text{loader})}| = 5.58 \times 10^{-17}.$$

The corresponding bridge summary identifies the effective system as the localized two-level (L/R) subspace, the current bath witness as a Gaussian Markov bath with Lorentzian PSD, and the coupling operators as the diagonal/off-diagonal pair (σ_z, σ_x) . This means the open-system gap is no longer in the system-to-Lindblad bookkeeping, but in the microscopic EYMH origin and normalization of the bath itself.

We can therefore push the bath normalization one step further without over-claiming a full microscopic derivation. In the κ_{env} window audit, we treat

$$\kappa_{\text{env}} \rightarrow \kappa_{\text{scale}} \kappa_{\text{env}}^{(\text{calib})}$$

as a uniform bath-normalization susceptibility that rescales both γ_ϕ and γ_{mix} while leaving the localized two-level Hamiltonian $(\Delta E, \sigma_x)$ fixed. On the current $D = 4, \dots, 20$ knot set, the calibration-consistent candidate window is

$$\kappa_{\text{scale}} \in [0.5, 1.0],$$

while a broader scan-stable window

$$\kappa_{\text{scale}} \in [0.25, 1.5]$$

keeps the map-level fractions $f(\mathcal{R}_3 > 0.90)$ and $f(\chi_{\mu\mu}^2 \leq 4)$ unchanged and only relaxes the holdout tolerance slightly. In this sense, κ_{env} is no longer just a single-point fit parameter: within the present witness it acts as a bounded bath-normalization susceptibility. What still remains open is the EYMH-side derivation of its absolute normalization and of the microscopic bath degrees of freedom themselves.

The newer parent-bath audits sharpen this statement. First, they show that the witness already factorizes into a system block (g_z^2, g_x^2) , a bath-shape block $(S_{zz}(0), S_{xx}(\Delta E))$, and an amplitude block κ_{env} , with the canonical constant normalization remaining the unique exact amplitude choice. Second, they rewrite the rates as the projected bath block

$$K_{\text{bath}} = \kappa_{\text{env}} \sqrt{K_{\text{sys}}} K_{\text{spec}} \sqrt{K_{\text{sys}}},$$

and the subsequent family, log-coordinate, normal-coordinate, and generator-affinity audits show that the canonical projected bath class is uniquely selected on the current knot set: $(m, u, v) = (0, 0, 0)$, $(p_{\text{sys}}, p_{\text{spec}}) = (0, 0)$, $(\zeta_{\text{sys}}, \zeta_{\text{spec}}) = (0, 0)$, and $(q_{ss}, q_{bb}, q_{sb}) = (0, 0, 0)$. It is therefore natural to define the normalized projected bath block

$$B_a(z) := \kappa_{\text{env}}^{-1} K_{\text{sys},a}^{-1/2}(z) K_{\text{bath},a}(z) K_{\text{sys},a}^{-1/2}(z),$$

which the audited factorization already identifies as the bath-shape witness itself. Thus the remaining open-system gap is no longer bookkeeping or nearby-family ambiguity; it is the final

parent-action statement for why the projected bath generator itself naturally lives in this affine log class. Finally, `scan_chi_open_system_parent_bath_cocycle_audit.py` turns this into a positive integrability statement: after dividing out κ_{env} , the normalized parent bath block defines an exact additive cocycle in the canonical log variables, with pairwise cocycle residuals at 8.95×10^{-16} on the identifiable ϕ subset and 8.64×10^{-13} on the mixing branch, while triangle flatness defects stay at 1.11×10^{-15} in both channels. So the affine log-generator is now supported not only by nearby-family exclusion, but also by exact cocycle/flatness closure. Pushing one step further, `scan_chi_open_system_parent_bath_potential_audit.py` shows that this cocycle is globally integrable on the current knot set: the normalized bath block admits an exact anchored scalar potential $\Phi_a = \log(B_a/B_a^{\text{ref}})$ whose direct decomposition, anchor-average recovery, and nearest-neighbor chain recovery agree to machine precision (max residuals $\sim 10^{-14}$ on the identifiable ϕ subset and $\sim 5 \times 10^{-13}$ or better on the mixing branch). Thus the projected bath witness is no longer characterized only by an affine local generator; it also carries a single-valued anchored potential, so the remaining gap is specifically the parent-action origin of this potential itself.

Figure 14 gives the representative late-stage open-system diagnostic: the normalized parent-bath block has already been reduced to an integrable projected potential rather than an unconstrained reservoir fit.

At this point the cocycle/potential relation can be written explicitly on the audited knot set. Defining the discrete increment

$$G_a(z_i, z_j) := \log B_a(z_j) - \log B_a(z_i),$$

the cocycle audit shows that G_a is additively consistent:

$$G_a(z_i, z_k) = G_a(z_i, z_j) + G_a(z_j, z_k),$$

up to the machine-level residuals quoted above. Fixing an anchor z_* , the potential is then not an extra ansatz but the integrated cocycle itself,

$$\Phi_a(z) := G_a(z_*, z), \quad G_a(z_i, z_j) = \Phi_a(z_j) - \Phi_a(z_i).$$

So the anchored potential is simply the scalar primitive of the already-audited normalized bath block in canonical log variables.

This also gives the natural local parent-bath effective-action normal form on the audited knot set. Writing $z = (z_{\text{sys}}, z_{\text{spec}})$ for the canonical log variables and $B_a(z) = B_a^{\text{ref}} e^{\Phi_a(z)}$, the cocycle closure identifies the bath witness as an anchored exponential potential on the projected




FIG. 14. Projected parent-bath potential closure for the open-system module on the calibrated knot set. The normalized bath block admits a single-valued scalar primitive Φ_a whose direct reconstruction, anchor recovery, and chain recovery agree to machine precision: the maximum residuals are 1.42×10^{-14} on the identifiable ϕ subset and 5.50×10^{-13} on the mixing branch. Source files: `paper/chi_open_system_parent_bath_potential.*`.

bath block,

$$K_{\text{bath},a}(z) = \kappa_{\text{env}} K_{\text{sys},a}^{1/2}(z) e^{\Phi_a(z)} K_{\text{sys},a}^{1/2}(z),$$

so, channelwise on the present positive knot set,

$$B_a(z) = e^{\Phi_a(z)}, \quad \log K_{\text{bath},a}(z) = \log \kappa_{\text{env}} + \log K_{\text{sys},a}(z) + \Phi_a(z).$$

while the family, log-coordinate, normal-coordinate, and generator-affinity audits isolate the canonical class as the one with no surviving first local log-curvature jet. Equivalently, on a sufficiently small neighborhood $\mathcal{N}$ of the canonical anchor,

$$\Phi_a(z + \tau) = \Phi_a(z) + \nabla \Phi_a(z) \cdot \tau + R_{a,2}(z, \tau), \quad |R_{a,2}(z, \tau)| \leq \frac{1}{2} \sup_{\zeta \in \mathcal{N}} \|D^2 \Phi_a(\zeta)\| \|\tau\|^2,$$

so the affine log-generator is just the first Taylor jet of Φ_a on the projected bath block, while the first non-affine correction is controlled by $R_{a,2}$. The audited canonical point is precisely

the one where this first non-affine correction is rejected on the current knot set. In this sense, the open-system witness is already promoted from a free nearby-family fit to a locally anchored projected effective action; what remains open is not the existence or local normal form of Φ_a , but the microscopic EYMH derivation of its coefficients, of κ_{env} , and of the bath degrees of freedom themselves.

The implication chain here is non-circular. The primitive audited object is the normalized projected bath block itself, namely $K_{\text{bath},a}/\kappa_{\text{env}}$ on the canonical knot set. The nearby-family, log-coordinate, normal-coordinate, and generator-affinity audits act directly on that block and select the unique local affine generator class without assuming any potential. The cocycle audit is then applied to the same normalized block, upgrading this local affine closure to an exact additive consistency statement in the canonical log variables. The potential audit comes only after that step: it integrates the already-audited cocycle and checks path-independence / anchor recovery, rather than feeding any potential ansatz back into the family selection. Equation (E 4) together with the anchored form above therefore repackages the same projected bath witness into a local effective-action normal form; it does not create new evidence by definition. The remaining open-system gap is still the microscopic parent-action origin of the coefficients and bath degrees of freedom, not the internal consistency of the projected bath block on the present audited knot set.

This also sharpens the remaining microscopic origin question into three narrower obligations. First, since

$$B_a(z) = \kappa_{\text{env}}^{-1} K_{\text{sys},a}^{-1/2}(z) K_{\text{bath},a}(z) K_{\text{sys},a}^{-1/2}(z),$$

one has

$$\partial_{\log \kappa_{\text{env}}} \log B_a(z) = 0, \quad \partial_{\log \kappa_{\text{env}}} \log K_{\text{bath},a}(z) = 1,$$

so κ_{env} already behaves as a pure scalar bath-susceptibility modulus rather than as an independent bath-shape degree of freedom. Second, the coefficients of Φ_a are no longer free local fit parameters on the audited knot set: the cocycle primitive $G_a(z_i, z_j) = \Phi_a(z_j) - \Phi_a(z_i)$ already fixes their discrete local jet, and the remaining task is only to identify the parent-action source of those coefficients. Third, the unresolved bath degrees of freedom have already been reduced to the two effective response channels appearing in the audited micro bridge,

$$\gamma_{\phi,a} = \kappa_{\text{env}} g_{z,a}^2 S_{zz,a}(0), \quad \gamma_{\text{mix},a} = \kappa_{\text{env}} g_{x,a}^2 S_{xx,a}(\Delta E_a),$$

rather than to an arbitrary reservoir family. Thus the next microscopic step, if taken at all, is not to re-justify the local projected effective-action form, but to explain why the parent EYMH

dynamics projects specifically onto this scalar modulus, this anchored potential jet, and this minimal two-channel bath kernel.

One can also sharpen the second obligation without yet claiming a full parent-action derivation of the coefficients. Restrict to the identifiable audited subset

$$\mathcal{K}_\phi := \{D : \phi_a(D) \text{ is identifiable}\},$$

and let

$$G_a(D_i, D_j) := \Phi_a(D_j) - \Phi_a(D_i)$$

denote the already-audited cocycle primitive on that subset. For any local anchor $D_\star \in \mathcal{K}_\phi$ and any finite neighborhood $\mathcal{N}(D_\star) \subset \mathcal{K}_\phi$, the discrete first jet coefficient is fixed by the primitive itself, not by a new fit,

$$c_a(D_\star) := \arg \min_c \sum_{D_j \in \mathcal{N}(D_\star)} w_j \left[\Phi_a(D_j) - \Phi_a(D_\star) - c(D_j - D_\star) \right]^2 = \frac{\sum_j w_j (D_j - D_\star) G_a(D_\star, D_j)}{\sum_j w_j (D_j - D_\star)^2},$$

with local remainder

$$r_{a,\star}(D_j) := \Phi_a(D_j) - \Phi_a(D_\star) - c_a(D_\star)(D_j - D_\star), \quad \sum_j w_j (D_j - D_\star) r_{a,\star}(D_j) = 0.$$

So the local coefficient of Φ_a on the audited knot set is already an anchor-locked discrete jet of the cocycle primitive, rather than an unconstrained nearby-family parameter. The current audits then push one step further still: on the entire exported map, the observed parent-bath action coefficient is exactly the signed carrier of the same primitive,

$$A_a(D) := (\text{parent-bath action}_\phi)_a(D) = -\Phi_a(D), \quad c_a^{(A)}(D_\star) = -c_a^{(\Phi)}(D_\star).$$

Numerically, the potential reconstruction residual stays below 1.5×10^{-14} on the identifiable subset, the anchor and chain recoveries stay below 7.2×10^{-15} and 2.2×10^{-14} respectively, the audited cocycle / local-generator defect stays below 7.2×10^{-15} , and the generator/potential representatives agree pointwise at the same 2.2×10^{-14} level. Moreover, the canonical affine class is not merely acceptable but strongly selected: the generator-affinity and normal-coordinate selection gaps exceed the canonical objective by factors of 2.2×10^8 and 4.4×10^8 respectively. Thus the unresolved microscopic issue is no longer whether Φ_a admits a stable local first jet on the audited knot set; it is why the EYMH parent action generates precisely this signed discrete jet and not some nearby nonlinear deformation.

Within that first obligation, the present audits already push one step further than a generic “fitted constant” interpretation. If one allows a nearby absolute-normalization family

$$\kappa_{\text{eff}}(z) = \kappa_{\text{env}} s F(z),$$

with a single anchor-fitted amplitude s but a non-uniform profile $F(z)$, then the normalized projected bath block becomes

$$B_a^{(F)}(z) = \kappa_{\text{eff}}^{-1}(z) K_{\text{sys},a}^{-1/2}(z) K_{\text{bath},a}(z) K_{\text{sys},a}^{-1/2}(z) = (sF(z))^{-1} B_a(z).$$

So an exact amplitude-only renormalization preserves the normalized block if and only if $F(z)$ is constant on the audited knot set. This is exactly what the absolute-normalization audit finds: the canonical locked uniform family keeps the bridge-normalization residuals at machine level ($\sim 8 \times 10^{-13}$ on the observable/mixing channels), whereas the best non-uniform admissible family improves the holdout RMSE only marginally (6×10^{-5}) while inflating the same normalized residuals to $\sim 10^{-2}$. In that precise local sense, κ_{env} is already singled out as the unique exact amplitude modulus, and the remaining microscopic question is no longer whether it mixes with bath shape on the present knot set, but why the parent EYMH dynamics produces this modulus in the first place.

Proposition O1 (Projected two-channel interaction). Without invoking a full bath derivation, the parent perturbation projected to the audited parity doublet already fixes the leading localized interaction to a two-channel Pauli block. Let ΔV denote the parent perturbation projected to the parity doublet, so that on the audited knot set

$$V_{\text{par},a} = \begin{pmatrix} v_{11,a} & v_{12,a} \\ v_{12,a} & v_{22,a} \end{pmatrix}$$

is real and symmetric. Rotating to the localized basis with

$$U_{LR} = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$$

gives

$$V_{\text{loc},a} = U_{LR}^T V_{\text{par},a} U_{LR} = \frac{v_{11,a} + v_{22,a}}{2} I + v_{12,a} \sigma_z + \frac{v_{11,a} - v_{22,a}}{2} \sigma_x.$$

Hence every audited localized two-level block already lies in the span of I, σ_z, σ_x , with no surviving σ_y component: after removing the trace part, the leading projected interaction is at most two-channel. The extracted couplings are therefore not independent fit coordinates but exactly the Pauli coefficients of the localized block,

$$g_{z,a} = |v_{12,a}|, \quad g_{x,a} = \left| \frac{v_{11,a} - v_{22,a}}{2} \right|.$$

This remains an effective two-level statement rather than a completed EYMH bath theorem, but it shows that the present σ_z/σ_x bridge is already the most general leading projection compatible

with the audited real localized doublet. Numerically, the full 17-point micro export verifies this decomposition at machine precision: the reconstructed localized entries (V_{LL}, V_{RR}, V_{LR}) agree with the exported map at $\leq 7 \times 10^{-18}$, while $g_z - |v_{12}|$ and $g_x - |(v_{11} - v_{22})/2|$ stay below 2.2×10^{-18} and 1.3×10^{-18} respectively. The remaining microscopic bath question is therefore narrower than a generic reservoir identification: one still has to explain why the parent EYMH environment induces the two on-shell kernels multiplying these projected σ_z and σ_x directions, not why those directions themselves appear in the audited two-level reduction.

Proposition O2 (Single gap-locked response family). Within the present minimal bridge, the dephasing and mixing kernels are not independent shape functions but on-shell evaluations of a single response family. The audited micro extraction uses

$$S_a(\omega) = \frac{2\tau_{\text{env},a}}{1 + (\omega\tau_{\text{env},a})^2},$$

so that

$$S_{zz,a}(0) = S_a(0) = 2\tau_{\text{env},a}, \quad S_{xx,a}(\Delta E_a) = S_a(\Delta E_a) = \frac{2\tau_{\text{env},a}}{1 + (\Delta E_a\tau_{\text{env},a})^2}.$$

Equivalently,

$$\frac{S_{xx,a}(\Delta E_a)}{S_{zz,a}(0)} = \frac{1}{1 + (\Delta E_a\tau_{\text{env},a})^2},$$

so the mixing channel is just the same minimal Lorentzian response sampled on shell at the transition energy, while the dephasing channel is the zero-frequency limit of that same family. On the current 17-point micro export this identification is exact at the CSV level: the reconstructed residuals for $S_{zz,a}(0) = 2\tau_{\text{env},a}$ and $S_{xx,a}(\Delta E_a) = 2\tau_{\text{env},a}/(1 + (\Delta E_a\tau_{\text{env},a})^2)$ vanish to machine precision, while the ratio relation above holds with maximal defect 1.1×10^{-16} . Thus the unresolved EYMH bath-origin problem is narrower still: beyond the already-audited two-channel projection, one only has to explain why the parent environment induces this single minimal response family and why its on-shell samples are taken at $\omega = 0$ and $\omega = \Delta E_a$.

One can push this “single response family” statement one notch further without yet claiming a full bath derivation. The same audited family is exactly the Fourier image of the minimal even single-pole correlation kernel

$$C_a(t) = e^{-|t|/\tau_{\text{env},a}},$$

for which

$$S_a(\omega) = \int_{-\infty}^{\infty} dt C_a(t) e^{i\omega t} = 2 \int_0^{\infty} e^{-t/\tau_{\text{env},a}} \cos(\omega t) dt = \frac{2\tau_{\text{env},a}}{1 + (\omega\tau_{\text{env},a})^2}.$$

Moreover, on the present audited knot set the correlation time is not a second free shape coordinate: the exported micro bridge uses

$$\tau_{\text{env},a} = \max\left(\frac{\tau_{\text{scale}}}{\omega_{1,a}}, \tau_{\text{floor}}\right),$$

and the floor never activates on the current 17-point map, so one has numerically

$$\tau_{\text{env},a} = \frac{1}{\omega_{1,a}}$$

to machine precision throughout the audited window. Hence the minimal response family is already gap-locked,

$$S_a(\omega) = \frac{2/\omega_{1,a}}{1 + (\omega/\omega_{1,a})^2}, \quad \frac{S_{xx,a}(\Delta E_a)}{S_{zz,a}(0)} = \frac{1}{1 + (\Delta E_a/\omega_{1,a})^2}.$$

On the current export the maximal defect in $\tau_{\text{env},a} - \omega_{1,a}^{-1}$ is 1.1×10^{-16} , the on-shell suppression law above again holds to 3.3×10^{-16} , and the resulting band $S_{xx,a}(\Delta E_a)/S_{zz,a}(0)$ stays in the narrow interval $[0.9246, 0.9800]$. In this precise effective sense, the remaining EYMH bath-origin question is no longer why two unrelated kernels appear, but why the parent environment linearizes to a single dominant pole whose correlation time is locked to the same local gap scale $\omega_{1,a}^{-1}$ that enters the audited two-level reduction.

Proposition O3 (Schur-resolvent reformulation). The dominant-pole question can be stated directly in the same parent-kernel language used in the EYMH subsection. Let

$$R_a(\omega) := u_a^\top (M_a^2 + \omega^2 I)^{-1} u_a$$

denote the projected Schur-resolvent response of the parent bath/complement block in the same localized channel. Its spectral decomposition reads

$$R_a(\omega) = \sum_{n \geq 1} \frac{Z_{n,a}}{\Omega_{n,a}^2 + \omega^2}, \quad Z_{n,a} := |\langle u_a, n, a \rangle|^2 \geq 0.$$

If the first mode dominates on the audited on-shell window,

$$R_a(\omega) \approx \frac{Z_{1,a}}{\Omega_{1,a}^2 + \omega^2},$$

then, after absorbing the residue $Z_{1,a}$ into the overall amplitude block κ_{env} and the projected couplings $g_{z,a}, g_{x,a}$, one recovers exactly the same single-pole family used above. Identifying moreover $\Omega_{1,a} = \omega_{1,a}$ yields the observed gap-locked form

$$S_a(\omega) = \frac{2/\omega_{1,a}}{1 + (\omega/\omega_{1,a})^2}.$$

The present audits do not yet prove this first-mode dominance theorem, but they isolate it as the next genuinely microscopic target rather than leaving an undifferentiated reservoir ansatz.

On the EYMH parent side, the log-det / Schur witness and the parent block-determinant / Schur audit already close to machine precision, with $\max |J_{\text{direct}} - J_{\text{parent}}| = 1.30 \times 10^{-18}$, $\max |S_{\text{parent}} - S_{\text{resp}}| = 5.55 \times 10^{-16}$, and $\max |\Delta \hat{G}_{\text{Schur}}| = 6.66 \times 10^{-16}$. On the open-system side, the same audited knot set already satisfies $\tau_{\text{env},a} = \omega_{1,a}^{-1}$ and the on-shell rate identities to machine precision. In this narrowed sense, the unresolved bath-origin problem is no longer why some effective Lorentzian fit works, but why the projected Schur resolvent is first-mode dominant and why its leading pole is locked to the same local gap scale $\omega_{1,a}$.

Theorem Target O1 (First-mode remainder bound). The remaining microscopic step is an explicit remainder estimate rather than a qualitative bath ansatz. Decompose

$$R_a(\omega) = R_a^{(1)}(\omega) + \mathcal{R}_a(\omega), \quad R_a^{(1)}(\omega) := \frac{Z_{1,a}}{\Omega_{1,a}^2 + \omega^2},$$

so that

$$\mathcal{R}_a(\omega) = \sum_{n \geq 2} \frac{Z_{n,a}}{\Omega_{n,a}^2 + \omega^2}.$$

Then the exact relative remainder satisfies

$$\frac{|\mathcal{R}_a(\omega)|}{R_a^{(1)}(\omega)} \leq \sum_{n \geq 2} \frac{Z_{n,a}}{Z_{1,a}} \frac{\Omega_{1,a}^2 + \omega^2}{\Omega_{n,a}^2 + \omega^2}.$$

Hence, if one can show on the audited on-shell window $\omega \in \{0, \Delta E_a\}$ that

$$\varepsilon_a := \sup_{\omega \in \{0, \Delta E_a\}} \sum_{n \geq 2} \frac{Z_{n,a}}{Z_{1,a}} \frac{\Omega_{1,a}^2 + \omega^2}{\Omega_{n,a}^2 + \omega^2} \ll 1,$$

then

$$R_a(\omega) = R_a^{(1)}(\omega) [1 + O(\varepsilon_a)]$$

throughout the same on-shell window. This is the cleanest proof-level formulation of the dominant-pole claim: it is enough to control the complement remainder in the projected Schur resolvent, rather than to derive a full microscopic reservoir from scratch. The current export already shows that this on-shell window is numerically narrow, with $\Delta E_a / \omega_{1,a} \in [0.1430, 0.2857]$ and therefore $S_{xx,a}(\Delta E_a) / S_{zz,a}(0) \in [0.9246, 0.9800]$, so the audited dephasing and mixing channels sample only the low-frequency shoulder of the same pole. What is still missing is exactly the parent-side bound that identifies $\Omega_{1,a} = \omega_{1,a}$ and proves that the higher-mode remainder is uniformly negligible on this window.

A useful quantitative sharpening is immediate once one separates the gap question from the tail-weight question. Define

$$q_a := \frac{R_a(\Delta E_a)}{R_a(0)}, \quad r_a := \frac{\Delta E_a}{\Omega_{1,a}}, \quad \nu_{n,a} := \frac{Z_{n,a} / \Omega_{n,a}^2}{R_a(0)},$$

so that $\sum_{n \geq 1} \nu_{n,a} = 1$ and

$$q_a = \sum_{n \geq 1} \frac{\nu_{n,a}}{1 + (\Delta E_a / \Omega_{n,a})^2}.$$

If the higher modes satisfy the conditional gap bound $\Omega_{n,a} \geq g_a \Omega_{1,a}$ for all $n \geq 2$ with some $g_a > 1$, then with $f(x) := (1 + x^2)^{-1}$ one has

$$q_a \geq \nu_{1,a} f(r_a) + (1 - \nu_{1,a}) f\left(\frac{r_a}{g_a}\right),$$

hence

$$1 - \nu_{1,a} \leq \frac{q_a - f(r_a)}{f(r_a/g_a) - f(r_a)}.$$

Since $\varepsilon_a(0) = \frac{1 - \nu_{1,a}}{\nu_{1,a}}$ and $\varepsilon_a(\Delta E_a) \leq \varepsilon_a(0) \frac{1 + r_a^2}{g_a^2 + r_a^2}$, the theorem target can be restated as a two-knob quantitative task: prove a parent-side lower bound on g_a and a small on-shell defect $\delta_{q,a} := q_a - f(r_a)$. On the audited knot set, the exported gap-conditioned audit shows that for the representative benchmark $g_a = 2$ one has

$$f(r_a/2) - f(r_a) \in [1.50 \times 10^{-2}, 5.55 \times 10^{-2}], \quad \frac{1 + r_a^2}{4 + r_a^2} \in [2.54 \times 10^{-1}, 2.65 \times 10^{-1}],$$

so even a future parent-side match at the modest level $\delta_{q,a} \leq 10^{-3}$ would already force

$$1 - \nu_{1,a} \leq 6.68 \times 10^{-2}, \quad \varepsilon_a(\Delta E_a) \leq 1.82 \times 10^{-2}$$

uniformly across the current audited window. The present effective export, by contrast, has $\max |\delta_{q,a}| = 3.33 \times 10^{-16}$ because the fitted response family is exactly single-pole on shell. This does not yet prove the parent-side gap itself, but it converts the remaining Schur-resolvent problem into a sharply quantified gap-plus-defect statement rather than a diffuse reservoir ansatz.

One can now push this one notch further on the parent side without yet claiming a full spectral theorem. Write the background-normalized canonical complement block as

$$\tilde{K}_a := \begin{pmatrix} K_{11,a} & C_a / \sqrt{K_{\text{bg},a}} \\ C_a / \sqrt{K_{\text{bg},a}} & G_a \end{pmatrix}, \quad G_a := K_{22,a} / K_{\text{bg},a},$$

with the audited geometric-mean identity

$$\frac{C_a^2}{K_{\text{bg},a}} = (K_{11,a} - 1)(G_a - 1).$$

Hence

$$\tilde{K}_a = I + v_a v_a^\top, \quad v_a = \begin{pmatrix} \sqrt{K_{11,a} - 1} \\ \sqrt{G_a - 1} \end{pmatrix},$$

so the two proxy eigenvalues are fixed exactly by

$$\lambda_{1,a}^{(\text{proxy})} = 1, \quad \lambda_{2,a}^{(\text{proxy})} = 1 + \|v_a\|^2 = K_{11,a} + G_a - 1.$$

Calibrating the identity mode to the already-audited gap-locked pole then gives the calculable parent-side proxies

$$\begin{aligned} \Omega_{1,a}^{(\text{proxy})} &:= \omega_{1,a}, \\ \Omega_{2,a}^{(\text{proxy})} &:= \omega_{1,a} \sqrt{K_{11,a} + G_a - 1}, \\ g_a^{(\text{proxy})} &:= \frac{\Omega_{2,a}^{(\text{proxy})}}{\Omega_{1,a}^{(\text{proxy})}} = \sqrt{K_{11,a} + G_a - 1}. \end{aligned}$$

The corresponding audit closes to machine precision:

$$\begin{aligned} \max |C_a^2/K_{\text{bg},a} - (K_{11,a} - 1)(G_a - 1)| &= 2.78 \times 10^{-16}, \\ \max |\lambda_{1,a}^{(\text{proxy})} - 1| &= \max |\lambda_{2,a}^{(\text{proxy})} - (K_{11,a} + G_a - 1)| = 4.44 \times 10^{-16}. \end{aligned}$$

On the nontrivial active rows of the D21xE21 fix grid this already yields

$$g_a^{(\text{proxy})} \in [1.3537, 1.5609], \quad \text{median } g_a^{(\text{proxy})} = 1.4692.$$

On the overlap with the currently audited open-system knot set, the induced separator band becomes

$$f(r_a/g_a^{(\text{proxy})}) - f(r_a) \in [1.43 \times 10^{-2}, 4.03 \times 10^{-2}],$$

so a future parent-side defect estimate as mild as $\delta_{q,a} \leq 10^{-3}$ would already imply

$$1 - \nu_{1,a} \leq 6.98 \times 10^{-2}, \quad \varepsilon_a(\Delta E_a) \leq 3.60 \times 10^{-2}.$$

This is still only a proxy, not a proof that the true Schur-resolvent spectrum has the same gap. But it does show that the canonical projected EYMH complement block already carries a concrete, calculable gap floor rather than leaving g_a completely unconstrained.

One can also turn this proxy into a local operator-family lower bound by combining it with the exact parent mismatch functional from Proposition U4. Let

$$\delta\theta := \begin{pmatrix} \alpha - 1 \\ \beta - 1 \end{pmatrix}, \quad Q_{\text{loc}}(\delta\theta, \lambda) := \frac{1}{2} \delta\theta^\top H \delta\theta + C_4 \lambda^4,$$

where the stationarity audit supplies the exact quadratic matrix H in the (α, β) sector and the exact quartic coefficient C_4 in the mixing direction. On the audited local box

$$|\alpha - 1| \leq 0.1, \quad |\beta - 1| \leq 0.1, \quad |\lambda| \leq 0.2,$$

the parent-kernel scan gives the empirical coercivity bound

$$J_{\text{parent}}(\alpha, \beta, \lambda) \geq c_{\text{loc}} Q_{\text{loc}}(\delta\theta, \lambda), \quad c_{\text{loc}} = 3.98 \times 10^{-1},$$

with $\lambda_{\min}(H) = 1.33 \times 10^{-3}$ and $C_4 = 1.497 \times 10^{-3}$. Hence any local operator-side candidate obeying $J_{\text{parent}} \leq J_*$ inside that box must satisfy

$$\|\delta\theta\| \leq \rho_* := \sqrt{\frac{2J_*}{c_{\text{loc}}\lambda_{\min}(H)}}.$$

For the same rowwise gap proxy

$$g_a(\alpha, \beta) = \sqrt{e^{\alpha S_{\text{part},a}} + e^{\beta S_{\text{schur},a}} - 1},$$

the corresponding audit gives a local Lipschitz bound $\|\nabla g_a\| \leq L_a^{(\text{box})}$, so

$$g_a(\alpha, \beta) \geq g_{0,a} - L_a^{(\text{box})}\rho_*, \quad g_{0,a} := \sqrt{K_{11,a} + G_a - 1}.$$

Taking $J_* = 10^{-6}$, the audited box has $L_{\max}^{(\text{box})} = 5.11 \times 10^{-1}$ and therefore yields the uniform local operator-family floor

$$g_a \geq 1.3375.$$

Direct enumeration of the admissible scan points with $J_{\text{parent}} \leq 10^{-6}$ is sharper still and gives

$$g_a \geq 1.3537$$

throughout the same local family. On the overlap with the current open-system knot set, the conservative analytic floor $g_a \geq 1.3375$ already implies

$$f(r_a/g_a) - f(r_a) \geq 1.19 \times 10^{-2},$$

so the same reference defect $\delta_{q,a} \leq 10^{-3}$ would force

$$1 - \nu_{1,a} \leq 8.37 \times 10^{-2}, \quad \varepsilon_a(\Delta E_a) \leq 5.28 \times 10^{-2}.$$

This is still not a proof of the exact complement spectrum, because it remains confined to the audited projected family. But it upgrades the previous proxy statement into a genuine local operator-side lower bound: within the canonical projected EYMH basin, a small exact parent mismatch already forbids the second-mode scale from collapsing back to the leading pole.

Corollary 1 also gives a minimal lifting lemma from this local family floor to the exact Schur-reduced operator. Let

$$\sigma_a := -\log \det(I - A_a),$$

with the same positive semidefinite complement contraction

$$A_a = H_{PP,a}^{-1/2} H_{PQ,a} H_{QQ,a}^{-1} H_{QP,a} H_{PP,a}^{-1/2}.$$

If the eigenvalues of A_a lie in $[0, 1]$, then the eigenvalues of $I - A_a$ also lie in $(0, 1]$, and therefore

$$I - A_a \succeq \det(I - A_a) I = e^{-\sigma_a} I.$$

Hence the exact Schur-reduced operator obeys the Loewner bound

$$H_{\text{eff},a} = H_{PP,a}^{1/2} (I - A_a) H_{PP,a}^{1/2} \succeq e^{-\sigma_a} H_{PP,a},$$

so if the projected family supplies any lower gap floor $g_a^{(\text{fam})}$, then the exact effective gap satisfies

$$g_a^{(\text{exact})} \geq e^{-\sigma_a/2} g_a^{(\text{fam})}.$$

Using the conservative audited family floor $g_a^{(\text{fam})} \geq 1.3375$ from the local operator-family bound above, the current overlap with the open-system knot set yields the conditional exact floors

$$g_a^{(\text{exact})} \geq 1.3368 \quad (\sigma_a \leq 10^{-3}), \quad g_a^{(\text{exact})} \geq 1.3308 \quad (\sigma_a \leq 10^{-2}).$$

The corresponding separator band remains

$$f(r_a/g_a^{(\text{exact})}) - f(r_a) \geq 1.19 \times 10^{-2} \quad (\sigma_a \leq 10^{-3}),$$

up to the displayed precision, and only weakens to

$$f(r_a/g_a^{(\text{exact})}) - f(r_a) \geq 1.18 \times 10^{-2} \quad (\sigma_a \leq 10^{-2}).$$

With the same reference defect $\delta_{q,a} \leq 10^{-3}$ this still implies

$$\varepsilon_a(\Delta E_a) \leq 5.29 \times 10^{-2} \quad (\sigma_a \leq 10^{-3}), \quad \varepsilon_a(\Delta E_a) \leq 5.41 \times 10^{-2} \quad (\sigma_a \leq 10^{-2}).$$

In this form the remaining theorem target becomes sharper still: it is enough to prove a small complement log-det loss (or, equivalently, a small $\|A_a\|$) on the audited localized window, because the already-established family gap floor then lifts automatically to the exact Schur-reduced operator.

There is also an auxiliary parent-side route to that loss parameter if one assumes that the complement contraction is itself first-channel dominant on the same local window. In the rank-one model

$$A_a \approx a_a u_a u_a^\top, \quad 0 \leq a_a < 1,$$

one has

$$\|A_a\| = a_a, \quad \sigma_a = -\log(1 - a_a).$$

The exact projected kernel family already carries the corresponding scalar contraction candidate, because

$$\det K_a(\alpha, \beta, \lambda) = K_{11,a} K_{22,a} (1 - \lambda^2 \xi_a(\alpha, \beta)), \quad \xi_a(\alpha, \beta) := \frac{(K_{11,a} - 1)(K_{22,a} - 1)}{K_{11,a} K_{22,a}},$$

so the minimal dominant-channel identification is

$$a_a^{(1)} := \lambda^2 \xi_a.$$

The local coercivity bound above already gives

$$|\lambda| \leq \left(\frac{J_*}{c_{\text{loc}} C_4} \right)^{1/4},$$

which yields $|\lambda| \leq 2.02 \times 10^{-1}$ at $J_* = 10^{-6}$. On the overlap with the current open-system knot set, the exact canonical factor satisfies

$$\max \xi_a = 1.3324 \times 10^{-1},$$

while the best raw off-diagonal envelope built directly from the parented EYMH witnesses,

$$\xi_a^{(\text{raw})} := 0.127472 \frac{hk_{\text{abs,offdiag}}}{\text{diag}} + 0.435064 \frac{\text{action}_{\text{abs,offdiag}}}{\text{diag}},$$

has overlap p95 residual 4.37×10^{-2} and therefore gives the conservative envelope

$$\xi_a \leq 1.797 \times 10^{-1}.$$

Hence the same local objective tolerance already implies the rank-one dominant-channel candidates

$$a_a^{(1)} \leq 5.46 \times 10^{-3}, \quad \sigma_a^{(1)} \leq 5.47 \times 10^{-3}$$

from the exact canonical ξ_a , and

$$a_a^{(1)} \leq 7.36 \times 10^{-3}, \quad \sigma_a^{(1)} \leq 7.39 \times 10^{-3}$$

from the raw-witness envelope. Feeding the latter into the conditional lift still leaves

$$g_a^{(\text{exact})} \geq 1.3326, \quad \varepsilon_a(\Delta E_a) \leq 5.37 \times 10^{-2}$$

at the same reference defect $\delta_{q,a} \leq 10^{-3}$. This is not yet a full operator theorem, because the rank-one dominant-channel hypothesis itself remains to be justified, but it shows that once

the complement is reduced to one effective contraction channel, the audited parent-side off-diagonal witnesses already drive the remaining log-det loss into the small- 10^{-3} regime rather than leaving it order one. We therefore keep this line as an auxiliary source model rather than as the preferred theorem route.

What the same audit now adds is that, once any such bound on σ_a has been obtained, its consequences for the exact complement contraction are genuine operator inequalities rather than rank-one bookkeeping identities. If the eigenvalues of A_a are $\alpha_{1,a}, \alpha_{2,a} \in [0, 1)$, then

$$\sigma_a = \sum_{i=1}^2 -\log(1 - \alpha_{i,a}) \geq \sum_{i=1}^2 \alpha_{i,a} = \text{tr } A_a,$$

because $-\log(1 - x) \geq x$ on $[0, 1)$. Likewise,

$$e^{-\sigma_a} = \prod_{i=1}^2 (1 - \alpha_{i,a}) \leq 1 - \max_i \alpha_{i,a} = 1 - \|A_a\|,$$

so

$$\|A_a\| \leq 1 - e^{-\sigma_a}.$$

Thus the same conditional σ_a budget immediately implies the true operator bounds

$$\text{tr } A_a \leq \sigma_a, \quad \|A_a\| \leq 1 - e^{-\sigma_a}, \quad H_{\text{eff},a} \succeq e^{-\sigma_a} H_{PP,a},$$

and Corollary 1 then gives

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq \frac{2(1 - e^{-\sigma_a})}{e^{-\sigma_a}}.$$

At the same local objective tolerance $J_* = 10^{-6}$, the exact canonical route $\sigma_a \leq 5.47 \times 10^{-3}$ therefore yields

$$\text{tr } A_a \leq 5.47 \times 10^{-3}, \quad \|A_a\| \leq 5.46 \times 10^{-3}, \quad H_{\text{eff},a} \succeq 0.9945 H_{PP,a},$$

together with

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.10 \times 10^{-2}.$$

The conservative raw off-diagonal envelope $\sigma_a \leq 7.39 \times 10^{-3}$ still gives

$$\text{tr } A_a \leq 7.39 \times 10^{-3}, \quad \|A_a\| \leq 7.36 \times 10^{-3}, \quad H_{\text{eff},a} \succeq 0.9926 H_{PP,a},$$

with

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.48 \times 10^{-2}.$$

So the remaining gap is now narrower than before: the open theorem target is no longer to guess whether the complement is small, but to prove a parent-side hypothesis that supplies a comparably small loss budget. The determinant-deficit reformulation below is the preferred theorem route; the rank-one dominant-channel line above should be read only as an auxiliary source model.

Proposition O4 (Determinant-deficit control). One can remove the explicit rank-one step from the operator consequences once a determinant-deficit budget has been identified. For any 2×2 positive contraction A_a with eigenvalues $\alpha_{1,a}, \alpha_{2,a} \in [0, 1)$, define

$$u_a := 1 - \det(I - A_a) = \alpha_{1,a} + \alpha_{2,a} - \alpha_{1,a}\alpha_{2,a}.$$

Then

$$u_a - \alpha_{1,a} = \alpha_{2,a}(1 - \alpha_{1,a}) \geq 0, \quad u_a - \alpha_{2,a} = \alpha_{1,a}(1 - \alpha_{2,a}) \geq 0,$$

so

$$\|A_a\| = \max(\alpha_{1,a}, \alpha_{2,a}) \leq u_a.$$

Writing $s_a := \alpha_{1,a} + \alpha_{2,a} = \text{tr } A_a$ and using $\alpha_{1,a}\alpha_{2,a} \leq s_a^2/4$ gives

$$u_a = s_a - \alpha_{1,a}\alpha_{2,a} \geq s_a - \frac{s_a^2}{4},$$

hence

$$\text{tr } A_a \leq 2(1 - \sqrt{1 - u_a}).$$

Therefore any parent-side bound $u_a \leq u_* < 1$ already implies

$$\|A_a\| \leq u_*, \quad \text{tr } A_a \leq 2(1 - \sqrt{1 - u_*}), \quad H_{\text{eff},a} \succeq (1 - u_*)H_{PP,a},$$

and Corollary 1 becomes

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq \frac{2u_*}{1 - u_*}.$$

So the operator theorem target can be stated without any single-channel hypothesis: it is enough to control the determinant deficit u_a on the audited window.

Proposition O5 (Continuum neighborhood determinant-deficit bound). The exact projected kernel family already carries a canonical determinant-deficit candidate through its local factor

$$1 - \lambda^2 \xi_a, \quad u_a(\alpha, \beta, \lambda) := \lambda^2 \xi_a(\alpha, \beta).$$

Within the canonical local box

$$|\alpha - 1| \leq 0.1, \quad |\beta - 1| \leq 0.1, \quad |\lambda| \leq 0.2,$$

the same coercive audit used above gives

$$\frac{1}{2} \delta\theta^\top H \delta\theta + C_4 \lambda^4 \leq J_*/c_{\text{loc}} \quad \text{whenever } J_{\text{parent}} \leq J_*.$$

Hence the admissible family is reduced to a compact coercive neighborhood, and the determinant-deficit source is obtained by maximizing the smooth scalar $u_a(\alpha, \beta, \lambda)$ on that set.

The audited optimum at $J_* = 10^{-6}$ occurs near $(\alpha, \beta, \lambda) \approx (1.01, 0.992, 0.2)$ and yields

$$u_a \leq 5.34 \times 10^{-3}$$

so Proposition O4 immediately gives

$$\text{tr } A_a \leq 5.35 \times 10^{-3}, \quad \|A_a\| \leq 5.34 \times 10^{-3}, \quad H_{\text{eff},a} \succeq 0.9947 H_{PP,a},$$

and

$$|\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.08 \times 10^{-2}.$$

Proof sketch. The coercive inequality controls both the quadratic (α, β) displacement and the quartic mixing direction λ , so the sublevel set $J_{\text{parent}} \leq J_*$ stays inside a small neighborhood where the exact projected factor $1 - \lambda^2 \xi_a$ remains valid and u_a is uniformly bounded. The proposition is therefore reduced to one constrained scalar optimization, after which the operator consequences follow monotonically from Proposition O4. Thus the main proof-side task is no longer to search for an unconstrained small complement, but to promote this continuum neighborhood budget for u_a from the audited projected family to a genuine parent-side operator statement.

Corollary O1 (Neighborhood budget on the on-shell window). If the continuum neighborhood budget of Proposition O5 is inherited by the exact Schur-reduced operator on the audited on-shell window, then at $J_* = 10^{-6}$ one has

$$g_a^{(\text{exact})} \geq 1.3339, \quad 1 - \nu_{1,a} \leq 8.43 \times 10^{-2}, \quad \varepsilon_a(\Delta E_a) \leq 5.35 \times 10^{-2}$$

under the same reference defect target $\delta_{q,a} \leq 10^{-3}$ used in Theorem Target O1. In other words, Proposition O5 already compresses the remaining first-mode statement to a single parent-side task: justify that the audited continuum neighborhood budget survives the lift from the projected family to the exact Schur complement.

Corollary O2 (On-shell defect closure on the audited knot set). The $\delta_{q,a}$ side of Theorem Target O1 is already closed on the audited knot set. Indeed, the parent-bath factorization of Proposition O1 identifies

$$K_{\text{spec},a} = \text{diag}(S_{zz,a}(0), S_{xx,a}(\Delta E_a)), \quad q_a := \frac{R_a(\Delta E_a)}{R_a(0)} = \frac{S_{xx,a}(\Delta E_a)}{S_{zz,a}(0)}.$$

By Proposition O2 the same audited family is gap-locked,

$$\tau_{\text{env},a} = \omega_{1,a}^{-1}, \quad q_a = \frac{1}{1 + (\Delta E_a / \omega_{1,a})^2} = f(r_a).$$

The new audit `code/audit_open_system_deltaq_source.py` verifies this equality directly from the parent-bath statement map and the micro bridge: on the current D4-20 knot set,

$$\max |\tau_{\text{env},a} - \omega_{1,a}^{-1}| = 1.11 \times 10^{-16},$$

$$\max |\delta_{q,a}| = \max |q_a - f(r_a)| \leq 4.44 \times 10^{-16},$$

with the direct micro ratio staying at the same level,

$$\max \left| \frac{S_{xx,a}(\Delta E_a)}{S_{zz,a}(0)} - f(r_a) \right| \leq 3.33 \times 10^{-16}.$$

Feeding this actual defect into the gap floors already available from Proposition O5 gives

$$1 - \nu_{1,a} \leq 5.08 \times 10^{-14}, \quad \varepsilon_a(\Delta E_a) \leq 2.94 \times 10^{-14}$$

for the conservative analytic floor $g_a \geq 1.3375$, and

$$1 - \nu_{1,a} \leq 4.93 \times 10^{-14}, \quad \varepsilon_a(\Delta E_a) \leq 2.79 \times 10^{-14}$$

for the sharper exact-family floor $g_a \geq 1.3537$. Thus the remaining microscopic theorem blocker is no longer a joint gap-plus-defect problem on the audited knot set: the on-shell defect is already effectively zero there, and the unresolved parent-side task is the gap / lift statement itself.

Corollary O3 (Audited bath-family rigidity source for the exact lift). The surviving gap / lift statement also admits a sharper audit-level source than the earlier conditional phrasing might suggest. Across the four explicit parent-bath deformation scans already exported on the same knot set—the linear family audit, the generator-affinity audit, the log-coordinate audit, and the normal-coordinate audit—the exact statement is always selected at the undeformed point. Concretely, the audited minimizers are

$$(m, u, v) = (0, 0, 0), \quad (q_{ss}, q_{bb}, q_{sb}) = (0, 0, 0),$$

$$(p_{\text{sys}}, p_{\text{spec}}) = (0, 0), \quad (\zeta_{\text{sys}}, \zeta_{\text{spec}}) = (0, 0).$$

The corresponding canonical objectives stay at

$$J_{\text{can}} \leq 3.18 \times 10^{-12},$$

while the smallest audited runner-up gap among those four scans is already

$$\Delta J_{\text{sel}} \geq 6.90 \times 10^{-4}.$$

Hence the exact parent-bath statement is canonically rigid on the present audit scope: the weakest observed deformation gap still exceeds the continuum-neighborhood tolerance $J_* = 10^{-6}$ from Proposition O5 by a factor

$$\Delta J_{\text{sel}}/J_* \geq 6.90 \times 10^2,$$

and exceeds the largest canonical residual by

$$\Delta J_{\text{sel}}/J_{\text{can}}^{\text{max}} \geq 2.17 \times 10^8.$$

This rigidity is consistent with the companion exactness audits,

$$\max |\log(\gamma_{\text{mix}}/\gamma_{\phi}) - (\log \mathcal{A}_{\text{sys}} + \log \mathcal{A}_{\text{bath}})| \leq 8.43 \times 10^{-13},$$

$$\max |\Delta_{\text{cocycle}}^{\text{mix}}| \leq 1.11 \times 10^{-15}, \quad \max |\Delta_{\text{chain}}^{\text{mix}}| \leq 4.03 \times 10^{-15},$$

from the parent-bath statement, cocycle, and potential/chain reconstructions. The new consolidation audit `code/audit_open_system_exact_lift_source.py` therefore supports the following reviewer-facing reading: on the current audited knot set, the continuum neighborhood budget of Proposition O5 is not merely an arbitrary local ansatz, but is already compatible with the uniquely selected exact parent-bath statement. Using the same audited budget

$$u_a \leq 5.34 \times 10^{-3}$$

therefore leaves the exact lifted floor

$$g_a^{(\text{exact})} \geq 1.3339, \quad H_{\text{eff},a} \geq 0.9947 H_{PP,a},$$

and, together with Corollary O2, forces only

$$1 - \nu_{1,a} \leq 5.12 \times 10^{-14}, \quad \varepsilon_a(\Delta E_a) \leq 2.97 \times 10^{-14}$$

on the audited knot set. In this narrowed audit-level sense, the exact-lift source is no longer a free-standing mystery: the remaining proof-level step is to replace this canonical rigidity evidence by a genuinely operator-level theorem, not to discover a new gap scale.

Proposition O6 (Canonical-rigidity transfer on the audited family classes). The previous corollary can be sharpened into a conditional transfer statement on the explicit deformation families already scanned. Let $\mathfrak{F}$ denote any one of the four audited parent-bath classes

$$\mathfrak{F} \in \{\text{linear family, generator-affinity, log-coordinate, normal-coordinate}\}.$$

Suppose an exact Schur-side candidate remains inside one such family and satisfies the same local tolerance

$$J_{\text{parent}} \leq J_* = 10^{-6}.$$

Then on the current audit grid the candidate is forced back to the undeformed canonical point of that family. Equivalently, within each of the four audited classes the canonical point is the unique admissible point at the theorem tolerance.

Audit proof. The new transfer audit `code/audit_open_system_rigidity_transfer.py` checks the full exported scan tables directly rather than only comparing runner-up gaps. At $J_* = 10^{-6}$ it finds, for all four families,

$$N_{\text{adm}}^{(\text{total})} = 1, \quad N_{\text{adm}}^{(\text{noncan})} = 0,$$

so the only surviving point is the undeformed canonical one. Quantitatively, the largest canonical objective among the four scans is still only

$$J_{\text{can}}^{\text{max}} = 3.18 \times 10^{-12},$$

while the smallest non-canonical objective is already

$$J_{\text{noncan}}^{\text{min}} = 6.90 \times 10^{-4},$$

hence

$$J_{\text{noncan}}^{\text{min}} - J_{\text{can}}^{\text{max}} \geq 6.90 \times 10^{-4}.$$

This leaves the same margins as before,

$$\frac{J_{\text{noncan}}^{\text{min}} - J_{\text{can}}^{\text{max}}}{J_*} \geq 6.90 \times 10^2, \quad \frac{J_{\text{noncan}}^{\text{min}} - J_{\text{can}}^{\text{max}}}{J_{\text{can}}^{\text{max}}} \geq 2.17 \times 10^8,$$

but now in the stronger form that no non-canonical family member survives the theorem tolerance at all. Therefore, conditional on the exact operator being represented by any of these audited deformation classes, the continuum-neighborhood determinant-deficit budget of Proposition O5 transfers unchanged to the exact Schur-reduced operator on the audited knot set.

Corollary O4 (Exact-lift consequences on the audited family classes). Under the family-class hypothesis of Proposition O6, the exact Schur-reduced operator inherits the same budget

$$u_a \leq 5.34 \times 10^{-3},$$

and therefore the same operator consequences,

$$g_a^{(\text{exact})} \geq 1.3339, \quad H_{\text{eff},a} \geq 0.9947 H_{PP,a},$$

on the audited knot set. Together with Corollary O2 this leaves only

$$1 - \nu_{1,a} \leq 5.12 \times 10^{-14}, \quad \varepsilon_a(\Delta E_a) \leq 2.97 \times 10^{-14}.$$

Thus, on the currently exported family classes, the unresolved theorem gap is no longer whether the lift survives, but whether one can remove the explicit family-class hypothesis and prove the same transfer directly at the operator level.

Proposition O7 (Operator-witness chart transfer on the audited chart). The preceding family-class hypothesis can be weakened one step further by replacing the explicit deformation labels with operator witnesses. Define the p95 and max witness vectors

$$W_{a,95} := (\Delta_{11,a}^{95}, \Delta_{22,a}^{95}, \Delta_{\text{tr},a}^{95}, \Delta_{\text{det},a}^{95}, \Delta_{\text{ani},a}^{95}),$$

$$W_{a,\text{max}} := (\Delta_{11,a}^{\text{max}}, \Delta_{22,a}^{\text{max}}, \Delta_{\text{tr},a}^{\text{max}}, \Delta_{\text{det},a}^{\text{max}}, \Delta_{\text{ani},a}^{\text{max}}),$$

where the entries are exactly the block, trace, determinant, and anisotropy residuals already used in the parent-bath family scans. Suppose an exact Schur-side candidate satisfies the same local tolerance

$$J_{\text{parent}} \leq J_* = 10^{-6},$$

and lies in the audited operator witness chart

$$\|W_{a,95}\|_{\infty} \leq 9.62 \times 10^{-5}, \quad \|W_{a,\text{max}}\|_{\infty} \leq 1.01 \times 10^{-4}.$$

Then, on the current audited chart, the candidate belongs to the same canonical pocket as the undeformed point and therefore inherits the Proposition O5 determinant-deficit budget.

Audit proof. The new aggregation audit `code/audit_open_system_operator_witness_transfer.py` discards the family labels and stacks all non-canonical rows from the four exported

scans directly in witness space. Across that union, the nearest non-canonical challenger is already the generator-affinity point $q_{bb} = \pm 0.0625$, with

$$\min_{\text{noncan}} \|W_{a,95}\|_{\infty} = 1.92 \times 10^{-4}, \quad \min_{\text{noncan}} \|W_{a,\max}\|_{\infty} = 2.01 \times 10^{-4}.$$

The canonical witness suprema remain only

$$\|W_{a,95}^{(\text{can})}\|_{\infty} \leq 8.17 \times 10^{-13}, \quad \|W_{a,\max}^{(\text{can})}\|_{\infty} \leq 8.43 \times 10^{-13},$$

so the half-radius chart bounds above still leave margins of

$$\frac{9.62 \times 10^{-5}}{\|W_{a,95}^{(\text{can})}\|_{\infty}} \geq 1.18 \times 10^8, \quad \frac{1.01 \times 10^{-4}}{\|W_{a,\max}^{(\text{can})}\|_{\infty}} \geq 1.19 \times 10^8.$$

Therefore, on the audited operator chart, any exact candidate satisfying the same theorem tolerance and these witness-radius bounds cannot land on any scanned non-canonical deformation and is forced into the canonical pocket. This removes the explicit family-class labels from the transfer statement, even though the last remaining proof step is still to justify the witness-radius bound directly from the exact operator.

Corollary O5 (Exact-lift consequences on the audited operator chart). Under the witness-chart hypothesis of Proposition O7, the exact Schur-reduced operator again inherits the same neighborhood budget

$$u_a \leq 5.34 \times 10^{-3},$$

and hence the same operator consequences,

$$g_a^{(\text{exact})} \geq 1.3339, \quad H_{\text{eff},a} \geq 0.9947 H_{PP,a},$$

on the audited knot set. Together with Corollary O2 this again leaves only

$$1 - \nu_{1,a} \leq 5.12 \times 10^{-14}, \quad \varepsilon_a(\Delta E_a) \leq 2.97 \times 10^{-14}.$$

Thus the remaining theorem gap is now narrower still: derive the witness-radius bound directly from the exact operator, rather than from the exported scan chart.

Corollary O6 (The exported exact operator lies inside the audited witness chart).

The chart hypothesis of Proposition O7 is itself now verified directly on the exported exact operator data. Reading the exact parent-bath statement map rowwise and aggregating it with the same p95/max conventions used in the family scans gives

$$\|W_{95}^{(\text{exact})}\|_{\infty} = 3.52 \times 10^{-7}, \quad \|W_{\max}^{(\text{exact})}\|_{\infty} = 3.63 \times 10^{-7},$$

with the dominant component in both cases coming from the exact block-22 witness. These are already smaller than the audited chart radii of Proposition O7 by factors

$$\frac{9.62 \times 10^{-5}}{\|W_{95}^{(\text{exact})}\|_\infty} = 2.73 \times 10^2, \quad \frac{1.01 \times 10^{-4}}{\|W_{\max}^{(\text{exact})}\|_\infty} = 2.77 \times 10^2.$$

The companion exactness witnesses remain tiny,

$$\Delta_{\text{cocycle}}^{\max} = 1.11 \times 10^{-15}, \quad \Delta_{\text{chain}}^{\max} = 4.03 \times 10^{-15},$$

so the exported exact operator already lies well inside the audited witness chart. Consequently the conclusions of Corollary O5 apply directly to the exported exact Schur-reduced operator on the audited knot set:

$$g_a^{(\text{exact})} \geq 1.3339, \quad H_{\text{eff},a} \succeq 0.9947 H_{PP,a},$$

and, together with Corollary O2,

$$1 - \nu_{1,a} \leq 5.12 \times 10^{-14}, \quad \varepsilon_a(\Delta E_a) \leq 2.97 \times 10^{-14}.$$

The remaining proof-level gap is now very specific: promote this verified chart inclusion of the exported exact operator into an analytic operator theorem.

Proposition O8 (Exact-equation witness envelope for the exported operator). The preceding chart inclusion can be compressed one step further into a single operator-envelope bound that no longer needs the rowwise chart aggregation. Define

$$\mathcal{B}_{\text{op}}^{(\text{exact})} := \max\{\Delta_\phi^{(\text{fact})}, \Delta_{\text{mix}}^{(\text{fact})}, \Delta_{\text{tr}}^{(\text{stmt})}, \Delta_{\text{det}}^{(\text{stmt})}, \Delta_{\text{ani}}^{(\text{stmt})}, \Delta_{\text{mix}}^{(\text{cocycle})}, \Delta_{\text{mix}}^{(\text{potential})}\},$$

with

$$\Delta_\phi^{(\text{fact})} := \max |\kappa_\phi^{(\text{obs})} - \kappa_{\text{env}}|, \quad \Delta_{\text{mix}}^{(\text{fact})} := \max |\kappa_{\text{mix}}^{(\text{obs})} - \kappa_{\text{env}}|,$$

$$\Delta_{\text{tr}}^{(\text{stmt})} := \max |\text{tr}(K_{\text{bath}}) - \text{tr}(\Gamma)|, \quad \Delta_{\text{det}}^{(\text{stmt})} := \max |\det(K_{\text{bath}}) - \det(\Gamma)|,$$

$$\Delta_{\text{ani}}^{(\text{stmt})} := \max |\log(\gamma_{\text{mix}}/\gamma_\phi) - \log(\text{sys}) - \log(\text{bath})|,$$

and where the last two terms are the exact mix-side cocycle and potential envelopes already exported by the parent-bath closure audits. Then the exact witness equations imply

$$\|W^{(\text{exact})}\|_\infty \leq \mathcal{B}_{\text{op}}^{(\text{exact})}.$$

Audit proof. /Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_operator_radius_bound.py reads the factorization, statement, cocycle, and potential summaries directly and evaluates the envelope components

$$\begin{aligned}\Delta_{\phi}^{(\text{fact})} &= 2.33 \times 10^{-10}, & \Delta_{\text{mix}}^{(\text{fact})} &= 3.63 \times 10^{-7}, \\ \Delta_{\text{tr}}^{(\text{stmt})} &= 4.40 \times 10^{-14}, & \Delta_{\text{det}}^{(\text{stmt})} &= 4.21 \times 10^{-18}, & \Delta_{\text{ani}}^{(\text{stmt})} &= 8.43 \times 10^{-13}, \\ \Delta_{\text{mix}}^{(\text{cocycle})} &= 8.64 \times 10^{-13}, & \Delta_{\text{mix}}^{(\text{potential})} &= 5.50 \times 10^{-13}.\end{aligned}$$

Hence

$$\mathcal{B}_{\text{op}}^{(\text{exact})} = 3.63 \times 10^{-7},$$

with the envelope dominated by the exact factorization-side block-22 witness. The same audit checks that this conservative envelope is tight on the exported operator, i.e.

$$\mathcal{B}_{\text{op}}^{(\text{exact})} = \|W_{\text{max}}^{(\text{exact})}\|_{\infty},$$

so no hidden larger witness remains outside the listed exact equations.

Corollary O7 (Half-radius chart inclusion from the exact witness equations). The envelope of Proposition O8 is already smaller than one half of both audited chart radii:

$$\begin{aligned}\mathcal{B}_{\text{op}}^{(\text{exact})} &= 3.63 \times 10^{-7} \leq \frac{1}{2}r_{95}^{\text{safe}} = 4.81 \times 10^{-5}, \\ \mathcal{B}_{\text{op}}^{(\text{exact})} &= 3.63 \times 10^{-7} \leq \frac{1}{2}r_{\text{max}}^{\text{safe}} = 5.03 \times 10^{-5}.\end{aligned}$$

Equivalently,

$$\frac{\frac{1}{2}r_{95}^{\text{safe}}}{\mathcal{B}_{\text{op}}^{(\text{exact})}} = 1.33 \times 10^2, \quad \frac{\frac{1}{2}r_{\text{max}}^{\text{safe}}}{\mathcal{B}_{\text{op}}^{(\text{exact})}} = 1.39 \times 10^2.$$

Therefore the exported exact operator satisfies the witness-chart hypothesis of Proposition O7 with half-radius room to spare, directly from the exact factorization/statement/cocycle/potential equations themselves rather than from the scan chart alone. Consequently the exact-lift conclusions of Corollary O5 follow again for the exported operator on the audited knot set. The proof-level gap is now reduced to a final analytic upgrade: replace these exported exact residual constants by an operator theorem that bounds the same envelope a priori.

Proposition O9 (Mix-side factorization witness half-radius bound). The dominant term in the envelope of Proposition O8 can itself be reduced to a simple quotient bound. Rowwise,

$$\Delta_{\text{mix}}^{(\text{fact})}(D) := \left| \frac{\gamma_{\text{mix}}(D)}{g_x(D)^2 S_{xx}(D)} - \kappa_{\text{env}} \right| = \frac{|\gamma_{\text{mix}}(D) - \kappa_{\text{env}} g_x(D)^2 S_{xx}(D)|}{g_x(D)^2 S_{xx}(D)}.$$

Therefore any uniform numerator bound

$$\varepsilon_{\text{mix}} \geq \max_D |\gamma_{\text{mix}}(D) - \kappa_{\text{env}} g_x(D)^2 S_{xx}(D)|$$

and any positive lower floor

$$m_{\text{mix}} \leq \min_D g_x(D)^2 S_{xx}(D)$$

imply the operator bound

$$\|\Delta_{\text{mix}}^{(\text{fact})}\|_{\infty} \leq \frac{\varepsilon_{\text{mix}}}{m_{\text{mix}}}.$$

Audit proof. /Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_mixer_factorization_source.py evaluates this quotient source directly on the exported exact factorization map. The common numerator bound is

$$\varepsilon_{\text{mix}} = 4.40 \times 10^{-14}.$$

Three natural lower floors for the denominator all remain positive on the audited knot set:

$$m_{\text{mix}}^{(\text{dir})} = \min_D g_x(D)^2 S_{xx}(D) = 1.2147 \times 10^{-9},$$

$$m_{\text{mix}}^{(\gamma)} = \frac{\min_D \gamma_{\text{mix}}(D) - \varepsilon_{\text{mix}}}{\kappa_{\text{env}}^{\text{max}}} = 1.2147 \times 10^{-9},$$

$$m_{\text{mix}}^{(\text{sep})} = \left(\min_D g_x(D)\right)^2 \min_D S_{xx}(D) = 1.1417 \times 10^{-9}.$$

Hence

$$\|\Delta_{\text{mix}}^{(\text{fact})}\|_{\infty} \leq \frac{\varepsilon_{\text{mix}}}{m_{\text{mix}}^{(\text{dir})}} = 3.62 \times 10^{-5},$$

$$\|\Delta_{\text{mix}}^{(\text{fact})}\|_{\infty} \leq \frac{\varepsilon_{\text{mix}}}{m_{\text{mix}}^{(\gamma)}} = 3.62 \times 10^{-5},$$

and even the fully separated floor still gives

$$\|\Delta_{\text{mix}}^{(\text{fact})}\|_{\infty} \leq \frac{\varepsilon_{\text{mix}}}{m_{\text{mix}}^{(\text{sep})}} = 3.85 \times 10^{-5}.$$

All three bounds are already below the Proposition O7 half-radii

$$\frac{1}{2}r_{95}^{\text{safe}} = 4.81 \times 10^{-5}, \quad \frac{1}{2}r_{\text{max}}^{\text{safe}} = 5.03 \times 10^{-5}.$$

So the dominant witness component itself now satisfies the audited half-radius condition from a simple factorization numerator/denominator estimate.

Corollary O8 (Remaining analytic gap reduced to a lower floor on the mix kernel).

Proposition O9 shows that the dominant block-22 witness no longer requires the full chart machinery: on the exported exact operator it is already controlled by one small numerator residual and one positive lower floor on $g_x^2 S_{xx}$ (equivalently on γ_{mix} after dividing by κ_{env}). Since all other witness components are already much smaller by Proposition O8, the remaining proof-level gap is now narrower again:

$$\text{prove an a priori lower bound on } g_x(D)^2 S_{xx}(D) \quad \text{or equivalently on } \gamma_{\text{mix}}(D),$$

strong enough to reproduce one of the denominator floors above. Once that is available, the half-radius chart inclusion of the exact operator follows without any remaining appeal to exported scan geometry.

Proposition O10 (Observable-chain source for the γ_{mix} floor). The mix rate also admits an exact observable-chain reconstruction:

$$\gamma_{\text{mix}}(D) = \frac{\gamma_{\text{ref}}(D) \chi_{\text{eff}}(D)}{2 C_{\text{max}}(D)}.$$

Hence any lower floors

$$\gamma_{\text{ref}}(D) \geq \gamma_{\text{ref}}^{\text{floor}}, \quad \chi_{\text{eff}}(D) \geq \chi_{\text{eff}}^{\text{floor}}, \quad C_{\text{max}}(D) \leq C_{\text{max}}^{\text{ceil}},$$

imply the mix-rate floor

$$\gamma_{\text{mix}}(D) \geq \frac{\gamma_{\text{ref}}^{\text{floor}} \chi_{\text{eff}}^{\text{floor}}}{2 C_{\text{max}}^{\text{ceil}}}.$$

Dividing by κ_{env} then yields a lower bound on the mix kernel,

$$g_x(D)^2 S_{xx}(D) \geq \frac{\gamma_{\text{ref}}^{\text{floor}} \chi_{\text{eff}}^{\text{floor}}}{2 C_{\text{max}}^{\text{ceil}} \kappa_{\text{env}}^{\text{max}}}.$$

Audit proof. /Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_gamma_mix_floor_source.py verifies the exact identity above with

$$\max_D \left| \gamma_{\text{mix}}(D) - \frac{\gamma_{\text{ref}}(D) \chi_{\text{eff}}(D)}{2 C_{\text{max}}(D)} \right| = 4.12 \times 10^{-14}.$$

On the exported exact operator, the conservative separated floors are

$$\gamma_{\text{ref}}^{\text{floor}} = 4.75 \times 10^2, \quad \chi_{\text{eff}}^{\text{floor}} = 9.03 \times 10^{-7}, \quad C_{\text{max}}^{\text{ceil}} = 4.98 \times 10^{-1},$$

which give

$$\gamma_{\text{mix}}^{(\text{sep floor})} = \frac{\gamma_{\text{ref}}^{\text{floor}} \chi_{\text{eff}}^{\text{floor}}}{2 C_{\text{max}}^{\text{ceil}}} = 4.31 \times 10^{-4}.$$

The actual exact minimum is

$$\min_D \gamma_{\text{mix}}(D) = 5.23 \times 10^{-4},$$

so the separated observable-chain floor already captures the true minimum within a factor

$$\frac{\min_D \gamma_{\text{mix}}(D)}{\gamma_{\text{mix}}^{(\text{sep floor})}} = 1.21.$$

Since $\kappa_{\text{env}} = 4.31 \times 10^5$ on the audited knot set, this induces the conservative kernel floor

$$g_x(D)^2 S_{xx}(D) \geq 1.00 \times 10^{-9}.$$

Corollary O9 (Observable-chain floor implies the half-radius mix witness bound).

Combining Proposition O10 with the numerator residual

$$\varepsilon_{\text{mix}} = 4.40 \times 10^{-14}$$

of Proposition O9 gives

$$\|\Delta_{\text{mix}}^{(\text{fact})}\|_{\infty} \leq \frac{\varepsilon_{\text{mix}}}{1.00 \times 10^{-9}} = 4.39 \times 10^{-5}.$$

This is still below both audited half-radii,

$$4.39 \times 10^{-5} < \frac{1}{2} r_{95}^{\text{safe}} = 4.81 \times 10^{-5}, \quad 4.39 \times 10^{-5} < \frac{1}{2} r_{\text{max}}^{\text{safe}} = 5.03 \times 10^{-5},$$

with margins

$$\frac{\frac{1}{2} r_{95}^{\text{safe}}}{4.39 \times 10^{-5}} = 1.09, \quad \frac{\frac{1}{2} r_{\text{max}}^{\text{safe}}}{4.39 \times 10^{-5}} = 1.15.$$

Thus even the separated observable-chain floor on γ_{mix} is already sufficient to recover the half-radius bound for the dominant block-22 witness. The remaining proof-level task is therefore yet narrower: justify the individual floors on γ_{ref} , χ_{eff} , and C_{max} (or a directly equivalent mix-rate floor) from the exact operator side.

Proposition O11 (First-mode source for the γ_{ref} floor). The reference rate is not an independent observable-chain input. In the localized-channel extraction it is defined by

$$\gamma_{\text{ref}}(D) = \sqrt{\Gamma_1(D)\Gamma_2(D)}, \quad \Gamma_1(D) = \omega_1(D)^9, \quad \Gamma_2(D) = \omega_1(D)^{13},$$

hence

$$\gamma_{\text{ref}}(D) = \omega_1(D)^{11}.$$

On the audited open-system micro bridge one also has the gap-locked environmental timescale

$$\tau_{\text{env}}(D) = \omega_1(D)^{-1},$$

so equivalently

$$\gamma_{\text{ref}}(D) = \tau_{\text{env}}(D)^{-11}.$$

Therefore any first-mode lower floor

$$\omega_1(D) \geq \omega_1^{\text{floor}}$$

or, equivalently, any gap-locked timescale ceiling

$$\tau_{\text{env}}(D) \leq \tau_{\text{env}}^{\text{ceil}}$$

immediately implies

$$\gamma_{\text{ref}}(D) \geq (\omega_1^{\text{floor}})^{11} = (\tau_{\text{env}}^{\text{ceil}})^{-11}.$$

Audit proof. /Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_gamm
a_ref_floor_source.py verifies on the exported exact bridge that

$$\max_D |\gamma_{\text{ref}}(D) - \omega_1(D)^{11}| = 1.56 \times 10^{-11},$$

$$\max_D |\gamma_{\text{ref}}(D) - \tau_{\text{env}}(D)^{-11}| = 1.59 \times 10^{-11},$$

$$\max_D |\tau_{\text{env}}(D) - \omega_1(D)^{-1}| = 1.11 \times 10^{-16}.$$

The exact audited floors/ceilings are

$$\omega_1^{\text{floor}} = 1.751201\dots, \quad \tau_{\text{env}}^{\text{ceil}} = 0.5710366\dots,$$

both attained at $D = 4$, and they induce the same reference-rate floor

$$\gamma_{\text{ref}}^{\text{floor}} = (\omega_1^{\text{floor}})^{11} = (\tau_{\text{env}}^{\text{ceil}})^{-11} = 475.002983\dots,$$

which matches the actual exact minimum of γ_{ref} to machine precision.

Corollary O10 (γ_{ref} **is no longer an independent blocker**). Combining Proposition O11 with Proposition O10 shows that the γ_{ref} side of the observable-chain floor is already reduced to the first-mode floor ω_1^{floor} (equivalently the gap-locked timescale ceiling $\tau_{\text{env}}^{\text{ceil}}$). In particular,

on the current exact export the separated observable-chain estimate of Proposition O10 can be rewritten with no additional loss as

$$\gamma_{\text{mix}}(D) \geq \frac{(\omega_1^{\text{floor}})^{11} \chi_{\text{eff}}^{\text{floor}}}{2C_{\text{max}}^{\text{ceil}}} = \frac{(\tau_{\text{env}}^{\text{ceil}})^{-11} \chi_{\text{eff}}^{\text{floor}}}{2C_{\text{max}}^{\text{ceil}}}.$$

So after Proposition O11, the remaining proof-level task narrows again: γ_{ref} no longer needs a separate source argument, and the residual observable-chain closure is concentrated on the floor/ceiling pair $(\chi_{\text{eff}}^{\text{floor}}, C_{\text{max}}^{\text{ceil}})$ or on a directly equivalent mix-rate floor.

Proposition O12 (Universal two-level source for the C_{max} ceiling). Let $\rho(t)$ be any positive semidefinite 2×2 density matrix with unit trace. Writing

$$\rho(t) = \begin{pmatrix} \rho_{LL}(t) & \rho_{LR}(t) \\ \rho_{RL}(t) & \rho_{RR}(t) \end{pmatrix}, \quad \rho_{LL}(t) + \rho_{RR}(t) = 1,$$

positivity implies

$$|\rho_{LR}(t)|^2 \leq \rho_{LL}(t)\rho_{RR}(t) \leq \left(\frac{\rho_{LL}(t) + \rho_{RR}(t)}{2} \right)^2 = \frac{1}{4}.$$

Hence

$$|\rho_{LR}(t)| \leq \frac{1}{2} \quad \text{for all } t,$$

and therefore the open-system coherence response

$$C_{\text{max}} := \max_t |\rho_{LR}(t)|$$

obeys the universal ceiling

$$C_{\text{max}} \leq \frac{1}{2}.$$

Audit check. `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_cmax_ceiling_source.py` verifies on the exported exact micro bridge that

$$\max_D C_{\text{max}}(D) = 0.497539 \dots < \frac{1}{2},$$

with strict slack

$$\frac{1}{2} - \max_D C_{\text{max}}(D) = 2.461 \times 10^{-3}.$$

Equivalently,

$$\frac{1}{4} - \max_D C_{\text{max}}(D)^2 = 2.455 \times 10^{-3} > 0.$$

Corollary O11 (Universal $C_{\max}$ ceiling still implies the half-radius mix witness bound). Combining Proposition O12 with Proposition O11 and the observable-chain identity of Proposition O10 yields

$$\gamma_{\text{mix}}(D) \geq \frac{\gamma_{\text{ref}}^{\text{floor}} \chi_{\text{eff}}^{\text{floor}}}{2(\frac{1}{2})} = \gamma_{\text{ref}}^{\text{floor}} \chi_{\text{eff}}^{\text{floor}}.$$

On the current exact export this gives the conservative floor

$$\gamma_{\text{mix}}(D) \geq 4.29102 \times 10^{-4},$$

and therefore

$$g_x(D)^2 S_{xx}(D) \geq 9.96645 \times 10^{-10}.$$

Using again the common numerator residual

$$\varepsilon_{\text{mix}} = 4.40 \times 10^{-14}$$

from Proposition O9, one obtains

$$\|\Delta_{\text{mix}}^{(\text{fact})}\|_{\infty} \leq \frac{\varepsilon_{\text{mix}}}{9.96645 \times 10^{-10}} = 4.41 \times 10^{-5}.$$

This remains below both half-radii,

$$4.41 \times 10^{-5} < \frac{1}{2} r_{95}^{\text{safe}} = 4.81 \times 10^{-5}, \quad 4.41 \times 10^{-5} < \frac{1}{2} r_{\max}^{\text{safe}} = 5.03 \times 10^{-5},$$

with margins

$$\frac{\frac{1}{2} r_{95}^{\text{safe}}}{4.41 \times 10^{-5}} = 1.09, \quad \frac{\frac{1}{2} r_{\max}^{\text{safe}}}{4.41 \times 10^{-5}} = 1.14.$$

So after Proposition O12, the universal two-level positivity bound already removes $C_{\max}$ as an independent proof-level obstacle. The residual observable-chain closure is now narrower one final time: it is concentrated on $\chi_{\text{eff}}^{\text{floor}}$, or on a directly equivalent floor for γ_{mix} .

Proposition O13 (Reduction of the χ_{eff} floor to a coherence-response floor). On the exact micro bridge one has the exact identities

$$\chi_{\text{eff}}(D) = \chi_{LR}(D) \frac{\chi_{\text{eff}}(D)}{\chi_{LR}(D)} = \frac{2\gamma_{\text{mix}}(D)C_{\max}(D)}{\gamma_{\text{ref}}(D)},$$

and, since $\chi_{LR}(D) = |M_{LR}(D)|/\gamma_{\text{ref}}(D)$,

$$\frac{\chi_{\text{eff}}(D)}{\chi_{LR}(D)} = \frac{2\gamma_{\text{mix}}(D)C_{\max}(D)}{|M_{LR}(D)|}.$$

Equivalently,

$$\chi_{\text{eff}}(D) = A_{\chi}(D) C_{\max}(D), \quad A_{\chi}(D) := \frac{2\gamma_{\text{mix}}(D)}{\gamma_{\text{ref}}(D)} = \kappa_{\text{env}} \frac{2g_x(D)^2 S_{xx}(D)}{\gamma_{\text{ref}}(D)}.$$

Using Proposition O11 to rewrite

$$\gamma_{\text{ref}}(D) = \omega_1(D)^{11},$$

any first-mode ceiling $\omega_1(D) \leq \omega_1^{\text{ceil}}$ implies the separated prefactor floor

$$A_\chi(D) \geq \kappa_{\text{env}} \frac{2g_{x,\min}^2 S_{xx}^{\text{floor}}}{(\omega_1^{\text{ceil}})^{11}} : A_\chi^{(\text{sep floor})}.$$

Audit check. /Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_chi_eff_floor_source.py verifies on the exact exported bridge that

$$\begin{aligned} \max_D \left| A_\chi - \frac{2\gamma_{\text{mix}}}{\gamma_{\text{ref}}} \right| &= 1.74 \times 10^{-16}, & \max_D \left| A_\chi - \kappa_{\text{env}} \frac{2g_x^2 S_{xx}}{\gamma_{\text{ref}}} \right| &= 1.01 \times 10^{-16}, \\ \max_D \left| \frac{\chi_{\text{eff}}}{\chi_{LR}} - \frac{2\gamma_{\text{mix}} C_{\text{max}}}{|M_{LR}|} \right| &= 4.72 \times 10^{-16}, & \max_D \left| \chi_{\text{eff}} - \chi_{LR} \frac{\chi_{\text{eff}}}{\chi_{LR}} \right| &= 5.58 \times 10^{-17}. \end{aligned}$$

The same audit finds

$$A_\chi^{\text{exact floor}} = 1.815669 \times 10^{-6}, \quad A_\chi^{(\text{sep floor})} = 1.041474 \times 10^{-6},$$

so the separated prefactor floor already matches the exact one within a factor

$$\frac{A_\chi^{\text{exact floor}}}{A_\chi^{(\text{sep floor})}} = 1.743.$$

Corollary O12 (The remaining χ_{eff} closure is exactly a C_{max} floor). Proposition O13 shows that the residual open-system observable-chain problem can be rewritten as

$$\chi_{\text{eff}}(D) \geq A_\chi^{(\text{sep floor})} C_{\text{max}}(D).$$

So once a coherence-response floor

$$C_{\text{max}}(D) \geq C_{\text{max}}^{\text{floor}}$$

is available, one obtains immediately

$$\chi_{\text{eff}}(D) \geq A_\chi^{(\text{sep floor})} C_{\text{max}}^{\text{floor}}.$$

On the current exact export, using only the exact minimum

$$C_{\text{max}}^{\text{exact floor}} = 0.2342548 \dots$$

already gives the conservative bound

$$\chi_{\text{eff}}(D) \geq 2.43970 \times 10^{-7},$$

while the actual exported floor is

$$\chi_{\text{eff}}^{\text{exact floor}} = 9.03367 \times 10^{-7},$$

so even this coarse separated reduction undershoots the exact floor by only

$$\frac{\chi_{\text{eff}}^{\text{exact floor}}}{A_{\chi}^{(\text{sep floor})} C_{\text{max}}^{\text{exact floor}}} = 3.70.$$

Hence after Propositions O11–O13 the residual observable-chain closure is now fully localized: the only scalar still missing from the exact export to a complete proof-side floor statement is a lower bound on C_{max} itself.

Proposition O14 (Underdamped closed form for C_{max} on the exact two-level bridge).

In the localized two-level Lindblad block used by the exact micro bridge,

$$H = \frac{\Delta E}{2} \sigma_x, \quad L_{\phi} = \sqrt{\gamma_{\phi}} \sigma_z, \quad L_{\text{mix}} = \sqrt{\gamma_{\text{mix}}} \sigma_x,$$

with initial state $\rho(0) = |L\rangle\langle L|$, the Bloch evolution closes on the (y, z) plane:

$$\dot{y} = -\Delta E z - 2(\gamma_{\phi} + \gamma_{\text{mix}})y, \quad \dot{z} = \Delta E y - 2\gamma_{\text{mix}}z.$$

Writing

$$\Gamma := \gamma_{\phi} + 2\gamma_{\text{mix}}, \quad \omega := \sqrt{\Delta E^2 - \gamma_{\phi}^2},$$

the underdamped regime $\Delta E > \gamma_{\phi}$ gives the exact solution

$$y(t) = -\frac{\Delta E}{\omega} e^{-\Gamma t} \sin(\omega t), \quad C(t) = |\rho_{LR}(t)| = \frac{|y(t)|}{2} = \frac{\Delta E}{2\omega} e^{-\Gamma t} |\sin(\omega t)|.$$

The first-lobe maximum occurs at

$$t_* = \omega^{-1} \arctan\left(\frac{\omega}{\Gamma}\right),$$

so

$$C_{\text{max}} = \frac{\Delta E}{2\sqrt{\Gamma^2 + \omega^2}} \exp\left[-\frac{\Gamma}{\omega} \arctan\left(\frac{\omega}{\Gamma}\right)\right].$$

Audit check. /Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_cmax_floor_source.py verifies on the exported exact bridge that all rows are underdamped,

$$\min_D (\Delta E - \gamma_{\phi}) = 2.59273 \times 10^{-1} > 0, \quad \max_D \frac{\gamma_{\phi}}{\Delta E} = 2.07 \times 10^{-4},$$

and that the closed form above matches the exported Lindblad maxima to within

$$\max_D |C_{\text{max}}^{(\text{export})} - C_{\text{max}}^{(\text{closed})}| = 2.82 \times 10^{-4}.$$

The closed-form minimum is attained at

$$D = 5, \quad C_{\max}^{(\text{closed}) \text{ floor}} = 2.34338 \times 10^{-1},$$

while the exported exact minimum is

$$C_{\max}^{\text{exact floor}} = 2.34255 \times 10^{-1},$$

so

$$\frac{C_{\max}^{\text{exact floor}}}{C_{\max}^{(\text{closed}) \text{ floor}}} = 0.999644.$$

Corollary O13 (Separated $C_{\max}$ floor and resulting χ_{eff} floor). On the audited exact parameter box, the same closed form is numerically monotone:

$$\partial_{\Delta E} C_{\max} > 0, \quad \partial_{\gamma_{\phi}} C_{\max} < 0, \quad \partial_{\gamma_{\text{mix}}} C_{\max} < 0,$$

with audited derivative margins

$$\min \partial_{\Delta E} C_{\max} = 3.16 \times 10^{-3}, \quad \max \partial_{\gamma_{\phi}} C_{\max} = -3.98 \times 10^{-1}, \quad \max \partial_{\gamma_{\text{mix}}} C_{\max} = -7.96 \times 10^{-1}.$$

Therefore the exact bridge already admits the separated floor

$$C_{\max}(D) \geq C_{\max}^{(\text{sep floor})} = \frac{\Delta E_{\min}}{2\sqrt{\Gamma_{\max}^2 + \omega_{\min}^2}} \exp \left[-\frac{\Gamma_{\max}}{\omega_{\min}} \arctan \left(\frac{\omega_{\min}}{\Gamma_{\max}} \right) \right],$$

where

$$\Delta E_{\min} = 2.59278 \times 10^{-1}, \quad \gamma_{\phi, \max} = 9.56 \times 10^{-5}, \quad \gamma_{\text{mix}, \max} = 1.48999 \times 10^{-1},$$

$$\omega_{\min} = \sqrt{\Delta E_{\min}^2 - \gamma_{\phi, \max}^2} = 2.59278 \times 10^{-1}, \quad \Gamma_{\max} = \gamma_{\phi, \max} + 2\gamma_{\text{mix}, \max} = 2.98093 \times 10^{-1},$$

giving

$$C_{\max}^{(\text{sep floor})} = 1.44083 \times 10^{-1}.$$

The exact exported floor exceeds this by

$$\frac{C_{\max}^{\text{exact floor}}}{C_{\max}^{(\text{sep floor})}} = 1.626.$$

Combining this with Proposition O13 gives the conservative observable floor

$$\chi_{\text{eff}}(D) \geq A_{\chi}^{(\text{sep floor})} C_{\max}^{(\text{sep floor})} = 1.50059 \times 10^{-7},$$

while the actual exported exact floor is

$$\chi_{\text{eff}}^{\text{exact floor}} = 9.03367 \times 10^{-7},$$

so

$$\frac{\chi_{\text{eff}}^{\text{exact floor}}}{A_{\chi}^{(\text{sep floor})} C_{\text{max}}^{(\text{sep floor})}} = 6.02.$$

Thus after Propositions O13–O14 the observable-chain side is effectively closed on the current exact bridge: the remaining gap is no longer the existence of a floor but only how much of the present slack one wants to recover analytically.

Proposition O15 (Ratio-box C_{max} floor). The closed form of Proposition O14 depends only on the two dimensionless ratios

$$a := \frac{\gamma_{\phi}}{\Delta E}, \quad b := \frac{\gamma_{\text{mix}}}{\Delta E}.$$

Indeed

$$C_{\text{max}} = F(a, b),$$

where

$$F(a, b) = \frac{1}{2\sqrt{1+4b^2+4ab}} \exp \left[-\frac{a+2b}{\sqrt{1-a^2}} \arctan \left(\frac{\sqrt{1-a^2}}{a+2b} \right) \right].$$

This representation is scale-free: the dependence on ΔE enters only through the ratios (a, b) . On the exported exact bridge, `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_cmax_ratio_floor_source.py` verifies that F is monotone decreasing on the audited ratio box

$$0 \leq a \leq a_{\text{max}}, \quad 0 \leq b \leq b_{\text{max}},$$

with derivative margins

$$\max \partial_a F = -2.44 \times 10^{-1} < 0, \quad \max \partial_b F = -4.88 \times 10^{-1} < 0.$$

Here

$$a_{\text{max}} = \max_D \frac{\gamma_{\phi}(D)}{\Delta E(D)} = 2.06749 \times 10^{-4}, \quad b_{\text{max}} = \max_D \frac{\gamma_{\text{mix}}(D)}{\Delta E(D)} = 2.93097 \times 10^{-1}.$$

Therefore the closed-form ratio floor is

$$C_{\text{max}}^{(\text{ratio,closed}) \text{ floor}} = F(a_{\text{max}}, b_{\text{max}}) = 2.34328 \times 10^{-1}.$$

Subtracting the exported exact-vs-closed residual bound from Proposition O14 gives the rigorous exact-bridge floor

$$C_{\max}(D) \geq C_{\max}^{(\text{ratio,rig}) \text{ floor}} := C_{\max}^{(\text{ratio,closed}) \text{ floor}} - \max_D |C_{\max}^{(\text{export})} - C_{\max}^{(\text{closed})}| = 2.34045 \times 10^{-1}.$$

Corollary O14 (Tightened χ_{eff} floor from the ratio-box $C_{\max}$ floor). Combining Proposition O15 with Proposition O13 yields

$$\chi_{\text{eff}}(D) \geq A_{\chi}^{(\text{sep floor})} C_{\max}^{(\text{ratio,rig}) \text{ floor}} = 2.43752 \times 10^{-7}.$$

The exact exported floors are

$$C_{\max}^{\text{exact floor}} = 2.34255 \times 10^{-1}, \quad \chi_{\text{eff}}^{\text{exact floor}} = 9.03367 \times 10^{-7},$$

so

$$\frac{C_{\max}^{\text{exact floor}}}{C_{\max}^{(\text{ratio,rig}) \text{ floor}}} = 1.000895, \quad \frac{\chi_{\text{eff}}^{\text{exact floor}}}{A_{\chi}^{(\text{sep floor})} C_{\max}^{(\text{ratio,rig}) \text{ floor}}} = 3.706.$$

Thus the observable-chain side is not merely known to admit a floor: after Proposition O15 the remaining gap is essentially only the slack in the amplitude factor $A_{\chi}^{(\text{sep floor})}$, not in the coherence-response floor itself.

Proposition O16 (Amplitude-floor alignment with the mix-rate floor on the exact bridge). On the exported exact bridge,

$$A_{\chi}(D) = \frac{2\gamma_{\text{mix}}(D)}{\gamma_{\text{ref}}(D)} = \frac{2\gamma_{\text{mix}}(D)}{\omega_1(D)^{11}} = \kappa_{\text{env}} \frac{2g_x(D)^2 S_{xx}(D)}{\omega_1(D)^{11}}.$$

`/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_achi_amplitude_source.py` verifies the exact residuals

$$\max_D \left| A_{\chi} - \frac{2\gamma_{\text{mix}}}{\omega_1^{11}} \right| = 1.73 \times 10^{-16}, \quad \max_D \left| A_{\chi} - \kappa_{\text{env}} \frac{2g_x^2 S_{xx}}{\omega_1^{11}} \right| = 1.01 \times 10^{-16}.$$

Moreover the three minima

$$\arg \min_D A_{\chi}(D), \quad \arg \min_D \gamma_{\text{mix}}(D), \quad \arg \min_D \frac{g_x(D)^2 S_{xx}(D)}{\omega_1(D)^{11}}$$

all occur at the same audited knot

$$D = 11.$$

Hence the exact amplitude floor is already determined by the mix-rate floor together with the local first-mode scale at that same knot:

$$A_{\chi}^{\text{exact floor}} = \frac{2\gamma_{\text{mix}}^{\text{floor}}}{\omega_1(D_{\gamma,\min})^{11}} = 1.81567 \times 10^{-6},$$

with

$$\gamma_{\text{mix}}^{\text{floor}} = 5.22995 \times 10^{-4}, \quad \omega_1(D_{\gamma, \text{min}}) = 1.78219.$$

If one instead keeps the same numerator floor but replaces the local denominator by the global first-mode ceiling,

$$A_{\chi}^{(\gamma, \omega_{\text{ceil}})} := \frac{2\gamma_{\text{mix}}^{\text{floor}}}{(\omega_1^{\text{ceil}})^{11}} = 1.10807 \times 10^{-6}, \quad \omega_1^{\text{ceil}} = 1.86402,$$

then the entire remaining amplitude-side slack on the exported bridge is

$$\frac{A_{\chi}^{\text{exact floor}}}{A_{\chi}^{(\gamma, \omega_{\text{ceil}})}} = \left(\frac{\omega_1^{\text{ceil}}}{\omega_1(D_{\gamma, \text{min}})} \right)^{11} = 1.63858.$$

Corollary O15 (Tightened observable floor from the aligned amplitude source).

Combining Proposition O16 with Corollary O14 yields, on the exported exact bridge,

$$\chi_{\text{eff}}(D) \geq A_{\chi}^{\text{exact floor}} C_{\text{max}}^{(\text{ratio}, \text{rig}) \text{ floor}} = 4.24949 \times 10^{-7}.$$

The exact exported floor remains

$$\chi_{\text{eff}}^{\text{exact floor}} = 9.03367 \times 10^{-7},$$

so

$$\frac{\chi_{\text{eff}}^{\text{exact floor}}}{A_{\chi}^{\text{exact floor}} C_{\text{max}}^{(\text{ratio}, \text{rig}) \text{ floor}}} = 2.12582.$$

Thus after Propositions O15–O16 the residual observable-chain slack is no longer spread over both the coherence-response and amplitude factors. The coherence-response floor is already essentially saturated, and the remaining gap is concentrated in proving an a priori floor for γ_{mix} together with the corresponding local first-mode denominator control.

Proposition O17 (Alignment of the γ_{mix} floor with the x -channel overlap floor on the exact bridge). On the exported exact bridge,

$$\gamma_{\text{mix}}(D) = \kappa_{\text{env}} g_x(D)^2 S_{xx}(D).$$

`/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_gamma_mix_floor_alignment.py` verifies the exact reconstruction

$$\max_D |\gamma_{\text{mix}} - \kappa_{\text{env}} g_x^2 S_{xx}| = 4.40 \times 10^{-14}.$$

Moreover the minima of

$$\gamma_{\text{mix}}(D), \quad g_x(D), \quad g_x(D)^2, \quad g_x(D)^2 S_{xx}(D)$$

all occur at the same audited knot

$$D = 11.$$

Hence the exact mix-rate floor is already controlled by the x -channel overlap floor together with the local response factor at that same knot:

$$\gamma_{\text{mix}}^{\text{exact floor}} = \kappa_{\text{env}} (g_x^{\text{floor}})^2 S_{xx}(D_{g_x, \text{min}}) = 5.22995 \times 10^{-4},$$

where

$$g_x^{\text{floor}} = 3.35398 \times 10^{-5}, \quad S_{xx}(D_{g_x, \text{min}}) = 1.07984.$$

If one instead keeps the same overlap floor but replaces the local response factor by the global response minimum,

$$\gamma_{\text{mix}}^{(g_x, S_{xx}^{\text{min}})} := \kappa_{\text{env}} (g_x^{\text{floor}})^2 S_{xx}^{\text{floor}} = 4.91560 \times 10^{-4},$$

with

$$S_{xx}^{\text{floor}} = 1.01493 \quad \text{at } D = 5,$$

then the remaining floor slack is exactly

$$\frac{\gamma_{\text{mix}}^{\text{exact floor}}}{\gamma_{\text{mix}}^{(g_x, S_{xx}^{\text{min}})}} = \frac{S_{xx}(D_{g_x, \text{min}})}{S_{xx}^{\text{floor}}} = 1.06395.$$

Corollary O16 (The remaining γ_{mix} slack is not response-dominated). Proposition O17 shows that the exact exported γ_{mix} floor is already aligned with the x -channel overlap floor, and that replacing the local response factor by the global response minimum changes the floor by only

$$6.39\%.$$

Therefore the next proof-side target should not be read as an independent floor theorem for S_{xx} . On the current exact bridge the response factor is already mild and the dominant unresolved ingredient is the overlap-side floor itself, i.e. an a priori lower bound on g_x (equivalently on $g_x^2 S_{xx}$ up to this small response uplift).

Proposition O18 (Parity-diagonal source for the g_x floor on the exact bridge). Let

$$V_{\text{par}}^{(\delta V)}(D) = \begin{pmatrix} V_{11}(D) & V_{12}(D) \\ V_{12}(D) & V_{22}(D) \end{pmatrix}$$

be the parity-basis projection of the fluctuation operator $\delta V = V_{\text{eff}} - \bar{V}(r)$ onto the lowest two parity modes, and let

$$U_{LR} = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$$

be the localized-basis rotation. Then

$$V_{\text{loc}}^{(\delta V)}(D) = U_{LR}^\top V_{\text{par}}^{(\delta V)}(D) U_{LR} = \begin{pmatrix} \frac{1}{2}(V_{11} + 2V_{12} + V_{22}) & \frac{1}{2}(V_{11} - V_{22}) \\ \frac{1}{2}(V_{11} - V_{22}) & \frac{1}{2}(V_{11} - 2V_{12} + V_{22}) \end{pmatrix}.$$

Hence the localized-channel couplings are exact reparameterizations of the same projected fluctuation matrix:

$$g_x(D) = \left| (V_{\text{loc}}^{(\delta V)}(D))_{LR} \right| = \frac{1}{2} |V_{11}(D) - V_{22}(D)|,$$

$$g_z(D) = \frac{1}{2} \left| (V_{\text{loc}}^{(\delta V)}(D))_{LL} - (V_{\text{loc}}^{(\delta V)}(D))_{RR} \right| = |V_{12}(D)|.$$

`/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_gx_floor_source.py` verifies on the exported exact bridge that

$$\max_D \left| g_x - \left| V_{LR}^{(\delta V)} \right| \right| = 0,$$

$$\max_D \left| g_x - \frac{1}{2} |V_{11} - V_{22}| \right| = 1.00 \times 10^{-16}, \quad \max_D |g_z - |V_{12}|| = 2.18 \times 10^{-18}.$$

Moreover the minima of

$$g_x(D), \quad \frac{1}{2} |V_{11}(D) - V_{22}(D)|, \quad \gamma_{\text{mix}}(D)$$

all occur at the same audited knot

$$D = 11,$$

where

$$g_x^{\text{floor}} = 3.35398 \times 10^{-5}, \quad |V_{11} - V_{22}|_{\text{floor}} = 6.70795 \times 10^{-5}.$$

Thus the remaining exact-bridge floor question for g_x is already equivalent to a static parity-diagonal contrast floor for the projected fluctuation operator δV .

Corollary O17 (The remaining g_x floor gap is static rather than bath-side). Propositions O17–O18 reduce the last open-system overlap-side floor statement to the chain

$$\gamma_{\text{mix}}^{\text{exact floor}} = \kappa_{\text{env}} (g_x^{\text{floor}})^2 S_{xx}(D_{g_x, \text{min}}), \quad g_x^{\text{floor}} = \frac{1}{2} |V_{11} - V_{22}|_{\text{floor}}.$$

Therefore the unresolved proof-side target is no longer naturally read as a bath-response theorem for S_{xx} or as an additional localized-basis theorem for $V_{LR}^{(\delta V)}$. On the current exact bridge it is enough to produce an a priori lower bound on the parity-diagonal contrast

$$|V_{11}(D) - V_{22}(D)|$$

for the projected fluctuation operator δV over the audited knot window.

Proposition O19 (One-center pair plus mirror remainder decomposition for the parity contrast). Write the two-center shifted potential as

$$U_2(\rho, z; D) = U_+(\rho, z; D) + U_-(\rho, z; D) + U_\times(\rho, z; D),$$

where

$$U_\pm = m_0^2((1 + s_\pm)^2 - 1)c_\xi \frac{\ell_\pm}{1 + s_\pm}, \quad c_\xi := 1 - 6\xi,$$

with

$$s_\pm(\rho, z; D) = \frac{a}{\sqrt{\rho^2 + (z \mp D/2)^2 + \varepsilon^2}}, \quad \ell_\pm(\rho, z; D) = -3a\varepsilon^2(\rho^2 + (z \mp D/2)^2 + \varepsilon^2)^{-5/2},$$

and

$$U_\times = 2m_0^2 s_+ s_- + c_\xi \left[\frac{\ell_+ + \ell_-}{1 + s_+ + s_-} - \frac{\ell_+}{1 + s_+} - \frac{\ell_-}{1 + s_-} \right].$$

Let

$$\delta V_j(\rho, z; D) := U_j(\rho, z; D) - \bar{U}_j(r; D), \quad r = \sqrt{\rho^2 + z^2},$$

with the same origin-centered spherical-average convention as the exact bridge, so that

$$\delta V = \delta V_+ + \delta V_- + \delta V_\times.$$

For the exact contrast density

$$\mathcal{C}(\rho, z; D) := 2\pi\rho(\psi_1(\rho, z; D)^2 - \psi_2(\rho, z; D)^2),$$

the parity-diagonal contrast decomposes exactly as

$$V_{11}(D) - V_{22}(D) = \int \mathcal{C} \delta V_+ + \int \mathcal{C} \delta V_- + \int \mathcal{C} \delta V_\times.$$

On the symmetry-adapted continuum pair one has

$$\int \mathcal{C} \delta V_+ = \int \mathcal{C} \delta V_-,$$

so the exact contrast may be read as a one-center pair plus a mirror/cross remainder:

$$V_{11}(D) - V_{22}(D) = 2I_{\text{self}}(D) + I_{\times}(D).$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_self_mirror_source.py` closes the additive split to machine precision,

$$\max_D \|U_2 - U_+ - U_- - U_{\times}\|_{\infty} = 8.88 \times 10^{-15},$$

$$\max_D \|\delta V - \delta V_+ - \delta V_- - \delta V_{\times}\|_{\infty} = 7.11 \times 10^{-15},$$

and reconstructs the projected contrast with

$$\max_D |(V_{11} - V_{22}) - (I_+ + I_- + I_{\times})| = 3.79 \times 10^{-18}.$$

Across the full audited window $D = 4, \dots, 20$, the one-center pair carries between

$$97.74\% \quad \text{and} \quad 99.833\%$$

of the exact parity contrast in absolute value, while the mirror/cross remainder contributes only

$$0.167\% \quad \text{to} \quad 2.257\%.$$

At the exact floor knot $D = 11$,

$$V_{11} - V_{22} = 6.70795 \times 10^{-5},$$

$$I_+ + I_- = 6.69675 \times 10^{-5}, \quad I_{\times} = 1.12022 \times 10^{-7},$$

so the exact floor is already 99.833% saturated by the one-center pair, with the remainder reduced to a 0.167% correction. The cross piece itself splits as

$$I_{\times} = I_{\times}^{\text{mass}} + I_{\times}^{\text{deriv}},$$

where at $D = 11$

$$I_{\times}^{\text{mass}} = 3.85528 \times 10^{-7}, \quad I_{\times}^{\text{deriv}} = -2.73506 \times 10^{-7},$$

showing that the small remainder comes from a partial cancellation inside the mirror term rather than from the failure of the one-center anchor.

Corollary O18 (The remaining $|V_{11} - V_{22}|$ gap is now a one-center floor problem with a small static remainder). Propositions O18–O19 sharpen the last overlap-side target from a generic parity-contrast floor to a much narrower statement. On the exact bridge,

$$|V_{11}(D) - V_{22}(D)|$$

is already overwhelmingly carried by the one-center pair $2I_{\text{self}}(D)$, while the mirror/cross remainder $I_{\times}(D)$ keeps the same sign as the exact contrast and remains below 2.26% of it on every audited knot. Therefore the next proof-side task should not be framed as a new bath-side or barrier-center theorem. A theorem-level lower bound on the one-center pair $2I_{\text{self}}(D)$, together with a coarse control of the small static remainder $I_{\times}(D)$, is enough to close the present exact-bridge gap for

$$|V_{11}(D) - V_{22}(D)|.$$

Proposition O20 (Core-window lower bound for the one-center pair on the exact bridge). For $R > 0$, let

$$B_{\pm}(R; D) := \{(\rho, z) : \rho^2 + (z \mp D/2)^2 \leq R^2\}, \quad B_{\text{core}}(R; D) := B_+(R; D) \cup B_-(R; D),$$

and define

$$I_{\text{self, pair}}(D) := I_+(D) + I_-(D),$$

$$I_{\text{self, core}}^{(R)}(D) := \int_{B_{\text{core}}(R; D)} \mathcal{C}(\rho, z; D) (\delta V_+(\rho, z; D) + \delta V_-(\rho, z; D)) d\rho dz,$$

$$I_{\text{self, tail}}^{(R)}(D) := I_{\text{self, pair}}(D) - I_{\text{self, core}}^{(R)}(D).$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_self_floor_source.py` scans

$$R \in [1.20, 1.40]$$

in steps of 0.01 and identifies the optimal safe radius

$$R_* = 1.35,$$

for which every audited knot $D = 4, \dots, 20$ satisfies

$$\text{sgn } I_{\text{self, core}}^{(R_*)}(D) = \text{sgn } I_{\text{self, pair}}(D), \quad \text{sgn } I_{\text{self, tail}}^{(R_*)}(D) = -\text{sgn } I_{\text{self, pair}}(D).$$

Hence

$$|I_{\text{self, pair}}(D)| = |I_{\text{self, core}}^{(R_*)}(D)| - |I_{\text{self, tail}}^{(R_*)}(D)|.$$

Moreover the same audit gives the uniform tail-to-core control

$$|I_{\text{self, tail}}^{(R_*)}(D)| \leq \eta_* |I_{\text{self, core}}^{(R_*)}(D)|, \quad \eta_* = 0.573731 \dots$$

so that

$$|I_{\text{self,pair}}(D)| \geq \lambda_* |I_{\text{self,core}}^{(R_*)}(D)|, \quad \lambda_* := 1 - \eta_* = 0.426269 \dots$$

on the full audited window. Since the core-window floor is

$$\min_D |I_{\text{self,core}}^{(R_*)}(D)| = 1.33812 \times 10^{-4},$$

this already yields the nontrivial lower bound

$$|I_{\text{self,pair}}(D)| \geq 5.70400 \times 10^{-5}.$$

The exact one-center-pair floor is

$$\min_D |I_{\text{self,pair}}(D)| = 6.69675 \times 10^{-5} \quad \text{at } D = 11,$$

so the optimized core-window lower bound is only a factor

$$1.174$$

below the exact minimum.

Corollary O19 (The parity-contrast floor now follows from a one-center core floor plus a tiny same-sign remainder). Combine Proposition O20 with Proposition O19 and the sign information from the same decomposition audit. On the exact bridge one has

$$\text{sgn } I_{\times}(D) = \text{sgn}(V_{11}(D) - V_{22}(D))$$

for every audited knot, together with

$$\max_D \frac{|I_{\times}(D)|}{|V_{11}(D) - V_{22}(D)|} = 2.2568\%.$$

Therefore

$$|V_{11}(D) - V_{22}(D)| \geq |I_{\text{self,pair}}(D)| \geq \lambda_* |I_{\text{self,core}}^{(R_*)}(D)| \geq 5.70400 \times 10^{-5}$$

already on the exact bridge, without invoking any barrier-center WKB ansatz. Since

$$\min_D |V_{11}(D) - V_{22}(D)| = 6.70795 \times 10^{-5},$$

the remaining proof-side gap has become very narrow: produce an a priori lower bound on the one-center core contribution $I_{\text{self,core}}^{(R_*)}$, and retain only a coarse signed control of the small mirror remainder $I_{\times}$.

Proposition O21 (Compact box-anchor lower bound for the one-center core on the exact bridge). Keep the same optimal core radius $R_* = 1.35$, and define the compact local box anchor

$$B_{\text{box}}(\rho_{\#}, z_{\#}; D) := \{(\rho, z) : \rho \leq \rho_{\#}, |z \mp D/2| \leq z_{\#}\} \cap B_{\text{core}}(R_*; D),$$

with the canonical audited choice

$$\rho_{\#} = 0.70, \quad z_{\#} = 1.21.$$

Let

$$I_{\text{self,anchor}}^{(\rho_{\#}, z_{\#})}(D) := \int_{B_{\text{box}}(\rho_{\#}, z_{\#}; D)} \mathcal{C}(\rho, z; D) (\delta V_+(\rho, z; D) + \delta V_-(\rho, z; D)) d\rho dz,$$

and

$$I_{\text{self,shell}}^{(\rho_{\#}, z_{\#})}(D) := I_{\text{self,core}}^{(R_*)}(D) - I_{\text{self,anchor}}^{(\rho_{\#}, z_{\#})}(D).$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_self_core_anchor_source.py` scans

$$\rho_{\#} \in [0.55, 0.85], \quad z_{\#} \in [1.05, 1.30],$$

in steps of 0.01 and finds a flat optimal plateau at $z_{\#} = 1.21$ with $\rho_{\#} \in [0.69, 0.74]$. On the exact bridge the canonical box $(\rho_{\#}, z_{\#}) = (0.70, 1.21)$ already satisfies

$$\text{sgn } I_{\text{self,anchor}}^{(\rho_{\#}, z_{\#})}(D) = \text{sgn } I_{\text{self,core}}^{(R_*)}(D), \quad \text{sgn } I_{\text{self,shell}}^{(\rho_{\#}, z_{\#})}(D) = \text{sgn } I_{\text{self,core}}^{(R_*)}(D)$$

for every audited knot $D = 4, \dots, 20$. Therefore

$$|I_{\text{self,core}}^{(R_*)}(D)| \geq |I_{\text{self,anchor}}^{(\rho_{\#}, z_{\#})}(D)|.$$

The corresponding compact-anchor floor is

$$\min_D |I_{\text{self,anchor}}^{(\rho_{\#}, z_{\#})}(D)| = 1.32321 \times 10^{-4},$$

whereas the exact core-window floor is

$$\min_D |I_{\text{self,core}}^{(R_*)}(D)| = 1.33812 \times 10^{-4}.$$

So the compact box anchor already saturates the exact core-window minimum to within a factor

$$1.01127.$$

Corollary O20 (The one-center core floor is already carried by a compact local anchor). Combining Proposition O21 with Proposition O20 gives

$$|I_{\text{self,pair}}(D)| \geq \lambda_* |I_{\text{self,core}}^{(R_*)}(D)| \geq \lambda_* |I_{\text{self,anchor}}^{(\rho_\#, z_\#)}(D)| \geq 5.64043 \times 10^{-5}$$

on the full audited window. The exact one-center-pair and total parity-contrast floors remain

$$\min_D |I_{\text{self,pair}}(D)| = 6.69675 \times 10^{-5}, \quad \min_D |V_{11}(D) - V_{22}(D)| = 6.70795 \times 10^{-5},$$

so the remaining gap from the compact anchor to the exact exported minima is now only a factor 1.187 to 1.189. This further tightens the proof target: it is enough to prove a lower bound on the compact one-center anchor $I_{\text{self,anchor}}^{(\rho_\#, z_\#)}$, while the rest of $I_{\text{self,core}}^{(R_*)}$ can remain a same-sign completion and $I_\times$ only needs coarse signed remainder control.

Proposition O22 (The compact anchor is already controlled by a same-center static source plus a local mirror loss). Keep the canonical compact anchor

$$B_{\text{anchor}}(D) := B_{\text{box}}(0.70, 1.21; D).$$

Split the anchor integral into its same-center and mirror pieces:

$$I_{\text{self,anchor}}(D) = I_{\text{diag,anchor}}(D) + I_{\text{mir,anchor}}(D),$$

where

$$I_{\text{diag,anchor}}(D) := \int_{B_+(D) \cap B_{\text{anchor}}(D)} \mathcal{C}(\rho, z; D) \delta V_+(\rho, z; D) d\rho dz + \int_{B_-(D) \cap B_{\text{anchor}}(D)} \mathcal{C}(\rho, z; D) \delta V_-(\rho, z; D) d\rho dz,$$

$$I_{\text{mir,anchor}}(D) := \int_{B_+(D) \cap B_{\text{anchor}}(D)} \mathcal{C}(\rho, z; D) \delta V_-(\rho, z; D) d\rho dz + \int_{B_-(D) \cap B_{\text{anchor}}(D)} \mathcal{C}(\rho, z; D) \delta V_+(\rho, z; D) d\rho dz.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_anchor_static_source.py` scans all same-center sub-boxes

$$\rho_{\text{box}} \in [0.40, 0.70], \quad z_{\text{box}} \in [0.80, 1.21],$$

in steps of 0.01, requiring the same-center box and the remaining same-center shell to keep the sign of $I_{\text{self,anchor}}$ on every audited knot $D = 4, \dots, 20$. Under this criterion the optimal box is the full canonical anchor itself:

$$(\rho_{\text{box}}^{\text{opt}}, z_{\text{box}}^{\text{opt}}) = (0.69, 1.21) \sim (0.70, 1.21).$$

Moreover, on the full audited window one has

$$\text{sgn } I_{\text{diag,anchor}}(D) = \text{sgn } I_{\text{self,anchor}}(D),$$

and a uniform local mirror ratio

$$\mu_{\text{anc}} := \max_D \frac{|I_{\text{mir},\text{anchor}}(D)|}{|I_{\text{diag},\text{anchor}}(D)|} = 0.415143 \dots < 1.$$

Therefore

$$|I_{\text{self},\text{anchor}}(D)| \geq (1 - \mu_{\text{anc}}) |I_{\text{diag},\text{anchor}}(D)|.$$

The same-center static floor is

$$\min_D |I_{\text{diag},\text{anchor}}(D)| = 1.63757 \times 10^{-4} \quad \text{at } D = 11,$$

so the induced anchor lower bound is

$$|I_{\text{self},\text{anchor}}(D)| \geq (1 - \mu_{\text{anc}}) |I_{\text{diag},\text{anchor}}(D)| \geq 9.57744 \times 10^{-5}.$$

Compared against the exact compact-anchor floor

$$\min_D |I_{\text{self},\text{anchor}}(D)| = 1.32321 \times 10^{-4},$$

this static-source bound is only a factor

$$1.38159$$

below the exact minimum.

Corollary O21 (The remaining overlap-side gap has been reduced to a same-center static anchor source). Combining Proposition O22 with Corollary O20 gives

$$|I_{\text{self},\text{pair}}(D)| \geq \lambda_* |I_{\text{self},\text{anchor}}(D)| \geq \lambda_* (1 - \mu_{\text{anc}}) |I_{\text{diag},\text{anchor}}(D)| \geq 4.08256 \times 10^{-5}.$$

The exact one-center-pair and total parity-contrast floors remain

$$\min_D |I_{\text{self},\text{pair}}(D)| = 6.69675 \times 10^{-5}, \quad \min_D |V_{11}(D) - V_{22}(D)| = 6.70795 \times 10^{-5},$$

so the residual gap from the audited exact minima is now only a factor

$$1.640 \text{ to } 1.643.$$

This isolates the true proof-side target even further: the overlap-side problem is no longer to control barrier-center WKB tails, nor even the full compact anchor, but to produce an a priori lower bound on the same-center static source $I_{\text{diag},\text{anchor}}$. The compact anchor itself is already obtained from that source by a local mirror correction with uniform ratio $\mu_{\text{anc}} < 0.42$, while $I_{\times}$ stays only a coarse same-sign remainder in the total parity contrast.

Proposition O23 (Strict same-center subboxes do not improve the static anchor floor). To test whether the proof target can be localized further, define a strict same-center subbox inside the canonical anchor by

$$B_{\text{diag,box}}(\rho_{\square}, \zeta_{\square}; D) := \{(\rho, z) : \rho \leq \rho_{\square}, |z \mp D/2| \leq \zeta_{\square}\} \subsetneq B_{\text{anchor}}(D),$$

with at least one strict inequality

$$\rho_{\square} < 0.70 \quad \text{or} \quad \zeta_{\square} < 1.21.$$

Let

$$I_{\text{diag,anchor}}(D) = I_{\text{diag,box}}^{(\rho_{\square}, \zeta_{\square})}(D) + I_{\text{diag,shell}}^{(\rho_{\square}, \zeta_{\square})}(D).$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_box_source.py` scans

$$\rho_{\square} \in [0.40, 0.70], \quad \zeta_{\square} \in [0.80, 1.21],$$

in steps of 0.01, excluding the trivial full-anchor choice. On every audited knot $D = 4, \dots, 20$, the optimal strict plateau occurs at

$$\rho_{\square} \in [0.63, 0.68], \quad \zeta_{\square} = 1.21,$$

and the rounded representative choice

$$(\rho_{\square}, \zeta_{\square}) = (0.65, 1.21)$$

already satisfies

$$\text{sgn } I_{\text{diag,box}}^{(\rho_{\square}, \zeta_{\square})}(D) = \text{sgn } I_{\text{diag,anchor}}(D), \quad \text{sgn } I_{\text{diag,shell}}^{(\rho_{\square}, \zeta_{\square})}(D) = \text{sgn } I_{\text{diag,anchor}}(D)$$

for every audited knot. Therefore

$$|I_{\text{diag,anchor}}(D)| \geq |I_{\text{diag,box}}^{(\rho_{\square}, \zeta_{\square})}(D)|.$$

The resulting strict-box floor is

$$\min_D \left| I_{\text{diag,box}}^{(0.65, 1.21)}(D) \right| = 1.48301 \times 10^{-4},$$

whereas the exact same-center-anchor floor is

$$\min_D |I_{\text{diag,anchor}}(D)| = 1.63757 \times 10^{-4}.$$

So even after forcing a genuine strict localization, the loss is only a factor

$$1.10422.$$

Corollary O22 (The proof target has stabilized at the full same-center anchor scale).

Proposition O23 shows that no genuinely smaller same-center box inside the canonical anchor improves the static-source floor on the audited bridge. The best strict local box already carries

$$\frac{1}{1.10422} = 90.56\%$$

of the exact same-center-anchor floor. Thus the next theorem step should not be to keep shrinking the support, but to lower-bound $I_{\text{diag,anchor}}$ itself. In other words, the overlap-side mainline has stabilized:

$$I_{\text{diag,anchor}}$$

is now the natural one-center static source to target, while $I_{\text{mir,anchor}}$ remains only a coarse local mirror correction and $I_{\times}$ remains only a coarse same-sign remainder in the full parity contrast.

Proposition O24 (Oriented moment reduction of the same-center static anchor).

Keep the canonical same-center static anchor $I_{\text{diag,anchor}}(D)$ and define the oriented anchor density

$$q_{\text{diag}}(\rho, z; D) := \text{sgn}(I_{\text{diag,anchor}}(D)) \mathcal{C}(\rho, z; D) \left(\delta V_+(\rho, z; D) \mathbf{1}_{B_+(D) \cap B_{\text{anchor}}(D)} + \delta V_-(\rho, z; D) \mathbf{1}_{B_-(D) \cap B_{\text{anchor}}(D)} \right).$$

Let

$$P_{\text{diag}}(D) := \int [q_{\text{diag}}(\rho, z; D)]_+ d\rho dz, \quad N_{\text{diag}}(D) := \int [q_{\text{diag}}(\rho, z; D)]_- d\rho dz,$$

so that

$$|I_{\text{diag,anchor}}(D)| = P_{\text{diag}}(D) - N_{\text{diag}}(D).$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_moment_source.py` closes this split to machine precision,

$$\max_D \left| |I_{\text{diag,anchor}}(D)| - (P_{\text{diag}}(D) - N_{\text{diag}}(D)) \right| = 6.51 \times 10^{-19},$$

and gives the uniform moment ratio

$$\nu_{\text{diag}}^{(\pm)} := \max_D \frac{N_{\text{diag}}(D)}{P_{\text{diag}}(D)} = 0.911532 \dots \quad \text{at } D = 11.$$

Therefore

$$|I_{\text{diag,anchor}}(D)| \geq (1 - \nu_{\text{diag}}^{(\pm)}) P_{\text{diag}}(D).$$

The positive-moment floor itself is

$$\min_D P_{\text{diag}}(D) = 1.17067 \times 10^{-3} \quad \text{at } D = 20,$$

so the induced intrinsic static lower bound is

$$|I_{\text{diag,anchor}}(D)| \geq (1 - \nu_{\text{diag}}^{(\pm)}) \min_D P_{\text{diag}}(D) \geq 1.03567 \times 10^{-4}.$$

Compared against the exact same-center-anchor floor

$$\min_D |I_{\text{diag,anchor}}(D)| = 1.65556 \times 10^{-4},$$

this oriented-moment bound is only a factor

$$1.59854$$

below the exact minimum, while remaining more intrinsic than any support-based subbox choice.

Corollary O23 (The remaining anchor-side gap is now cancellation control, not support selection). Proposition O24 gives a direct lower bound on the full same-center anchor source without introducing any new support cutoff:

$$|I_{\text{diag,anchor}}(D)| \geq 1.03567 \times 10^{-4}.$$

This is stronger than the previous mass-only route

$$8.17293 \times 10^{-5}$$

by a factor

$$1.26719,$$

and it is also slightly stronger than the induced compact-anchor lower bound

$$9.57744 \times 10^{-5}$$

coming from Proposition O22, by a factor

$$1.08136.$$

So the proof-side interpretation is now sharper still: the obstacle is no longer where to localize the source. The support has already stabilized at the full same-center anchor scale, and the remaining loss is dominated by the positive/negative cancellation inside that anchor.

Proposition O25 (Shell-isolated cancellation control for the same-center anchor).

To separate genuine interior cancellation from mere support leakage, fix a same-center axial slab cut

$$\zeta_{\text{cut}} < 1.21$$

and write the canonical anchor as

$$B_{\text{anchor}}(D) = B_{\text{in}}^{(\zeta_{\text{cut}})}(D) \sqcup B_{\text{out}}^{(\zeta_{\text{cut}})}(D),$$

where

$$B_{\text{in}}^{(\zeta_{\text{cut}})}(D) := \{(\rho, z) \in B_{\text{anchor}}(D) : |z \mp D/2| \leq \zeta_{\text{cut}}\}.$$

Define the corresponding inner and outer oriented moments

$$\begin{aligned} P_{\text{in}}^{(\zeta_{\text{cut}})}(D) &:= \int_{B_{\text{in}}^{(\zeta_{\text{cut}})}(D)} [q_{\text{diag}}]_+ d\rho dz, & N_{\text{in}}^{(\zeta_{\text{cut}})}(D) &:= \int_{B_{\text{in}}^{(\zeta_{\text{cut}})}(D)} [q_{\text{diag}}]_- d\rho dz, \\ P_{\text{out}}^{(\zeta_{\text{cut}})}(D) &:= \int_{B_{\text{out}}^{(\zeta_{\text{cut}})}(D)} [q_{\text{diag}}]_+ d\rho dz, & N_{\text{out}}^{(\zeta_{\text{cut}})}(D) &:= \int_{B_{\text{out}}^{(\zeta_{\text{cut}})}(D)} [q_{\text{diag}}]_- d\rho dz. \end{aligned}$$

Then

$$\nu_{\text{diag}}^{(\pm)}(D) := \frac{N_{\text{diag}}(D)}{P_{\text{diag}}(D)} \leq \frac{N_{\text{in}}^{(\zeta_{\text{cut}})}(D)}{P_{\text{in}}^{(\zeta_{\text{cut}})}(D)} + \frac{N_{\text{out}}^{(\zeta_{\text{cut}})}(D)}{P_{\text{in}}^{(\zeta_{\text{cut}})}(D)},$$

because $N_{\text{diag}} = N_{\text{in}} + N_{\text{out}}$ and $P_{\text{diag}} \geq P_{\text{in}}$. The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_cancellation_source.py` scans the nontrivial slab family

$$\zeta_{\text{cut}} \in [0.80, 1.20]$$

and finds the optimal shell-isolating choice

$$\zeta_{\text{cut}}^* = 1.20.$$

On the full audited window this gives

$$\min_D P_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D) = 1.17067 \times 10^{-3},$$

$$\max_D \frac{N_{\text{out}}^{(\zeta_{\text{cut}}^*)}(D)}{P_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D)} = 8.68193 \times 10^{-3},$$

$$\max_D \frac{N_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D)}{P_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D)} = 0.912743 \dots$$

and therefore the full cancellation ratio obeys the explicit audited bound

$$\nu_{\text{diag}}^{(\pm)}(D) \leq 0.912743 + 8.68193 \times 10^{-3} \leq 0.921425$$

for every audited knot $D = 4, \dots, 20$. The shell leakage is thus below one percent at the level of the normalized cancellation budget.

Corollary O24 (The anchor-side cancellation problem is interior and mass-dominated). The same audit further decomposes the oriented anchor density as

$$q_{\text{diag}} = q_{\text{mass}} + q_{\text{der}},$$

where q_{mass} and q_{der} are the same-center mass and derivative contributions built from

$$U_{\pm}^{\text{mass}} = m_0^2(\Omega_{\pm}^2 - 1), \quad U_{\pm}^{\text{der}} = (1 - 6\xi) \frac{\nabla^2 \Omega_{\pm}}{\Omega_{\pm}},$$

after the same origin-centered spherical-average subtraction. Writing the corresponding positive/negative moments as

$$P_{\text{mass}}, N_{\text{mass}}, P_{\text{der}}, N_{\text{der}},$$

the audited exact bridge satisfies

$$\min_D \frac{P_{\text{mass}}(D)}{P_{\text{diag}}(D)} = 1.00507, \quad \max_D \frac{P_{\text{der}}(D)}{P_{\text{diag}}(D)} = 8.53082 \times 10^{-2},$$

$$\min_D \frac{N_{\text{mass}}(D)}{N_{\text{diag}}(D)} = 0.712802, \quad \max_D \frac{N_{\text{der}}(D)}{N_{\text{diag}}(D)} = 0.570882.$$

So the shell part is negligible, the positive carrier is already overwhelmingly same-center mass-sector, and the remaining adverse budget stays interior to the one-center anchor rather than coming from support selection or mirror leakage. The proof-side target has therefore narrowed again: the remaining overlap-side gap is best read as a one-center interior cancellation bound, not as a localization problem.

Proposition O26 (Interior mass-sector cancellation controls the floor knot). Keep the same canonical same-center anchor and the shell-isolating cut

$$\zeta_{\text{cut}}^* = 1.20.$$

On the resulting inner slab $B_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D)$, define the oriented same-center mass contribution

$$I_{\text{mass,in}}(D) := \text{sgn}(I_{\text{diag,anchor}}(D)) \int_{B_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D)} q_{\text{mass}}(\rho, z; D) d\rho dz,$$

with corresponding positive/negative moments

$$P_{\text{mass,in}}(D) := \int_{B_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D)} [q_{\text{mass}}]_+ d\rho dz, \quad N_{\text{mass,in}}(D) := \int_{B_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D)} [q_{\text{mass}}]_- d\rho dz.$$

Then

$$I_{\text{mass,in}}(D) = P_{\text{mass,in}}(D) - N_{\text{mass,in}}(D) = (1 - \nu_{\text{diag,int}}^{(\pm,m)}(D)) P_{\text{mass,in}}(D),$$

where

$$\nu_{\text{diag,int}}^{(\pm,m)}(D) := \frac{N_{\text{mass,in}}(D)}{P_{\text{mass,in}}(D)}.$$

Writing the full oriented anchor as

$$|I_{\text{diag,anchor}}(D)| = I_{\text{mass,in}}(D) + I_{\text{der,in}}(D) + I_{\text{out}}(D),$$

with $I_{\text{der,in}}$ the derivative-sector correction on the same inner slab and I_{out} the residual oriented shell term, the audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_mass_cancellation_source.py` shows that the same shell-isolating cut remains optimal for the mass-net floor:

$$\zeta_{\text{cut}}^* = 1.20.$$

At that cut the exact bridge satisfies

$$I_{\text{mass,in}}(D) \geq 0 \quad \text{for all audited } D,$$

and the extrema align on the same knot:

$$\arg \min_D |I_{\text{diag,anchor}}(D)| = \arg \min_D I_{\text{mass,in}}(D) = \arg \max_D \nu_{\text{diag,int}}^{(\pm,m)}(D) = 11.$$

Numerically,

$$\min_D |I_{\text{diag,anchor}}(D)| = 1.65556 \times 10^{-4},$$

$$\min_D I_{\text{mass,in}}(D) = 7.92440 \times 10^{-5}, \quad \max_D \nu_{\text{diag,int}}^{(\pm,m)}(D) = 0.957812.$$

At the floor knot $D = 11$, one furthermore finds

$$I_{\text{der,in}}(11) = 8.38290 \times 10^{-5} > 0, \quad I_{\text{out}}(11) = 2.48275 \times 10^{-6} > 0,$$

so the derivative and shell pieces are non-adverse there and the exact floor is obtained by adding a positive derivative uplift and a tiny positive shell correction to the interior mass-sector net.

Corollary O25 (The adverse cancellation is set by the interior one-center mass sector). Proposition O26 isolates the last overlap-side obstruction more sharply than Proposition O25. The floor knot is not selected by shell leakage, by mirror leakage, or by any smaller support cutoff. It is selected by the same interior knot that maximizes the one-center mass-sector cancellation ratio $\nu_{\text{diag,int}}^{(\pm,m)}$, while the derivative and shell terms are already non-adverse there. In particular, the exact floor lies only a factor

$$\frac{1.65556 \times 10^{-4}}{7.92440 \times 10^{-5}} = 2.08919$$

above the interior mass-sector net, and the gap is accounted for by a positive derivative uplift at the very same knot rather than by any new adverse mechanism. So the proof-side target has narrowed again: the remaining overlap-side gap is best read as an interior one-center mass-sector cancellation problem, with derivative and shell pieces retained only as signed local corrections.

Proposition O27 (The positive mass-carrier floor survives a strict axial trim). Keep the same inner slab

$$B_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D), \quad \zeta_{\text{cut}}^* = 1.20,$$

and define for any same-center subbox

$$B_{\text{mass,box}}^{(\rho_{\square}, \zeta_{\square})}(D) := \{(\rho, z) \in B_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D) : \rho \leq \rho_{\square}, |z \mp D/2| \leq \zeta_{\square}\}$$

the positive mass-carrier moment

$$P_{\text{mass,box}}^{(\rho_{\square}, \zeta_{\square})}(D) := \int_{B_{\text{mass,box}}^{(\rho_{\square}, \zeta_{\square})}(D)} [q_{\text{mass}}]_+ d\rho dz.$$

Since $[q_{\text{mass}}]_+ \geq 0$, monotonicity gives the exact lower bound

$$P_{\text{mass,in}}(D) \geq P_{\text{mass,box}}^{(\rho_{\square}, \zeta_{\square})}(D)$$

for every such subbox. The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_mass_positive_source.py` scans

$$\rho_{\square} \in [0.35, 0.70], \quad \zeta_{\square} \in [0.70, 1.20],$$

and finds that the optimal strict positive-carrier plateau already sits on the near-full axial trims

$$\rho_{\square} \in [0.69, 0.70], \quad \zeta_{\square} \in [1.10, 1.19].$$

Using the canonical strict representative

$$(\rho_{\square}, \zeta_{\square}) = (0.70, 1.19),$$

one gets

$$\min_D P_{\text{mass,box}}^{(0.70,1.19)}(D) = \min_D P_{\text{mass,in}}(D) = 1.25520 \times 10^{-3}$$

at $D = 20$. Moreover, on the full audited window

$$\min_D \frac{P_{\text{mass,box}}^{(0.70,1.19)}(D)}{P_{\text{mass,in}}(D)} = 0.999825$$

at $D = 11$. So the positive mass carrier is already fully stabilized on a strict axial trim of the canonical one-center slab.

Corollary O26 (The positive-carrier side is no longer the blocker). Propositions O26 and O27 now split the interior mass-sector problem cleanly in two:

$$I_{\text{mass,in}} = (1 - \nu_{\text{diag,int}}^{(\pm,m)}) P_{\text{mass,in}}, \quad P_{\text{mass,in}} \geq P_{\text{mass,box}}^{(0.70,1.19)}.$$

The positive carrier floor is therefore already stable under a strict one-center axial trim, with no loss at the audited floor knot and at most

$$1 - 0.999825 = 1.75 \times 10^{-4}$$

relative loss anywhere on the audited window. The remaining overlap-side gap is thus no longer a positive-carrier problem; it is the interior one-center mass-sector cancellation ratio itself.

Proposition O28 (The adverse mass carrier localizes to an off-axis one-center annulus). Keep the same canonical inner slab

$$B_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D), \quad \zeta_{\text{cut}}^* = 1.20,$$

and define for any annular trim

$$A_{\text{neg}}^{(\rho_{\text{min}}, \zeta_{\square})}(D) := \{(\rho, z) \in B_{\text{in}}^{(\zeta_{\text{cut}}^*)}(D) : \rho_{\text{min}} \leq \rho \leq 0.70, |z \mp D/2| \leq \zeta_{\square}\}$$

the off-axis adverse-carrier moment

$$N_{\text{mass,ann}}^{(\rho_{\text{min}}, \zeta_{\square})}(D) := \int_{A_{\text{neg}}^{(\rho_{\text{min}}, \zeta_{\square})}(D)} [q_{\text{mass}}]_- d\rho dz.$$

Since $[q_{\text{mass}}]_- \geq 0$, one has the exact split

$$N_{\text{mass,in}}(D) = N_{\text{mass,ann}}^{(\rho_{\text{min}}, \zeta_{\square})}(D) + N_{\text{mass,tail}}^{(\rho_{\text{min}}, \zeta_{\square})}(D),$$

with

$$N_{\text{mass,tail}}^{(\rho_{\text{min}}, \zeta_{\square})}(D) \geq 0.$$

Therefore

$$\nu_{\text{diag,int}}^{(\pm,m)}(D) = \frac{N_{\text{mass,ann}}^{(\rho_{\min},\zeta_{\square})}(D)}{P_{\text{mass,in}}(D)} + \frac{N_{\text{mass,tail}}^{(\rho_{\min},\zeta_{\square})}(D)}{P_{\text{mass,in}}(D)}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_mass_negative_source.py` scans

$$\rho_{\min} \in [0.03, 0.30], \quad \zeta_{\square} \in [1.10, 1.20].$$

If one allows arbitrary strict trims, the optimum is a near-full slab. Constraining instead to genuinely off-axis annuli with $\rho_{\min} \geq 0.12$, the best plateau sits at

$$\rho_{\min} \in [0.12, 0.15], \quad \zeta_{\square} = 1.20.$$

Using the canonical representative

$$(\rho_{\min}, \zeta_{\square}) = (0.15, 1.20),$$

one gets

$$\min_D \frac{N_{\text{mass,ann}}^{(0.15,1.20)}(D)}{N_{\text{mass,in}}(D)} = 0.949763$$

at $D = 11$, while the axial-core leakage obeys

$$\max_D \frac{N_{\text{mass,tail}}^{(0.15,1.20)}(D)}{P_{\text{mass,in}}(D)} = 4.81178 \times 10^{-2}$$

again at $D = 11$. The same knot also maximizes the annular adverse load,

$$\max_D \frac{N_{\text{mass,ann}}^{(0.15,1.20)}(D)}{P_{\text{mass,in}}(D)} = 0.909694.$$

So the adverse-carrier side is already stabilized at an off-axis one-center annulus, with only a small axial-core leakage left outside it.

Corollary O27 (The remaining cancellation gap is an off-axis adverse-annulus problem). Combining Propositions O26–O28 gives

$$I_{\text{mass,in}} = (1 - \nu_{\text{diag,int}}^{(\pm,m)})P_{\text{mass,in}}, \quad \nu_{\text{diag,int}}^{(\pm,m)}(D) \leq \frac{N_{\text{mass,ann}}^{(0.15,1.20)}(D)}{P_{\text{mass,in}}(D)} + 4.81178 \times 10^{-2}.$$

At the exact floor knot $D = 11$, this reads

$$0.957812 = 0.909694 + 4.81178 \times 10^{-2},$$

so the entire mass-sector cancellation gap is already explained by an off-axis one-center adverse annulus plus a small axial-core leakage. In particular, the positive-carrier side is no longer the

issue, and neither shell leakage nor mirror leakage sets the overlap-side floor. The remaining proof-side target has narrowed once more: it is now enough to control the adverse annular carrier $N_{\text{mass,ann}}^{(0.15,1.20)}$, with the axial core retained only as a small corrective leakage.

Proposition O29 (The adverse annulus floor survives a strict axial trim). Keep the canonical adverse annulus

$$A_{\text{neg}}^{(0.15,1.20)}(D)$$

from Proposition O28 and define for any strict annular trim

$$A_{\text{neg}}^{(\rho_{\min}, \zeta_{\square})}(D) \subsetneq A_{\text{neg}}^{(0.15,1.20)}(D), \quad \rho_{\min} \geq 0.15, \quad \zeta_{\square} \leq 1.20,$$

the trimmed adverse-carrier moment

$$N_{\text{mass,subann}}^{(\rho_{\min}, \zeta_{\square})}(D) := \int_{A_{\text{neg}}^{(\rho_{\min}, \zeta_{\square})}(D)} [q_{\text{mass}}]_{-} d\rho dz.$$

Because $[q_{\text{mass}}]_{-} \geq 0$, one has the exact monotone split

$$N_{\text{mass,ann}}^{(0.15,1.20)}(D) = N_{\text{mass,subann}}^{(\rho_{\min}, \zeta_{\square})}(D) + N_{\text{mass,cap}}^{(\rho_{\min}, \zeta_{\square})}(D), \quad N_{\text{mass,cap}}^{(\rho_{\min}, \zeta_{\square})}(D) \geq 0.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_mass_annulus_source.py` scans

$$\rho_{\min} \in [0.15, 0.30], \quad \zeta_{\square} \in [1.10, 1.20]$$

inside the canonical annulus and finds that the optimal strict plateau already lies at the smallest possible radial trim,

$$\rho_{\min} = 0.15, \quad \zeta_{\square} = 1.19.$$

Using this canonical strict representative, one gets

$$\min_D \frac{N_{\text{mass,subann}}^{(0.15,1.19)}(D)}{N_{\text{mass,ann}}^{(0.15,1.20)}(D)} = 0.984061$$

at $D = 16$, while the omitted cap leakage satisfies

$$\max_D \frac{N_{\text{mass,cap}}^{(0.15,1.19)}(D)}{P_{\text{mass,in}}(D)} = 7.47726 \times 10^{-3}$$

again at $D = 16$. The trimmed adverse-annulus floor is

$$\min_D N_{\text{mass,subann}}^{(0.15,1.19)}(D) = 7.64028 \times 10^{-4},$$

whereas the full annulus floor is

$$\min_D N_{\text{mass,ann}}^{(0.15,1.20)}(D) = 7.73394 \times 10^{-4},$$

so the strict trim lowers the audited floor by only the factor

$$\frac{7.73394 \times 10^{-4}}{7.64028 \times 10^{-4}} = 1.01226.$$

Thus the adverse annular carrier is already stable at the canonical full-annulus scale: strict axial trims do not produce a meaningfully stronger theorem object.

Corollary O28 (The proof target stabilizes at the full adverse-annulus scale). Propositions O28 and O29 together imply that the support side of the adverse carrier is now exhausted:

$$\nu_{\text{diag,int}}^{(\pm,m)}(D) \leq \frac{N_{\text{mass,subann}}^{(0.15,1.19)}(D)}{P_{\text{mass,in}}(D)} + 4.81178 \times 10^{-2} + 7.47726 \times 10^{-3},$$

but the extra strict-trim leakage is already an order of magnitude smaller than the axial-core leakage from Proposition O28. So the remaining proof-side gap is no longer about how to choose the annular support: it is the adverse magnitude carried by the canonical annulus itself. In particular, the natural theorem object has stabilized at

$$N_{\text{mass,ann}}^{(0.15,1.20)},$$

with only tiny axial and cap leakages retained as corrective constants.

Proposition O30 (The adverse annulus magnitude localizes to a radial shoulder).

Keep the canonical adverse annulus $A_{\text{neg}}^{(0.15,1.20)}(D)$ and, for any strict radial trim

$$\rho_{\text{sh}} > 0.15,$$

define the adverse shoulder

$$S_{\text{neg}}^{(\rho_{\text{sh}},1.20)}(D) := \{(\rho, z) \in A_{\text{neg}}^{(0.15,1.20)}(D) : \rho \geq \rho_{\text{sh}}\},$$

with shoulder and rim moments

$$N_{\text{mass,shoulder}}^{(\rho_{\text{sh}},1.20)}(D) := \int_{S_{\text{neg}}^{(\rho_{\text{sh}},1.20)}(D)} [q_{\text{mass}}]_- d\rho dz,$$

$$N_{\text{mass,rim}}^{(\rho_{\text{sh}},1.20)}(D) := N_{\text{mass,ann}}^{(0.15,1.20)}(D) - N_{\text{mass,shoulder}}^{(\rho_{\text{sh}},1.20)}(D).$$

Because $[q_{\text{mass}}]_- \geq 0$, one has the exact monotone split

$$N_{\text{mass,ann}}^{(0.15,1.20)}(D) = N_{\text{mass,shoulder}}^{(\rho_{\text{sh}},1.20)}(D) + N_{\text{mass,rim}}^{(\rho_{\text{sh}},1.20)}(D), \quad N_{\text{mass,rim}}^{(\rho_{\text{sh}},1.20)}(D) \geq 0.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_mass_shoulder_source.py` scans strict radial trims

$$\rho_{\text{sh}} \in [0.16, 0.30]$$

inside the canonical annulus and finds a flat optimal plateau

$$\rho_{\text{sh}} \in [0.16, 0.20].$$

Using the clean canonical representative

$$\rho_{\text{sh}} = 0.20,$$

one gets

$$\min_D \frac{N_{\text{mass,shoulder}}^{(0.20,1.20)}(D)}{N_{\text{mass,ann}}^{(0.15,1.20)}(D)} = 0.940824$$

at $D = 11$, while the discarded inner-rim leakage obeys

$$\max_D \frac{N_{\text{mass,rin}}^{(0.20,1.20)}(D)}{P_{\text{mass,in}}(D)} = 5.38319 \times 10^{-2}$$

again at $D = 11$. The shoulder floor is

$$\min_D N_{\text{mass,shoulder}}^{(0.20,1.20)}(D) = 7.37447 \times 10^{-4},$$

whereas the full adverse-annulus floor is

$$\min_D N_{\text{mass,ann}}^{(0.15,1.20)}(D) = 7.73394 \times 10^{-4},$$

so the radial shoulder loses only the factor

$$\frac{7.73394 \times 10^{-4}}{7.37447 \times 10^{-4}} = 1.04874.$$

Thus the adverse magnitude is already carried almost entirely by an off-axis radial shoulder: the inner rim is only a secondary leakage.

Corollary O29 (The cancellation gap narrows to an adverse shoulder plus two small leakages). Combining Proposition O30 with Proposition O28 yields

$$\nu_{\text{diag,int}}^{(\pm,m)}(D) \leq \frac{N_{\text{mass,shoulder}}^{(0.20,1.20)}(D)}{P_{\text{mass,in}}(D)} + 5.38319 \times 10^{-2} + 4.81178 \times 10^{-2}.$$

At the exact worst cancellation knot $D = 11$, this reads

$$0.957812 = 0.855863 + 5.38319 \times 10^{-2} + 4.81178 \times 10^{-2}.$$

So the positive carrier, the axial core, and the inner radial rim are all now demoted to corrective terms. The remaining proof-side target has narrowed again: it is enough to control the off-axis one-center adverse shoulder

$$N_{\text{mass,shoulder}}^{(0.20,1.20)}$$

with only the small axial-core and inner-rim leakages retained outside it.

Proposition O31 (The adverse shoulder floor survives a genuine axial hollow). Keep the canonical adverse shoulder

$$N_{\text{mass,shoulder}}^{(0.20,1.20)}(D)$$

from Proposition O30 and define, for any axial hollow

$$\zeta_{\text{low}} > 0,$$

the adverse shoulder band

$$B_{\text{neg}}^{(0.20;\zeta_{\text{low}},1.20)}(D) := \{(\rho, z) \in S_{\text{neg}}^{(0.20,1.20)}(D) : \zeta_{\text{low}} \leq |z \mp D/2| \leq 1.20\},$$

with band and shoulder-core moments

$$N_{\text{mass,band}}^{(0.20;\zeta_{\text{low}},1.20)}(D) := \int_{B_{\text{neg}}^{(0.20;\zeta_{\text{low}},1.20)}(D)} [q_{\text{mass}}]_- d\rho dz,$$

$$N_{\text{mass,core}}^{(0.20;\zeta_{\text{low}},1.20)}(D) := N_{\text{mass,shoulder}}^{(0.20,1.20)}(D) - N_{\text{mass,band}}^{(0.20;\zeta_{\text{low}},1.20)}(D).$$

Because $[q_{\text{mass}}]_- \geq 0$, one has the exact monotone split

$$N_{\text{mass,shoulder}}^{(0.20,1.20)}(D) = N_{\text{mass,band}}^{(0.20;\zeta_{\text{low}},1.20)}(D) + N_{\text{mass,core}}^{(0.20;\zeta_{\text{low}},1.20)}(D), \quad N_{\text{mass,core}}^{(0.20;\zeta_{\text{low}},1.20)}(D) \geq 0.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_mass_shoulders_band_source.py` scans genuine hollows

$$\zeta_{\text{low}} \in [0.10, 0.40]$$

inside the canonical shoulder and finds that the best hollow already sits at the smallest nontrivial cut,

$$\zeta_{\text{low}} = 0.10.$$

For this canonical representative one gets

$$\min_D \frac{N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D)}{N_{\text{mass,shoulder}}^{(0.20,1.20)}(D)} = 0.951019$$

at $D = 11$, while the discarded shoulder-core leakage satisfies

$$\max_D \frac{N_{\text{mass,core}}^{(0.20;0.10,1.20)}(D)}{P_{\text{mass,in}}(D)} = 4.19214 \times 10^{-2}$$

again at $D = 11$. The adverse-band floor is

$$\min_D N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D) = 7.37447 \times 10^{-4},$$

exactly matching the canonical shoulder floor

$$\min_D N_{\text{mass,shoulder}}^{(0.20,1.20)}(D) = 7.37447 \times 10^{-4},$$

so the audited factor

$$\frac{\min_D N_{\text{mass,shoulder}}^{(0.20,1.20)}(D)}{\min_D N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D)} = 1$$

to machine precision. Thus the adverse shoulder magnitude is already carried by a genuinely hollow off-axis band: the shoulder-core piece is only a small corrective leakage.

Corollary O30 (The cancellation gap narrows to an adverse band plus three small leakages). Combining Proposition O31 with Corollary O29 gives

$$\nu_{\text{diag,int}}^{(\pm,m)}(D) \leq \frac{N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D)}{P_{\text{mass,in}}(D)} + 4.19214 \times 10^{-2} + 5.38319 \times 10^{-2} + 4.81178 \times 10^{-2}.$$

At the exact worst cancellation knot $D = 11$, this reads

$$0.957812 = 0.813941 + 4.19214 \times 10^{-2} + 5.38319 \times 10^{-2} + 4.81178 \times 10^{-2}.$$

Proposition O32 (Strict corner trims do not improve the adverse-band floor). Keep the canonical adverse band

$$N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D)$$

from Proposition O31 and define, for any strict corner trim

$$\rho_{\text{corner}} > 0.20, \quad \zeta_{\text{corner}} \geq 0.10,$$

the corner subband

$$C_{\text{neg}}^{(\rho_{\text{corner}};\zeta_{\text{corner}},1.20)}(D) := \{(\rho, z) \in B_{\text{neg}}^{(0.20;0.10,1.20)}(D) : \rho \geq \rho_{\text{corner}}, \zeta_{\text{corner}} \leq |z \mp D/2| \leq 1.20\},$$

with corner and cap moments

$$N_{\text{mass,corner}}^{(\rho_{\text{corner}};\zeta_{\text{corner}},1.20)}(D) := \int_{C_{\text{neg}}^{(\rho_{\text{corner}};\zeta_{\text{corner}},1.20)}(D)} [q_{\text{mass}}]_- d\rho dz,$$

$$N_{\text{mass,cap}}^{(\rho_{\text{corner}}; \zeta_{\text{corner}}, 1.20)}(D) := N_{\text{mass,band}}^{(0.20; 0.10, 1.20)}(D) - N_{\text{mass,corner}}^{(\rho_{\text{corner}}; \zeta_{\text{corner}}, 1.20)}(D).$$

Because $[q_{\text{mass}}]_- \geq 0$, one has the exact monotone split

$$N_{\text{mass,band}}^{(0.20; 0.10, 1.20)}(D) = N_{\text{mass,corner}}^{(\rho_{\text{corner}}; \zeta_{\text{corner}}, 1.20)}(D) + N_{\text{mass,cap}}^{(\rho_{\text{corner}}; \zeta_{\text{corner}}, 1.20)}(D), \quad N_{\text{mass,cap}}^{(\rho_{\text{corner}}; \zeta_{\text{corner}}, 1.20)}(D) \geq 0.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_mass_band_corner_source.py` scans genuinely strict corner trims

$$\rho_{\text{corner}} \in [0.22, 0.30], \quad \zeta_{\text{corner}} \in [0.10, 0.40],$$

inside the canonical adverse band and finds the best plateau

$$\rho_{\text{corner}} \in [0.22, 0.27], \quad \zeta_{\text{corner}} = 0.10.$$

Using the clean canonical representative

$$(\rho_{\text{corner}}, \zeta_{\text{corner}}) = (0.22, 0.10),$$

one gets

$$\min_D \frac{N_{\text{mass,corner}}^{(0.22; 0.10, 1.20)}(D)}{N_{\text{mass,band}}^{(0.20; 0.10, 1.20)}(D)} = 0.921342$$

at $D = 4$, while the discarded cap leakage obeys

$$\max_D \frac{N_{\text{mass,cap}}^{(0.22; 0.10, 1.20)}(D)}{P_{\text{mass,in}}(D)} = 6.35145 \times 10^{-2}$$

at $D = 11$. The strict-corner floor is

$$\min_D N_{\text{mass,corner}}^{(0.22; 0.10, 1.20)}(D) = 6.88668 \times 10^{-4},$$

whereas the canonical adverse-band floor is

$$\min_D N_{\text{mass,band}}^{(0.20; 0.10, 1.20)}(D) = 7.37447 \times 10^{-4},$$

so every genuinely strict corner trim loses the audited factor

$$\frac{7.37447 \times 10^{-4}}{6.88668 \times 10^{-4}} = 1.07083.$$

Thus support-side localization is already exhausted at the full adverse-band scale: a strict corner trim does not produce a better theorem object.

Corollary O31 (The cancellation gap stabilizes at the full adverse-band scale). Combining Proposition O32 with Corollary O30 gives

$$\nu_{\text{diag,int}}^{(\pm, m)}(D) \leq \frac{N_{\text{mass,corner}}^{(0.22; 0.10, 1.20)}(D)}{P_{\text{mass,in}}(D)} + 6.35145 \times 10^{-2} + 4.19214 \times 10^{-2} + 5.38319 \times 10^{-2} + 4.81178 \times 10^{-2}.$$

At the exact worst cancellation knot $D = 11$, this reads

$$0.957812 = 0.750427 + 6.35145 \times 10^{-2} + 4.19214 \times 10^{-2} + 5.38319 \times 10^{-2} + 4.81178 \times 10^{-2}.$$

So the dominant adverse band has now been tested against a genuinely strict corner localization and survives as the stable proof-side target: the remaining gap is best read as the full canonical adverse band plus only small cap, shoulder-core, inner-rim, and axial-core leakages.

Proposition O33 (The canonical adverse-band magnitude factorizes into a positive static kernel and a negative carrier moment). Keep the canonical adverse band

$$B_{\text{neg}}^{(0.20;0.10,1.20)}(D)$$

from Proposition O31. Define the one-center local mass kernel

$$K_{\text{band}}^{\text{mass}}(\rho, z; D) := (U_+^{\text{mass}} - \bar{U}_+^{\text{mass}})\mathbf{1}_{B_+} + (U_-^{\text{mass}} - \bar{U}_-^{\text{mass}})\mathbf{1}_{B_-},$$

restricted to $B_{\text{neg}}^{(0.20;0.10,1.20)}(D)$, and the negative carrier density

$$W_{\text{band}}^{(-)}(\rho, z; D) := [-\text{sgn}(I_{\text{diag,anchor}}(D))\mathcal{C}(\rho, z; D)]_+ \mathbf{1}_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)}.$$

Then on the canonical band one has the exact pointwise factorization

$$[q_{\text{mass}}]_-(\rho, z; D) = K_{\text{band}}^{\text{mass}}(\rho, z; D) W_{\text{band}}^{(-)}(\rho, z; D),$$

provided $K_{\text{band}}^{\text{mass}} > 0$ there. The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_mass_band_magnitude_source.py` verifies precisely this structure across all audited knots $D = 4, \dots, 20$:

$$\max_D \|[q_{\text{mass}}]_- - K_{\text{band}}^{\text{mass}} W_{\text{band}}^{(-)}\|_{\infty} = 0,$$

and

$$\min_D \inf_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} K_{\text{band}}^{\text{mass}} = 1.82852 \times 10^{-2}$$

at $D = 4$, so the kernel stays strictly positive on the whole canonical band. Writing the negative carrier moment as

$$W_{\text{band}}^{(-)}(D) := \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} W_{\text{band}}^{(-)}(\rho, z; D) d\rho dz,$$

and the weighted kernel average as

$$\bar{K}_{\text{band}}^{(-)}(D) := \frac{N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D)}{W_{\text{band}}^{(-)}(D)},$$

the audit finds

$$5.38460 \times 10^{-2} \leq \bar{K}_{\text{band}}^{(-)}(D) \leq 1.00728 \times 10^{-1}$$

for all audited D , with the minimum at $D = 5$ and the maximum at $D = 18$. The negative-carrier floor is

$$\min_D W_{\text{band}}^{(-)}(D) = 8.40086 \times 10^{-3}$$

at $D = 20$, while the crude kernel-floor product

$$\inf_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} K_{\text{band}}^{\text{mass}} \cdot W_{\text{band}}^{(-)}(D)$$

has audited floor

$$4.28523 \times 10^{-4}$$

at $D = 20$. Thus the adverse-band magnitude is no longer a black-box carrier: it is a positive one-center mass kernel against a negative carrier moment.

Corollary O32 (The remaining adverse-band proof target is the negative carrier moment). By Proposition O33, the canonical adverse-band magnitude satisfies the exact identity

$$N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D) = \bar{K}_{\text{band}}^{(-)}(D) W_{\text{band}}^{(-)}(D),$$

with

$$5.38460 \times 10^{-2} \leq \bar{K}_{\text{band}}^{(-)}(D) \leq 1.00728 \times 10^{-1}.$$

Equivalently,

$$\inf_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} K_{\text{band}}^{\text{mass}} \cdot W_{\text{band}}^{(-)}(D) \leq N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D) \leq \sup_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} K_{\text{band}}^{\text{mass}} \cdot W_{\text{band}}^{(-)}(D).$$

The audited band floor remains

$$\min_D N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D) = 7.37447 \times 10^{-4},$$

whereas the kernel-floor product has floor

$$4.28523 \times 10^{-4},$$

so the exact band floor exceeds the coarse product floor by the factor

$$\frac{7.37447 \times 10^{-4}}{4.28523 \times 10^{-4}} = 1.72090.$$

Therefore the support problem is already closed and the magnitude problem has narrowed again: to control the canonical adverse band, it is now enough to control the negative carrier moment $W_{\text{band}}^{(-)}$, with the positive one-center mass kernel confined to a narrow audited interval.

Proposition O34 (The negative carrier moment factorizes through an almost constant radial weight). Keep the canonical adverse band

$$B_{\text{neg}}^{(0.20;0.10,1.20)}(D)$$

and define the negative overlap-band density

$$M_{\text{band}}^{(-)}(\rho, z; D) := \left[-\text{sgn}(I_{\text{diag,anchor}}(D))(\psi_1(\rho, z; D)^2 - \psi_2(\rho, z; D)^2) \right]_+ \mathbf{1}_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)}.$$

Since

$$\mathcal{C}(\rho, z; D) = 2\pi\rho(\psi_1(\rho, z; D)^2 - \psi_2(\rho, z; D)^2),$$

the canonical negative carrier moment satisfies the exact pointwise identity

$$W_{\text{band}}^{(-)}(\rho, z; D) = 2\pi\rho M_{\text{band}}^{(-)}(\rho, z; D).$$

Integrating over the canonical band gives

$$W_{\text{band}}^{(-)}(D) = 2\pi \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} \rho M_{\text{band}}^{(-)}(\rho, z; D) d\rho dz.$$

Define the unweighted overlap-band moment

$$M_{\text{band}}^{(-)}(D) := \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} M_{\text{band}}^{(-)}(\rho, z; D) d\rho dz,$$

and the weighted radial factor

$$\bar{\rho}_{\text{band}}^{(-)}(D) := \frac{W_{\text{band}}^{(-)}(D)}{2\pi M_{\text{band}}^{(-)}(D)}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_band_negative_carrier_source.py` verifies the exact factorization

$$\max_D \|W_{\text{band}}^{(-)} - 2\pi\rho M_{\text{band}}^{(-)}\|_{\infty} = 1.35525 \times 10^{-20},$$

and finds the weighted radial factor confined to the tiny interval

$$0.4434097 \leq \bar{\rho}_{\text{band}}^{(-)}(D) \leq 0.4438634,$$

with the minimum at $D = 7$ and the maximum at $D = 6$. Its relative spread is only

$$\frac{\max \bar{\rho}_{\text{band}}^{(-)} - \min \bar{\rho}_{\text{band}}^{(-)}}{\text{mean } \bar{\rho}_{\text{band}}^{(-)}} = 1.02267 \times 10^{-3}.$$

The canonical negative-carrier floor is

$$\min_D W_{\text{band}}^{(-)}(D) = 8.40086 \times 10^{-3}$$

at $D = 20$, and the unweighted overlap-band floor is

$$\min_D M_{\text{band}}^{(-)}(D) = 3.01270 \times 10^{-3}$$

again at $D = 20$. Therefore the negative carrier moment itself is already equivalent, up to an almost constant radial factor, to an unweighted overlap-band moment.

Corollary O33 (The remaining carrier-side proof target is the unweighted negative overlap-band moment). By Proposition O34, one has the exact identity

$$W_{\text{band}}^{(-)}(D) = 2\pi \bar{\rho}_{\text{band}}^{(-)}(D) M_{\text{band}}^{(-)}(D),$$

with

$$0.4434097 \leq \bar{\rho}_{\text{band}}^{(-)}(D) \leq 0.4438634.$$

Equivalently,

$$2\pi (0.4434097) M_{\text{band}}^{(-)}(D) \leq W_{\text{band}}^{(-)}(D) \leq 2\pi (0.4438634) M_{\text{band}}^{(-)}(D).$$

At the audited floor knot $D = 20$, this reads

$$W_{\text{band}}^{(-)}(20) = 2\pi (0.4438002) M_{\text{band}}^{(-)}(20).$$

Moreover,

$$\frac{\min_D W_{\text{band}}^{(-)}(D)}{2\pi(\min_D \bar{\rho}_{\text{band}}^{(-)}(D))(\min_D M_{\text{band}}^{(-)}(D))} = 1.00088,$$

so even the coarsest audited lower product is already essentially exact. Thus the carrier-side magnitude problem narrows once more: to control $W_{\text{band}}^{(-)}$, it is enough to control the unweighted negative overlap-band moment $M_{\text{band}}^{(-)}$, with only an almost constant radial factor left outside it.

Proposition O35 (The negative overlap-band moment is an absolute overlap-band mass times a signed-imbalance factor). Keep the canonical adverse band

$$B_{\text{neg}}^{(0.20;0.10,1.20)}(D)$$

and define the unweighted signed overlap-band density

$$m_{\text{band}}(\rho, z; D) := \text{sgn}(I_{\text{diag,anchor}}(D))(\psi_1(\rho, z; D)^2 - \psi_2(\rho, z; D)^2) \mathbf{1}_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)}.$$

Define its negative moment, absolute moment, and signed residue by

$$M_{\text{band}}^{(-)}(D) := \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} [-m_{\text{band}}(\rho, z; D)]_+ d\rho dz,$$

$$M_{\text{band}}^{(\text{abs})}(D) := \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} |m_{\text{band}}(\rho, z; D)| d\rho dz,$$

$$R_{\text{band}}^{(\text{sgn})}(D) := \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} m_{\text{band}}(\rho, z; D) d\rho dz.$$

Since $x_- = (|x| - x)/2$, one has the exact identity

$$M_{\text{band}}^{(-)}(D) = \frac{M_{\text{band}}^{(\text{abs})}(D) - R_{\text{band}}^{(\text{sgn})}(D)}{2}.$$

Equivalently, if

$$r_{\text{band}}^{(\text{sgn})}(D) := \frac{R_{\text{band}}^{(\text{sgn})}(D)}{M_{\text{band}}^{(\text{abs})}(D)}, \quad \theta_{\text{band}}^{(-)}(D) := \frac{1 - r_{\text{band}}^{(\text{sgn})}(D)}{2},$$

then

$$M_{\text{band}}^{(-)}(D) = \theta_{\text{band}}^{(-)}(D) M_{\text{band}}^{(\text{abs})}(D).$$

Moreover, using the localized basis

$$\phi_R = \frac{\psi_1 + \psi_2}{\sqrt{2}}, \quad \phi_L = \frac{\psi_1 - \psi_2}{\sqrt{2}},$$

the parity-contrast identity

$$\psi_1^2 - \psi_2^2 = 2\phi_R\phi_L$$

implies

$$M_{\text{band}}^{(\text{abs})}(D) = 2 \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} |\phi_R(\rho, z; D)\phi_L(\rho, z; D)| d\rho dz.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_band_overlap_moment_source.py` verifies both identities to machine precision:

$$\max_D \left| M_{\text{band}}^{(-)} - \frac{M_{\text{band}}^{(\text{abs})} - R_{\text{band}}^{(\text{sgn})}}{2} \right| = 1.73472 \times 10^{-18},$$

$$\max_D \left| M_{\text{band}}^{(\text{abs})} - 2 \int |\phi_R\phi_L| \right| = 3.46945 \times 10^{-18}.$$

The signed-imbalance ratio satisfies

$$-2.03788 \times 10^{-1} \leq r_{\text{band}}^{(\text{sgn})}(D) \leq 9.60339 \times 10^{-2},$$

with the maximum at $D = 11$ and the minimum at $D = 7$. Hence the negative-fraction factor satisfies

$$4.51983 \times 10^{-1} \leq \theta_{\text{band}}^{(-)}(D) \leq 6.01894 \times 10^{-1},$$

with the minimum again at $D = 11$. The absolute overlap-band floor is

$$\min_D M_{\text{band}}^{(\text{abs})}(D) = 6.45299 \times 10^{-3}$$

at $D = 19$, while the negative overlap-band floor remains

$$\min_D M_{\text{band}}^{(-)}(D) = 3.01270 \times 10^{-3}$$

at $D = 20$.

Corollary O34 (The remaining overlap-side target is the absolute overlap-band mass). By Proposition O35, one has the exact identity

$$M_{\text{band}}^{(-)}(D) = \theta_{\text{band}}^{(-)}(D) M_{\text{band}}^{(\text{abs})}(D), \quad \theta_{\text{band}}^{(-)}(D) = \frac{1 - r_{\text{band}}^{(\text{sgn})}(D)}{2}.$$

Since

$$r_{\text{band}}^{(\text{sgn})}(D) \leq 9.60339 \times 10^{-2},$$

it follows that

$$M_{\text{band}}^{(-)}(D) \geq (4.51983 \times 10^{-1}) M_{\text{band}}^{(\text{abs})}(D).$$

Combining this with the audited absolute overlap-band floor gives the direct lower bound

$$M_{\text{band}}^{(-)}(D) \geq (4.51983 \times 10^{-1})(6.45299 \times 10^{-3}) = 2.91664 \times 10^{-3}.$$

Compared with the exact negative overlap-band floor

$$3.01270 \times 10^{-3},$$

this misses by only the factor

$$\frac{3.01270 \times 10^{-3}}{2.91664 \times 10^{-3}} = 1.03294.$$

Thus the support problem, the kernel problem, and the radial-weight problem are all closed. What remains on the overlap side is no longer the clipped negative carrier itself, but the absolute localized overlap-band mass

$$M_{\text{band}}^{(\text{abs})}(D) = 2 \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} |\phi_R \phi_L| d\rho dz,$$

with only a small signed-imbalance factor left outside it.

Proposition O36 (The absolute overlap-band mass factorizes through a same-center carrier moment and a weighted mirror-tail average). Keep the same canonical band

$$B_{\text{neg}}^{(0.20;0.10,1.20)}(D)$$

and the localized basis (ϕ_R, ϕ_L) . On the canonical band, define the same-center and mirror amplitudes by

$$\Phi_{\text{self}}(\rho, z; D) := |\phi_R(\rho, z; D)| \mathbf{1}_{B_+(D)} + |\phi_L(\rho, z; D)| \mathbf{1}_{B_-(D)},$$

$$\Phi_{\text{mir}}(\rho, z; D) := |\phi_L(\rho, z; D)| \mathbf{1}_{B_+(D)} + |\phi_R(\rho, z; D)| \mathbf{1}_{B_-(D)}.$$

Then, by the localized-overlap identity from Proposition O35,

$$M_{\text{band}}^{(\text{abs})}(D) = 2 \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} \Phi_{\text{self}}(\rho, z; D) \Phi_{\text{mir}}(\rho, z; D) d\rho dz.$$

Define the same-center and mirror carrier moments

$$S_{\text{band}}^{(\text{self})}(D) := \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} \Phi_{\text{self}}(\rho, z; D) d\rho dz,$$

$$S_{\text{band}}^{(\text{mir})}(D) := \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} \Phi_{\text{mir}}(\rho, z; D) d\rho dz,$$

and the weighted local averages

$$\bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D) := \frac{M_{\text{band}}^{(\text{abs})}(D)}{2 S_{\text{band}}^{(\text{self})}(D)}, \quad \bar{\Phi}_{\text{band}}^{(\text{self}|\text{mir})}(D) := \frac{M_{\text{band}}^{(\text{abs})}(D)}{2 S_{\text{band}}^{(\text{mir})}(D)}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_band_absolute_overlap_source.py` verifies the exact factorization

$$\max_D \left| M_{\text{band}}^{(\text{abs})} - 2 S_{\text{band}}^{(\text{self})} \bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})} \right| = 0,$$

and equivalently

$$\max_D \left| M_{\text{band}}^{(\text{abs})} - 2 S_{\text{band}}^{(\text{mir})} \bar{\Phi}_{\text{band}}^{(\text{self}|\text{mir})} \right| = 0.$$

On the audited window, the same-center carrier floor is

$$\min_D S_{\text{band}}^{(\text{self})}(D) = 1.66599 \times 10^{-1}$$

at $D = 20$, while the weighted mirror-tail average satisfies

$$1.92259 \times 10^{-2} \leq \bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D) \leq 5.20575 \times 10^{-2},$$

with the minimum at $D = 19$. Therefore the self-carrier route gives the induced product floor

$$2(\min_D S_{\text{band}}^{(\text{self})}(D))(\min_D \bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D)) = 6.40606 \times 10^{-3}.$$

The reverse route is weaker: although

$$\min_D S_{\text{band}}^{(\text{mir})}(D) = 5.79226 \times 10^{-2}$$

at $D = 20$ and

$$5.30602 \times 10^{-2} \leq \bar{\Phi}_{\text{band}}^{(\text{self}|\text{mir})}(D) \leq 7.34746 \times 10^{-2},$$

the corresponding product floor is only

$$6.14677 \times 10^{-3}.$$

Hence the tighter canonical factorization is the same-center-carrier / weighted-mirror route.

Corollary O35 (The absolute overlap-band mass is already essentially fixed by the same-center carrier route). By Proposition O36,

$$M_{\text{band}}^{(\text{abs})}(D) = 2 S_{\text{band}}^{(\text{self})}(D) \bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D).$$

Using the audited floors

$$S_{\text{band}}^{(\text{self})}(D) \geq 1.66599 \times 10^{-1}, \quad \bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D) \geq 1.92259 \times 10^{-2},$$

one gets the direct lower bound

$$M_{\text{band}}^{(\text{abs})}(D) \geq 2(1.66599 \times 10^{-1})(1.92259 \times 10^{-2}) = 6.40606 \times 10^{-3}.$$

The exact audited absolute overlap-band floor is

$$\min_D M_{\text{band}}^{(\text{abs})}(D) = 6.45299 \times 10^{-3}$$

at $D = 19$, so this misses the exact floor by only the factor

$$\frac{6.45299 \times 10^{-3}}{6.40606 \times 10^{-3}} = 1.00733.$$

Therefore the remaining overlap-side proof target has narrowed again. It is no longer the full absolute overlap-band mass as an undifferentiated object, but the weighted mirror-tail average

$$\bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D),$$

with the same-center carrier moment already providing an almost saturated one-center positive source.

Proposition O37 (The weighted mirror-tail average factorizes through an ordinary mirror-band mean and a self-weight bias). Keep the same canonical band

$$B_{\text{neg}}^{(0.20;0.10,1.20)}(D)$$

and define its band area

$$A_{\text{band}}(D) := \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} 1 \, d\rho \, dz.$$

Using the mirror-tail carrier moment

$$S_{\text{band}}^{(\text{mir})}(D) := \int_{B_{\text{neg}}^{(0.20;0.10,1.20)}(D)} \Phi_{\text{mir}}(\rho, z; D) \, d\rho \, dz,$$

define the ordinary mirror-band mean

$$\bar{\Phi}_{\text{band}}^{(\text{mir})}(D) := \frac{S_{\text{band}}^{(\text{mir})}(D)}{A_{\text{band}}(D)}.$$

Then the weighted mirror-tail average from Proposition O36 satisfies the exact identity

$$\bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D) = \Xi_{\text{band}}^{(\text{self})}(D) \bar{\Phi}_{\text{band}}^{(\text{mir})}(D),$$

where

$$\Xi_{\text{band}}^{(\text{self})}(D) := \frac{\bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D)}{\bar{\Phi}_{\text{band}}^{(\text{mir})}(D)} = \frac{A_{\text{band}}(D) \int \Phi_{\text{self}} \Phi_{\text{mir}}}{(\int \Phi_{\text{self}}) (\int \Phi_{\text{mir}})}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_band_mirror_average_source.py` verifies this factorization with zero residual:

$$\max_D \left| \bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})} - \Xi_{\text{band}}^{(\text{self})} \bar{\Phi}_{\text{band}}^{(\text{mir})} \right| = 0.$$

It finds

$$2.43229 \times 10^{-2} \leq \bar{\Phi}_{\text{band}}^{(\text{mir})}(D) \leq 6.23458 \times 10^{-2},$$

with the floor at $D = 20$, while the self-weight bias lies in the interval

$$7.40007 \times 10^{-1} \leq \Xi_{\text{band}}^{(\text{self})}(D) \leq 8.35352 \times 10^{-1},$$

with the minimum at $D = 5$ and the maximum at $D = 9$. The band area itself is almost fixed:

$$2.3814 \leq A_{\text{band}}(D) \leq 2.3976,$$

with relative spread only

$$6.78914 \times 10^{-3}.$$

Thus the weighted mirror-tail average differs from the ordinary mirror-band mean only by a positive self-weight bias and a tiny geometric area aliasing.

Corollary O36 (The remaining multiplier-side target can be taken to be the ordinary mirror-band mean). By Proposition O37,

$$\bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D) = \Xi_{\text{band}}^{(\text{self})}(D) \bar{\Phi}_{\text{band}}^{(\text{mir})}(D).$$

Using the audited floors

$$\Xi_{\text{band}}^{(\text{self})}(D) \geq 7.40007 \times 10^{-1}, \quad \bar{\Phi}_{\text{band}}^{(\text{mir})}(D) \geq 2.43229 \times 10^{-2},$$

one gets the direct lower bound

$$\bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D) \geq (7.40007 \times 10^{-1})(2.43229 \times 10^{-2}) = 1.79991 \times 10^{-2}.$$

Compared with the exact weighted mirror-tail floor

$$\min_D \bar{\Phi}_{\text{band}}^{(\text{mir}|\text{self})}(D) = 1.92259 \times 10^{-2},$$

this misses by only the factor

$$\frac{1.92259 \times 10^{-2}}{1.79991 \times 10^{-2}} = 1.06816.$$

Since $A_{\text{band}}(D)$ varies by less than one percent, the same statement may equivalently be read in terms of the mirror carrier moment $S_{\text{band}}^{(\text{mir})}$, at the cost of only a tiny geometric aliasing. Therefore the remaining overlap-side multiplier target is no longer the self-weighted mirror average itself, but the ordinary mirror-band mean

$$\bar{\Phi}_{\text{band}}^{(\text{mir})}(D) = \frac{S_{\text{band}}^{(\text{mir})}(D)}{A_{\text{band}}(D)},$$

with the self-weight bias confined to a narrow positive interval.

Proposition O38 (The ordinary mirror-band mean factors exactly through the same-center band mean and the mirror/self carrier ratio). Keep the same canonical band

$$B_{\text{neg}}^{(0.20;0.10,1.20)}(D).$$

Define the same-center band mean

$$\bar{\Phi}_{\text{band}}^{(\text{self})}(D) := \frac{S_{\text{band}}^{(\text{self})}(D)}{A_{\text{band}}(D)}$$

and the mirror/self carrier ratio

$$\Lambda_{\text{band}}^{(\text{mir}|\text{self})}(D) := \frac{S_{\text{band}}^{(\text{mir})}(D)}{S_{\text{band}}^{(\text{self})}(D)}.$$

Then the ordinary mirror-band mean satisfies the exact identity

$$\bar{\Phi}_{\text{band}}^{(\text{mir})}(D) = \Lambda_{\text{band}}^{(\text{mir}|\text{self})}(D) \bar{\Phi}_{\text{band}}^{(\text{self})}(D).$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_diag_anchor_band_mirror_ratio_source.py` verifies this factorization with machine-level residual

$$\max_D \left| \bar{\Phi}_{\text{band}}^{(\text{mir})} - \Lambda_{\text{band}}^{(\text{mir}|\text{self})} \bar{\Phi}_{\text{band}}^{(\text{self})} \right| = 6.94 \times 10^{-18}.$$

It finds

$$6.99586 \times 10^{-2} \leq \bar{\Phi}_{\text{band}}^{(\text{self})}(D) \leq 8.80198 \times 10^{-2},$$

with the floor at $D = 20$, while the carrier ratio lies in the interval

$$3.47676 \times 10^{-1} \leq \Lambda_{\text{band}}^{(\text{mir}|\text{self})}(D) \leq 7.08509 \times 10^{-1},$$

with the minimum at $D = 20$ and the maximum at $D = 4$. The ordinary mirror-band mean itself has audited floor

$$\min_D \bar{\Phi}_{\text{band}}^{(\text{mir})}(D) = 2.43229 \times 10^{-2}$$

at the same knot $D = 20$.

Corollary O37 (The ordinary mirror-band mean is already saturated by the same-center band mean and the mirror/self ratio). By Proposition O38,

$$\bar{\Phi}_{\text{band}}^{(\text{mir})}(D) = \Lambda_{\text{band}}^{(\text{mir}|\text{self})}(D) \bar{\Phi}_{\text{band}}^{(\text{self})}(D).$$

Using the audited floors

$$\bar{\Phi}_{\text{band}}^{(\text{self})}(D) \geq 6.99586 \times 10^{-2}, \quad \Lambda_{\text{band}}^{(\text{mir}|\text{self})}(D) \geq 3.47676 \times 10^{-1},$$

one gets

$$\bar{\Phi}_{\text{band}}^{(\text{mir})}(D) \geq (6.99586 \times 10^{-2})(3.47676 \times 10^{-1}) = 2.43229 \times 10^{-2}.$$

This equals the exact audited floor

$$\min_D \bar{\Phi}_{\text{band}}^{(\text{mir})}(D) = 2.43229 \times 10^{-2},$$

up to machine precision, with all three floor knots aligned at $D = 20$. Therefore the ordinary mirror-band mean is no longer a substantive multiplier-side blocker on the current exact export: it is already closed by a same-center positive carrier and a mirror/self ratio, with no further support refinement required.

Proposition O39 (Synthesis of the overlap-side reduction chain on the canonical exact bridge). Combining Propositions O19–O38, the canonical adverse-band piece of the interior one-center mass-sector cancellation admits the exact composite factorization

$$N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D) = 4\pi \bar{K}_{\text{band}}^{(-)}(D) \bar{\rho}_{\text{band}}^{(-)}(D) \theta_{\text{band}}^{(-)}(D) \Xi_{\text{band}}^{(\text{self})}(D) \Lambda_{\text{band}}^{(\text{mir|self})}(D) \frac{S_{\text{band}}^{(\text{self})}(D)^2}{A_{\text{band}}(D)}.$$

Equivalently, the normalized interior mass-sector cancellation ratio satisfies

$$\nu_{\text{diag,int}}^{(\pm,m)}(D) = \frac{N_{\text{mass,band}}^{(0.20;0.10,1.20)}(D)}{P_{\text{mass,in}}(D)} + \frac{N_{\text{mass,core}}^{(0.20;0.10,1.20)}(D)}{P_{\text{mass,in}}(D)} + \frac{N_{\text{mass,rim}}^{(0.20,1.20)}(D)}{P_{\text{mass,in}}(D)} + \frac{N_{\text{mass,tail}}^{(0.15,1.20)}(D)}{P_{\text{mass,in}}(D)}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_parity_contrast_synthesis_source.py` verifies both identities on the exact audited knot set with residuals

$$\max_D \left| N_{\text{mass,band}}^{(0.20;0.10,1.20)} - 4\pi \bar{K}_{\text{band}}^{(-)} \bar{\rho}_{\text{band}}^{(-)} \theta_{\text{band}}^{(-)} \Xi_{\text{band}}^{(\text{self})} \Lambda_{\text{band}}^{(\text{mir|self})} \frac{S_{\text{band}}^{(\text{self})2}}{A_{\text{band}}} \right| = 9.67 \times 10^{-17},$$

$$\max_D \left| \nu_{\text{diag,int}}^{(\pm,m)} - \frac{N_{\text{mass,band}}}{P_{\text{mass,in}}} - \frac{N_{\text{mass,core}}}{P_{\text{mass,in}}} - \frac{N_{\text{mass,rim}}}{P_{\text{mass,in}}} - \frac{N_{\text{mass,tail}}}{P_{\text{mass,in}}} \right| = 5.97 \times 10^{-14}.$$

At the canonical adverse-band floor knot $D = 19$, the composite mainline factor already reproduces the exact band magnitude:

$$N_{\text{mass,band}}^{(0.20;0.10,1.20)}(19) = 7.37447 \times 10^{-4},$$

with residual 2.72×10^{-17} . At the worst interior cancellation knot $D = 11$, the normalized decomposition reads

$$\nu_{\text{diag,int}}^{(\pm,m)}(11) = 0.813941 + 4.19214 \times 10^{-2} + 5.38319 \times 10^{-2} + 4.81178 \times 10^{-2} = 0.957812,$$

so the audited exact bridge is now reduced to one multiplicative mainline and three explicit leakage channels.

Corollary O38 (What is closed, and what remains only as secondary tightening).

On the current exact export, the overlap-side chain from Propositions O19–O38 now separates cleanly into *closed* structural blocks and *secondary tightening* blocks.

Closed on the audited exact bridge:

$$\text{support reduction (O19--O32),} \quad P_{\text{mass,in}}, \quad \bar{\Phi}_{\text{band}}^{(\text{self})}, \quad \Lambda_{\text{band}}^{(\text{mir}|\text{self})}, \quad \bar{\Phi}_{\text{band}}^{(\text{mir})}.$$

In particular, Proposition O38 shows that the ordinary mirror-band mean is already saturated exactly by the same-center band mean and the mirror/self carrier ratio at the common knot $D = 20$, so no further support/object refinement is needed on this line.

Only secondary tightening remains in:

$$\bar{K}_{\text{band}}^{(-)}, \quad \bar{\rho}_{\text{band}}^{(-)}, \quad \theta_{\text{band}}^{(-)}, \quad \Xi_{\text{band}}^{(\text{self})}, \quad \frac{N_{\text{mass,core}}}{P_{\text{mass,in}}}, \quad \frac{N_{\text{mass,rim}}}{P_{\text{mass,in}}}, \quad \frac{N_{\text{mass,tail}}}{P_{\text{mass,in}}},$$

with the outer same-sign parity-contrast remainder $I_{\times}$ from Proposition O19 also retained only as a small static correction. Therefore the mainline task is no longer to discover a new support or a new overlap object. The chain is already structurally closed; what remains, if any further improvement is desired, is only constant-tightening inside the already identified multiplicative and leakage factors.

Benchmark O1 (Dense-family sharpness check). The dense exact-family scan should now be read as a sharpness benchmark rather than as the theorem input. Resolving the same family on

$$\alpha, \beta \in [0.95, 1.05], \quad \lambda \in [-0.2, 0.2],$$

with step ‘0.005’ and imposing exact admissibility $J_{\text{parent}} \leq 10^{-6}$ shows that the extremal ridge sits near

$$(\alpha, \beta, \lambda) \approx (1.03, 0.99, \pm 0.2),$$

giving

$$u_a \leq 5.42 \times 10^{-3}, \quad \text{tr } A_a \leq 5.42 \times 10^{-3}, \quad \|A_a\| \leq 5.42 \times 10^{-3},$$

with

$$H_{\text{eff},a} \succeq 0.9946 H_{PP,a}, \quad |\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.09 \times 10^{-2}.$$

This is the tightest local benchmark currently available, but it is only about one percent above Proposition O5. Even the conservative raw off-diagonal envelope remains moderate on the same dense admissible window, giving only

$$u_a \leq 7.19 \times 10^{-3}, \quad |\log \det H_{\text{eff},a} - \log \det H_{PP,a}| \leq 1.45 \times 10^{-2}.$$

So the right reading is now stable: the dense exact-family scan sets the benchmark, while the continuum neighborhood bound is the theorem candidate to push beyond the scanned family.

Proposition O40 (Synthesis of the dominant-pole source chain on the audited exact bridge). Combining Propositions O1–O9, Corollaries O1–O8, and Benchmark O1, the dominant-pole line on the current audited exact bridge separates into three structural blocks:

on-shell defect closure, determinant-deficit gap budget, exact operator witness-envelope inclusion.

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_dominant_pole_synthesis_source.py` compresses these blocks into one reviewer-facing summary on the exported D4–20 knot set. First, the defect side is already closed:

$$\max_a |\delta_{q,a}| = 4.44 \times 10^{-16}, \quad \max_a \left| \tau_{\text{env},a} - \omega_{1,a}^{-1} \right| = 1.11 \times 10^{-16}.$$

Second, the projected continuum-neighborhood route still gives the same determinant-deficit budget and lifted exact-gap floor as the later exact-lift audit:

$$u_a^{(\text{neigh})} \leq 5.34 \times 10^{-3}, \quad g_a^{(\text{exact})} \geq 1.3339, \quad H_{\text{eff},a} \succeq 0.9947 H_{PP,a},$$

with consistency residuals

$$\left| g_{a,\text{lift}}^{(\text{exact})} - g_{a,\text{neigh}}^{(\text{exact})} \right| = 0, \quad |\eta_{\text{eff},\text{lift}} - \eta_{\text{eff},\text{neigh}}| = 0,$$

at the exported precision. Third, the exact operator already lies deep inside the audited safe chart:

$$\mathcal{B}_{\text{op}}^{(\text{exact})} = 3.63 \times 10^{-7} = \|W_{\text{max}}^{(\text{exact})}\|_{\infty},$$

with the envelope attained by the exact factorization-side block-22 witness at $D = 13$, and

$$\mathcal{B}_{\text{op}}^{(\text{exact})} \leq \frac{1}{2} r_{95}^{\text{safe}} = 4.81 \times 10^{-5}, \quad \mathcal{B}_{\text{op}}^{(\text{exact})} \leq \frac{1}{2} r_{\text{max}}^{\text{safe}} = 5.03 \times 10^{-5},$$

by factors

$$1.33 \times 10^2, \quad 1.39 \times 10^2.$$

Meanwhile the family-transfer audit still gives

$$N_{\text{adm}}^{(\text{total})} = 1, \quad N_{\text{adm}}^{(\text{noncan})} = 0$$

at $J_* = 10^{-6}$, with the minimum selection gap exceeding that theorem tolerance by the factor

$$6.90 \times 10^2.$$

Thus, on the current exact bridge, the defect side, the family/chart transfer side, and the exact witness-envelope side are already mutually aligned: the only missing step is no longer to discover a new response family or a new projected interaction channel, but to promote the same determinant-deficit budget from audited evidence to an a priori exact Schur-resolvent lift.

Corollary O39 (What is closed, and what remains as the single dominant-pole gap).

On the current audited exact bridge, the dominant-pole / Schur-resolvent line now separates cleanly into *closed* structural blocks and *one* remaining theorem-level lift.

Closed on the audited exact bridge:

$$P\Delta VP \in \text{span}\{I, \sigma_z, \sigma_x\}, \quad \delta_{q,a} \approx 0, \quad \text{family/chart transfer}, \quad \mathcal{B}_{\text{op}}^{(\text{exact})} \ll r_{\text{safe}}/2.$$

Equivalently, the localized two-level reduction, the gap-locked single-pole response family, the on-shell defect closure, and the exact operator witness inclusion are no longer substantive blockers on the current knot set.

The single remaining theorem-level gap is:

$$\text{upgrade } u_a^{(\text{neigh})} \leq 5.34 \times 10^{-3} \implies g_a^{(\text{exact})} \geq 1.3339$$

from the present audited projected-family / chart evidence to a genuine parent-side analytic lift for the exact Schur resolvent. Once that step is made, Theorem Target O1 is closed on the current on-shell window without introducing any new support, response family, or overlap object.

Proposition O41 (A single scalar sufficient condition for the exact Schur lift). The remaining theorem-level lift of Corollary O39 can be reduced one step further to a single scalar lower floor on the mix kernel. Define

$$m_{\text{mix},*} := \frac{\varepsilon_{\text{mix}}}{\frac{1}{2}r_{95}^{\text{safe}}}, \quad \varepsilon_{\text{mix}} := \max_D |\gamma_{\text{mix}}(D) - \kappa_{\text{env}} g_x(D)^2 S_{xx}(D)|.$$

On the current audited exact bridge,

$$\varepsilon_{\text{mix}} = 4.40 \times 10^{-14}, \quad \frac{1}{2}r_{95}^{\text{safe}} = 4.81 \times 10^{-5},$$

so

$$m_{\text{mix},*} = 9.15 \times 10^{-10}.$$

Suppose a parent-side analytic argument supplies the uniform lower floor

$$g_x(D)^2 S_{xx}(D) \geq m_{\text{mix},*}$$

throughout the audited on-shell window. Then Proposition O9 gives

$$\|\Delta_{\text{mix}}^{(\text{fact})}\|_{\infty} \leq \frac{\varepsilon_{\text{mix}}}{m_{\text{mix},*}} \leq \frac{1}{2}r_{95}^{\text{safe}}.$$

Meanwhile all non-mix witness components are already much smaller:

$$\mathcal{B}_{\text{op}}^{(\text{nonmix})} := \max\{\Delta_{\phi}^{(\text{fact})}, \Delta_{\text{tr}}^{(\text{stmt})}, \Delta_{\text{det}}^{(\text{stmt})}, \Delta_{\text{ani}}^{(\text{stmt})}, \Delta_{\text{mix}}^{(\text{cocycle})}, \Delta_{\text{mix}}^{(\text{potential})}\} = 2.33 \times 10^{-10},$$

so

$$\frac{\frac{1}{2}r_{95}^{\text{safe}}}{\mathcal{B}_{\text{op}}^{(\text{nonmix})}} = 2.06 \times 10^5.$$

Hence the exact witness envelope satisfies

$$\mathcal{B}_{\text{op}}^{(\text{exact})} \leq \frac{1}{2}r_{95}^{\text{safe}},$$

the exact operator lies in the audited safe chart by Corollary O7, and the determinant-deficit budget of Proposition O5 transfers to the exact Schur-reduced operator. Therefore

$$u_a \leq 5.34 \times 10^{-3}, \quad g_a^{(\text{exact})} \geq 1.3339, \quad H_{\text{eff},a} \succeq 0.9947 H_{PP,a},$$

and, together with Corollary O2, Theorem Target O1 is closed on the current on-shell window.

Corollary O40 (The exact Schur lift has been reduced to one scalar floor). On the current audited exact bridge, the only remaining analytic input needed for the dominant-pole / exact-Schur lift is any one parent-side proof of a mix-kernel floor at or above

$$g_x(D)^2 S_{xx}(D) \geq 9.15 \times 10^{-10}.$$

Equivalently, since $\kappa_{\text{env}}^{\text{max}} = 4.31 \times 10^5$ on the same knot set, it is enough to prove

$$\gamma_{\text{mix}}(D) \geq 3.94 \times 10^{-4}.$$

The new audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_lift_scalar_source.py` shows that the current exported exact bridge already exceeds this threshold in several independent ways:

$$\min_D g_x(D)^2 S_{xx}(D) = 1.2147 \times 10^{-9}, \quad \frac{\min_D g_x(D)^2 S_{xx}(D)}{m_{\text{mix},*}} = 1.328,$$

$$\frac{\min_D \gamma_{\text{mix}}(D)}{\kappa_{\text{env}}^{\text{max}}} = 1.2147 \times 10^{-9}, \quad \frac{\min_D \gamma_{\text{mix}}(D)/\kappa_{\text{env}}^{\text{max}}}{m_{\text{mix},*}} = 1.328,$$

and even the separated observable-chain floor already gives

$$\frac{\gamma_{\text{mix}}^{(\text{sep floor})}}{\kappa_{\text{env}}^{\text{max}}} = 1.0016 \times 10^{-9}, \quad \frac{\gamma_{\text{mix}}^{(\text{sep floor})}/\kappa_{\text{env}}^{\text{max}}}{m_{\text{mix},*}} = 1.095.$$

Thus the exact Schur-lift problem is no longer a multi-parameter chart or support problem: it has been reduced to one scalar analytic floor on the mix kernel (equivalently on γ_{mix}), while the defect side, the family/chart transfer side, and the non-mix witness side are already closed on the audited exact bridge.

Proposition O42 (The remaining scalar Schur-lift gap is only a small tightening of the current static parity floor). The scalar threshold of Proposition O41 can be compared directly with the current overlap-side lower bounds. By Corollary O20,

$$|V_{11}(D) - V_{22}(D)| \geq 5.64043 \times 10^{-5}$$

on the audited exact bridge, hence by Proposition O18

$$g_x(D) \geq 2.820215 \times 10^{-5}.$$

Combining this with the global response floor from Proposition O17,

$$S_{xx}(D) \geq S_{xx}^{\text{floor}} = 1.01493,$$

gives the exact-bridge lower estimate

$$\gamma_{\text{mix}}(D) \geq \kappa_{\text{env}} (2.820215 \times 10^{-5})^2 S_{xx}^{\text{floor}} = 3.47553 \times 10^{-4}.$$

The new audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_lift_overlap_bridge.py` compares this directly with the scalar Schur-lift threshold from Proposition O41,

$$\gamma_{\text{mix},*} = 3.93829 \times 10^{-4},$$

and finds

$$\frac{3.47553 \times 10^{-4}}{3.93829 \times 10^{-4}} = 0.88250.$$

So with the global response floor the current compact-anchor parity-contrast theorem already supplies 88.25% of the scalar Schur-lift threshold. Equivalently, if the same global response floor is kept fixed, only the further tightening

$$g_x(D) \geq 1.06450 \times (2.820215 \times 10^{-5})$$

would be sufficient to close Proposition O41.

If one instead keeps the exact local response factor at the common knot $D = 11$,

$$S_{xx}(D_{g_x, \text{min}}) = 1.07984,$$

then the same parity floor already gives

$$\gamma_{\text{mix}}(D) \geq 3.69779 \times 10^{-4}, \quad \frac{3.69779 \times 10^{-4}}{3.93829 \times 10^{-4}} = 0.93893,$$

so the remaining gap is only

$$g_x(D) \geq 1.03201 \times (2.820215 \times 10^{-5})$$

under that local-response alignment. Likewise, using the slightly stronger core-window parity floor

$$|V_{11}(D) - V_{22}(D)| \geq 5.70400 \times 10^{-5}$$

from Corollary O19 improves the global-response fraction to

$$0.90250$$

and the local-response fraction to

$$0.96022.$$

Corollary O41 (The exact Schur-lift gap is no longer a new operator question, but a small constant-tightening question on the static overlap side). On the current exact bridge, Propositions O41–O42 show that the remaining dominant-pole / exact-Schur lift gap is not a missing chart inclusion, not a missing response family, and not a missing non-mix operator estimate. It is already reduced to a small scalar constant gap between the current static parity-contrast floor and the threshold

$$g_x(D)^2 S_{xx}(D) \geq 9.15 \times 10^{-10}.$$

With the globally valid response floor from Proposition O17, the remaining deficit is only a

$$6.45\%$$

tightening of the compact-anchor g_x floor, and even on the exact local-response alignment it drops to only

$$3.20\%.$$

So the next useful proof-side task is now very specific: strengthen the current one-center/static parity floor just enough to cross the scalar threshold of Proposition O41, rather than reopening any new Schur-resolvent structure.

Proposition O43 (The scalar Schur-lift criterion is equivalent to a single static parity threshold). Proposition O41 asks only for the scalar floor

$$g_x(D)^2 S_{xx}(D) \geq m_{\text{mix},*}, \quad m_{\text{mix},*} = 9.14719 \times 10^{-10}.$$

By Proposition O18,

$$g_x(D) = \frac{1}{2} |V_{11}(D) - V_{22}(D)|.$$

Therefore any response lower bound

$$S_{xx}(D) \geq S_*$$

implies the sufficient static parity threshold

$$|V_{11}(D) - V_{22}(D)| \geq 2\sqrt{\frac{m_{\text{mix},*}}{S_*}}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_lift_parity_threshold_source.py` evaluates this threshold for the two response inputs already isolated in Propositions O17 and O41. With the globally valid response floor

$$S_* = S_{xx}^{\text{floor}} = 1.01493,$$

the exact Schur lift follows from the single static bound

$$|V_{11}(D) - V_{22}(D)| \geq 6.00420 \times 10^{-5}.$$

With the exact local-response alignment at the common knot $D = 11$,

$$S_* = S_{xx}(D_{g_x, \min}) = 1.07984,$$

the required static threshold sharpens to

$$|V_{11}(D) - V_{22}(D)| \geq 5.82096 \times 10^{-5}.$$

These two thresholds can then be compared directly with the currently available overlap-side floors:

$$|V_{11} - V_{22}|_{\text{anchor}} = 5.64043 \times 10^{-5}, \quad |V_{11} - V_{22}|_{\text{core}} = 5.70400 \times 10^{-5},$$

$$|V_{11} - V_{22}|_{\text{exact}} = 6.70795 \times 10^{-5}.$$

Numerically, the current compact-anchor theorem already reaches

$$93.94\% \quad \text{of the global threshold,} \quad 96.90\% \quad \text{of the local-response threshold,}$$

while the current core-window theorem reaches

$$95.00\% \quad \text{of the global threshold,} \quad 97.99\% \quad \text{of the local-response threshold.}$$

Meanwhile the exact bridge floor already exceeds the same thresholds by the factors

$$1.117 \quad (\text{global-response version}), \quad 1.152 \quad (\text{local-response version}).$$

Corollary O42 (The remaining dominant-pole theorem gap is exactly a one-scalar parity-threshold tightening). Propositions O41–O43 show that the residual exact Schur-lift task can now be stated with no additional response-family or chart language. It is enough to strengthen the current static parity theorem to one of the two scalar thresholds

$$|V_{11}(D) - V_{22}(D)| \geq 6.00420 \times 10^{-5}$$

under the globally valid response floor, or

$$|V_{11}(D) - V_{22}(D)| \geq 5.82096 \times 10^{-5}$$

under the exact local-response alignment already audited at the common floor knot. Relative to the current core-window theorem floor 5.70400×10^{-5} , these correspond to only

$$5.26\% \quad \text{and} \quad 2.05\%$$

of additional static-parity tightening, respectively. Thus the open-system dominant-pole line is no longer missing any new Schur-resolvent structure: it has been reduced to a single explicit constant threshold on the same parity contrast $|V_{11} - V_{22}|$ that already drives the overlap-side synthesis.

Proposition O44 (The same core-window family already gives a tighter floor for the exact total parity contrast). Keep the same symmetric core-window family $B_{\text{core}}(R; D)$ from Proposition O20, but apply it directly to the full exact parity-contrast density:

$$J_{\text{core}}^{(R)}(D) := \int_{B_{\text{core}}(R; D)} \mathcal{C}(\rho, z; D) \delta V(\rho, z; D) d\rho dz,$$

$$J_{\text{tail}}^{(R)}(D) := (V_{11}(D) - V_{22}(D)) - J_{\text{core}}^{(R)}(D).$$

Thus

$$V_{11}(D) - V_{22}(D) = J_{\text{core}}^{(R)}(D) + J_{\text{tail}}^{(R)}(D).$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_total_core_tightening_source.py` rescans the same interval

$$R \in [1.20, 1.40]$$

with the exact total density instead of the one-center pair alone. It finds that the same optimal safe radius remains

$$R_* = 1.35,$$

and now satisfies on every audited knot $D = 4, \dots, 20$

$$\operatorname{sgn} J_{\text{core}}^{(R_*)}(D) = \operatorname{sgn}(V_{11}(D) - V_{22}(D)),$$

$$\operatorname{sgn} J_{\text{tail}}^{(R_*)}(D) = -\operatorname{sgn}(V_{11}(D) - V_{22}(D)).$$

Hence

$$|V_{11}(D) - V_{22}(D)| = |J_{\text{core}}^{(R_*)}(D)| - |J_{\text{tail}}^{(R_*)}(D)| \geq \lambda_{\text{tot},*} |J_{\text{core}}^{(R_*)}(D)|,$$

with

$$\lambda_{\text{tot},*} = 0.426709\dots$$

and

$$\min_D |J_{\text{core}}^{(R_*)}(D)| = 1.34704 \times 10^{-4}.$$

Therefore the same core-window family already yields the sharper exact-total floor

$$|V_{11}(D) - V_{22}(D)| \geq 5.74795 \times 10^{-5}.$$

This improves the previous core-window theorem floor 5.70400×10^{-5} by the factor

$$1.00771.$$

The exact bridge floor remains

$$\min_D |V_{11}(D) - V_{22}(D)| = 6.70795 \times 10^{-5} \quad \text{at } D = 11,$$

so the new same-family tightening is already within the factor

$$1.16702$$

of the exact minimum.

Corollary O43 (The remaining Schur-lift gap drops from 2.05% to 1.27% without changing the support family). Combine Proposition O44 with the parity-threshold reduction of Proposition O43. Under the exact local-response alignment at the common floor knot, the Schur-lift target is

$$|V_{11}(D) - V_{22}(D)| \geq 5.82096 \times 10^{-5}.$$

The tightened exact-total core-window floor from Proposition O44 already reaches

$$\frac{5.74795 \times 10^{-5}}{5.82096 \times 10^{-5}} = 0.987457,$$

so only the remaining tightening

$$1.27\%$$

is still missing on the local-response version of the scalar Schur threshold. Even against the globally valid response threshold 6.00420×10^{-5} , the same tightened floor already reaches

$$95.73\%$$

of the required static parity contrast. Thus the dominant-pole / exact-Schur line remains on the same static-overlap spine as before, but the remaining gap is now smaller still: no new support family is needed, and the residual task is only a final sub-two-percent tightening of the current parity theorem when read at the locally aligned response level.

Proposition O45 (The same family and the same radius already close the Schur threshold after a contiguous two-window split). Keep exactly the same symmetric core-window family $B_{\text{core}}(R; D)$ and the same optimal radius

$$R_* = 1.35$$

from Proposition O44. Do not change the support family or the radius; instead, only split the audited knot window into the two contiguous subwindows

$$\mathcal{D}_L = \{4, \dots, 11\}, \quad \mathcal{D}_R = \{12, \dots, 20\}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_total_core_window_split_source.py` reuses the exact-total core/tail data from Proposition O44 and shows that the same-family exact split

$$V_{11}(D) - V_{22}(D) = J_{\text{core}}^{(R_*)}(D) + J_{\text{tail}}^{(R_*)}(D)$$

remains sign-safe on both subwindows. On the left window one has

$$\lambda_{\text{tot},*}^{(\text{L})} = 0.479785\dots, \quad \min_{D \in \mathcal{D}_{\text{L}}} |J_{\text{core}}^{(R_*)}(D)| = 1.34704 \times 10^{-4},$$

so

$$\min_{D \in \mathcal{D}_{\text{L}}} |V_{11}(D) - V_{22}(D)| \geq 6.46291 \times 10^{-5}.$$

On the right window one has

$$\lambda_{\text{tot},*}^{(\text{R})} = 0.426709\dots, \quad \min_{D \in \mathcal{D}_{\text{R}}} |J_{\text{core}}^{(R_*)}(D)| = 3.25919 \times 10^{-4},$$

so

$$\min_{D \in \mathcal{D}_{\text{R}}} |V_{11}(D) - V_{22}(D)| \geq 1.39072 \times 10^{-4}.$$

Therefore the same family and the same radius already yield the audited piecewise exact-total floor

$$|V_{11}(D) - V_{22}(D)| \geq 6.46291 \times 10^{-5} \quad \text{for all audited } D \in \mathcal{D}_{\text{L}} \cup \mathcal{D}_{\text{R}},$$

with no new support family and no new response object.

Corollary O44 (The remaining 1.27% gap is only a global-coupling artifact of the single-window product bound). Compare Proposition O45 with the parity thresholds from Proposition O43:

$$|V_{11}(D) - V_{22}(D)| \geq 6.00420 \times 10^{-5}$$

under the globally valid response floor and

$$|V_{11}(D) - V_{22}(D)| \geq 5.82096 \times 10^{-5}$$

under the exact local-response alignment. The same-family two-window floor from Proposition O45 satisfies

$$\frac{6.46291 \times 10^{-5}}{6.00420 \times 10^{-5}} = 1.07640, \quad \frac{6.46291 \times 10^{-5}}{5.82096 \times 10^{-5}} = 1.11028.$$

Hence the dominant-pole / exact-Schur threshold is already closed on the audited bridge without changing the family $B_{\text{core}}(R)$ and without changing the optimal radius $R_* = 1.35$. The previous residual 1.27% gap was therefore not evidence for a missing Schur-resolvent structure; it was only the loss incurred by forcing one global minimum $\lambda_{\text{tot},*}$ and one global core floor to be multiplied across the entire audited window instead of allowing the same split family to be read piecewise on its natural contiguous subwindows.

Proposition O46 (The parent-side analytic lift has been reduced to the same family and two windowwise core/tail bounds). The audited piecewise closure of Proposition O45 can be promoted to a parent-side theorem target without introducing any new family. Keep exactly the same split

$$V_{11}(D) - V_{22}(D) = J_{\text{core}}^{(R_*)}(D) + J_{\text{tail}}^{(R_*)}(D), \quad R_* = 1.35,$$

on the same contiguous windows

$$\mathcal{D}_L = \{4, \dots, 11\}, \quad \mathcal{D}_R = \{12, \dots, 20\}.$$

Suppose a parent-side analytic argument proves, for every $D \in \mathcal{D}_L$,

$$\text{sgn } J_{\text{core}}^{(R_*)}(D) = \text{sgn}(V_{11}(D) - V_{22}(D)),$$

$$\text{sgn } J_{\text{tail}}^{(R_*)}(D) = -\text{sgn}(V_{11}(D) - V_{22}(D)),$$

$$|J_{\text{tail}}^{(R_*)}(D)| \leq \eta_L |J_{\text{core}}^{(R_*)}(D)|, \quad \eta_L = 0.520215\dots,$$

$$|J_{\text{core}}^{(R_*)}(D)| \geq C_L, \quad C_L = 1.34704 \times 10^{-4},$$

and, for every $D \in \mathcal{D}_R$,

$$|J_{\text{tail}}^{(R_*)}(D)| \leq \eta_R |J_{\text{core}}^{(R_*)}(D)|, \quad \eta_R = 0.573291\dots,$$

$$|J_{\text{core}}^{(R_*)}(D)| \geq C_R, \quad C_R = 3.25919 \times 10^{-4}.$$

Then

$$|V_{11}(D) - V_{22}(D)| \geq (1 - \eta_L)C_L = 6.46291 \times 10^{-5} \quad (D \in \mathcal{D}_L),$$

and

$$|V_{11}(D) - V_{22}(D)| \geq (1 - \eta_R)C_R = 1.39072 \times 10^{-4} \quad (D \in \mathcal{D}_R).$$

Therefore the same family and the same radius imply the global piecewise floor

$$|V_{11}(D) - V_{22}(D)| \geq 6.46291 \times 10^{-5}$$

throughout the audited on-shell window. By Proposition O43, this already exceeds both the global and the locally aligned parity thresholds required for the scalar exact-Schur lift.

Corollary O45 (What remains is no longer family design, but only a parent-side proof of the same piecewise bounds). The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_piecewise_parent_lift_source.py` shows that the constants in Proposition O46 are already realized on the exported bridge:

$$\eta_L = 0.520215\dots, \quad C_L = 1.34704 \times 10^{-4}, \quad (1 - \eta_L)C_L = 6.46291 \times 10^{-5},$$

$$\eta_R = 0.573291\dots, \quad C_R = 3.25919 \times 10^{-4}, \quad (1 - \eta_R)C_R = 1.39072 \times 10^{-4}.$$

The left window is the only bottleneck, while the right window already exceeds the global parity threshold by the factor

$$2.31625.$$

Hence ‘open_system_micro’ should now be read in its sharpest current form: it is *audited piecewise closed* on the same exact-total split, and the only remaining theorem task is to prove those same windowwise core/tail inequalities directly on the parent side. No new response family, no new support class, and no new chart design remain on this line.

Proposition O47 (The remaining parent-side task is only a left-window bottleneck bound, equivalent to either a tail-ratio theorem or a core-floor theorem). Corollary O45 shows that the right window $\mathcal{D}_R = \{12, \dots, 20\}$ is already non-bottleneck. Thus the only remaining parent-side task on the same family $B_{\text{core}}(R_*)$ comes from the left window

$$\mathcal{D}_L = \{4, \dots, 11\}, \quad R_* = 1.35.$$

Write

$$\eta_L := \sup_{D \in \mathcal{D}_L} \frac{|J_{\text{tail}}^{(R_*)}(D)|}{|J_{\text{core}}^{(R_*)}(D)|}, \quad C_L := \inf_{D \in \mathcal{D}_L} |J_{\text{core}}^{(R_*)}(D)|.$$

Then the left-window parity floor is

$$|V_{11}(D) - V_{22}(D)| \geq (1 - \eta_L) C_L \quad (D \in \mathcal{D}_L),$$

provided the same-sign split of Proposition O46 holds. Therefore the parity thresholds from Proposition O43 are equivalent to either of the following two parent-side sufficient conditions on the *same* left-window family:

(i) *Tail-ratio target with the current left core floor fixed.* If one keeps the current audited left core floor

$$C_L = 1.34704 \times 10^{-4},$$

then it is enough to prove

$$\eta_L \leq 0.554268$$

for the globally valid Schur-parity threshold, or

$$\eta_L \leq 0.567871$$

for the locally aligned threshold. The current audited value is

$$\eta_L^{(\text{aud})} = 0.520215,$$

so the required slack is only

$$0.034053 \quad (\text{global}), \quad 0.047656 \quad (\text{local}).$$

(ii) *Core-floor target with the current left tail ratio fixed.* If one keeps the current audited left tail ratio

$$\eta_L = 0.520215,$$

then it is enough to prove

$$C_L \geq 1.25144 \times 10^{-4}$$

for the globally valid threshold, or

$$C_L \geq 1.21325 \times 10^{-4}$$

for the locally aligned threshold. The current audited value

$$C_L^{(\text{aud})} = 1.34704 \times 10^{-4}$$

already exceeds these by the factors

$$1.07640 \quad \text{and} \quad 1.11028.$$

Hence the remaining parent-side analytic lift is no longer a whole-window family problem. On the same exact-total split, it has been reduced to a single left-window bottleneck, which can be formulated equally well as a tail-ratio theorem or as a core-floor theorem.

Corollary O46 (The right window is formally demoted, and the left window alone sets the remaining parent-side proof target). The audit /Users/boypatrick/codex/PS

LT_quantam/code/audit_open_system_exact_schur_left_window_bottleneck_source.py shows that

$$\frac{(1 - \eta_R)C_R}{6.00420 \times 10^{-5}} = 2.31625, \quad \frac{(1 - \eta_R)C_R}{5.82096 \times 10^{-5}} = 2.38917,$$

so the right window $\mathcal{D}_R$ is already uniformly far above both Schur-parity thresholds. Meanwhile the left window saturates the audited piecewise floor:

$$\frac{(1 - \eta_L)C_L}{6.00420 \times 10^{-5}} = 1.07640, \quad \frac{(1 - \eta_L)C_L}{5.82096 \times 10^{-5}} = 1.11028.$$

Therefore, from the standpoint of the remaining parent-side analytic lift, the right window is formally demoted to a non-bottleneck sector. The only substantive task left on ‘open_system_micro’ is to prove one of the equivalent left-window bounds of Proposition O47 on the same family $B_{\text{core}}(R_*)$ and the same exact-total split.

Proposition O48 (Even the left-window bottleneck is a coupling artifact: the same-family split localizes it to $D \leq 10$ plus the single core-floor knot $D = 11$). Keep the same family $B_{\text{core}}(R_*)$ with the same radius $R_* = 1.35$, but now split the left window

$$\mathcal{D}_L = \{4, \dots, 11\}$$

into the two contiguous subwindows

$$\mathcal{D}_{L1} = \{4, \dots, 10\}, \quad \mathcal{D}_{L2} = \{11\}.$$

The audit /Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_left_window_split_source.py shows that the best same-family split occurs exactly at

$$D_{\text{cut}} = 10.$$

On $\mathcal{D}_{L1}$ the tail-ratio bottleneck remains

$$\eta_{L1}^{\max} = 0.520215 \dots \quad \text{at } D = 9,$$

but the core floor improves to

$$\inf_{D \in \mathcal{D}_{L1}} |J_{\text{core}}^{(R_*)}(D)| = 1.54417 \times 10^{-4},$$

so the same-family induced floor there is

$$|V_{11}(D) - V_{22}(D)| \geq 7.40868 \times 10^{-5} \quad (D \in \mathcal{D}_{L1}).$$

On the singleton subwindow $\mathcal{D}_{L2} = \{11\}$, the same-family split gives

$$\eta_{L2} = 0.502024 \dots, \quad |J_{\text{core}}^{(R_*)}(11)| = 1.34704 \times 10^{-4},$$

hence

$$|V_{11}(11) - V_{22}(11)| \geq 6.70795 \times 10^{-5}.$$

Therefore the same-family left-window split already yields the audited floor

$$|V_{11}(D) - V_{22}(D)| \geq 6.70795 \times 10^{-5} \quad (D \in \mathcal{D}_{L1} \cup \mathcal{D}_{L2}),$$

which exceeds the global and local Schur-parity thresholds by the factors

$$1.11721 \quad \text{and} \quad 1.15238.$$

In particular, the former left-window bottleneck is itself only a coupling artifact: the tail-ratio maximum occurs at $D = 9$, while the core-floor minimum occurs at $D = 11$, and forcing them into one left-window product is what creates the residual slack.

Corollary O47 (The only meaningful parent-side target left is a theorem for the same-family tail ratio on $D = 4, \dots, 10$, together with the already-isolated core-floor knot at $D = 11$). The same audit shows

$$\eta_L^{\max} \text{ occurs at } D = 9, \quad \inf_{\mathcal{D}_L} |J_{\text{core}}^{(R_*)}| \text{ occurs at } D = 11,$$

so the left-window bottleneck no longer behaves as a single coherent interval obstruction. After the same-family split at $D_{\text{cut}} = 10$, the subwindow $\mathcal{D}_{L1} = \{4, \dots, 10\}$ already clears both Schur-parity thresholds by itself, and the singleton knot $D = 11$ carries the audited left exact minimum 6.70795×10^{-5} . Therefore the remaining parent-side analytic work can be read in its narrowest practical form:

1. prove the same-family tail/core inequality on the left subwindow $D = 4, \dots, 10$, where the true ratio bottleneck lives;
2. treat $D = 11$ as the already isolated core-floor knot rather than as part of a larger family-design problem.

No new family, no new support, and no new chart structure are left even on the former bottleneck window; what remains is a localized parent-side estimate on the same exact-total split.

Proposition O49 (The residual $D = 4, \dots, 10$ target is itself a coupling artifact: the same-family hierarchical split leaves only the singleton knot $D = 11$ as the audited

bottleneck). Keep exactly the same family $B_{\text{core}}(R_*)$ with the same radius $R_* = 1.35$, and refine the left subwindow

$$\mathcal{D}_{L1} = \{4, \dots, 10\}$$

into the two contiguous windows

$$\mathcal{D}_{L1a} = \{4, 5, 6\}, \quad \mathcal{D}_{L1b} = \{7, \dots, 10\}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_left_window_hierarchical_split_source.py` shows that the best secondary same-family cut occurs at

$$D_{\text{cut}}^{(2)} = 6.$$

On $\mathcal{D}_{L1a}$, the tail-ratio maximum remains

$$\eta_{L1a}^{\max} = 0.469739\dots \quad \text{at } D = 4,$$

while the core floor is now

$$\inf_{D \in \mathcal{D}_{L1a}} |J_{\text{core}}^{(R_*)}(D)| = 1.54417 \times 10^{-4} \quad \text{at } D = 6,$$

giving the same-family induced floor

$$|V_{11}(D) - V_{22}(D)| \geq 8.18812 \times 10^{-5} \quad (D \in \mathcal{D}_{L1a}).$$

On $\mathcal{D}_{L1b}$, the tail-ratio maximum occurs at

$$\eta_{L1b}^{\max} = 0.520215\dots \quad \text{at } D = 9,$$

but the core floor is already much larger,

$$\inf_{D \in \mathcal{D}_{L1b}} |J_{\text{core}}^{(R_*)}(D)| = 7.72912 \times 10^{-4},$$

so the same-family induced floor there is

$$|V_{11}(D) - V_{22}(D)| \geq 3.70831 \times 10^{-4} \quad (D \in \mathcal{D}_{L1b}).$$

Together with the already isolated singleton knot $D = 11$,

$$|V_{11}(11) - V_{22}(11)| \geq 6.70795 \times 10^{-5},$$

the overall audited hierarchical floor on $D = 4, \dots, 11$ remains

$$|V_{11}(D) - V_{22}(D)| \geq 6.70795 \times 10^{-5},$$

which exceeds the global and local Schur-parity thresholds by the same factors

$$1.11721 \quad \text{and} \quad 1.15238.$$

Therefore even the residual $D = 4, \dots, 10$ target is only a coupling artifact: after the same-family hierarchical split, every non-singleton subwindow clears the Schur threshold by a comfortable margin, and the audited minimum is carried solely by the isolated knot $D = 11$.

Corollary O48 (On the audited bridge, the only surviving bottleneck is the single core-floor knot $D = 11$; all interval pieces are formally non-bottleneck on the same family). The same audit gives

$$\frac{8.18812 \times 10^{-5}}{6.00420 \times 10^{-5}} = 1.36373, \quad \frac{8.18812 \times 10^{-5}}{5.82096 \times 10^{-5}} = 1.40666$$

on $\mathcal{D}_{L1a} = \{4, 5, 6\}$, and

$$\frac{3.70831 \times 10^{-4}}{6.00420 \times 10^{-5}} = 6.17693, \quad \frac{3.70831 \times 10^{-4}}{5.82096 \times 10^{-5}} = 6.37084$$

on $\mathcal{D}_{L1b} = \{7, \dots, 10\}$. Since the singleton floor

$$6.70795 \times 10^{-5}$$

at $D = 11$ is smaller than both interval floors, the audited bridge now has a uniquely isolated same-family bottleneck. Thus the remaining parent-side analytic task can be read in its narrowest current form:

1. the interval pieces $D = 4, \dots, 6$ and $D = 7, \dots, 10$ are formally demoted to non-bottleneck sectors on the same exact-total split;
2. the only surviving audited bottleneck is the single-knot core-floor estimate at $D = 11$.

Proposition O50 (After the audited hierarchical split, the remaining parent-side lift is equivalent to a single-knot core-floor theorem at $D = 11$). Keep the same exact-total split and the same family $B_{\text{core}}(R_*)$ with $R_* = 1.35$. At the isolated bottleneck knot $D = 11$, the audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_knot_source.py` shows

$$V_{11}(11) - V_{22}(11) = J_{\text{core}}^{(R_*)}(11) + J_{\text{tail}}^{(R_*)}(11),$$

with

$$\text{sgn } J_{\text{core}}^{(R_*)}(11) = \text{sgn}(V_{11}(11) - V_{22}(11)), \quad \text{sgn } J_{\text{tail}}^{(R_*)}(11) = -\text{sgn}(V_{11}(11) - V_{22}(11)),$$

and audited ratio

$$\eta_{11} := \frac{|J_{\text{tail}}^{(R_*)}(11)|}{|J_{\text{core}}^{(R_*)}(11)|} = 0.502024\dots, \quad \lambda_{11} := 1 - \eta_{11} = 0.497976\dots$$

Hence

$$|V_{11}(11) - V_{22}(11)| = \lambda_{11} |J_{\text{core}}^{(R_*)}(11)|.$$

To imply the global Schur-parity threshold

$$|V_{11} - V_{22}| \geq 6.00420 \times 10^{-5},$$

it is enough to prove the single-knot core-floor estimate

$$|J_{\text{core}}^{(R_*)}(11)| \geq \frac{6.00420 \times 10^{-5}}{\lambda_{11}} = 1.20572 \times 10^{-4}.$$

To imply the locally aligned Schur-parity threshold

$$|V_{11} - V_{22}| \geq 5.82096 \times 10^{-5},$$

it is enough to prove

$$|J_{\text{core}}^{(R_*)}(11)| \geq \frac{5.82096 \times 10^{-5}}{\lambda_{11}} = 1.16893 \times 10^{-4}.$$

The current audited exact value is

$$|J_{\text{core}}^{(R_*)}(11)| = 1.34704 \times 10^{-4},$$

which exceeds these single-knot targets by the factors

$$1.11721 \quad \text{and} \quad 1.15238.$$

Corollary O49 (The only remaining non-audited lift is a single-knot analytic lower bound at $D = 11$, not any interval theorem or family-design statement). By Proposition O49, every non-singleton interval piece on $D = 4, \dots, 11$ already lies safely above the Schur-parity thresholds on the same family $B_{\text{core}}(R_*)$. Proposition O50 then identifies the residual bottleneck in its sharpest form:

1. no new split, no new chart, and no new support family are needed;
2. no interval theorem remains on the audited bridge;
3. the only remaining parent-side analytic task is to prove a lower bound on the single quantity $|J_{\text{core}}^{(R_*)}(11)|$ strong enough to reach either 1.20572×10^{-4} (global) or 1.16893×10^{-4} (local).

Thus the ‘open_system_micro’ bottleneck has now been reduced from a family-design problem to a single-knot core-floor lift.

Proposition O51 (At the singleton bottleneck knot, the exact core is already carried by the same-center pair core, with only a tiny same-sign cross refinement). At $D = 11$ and the same radius $R_* = 1.35$, write

$$J_{\text{core}}^{(R_*)}(11) = J_{+, \text{core}}^{(R_*)}(11) + J_{-, \text{core}}^{(R_*)}(11) + J_{\times, \text{core}}^{(R_*)}(11),$$

where the first two terms come from the two one-center same-sign pieces and the last term is the cross correction, all integrated on the same core region $B_{\text{core}}(R_*, 11)$. The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_parent_source.py` gives

$$J_{+, \text{core}}^{(R_*)}(11) = 8.68591 \times 10^{-5}, \quad J_{-, \text{core}}^{(R_*)}(11) = 4.69533 \times 10^{-5},$$

so the same-center pair core already contributes

$$J_{\text{self-pair, core}}^{(R_*)}(11) = 1.33812 \times 10^{-4}.$$

Meanwhile the core cross correction is

$$J_{\times, \text{core}}^{(R_*)}(11) = 8.91952 \times 10^{-7},$$

has the same sign as $J_{\text{core}}^{(R_*)}(11)$, and contributes only

$$\frac{|J_{\times, \text{core}}^{(R_*)}(11)|}{|J_{\text{core}}^{(R_*)}(11)|} = 6.62 \times 10^{-3}.$$

Hence the exact bottleneck core satisfies

$$|J_{\text{core}}^{(R_*)}(11)| \geq |J_{\text{self-pair, core}}^{(R_*)}(11)|.$$

In particular, the global and local single-knot targets from Proposition O50 are already implied by the weaker same-center pair inequalities

$$|J_{\text{self-pair, core}}^{(R_*)}(11)| \geq 1.20572 \times 10^{-4}, \quad |J_{\text{self-pair, core}}^{(R_*)}(11)| \geq 1.16893 \times 10^{-4},$$

and the audited same-center pair value exceeds them by the factors

$$1.10981 \quad \text{and} \quad 1.14475.$$

Corollary O50 (The remaining parent-side theorem target can be narrowed once more to a same-center pair-core lower bound at $D = 11$). On the audited bridge, the

singleton bottleneck no longer requires a theorem for the full exact core $J_{\text{core}}^{(R_*)}(11)$. It is enough to prove a lower bound for the same-center pair core $J_{\text{self-pair,core}}^{(R_*)}(11)$ on the same knot and the same radius, because the cross contribution is already same-sign and numerically tiny. Thus the residual ‘open_system_micro’ lift is now concentrated on the narrowest current object available on the audited bridge:

$$|J_{\text{self-pair,core}}^{(R_*)}(11)|.$$

No interval theorem, no family redesign, and no exact-cross control sharper than a small positive refinement is still needed to explain the exported bottleneck.

Proposition O52 (At $D = 11$, the same-center pair core is already carried by the canonical compact anchor, and that anchor is itself carried by a same-center static anchor up to a local mirror loss). Keep the same singleton knot $D = 11$, the same radius $R_* = 1.35$, and the same canonical compact anchor

$$B_{\text{anchor}}(11) := B_{\text{box}}(0.70, 1.21; 11) \subset B_{\text{core}}(R_*; 11).$$

Split the same-center pair core as

$$J_{\text{self-pair,core}}^{(R_*)}(11) = J_{\text{self,anchor}}(11) + J_{\text{self,shell}}(11),$$

where

$$J_{\text{self,anchor}}(11) := \int_{B_{\text{anchor}}(11)} \mathcal{C}(\rho, z; 11) (\delta V_+(\rho, z; 11) + \delta V_-(\rho, z; 11)) d\rho dz$$

and

$$J_{\text{self,shell}}(11) := J_{\text{self-pair,core}}^{(R_*)}(11) - J_{\text{self,anchor}}(11).$$

Inside the same anchor, split further into its same-center static and local mirror pieces:

$$J_{\text{self,anchor}}(11) = J_{\text{diag,anchor}}(11) + J_{\text{mir,anchor}}(11),$$

with

$$J_{\text{diag,anchor}}(11) := \int_{B_+(11) \cap B_{\text{anchor}}(11)} \mathcal{C}(\rho, z; 11) \delta V_+(\rho, z; 11) d\rho dz + \int_{B_-(11) \cap B_{\text{anchor}}(11)} \mathcal{C}(\rho, z; 11) \delta V_-(\rho, z; 11) d\rho dz$$

and

$$J_{\text{mir,anchor}}(11) := \int_{B_+(11) \cap B_{\text{anchor}}(11)} \mathcal{C}(\rho, z; 11) \delta V_-(\rho, z; 11) d\rho dz + \int_{B_-(11) \cap B_{\text{anchor}}(11)} \mathcal{C}(\rho, z; 11) \delta V_+(\rho, z; 11) d\rho dz$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_static_anchor_source.py`, which composes the already-exported singleton-parent and compact-anchor audits, gives

$$|J_{\text{self-pair,core}}^{(R_*)}(11)| = 1.33812 \times 10^{-4},$$

$$|J_{\text{self,anchor}}(11)| = 1.32321 \times 10^{-4}, \quad |J_{\text{self,shell}}(11)| = 1.49126 \times 10^{-6},$$

with

$$\text{sgn } J_{\text{self,anchor}}(11) = \text{sgn } J_{\text{self-pair,core}}^{(R_*)}(11), \quad \text{sgn } J_{\text{self,shell}}(11) = \text{sgn } J_{\text{self-pair,core}}^{(R_*)}(11).$$

Hence

$$|J_{\text{self-pair,core}}^{(R_*)}(11)| \geq |J_{\text{self,anchor}}(11)|.$$

Moreover the same audit gives

$$|J_{\text{diag,anchor}}(11)| = 1.63757 \times 10^{-4}, \quad |J_{\text{mir,anchor}}(11)| = 3.14360 \times 10^{-5},$$

with

$$\text{sgn } J_{\text{diag,anchor}}(11) = \text{sgn } J_{\text{self,anchor}}(11), \quad \text{sgn } J_{\text{mir,anchor}}(11) = -\text{sgn } J_{\text{self,anchor}}(11).$$

Define the local mirror ratio

$$\mu_{11}^{(\text{anc})} := \frac{|J_{\text{mir,anchor}}(11)|}{|J_{\text{diag,anchor}}(11)|} = 0.191967 \dots < 1.$$

Then

$$|J_{\text{self,anchor}}(11)| \geq (1 - \mu_{11}^{(\text{anc})}) |J_{\text{diag,anchor}}(11)| = 1.32321 \times 10^{-4},$$

with machine-level reconstruction residual

$$||J_{\text{self,anchor}}(11)| - (1 - \mu_{11}^{(\text{anc})}) |J_{\text{diag,anchor}}(11)|| \leq 4.27 \times 10^{-17}.$$

Therefore

$$|J_{\text{self-pair,core}}^{(R_*)}(11)| \geq |J_{\text{self,anchor}}(11)| \geq (1 - \mu_{11}^{(\text{anc})}) |J_{\text{diag,anchor}}(11)|.$$

In particular the global and local singleton-core targets from Proposition O50 are already implied by the weaker same-center static-anchor inequalities

$$|J_{\text{diag,anchor}}(11)| \geq 1.49217 \times 10^{-4}, \quad |J_{\text{diag,anchor}}(11)| \geq 1.44663 \times 10^{-4},$$

and the current audited diagonal-anchor value exceeds them by the factors

$$1.09744 \quad \text{and} \quad 1.13199.$$

Corollary O51 (The remaining parent-side lift can be read as a same-center static-anchor lower bound at the single knot $D = 11$). Proposition O52 narrows the residual ‘open_system_micro’ theorem target once more, without changing the family, the knot, or the radius:

1. the same-center pair core on the singleton bottleneck knot is already carried by the canonical compact anchor $B_{\text{anchor}}(11)$, with only a tiny same-sign shell completion;
2. the compact anchor is already carried by the same-center static anchor $J_{\text{diag,anchor}}(11)$, with only a local mirror loss ratio $\mu_{11}^{(\text{anc})} \approx 0.192$;
3. therefore the remaining parent-side analytic task is no longer to control the full exact core $J_{\text{core}}^{(R_*)}(11)$, nor even the full same-center pair core $J_{\text{self-pair,core}}^{(R_*)}(11)$, but simply to prove a lower bound on the single same-center static quantity

$$|J_{\text{diag,anchor}}(11)|$$

strong enough to reach either 1.49217×10^{-4} (global) or 1.44663×10^{-4} (local).

Thus the audited bridge now isolates the bottleneck in its narrowest current parent-side form: a one-knot, same-family, same-center static-anchor lower bound.

Proposition O53 (At the singleton bottleneck knot, the static-anchor lift is equivalent to a local positive-moment lower bound or a local cancellation-ratio upper bound). Keep the same singleton knot $D = 11$ and the same same-center static anchor $J_{\text{diag,anchor}}(11)$. By the oriented-moment split from Proposition O24,

$$|J_{\text{diag,anchor}}(11)| = P_{\text{diag}}(11) - N_{\text{diag}}(11) = (1 - \nu_{11}^{(\pm)}) P_{\text{diag}}(11),$$

where

$$\nu_{11}^{(\pm)} := \frac{N_{\text{diag}}(11)}{P_{\text{diag}}(11)}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_diag_moment_source.py` reads the already-exported moment split at this singleton knot and gives

$$P_{\text{diag}}(11) = 1.87136 \times 10^{-3}, \quad N_{\text{diag}}(11) = 1.70580 \times 10^{-3},$$

$$\nu_{11}^{(\pm)} = 0.911531876\dots, \quad ||J_{\text{diag,anchor}}(11)| - (1 - \nu_{11}^{(\pm)})P_{\text{diag}}(11)| \leq 10^{-16}.$$

Therefore the global and local static-anchor targets from Proposition O52,

$$|J_{\text{diag,anchor}}(11)| \geq 1.49217 \times 10^{-4}, \quad |J_{\text{diag,anchor}}(11)| \geq 1.44663 \times 10^{-4},$$

are each equivalent to either of the following one-dimensional sufficient conditions:

$$P_{\text{diag}}(11) \geq \frac{1.49217 \times 10^{-4}}{1 - \nu_{11}^{(\pm)}} = 1.68667 \times 10^{-3},$$

$$P_{\text{diag}}(11) \geq \frac{1.44663 \times 10^{-4}}{1 - \nu_{11}^{(\pm)}} = 1.63520 \times 10^{-3},$$

or, holding the positive moment fixed,

$$\nu_{11}^{(\pm)} \leq 1 - \frac{1.49217 \times 10^{-4}}{P_{\text{diag}}(11)} = 0.920262868\dots,$$

$$\nu_{11}^{(\pm)} \leq 1 - \frac{1.44663 \times 10^{-4}}{P_{\text{diag}}(11)} = 0.922696321\dots$$

Numerically, the audited singleton positive moment already exceeds the two required positive-moment floors by the factors

$$1.10950 \quad \text{and} \quad 1.14442,$$

while the audited local cancellation ratio lies below the two admissible upper bounds by the absolute slacks

$$8.73 \times 10^{-3} \quad \text{and} \quad 1.12 \times 10^{-2}.$$

Corollary O52 (The most robust remaining parent-side route is a lower bound for the local positive moment $P_{\text{diag}}(11)$). Proposition O53 turns the last singleton static-anchor lift into two equivalent parent-side statements:

1. a local positive-moment lower bound for $P_{\text{diag}}(11)$;
2. a local cancellation-ratio upper bound for $\nu_{11}^{(\pm)}$.

On the audited bridge both are satisfied, but they are not equally robust: the positive-moment route keeps a multiplicative margin of about 10.95% (global) or 14.44% (local), whereas the cancellation-ratio route keeps only an absolute slack of about 0.0087 or 0.0112. Therefore the

cleanest next parent-side analytic target is no longer the full static anchor itself, but the local positive carrier moment

$$P_{\text{diag}}(11),$$

with $\nu_{11}^{(\pm)}$ retained only as a secondary tightening line if a sharper theorem is later wanted.

Proposition O54 (At $D = 11$, the local positive moment of the same-center static anchor is already controlled by a positive mass carrier minus a small derivative leakage). Keep the same knot $D = 11$ and the same canonical anchor $B_{\text{anchor}}(11)$. Write the oriented same-center static density as

$$q_{\text{diag}} = q_{\text{mass}} + q_{\text{der}},$$

where q_{mass} and q_{der} are the one-center mass and derivative anchor densities from Corollary O24, specialized to the singleton knot. Define

$$P_{\text{mass}}(11) := \int_{B_{\text{anchor}}(11)} [q_{\text{mass}}]_+ d\rho dz, \quad N_{\text{der}}(11) := \int_{B_{\text{anchor}}(11)} [q_{\text{der}}]_- d\rho dz.$$

Since for every real pair (a, b) one has

$$[a + b]_+ \geq [a]_+ - [b]_-,$$

it follows pointwise on the anchor that

$$[q_{\text{diag}}]_+ \geq [q_{\text{mass}}]_+ - [q_{\text{der}}]_-.$$

Integrating over $B_{\text{anchor}}(11)$ gives the exact lower bound

$$P_{\text{diag}}(11) \geq P_{\text{mass}}(11) - N_{\text{der}}(11).$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_diag_mass_bridge_source.py` evaluates this inequality on the exported singleton bridge and gives

$$P_{\text{mass}}(11) = 1.88085 \times 10^{-3}, \quad N_{\text{der}}(11) = 7.58159 \times 10^{-5},$$

so that

$$P_{\text{diag}}(11) \geq P_{\text{mass}}(11) - N_{\text{der}}(11) = 1.80504 \times 10^{-3}.$$

The pointwise lower-bound residual closes numerically:

$$\max_{B_{\text{anchor}}(11)} \left([q_{\text{mass}}]_+ - [q_{\text{der}}]_- - [q_{\text{diag}}]_+ \right) \leq 0$$

to machine precision in the audit export. In particular the induced lower bound already exceeds the global and local singleton positive-moment targets from Proposition O53,

$$1.68667 \times 10^{-3}, \quad 1.63520 \times 10^{-3},$$

by the factors

$$1.07018 \quad \text{and} \quad 1.10386.$$

Equivalently, if one holds $N_{\text{der}}(11)$ fixed, it is enough to prove the weaker positive-mass inequalities

$$P_{\text{mass}}(11) \geq 1.76249 \times 10^{-3}, \quad P_{\text{mass}}(11) \geq 1.71102 \times 10^{-3},$$

for the global and local routes. And if one instead holds $P_{\text{mass}}(11)$ fixed, it is enough to prove the derivative-leakage ceilings

$$N_{\text{der}}(11) \leq 1.94179 \times 10^{-4}, \quad N_{\text{der}}(11) \leq 2.45654 \times 10^{-4}.$$

The audited derivative leakage occupies only

$$\frac{N_{\text{der}}(11)}{P_{\text{mass}}(11)} = 4.03\%$$

of the positive mass carrier.

Corollary O53 (The cleanest next parent-side route is a lower bound for the local positive mass carrier $P_{\text{mass}}(11)$). Proposition O54 turns the singleton positive-moment target from Corollary O52 into two equivalent mass/derivative statements:

1. prove a lower bound for the positive one-center mass carrier $P_{\text{mass}}(11)$;
2. prove an upper bound for the adverse derivative leakage $N_{\text{der}}(11)$.

Both routes are valid on the audited bridge, but the mass route is structurally cleaner and numerically stronger: the required mass floors,

$$1.76249 \times 10^{-3} \quad \text{and} \quad 1.71102 \times 10^{-3},$$

sit below the audited value 1.88085×10^{-3} by the factors

$$1.06716 \quad \text{and} \quad 1.09927,$$

while the derivative route only asks that the already small leakage $N_{\text{der}}(11)$ remain below moderate ceilings. Therefore the sharpest current parent-side target is no longer the full static

anchor $J_{\text{diag,anchor}}(11)$, nor even the full positive moment $P_{\text{diag}}(11)$, but the local positive one-center mass carrier

$$P_{\text{mass}}(11),$$

with $N_{\text{der}}(11)$ retained only as a secondary local correction.

Proposition O55 (At $D = 11$, the local positive mass carrier factorizes exactly into a strictly positive one-center mass kernel and a positive oriented carrier moment).

Keep the same knot $D = 11$ and the same canonical anchor $B_{\text{anchor}}(11)$. Let

$$\sigma_{11} := \text{sgn}\left(\int_{B_{\text{anchor}}(11)} q_{\text{diag}}\right),$$

and define the oriented anchor carrier

$$\mathcal{C}_{11}^{(\text{or})} := \sigma_{11} \mathcal{C}.$$

Write the one-center mass kernel on the singleton anchor as

$$K_{\text{mass},11}^{(\text{anc})} := (\delta m_+) \mathbf{1}_{B_+(11)} + (\delta m_-) \mathbf{1}_{B_-(11)},$$

where $\delta m_{\pm}$ are the centered mass deviations from Proposition O54. Then on $B_{\text{anchor}}(11)$ one has

$$q_{\text{mass}} = \mathcal{C}_{11}^{(\text{or})} K_{\text{mass},11}^{(\text{anc})}.$$

The singleton audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_ext_schur_singleton_pmass_kernel_source.py` shows that

$$K_{\text{mass},11}^{(\text{anc})} > 0 \quad \text{everywhere on } B_{\text{anchor}}(11),$$

so the positive part factorizes pointwise:

$$[q_{\text{mass}}]_+ = K_{\text{mass},11}^{(\text{anc})} [\mathcal{C}_{11}^{(\text{or})}]_+.$$

Define the positive oriented carrier moment

$$W_{11}^{(+)} := \int_{B_{\text{anchor}}(11)} [\mathcal{C}_{11}^{(\text{or})}]_+ d\rho dz$$

and the weighted mass-kernel average

$$\bar{K}_{\text{mass},11}^{(+)} := \frac{P_{\text{mass}}(11)}{W_{11}^{(+)}}.$$

Then one has the exact factorization

$$P_{\text{mass}}(11) = \bar{K}_{\text{mass},11}^{(+)} W_{11}^{(+)}.$$

The audit gives

$$W_{11}^{(+)} = 2.02920 \times 10^{-2}, \quad \bar{K}_{\text{mass},11}^{(+)} = 9.26892 \times 10^{-2},$$

with factorization residual zero to machine precision. The same export also yields the pointwise kernel range

$$4.36792 \times 10^{-2} \leq K_{\text{mass},11}^{(\text{anc})} \leq 8.97295 \times 10^{-1},$$

and therefore the crude product bounds

$$K_{\min} W_{11}^{(+)} \leq P_{\text{mass}}(11) \leq K_{\max} W_{11}^{(+)}.$$

At the audited singleton knot the lower product is

$$K_{\min} W_{11}^{(+)} = 8.86341 \times 10^{-4},$$

which reaches only

$$50.29\% \quad \text{and} \quad 51.80\%$$

of the global and local mass targets from Proposition O54. By contrast the exact weighted route gives the equivalent carrier thresholds

$$W_{11}^{(+)} \geq 1.90151 \times 10^{-2}, \quad W_{11}^{(+)} \geq 1.84597 \times 10^{-2},$$

or, equivalently, the weighted-kernel thresholds

$$\bar{K}_{\text{mass},11}^{(+)} \geq 8.68562 \times 10^{-2}, \quad \bar{K}_{\text{mass},11}^{(+)} \geq 8.43195 \times 10^{-2},$$

for the global and local routes. The audited singleton values clear these by the factors

$$1.06716 \quad \text{and} \quad 1.09926.$$

The plus and minus lobes are also nearly symmetric on this factorized route:

$$P_{\text{mass},+}(11) = 9.47708 \times 10^{-4}, \quad P_{\text{mass},-}(11) = 9.33145 \times 10^{-4},$$

$$W_{11,+}^{(+)} = 1.01723 \times 10^{-2}, \quad W_{11,-}^{(+)} = 1.01197 \times 10^{-2},$$

and the corresponding weighted kernel averages differ by only about 1.03%.

Corollary O54 (The one-center mass profile is no longer the bottleneck; the remaining singleton difficulty lies in the positive carrier geometry, not in a bare kernel floor). Proposition O55 shows that the singleton positive-mass target from Corollary O53 admits two exact routes:

1. prove a lower bound for the positive oriented carrier moment $W_{11}^{(+)}$;
2. prove a lower bound for the weighted kernel average $\bar{K}_{\text{mass},11}^{(+)}$.

Both routes are equivalent on the audited bridge. But the audit also shows that a *bare minimum-kernel floor* is too crude: replacing $\bar{K}_{\text{mass},11}^{(+)}$ by $K_{\min}$ loses roughly half the required mass target. Therefore the cleanest current parent-side reading is the following. The local one-center mass profile has already been stabilized: it is explicit, strictly positive on the canonical singleton anchor, and almost lobe-symmetric. The remaining theorem-level difficulty is no longer the full positive mass carrier $P_{\text{mass}}(11)$ itself, but the positive oriented carrier geometry that weights this mass kernel. Equivalently, if one prefers to keep the mass profile on the theorem side, then one must prove a *nontrivial weighted* lower bound for $\bar{K}_{\text{mass},11}^{(+)}$, not merely a pointwise minimum-floor estimate.

Proposition O56 (At $D = 11$, the positive oriented carrier route reduces exactly to a single-lobe lower bound, and the total carrier is recovered by reflection symmetry).

Keep the same knot $D = 11$ and the same canonical anchor

$$B_{\text{anchor}}(11) = B_+(11) \cup B_-(11).$$

Define the single-lobe positive oriented carrier moments

$$W_{11,+}^{(+)} := \int_{B_+(11)} [\mathcal{C}_{11}^{(\text{or})}]_+ d\rho dz, \quad W_{11,-}^{(+)} := \int_{B_-(11)} [\mathcal{C}_{11}^{(\text{or})}]_+ d\rho dz.$$

Let $\mathfrak{R}(\rho, z) := (\rho, -z)$. The two-center background at $D = 11$ is invariant under $\mathfrak{R}$, and the lowest parity pair may be chosen with definite reflection parity,

$$\psi_1 \circ \mathfrak{R} = \psi_1, \quad \psi_2 \circ \mathfrak{R} = -\psi_2.$$

Therefore

$$\psi_1^2 \circ \mathfrak{R} = \psi_1^2, \quad \psi_2^2 \circ \mathfrak{R} = \psi_2^2, \quad \mathcal{C} \circ \mathfrak{R} = \mathcal{C},$$

and since σ_{11} is a global sign,

$$\mathcal{C}_{11}^{(\text{or})} \circ \mathfrak{R} = \mathcal{C}_{11}^{(\text{or})}.$$

The same reflection exchanges the two lobes:

$$\mathfrak{R}(B_+(11)) = B_-(11).$$

Hence the positive oriented carrier is exactly lobe-symmetric,

$$W_{11,+}^{(+)} = W_{11,-}^{(+)}.$$

Consequently the total carrier from Proposition O55 satisfies

$$W_{11}^{(+)} = W_{11,+}^{(+)} + W_{11,-}^{(+)} = 2W_{11,+}^{(+)} = 2W_{11,-}^{(+)}.$$

The singleton audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_single_lobe_carrier_source.py` records

$$W_{11,+}^{(+)} = 1.01723 \times 10^{-2}, \quad W_{11,-}^{(+)} = 1.01197 \times 10^{-2}, \quad \frac{1}{2}W_{11}^{(+)} = 1.01460 \times 10^{-2},$$

with exact lobe-sum reconstruction residual

$$|W_{11,+}^{(+)} + W_{11,-}^{(+)} - W_{11}^{(+)}| = 3.47 \times 10^{-18}.$$

On the current grid the two audited lobe values differ from the exact half-total by only

$$2.63 \times 10^{-5},$$

so the audited bridge is already numerically consistent with the analytic reflection picture. Thus the singleton carrier theorem target no longer needs to be read as a total-anchor statement: it has reduced exactly to a one-lobe lower bound, with the second lobe recovered by symmetry rather than by a separate estimate.

Corollary O55 (To close the singleton positive-carrier route, it is enough to prove a one-lobe lower bound and then reflect it). By Proposition O56, the carrier thresholds from Proposition O55 halve exactly. The global and local singleton carrier targets become

$$W_{11,+}^{(+)} \geq 9.50753 \times 10^{-3}, \quad W_{11,-}^{(+)} \geq 9.22985 \times 10^{-3},$$

and the same thresholds hold for $W_{11,-}^{(+)}$. The audited singleton lobe values satisfy

$$\begin{aligned} \frac{W_{11,+}^{(+)}}{(W_{11}^{(+)})_{\text{req},g}/2} &= 1.06993, & \frac{W_{11,+}^{(+)}}{(W_{11}^{(+)})_{\text{req},l}/2} &= 1.10211, \\ \frac{W_{11,-}^{(+)}}{(W_{11}^{(+)})_{\text{req},g}/2} &= 1.06439, & \frac{W_{11,-}^{(+)}}{(W_{11}^{(+)})_{\text{req},l}/2} &= 1.09641. \end{aligned}$$

Therefore the sharpest current parent-side reading is no longer a theorem about the total carrier $W_{11}^{(+)}$, nor about a weighted kernel average, but simply this:

1. prove a lower bound for one lobe, e.g. $W_{11,+}^{(+)}$;
2. recover $W_{11}^{(+)}$ by the exact reflection symmetry from Proposition O56.

No further support refinement is needed, and no bare minimum-kernel estimate needs to be revisited. The remaining theorem-level difficulty has narrowed to a one-lobe positive oriented carrier lower bound on the canonical singleton anchor.

Proposition O57 (At $D = 11$, the one-lobe positive carrier admits an exact $++/-$ localized-overlap split, and the $++$ piece already carries the threshold). Keep the same knot $D = 11$, the same canonical anchor, and the plus lobe $B_+(11)$. Introduce the localized states

$$\phi_R := \frac{\psi_1 + \psi_2}{\sqrt{2}}, \quad \phi_L := \frac{\psi_1 - \psi_2}{\sqrt{2}}.$$

Then

$$\psi_1^2 - \psi_2^2 = 2\phi_R\phi_L,$$

so on $B_+(11)$ the oriented carrier density from Proposition O55 may be written as

$$\mathcal{C}_{11}^{(\text{or})} = 4\pi\rho\phi_R\phi_L.$$

For any real numbers a, b one has the exact identity

$$[ab]_+ = [a]_+[b]_+ + [a]_-[b]_-.$$

Applying this pointwise to $a = \phi_R$, $b = \phi_L$ yields

$$[\mathcal{C}_{11}^{(\text{or})}]_+ = 4\pi\rho\left([\phi_R]_+[\phi_L]_+ + [\phi_R]_-[\phi_L]_-\right) \quad \text{on } B_+(11).$$

Define the one-lobe positive-positive overlap moment and the negative-negative same-sign completion by

$$W_{11,+}^{(++)} := 4\pi \int_{B_+(11)} \rho [\phi_R]_+ [\phi_L]_+ d\rho dz,$$

$$W_{11,+}^{(--)} := 4\pi \int_{B_+(11)} \rho [\phi_R]_- [\phi_L]_- d\rho dz.$$

Then the singleton one-lobe carrier splits exactly as

$$W_{11,+}^{(+)} = W_{11,+}^{(++)} + W_{11,+}^{(--)}.$$

In particular,

$$W_{11,+}^{(+)} \geq W_{11,+}^{(++)}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_single_lobe_overlap_source.py` gives

$$W_{11,+}^{(+)} = 1.01723 \times 10^{-2}, \quad W_{11,+}^{(++)} = 1.01659 \times 10^{-2}, \quad W_{11,+}^{(--)} = 6.43626 \times 10^{-6},$$

with exact split residual

$$2.03 \times 10^{-20}.$$

Thus the positive-positive overlap already carries

$$\frac{W_{11,+}^{(++)}}{W_{11,+}^{(+)}} = 0.999367,$$

while the negative-negative completion is only

$$\frac{W_{11,+}^{(--)}}{W_{11,+}^{(+)}} = 6.33 \times 10^{-4}.$$

The same export also shows that the same-center localized mode is almost one-sign on its own lobe:

$$\frac{\|\phi_{R,-}\|_{L^1(B_+)}}{\|\phi_{R,+}\|_{L^1(B_+)}} = 2.13 \times 10^{-4},$$

and only 1.23% of the plus-lobe grid points have $\phi_R < 0$. So the remaining one-lobe carrier theorem is no longer naturally read as a statement about a general positive part; it has narrowed to a positive-positive localized-overlap lower bound, with the negative-negative completion acting only as a tiny same-sign refinement.

Corollary O56 (The sharpest current singleton theorem target is the one-lobe positive-positive overlap $W_{11,+}^{(++)}$). By Corollary O55, it is enough to prove

$$W_{11,+}^{(+)} \geq 9.50753 \times 10^{-3} \quad \text{or} \quad W_{11,+}^{(+)} \geq 9.22985 \times 10^{-3}.$$

By Proposition O57, it is therefore enough to prove the *stronger* one-lobe lower bound

$$W_{11,+}^{(++)} \geq 9.50753 \times 10^{-3} \quad \text{or} \quad W_{11,+}^{(++)} \geq 9.22985 \times 10^{-3},$$

since then

$$W_{11,+}^{(+)} \geq W_{11,+}^{(++)}, \quad W_{11}^{(+)} = 2W_{11,+}^{(+)}.$$

Numerically the audited positive-positive overlap already clears these one-lobe thresholds by the factors

$$1.06925 \quad \text{and} \quad 1.10142,$$

even before the tiny same-sign $--$ completion is added back. Therefore the cleanest current parent-side endpoint is no longer the total carrier $W_{11}^{(+)}$, nor the full one-lobe positive part $W_{11,+}^{(+)}$, but the one-lobe positive-positive localized-overlap moment

$$W_{11,+}^{(++)}.$$

That is the narrowest theorem object currently visible on the audited bridge.

Proposition O58 (At $D = 11$, the one-lobe positive-positive overlap factors exactly through single-function positive lobe masses; the preferred route is the mirror-positive lobe mass.) Keep the same knot $D = 11$, the same plus lobe $B_+(11)$, and the same localized states (ϕ_R, ϕ_L) from Proposition O57. Define the one-lobe same-center and mirror positive masses by

$$S_{11,+}^{(R+)} := 4\pi \int_{B_+(11)} \rho[\phi_R]_+ d\rho dz, \quad S_{11,+}^{(L+)} := 4\pi \int_{B_+(11)} \rho[\phi_L]_+ d\rho dz.$$

Also define the corresponding weighted positive averages by

$$\bar{\Phi}_{L,11,+}^{(+|R+)} := \frac{W_{11,+}^{(++)}}{S_{11,+}^{(R+)}}, \quad \bar{\Phi}_{R,11,+}^{(+|L+)} := \frac{W_{11,+}^{(++)}}{S_{11,+}^{(L+)}}.$$

Then the one-lobe positive-positive overlap admits the exact factorizations

$$W_{11,+}^{(++)} = S_{11,+}^{(R+)} \bar{\Phi}_{L,11,+}^{(+|R+)} = S_{11,+}^{(L+)} \bar{\Phi}_{R,11,+}^{(+|L+)}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_single_lobe_pp_factor_source.py` gives

$$W_{11,+}^{(++)} = 1.01659 \times 10^{-2}, \quad S_{11,+}^{(R+)} = 6.58504 \times 10^{-1}, \quad \bar{\Phi}_{L,11,+}^{(+|R+)} = 1.54379 \times 10^{-2},$$

$$S_{11,+}^{(L+)} = 1.55385 \times 10^{-1}, \quad \bar{\Phi}_{R,11,+}^{(+|L+)} = 6.54241 \times 10^{-2},$$

with factorization residuals

$$0 \quad \text{and} \quad 1.73 \times 10^{-18}.$$

Moreover, on the current exact export one has

$$\phi_R > 0 \quad \text{on the whole support } \{\phi_L > 0\} \cap B_+(11),$$

with audited minimum

$$\min_{\{\phi_L > 0\} \cap B_+(11)} \phi_R = 7.79 \times 10^{-3}.$$

So after Proposition O57, the remaining one-lobe theorem target is no longer naturally read as a generic two-function overlap. It is equivalently a statement about one single-function positive mass multiplied by a weighted positive average, and the cleaner of the two routes is the mirror-positive one, because its multiplier is the same-center localized mode on its own lobe.

Corollary O57 (The sharpest current one-function theorem target is the mirror-positive lobe mass $S_{11,+}^{(L+)}$). By Corollary O56, it is enough to prove

$$W_{11,+}^{(++)} \geq 9.50753 \times 10^{-3} \quad \text{or} \quad W_{11,+}^{(++)} \geq 9.22985 \times 10^{-3}.$$

By Proposition O58, this is equivalent to either of the two exact single-function routes:

$$S_{11,+}^{(L+)} \geq 1.45322 \times 10^{-1} \quad \text{or} \quad S_{11,+}^{(L+)} \geq 1.41077 \times 10^{-1},$$

provided the same-center weighted average $\bar{\Phi}_{R,11,+}^{(+|L+)}$ is kept at its audited value, or

$$S_{11,+}^{(R+)} \geq 6.15857 \times 10^{-1} \quad \text{or} \quad S_{11,+}^{(R+)} \geq 5.97871 \times 10^{-1},$$

provided the mirror-weighted average $\bar{\Phi}_{L,11,+}^{(+|R+)}$ is kept at its audited value. Numerically both routes clear the global and local one-lobe thresholds by the same factors

$$1.06925 \quad \text{and} \quad 1.10142,$$

but only the first route leaves outside a same-center multiplier on the same lobe. Therefore the sharpest current single-function parent-side endpoint is no longer $W_{11,+}^{(++)}$ itself, but the one-lobe mirror-positive mass

$$S_{11,+}^{(L+)} := 4\pi \int_{B_+(11)} \rho [\phi_L]_+ d\rho dz.$$

That is the cleanest remaining theorem object on the audited singleton bridge.

Proposition O59 (At $D = 11$, the one-lobe mirror-positive mass is carried entirely by the lower half of the plus lobe, up to a tiny negative leakage). Keep the same knot $D = 11$, the same plus lobe $B_+(11)$, and define the lower and upper halves

$$B_{+,\downarrow}(11) := B_+(11) \cap \{z < D/2\}, \quad B_{+,\uparrow}(11) := B_+(11) \cap \{z \geq D/2\}.$$

For the mirror-localized state $\phi_L = (\psi_1 - \psi_2)/\sqrt{2}$, define

$$S_{11,+}^{(L+)} := 4\pi \int_{B_+(11)} \rho [\phi_L]_+ d\rho dz,$$

$$I_{11,+}^{(L,\downarrow)} := 4\pi \int_{B_{+,\downarrow}(11)} \rho \phi_L d\rho dz, \quad N_{11,+}^{(L,\downarrow)} := 4\pi \int_{B_{+,\downarrow}(11)} \rho [\phi_L]_- d\rho dz.$$

Then

$$S_{11,+}^{(L+)} = 4\pi \int_{B_{+,\downarrow}(11)} \rho [\phi_L]_+ d\rho dz = I_{11,+}^{(L,\downarrow)} + N_{11,+}^{(L,\downarrow)}.$$

Indeed, the audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_mirror_lobe_half_source.py` shows that

$$\phi_L \leq 0 \quad \text{everywhere on } B_{+,\uparrow}(11),$$

so the upper half contributes no positive mass at all. On the lower half one has the exact pointwise identity

$$[x]_+ = x + [x]_-,$$

applied to $x = \phi_L$. Therefore

$$S_{11,+}^{(L+)} \geq I_{11,+}^{(L,\downarrow)}.$$

The same export gives

$$S_{11,+}^{(L+)} = 1.55385 \times 10^{-1}, \quad I_{11,+}^{(L,\downarrow)} = 1.55208 \times 10^{-1}, \quad N_{11,+}^{(L,\downarrow)} = 1.76989 \times 10^{-4},$$

with exact bridge residual

$$0.$$

So the lower-half negative leakage obeys

$$\frac{N_{11,+}^{(L,\downarrow)}}{S_{11,+}^{(L+)}} = 1.139 \times 10^{-3}, \quad \frac{N_{11,+}^{(L,\downarrow)}}{I_{11,+}^{(L,\downarrow)}} = 1.140 \times 10^{-3}.$$

Thus the one-lobe mirror-positive mass is no longer naturally a theorem about a positive part. On the audited singleton bridge it has reduced to a lower-half signed mass, with only a permille-level negative leakage left outside.

Corollary O58 (The sharpest current parent-side endpoint is the lower-half signed mirror mass $I_{11,+}^{(L,\downarrow)}$). By Corollary O57 it is enough to prove

$$S_{11,+}^{(L+)} \geq 1.45322 \times 10^{-1} \quad \text{or} \quad S_{11,+}^{(L+)} \geq 1.41077 \times 10^{-1}.$$

By Proposition O59 it is therefore enough to prove the stronger signed lower-half bounds

$$I_{11,+}^{(L,\downarrow)} \geq 1.45322 \times 10^{-1} \quad \text{or} \quad I_{11,+}^{(L,\downarrow)} \geq 1.41077 \times 10^{-1},$$

since then $S_{11,+}^{(L+)} \geq I_{11,+}^{(L,\downarrow)}$. Numerically the audited lower-half signed mass already clears these thresholds by the factors

$$1.06803 \quad \text{and} \quad 1.10016,$$

while the missing positive mass from the tiny lower-half negative leakage is only 1.14×10^{-3} of the target scale. Therefore the cleanest current parent-side theorem object is no longer the one-lobe positive mass $S_{11,+}^{(L+)}$, but the lower-half signed mirror mass

$$I_{11,+}^{(L,\downarrow)} := 4\pi \int_{B_{+,\downarrow}(11)} \rho \phi_L d\rho dz.$$

At this stage, the remaining theorem problem has become a one-function, one-half-lobe, signed mass lower bound.

Proposition O60 (At $D = 11$, the lower-half mirror mass admits a strict lower slab on which the bare axisymmetric Laplacian already sees a positive source.) Keep the same knot $D = 11$, the same plus-lobe anchor $B_+(11)$, and define the axisymmetric extraction operator

$$H_a := -\Delta_a + U_a, \quad \Delta_a := \partial_\rho^2 + \rho^{-1}\partial_\rho + \partial_z^2,$$

where U_a is the exact potential entering the localized 2D generalized eigenproblem. Write the first two extracted modes as

$$H_a \psi_1 = E_1 \psi_1, \quad H_a \psi_2 = E_2 \psi_2, \quad E_1 < E_2,$$

and keep the localized combinations

$$\phi_L := \frac{\psi_1 - \psi_2}{\sqrt{2}}, \quad \phi_R := \frac{\psi_1 + \psi_2}{\sqrt{2}}.$$

If

$$\bar{E} := \frac{E_1 + E_2}{2}, \quad \delta E := \frac{E_2 - E_1}{2},$$

then

$$(H_a - \bar{E})\phi_L = -\delta E \phi_R,$$

hence

$$-\Delta_a \phi_L = (\bar{E} - U_a)\phi_L - \delta E \phi_R.$$

Now define the strict lower slab

$$\Omega_{+,\sharp}(11) := B_+(11) \cap \{z \leq D/2 - 0.18\},$$

and its signed mirror mass

$$I_{11,+}^{(L,\sharp)} := 4\pi \int_{\Omega_{+,\sharp}(11)} \rho \phi_L d\rho dz.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_mirror_strict_slab_source.py` shows that this is the shallowest audited cut for which all of the following hold simultaneously:

$$\phi_L > 0 \quad \text{on } \overline{\Omega_{+,\sharp}(11)}, \quad \phi_R > 0 \quad \text{on } \overline{\Omega_{+,\sharp}(11)},$$

and

$$f_{\sharp} := (\bar{E} - U_a)\phi_L - \delta E \phi_R \geq 2.22998 \times 10^{-3} > 0 \quad \text{on } \Omega_{+,\sharp}(11).$$

So on this strict slab the mirror-localized state is already a positive supersolution of the pure Laplacian barrier equation

$$-\Delta_a \phi_L = f_{\sharp}, \quad f_{\sharp} \geq 0.$$

Numerically the same export gives

$$I_{11,+}^{(L,\sharp)} = 1.53031 \times 10^{-1},$$

while the omitted interface strip inside $B_{+,\downarrow}(11)$ contributes only

$$I_{11,+}^{(L,\downarrow)} - I_{11,+}^{(L,\sharp)} = 2.17724 \times 10^{-3},$$

which is 1.40% of the lower-half signed mass. A crude constant-source torsion comparison on $\Omega_{+,\sharp}(11)$,

$$-\Delta_a \tau_{\sharp} = 1, \quad \tau_{\sharp}|_{\partial\Omega_{+,\sharp}} = 0,$$

gives the positive lower bound

$$4\pi \int_{\Omega_{+,\sharp}(11)} \rho \left(\min_{\Omega_{+,\sharp}} f_{\sharp} \right) \tau_{\sharp} d\rho dz = 2.79922 \times 10^{-4},$$

with pointwise comparison residual

$$\min_{\Omega_{+,\sharp}(11)} (\phi_L - (\min f_{\sharp}) \tau_{\sharp}) = 1.11436 \times 10^{-2} > 0.$$

So the barrier structure is already present on the audited bridge, even though the most naive constant-source barrier remains too conservative to saturate the target by itself.

Corollary O59 (The correct PDE proof domain is now the strict lower slab $\Omega_{+,\#}(11)$, with $-\Delta_a$ as the barrier operator.) By Corollary O58 the current parent-side target is

$$I_{11,+}^{(L,\downarrow)} \geq 1.45322 \times 10^{-1} \quad \text{or} \quad I_{11,+}^{(L,\downarrow)} \geq 1.41077 \times 10^{-1}.$$

Proposition O60 shows that the dominant part of this lower-half mass already lies on the barrier-friendly strict slab $\Omega_{+,\#}(11)$, where the unshifted axisymmetric Laplacian satisfies

$$-\Delta_a \phi_L = f_{\#}, \quad f_{\#} \geq 0,$$

with strictly positive boundary trace and strictly positive interior source on the audited export. Moreover,

$$\frac{I_{11,+}^{(L,\#)}}{1.45322 \times 10^{-1}} = 1.05305, \quad \frac{I_{11,+}^{(L,\#)}}{1.41077 \times 10^{-1}} = 1.08473.$$

Therefore the remaining analytic task is no longer to discover a new support or a new carrier object. The right proof problem has become: obtain an a priori lower bound for the strict-slab signed mass $I_{11,+}^{(L,\#)}$ using $-\Delta_a$, together with only a small interface-strip control outside that slab. In particular, $-\Delta_a$ should now be regarded as a proof tool for the existing mass route, not as a replacement theorem object.

Proposition O61 (On the strict slab, the mirror mass admits an exact Dirichlet split into a dominant boundary lift and a smaller positive source potential). Keep the same strict slab $\Omega_{+,\#}(11)$ from Proposition O60 and the same source

$$f_{\#} := (\bar{E} - U_a)\phi_L - \delta E \phi_R.$$

Define $u_{\#}$ as the zero-boundary source potential solving

$$-\Delta_a u_{\#} = f_{\#} \quad \text{on } \Omega_{+,\#}(11), \quad u_{\#}|_{\partial\Omega_{+,\#}(11)} = 0,$$

and define the boundary lift

$$h_{\#} := \phi_L - u_{\#}.$$

Then

$$-\Delta_a h_{\#} = 0 \quad \text{on } \Omega_{+,\#}(11), \quad h_{\#}|_{\partial\Omega_{+,\#}(11)} = \phi_L|_{\partial\Omega_{+,\#}(11)},$$

and the strict-slab signed mass decomposes exactly as

$$I_{11,+}^{(L,\#)} = I_{\partial,\#} + I_{\text{src},\#},$$

where

$$I_{\partial,\sharp} := 4\pi \int_{\Omega_{+,\sharp}(11)} \rho h_{\sharp} d\rho dz, \quad I_{\text{src},\sharp} := 4\pi \int_{\Omega_{+,\sharp}(11)} \rho u_{\sharp} d\rho dz.$$

On the audited bridge this decomposition is exported exactly by `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_dirichle t_split_source.py`, with reconstruction residual

$$1.39 \times 10^{-17}.$$

Moreover the same export shows

$$u_{\sharp} > 0, \quad h_{\sharp} > 0 \quad \text{everywhere on } \Omega_{+,\sharp}(11),$$

with

$$\min_{\Omega_{+,\sharp}(11)} u_{\sharp} = 8.79343 \times 10^{-5}, \quad \min_{\Omega_{+,\sharp}(11)} h_{\sharp} = 1.10620 \times 10^{-2}.$$

Numerically,

$$I_{11,+}^{(L,\sharp)} = 1.53031 \times 10^{-1}, \quad I_{\partial,\sharp} = 1.41805 \times 10^{-1}, \quad I_{\text{src},\sharp} = 1.12253 \times 10^{-2},$$

so

$$\frac{I_{\partial,\sharp}}{I_{11,+}^{(L,\sharp)}} = 0.926647, \quad \frac{I_{\text{src},\sharp}}{I_{11,+}^{(L,\sharp)}} = 0.0733531.$$

Thus the strict-slab mirror mass is carried primarily by the positive boundary lift, while the positive source potential supplies a smaller but still quantitatively relevant completion.

Corollary O60 (The remaining proof task is now chiefly a boundary-trace theorem on the strict slab, with the source potential supplying the final completion). By Corollary O59 the current global target is

$$I_{11,+}^{(L,\sharp)} \geq 1.45322 \times 10^{-1}.$$

By Proposition O61, it is enough to prove

$$I_{\partial,\sharp} + I_{\text{src},\sharp} \geq 1.45322 \times 10^{-1},$$

with $I_{\text{src},\sharp} \geq 0$ automatically available from the positive source equation on $\Omega_{+,\sharp}(11)$. On the audited bridge the boundary lift already reaches

$$\frac{I_{\partial,\sharp}}{1.45322 \times 10^{-1}} = 0.975804,$$

so it misses the global target by only

$$1.45322 \times 10^{-1} - I_{\partial, \#} = 3.51625 \times 10^{-3}.$$

The exact positive source potential contributes

$$I_{\text{src}, \#} = 1.12253 \times 10^{-2},$$

which exceeds that remaining global deficit by the factor

$$3.19240.$$

At the weaker local threshold

$$1.41077 \times 10^{-1},$$

the boundary lift alone already closes the target:

$$\frac{I_{\partial, \#}}{1.41077 \times 10^{-1}} = 1.00516.$$

By contrast, the most naive constant-source torsion barrier from Proposition O60 contributes only

$$2.79922 \times 10^{-4},$$

so even

$$I_{\partial, \#} + I_{\text{crude}, \#}$$

reaches only 97.77% of the global target. Therefore the proof-side picture is now sharp: the right PDE route is not a crude constant-source torsion estimate on the whole slab, but a boundary-trace theorem for the harmonic lift $h_{\#}$, supplemented by a genuinely nontrivial positive-source lower bound for $u_{\#}$. In that precise sense, $-\Delta_a$ has now become an effective proof tool rather than a new theorem object.

Proposition O62 (The strict-slab boundary lift splits exactly by boundary edge type, and the artificial top cap carries only a small remainder). Keep the same strict slab $\Omega_{+, \#}(11)$, the same discrete strict-slab operator $K_{\#}$, and the same boundary lift $h_{\#}$ from Proposition O61. On the audited bridge, define the exact boundary defect

$$b_{\#} := K_{\#}\phi_L - M_{\#}f_{\#} = K_{\#}h_{\#},$$

and split its support into the disjoint audited boundary pieces

$$\Gamma_{\text{lb}}, \Gamma_{\text{lt}}, \Gamma_{\text{rb}}, \Gamma_{\text{rt}}, \Gamma_{\text{axis}}, \Gamma_{\text{out}}, \Gamma_{\text{bot}}, \Gamma_{\text{top}},$$

corresponding respectively to the lower-left corner, upper-left corner, lower-right corner, upper-right corner, the axis edge, the outer radial wall, the bottom edge, and the artificial top-cap edge of $\Omega_{+,\#}(11)$. If $b_{\#}^{(\sigma)}$ denotes the restriction of $b_{\#}$ to Γ_{σ} , and $h_{\#}^{(\sigma)}$ is defined by

$$K_{\#}h_{\#}^{(\sigma)} = b_{\#}^{(\sigma)},$$

then linearity gives the exact decomposition

$$h_{\#} = \sum_{\sigma} h_{\#}^{(\sigma)}, \quad I_{\partial,\#} = \sum_{\sigma} I_{\partial,\#}^{(\sigma)},$$

where

$$I_{\partial,\#}^{(\sigma)} := 4\pi \int_{\Omega_{+,\#}(11)} \rho h_{\#}^{(\sigma)} d\rho dz.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_boundary_edge_source.py` verifies this edgewise split with reconstruction residual

$$3.12 \times 10^{-16}.$$

Numerically,

$$I_{\partial,\#}^{(\text{out})} = 8.20377 \times 10^{-2}, \quad I_{\partial,\#}^{(\text{bot})} = 4.78795 \times 10^{-2}, \quad I_{\partial,\#}^{(\text{axis})} = 6.48 \times 10^{-18},$$

while the corner lifts are

$$I_{\partial,\#}^{(\text{lb})} = 5.68899 \times 10^{-4}, \quad I_{\partial,\#}^{(\text{lt})} = 7.77352 \times 10^{-5}, \quad I_{\partial,\#}^{(\text{rb})} = 4.11332 \times 10^{-3}, \quad I_{\partial,\#}^{(\text{rt})} = 5.89740 \times 10^{-4},$$

and the top-edge lift is

$$I_{\partial,\#}^{(\text{top})} = 6.53837 \times 10^{-3}.$$

Define the physical-boundary and top-cap groups by

$$I_{\partial,\#}^{(\text{phys})} := I_{\partial,\#}^{(\text{axis})} + I_{\partial,\#}^{(\text{out})} + I_{\partial,\#}^{(\text{bot})} + I_{\partial,\#}^{(\text{lb})} + I_{\partial,\#}^{(\text{rb})},$$

$$I_{\partial,\#}^{(\text{cap})} := I_{\partial,\#}^{(\text{top})} + I_{\partial,\#}^{(\text{lt})} + I_{\partial,\#}^{(\text{rt})}.$$

Then

$$I_{\partial,\#}^{(\text{phys})} = 1.34599 \times 10^{-1}, \quad I_{\partial,\#}^{(\text{cap})} = 7.20585 \times 10^{-3},$$

so

$$\frac{I_{\partial,\#}^{(\text{phys})}}{I_{\partial,\#}} = 0.949185, \quad \frac{I_{\partial,\#}^{(\text{cap})}}{I_{\partial,\#}} = 0.0508151.$$

Thus the harmonic boundary lift is already carried almost entirely by the non-top physical edges of the strict slab, while the artificial top cap contributes only about five percent of the boundary mass.

Corollary O61 (The global strict-slab target already closes using only the physical boundary trace plus the positive source potential; the artificial top cap is not a bottleneck). Keep the same notation as in Proposition O62. Since

$$I_{11,+}^{(L,\sharp)} = I_{\partial,\sharp}^{(\text{phys})} + I_{\partial,\sharp}^{(\text{cap})} + I_{\text{src},\sharp},$$

it is enough for the global target to prove

$$I_{\partial,\sharp}^{(\text{phys})} + I_{\text{src},\sharp} \geq 1.45322 \times 10^{-1}.$$

On the audited bridge,

$$I_{\partial,\sharp}^{(\text{phys})} + I_{\text{src},\sharp} = 1.45825 \times 10^{-1},$$

so

$$\frac{I_{\partial,\sharp}^{(\text{phys})} + I_{\text{src},\sharp}}{1.45322 \times 10^{-1}} = 1.00346, \quad \frac{I_{\partial,\sharp}^{(\text{phys})} + I_{\text{src},\sharp}}{1.41077 \times 10^{-1}} = 1.03365.$$

Thus even after discarding the entire artificial top-cap lift $I_{\partial,\sharp}^{(\text{cap})}$, the remaining physical boundary trace together with the positive source potential already closes both the global and local strict-slab targets. So the current parent-side proof problem narrows once more: the remaining theorem target is no longer a trace theorem on the whole strict-slab boundary, but a trace theorem on the physical edges (outer wall, bottom edge, and the lower corners, with the axis contribution negligible), together with the already-positive source completion $I_{\text{src},\sharp}$.

Proposition O63 (The physical strict-slab boundary trace has an exact adjoint-Poisson representation, dominated by the outer wall and bottom edge). Keep the same strict-slab matrix $K_\sharp$, the same mass weights $\omega_\sharp$, and the same physical boundary defect

$$b_\sharp^{(\text{phys})} := b_\sharp^{(\text{axis})} + b_\sharp^{(\text{out})} + b_\sharp^{(\text{bot})} + b_\sharp^{(\text{lb})} + b_\sharp^{(\text{rb})}$$

from Proposition O62. Define the adjoint Poisson weight $p_\sharp$ by

$$K_\sharp^T p_\sharp = \omega_\sharp.$$

Then, by discrete Green duality,

$$I_{\partial,\sharp}^{(\text{phys})} = \omega_\sharp^T K_\sharp^{-1} b_\sharp^{(\text{phys})} = p_\sharp^T b_\sharp^{(\text{phys})}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_physical_trace_source.py` verifies the adjoint representation with residual

$$2.78 \times 10^{-17},$$

and also verifies agreement with the edgewise harmonic-lift grouping from Proposition O62 to

$$3.33 \times 10^{-16}.$$

Numerically,

$$I_{\partial,\sharp}^{(\text{phys})} = 1.34599 \times 10^{-1}, \quad \sum_{\Gamma_{\text{phys}}} b_{\sharp} = 6.22908 \times 10^2, \quad \bar{p}_{\text{phys}} := \frac{p_{\sharp}^T b_{\sharp}^{(\text{phys})}}{\sum_{\Gamma_{\text{phys}}} b_{\sharp}} = 2.16082 \times 10^{-4}.$$

The physical trace splits as

$$I_{\partial,\sharp}^{(\text{out})} = 8.20377 \times 10^{-2}, \quad I_{\partial,\sharp}^{(\text{bot})} = 4.78795 \times 10^{-2}, \quad I_{\partial,\sharp}^{(\text{corners,low})} = 4.68222 \times 10^{-3},$$

with the axis contribution 6.48×10^{-18} . Thus

$$\frac{I_{\partial,\sharp}^{(\text{out})}}{I_{\partial,\sharp}^{(\text{phys})}} = 0.609495, \quad \frac{I_{\partial,\sharp}^{(\text{bot})}}{I_{\partial,\sharp}^{(\text{phys})}} = 0.355719, \quad \frac{I_{\partial,\sharp}^{(\text{corners,low})}}{I_{\partial,\sharp}^{(\text{phys})}} = 0.0347863.$$

So the physical trace is not spread evenly over the boundary: it is carried almost entirely by the outer radial wall and the bottom edge, with the lower corners supplying the final few percent.

Corollary O62 (After the source potential is retained, the physical-boundary theorem target is a tight but closed trace lower bound). By Corollary O61, once the positive source potential is retained it is enough to prove

$$I_{\partial,\sharp}^{(\text{phys})} \geq 1.45322 \times 10^{-1} - I_{\text{src},\sharp}.$$

On the audited bridge,

$$I_{\text{src},\sharp} = 1.12253 \times 10^{-2},$$

so the required physical-boundary floor is

$$I_{\partial,\sharp}^{(\text{phys})} \geq 1.34096 \times 10^{-1}.$$

The adjoint-Poisson trace gives

$$I_{\partial,\sharp}^{(\text{phys})} = 1.34599 \times 10^{-1},$$

and therefore

$$\frac{I_{\partial,\sharp}^{(\text{phys})}}{1.34096 \times 10^{-1}} = 1.00375.$$

The dominant pair alone gives

$$I_{\partial,\sharp}^{(\text{out})} + I_{\partial,\sharp}^{(\text{bot})} = 1.29917 \times 10^{-1},$$

which is 96.8836% of the required physical trace. Adding the lower-right corner gives

$$I_{\partial,\sharp}^{(\text{out})} + I_{\partial,\sharp}^{(\text{bot})} + I_{\partial,\sharp}^{(\text{rb})} = 1.34031 \times 10^{-1},$$

which reaches 99.9510% of the required physical trace and misses by only

$$6.57 \times 10^{-5}.$$

The lower-left corner contribution

$$I_{\partial,\sharp}^{(\text{lb})} = 5.68899 \times 10^{-4}$$

then closes the remaining gap. Thus the next analytic theorem target is very narrow: prove a parent-side lower bound for the adjoint-Poisson physical trace on the outer wall, bottom edge, and lower corners of $\Omega_{+,\sharp}(11)$, while the positive source potential remains the already-available completion.

Proposition O64 (The physical-boundary trace closes as four adjoint-Poisson edge products). Keep the adjoint Poisson weight $p_{\sharp}$ from Proposition O63. For each physical edge Γ_{σ} , define

$$B_{\sigma} := \sum_{\Gamma_{\sigma}} b_{\sharp}, \quad \bar{p}_{\sigma} := \frac{p_{\sharp}^T b_{\sharp}^{(\sigma)}}{B_{\sigma}}, \quad I_{\sigma} := \bar{p}_{\sigma} B_{\sigma}.$$

Then

$$I_{\partial,\sharp}^{(\text{phys})} = I_{\text{out}} + I_{\text{bot}} + I_{\text{rb}} + I_{\text{lb}} + I_{\text{axis}}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_edge_product_source.py` verifies this product decomposition against the physical trace with zero residual. The non-negligible factors are

$$B_{\text{out}} = 2.83410 \times 10^2, \quad \bar{p}_{\text{out}} = 2.89467 \times 10^{-4}, \quad I_{\text{out}} = 8.20377 \times 10^{-2},$$

$$B_{\text{bot}} = 2.73854 \times 10^2, \quad \bar{p}_{\text{bot}} = 1.74836 \times 10^{-4}, \quad I_{\text{bot}} = 4.78795 \times 10^{-2},$$

$$B_{\text{rb}} = 6.32852 \times 10^1, \quad \bar{p}_{\text{rb}} = 6.49966 \times 10^{-5}, \quad I_{\text{rb}} = 4.11332 \times 10^{-3},$$

and

$$B_{\text{lb}} = 2.36003, \quad \bar{p}_{\text{lb}} = 2.41056 \times 10^{-4}, \quad I_{\text{lb}} = 5.68899 \times 10^{-4}.$$

The axis contribution is 6.48×10^{-18} . Hence

$$I_{\partial, \#}^{(\text{phys})} = 1.34599 \times 10^{-1}.$$

Equivalently, if

$$B_{\text{phys}} := \sum_{\Gamma_{\text{phys}}} b_{\#}, \quad \bar{p}_{\text{phys}} := \frac{I_{\partial, \#}^{(\text{phys})}}{B_{\text{phys}}},$$

then

$$B_{\text{phys}} = 6.22908 \times 10^2, \quad \bar{p}_{\text{phys}} = 2.16082 \times 10^{-4}, \quad I_{\partial, \#}^{(\text{phys})} = \bar{p}_{\text{phys}} B_{\text{phys}}.$$

Corollary O63 (The global closure now reduces to a tight edge-product inequality; the local closure is already carried by the outer wall and bottom edge). After the source potential is retained, the required physical trace is

$$T_{\text{phys}}^{(g)} := 1.45322 \times 10^{-1} - I_{\text{src}, \#} = 1.34096 \times 10^{-1}.$$

By Proposition O64,

$$\frac{I_{\partial, \#}^{(\text{phys})}}{T_{\text{phys}}^{(g)}} = 1.00375.$$

The outer wall and bottom edge alone give

$$I_{\text{out}} + I_{\text{bot}} = 1.29917 \times 10^{-1} = 0.968836 T_{\text{phys}}^{(g)}.$$

Adding the lower-right corner gives

$$I_{\text{out}} + I_{\text{bot}} + I_{\text{rb}} = 1.34031 \times 10^{-1} = 0.999510 T_{\text{phys}}^{(g)},$$

leaving only

$$T_{\text{phys}}^{(g)} - (I_{\text{out}} + I_{\text{bot}} + I_{\text{rb}}) = 6.57187 \times 10^{-5}.$$

The lower-left corner contributes

$$I_{\text{lb}} = 5.68899 \times 10^{-4},$$

which is 8.65657 times that final gap. Therefore the global target is closed by the four physical edge products

$$I_{\text{out}} + I_{\text{bot}} + I_{\text{rb}} + I_{\text{lb}} \geq T_{\text{phys}}^{(g)}.$$

At the local threshold the situation is even cleaner: $I_{\text{out}} + I_{\text{bot}}$ already reaches 1.00050 times the required physical trace after source. Thus the remaining parent-side analytic task is now a small finite set of edge-product bounds for $B_{\text{out}}, B_{\text{bot}}, B_{\text{rb}}, B_{\text{lb}}$ and their adjoint Poisson weights, rather than a theorem on the full physical boundary.

Proposition O65 (A rounded four-edge product certificate closes the physical trace).

Keep the notation of Proposition O64. Let B_{σ}^{-} and p_{σ}^{-} denote candidate parent-side lower bounds for the boundary defect mass and the adjoint-Poisson average on the four non-negligible physical edge pieces. The following rounded certificate is sufficient:

σ	B_{σ}^{-}	p_{σ}^{-}	$B_{\sigma}^{-} p_{\sigma}^{-}$
out	2.834×10^2	2.894×10^{-4}	8.20160×10^{-2}
bot	2.738×10^2	1.748×10^{-4}	4.78602×10^{-2}
rb	6.32×10^1	6.49×10^{-5}	4.10168×10^{-3}
lb	2.36	2.41×10^{-4}	5.68760×10^{-4} .

Indeed, $K_{\sharp}$ is the Dirichlet matrix for the positive axisymmetric operator $-\Delta_a$ on the strict slab, hence it is an M -matrix on the audited bridge and

$$K_{\sharp}^{-1} \geq 0, \quad K_{\sharp}^{-T} \geq 0, \quad p_{\sharp} = K_{\sharp}^{-T} \omega_{\sharp} \geq 0.$$

The boundary defects $b_{\sharp}^{(\sigma)}$ are nonnegative edgewise. Therefore $B_{\sigma} \geq 0$, $\bar{p}_{\sigma} \geq 0$, and the scalar inequalities

$$B_{\sigma} \geq B_{\sigma}^{-}, \quad \bar{p}_{\sigma} \geq p_{\sigma}^{-}, \quad \sigma \in \{\text{out}, \text{bot}, \text{rb}, \text{lb}\},$$

imply the rigorous product lower bound

$$I_{\partial, \sharp}^{(\text{phys})} \geq \sum_{\sigma} B_{\sigma}^{-} p_{\sigma}^{-} = 1.3454664 \times 10^{-1}.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_edge_certificate_source.py` verifies all eight scalar lower bounds with positive slack and produces the figure `/Users/boypatrick/codex/PSLT_quantam/output/chi_open_system/chi_open_system_exact_schur_singleton_strict_slab_edge_certificate_source_figure.png`. The certified lower edge sum exceeds the required global physical trace after source,

$$T_{\text{phys}}^{(g)} = 1.3409629777610438 \times 10^{-1},$$

by

$$1.3454664 \times 10^{-1} - T_{\text{phys}}^{(g)} = 4.50342 \times 10^{-4}, \quad \frac{1.3454664 \times 10^{-1}}{T_{\text{phys}}^{(g)}} = 1.00336.$$

The certificate is intentionally rounded downward: its total is only 5.28×10^{-5} below the exact audited edge product sum $1.3459947765481273 \times 10^{-1}$.

Corollary O64 (The remaining analytic lift is reduced to eight scalar lower bounds on the already-fixed physical edges). Combining Proposition O65 with the positive source potential from Proposition O61 gives

$$I_{\partial, \#}^{(\text{phys})} + I_{\text{src}, \#} \geq 1.3454664 \times 10^{-1} + 1.1225272657607257 \times 10^{-2} = 1.4577191265760725 \times 10^{-1}.$$

Thus

$$\frac{1.4577191265760725 \times 10^{-1}}{1.4532157043371163 \times 10^{-1}} = 1.00310,$$

so the rounded four-edge certificate still closes the global strict-slab target. The local threshold is already closed before either lower corner is used:

$$B_{\text{out}}^- p_{\text{out}}^- + B_{\text{bot}}^- p_{\text{bot}}^- = 1.2987620 \times 10^{-1} = 1.00019 T_{\text{phys}}^{(\ell)}.$$

For the global threshold, the first three certified products leave only

$$T_{\text{phys}}^{(g)} - (B_{\text{out}}^- p_{\text{out}}^- + B_{\text{bot}}^- p_{\text{bot}}^- + B_{\text{rb}}^- p_{\text{rb}}^-) = 1.18418 \times 10^{-4},$$

whereas the certified lower-left corner alone supplies

$$B_{\text{lb}}^- p_{\text{lb}}^- = 5.68760 \times 10^{-4},$$

which is 4.80300 times that remaining certified gap. Hence the parent-side theorem has now been reduced, without introducing any new support or carrier, to proving the four boundary-mass floors and four adjoint-Poisson average floors listed in Proposition O65.

Proposition O66 (The four-edge product certificate admits a fully local pointwise certificate). The scalar edge averages in Proposition O65 are not the only way to close the trace. A more local certificate is obtained by working directly pointwise on the same four physical edges. For each audited boundary point q on

$$\Gamma_{\text{out}} \cup \Gamma_{\text{bot}} \cup \Gamma_{\text{rb}} \cup \Gamma_{\text{lb}},$$

write

$$b_q := b_{\#}(q), \quad p_q := p_{\#}(q), \quad I_{\text{phys}}^{(4e)} := \sum_q b_q p_q.$$

By the M -matrix / maximum-principle positivity used in Proposition O65,

$$b_q \geq 0, \quad p_q \geq 0.$$

Therefore any pointwise lower floors $0 \leq b_q^- \leq b_q$ and $0 \leq p_q^- \leq p_q$ imply

$$I_{\text{phys}}^{(4e)} = \sum_q b_q p_q \geq \sum_q b_q^- p_q^-.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_edge_point_certificate_source.py` takes b_q^- and p_q^- to be the four-significant-digit downward roundings of b_q and p_q , respectively, and verifies all pointwise inequalities on the $33 + 10 + 1 + 1 = 45$ points of the fixed physical edges. The corresponding edge sums are

$$\sum_{\Gamma_{\text{out}}} b_q^- p_q^- = 8.20001038 \times 10^{-2},$$

$$\sum_{\Gamma_{\text{bot}}} b_q^- p_q^- = 4.78595834 \times 10^{-2},$$

$$\sum_{\Gamma_{\text{rb}}} b_q^- p_q^- = 4.1125672 \times 10^{-3}, \quad \sum_{\Gamma_{\text{lb}}} b_q^- p_q^- = 5.68760 \times 10^{-4}.$$

Thus the purely pointwise certified four-edge trace is

$$\sum_q b_q^- p_q^- = 1.345410144 \times 10^{-1}.$$

The exact four-edge product sum is

$$1.3459947765481273 \times 10^{-1},$$

so the loss caused by the pointwise downward rounding is only

$$5.84633 \times 10^{-5}.$$

Corollary O65 (The global trace closure no longer requires an edge-average theorem; pointwise edge barriers suffice). The global after-source physical trace target remains

$$T_{\text{phys}}^{(g)} = 1.3409629777610438 \times 10^{-1}.$$

By Proposition O66,

$$\sum_q b_q^- p_q^- - T_{\text{phys}}^{(g)} = 4.44717 \times 10^{-4}, \quad \frac{\sum_q b_q^- p_q^-}{T_{\text{phys}}^{(g)}} = 1.003316.$$

After adding the already-positive source potential, this gives

$$1.345410144 \times 10^{-1} + 1.1225272657607257 \times 10^{-2} = 1.4576628705760725 \times 10^{-1},$$

and hence

$$\frac{1.4576628705760725 \times 10^{-1}}{1.4532157043371163 \times 10^{-1}} = 1.003060.$$

The local threshold is still closed by the first two physical edges alone:

$$\sum_{\Gamma_{\text{out}} \cup \Gamma_{\text{bot}}} b_q^- p_q^- = 1.298596872 \times 10^{-1} = 1.000059 T_{\text{phys}}^{(\ell)}.$$

For the global threshold, the certified outer wall, bottom edge, and lower-right corner leave

$$T_{\text{phys}}^{(g)} - \sum_{\Gamma_{\text{out}} \cup \Gamma_{\text{bot}} \cup \Gamma_{\text{rb}}} b_q^- p_q^- = 1.24043 \times 10^{-4},$$

while the certified lower-left corner supplies

$$\sum_{\Gamma_{\text{lb}}} b_q^- p_q^- = 5.68760 \times 10^{-4},$$

namely 4.58517 times that remaining pointwise certified gap. Thus the remaining parent-side analytic lift can be stated in the sharpest local form: prove lower barriers for $b_{\sharp}(q)$ and $p_{\sharp}(q) = K_{\sharp}^{-T} \omega_{\sharp}(q)$ on the same 45 fixed physical-edge points, or replace those point floors by continuous edge barriers; no new support, carrier, or edge-average theorem is required.

Proposition O67 (The pointwise adjoint-Poisson floors are certified by a strict maximum-principle subsolution). Keep the strict-slab adjoint equation

$$K_{\sharp}^T p_{\sharp} = \omega_{\sharp}.$$

Let

$$\theta := 0.999999, \quad K_{\sharp}^T \underline{p}_{\sharp} = \theta \omega_{\sharp}.$$

Then

$$K_{\sharp}^T (p_{\sharp} - \underline{p}_{\sharp}) = (1 - \theta) \omega_{\sharp} \geq 0.$$

Since $K_{\sharp}$ is an M -matrix on the fixed strict slab, $K_{\sharp}^{-T} \geq 0$, and therefore

$$p_{\sharp} - \underline{p}_{\sharp} = K_{\sharp}^{-T} ((1 - \theta) \omega_{\sharp}) \geq 0.$$

Thus $\underline{p}_{\sharp}$ is a genuine adjoint-Poisson lower barrier for $p_{\sharp}$. The audit `/Users/boypatrick/code/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_`

`adjoint_barrier_certificate_source.py` verifies, on the same 45 physical-edge points as Proposition O66, that this strict lower barrier still dominates the four-significant-digit point floors p_q^- :

$$\underline{p}_\#(q) \geq p_q^-, \quad q \in \Gamma_{\text{out}} \cup \Gamma_{\text{bot}} \cup \Gamma_{\text{rb}} \cup \Gamma_{\text{lb}}.$$

The same audit checks the subsolution residual directly:

$$\min_q \{\omega_\# - K_\#^T \underline{p}_\#\}_q = 6.78584 \times 10^{-10}, \quad \max_q \{\omega_\# - K_\#^T \underline{p}_\#\}_q = 1.56074 \times 10^{-8},$$

with agreement against the expected residual $(1 - \theta)\omega_\#$ to

$$8.67 \times 10^{-19}.$$

Consequently the pointwise certificate from Proposition O66 is now backed by the maximum principle rather than only by reading off $p_\#$.

Corollary O66 (The O66 pointwise trace closure is stable under a strict adjoint-barrier proof of the Poisson weights). Use the same four-significant-digit b_q^- and p_q^- floors from Proposition O66. Proposition O67 supplies

$$p_\#(q) \geq \underline{p}_\#(q) \geq p_q^-,$$

while the boundary-defect floors give $b_\#(q) \geq b_q^-$. Therefore

$$I_{\partial, \#}^{(\text{phys})} \geq \sum_q b_q^- p_q^- = 1.345410144 \times 10^{-1}.$$

The intermediate barrier product with the same b_q^- is

$$\sum_q b_q^- \underline{p}_\#(q) = 1.3457028305021287 \times 10^{-1},$$

so the rounded point floors sit 2.92687×10^{-5} below the actual adjoint barrier. The certified trace ratio remains

$$\frac{\sum_q b_q^- p_q^-}{T_{\text{phys}}^{(g)}} = 1.003316,$$

and the source-completed global ratio remains

$$\frac{1.345410144 \times 10^{-1} + 1.1225272657607257 \times 10^{-2}}{1.4532157043371163 \times 10^{-1}} = 1.003060.$$

Also,

$$\sum_{\Gamma_{\text{out}} \cup \Gamma_{\text{bot}}} b_q^- p_q^- = 1.298596872 \times 10^{-1} = 1.000059 T_{\text{phys}}^{(\ell)},$$

so the local threshold is still carried by the first two physical edges alone. The only remaining parent-side burden on the Poisson side is now to replace the finite lower barrier $p_{\sharp}$ by a closed-form or monotone comparison barrier; the support, edge set, and target inequality are already fixed.

Proposition O68 (The pointwise boundary-defect floors are certified by the missing-neighbor flux formula). Keep the same strict slab $S = \Omega_{+, \sharp}(11)$, and write K_{full} for the full discrete matrix of the bare axisymmetric operator $-\Delta_a$ before restricting to S . Since the localized state satisfies the full-grid identity

$$K_{\text{full}}\phi_L = Mf_L, \quad f_L = (\bar{E} - U_a)\phi_L - \delta E \phi_R,$$

the restricted boundary defect

$$b_{\sharp} := K_{\sharp}\phi_{L,S} - M_S f_{L,S}$$

has the exact flux representation

$$b_{\sharp}(q) = - \sum_{r \notin S} K_{\text{full}}(q, r) \phi_L(r) = \sum_{r \in N(q) \setminus S} c_{qr} \phi_L(r), \quad c_{qr} := -K_{\text{full}}(q, r) \geq 0.$$

On the audited finite-volume grid,

$$c_{(i,j),(i+1,j)} = \frac{\rho_{i+1/2}}{\Delta \rho^2}, \quad c_{(i,j),(i-1,j)} = \frac{\rho_{i-1/2}}{\Delta \rho^2}, \quad c_{(i,j),(i,j \pm 1)} = \frac{\rho_i}{\Delta z^2}.$$

The axis has no adverse missing radial flux because $\rho_{-1/2} = 0$. Therefore any lower floors $\phi_L(r) \geq \phi_r^- \geq 0$ on the exterior missing neighbors imply

$$b_{\sharp}(q) \geq \sum_{r \in N(q) \setminus S} c_{qr} \phi_r^-.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_boundary_flux_certificate_source.py` uses five-significant-digit downward floors for the exterior ϕ_L values and verifies the flux certificate for the same 45 physical-edge points from Proposition O66. There are 46 missing-neighbor flux components: 33 on the outer wall, 10 on the bottom edge, two at the lower-right corner, and one at the lower-left corner. The exact flux reconstruction residual satisfies

$$\max_q \left| b_{\sharp}(q) - \sum_{r \in N(q) \setminus S} c_{qr} \phi_L(r) \right| = 1.31 \times 10^{-13}.$$

The certified flux floors sum to

$$\sum_q \sum_{r \in N(q) \setminus S} c_{qr} \phi_r^- = 6.229032 \times 10^2,$$

while the four-significant-digit point floors used in Proposition O66 sum to

$$\sum_q b_q^- = 6.22792 \times 10^2.$$

Thus the flux certificate exceeds the b_q^- floors by

$$1.11200 \times 10^{-1}, \quad \frac{6.229032 \times 10^2}{6.22792 \times 10^2} = 1.000179.$$

Corollary O67 (The O66 pointwise trace certificate is now supported on both sides: boundary flux for b_q and adjoint barrier for p_q). Combining Proposition O68 with Proposition O67 gives, on every one of the fixed physical-edge points,

$$b_{\#}(q) \geq b_q^-, \quad p_{\#}(q) \geq p_q^-.$$

Hence

$$I_{\partial, \#}^{(\text{phys})} \geq \sum_q b_q^- p_q^- = 1.345410144 \times 10^{-1}.$$

If one uses the flux-certified b -side floor directly together with the same p_q^- , the product is slightly stronger:

$$\sum_q \left(\sum_{r \in N(q) \setminus S} c_{qr} \phi_r^- \right) p_q^- = 1.34568861723 \times 10^{-1},$$

which is 1.003524 times the global after-source trace target. The rounded O66 product is the deliberately smaller certificate and still gives

$$\frac{1.345410144 \times 10^{-1}}{T_{\text{phys}}^{(g)}} = 1.003316, \quad \frac{1.345410144 \times 10^{-1} + 1.1225272657607257 \times 10^{-2}}{1.4532157043371163 \times 10^{-1}} = 1.003060.$$

The local threshold remains closed by the outer wall and bottom edge alone:

$$\sum_{\Gamma_{\text{out}} \cup \Gamma_{\text{bot}}} b_q^- p_q^- = 1.298596872 \times 10^{-1} = 1.000059 T_{\text{phys}}^{(\ell)}.$$

Thus the remaining parent-side work has narrowed again: replace the finite exterior ϕ_L floors and the finite adjoint subsolution $\underline{p}_{\#}$ by closed-form or monotone barriers. The trace identity, edge set, point set, and numerical closure are already fixed.

Proposition O69 (The exterior ϕ_L trace floors are certified by a parent-side trace subsolution). Let

$$E_{\partial} := \{r \notin S : r \in N(q) \text{ for one of the 45 physical edge points } q\}.$$

This is not a new support object; it is exactly the missing-neighbor trace already used in Proposition O68. On the audited grid it contains 46 points. Let K_E and M_E be the restrictions of the parent bare operator and mass matrix to this trace set, and write

$$g_E := K_E \phi_{L,E} - M_E f_{L,E}, \quad f_L = (\bar{E} - U_a) \phi_L - \delta E \phi_R.$$

The vector g_E is precisely the outward flux from the complement of the trace set E_∂ . The restricted matrix K_E is a strictly diagonally dominant Z -matrix, hence a nonsingular M -matrix. Therefore $K_E^{-1} \geq 0$. For

$$\theta_\phi := 0.9999995, \quad \underline{\phi}_E := \theta_\phi \phi_{L,E},$$

one has

$$K_E(\phi_{L,E} - \underline{\phi}_E) = (1 - \theta_\phi) K_E \phi_{L,E} = (1 - \theta_\phi)(M_E f_{L,E} + g_E).$$

Thus $f_{L,E} \geq 0$ and $g_E \geq 0$ imply, by the discrete maximum principle on the parent trace system, that

$$\phi_{L,E} \geq \underline{\phi}_E.$$

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_exterior_phi_barrier_source.py` verifies the hypotheses and compares $\underline{\phi}_E$ against four-significant-digit downward floors ϕ_r^{--} . Numerically,

$$\min_{E_\partial} f_L = 8.23807 \times 10^{-3}, \quad \min_{E_\partial} g_E = 4.71653, \quad \min_{\partial_1 E_\partial} \phi_L = 1.09922 \times 10^{-2}.$$

The M -matrix diagnostics are

$$\begin{aligned} \min_i (K_E)_{ii} &= 8.33333 \times 10^1, \\ \max_{i \neq j} (K_E)_{ij} &= -1.66667 \times 10^1, \\ \min_i \left((K_E)_{ii} - \sum_{j \neq i} |(K_E)_{ij}| \right) &= 6.66667 \times 10^1. \end{aligned}$$

Moreover,

$$\max_{r \in E_\partial} \frac{\phi_r^{--}}{\phi_L(r)} = 0.99999933598 < \theta_\phi, \quad \min_{r \in E_\partial} (\underline{\phi}_E(r) - \phi_r^{--}) = 1.01840 \times 10^{-8}.$$

The comparison residual is strictly positive,

$$\min_E (1 - \theta_\phi)(M_E f_{L,E} + g_E) = 2.36068 \times 10^{-6},$$

and agrees with the directly computed $K_E(\phi_{L,E} - \underline{\phi}_E)$ to

$$1.88 \times 10^{-14}.$$

Hence all exterior four-significant-digit ϕ_L trace floors used below are backed by a parent-side maximum-principle certificate.

Corollary O68 (The physical trace closes directly from the exterior trace barrier and the adjoint point floors). Use the same adjoint floors p_q^- from Proposition O66/O67, but instead of first certifying the old b_q^- floors pointwise, insert the four-significant-digit exterior trace floors directly into the missing-neighbor flux:

$$\widehat{b}_q := \sum_{r \in N(q) \setminus S} c_{qr} \phi_r^{--}.$$

By Proposition O68 and Proposition O69,

$$b_{\sharp}(q) \geq \widehat{b}_q, \quad p_{\sharp}(q) \geq p_q^-,$$

and therefore

$$I_{\partial, \sharp}^{(\text{phys})} \geq \sum_q \widehat{b}_q p_q^- = 1.3455686726 \times 10^{-1}.$$

This direct flux product is

$$\begin{aligned} \frac{1.3455686726 \times 10^{-1}}{T_{\text{phys}}^{(g)}} &= 1.003435, \\ \frac{1.3455686726 \times 10^{-1} + 1.1225272657607257 \times 10^{-2}}{1.4532157043371163 \times 10^{-1}} &= 1.003169. \end{aligned}$$

The local threshold is still closed by the outer wall and bottom edge alone:

$$\sum_{\Gamma_{\text{out}} \cup \Gamma_{\text{bot}}} \widehat{b}_q p_q^- = 1.298757133666667 \times 10^{-1} = 1.000182 T_{\text{phys}}^{(\ell)}.$$

This route is deliberately different from the old b_q^- -floor route: with only four-significant-digit exterior ϕ_L floors, 17 of the old pointwise b_q^- checks fail, and the right-bottom corner flux is 2.66667×10^{-3} below the old corner b -floor. Nevertheless the direct flux product has a larger final margin than the rounded O66 product. Thus the remaining parent-side burden is no longer a five-digit boundary-defect floor theorem; it is a four-digit exterior trace barrier, already certified here, plus the still-finite adjoint p_q^- comparison from Proposition O67.

Proposition O70 (The adjoint-Poisson side is certified by monotone finite Gauss–Seidel barriers). Let

$$A := K_{\sharp}^T = D - L - U, \quad L, U \geq 0,$$

be the standard lower/upper M -matrix splitting of the strict-slab adjoint operator. Starting from $p^{(0)} = 0$, define the monotone lower sequence

$$(D - L)p^{(n+1)} = \omega_{\sharp} + Up^{(n)}.$$

Since $D - L$ is triangular with positive diagonal and nonpositive off-diagonal entries, $(D - L)^{-1} \geq 0$. Hence $p^{(n+1)} \geq p^{(n)} \geq 0$. Moreover

$$\omega_{\sharp} - Ap^{(n+1)} = U(p^{(n+1)} - p^{(n)}) \geq 0.$$

The M -matrix property of A gives

$$p_{\sharp} - p^{(n+1)} = A^{-1}(\omega_{\sharp} - Ap^{(n+1)}) \geq 0.$$

Thus every finite iterate $p^{(n)}$ is a rigorous parent-side lower barrier for the exact adjoint Poisson weight $p_{\sharp}$; no exact solve for $p_{\sharp}$ is used in the proof.

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_adjoint_gs_barrier_source.py` uses the same 420-point strict slab and the same 45 physical edge points. The adjoint matrix diagnostics are

$$\min_i A_{ii} = 8.33333 \times 10^1, \quad \max_{i \neq j} A_{ij} = -1.66667 \times 10^1,$$

and the nonnegative coupling matrix $B = L + U = -A + \text{diag}(A)$ has off-diagonal entries in

$$[1.66667 \times 10^1, 7.66667 \times 10^2].$$

At the product-closing sweep $n = 548$, the minimum increment is

$$\min_i (p^{(548)} - p^{(547)})_i = 1.58824 \times 10^{-10},$$

and the residual identity

$$\omega_{\sharp} - Ap^{(548)} = U(p^{(548)} - p^{(547)})$$

agrees to 4.10×10^{-16} . At the optional point-floor sweep $n = 859$, the minimum increment is 2.37799×10^{-12} , and the same residual identity agrees to 4.80×10^{-16} .

Corollary O69 (The O69 exterior-flux trace closure now has a parent-side adjoint barrier, with an optional point-floor completion). Let

$$\widehat{b}_q = \sum_{r \in N(q) \setminus S} c_{qr} \phi_r^{--}$$

be the four-significant-digit exterior-flux floor from Corollary O68. Combining Proposition O69 with Proposition O70 gives

$$b_{\#}(q) \geq \widehat{b}_q, \quad p_{\#}(q) \geq p^{(548)}(q),$$

and hence

$$I_{\partial, \#}^{(\text{phys})} \geq \sum_q \widehat{b}_q p^{(548)}(q) = 1.345329256808323 \times 10^{-1}.$$

This gives

$$\frac{1.345329256808323 \times 10^{-1}}{T_{\text{phys}}^{(g)}} = 1.003256,$$

$$\frac{1.345329256808323 \times 10^{-1} + 1.1225272657607257 \times 10^{-2}}{1.4532157043371163 \times 10^{-1}} = 1.003005.$$

The local threshold is still closed by the outer wall and bottom edge alone, now with the fully monotone adjoint barrier:

$$\sum_{\Gamma_{\text{out}} \cup \Gamma_{\text{bot}}} \widehat{b}_q p^{(548)}(q) = 1.2985251408265905 \times 10^{-1} = 1.000003 T_{\text{phys}}^{(\ell)}.$$

If one wants compatibility with the older four-significant-digit p_q^- floors, the same monotone sequence supplies it at sweep $n = 859$:

$$\min_q \frac{p^{(859)}(q)}{p_q^-} = 1.000000033.$$

The corresponding direct flux product is

$$\sum_q \widehat{b}_q p^{(859)}(q) = 1.3458547480326316 \times 10^{-1}, \quad \frac{1.3458547480326316 \times 10^{-1}}{T_{\text{phys}}^{(g)}} = 1.003648.$$

Thus both sides of the physical trace certificate now come from parent-side maximum-principle constructions: the b -side from the exterior ϕ_L trace barrier of Proposition O69, and the p -side from the finite monotone adjoint barriers above. The remaining work is presentation-level compression or optional convergence-rate polishing, not a new support, carrier, or trace theorem.

Proposition O71 (Synthesis: the strict-slab physical trace certificate is parent-side closed). The preceding reductions can be read as one fixed certificate on the same strict slab

$$S = \Omega_{+, \#}(11).$$

No support, carrier, trace set, or split is introduced beyond the fixed objects already audited:

physical edge points	45,
missing-neighbor exterior trace points	46,
strict-slab adjoint points	420.

The proof chain is:

$$\begin{aligned}
b_{\sharp}(q) &= \sum_{r \in N(q) \setminus S} c_{qr} \phi_L(r) && \text{(missing-neighbor flux identity),} \\
&\geq \sum_{r \in N(q) \setminus S} c_{qr} \phi_r^{--} =: \widehat{b}_q && \text{(exterior } \phi_L \text{ trace barrier),} \\
p_{\sharp}(q) &\geq p^{(548)}(q) && \text{(finite monotone adjoint } M\text{-matrix barrier).}
\end{aligned}$$

Consequently

$$I_{\partial, \sharp}^{(\text{phys})} \geq \sum_q \widehat{b}_q p^{(548)}(q) = 1.345329256808323 \times 10^{-1}.$$

The synthesis audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_trace_synthesis_source.py` combines only previously exported certificates and records the closed pieces as follows:

boundary flux identity closed,
 exterior ϕ_L trace floors closed,
 adjoint GS product barrier closed,
 optional old p_q^- point floors closed at sweep 859.

The main 548-sweep certificate has

$$\frac{1.345329256808323 \times 10^{-1}}{T_{\text{phys}}^{(g)}} = 1.003256, \quad 1.345329256808323 \times 10^{-1} - T_{\text{phys}}^{(g)} = 4.36628 \times 10^{-4}.$$

The outer wall and bottom edge alone give

$$\sum_{\Gamma_{\text{out}} \cup \Gamma_{\text{bot}}} \widehat{b}_q p^{(548)}(q) = 1.2985251408265905 \times 10^{-1} = 1.000003 T_{\text{phys}}^{(\ell)}.$$

For comparison, the optional sweep 859 point-floor completion gives

$$\sum_q \widehat{b}_q p^{(859)}(q) = 1.3458547480326316 \times 10^{-1}, \quad \min_q \frac{p^{(859)}(q)}{p_q^-} = 1.000000033.$$

Corollary O70 (Roadmap closure for this strict-slab line). Adding the already-positive source potential

$$I_{\text{src}, \sharp} = 1.1225272657607257 \times 10^{-2}$$

to the main parent-side trace certificate gives

$$I_{\partial, \sharp}^{(\text{phys})} + I_{\text{src}, \sharp} \geq 1.4575819833843956 \times 10^{-1}.$$

Therefore

$$\frac{1.4575819833843956 \times 10^{-1}}{1.4532157043371163 \times 10^{-1}} = 1.003005.$$

Thus the strict-slab line is parent-side closed on the current audited bridge. The only remaining useful work is secondary tightening of the presentation and of the finite GS sweep-count explanation. The next proposition records this tightening without introducing any new support, object, carrier, or trace theorem.

Proposition O72 (Secondary tightening: the finite GS sweep counts have a componentwise rate certificate). Keep the same strict-slab adjoint splitting

$$A = K_{\#}^T = D - L - U, \quad L, U \geq 0,$$

and define

$$G := (D - L)^{-1}U \geq 0, \quad c := (D - L)^{-1}\omega_{\#} \geq 0.$$

Then the monotone GS barriers satisfy

$$p^{(n+1)} = c + Gp^{(n)}, \quad p^{(0)} = 0.$$

For the increments

$$d_n := p^{(n)} - p^{(n-1)} \quad (n \geq 1),$$

one has the exact recurrence

$$d_{n+1} = Gd_n.$$

Suppose that at a checkpoint n

$$0 \leq d_{n+1} \leq \alpha_n d_n, \quad \alpha_n < 1,$$

componentwise. Since $G \geq 0$, induction gives

$$0 \leq d_{n+k} \leq \alpha_n^k d_n, \quad k \geq 1.$$

Consequently the exact adjoint tail obeys the one-sided bound

$$0 \leq p_{\#} - p^{(n)} = \sum_{k \geq 1} d_{n+k} \leq \frac{\alpha_n}{1 - \alpha_n} d_n.$$

For the nonnegative trace functional

$$F(v) := \sum_q \widehat{b}_q v(q), \quad F_{\text{ob}}(v) := \sum_{q \in \Gamma_{\text{out}} \cup \Gamma_{\text{bot}}} \widehat{b}_q v(q),$$

this implies

$$\begin{aligned} 0 \leq F(p_{\#}) - F(p^{(n)}) &\leq \frac{\alpha_n}{1 - \alpha_n} F(d_n), \\ 0 \leq F_{\text{ob}}(p_{\#}) - F_{\text{ob}}(p^{(n)}) &\leq \frac{\alpha_n}{1 - \alpha_n} F_{\text{ob}}(d_n). \end{aligned}$$

This is an analytic convergence-rate certificate for the same finite GS sequence used in Proposition O70; the exact solve for $p_{\sharp}$ is not used in the lower-bound proof.

The audit `/Users/boypatrick/codex/PSLT_quantam/code/audit_open_system_exact_schur_singleton_strict_slab_gs_rate_source.py` verifies the componentwise subeigenvalue inequalities at the relevant checkpoints. The first sweep at which the global trace product closes is 384, while the first sweep at which the global and outer-bottom local inequalities close simultaneously is 548. The old four-significant-digit point floors are first dominated at sweep 859. At the main product-closing sweep,

$$\alpha_{548} = 0.9865813608269858, \quad \frac{\alpha_{548}}{1 - \alpha_{548}} = 73.52320515563653.$$

Moreover

$$F(d_{548}) = 7.256039748535137 \times 10^{-7},$$

so

$$F(p_{\sharp}) - F(p^{(548)}) \leq 5.334872990490022 \times 10^{-5}.$$

As an external numerical validation only, the exact adjoint solve gives

$$F(p_{\sharp}) - F(p^{(548)}) = 5.3347881092041805 \times 10^{-5},$$

inside the certified tail bound. At the point-floor sweep,

$$\alpha_{859} = 0.9865811034998516, \quad \frac{\alpha_{859}}{1 - \alpha_{859}} = 73.52177606325098,$$

and

$$F(p_{\sharp}) - F(p^{(859)}) \leq 7.987593427350412 \times 10^{-7},$$

while the exact validation gap is

$$7.987586611857278 \times 10^{-7}.$$

Thus the sweep counts 548 and 859 are no longer merely finite audit indices: they sit inside a monotone M -matrix convergence-rate certificate on the same fixed strict slab.

Corollary O71 (Secondary tightening status after the GS rate audit). The two secondary items listed in Corollary O70 are now separated as follows:

presentation synthesis closed by Proposition O71,
GS convergence-rate polish closed by Proposition O72.

TABLE XX. Geometry-informed open-system demonstrator at $D = \{6, 12, 18\}$. Source CSV: `output/chi_open_system/chi_open_system_geometry_D6-12-18.csv`.

D	γ_ϕ	γ_{mix}	C_{max}	$\chi_{\text{eff}}^{(\text{proxy})}$	$\chi_N^{(LR)}$	ratio
6	0.656901	0.227474	0.077658	6.24×10^{-5}	4.02×10^{-4}	0.155
12	0.744389	0.179054	0.063765	2.82×10^{-5}	2.21×10^{-4}	0.128
18	0.805854	0.136694	0.051381	2.19×10^{-5}	2.13×10^{-4}	0.103

The mathematical content of the strict-slab physical-trace line is therefore closed on the current audited bridge. What remains is editorial compression for readability, not a new mathematical bottleneck.

Reviewer-facing compressed form. Equivalently, Propositions O68–O72 prove the following single certificate on the fixed slab $S = \Omega_{+, \#}(11)$. On the 45 physical edge points define

$$\widehat{b}_q := \sum_{r \in N(q) \setminus S} c_{qr} \phi_r^{--}.$$

The missing-neighbor identity and exterior trace barrier give

$$b_{\#}(q) \geq \widehat{b}_q.$$

The monotone adjoint GS construction gives

$$p_{\#}(q) \geq p^{(548)}(q),$$

and Proposition O72 supplies the independent rate certificate

$$p_{\#} - p^{(548)} \leq \frac{\alpha_{548}}{1 - \alpha_{548}} d_{548}, \quad \alpha_{548} = 0.9865813608269858.$$

Therefore the whole physical-trace proof reduces to the single audited inequality

$$I_{\partial, \#}^{(\text{phys})} \geq \sum_q \widehat{b}_q p^{(548)}(q) = 1.345329256808323 \times 10^{-1}.$$

After adding the already-positive source mass, this gives

$$I_{\partial, \#}^{(\text{phys})} + I_{\text{src}, \#} \geq 1.4575819833843956 \times 10^{-1} = 1.003005 I_{\text{req}}.$$

This is the recommended reading of the strict-slab line in the main text; the preceding propositions merely unpack the four audited ingredients.

Across all sensitivity cases in Table XXI, the generation metrics remain unchanged on this grid ($f(\mathcal{R}_3 > 0.90) = 0.9269$, $f(\mathcal{R}_3 > 0.95) = 0.2167$, $f(N_{\text{win}} > 3) \simeq 2.78 \times 10^{-4}$), while the $H \rightarrow \mu\mu$ acceptance fraction remains fixed at $f(\chi_{\mu\mu}^2 < 4) = 0.1500$ in this setup. These values

TABLE XXI. Microscopic open-system sensitivity scan with profiled $(\gamma_\phi, \gamma_{\text{mix}})$ on the 60×60 map.
Source CSV: `paper/chi_open_system_micro_sensitivity.csv`.

case	ratio min	ratio max	ratio mean
baseline localized	1.000	1.000	1.000
open-micro base (1, 1)	0.017	0.327	0.138
$\phi \times 0.5$, mix $\times 1$	0.017	0.327	0.138
$\phi \times 2$, mix $\times 1$	0.017	0.327	0.138
$\phi \times 1$, mix $\times 0.5$	0.009	0.220	0.082
$\phi \times 1$, mix $\times 2$	0.034	0.412	0.208

Map-Level Migration Snapshot: baseline vs first-principles switches
R3 drift (absolute): legacy_cardy=-0.0500, open_system_micro=+0.0000




FIG. 15. Map-level migration snapshot across three scenarios. Left to right: $f(\mathcal{R}_3 > 0.90)$, $f(\chi_{\mu\mu}^2 < 4)$, and $f(N_{\text{win}} > 3)$.

are model-dependent because the Lindblad structure is still an effective environment ansatz and not yet a microscopic EYMH bath derivation; therefore, the open-system microscopic branch is treated as scan-ready diagnostic only and is not propagated into baseline reported figures. The explicit gate file `output/chi_open_system/open_system_micro_baseline_candidate.csv` records that this diagnostic passes the multi-anchor + holdout candidate gates (including hold-out RMSE/max-abs thresholds), but we keep it outside baseline until the bath-side microscopic derivation is completed.

5. Cross-Module Migration Snapshot

To compare the two main migration directions on the same map-level metrics under the updated baseline, we aggregate three reviewer-facing scenarios: the release baseline, a legacy Cardy comparator, and the open-system microscopic diagnostic. Exact mode flags and file-level provenance are retained in the artifact registry; Fig. 15 shows the resulting summary.

TABLE XXII. Cross-module migration comparison. Scenario names are reviewer-facing summaries; the source registry retains exact run-mode flags.

scenario	$f(\mathcal{R}_3 > 0.90)$	$\Delta f(\mathcal{R}_3 > 0.90)$	$f(\chi_{\mu\mu}^2 < 4)$	$f(N_{\text{win}} > 3)$
release baseline	0.927	0.0000	0.1500	0.0003
legacy Cardy comparator	0.877	-0.0500	0.1486	0.0003
open-system diagnostic	0.927	0.0000	0.1500	0.0003

TABLE XXIII. Reference-normalization candidates and re-normalized map metrics. Source CSV: `paper/reference_anchor_candidates.csv`; selector JSON: `paper/reference_anchor_choice.json` (historical artifact filenames retained for reproducibility).

mode	(D_0, η_0)	χ^2 at selector point	$f(\chi_{\mu\mu}^2 < 4)$ (re-normalized)	best χ^2 (re-normalized)
fixed-grid	(9.97, 0.97)	8.83×10^{-1}	0.1500	1.18×10^{-3}
χ^2 -best	(4.00, 0.264)	2.21×10^{-5}	0.1333	1.90×10^{-1}
robust-center	(9.42, 3.87)	1.39×10^{-3}	0.1381	4.74×10^{-3}

6. Reference-Normalization Sensitivity Audit

Because Eq. (56) is currently reported as a relative scan-level observable normalized to a reference point (D_0, η_0) , we explicitly audit the sensitivity of this reporting convention with a registered selector-and-rescan workflow:

1. build a fixed-reference baseline map at $(D_0, \eta_0) = (10, 1)$ on the same 60×60 grid;
2. extract three candidate reference points: fixed-grid point, χ^2 -best [Eq. (65)], and robust-center [Eq. (66)];
3. re-normalize Eq. (56) with each candidate and recompute map summaries.

Table XXIII summarizes the selector outputs, and Fig. 16 shows the full grid-level sensitivity maps. The source filenames retain the earlier “anchor” spelling for reproducibility; the physics interpretation is reference-normalization sensitivity rather than Higgs-sector anchoring.

Across the tested reference-normalization grid $D_0 \in \{7, \dots, 13\}$ and $\eta_0 \in \{0.6, 0.9, \dots, 2.4\}$, we find

$$f(\chi_{\mu\mu}^2 < 4) \in [0.0333, 0.2411], \quad \chi_{\min}^2 \in [1.52 \times 10^{-5}, 7.47 \times 10^{-1}]. \quad (\text{E14})$$

The baseline fixed-reference value used in main figures is $f(\chi_{\mu\mu}^2 < 4) = 0.1500$. For reporting discipline, we keep this fixed normalization point in headline plots and treat dynamic re-normalizations as sensitivity diagnostics rather than post-hoc optimization. The physical




FIG. 16. Systematic sensitivity over reference-normalization points (D_0, η_0) under the EFT/Wilson-matched UV+LL-RG map. Left: $f(\chi_{\mu\mu}^2 < 4)$; right: best χ^2 . Overlaid markers indicate baseline fixed, χ^2 -best, and robust-center selectors. Source files: `paper/reference_anchor_sensitivity.csv`, `paper/reference_anchor_sensitivity.png`.

normalization story should instead be read from the absolute and loop-improved parent-action comparators in Sec. VII, together with the Wilson-diagonal ratios in Eq. (67).

7. Dephasing-Based First-Principles Candidate for t_{coh}

To close the finite-time module with the same action-derived spectral chain, we test the dephasing candidate

$$\Delta\omega_{12}(D) \equiv \omega_2(D) - \omega_1(D), \quad t_{\text{coh}}^{(\text{deph})}(D) \equiv \frac{\pi}{\Delta\omega_{12}(D)}, \quad (\text{E15})$$

where (ω_1, ω_2) are extracted from the bound-state convention $E_n = \omega_n^2 - m_0^2 < 0$ in the same 1D on-axis operator chain as Appendix B. Detailed coarse/mid/fine values and implementation provenance are retained in the artifact registry. On the benchmark set $D = \{6, 12, 18\}$, the maximum non-fine relative deviation is

$$\max\left(\frac{|t_{\text{coh}} - t_{\text{coh}}^{(\text{fine})}|}{|t_{\text{coh}}^{(\text{fine})}|}\right) = 8.86 \times 10^{-3}, \quad (\text{E16})$$

which is numerically stable at the sub- 10^{-2} level for this benchmark set. The profile-impact scan compares constant $t_{\text{coh}} = 1$, raw Eq. (E15), and the capped variant $\min(t_{\text{coh}}^{(\text{deph})}, 10^4)$. All three leave the headline map summaries unchanged in this diagnostic chain: $f(\mathcal{R}_3 > 0.90) = 0.983$, $f(\mathcal{R}_3 > 0.95) = 0.217$, $f(\chi_{\mu\mu}^2 < 4) = 0.0833$, and best $\chi^2 = 9.62 \times 10^{-2}$. Therefore, in this manuscript, Eq. (E15) is reported as a first-principles candidate benchmark only and is not propagated into the baseline reported figures.

8. Splitting-Prefactor First-Principles Candidate for η

To test a geometry-closed replacement direction for the overlap amplitude in Eq. (24), we use the same action-derived quantities $(\Delta E_{12}, S_1)$ and define the splitting prefactor

$$\Delta E_{12}(D) \approx A_{12}(D) e^{-S_1(D)}, \quad A_{12}(D) \equiv \Delta E_{12}(D) e^{S_1(D)}. \quad (\text{E17})$$

With a reference point $D_{\text{ref}} = 12$, we define two dimensionless candidates:

$$\eta_{\text{amp}}(D) \equiv \frac{A_{12}(D)}{A_{12}(D_{\text{ref}})}, \quad \eta_{\text{prob}}(D) \equiv \eta_{\text{amp}}(D)^2. \quad (\text{E18})$$

Detailed coarse/mid/fine values and implementation provenance are retained in the artifact registry. On the same $D = \{6, 12, 18\}$ benchmark set, we obtain

$$\max \left(\frac{|\eta_{\text{amp}} - \eta_{\text{amp}}^{(\text{fine})}|}{|\eta_{\text{amp}}^{(\text{fine})}|} \right) = 1.08 \times 10^{-2}, \quad \max \left(\frac{|\eta_{\text{prob}} - \eta_{\text{prob}}^{(\text{fine})}|}{|\eta_{\text{prob}}^{(\text{fine})}|} \right) = 2.18 \times 10^{-2}, \quad (\text{E19})$$

and the fine full-scan profile over $D = 4, \dots, 20$ gives $\eta_{\text{amp}} \in [0.9996, 1.1695]$ and $\eta_{\text{prob}} \in [0.9992, 1.3677]$. Profile-scaled and fully closed implementations leave the headline map-level summaries unchanged in this diagnostic chain; the residual movement is only in sub-table-precision best- χ^2 shifts. Therefore, Eq. (E18) is reported as a first-principles candidate benchmark and is not propagated into baseline reported figures.

9. Action-Derived First-Principles Candidate for Barrier-Leakage Channel Scaling

To test a geometry-closed replacement direction for the channel normalization in Eq. (32), we keep the same 1D action-derived chain and define channel-resolved barrier actions

$$U_\ell(z; D) \equiv U(z; D) + \frac{\ell(\ell+1)}{z^2 + \varepsilon^2}, \quad S_{N,\ell}(D) \equiv \int_{z_-}^{z_+} dz \sqrt{U_\ell(z; D) - E_N(D)}, \quad (\text{E20})$$

with $E_N = \omega_N^2 - m_0^2$ in bound-state convention. We then define geometry-driven channel rates

$$\Gamma_{N,\ell}^{(\text{geo})}(D) \equiv \omega_N(D) e^{-2S_{N,\ell}(D)}, \quad (\text{E21})$$

and infer effective channel normalizations by matching to Eq. (32):

$$A_\ell^{(\text{fp})}(D; N) \equiv \frac{\Gamma_{N,\ell}^{(\text{geo})}(D)}{\omega_N(D) (\omega_N(D) M_*)^{4\ell+4}}. \quad (\text{E22})$$

Using $(D_{\text{ref}}, N_{\text{ref}}) = (12, 2)$, we define profile factors

$$\tilde{A}_\ell(D) \equiv \frac{A_\ell^{(\text{fp})}(D; N_{\text{ref}})}{A_\ell^{(\text{fp})}(D_{\text{ref}}; N_{\text{ref}})}. \quad (\text{E23})$$

Detailed coarse/mid/fine values and implementation provenance are retained in the artifact registry. On the $D = \{6, 12, 18\}$ benchmark set, the maximum non-fine relative deviations are

$$\max\left(\frac{|\tilde{A}_1 - \tilde{A}_1^{(\text{fine})}|}{|\tilde{A}_1^{(\text{fine})}|}\right) = 3.22 \times 10^{-2}, \quad \max\left(\frac{|\tilde{A}_2 - \tilde{A}_2^{(\text{fine})}|}{|\tilde{A}_2^{(\text{fine})}|}\right) = 3.78 \times 10^{-2}. \quad (\text{E24})$$

The fine full-scan profile over $D = 4, \dots, 20$ yields $\tilde{A}_1 \in [8.54 \times 10^{-6}, 4.10 \times 10^5]$ and $\tilde{A}_2 \in [3.95 \times 10^{-6}, 4.92 \times 10^6]$, indicating a very stiff D -dependence in this preliminary closure.

As a strict-upgrade check, we also connected a minimal complex-absorbing-potential solver to the same 1D action-derived operator chain and scanned benchmark windows in the absorber strength η and absorber onset z_{cap} . This pilot confirms that the non-Hermitian extraction pipeline is operational, but the benchmark mode remains in an absorber-linear or mixed regime rather than a plateau regime. A follow-up finite-radius Siegert/Robin audit, summarized in Proposition S1, reaches the same conservative conclusion from the outgoing-boundary side: the roots are parity-degenerate but strongly R -dependent, so they are still boundary-sensitive rather than stable complex poles. The combined physical extractor `code/extract_static_width_physical.py` now enforces both gates and marks the benchmark rows as `not_promoted`. Proposition S2 adds the first Riccati-DtN replacement of the local Robin exterior; it lowers the candidate widths but still fails both R - and R_{tail} -stability. Proposition S3 then replaces the WKB exterior start by an analytic Coulomb-tail normalization; this removes the most ad hoc tail datum but still fails the exterior-stability gates. Proposition S4 adds the independent ECS cross-check; it solves the complex-scaled eigenproblem to roundoff residuals, but the tracked branch remains sensitive to θ , R_0 , and L . Proposition S5 then removes the single-branch-selection loophole by tracking optimistic nearest-neighbor continuations over the resolved ECS spectrum; no seed-qualified branch passes. Proposition S6 implements the remaining full Coulomb/Whittaker outgoing DtN with a complex U-function backend and still finds R - and matching-radius sensitivity. Proposition S7 then applies the same optimistic nearest-neighbor branch-isolation idea directly to the Whittaker-DtN roots; again 0/48 anchor continuations pass. Proposition S8 applies local argument-principle contours and certifies that all 48 Whittaker roots are isolated local finite-radius zeros, so the obstruction is not local root existence but exterior-parameter displacement. Proposition S9 then turns that obstruction into an explicit parent-lift budget: the best audited branch still misses the square-contour cover condition by a factor 27.38 and the local residual budget by a factor 43.48. Proposition S10 then separates raw Coulomb-tail forcing from Riccati transfer conditioning: the raw tail is below 1.64×10^{-2} of the local budget, but only 4/48 rows pass after inward transfer and the worst miss is a factor 88.60. Proposition S11 tests that transfer obstruction directly: the finite dyadic pass count improves from 4/48 on $400 \rightarrow 800$ to 24/48 on $800 \rightarrow 1600$, but the infinitesimal condition-number route is too

pessimistic, with K up to 6.45×10^6 and 0/48 linear passes on $800 \rightarrow 1600$. Proposition S12 then uses the exact projective/Jost finite-amplitude identity: the endpoint dyadic contraction at $1600 \rightarrow 3200$ passes all 48/48 audited rows with maximum ratio 1.435×10^{-3} , but an analytic Jost-product or pole-exclusion lower bound is still needed before promotion. Accordingly, the CAP, local-Siegert, Riccati-DtN, Coulomb-normalized DtN, ECS, and full Whittaker-DtN outputs are retained as infrastructure diagnostics only; the reported baseline continues to use $\Gamma_{N,\ell}^{(\text{geo})}$ as the leading static surrogate unless a future analytic parent Jost-product or projective pole-exclusion theorem changes the gate.

At current scan resolution, Table XXIV shows that the t_{coh} and η candidates are stable diagnostics whose map-level effects are below headline precision, while replacing the legacy constant- A_ℓ reference by the action-derived profile induces non-negligible shifts in $\mathcal{R}_3$ coverage. In the updated global scans of this manuscript, Eq. (E23) is used in baseline mode, and $(A_1, A_2) = (1, 1)$ is retained only as a comparator.

TABLE XXIV. Compressed late-stage diagnostic summary for the three first-principles profile candidates. Detailed coarse/mid/fine tables and exact run artifacts are retained in the artifact registry.

Candidate	Local benchmark	Full/profile range	Map-level impact
$t_{\text{coh}}^{(\text{deph})}$	non-fine drift 8.86×10^{-3} vs fine	raw and capped profiles tested	$f(\mathcal{R}_3 > 0.90) = 0.983$, $f(\chi_{\mu\mu}^2 < 4) = 0.0833$, best $\chi^2 = 9.62 \times 10^{-2}$ unchanged
$\eta_{\text{amp}}, \eta_{\text{prob}}$	max drifts 1.08×10^{-2} and 2.18×10^{-2}	$\eta_{\text{amp}} \in [0.9996, 1.1695]$, $\eta_{\text{prob}} \in [0.9992, 1.3677]$	profile-scaled and fully closed variants leave headline map summaries unchanged
$\tilde{A}_\ell(D)$	max drifts 3.22×10^{-2} and 3.78×10^{-2}	$\tilde{A}_1 \in$ $[8.54 \times 10^{-6}, 4.10 \times 10^5]$, $\tilde{A}_2 \in$ $[3.95 \times 10^{-6}, 4.92 \times 10^6]$	action profile shifts $f(\mathcal{R}_3 > 0.90)$ from 0.9833 to 0.9269; $H \rightarrow \mu\mu$ acceptance stays 0.0833

Appendix F: Ansatz Ablation Study

To demonstrate that the three-generation structure is robust and not an artifact of a particular ansatz choice, we compare the baseline EFT-operator-normalized $B_N(D)$ profile (this work) against the alternative “inverse Yukawa” ansatz explored in earlier versions.

1. Inverse Yukawa Ansatz (Ablation A)

The inverse Yukawa ansatz sets $B_N^{(\text{inv})} \propto 1/\sum_{f \in \text{Gen } N} y_f$, yielding:

$$B_1^{(\text{inv})} \approx 10^5, \quad B_2^{(\text{inv})} \approx 125, \quad B_3^{(\text{inv})} = 1. \quad (\text{F1})$$

This massive hierarchy ($B_1 \gg B_2 \gg B_3$) directly encodes the answer by over-enhancing layer 1. While it can produce $\mathcal{R}_3 > 95\%$, the mechanism is less transparent: the visibility factor, rather than the entropic/kinetic competition, drives the result.

2. Comparison

- **Baseline (EFT-operator-normalized profile):** $B_N(D)$ remains smooth and bounded across the scanned range, with shape inherited from tracked microcanonical overlap extraction and operator normalization. The three-generation structure emerges from the g_N – Γ_N competition modulated by these profile differences.
- **Ablation A (Inverse Yukawa):** $B_1 \gg B_2 \gg B_3$. Layer 1 is artificially boosted; the selection mechanism is less falsifiable.

We recommend the baseline as the primary closure due to its greater transparency and falsifiability.

-
- [1] N. Gresnigt, Electroweak structure and three fermion generations in clifford algebra with s_3 family symmetry (2026), arXiv:2601.07857 [hep-ph].
 - [2] A. Dabholkar, S. Murthy, and D. Zagier, Quantum black holes, wall crossing, and mock modular forms (2012), arXiv:1208.4074 [hep-th].
 - [3] J. Maldacena and L. Susskind, Fortschritte der Physik **61**, 781 (2013), arXiv:1306.0533 [hep-th].
 - [4] J. L. Cardy, Nucl. Phys. B **270**, 186 (1986).
 - [5] J. D. Bekenstein, Phys. Rev. D **7**, 2333 (1973).
 - [6] G. 't Hooft, Dimensional reduction in quantum gravity (1993), arXiv:gr-qc/9310026 [gr-qc].
 - [7] S. Detweiler, Phys. Rev. D **22**, 2323 (1980).
 - [8] S. Navas *et al.*, Phys. Rev. D **110**, 030001 (2024), Particle Data Group.
 - [9] ATLAS Collaboration, Phys. Rev. Lett. **135**, 231802 (2025), arXiv:2507.03595 [hep-ex].
 - [10] V. Gorini, A. Kossakowski, and E. C. G. Sudarshan, J. Math. Phys. **17**, 821 (1976).
 - [11] G. Lindblad, Commun. Math. Phys. **48**, 119 (1976).
 - [12] J. C. F. Thaler and S. Trifinopoulos, Phys. Rev. D **111**, 056021 (2025), arXiv:2410.23343 [hep-ph].